

1.3 A Simple OS

- OS Structure
- Processes, Address Spaces, & Threads
- Managing Processes
- Loading Program into Processes
- Files

(first 13 slides overlap with "A Simple OS" slides)



Copyright © William C. Cheng



Copyright © William C. Cheng

- Directory system
 - shared by all processes running on a computer
 - although each process can have a different view
 - Unix provides a means to restrict a process to a subtree
 - by redefining what "root" means for the process
 - name space is outside the processes
 - a user process provides the name of a file to the OS
 - the OS returns a *handle* to be used to access the file
 - after it has verified that the process is allowed *access* along the *entire path*, starting from root
 - user process uses the handle to read/write the file
 - avoid access checks
- Using a handle to refer to an object managed by the kernel is an important concept
- handles are essentially an *extension* to the process's *address space*
 - can even survive execs!

Naming Files

Operating Systems - CSCI 402



Copyright © William C. Cheng

```
int fd;
char buffer[1024];
int count;
if ((fd = open("/home/bc/file", O_RDWR)) == -1) {
    perror("read");
    exit(1);
}
if ((count = read(fd, buffer, 1024)) == -1) {
    // the read failed
    perror("read");
    exit(1);
}
// buffer now contains count bytes read from the file
// what is O_RDWR?
// what does perror() do?
// cursor position in an opened file depends on what
// functions/system calls you use
// what about C++?
```

File Handles (File Descriptors)

Operating Systems - CSCI 402

Standard File Descriptors

- 0 is stdin (by default, the keyboard)
- 1 is stdout (by default, the display)
- 2 is stderr (by default, the display)

```
main() {
    char buf[BUFFSIZE];
    while ((n = read(0, buf, sizeof(buf))) > 0) {
        if (write(1, buf, n) != n) {
            void write(2, note, strlen(note));
            exit(EXIT_FAILURE);
        }
        return(EXIT_SUCCESS);
    }
}
```

Copyright © William C. Cheng

The File Abstraction

- A file is a simple array of bytes
- Files are made larger by writing beyond their current end
- Files are named by paths in a naming tree
- Files are named by paths in a naming tree
- System calls on files are *synchronous*
- File API
 - open(), read(), write(), close()
 - e.g., cat



Copyright © William C. Cheng



Copyright © William C. Cheng

- Our primes program wasn't too interesting
- it has no output!
- cannot even verify that it's doing the right thing
- other program cannot use its result
- how does a process write to someplace outside the process?
- The notion of a *file* is our Unix system's sole abstraction for this concept of "someplace outside the process"
- modern Unix systems have additional abstractions
- Files
 - abstraction of persistent data storage
 - means for fetching and storing data outside a process
 - including disks, another process, keyboard, display, etc.
 - need to *name* these different places
 - hierarchical naming structure
 - part of a process's *extended address space*

Files

Operating Systems - CSCI 402

I/O Redirection

The ">" parameter in a shell command that instructs the command shell to *redirect* the output to the given file

↳ If ">" weren't there, the output would go to the display

↳ Can also redirect input

```
% cat < /home/bc/output
```

↳ when the "cat" program reads from file descriptor 0, it would get the data bytes from the file "/home/bc/output"

Copyright © William C. Cheng

11

Allocation of File Descriptors

↳ Whenever a process requests a new file descriptor, the lowest numbered file descriptor not already associated with an open file is selected; thus

```
#include <fcntl.h>
#include <unistd.h>
...
close(0);
fd = open("file", O_RDONLY);
```

↳ will always associate "file" with file descriptor 0 (assuming that the open succeeds)

Copyright © William C. Cheng

9

Back to Primes

↳ Have our primes program write out the solution, i.e., the primes [array

```
int nprimes;
int *prime;
int main(int argc, char *argv[]) {
    ...
    for (i=1; i<nprimes; i++) {
        ...
        if (write(1, prime, nprimes*sizeof(int)) == -1) {
            perror("primes output");
            exit(1);
        }
        return(0);
    }
}
```

↳ the output is not readable by human

Copyright © William C. Cheng

7

File-Descriptor Table

↳ A file descriptor refers not just to a file

- ↳ it also refers to the *process's* current *context* for that file
- ↳ includes how the file is to be accesses (how open () invoked)
- ↳ *cursor* position

↳ *Context* information must be maintained by the OS and not directly by the user program

↳ let's say a user program opened a file with O_RDONLY later on it calls write () using the opened file descriptor

↳ how does the OS knows that it doesn't have write access?

- ↳ stores O_RDONLY in context
- ↳ if the user program can manipulate the context, it can change O_RDONLY to O_RDWR

↳ therefore, user program must not have access to context!

- ↳ all it can see is the handle
- ↳ the handle is an *index* into an array maintained for the process in kernel's address space

Copyright © William C. Cheng

12

Running it

```
if (fork() == 0) {
    /* set up file descriptor 1 in the child process */
    close(1);
    if (open("/home/bc/output", O_WRONLY) == -1) {
        perror("/home/bc/output");
        exit(1);
    }
    execl("/home/bc/bin/primes", "primes", "300", 0);
    exit(1);
}
/* parent continues here */
while (pid != wait(0)) /* ignore the return code */
;
close(1) removes file descriptor 1 from extended address space
file descriptors are allocated lowest first on open ()
↳ extended address space survives execs
new code is same as running
% primes 300 > /home/bc/output
```

Copyright © William C. Cheng

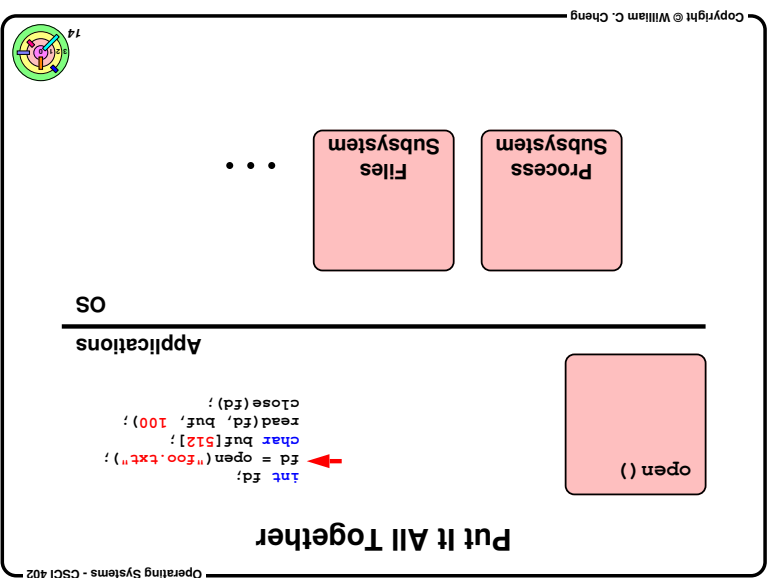
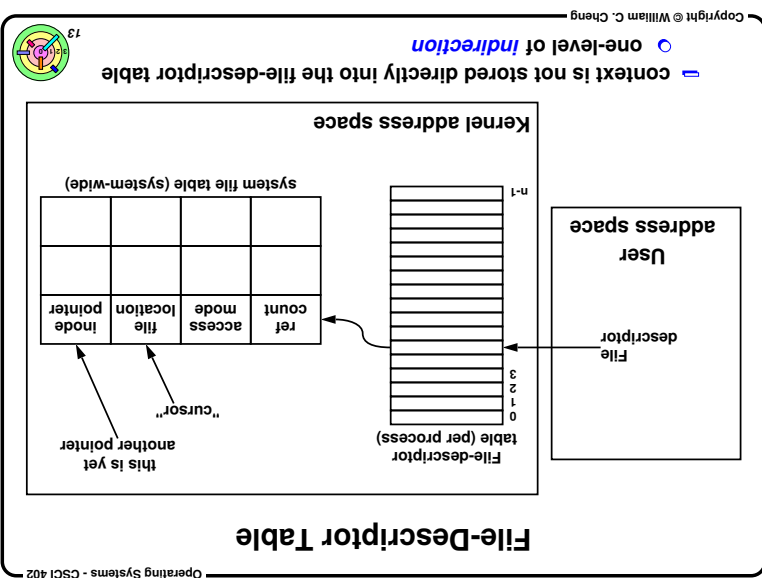
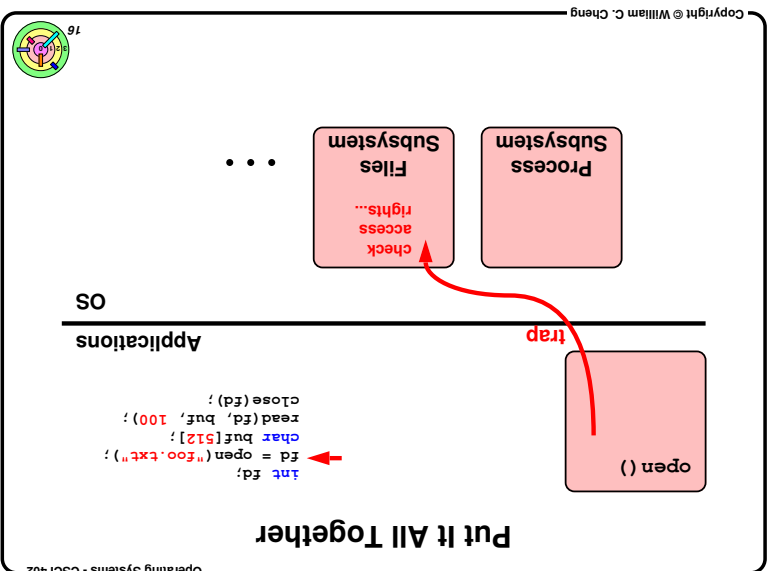
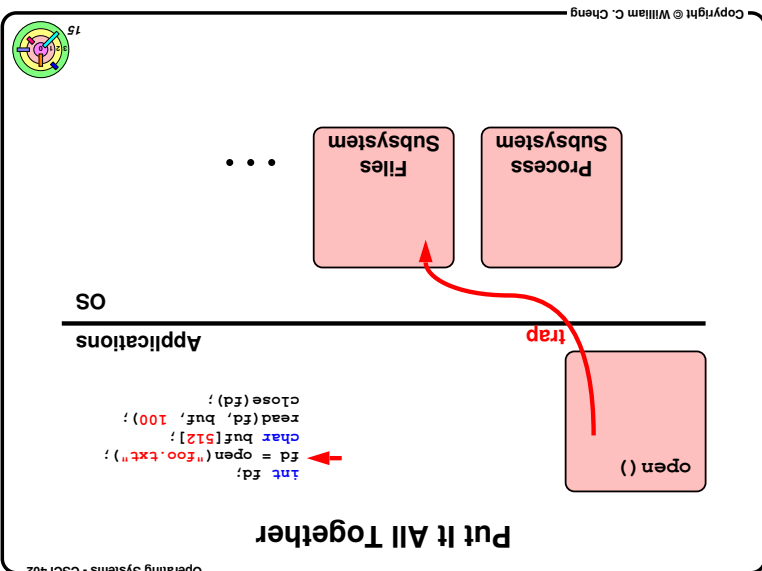
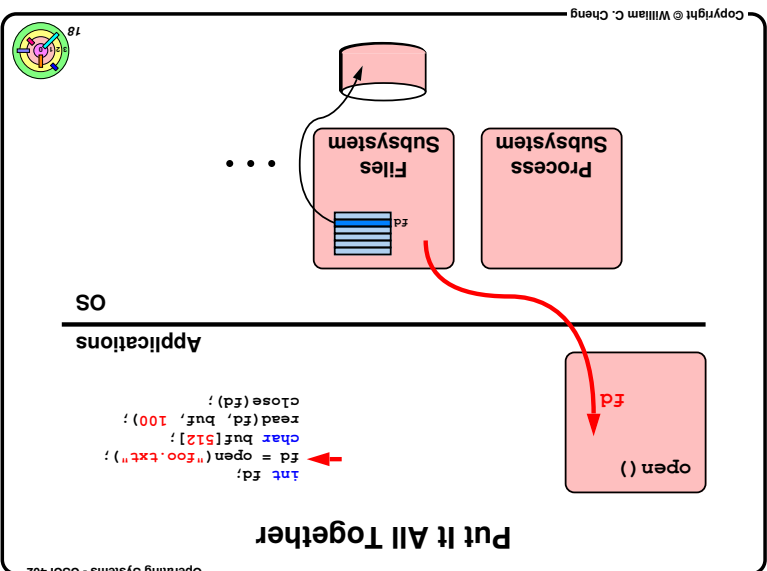
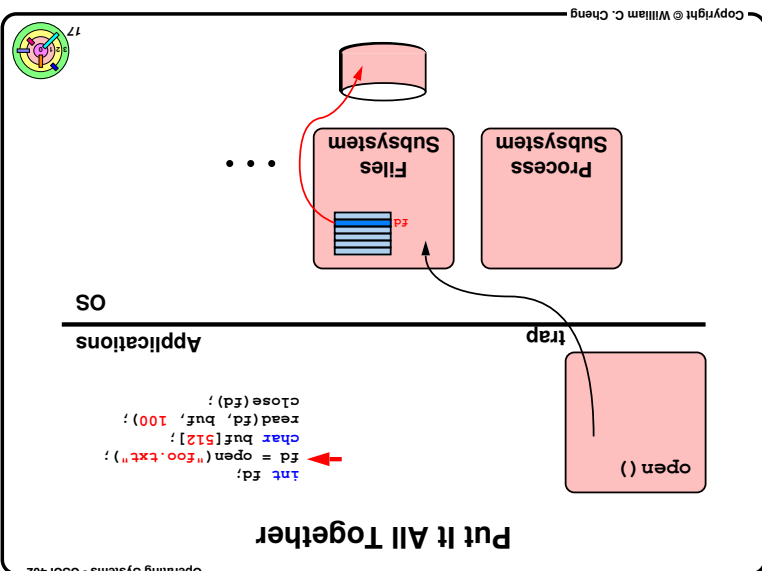
10

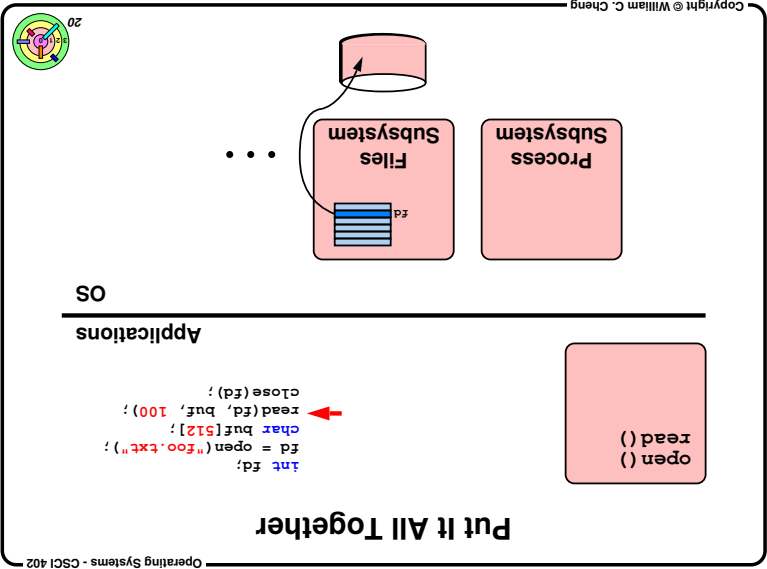
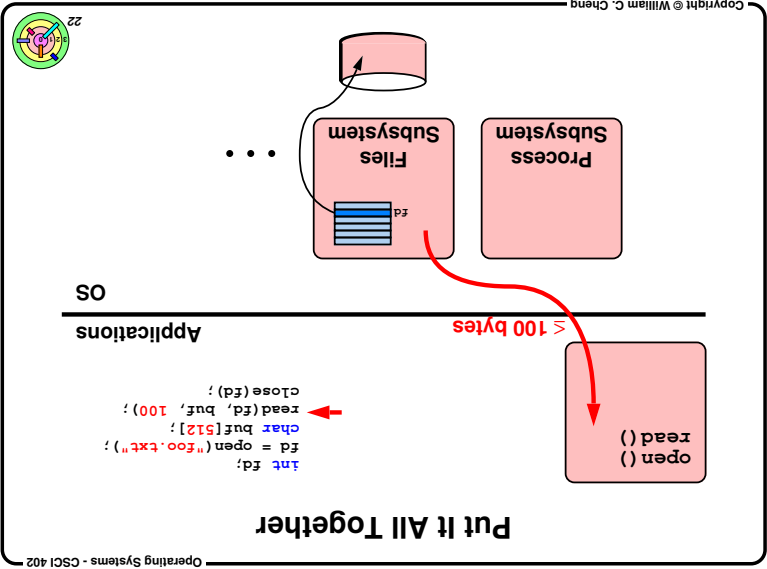
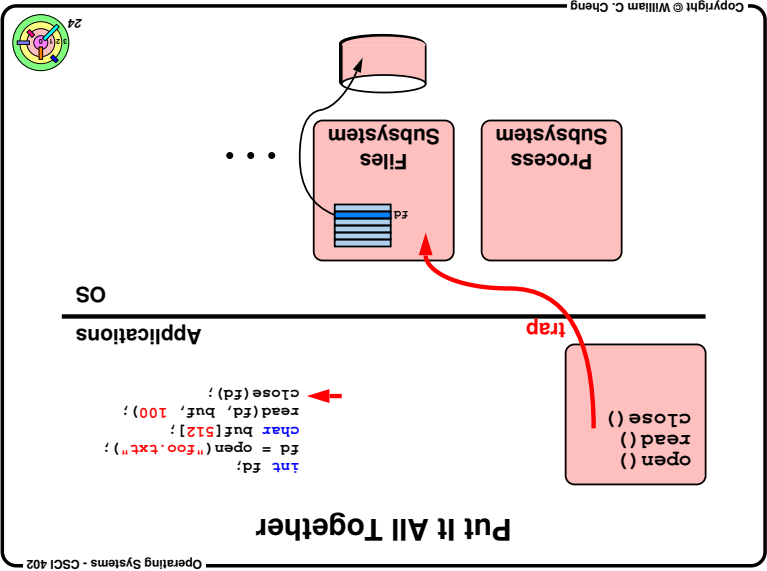
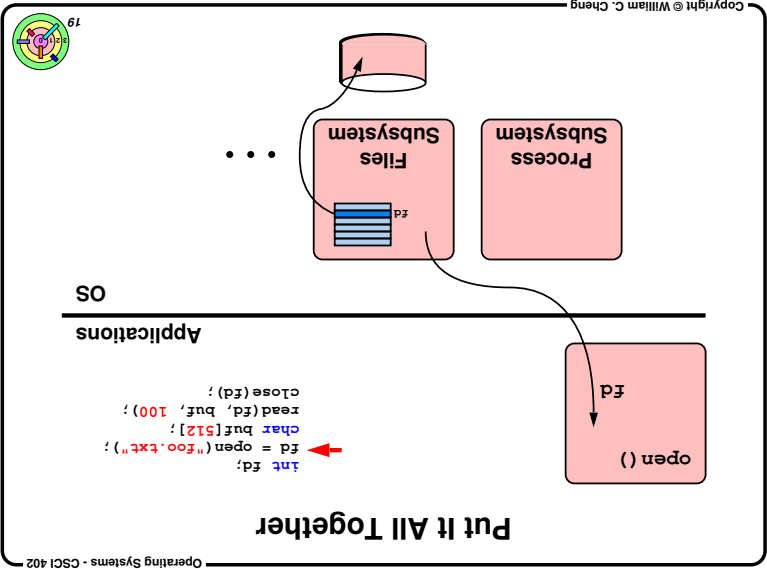
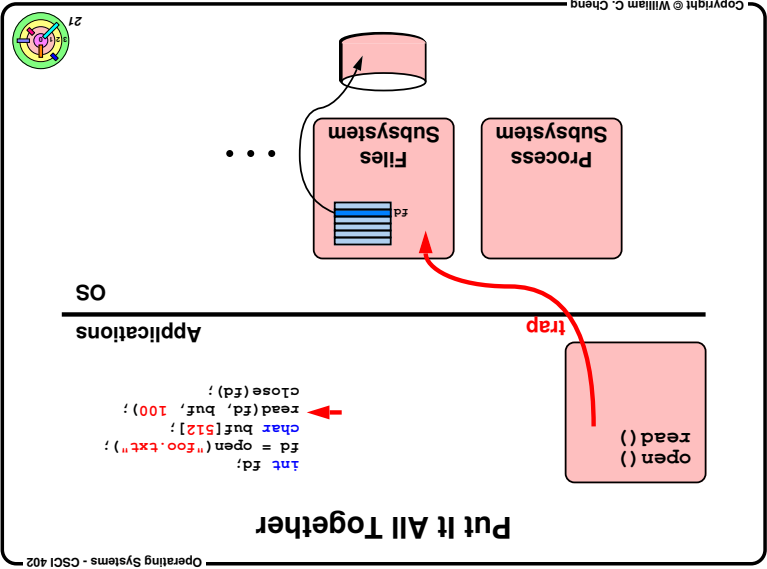
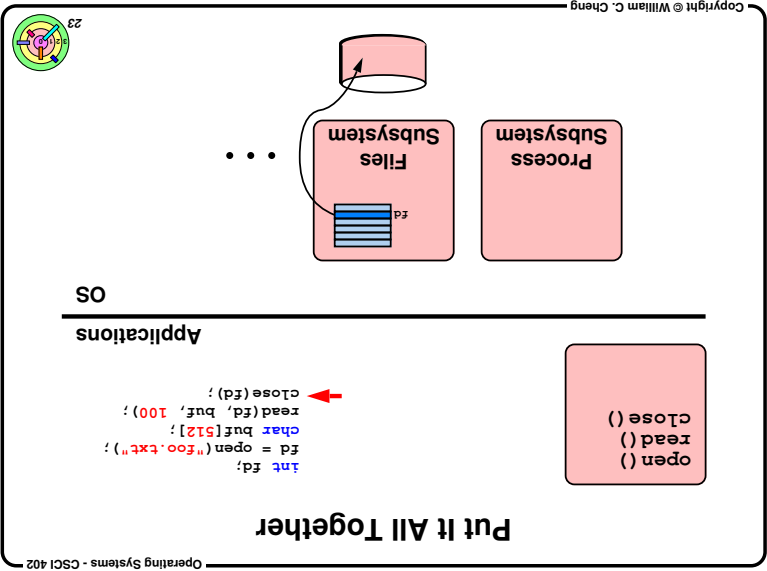
Human-Readable Output

```
int nprimes;
int *prime;
int main(int argc, char *argv[]) {
    ...
    for (i=1; i<nprimes; i++) {
        ...
        for (j=0; j<nprimes; j++) {
            printf("%d\n", prime[j]);
        }
        return(0);
    }
}
```

Copyright © William C. Cheng

8





Copyright © William C. Cheng

29

Child's new address space

File descriptor 1
File descriptor 2

File-descriptor table (per process)

Kernel address space

```

...
close(1);
...
if (open(
"/home/bc/output",
O_WRONLY) == -1) {
    exit(1);
}
if (open(
"/home/bc/output",
O_WRONLY) == -1) {
    exit(1);
}
exec1(...);
    
```

Redirected Output

Operating Systems - CSCI 402

Copyright © William C. Cheng

27

Every call to `open()` creates a new entry in the system file table

Redirecting Output ... Twice

```

if (fork() == 0) {
    /* set up file descriptors 1 and 2 in the child
    process */
    close(1);
    if (open("/home/bc/output", O_WRONLY) == -1) {
        exit(1);
    }
    if (open("/home/bc/output", O_WRONLY) == -1) {
        exit(1);
    }
    exec1("/home/bc/bin/program", "program", 0);
    exit(1);
}
/* parent continues here */
    
```

↳ `stdout` and `stderr` both go into the same file

○ would it cause any problem?

Operating Systems - CSCI 402

Copyright © William C. Cheng

25

Put it All Together

```

int fd;
fd = open("foo.txt");
char buf[512];
read(fd, buf, 100);
close(fd);
    
```

Applications

OS

Subsystem

Files

Process

Operating Systems - CSCI 402

Copyright © William C. Cheng

30

Child's new address space

File descriptor 1
File descriptor 2

File-descriptor table (per process)

Kernel address space

```

...
close(1);
...
if (open(
"/home/bc/output",
O_WRONLY) == -1) {
    exit(1);
}
if (open(
"/home/bc/output",
O_WRONLY) == -1) {
    exit(1);
}
exec1(...);
    
```

↳ remember, extended address space survives `execs`

Redirected Output

Operating Systems - CSCI 402

Copyright © William C. Cheng

28

Child's new address space

File descriptor 1

File-descriptor table (per process)

Kernel address space

```

...
close(1);
...
if (open(
"/home/bc/output",
O_WRONLY) == -1) {
    exit(1);
}
if (open(
"/home/bc/output",
O_WRONLY) == -1) {
    exit(1);
}
exec1(...);
    
```

Redirected Output

Operating Systems - CSCI 402

Copyright © William C. Cheng

26

Put it All Together

```

int fd;
fd = open("foo.txt");
char buf[512];
read(fd, buf, 100);
close(fd);
    
```

Applications

OS

Subsystem

Files

Process

Operating Systems - CSCI 402

Copyright © William C. Cheng

35

When `fork()` is called, the child process gets a **copy** of the parent's file descriptor table

→ remember, extended address space survives execs

also `fork()`

```

...
}
/* parent process computes something, then does: */
write(logfile, LogEntry, strlen(LogEntry));
exit(0);
/* child process computes something, then does: */
if (fork() == 0) {
    int logfile = open("log", O_WRONLY);
    write(logfile, LogEntry, strlen(LogEntry));
}

```

File Descriptors

Copyright © William C. Cheng

33

Redirected Output After Dup

```

...
if (open(
    "/home/bc/output",
    O_WRONLY) == -1) {
    dup(1);
    exit(1);
}
close(2);
close(1);
dup(1);
exit(1);

```

Copyright © William C. Cheng

31

Redirected Output After Writing 100 Bytes

→ `write()` to `fd=2` will wipe out data in the first 100 bytes that may not be the intent

Copyright © William C. Cheng

36

File Descriptors After Fork

→ parent and child processes get **separate** file descriptor table but **share** extended address space

```

...
if (fork() == 0) {
    int logfile = open("log",
        O_WRONLY);
    write(logfile, LogEntry,
        strlen(LogEntry));
    exit(0);
}
write(logfile, LogEntry,
    strlen(LogEntry));

```

Copyright © William C. Cheng

34

Redirected Output After Dup

```

...
if (open(
    "/home/bc/output",
    O_WRONLY) == -1) {
    dup(1);
    exit(1);
}
close(2);
close(1);
dup(1);
exit(1);

```

Copyright © William C. Cheng

32

Sharing Context Information

→ use the `dup()` system call to **share** context information

→ if that's what you want

```

/* parent continues here */
}
dup(1);
execl("/home/bc/bin/program", "program", 0);
exit(1);
if (open(
    "/home/bc/output",
    O_WRONLY) == -1) {
    close(1);
    close(2);
    process;
}
/* set up file descriptors 1 and 2 in the child */
if (fork() == 0) {
    dup(1);
    execl("/home/bc/bin/program", "program", 0);
    exit(1);
}

```

Copyright © William C. Cheng

41

parent and child processes get **separate** file descriptor table but **share** extended address space

Kernel address space

Child's address space

Parent's address space

Pipes

```

    ...
    write(p[1], buf, 80);
    for (;;) {
        close(p[0]);
        else {
            exit(0);
        }
        ...
        while (read(p[0], buf, 80) > 0) {
            if (fork() == 0) {
                pipe(p);
                close(p[1]);
                while (read(p[0], buf, 80) > 0) {
                    write(p[1], buf, 80);
                }
            }
        }
    }

```

Operating Systems - CSCI 402

Copyright © William C. Cheng

39

parent creates a pipe object in the kernel

Kernel address space

Parent's address space

Pipes

```

    int p[2]; // array to hold pipe's file descriptors
    // creates a pipe, assume no errors
    // p[0] refers to the read/output end of the pipe
    // p[1] refers to the write/input end of the pipe
    if (fork() == 0) {
        char buf[80];
        close(p[1]); // not needed by the child
        while (read(p[0], buf, 80) > 0) {
            // use data obtained from parent
            ...
            exit(0); // child done
        }
        else {
            char buf[80];
            close(p[0]); // not needed by the parent
            // prepare data for child
            write(p[1], buf, 80);
        }
    }
}

```

Operating Systems - CSCI 402

Copyright © William C. Cheng

37

parent and child processes get **separate** file descriptor table but **share** extended address space

Kernel address space

Child's address space

Parent's address space

Pipes

```

    int logfile = open("log", O_WRONLY);
    if (fork() == 0) {
        write(logfile, "Lognry", strlen(lognry));
        exit(0);
        write(logfile, "Lognry", strlen(lognry));
    }

```

File Descriptors After Fork

Operating Systems - CSCI 402

Copyright © William C. Cheng

42

parent and child processes get **separate** file descriptor table but **share** extended address space

Kernel address space

Child's address space

Parent's address space

Pipes

```

    int p[2];
    pipe(p);
    if (fork() == 0) {
        close(p[1]);
        while (read(p[0], buf, 80) > 0) {
            write(p[1], buf, 80);
        }
    }
    else {
        close(p[0]);
        exit(0);
    }
}

```

Operating Systems - CSCI 402

Copyright © William C. Cheng

40

parent creates a pipe object in the kernel

Kernel address space

Parent's address space

Pipes

```

    int p[2];
    pipe(p);
    if (fork() == 0) {
        close(p[1]);
        while (read(p[0], buf, 80) > 0) {
            write(p[1], buf, 80);
        }
    }
    else {
        close(p[0]);
        exit(0);
    }
}

```

Operating Systems - CSCI 402

Copyright © William C. Cheng

38

A pipe is a means for one process to send data to another directly, as if it were writing to a file

process

sender

receiver

Pipes

- the sending process behaves as if it has a file descriptor to a file that has been opened for writing
- the receiving process behaves as if it has a file descriptor to a file that has been opened for reading

The pipe() system call creates a pipe object in the kernel and returns (via an output parameter) the two file descriptors that refer to the pipe

- one, set for write-only, refers to the input side
- the other, set for read-only, refers to the output side
- a pipe has no name, cannot be passed to another process

Operating Systems - CSCI 402

Copyright © William C. Cheng 47

More On Naming

(Almost) everything has a path name

- files
- directories
- devices (known as *special files*)
- keyboards
- displays
- disks
- etc.

Uniformity

```

// opening a normal file
int file = open("/home/bc/data", O_RDWR);

// opening a device (one's terminal or window)
int device = open("/dev/tty", O_RDWR);

int bytes = read(file, buffer, sizeof(buffer));
write(device, buffer, bytes);

```

Operating Systems - CSCI 402

Copyright © William C. Cheng 45

Command Shell

Now you know enough to write a command shell

- execute a command
- redirect I/O
- pipe the output of one program to another
- cat f0 | ./warmup1 sort
- the shell needs to create a pipe
- create two child processes
- in the first child
- have stdout go to the write-end of the pipe
- close the read-end of the pipe
- exec "cat f0"
- in the 2nd child
- have stdin come from the read-end of the pipe
- close the write-end of the pipe
- exec "./warmup1 sort"
- run a program in the background
- primes 1000000 > primes.out &

Operating Systems - CSCI 402

Copyright © William C. Cheng 43

Pipes

parent and child processes get *separate* file descriptor table but *share* extended address space

```

int p[2];
if (fork() == 0) {
    close(p[1]);
    while (read(p[0], buf, 80) > 0) {
        // ...
    }
} else {
    close(p[0]);
    while (write(p[1], buf, 80) > 0) {
        // ...
    }
}

```

Operating Systems - CSCI 402

Copyright © William C. Cheng 46

Directories

A *directory* is a *file*

- interprets differently by the OS as containing *references* to other files/directories
- a file is represented as an *index node* (or *inode*) in the file system

A *directory* maps a *file* name to an *inode number*

- maps a string to an integer
- done inside Virtual File System in weenix

An *inode* maps an *inode number* to *sectors on disk*

- done inside Actual File System in weenix

Operating Systems - CSCI 402

Copyright © William C. Cheng 46

Random Access

"man lseek" gives

```

off_t lseek(int fd, off_t offset, int whence);

```

- whence can be SEEK_SET, SEEK_CUR, SEEK_END
- if succeeds, returns cursor position (always measured from the beginning of the file)
- otherwise, returns (-1)
- errno is set to indicate the error

We need to move the cursor position back 2 positions

```

read(fd, buf, 1);
write(1, buf, 1);
while (fptr != -1) {
    fptr = lseek(fd, (off_t)(-1), SEEK_END);
    // go to last char in file
    fptr = lseek(fd, (off_t)(-1), SEEK_END);
}
fptr = lseek(fd, (off_t)(-2), SEEK_CUR);

```

Operating Systems - CSCI 402

Copyright © William C. Cheng 44

Pipes

parent closes the read-end of the pipe

child closes the write-end of the pipe

```

int p[2];
if (fork() == 0) {
    close(p[1]);
    while (read(p[0], buf, 80) > 0) {
        // ...
    }
} else {
    close(p[0]);
    while (write(p[1], buf, 80) > 0) {
        // ...
    }
}

```

Operating Systems - CSCI 402



Operating Systems - CSC1402

57



Operating Systems - CSC1402



Operating Systems - CSCI 402



Operating Systems - CSCI 402



Operating Systems - CSC1402

-

Operating Systems - CSCI 402



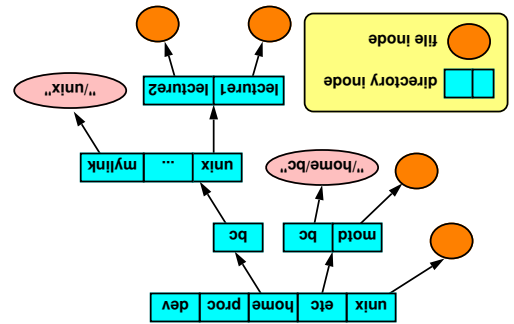
- ➡ OS needs to make sure that only authorized processes are allowed access to system resources
- ➡ various ways to provide this
- ➡ Unix (and many other systems, such as Windows) associates with files some indication of which *security principals* are allowed access
- ➡ along with what sort of access is allowed
- ➡ A *security principal* is normally a user or group of users
- ➡ a "user" can be an identity used by processes performing system functions
- ➡ each running process can have several security principals associated with it
- ➡ all processes have a user identification and a set of group identifications
- ➡ for Sixth-Edition Unix, only one user ID and one group ID

Access Protection



- Each file has associated with it a set of access permissions
 - there are 3 classes of security principals:
 - user**: owner of the file
 - group**: group owner of the file
 - others**: everyone else
 - for each of the 3 classes of principals, specify what sorts of operations on the file are allowed
 - the operations are grouped into 3 classes:
 - read**: can read a file or directory
 - write**: can write a file or directory
 - execute**: one must have **execute permission** for a **directory** in order to **follow a path through it**
- Rules for checking permissions**
- determines the **smallest class of principals the requester belongs to** (user being smallest and others being largest)
 - then it checks for appropriate permissions with that

Access Protection



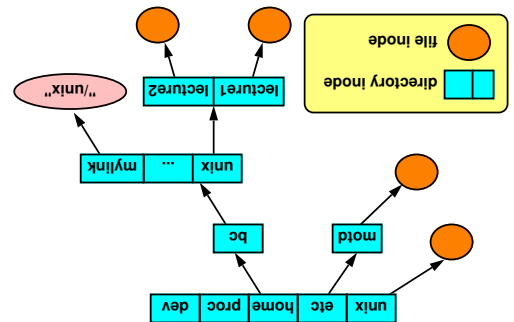
- ```
% ls -l /etc/bc/unix/lecture1
= same as "ls -l /home/bc/unix/lecture1", or is it?
o yes for the "root" account, may be no for the "bc" account
o see "access protection"
```

## Soft Links



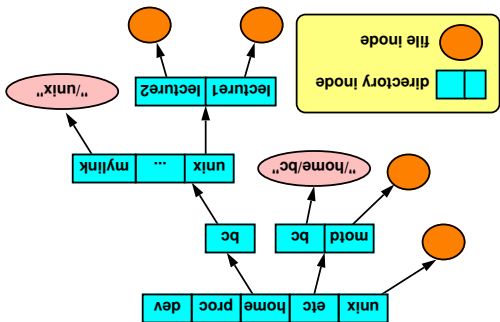
- ↳ Maintained *in kernel for each process*
- ↳ paths not starting from "/" start with the working directory
- ↳ get by using the `getcwd ()` system call
- ↳ set by using the `chdir ()` system call
- ↳ displayed (via shell) using "pwd"

## Working Directory



% 1n s- /unɪx/homə/bc/wɪlɪŋk  
 % 1n s- /homə/bc/etɪc/bc

## Soft Links



% ln -s /unix/home/bc/mylink  
% ln -s /home/bc/etc/bc

## Soft Links

Suppose that bill and trina are members of the adm group and

## Permissions Example

**Suppose that all and trina are members of the adm group and**

## Permissions Example

Suppose that bill and trina are members of the adm group and

## Permissions Example

**Suppose that all and trina are members of the adm group and**

## Permissions Example

Suppose that bill and trina are members of the adm group and

## Permissions Example

**Suppose that bill and trina are members of the adm group and**

## Permissions Example



and y is not  
6) Q: May b<sub>11</sub> read B/y?

Suppose that bill and trina are members of the adm group and

|                                |
|--------------------------------|
| % 1s -LR                       |
| :                              |
| total 2                        |
| dwxxr-x--x                     |
| 2 bill adm 1024 Dec 17 13:34 A |
| -----                          |
| 2 bill adm 1024 Dec 17 13:34 B |
| .                              |
| /A:                            |
| total 1                        |
| -rw-rw-rw-                     |
| 1 bill adm 593 Dec 17 13:34 x  |
| .                              |
| B:                             |
| total 2                        |
| --rw-rw-rw-                    |
| 1 bill adm 446 Dec 17 13:34 x  |
| --rw-rw-rw-                    |
| 1 trina adm 446 Dec 17 13:45 y |

## Permissions Example



6) Q: May bill read B/Y?  
A: No

Suppose that all and true are members of the adm group and  
andy is not

|                                           |
|-------------------------------------------|
| % ls -lR                                  |
| :                                         |
| total 2                                   |
| drwxr-x--x 2 bill adm 1024 Dec 17 13:34 A |
| drwxr----- 2 bill adm 1024 Dec 17 13:34 B |
| ./A:                                      |
| total 1                                   |
| -rw-rw-rw- 1 bill adm 593 Dec 17 13:34 x  |
| ./B:                                      |
| total 2                                   |
| -r--r--r-- 1 bill adm 446 Dec 17 13:34 y  |
| -rw----rw- 1 trina adm 446 Dec 17 13:45 v |

## Permissions Example



and y is not  
5) Q: May b:11 modify B/x?

Suppose that bill and trina are members of the adm group and

|                     |
|---------------------|
| % 1s -LR            |
| :                   |
| total 2             |
| dxwxr-x--x          |
| 2 bill              |
| adm                 |
| 1024 Dec 17 13:34 A |
| -----               |
| 2 bill              |
| adm                 |
| 1024 Dec 17 13:34 B |
| .                   |
| /A:                 |
| total 1             |
| -rw-rw-rw-          |
| 1 bill              |
| adm                 |
| 593 Dec 17 13:34 x  |
| .                   |
| B:                  |
| total 2             |
| -rw-rw-rw-          |
| 1 bill              |
| adm                 |
| 446 Dec 17 13:34 x  |
| -rw-rw-rw-          |
| 1 trina             |
| adm                 |
| 446 Dec 17 13:45 y  |

## Permissions Example



5) Q: May ball modify B/x?  
A: No

Suppose that all and true are members of the adm group and  
andy is not

|                     |
|---------------------|
| % Is -IR            |
| :                   |
| total 2             |
| dwrxr-x--x          |
| 2 bill              |
| adm                 |
| 1024 Dec 17 13:34 A |
| dwrxr-----          |
| 2 bill              |
| adm                 |
| 1024 Dec 17 13:34 B |
| /A:                 |
| total 1             |
| -rw-rw-rw-          |
| 1 bill              |
| adm                 |
| 593 Dec 17 13:34 x  |
| /B:                 |
| total 2             |
| --x--rw-rw-         |
| 1 bill              |
| adm                 |
| 446 Dec 17 13:34 x  |
| --rw-----           |
| 1 trina             |
| adm                 |
| 446 Dec 17 13:45 y  |

## Permissions Example



and is not  
4) Q: May traina modify B/Y?

- Suppose that bill and trina are members of the adm group and

|                     |
|---------------------|
| % 1s -1R            |
| :                   |
| total 2             |
| drwxr-x--x          |
| 2 blll              |
| adm                 |
| 1024 Dec 17 13:34 A |
| -----               |
| 2 blll              |
| adm                 |
| 1024 Dec 17 13:34 B |
| .                   |
| /A:                 |
| total 1             |
| -rw-rw-rw-          |
| 1 blll              |
| adm                 |
| 593 Dec 17 13:34 x  |
| .                   |
| total 2             |
| -rw-rw-rw-          |
| 1 blll              |
| adm                 |
| 446 Dec 17 13:34 y  |
| -rw-rw-rw-          |
| 1 trlna             |
| adm                 |
| 446 Dec 17 13:45 Y  |

## Permissions Example



4) Q: May traina modify B/y?  
A: No

Suppose that all and true are members of the adm group and  
andy is not

|                                           |
|-------------------------------------------|
| % ls -lR                                  |
| ..:                                       |
| total 2                                   |
| drwxr-x--x 2 bill adm 1024 Dec 17 13:34 A |
| drwxr-x--- 2 bill adm 1024 Dec 17 13:34 B |
| ./A:                                      |
| total 1                                   |
| -rw-rw-rw- 1 bill adm 593 Dec 17 13:34 x  |
| total 2                                   |
| -r-----rw- 1 trina adm 446 Dec 17 13:34 y |
| -r-----rw- 1 bill adm 446 Dec 17 13:34 x  |

## Permissions Example

# 1.4 Beyond A Simple OS

- ➡ Extensions
- ➡ New Functionality

## What Else?

- ➡ Beyond Sixth-Edition Unix (1975)
  - multiple threads per process
  - how is the process model affected?
  - virtual memory
  - in Sixth-Edition Unix, all currently running process had to fit into the computer's memory at once, along with the OS
  - virtual memory separates the address space from physical resources
  - ➡ name everything using directory-system path names
  - e.g., /proc
  - ➡ security
  - Unix solution is pretty elegant
  - new types of requirement such as permissions to add new software, perform backups, etc.

## Creating a File

- ➡ Use either open or creat
  - open (`const char *pathname, int flags, mode_t mode`)
  - flags must include `O_CREAT`
  - ➡ creat (`const char *pathname, mode_t mode`)
  - ➡ open is preferred
- ➡ The mode parameter helps specify the permissions of the newly created file
  - ➡ permissions = mode & ~umask

## Umask

- ➡ Standard programs create files with "maximum needed permissions" as `mode`
  - ➡ compilers: 0777
  - ➡ editors: 0666
- ➡ Per-process parameter, `umask`, used to *turn off* undesired permission bits
  - ➡ e.g., turn off all permissions for others, write permission for group: set umask to 027
  - ➡ compilers: permissions = 0777 & ~(027) = 0750
  - ➡ editors: permissions = 0666 & ~(027) = 0640
  - ➡ set with `umask ()` system call or (usually) `umask` shell command

## Open

- ```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int open(const char *path, int options [, mode_t mode])
options
➡ O_RDONLY open for reading only
➡ O_WRONLY open for writing only
➡ O_RDWR open for reading and writing
➡ O_APPEND set the file offset to end of file prior to each write
➡ O_CREAT if the file does not exist, then create it, setting its mode to mode adjusted by umask
➡ O_EXCL: if O_EXCL and O_CREAT are set, then open fails if the file exists
➡ O_TRUNC delete any previous contents of the file
➡ O_NONBLOCK don't wait if I/O cannot be done
immediately
```

Setting File Permissions

- ```
#include <sys/types.h>
#include <sys/stat.h>
int chmod(const char *path, mode_t mode)

➡ sets the file permissions of the given file to those specified in mode
➡ only the owner of a file and the superuser may change its permissions
➡ nine combinable possibilities for mode (read/write/execute for user, group, and others)
➡ S_IRUSR (0400), S_IWUSR (0200), S_IXUSR (0100)
➡ S_IRGRP (040), S_IWGRP (020), S_IXGRP (010)
➡ S_IROTH (04), S_IWOTH (02), S_IXOTH (01)
○ note: numeric prefix of 0 means the number is in octal format
```

## What Else?

### New functionalities

- networking
- much is beyond the scope of this class
- interactive, multimedia user interface
- make sure interactive user receives excellent response
- software complexity
- in Sixth-Edition Unix, to add a new device you need to:
  - write a "device driver" to handle the device
  - modify OS source code by adding references to the driver to a few tables
  - recompile the OS and reboot your computer
  - plug-and-play is desirable
  - need to support dynamic linking of modules into a running system
- microkernel
- how much of the OS functionality can be moved out of the kernel?

