



1.3 A Simple OS

- ➡ OS Structure
- ➡ Processes, Address Spaces, & Threads
- ➡ Managing Processes
- ➡ Loading Program Into Processes
- ➡ *Files*
 - (first 13 slides overlap with "A Simple OS" slides)



Files

- ➡ Our primes program wasn't too interesting
 - it has no output!
 - cannot even verify that it's doing the right thing
 - other program cannot use its result
 - how does a process write to someplace outside the process?
- ➡ The notion of a *file* is our Unix system's sole abstraction for this concept of "someplace outside the process"
 - modern Unix systems have additional abstractions
- ➡ Files
 - abstraction of persistent data storage
 - means for fetching and storing data outside a process
 - including disks, another process, keyboard, display, etc.
 - need to *name* these different places
 - ◆ hierarchical naming structure
 - part of a process's *extended address space*



Naming Files




Directory system

- shared by all processes running on a computer
 - although each process can have a different view
 - Unix provides a means to restrict a process to a subtree
 - ◆ by redefining what "root" means for the process
- name space is outside the processes
 - a user process provides the name of a file to the OS
 - the OS returns a *handle* to be used to access the file
 - ◆ after it has verified that the process is allowed *access* along the *entire path*, starting from root
 - user process uses the handle to read/write the file
 - ◆ avoid access checks




Using a handle to refer to an object managed by the kernel is an important concept

- handles are essentially an *extension* to the process's *address space*
 - can even survive execs!



The File Abstraction


- ➡ A file is a simple array of bytes
- ➡ Files are made larger by writing beyond their current end
- ➡ Files are named by paths in a naming tree
- ➡ System calls on files are *synchronous*
- ➡ File API
 - `open()`, `read()`, `write()`, `close()`
 - e.g., `cat`



File Handles (File Descriptors)

```
int fd;
char buffer[1024];
int count;
if ((fd = open("/home/bc/file", O_RDWR) == -1) {
    // the file couldn't be opened
    perror("/home/bc/file");
    exit(1);
}
if ((count = read(fd, buffer, 1024)) == -1) {
    // the read failed
    perror("read");
    exit(1);
}
// buffer now contains count bytes read from the file
```

- what is O_RDWR?
- what does perror() do?
- **cursor** position in an opened file depends on what functions/system calls you use
 - what about C++?




Standard File Descriptors



Standard File Descriptors

- 0 is stdin (by default, the keyboard)
- 1 is stdout (by default, the display)
- 2 is stderr (by default, the display)

```
main() {  
    char buf[BUFSIZE];  
    int n;  
    const char *note = "Write failed\n";  
  
    while ((n = read(0, buf, sizeof(buf))) > 0)  
        if (write(1, buf, n) != n) {  
            (void)write(2, note, strlen(note));  
            exit(EXIT_FAILURE);  
        }  
    return (EXIT_SUCCESS);  
}
```

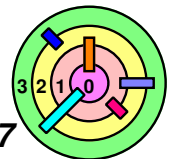


Back to Primes

➡ Have our primes program write out the solution, i.e., the `primes[]` array

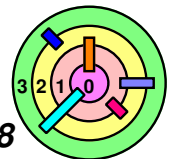
```
int nprimes;
int *prime;
int main(int argc, char *argv[]) {
    ...
    for (i=1; i<nprimes; i++) {
        ...
    }
    if (write(1, prime, nprimes*sizeof(int)) == -1) {
        perror("primes output");
        exit(1);
    }
    return(0);
}
```

➡ the output is not readable by human



Human-Readable Output

```
int nprimes;  
int *prime;  
int main(int argc, char *argv[]) {  
    ...  
    for (i=1; i<nprimes; i++) {  
        ...  
    }  
    for (i=0; i<nprimes; i++) {  
        printf("%d\n", prime[i]);  
    }  
    return(0);  
}
```




Allocation of File Descriptors

- ➡ Whenever a process requests a new file descriptor, the lowest numbered file descriptor not already associated with an open file is selected; thus

```
#include <fcntl.h>
#include <unistd.h>
...
close(0);
fd = open("file", O_RDONLY);
```

- ➡ will always associate "file" with file descriptor 0 (assuming that the open succeeds)



Running It


```

if (fork() == 0) {
    /* set up file descriptor 1 in the child process */
    close(1);
    if (open("/home/bc/Output", O_WRONLY) == -1) {
        perror("/home/bc/Output");
        exit(1);
    }
    execl("/home/bc/bin/primes", "primes", "300", 0);
    exit(1);
}
/* parent continues here */
while(pid != wait(0)) /* ignore the return code */
    ;

```

- ⇒ close(1) removes file descriptor 1 from *extended address space*
- ⇒ file descriptors are allocated *lowest first* on open()
- ⇒ *extended address space* survives *execs*
- ⇒ new code is same as running

```
% primes 300 > /home/bc/Output
```



I/O Redirection


```
% primes 300 > /home/bc/Output
```

- ➡ The ">" parameter in a shell command that instructs the command shell to *redirect* the output to the given file
- If ">" weren't there, the output would go to the display

- ➡ Can also redirect input

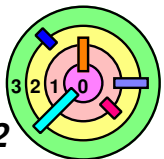
```
% cat < /home/bc/Output
```

- when the "cat" program reads from file descriptor 0, it would get the data bytes from the file "/home/bc/Output"

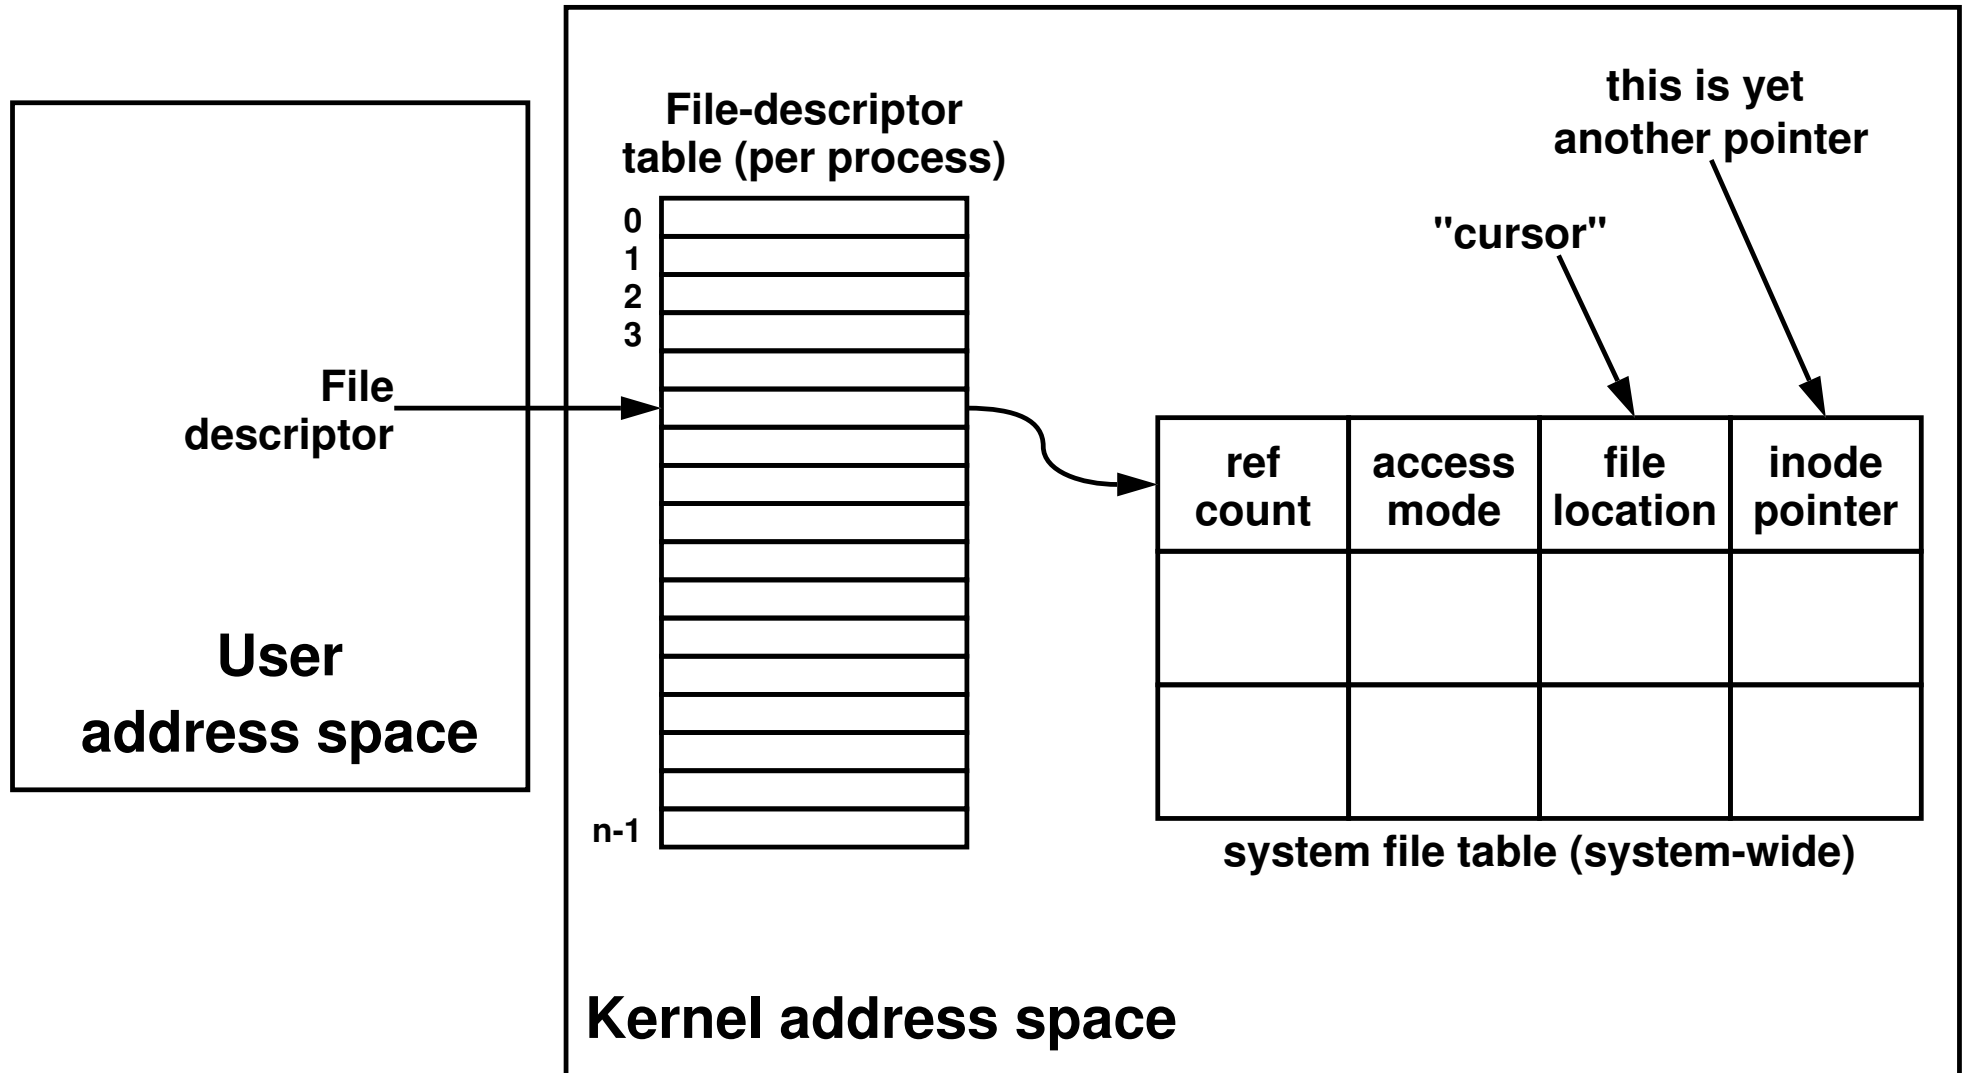


File-Descriptor Table

- ➡ A file descriptor refers not just to a file
 - ➡ it also refers to the *process's* current *context* for that file
 - includes how the file is to be accesses (how `open()` was invoked)
 - *cursor* position
- ➡ *Context* information must be maintained by the OS and not directly by the user program
 - ➡ let's say a user program opened a file with `O_RDONLY`
 - ➡ later on it calls `write()` using the opened file descriptor
 - ➡ how does the OS knows that it doesn't have write access?
 - stores `O_RDONLY` in context
 - ➡ if the user program can manipulate the context, it can change `O_RDONLY` to `O_RDWR`
 - ➡ therefore, user program must not have access to context!
 - all it can see is the handle
 - the handle is an *index* into an array maintained for the process in kernel's address space



File-Descriptor Table



- context is not stored directly into the file-descriptor table
- one-level of *indirection*

Put It All Together

open ()

```
→ int fd;  
   fd = open("foo.txt");  
   char buf[512];  
   read(fd, buf, 100);  
   close(fd);
```


Applications

OS

Process
Subsystem

Files
Subsystem

...



Put It All Together

open ()

```
→ int fd;  
   fd = open("foo.txt");  
   char buf[512];  
   read(fd, buf, 100);  
   close(fd);
```

trap


Applications

OS


Process
Subsystem

Files
Subsystem

...



Put It All Together





→
`int fd;
fd = open("foo.txt");
char buf[512];
read(fd, buf, 100);
close(fd);`


trap

Applications


OS



...



Put It All Together





```
int fd;  
→ fd = open("foo.txt");  
char buf[512];  
read(fd, buf, 100);  
close(fd);
```



trap

Applications


OS



...





Put It All Together




```
int fd;  
→ fd = open("foo.txt");  
char buf[512];  
read(fd, buf, 100);  
close(fd);
```

Applications


OS



...





Put It All Together




```
int fd;  
→ fd = open("foo.txt");  
char buf[512];  
read(fd, buf, 100);  
close(fd);
```

Applications

OS



...




Put It All Together

`open()`
`read()`


```
int fd;  
fd = open("foo.txt");  
char buf[512];  
→ read(fd, buf, 100);  
close(fd);
```

Applications

OS



Put It All Together





```
int fd;  
fd = open("foo.txt");  
char buf[512];  
→ read(fd, buf, 100);  
close(fd);
```


Applications

OS


trap



...



Put It All Together





```
int fd;  
fd = open("foo.txt");  
char buf[512];  
→ read(fd, buf, 100);  
close(fd);
```


≤ 100 bytes

Applications

OS



...




Put It All Together

```
open()  
read()  
close()
```

```
int fd;  
fd = open("foo.txt");  
char buf[512];  
read(fd, buf, 100);  
→ close(fd);
```

Applications

OS



Put It All Together

`open()`
`read()`
`close()`



```
int fd;  
fd = open("foo.txt");  
char buf[512];  
read(fd, buf, 100);  
close(fd);
```





trap

Applications

OS



...



Put It All Together

`open()`
`read()`
`close()`

```
int fd;  
fd = open("foo.txt");  
char buf[512];  
read(fd, buf, 100);  
close(fd);
```



Applications

OS

Process
Subsystem

Files
Subsystem

...



Put It All Together

```
open()  
read()  
close()
```

```
int fd;  
fd = open("foo.txt");  
char buf[512];  
read(fd, buf, 100);  
→ close(fd);
```



Applications

OS

Process
Subsystem

Files
Subsystem

...




Redirecting Output ... Twice

➡ Every calls to `open()` creates a new entry in the system file table

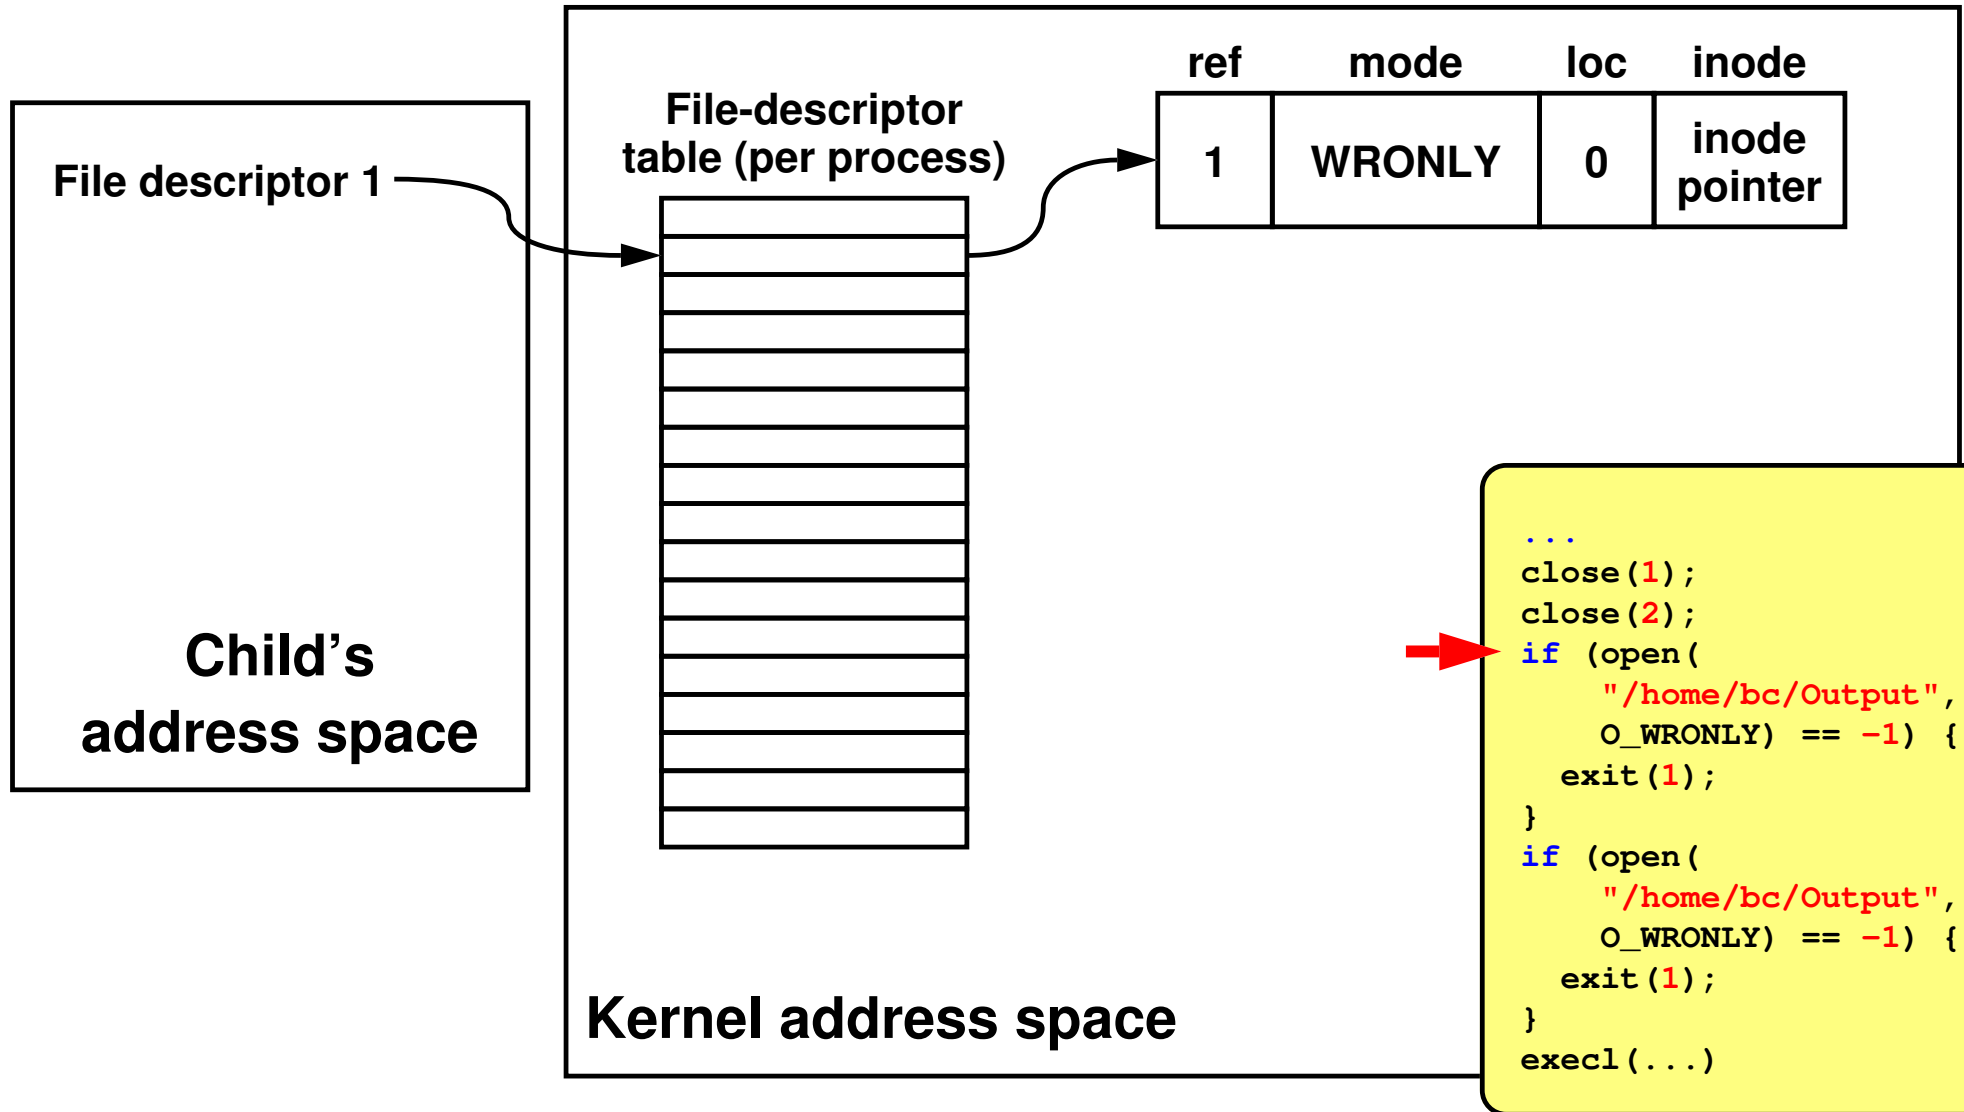
```
if (fork() == 0) {
    /* set up file descriptors 1 and 2 in the child
       process */
    close(1);
    close(2);
    if (open("/home/bc/Output", O_WRONLY) == -1) {
        exit(1);
    }
    if (open("/home/bc/Output", O_WRONLY) == -1) {
        exit(1);
    }
    execl("/home/bc/bin/program", "program", 0);
    exit(1);
}
/* parent continues here */
```

⇒ stdout and stderr both go into the same file


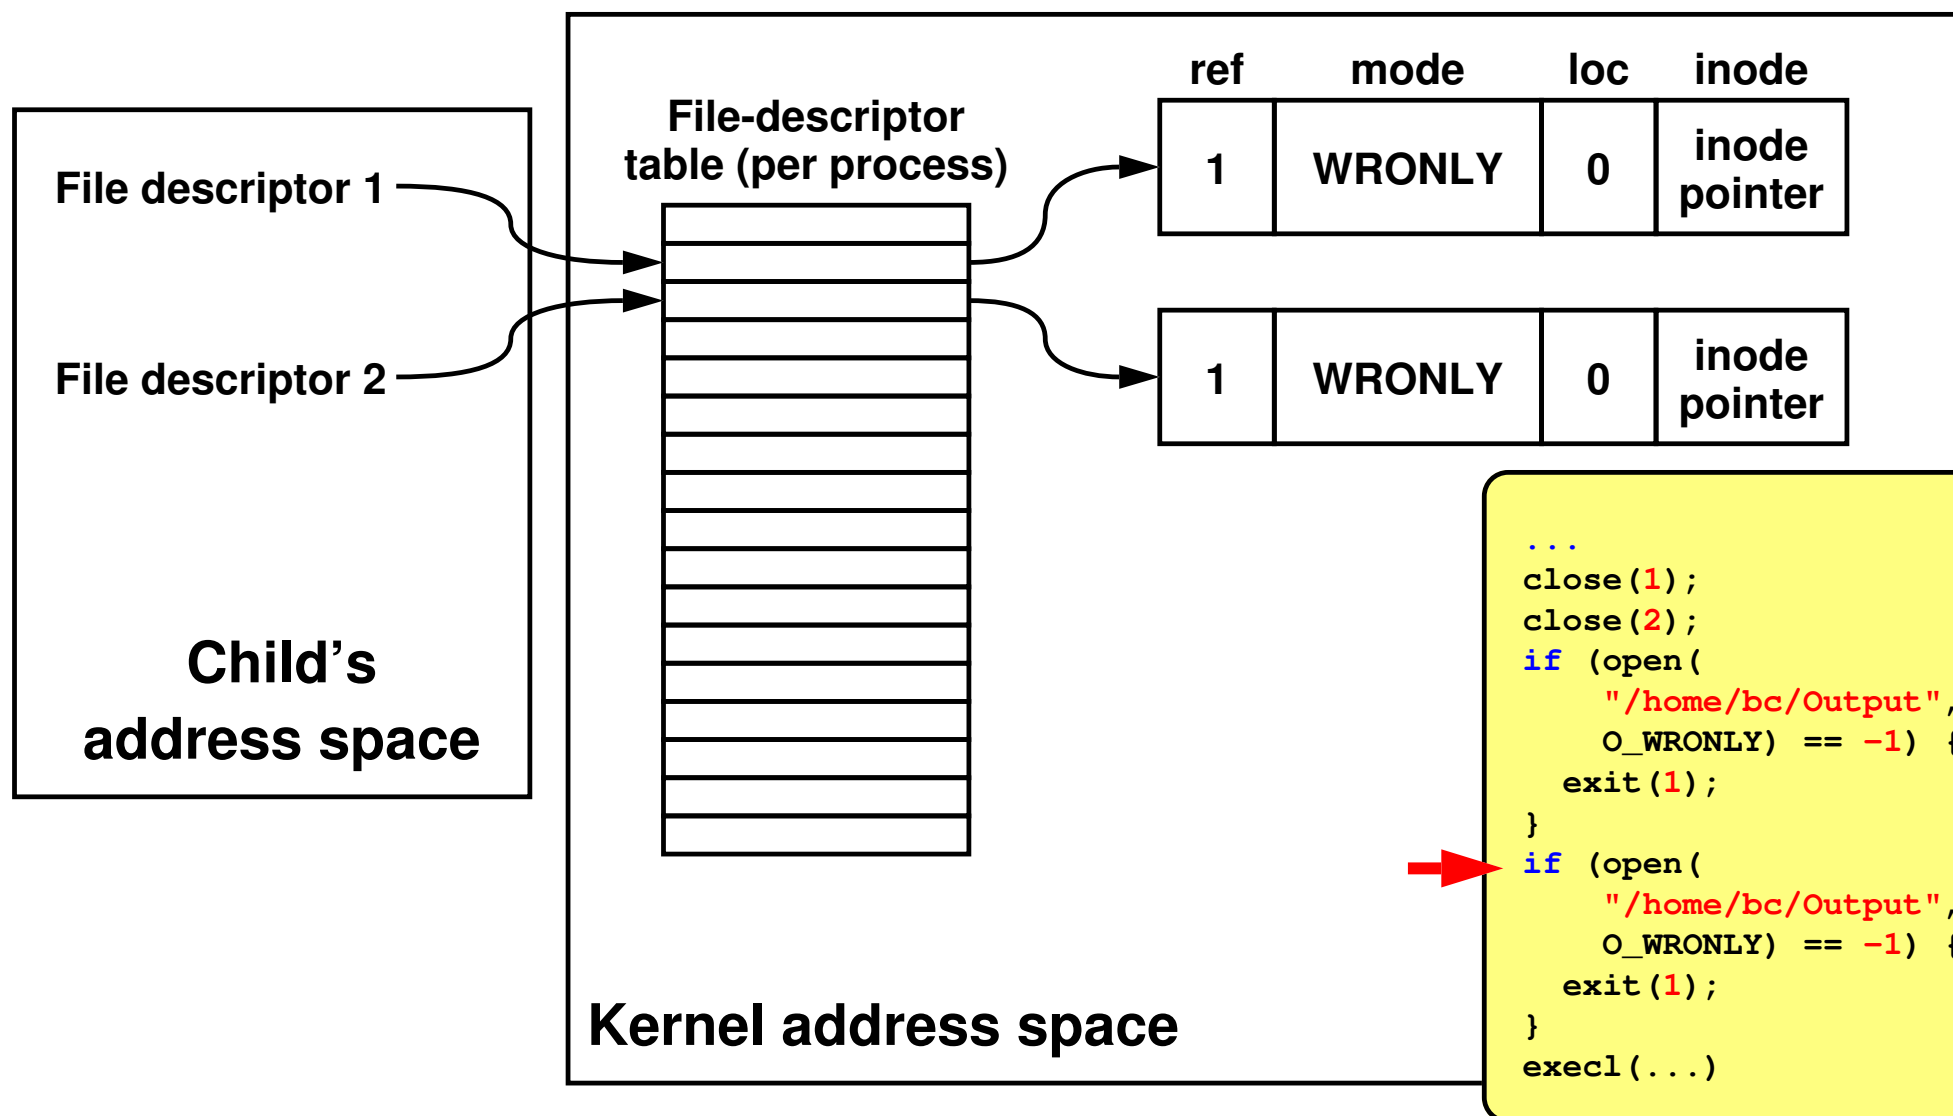
○ would it cause any problem?



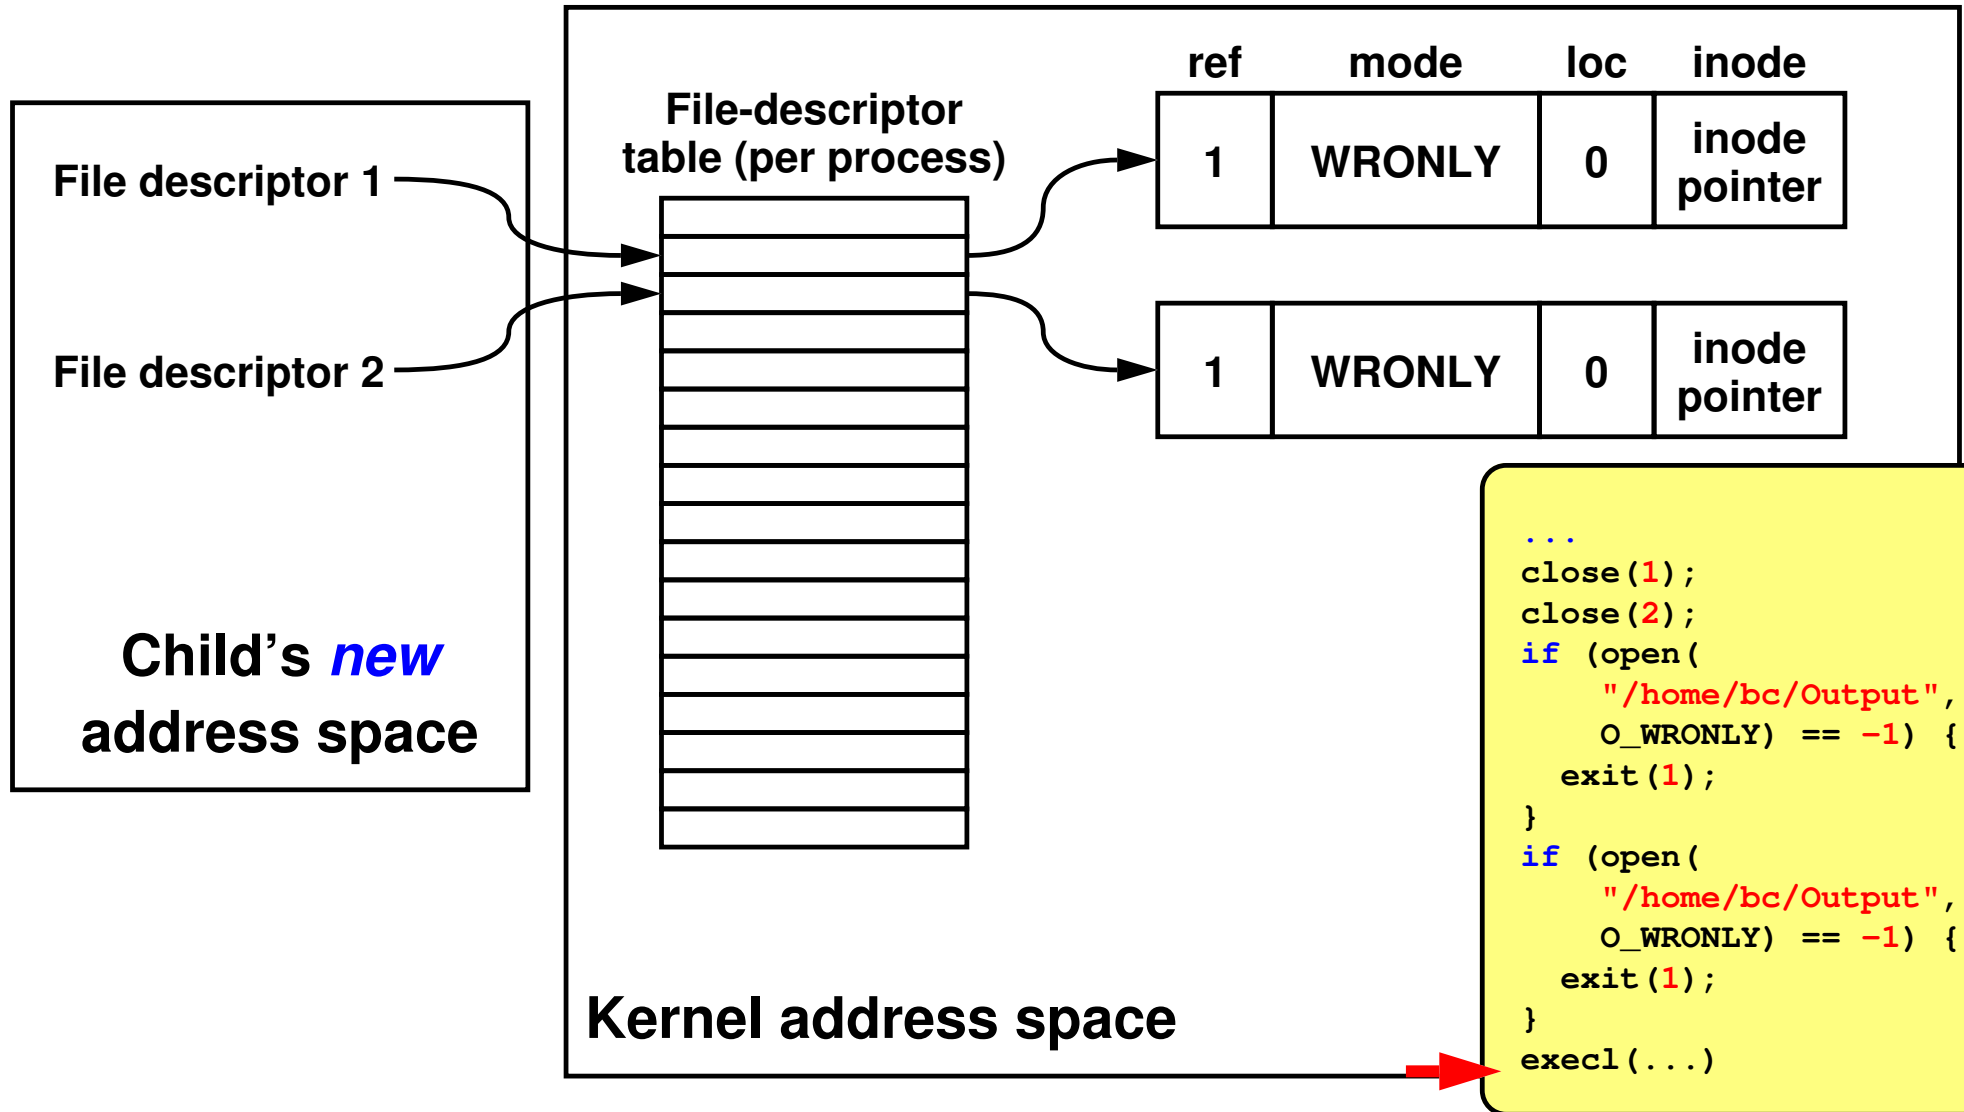
Redirected Output




Redirected Output



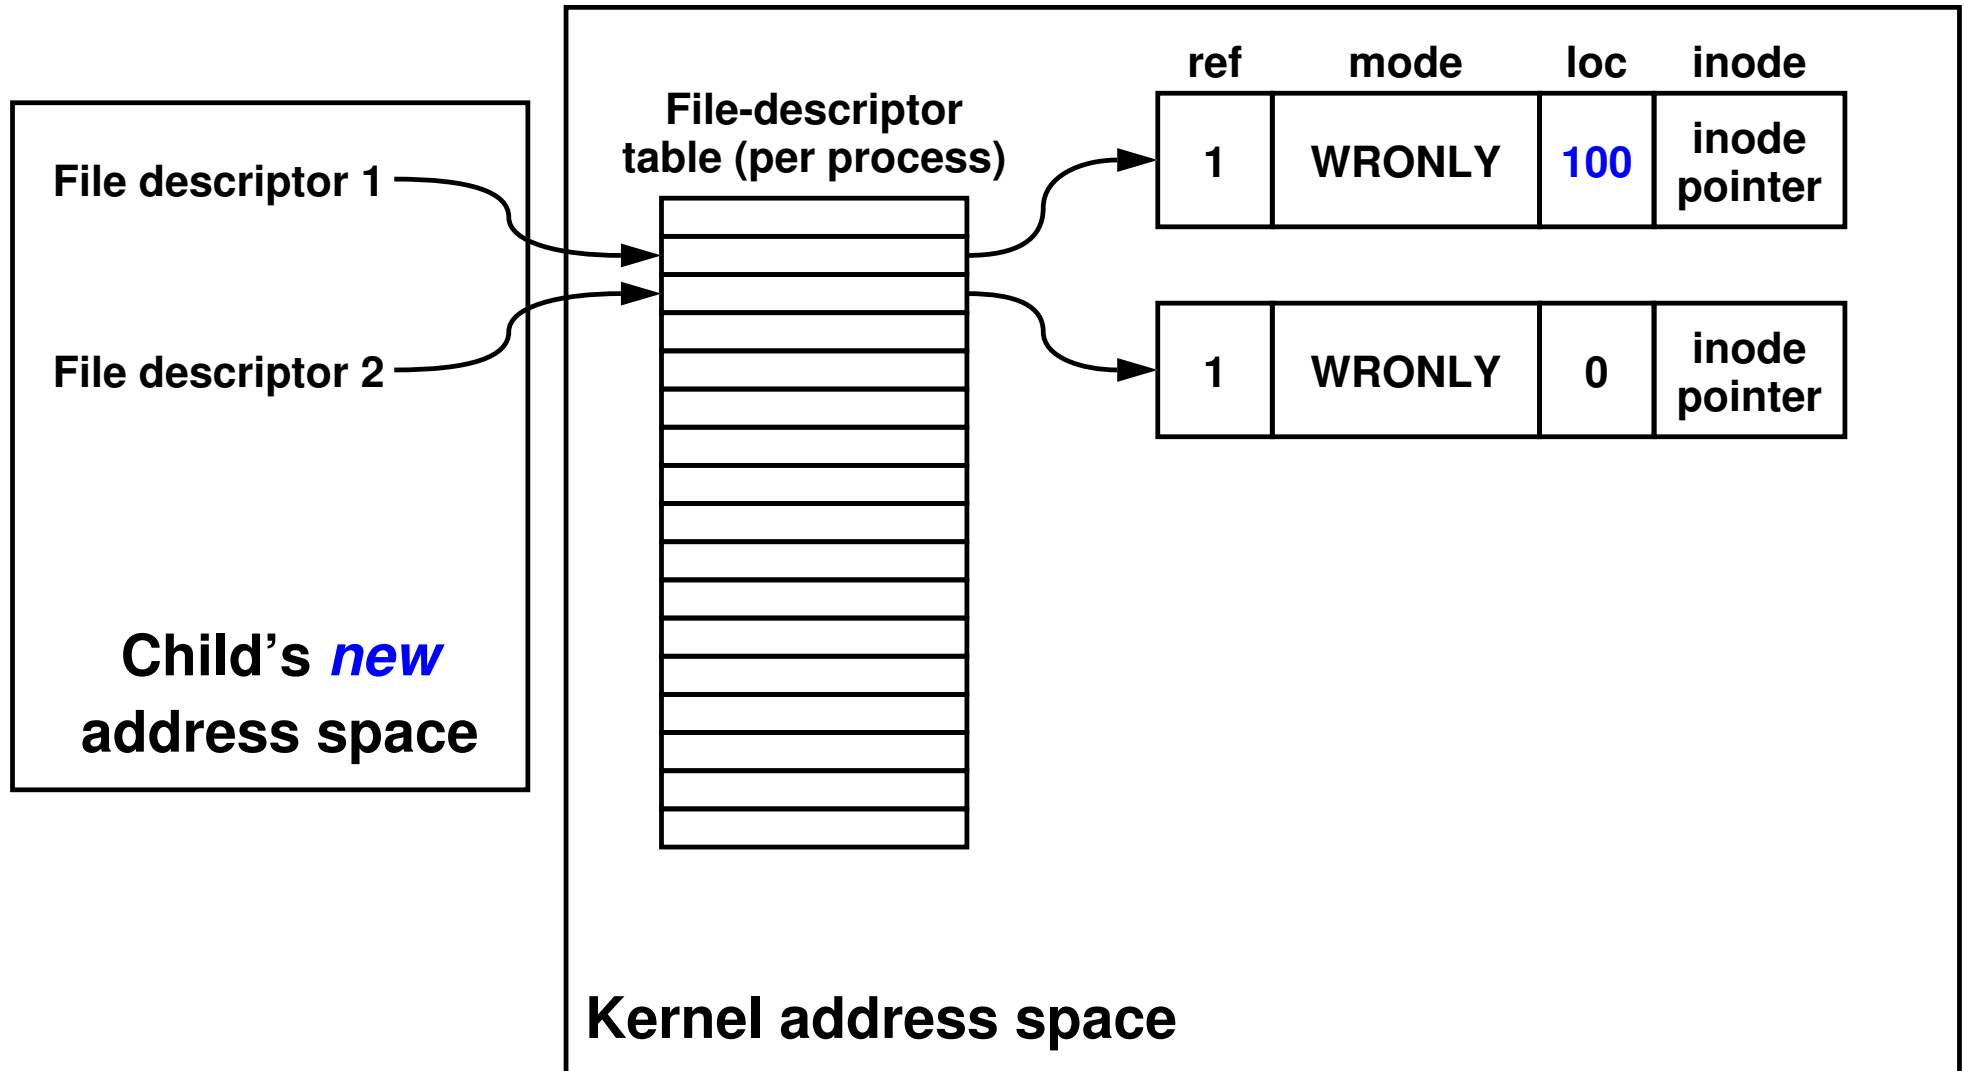
Redirected Output



— remember, extended address space survives execs



Redirected Output After Writing 100 Bytes




- `write()` to `fd=2` will wipe out data in the first 100 bytes
- that may not be the intent


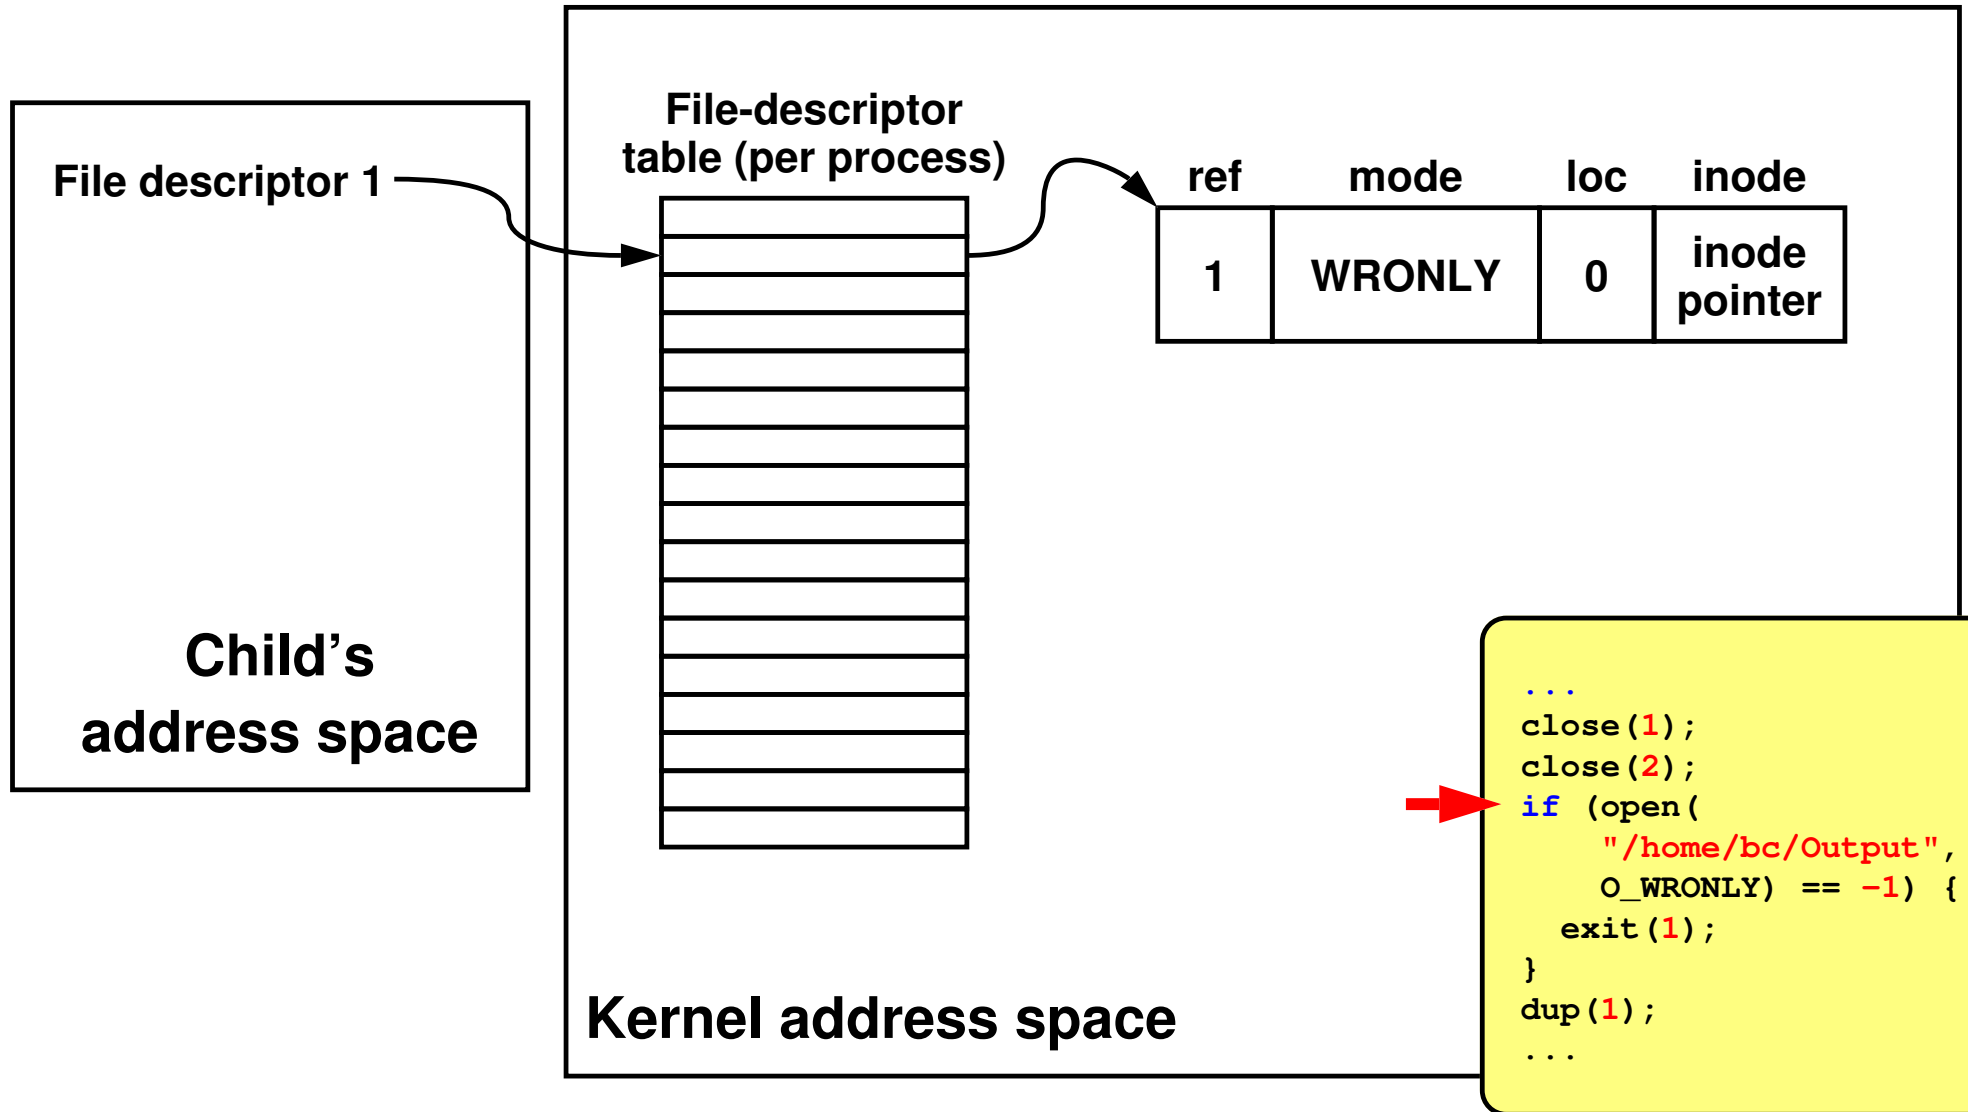
Sharing Context Information

```
if (fork() == 0) {  
    /* set up file descriptors 1 and 2 in the child  
       process */  
    close(1);  
    close(2);  
    if (open("/home/bc/Output", O_WRONLY) == -1) {  
        exit(1);  
    }  
    dup(1);  
    execl("/home/bc/bin/program", "program", 0);  
    exit(1);  
}  
/* parent continues here */
```

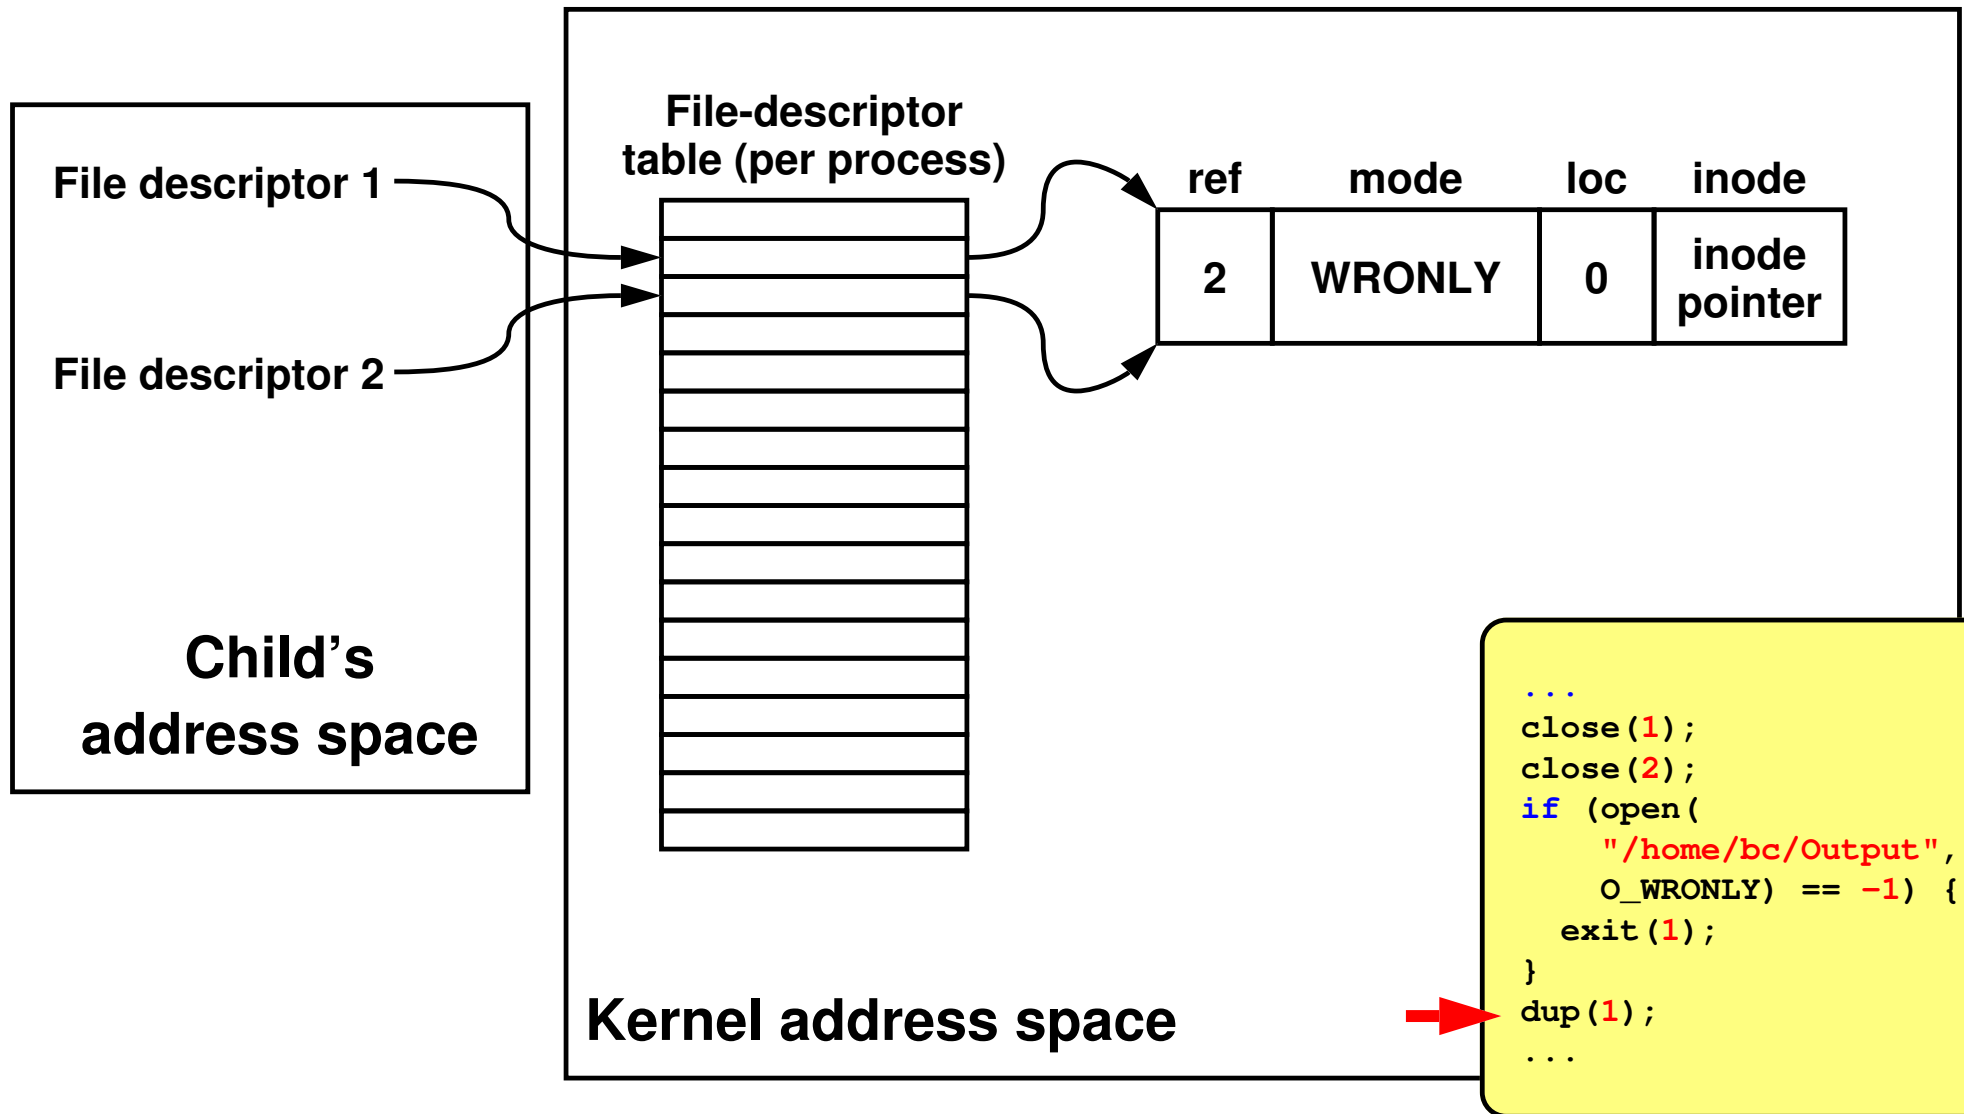
- use the `dup()` system call to *share* context information
 - if that's what you want



Redirected Output After Dup



Redirected Output After Dup




Fork and File Descriptors


➡ When `fork()` is called, the child process gets a *copy* of the parent's file descriptor table

```
int logfile = open("log", O_WRONLY);
if (fork() == 0) {
    /* child process computes something, then does: */
    write(logfile, LogEntry, strlen(LogEntry));
    ...
    exit(0);
}
/* parent process computes something, then does: */
write(logfile, LogEntry, strlen(LogEntry));
...
```


- remember, extended address space survives execs
 - also `fork()`



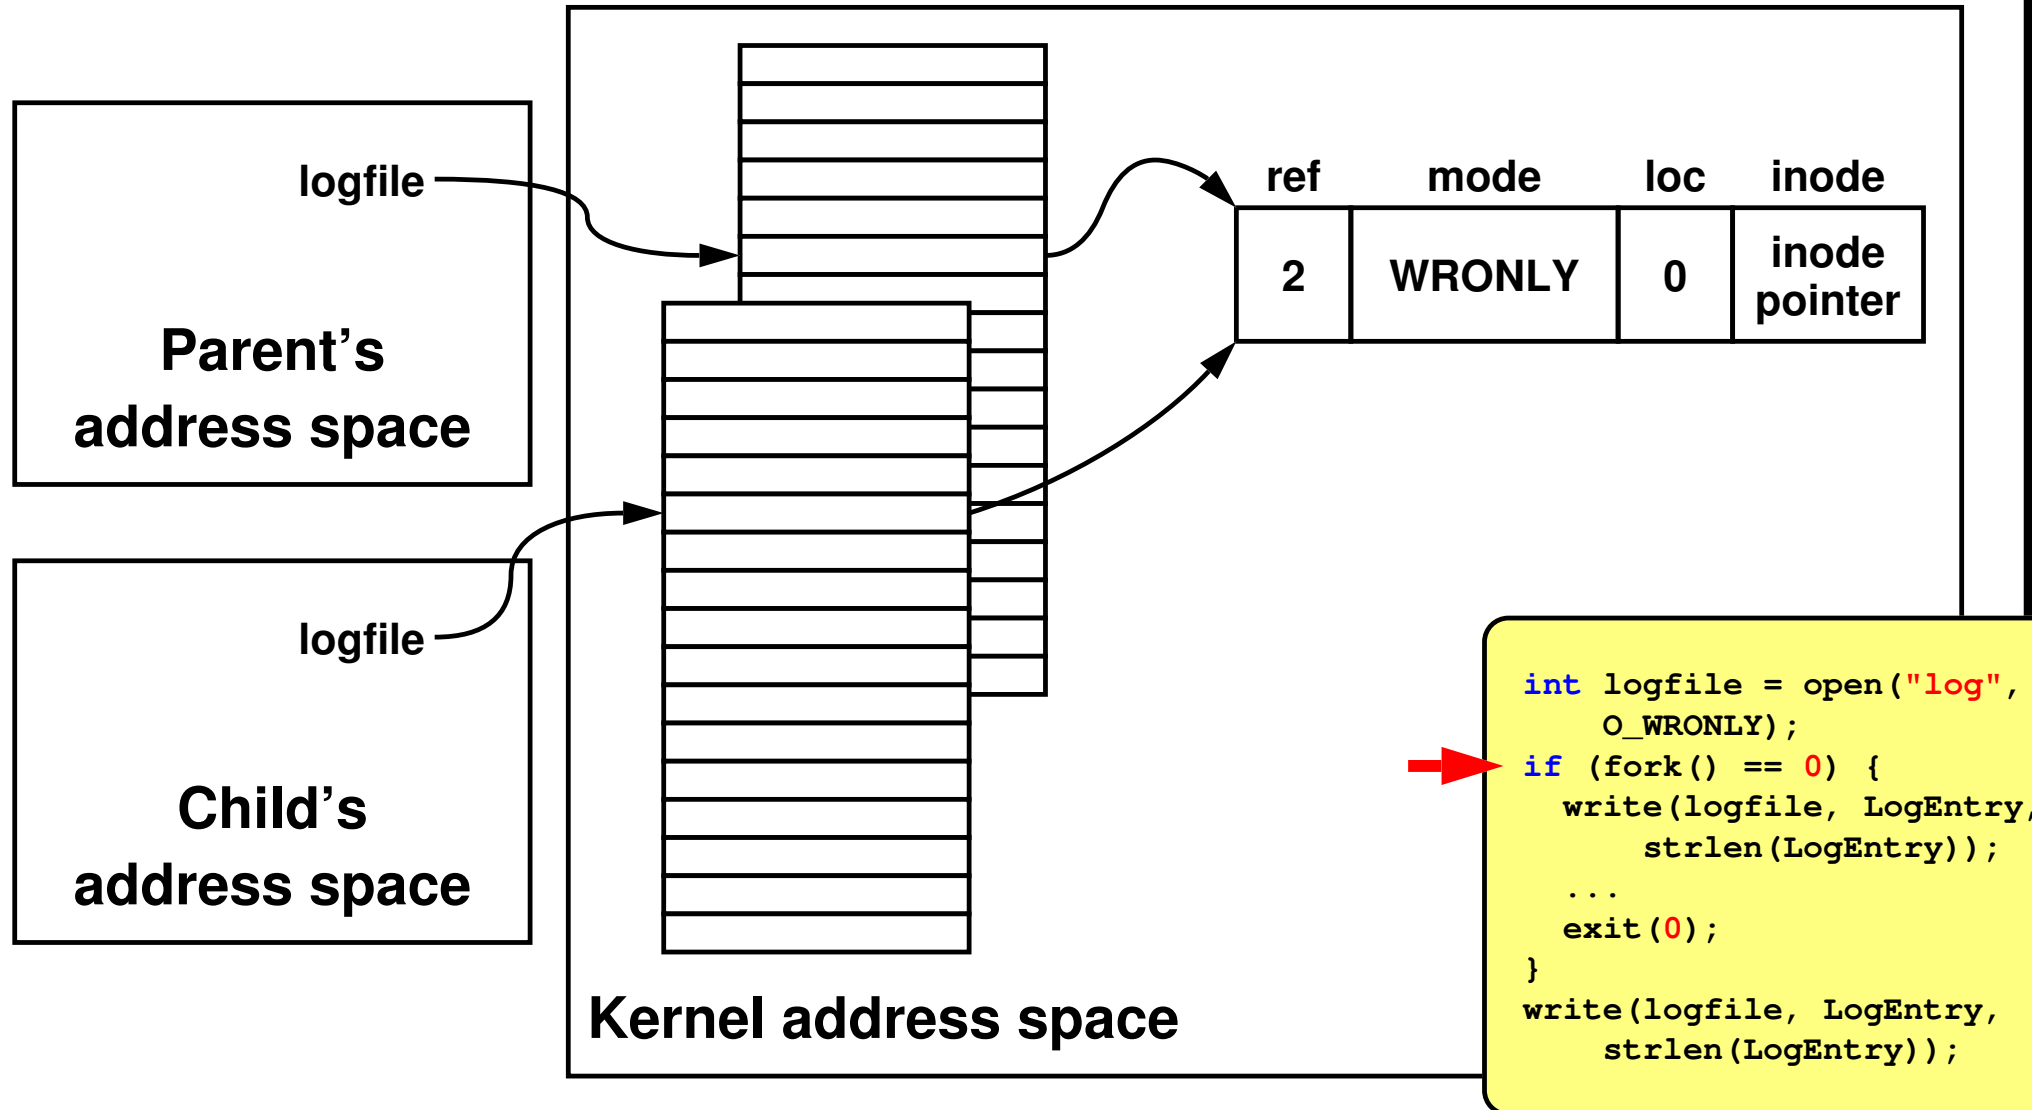
File Descriptors After Fork




- parent and child processes get *separate* file descriptor table but *share* extended address space



File Descriptors After Fork




- parent and child processes get *separate* file descriptor table but *share* extended address space



Pipes

➡ A pipe is a means for one process to send data to another directly, as if it were writing to a file




- the sending process behaves as if it has a file descriptor to a file that has been opened for writing
- the receiving process behaves as if it has a file descriptor to a file that has been opened for reading

➡ The `pipe()` system call creates a pipe object in the kernel and returns (via an output parameter) the two file descriptors that refer to the pipe


- one, set for write-only, refers to the input side
- the other, set for read-only, refers to the output side
- a pipe has no name, cannot be passed to another process

Pipes


```
int p[2]; // array to hold pipe's file descriptors
pipe(p); // creates a pipe, assume no errors
// p[0] refers to the read/output end of the pipe
// p[1] refers to the write/input end of the pipe
if (fork() == 0) {
    char buf[80];
    close(p[1]); // not needed by the child
    while (read(p[0], buf, 80) > 0) {
        // use data obtained from parent
        ...
    }
    exit(0); // child done
} else {
    char buf[80];
    close(p[0]); // not needed by the parent
    for (;;) {
        // prepare data for child
        ...
        write(p[1], buf, 80);
    }
}
```




Pipes




- parent creates a pipe object in the kernel



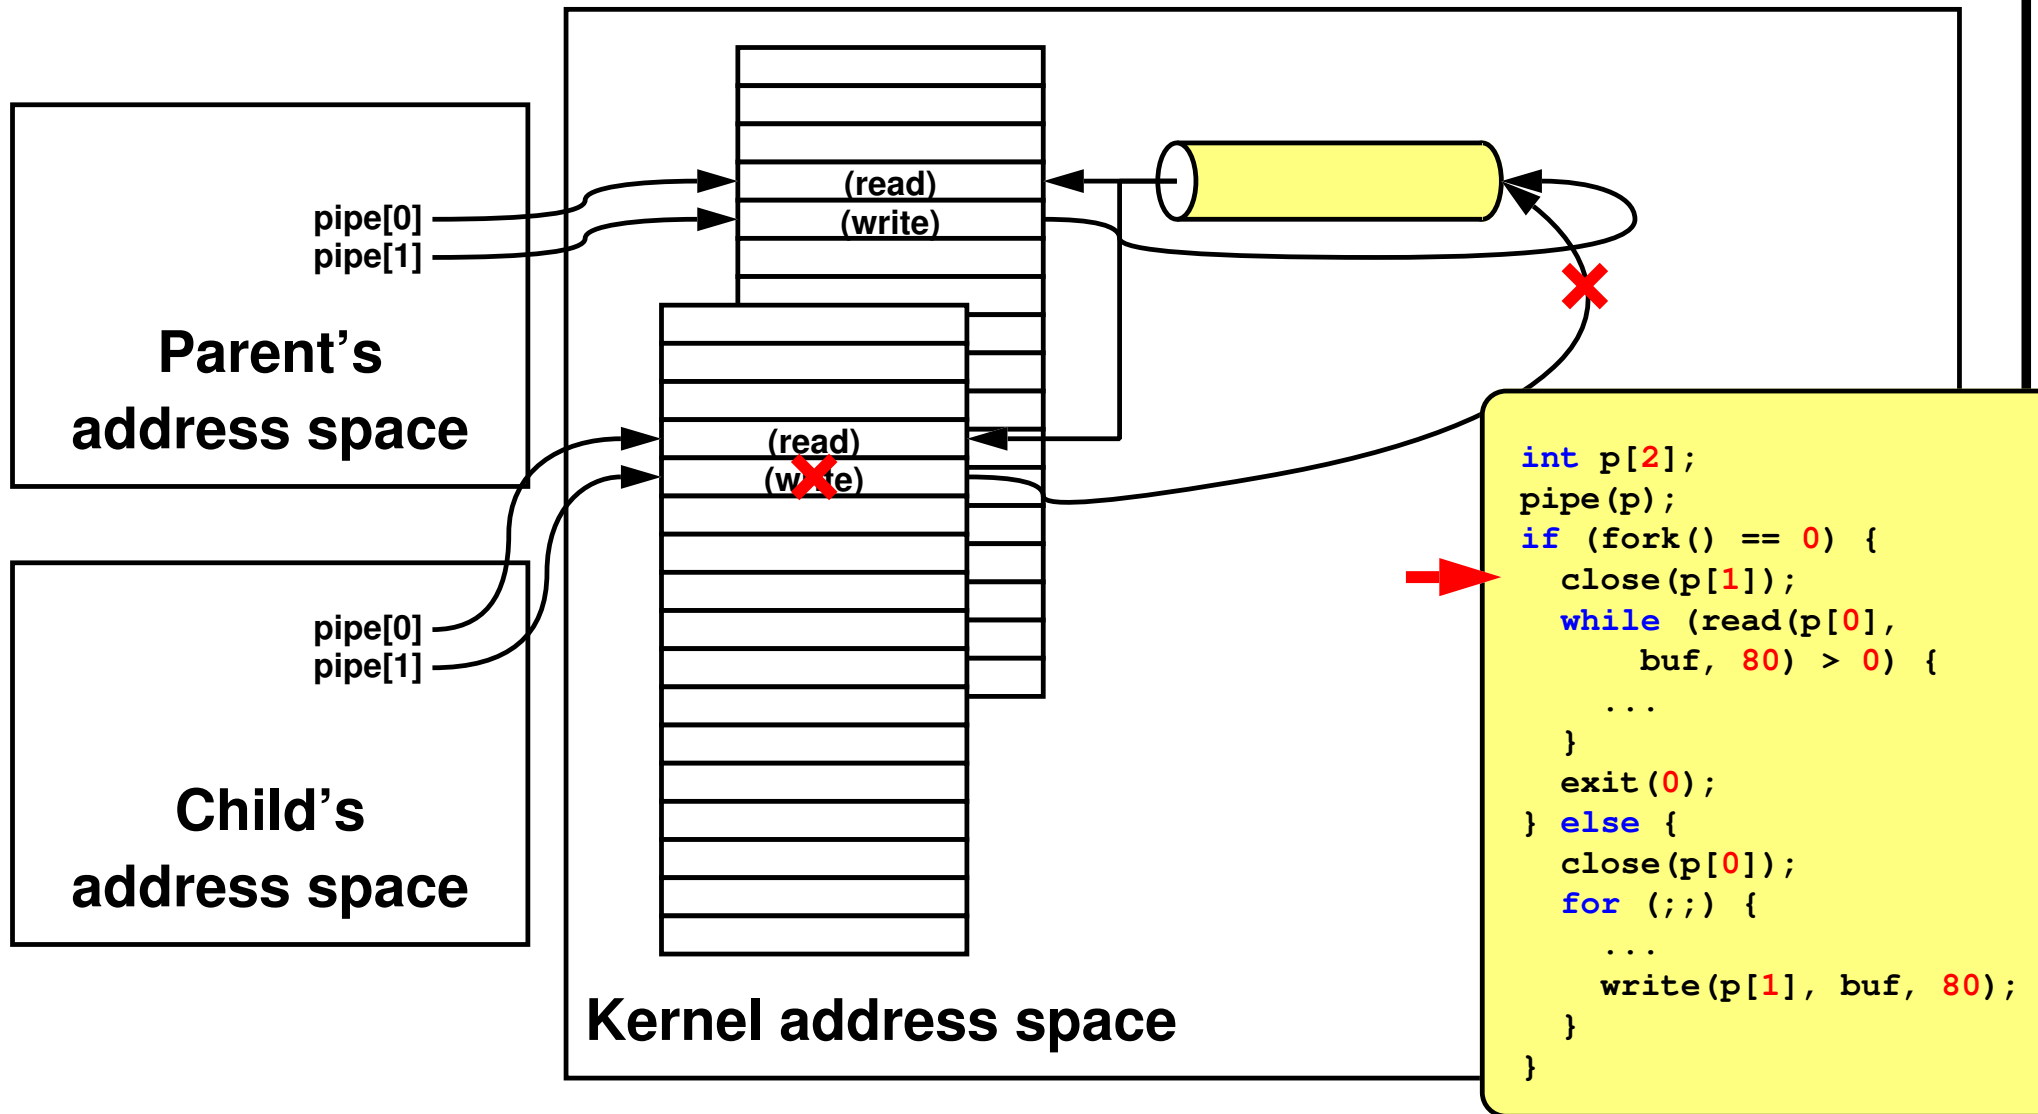
Pipes




- parent and child processes get *separate* file descriptor table but *share* extended address space




Pipes



- parent and child processes get *separate* file descriptor table but *share* extended address space

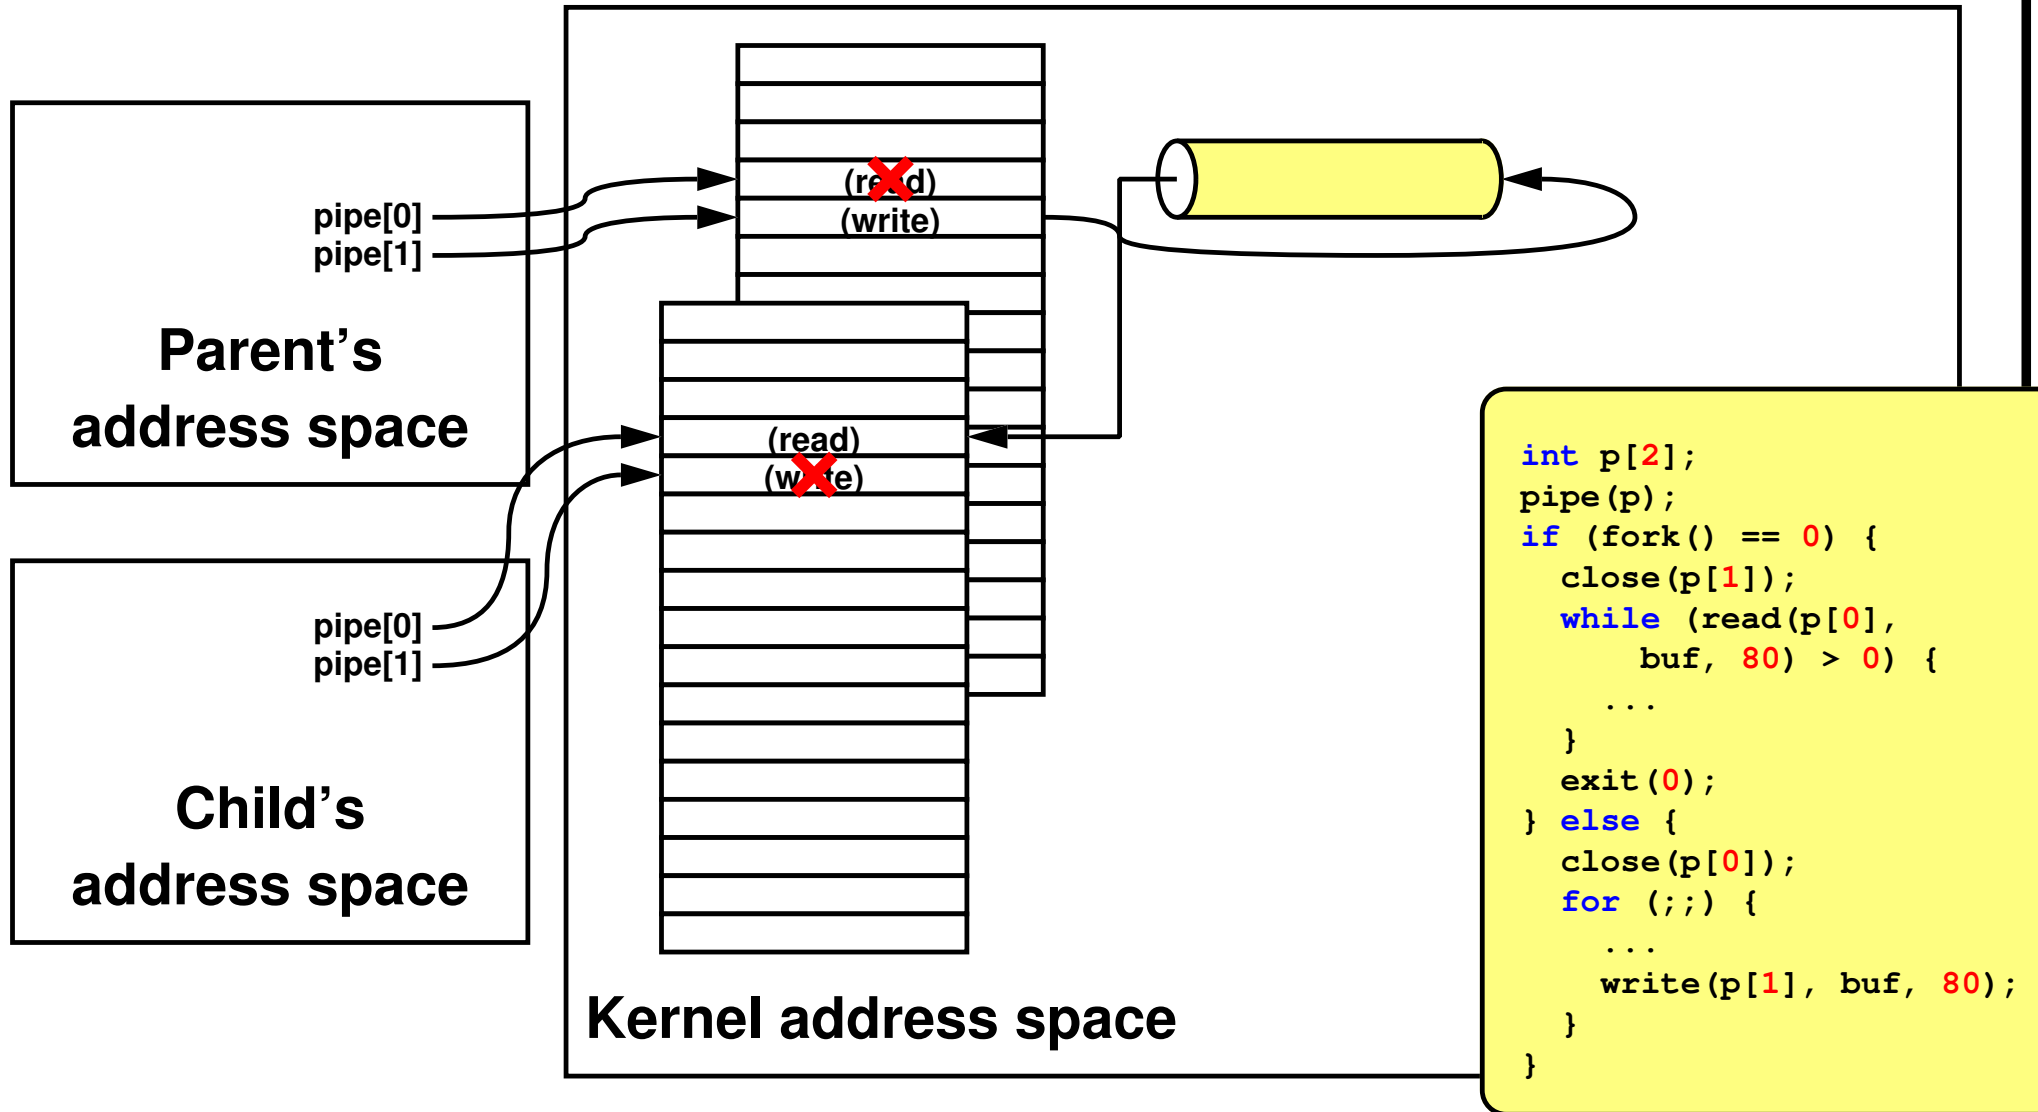


Pipes




- parent and child processes get *separate* file descriptor table but *share* extended address space

Pipes



- parent closes the read-end of the pipe
- child closes the write-end of the pipe



Command Shell


➡ Now you know enough to write a command shell

- execute a command
- redirect I/O
- pipe the output of one program to another

```
cat f0 | ./warmup1 sort
```

- the shell needs to create a pipe
- create two child processes
- in the first child
 - ◆ have `stdout` go to the write-end of the pipe
 - ◆ close the read-end of the pipe
 - ◆ `exec "cat f0"`
- in the 2nd child
 - ◆ have `stdin` come from the read-end of the pipe
 - ◆ close the write-end of the pipe
 - ◆ `exec "./warmup1 sort"`
- run a program in the background

```
primes 1000000 > primes.out &
```




Random Access

```
fd = open("textfile", O_RDONLY);  
// go to last char in file  
fptr = lseek(fd, (off_t)(-1), SEEK_END);  
while (fptr != -1) {  
    read(fd, buf, 1);  
    write(1, buf, 1);  
    fptr = lseek(fd, (off_t)(-2), SEEK_CUR);  
}
```

— "man lseek" gives

`off_t lseek(int fd, off_t offset, int whence);`

- whence can be `SEEK_SET`, `SEEK_CUR`, `SEEK_END`
- if succeeds, returns cursor position (always measured from the beginning of the file)
 - otherwise, returns (-1)
 - `errno` is set to indicate the error
- `read(fd, buf, 1)` advances the cursor position by 1, so we need to move the cursor position back 2 positions



More On Naming

➡ (Almost) everything has a path name

- files
- directories
- devices (known as *special files*)
- keyboards
- displays
- disks
- etc.

➡ Uniformity

```
// opening a normal file
```


```
int file = open("/home/bc/data", O_RDWR);
```

```
// opening a device (one's terminal or window)
```

```
int device = open("/dev/tty", O_RDWR);
```

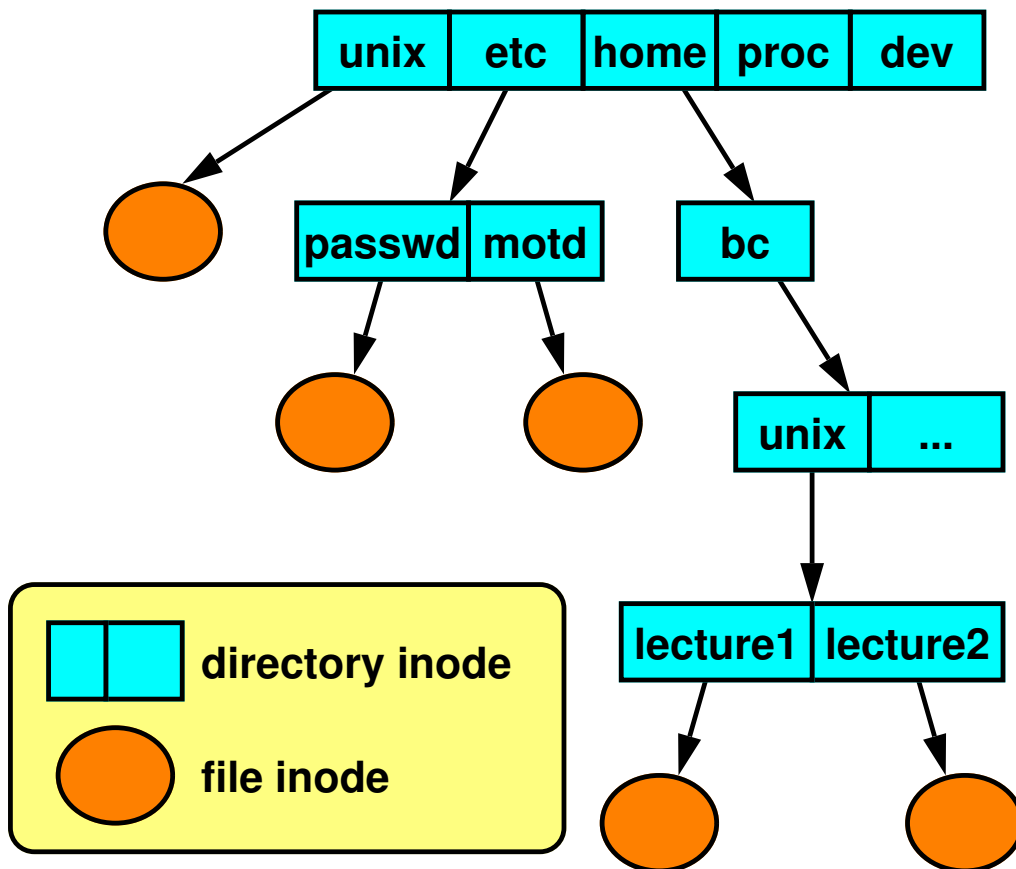
```
int bytes = read(file, buffer, sizeof(buffer));
```

```
write(device, buffer, bytes);
```




Directories

- ➡ A **directory** is a **file**
- interprets differently by the OS as containing **references** to other files/directories
 - a file is represented as an **index node** (or **inode**) in the file system

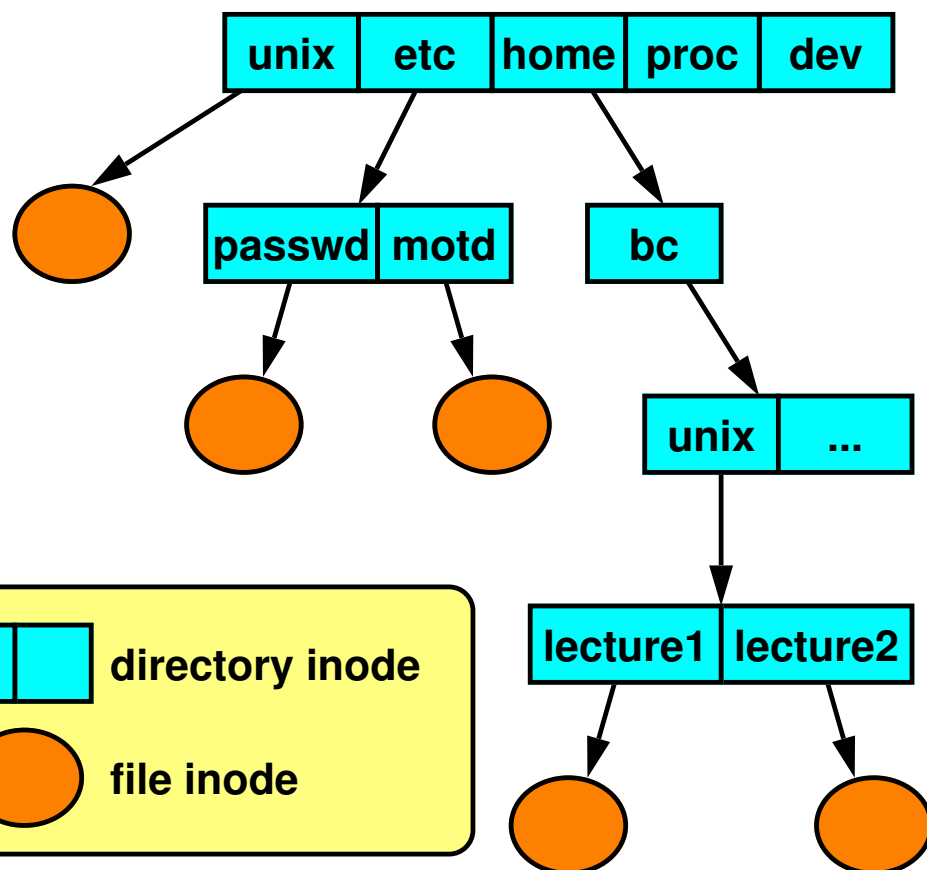


- ➡ A **directory** maps a **file name** to an **inode number**
- maps a string to an integer
 - done inside Virtual File System in `weenix`
- ➡ An **inode** maps an **inode number** to **sectors on disk**
- done inside Actual File System in `weenix`




Directory Representation

- ➡ A root directory entry example
- parent inode number = its own inode number



directory inode



file inode


- ➡ Tree structured hierarchy

Component Name	Inode number
----------------	--------------

directory entry

.	1
..	1
unix	117
etc	4
home	18
proc	36
dev	93

↑ this is what a S5FS directory looks like in weenix





Look Up Inode Number Of A Path

➡ Ex: how do figure out the inode number of `"/home/bc/foo.c"`?


/

home	




Look Up Inode Number Of A Path

➡ Ex: how do figure out the inode number of `"/home/bc/foo.c"`?




Look Up Inode Number Of A Path

➡ Ex: how do figure out the inode number of `"/home/bc/foo.c"`?



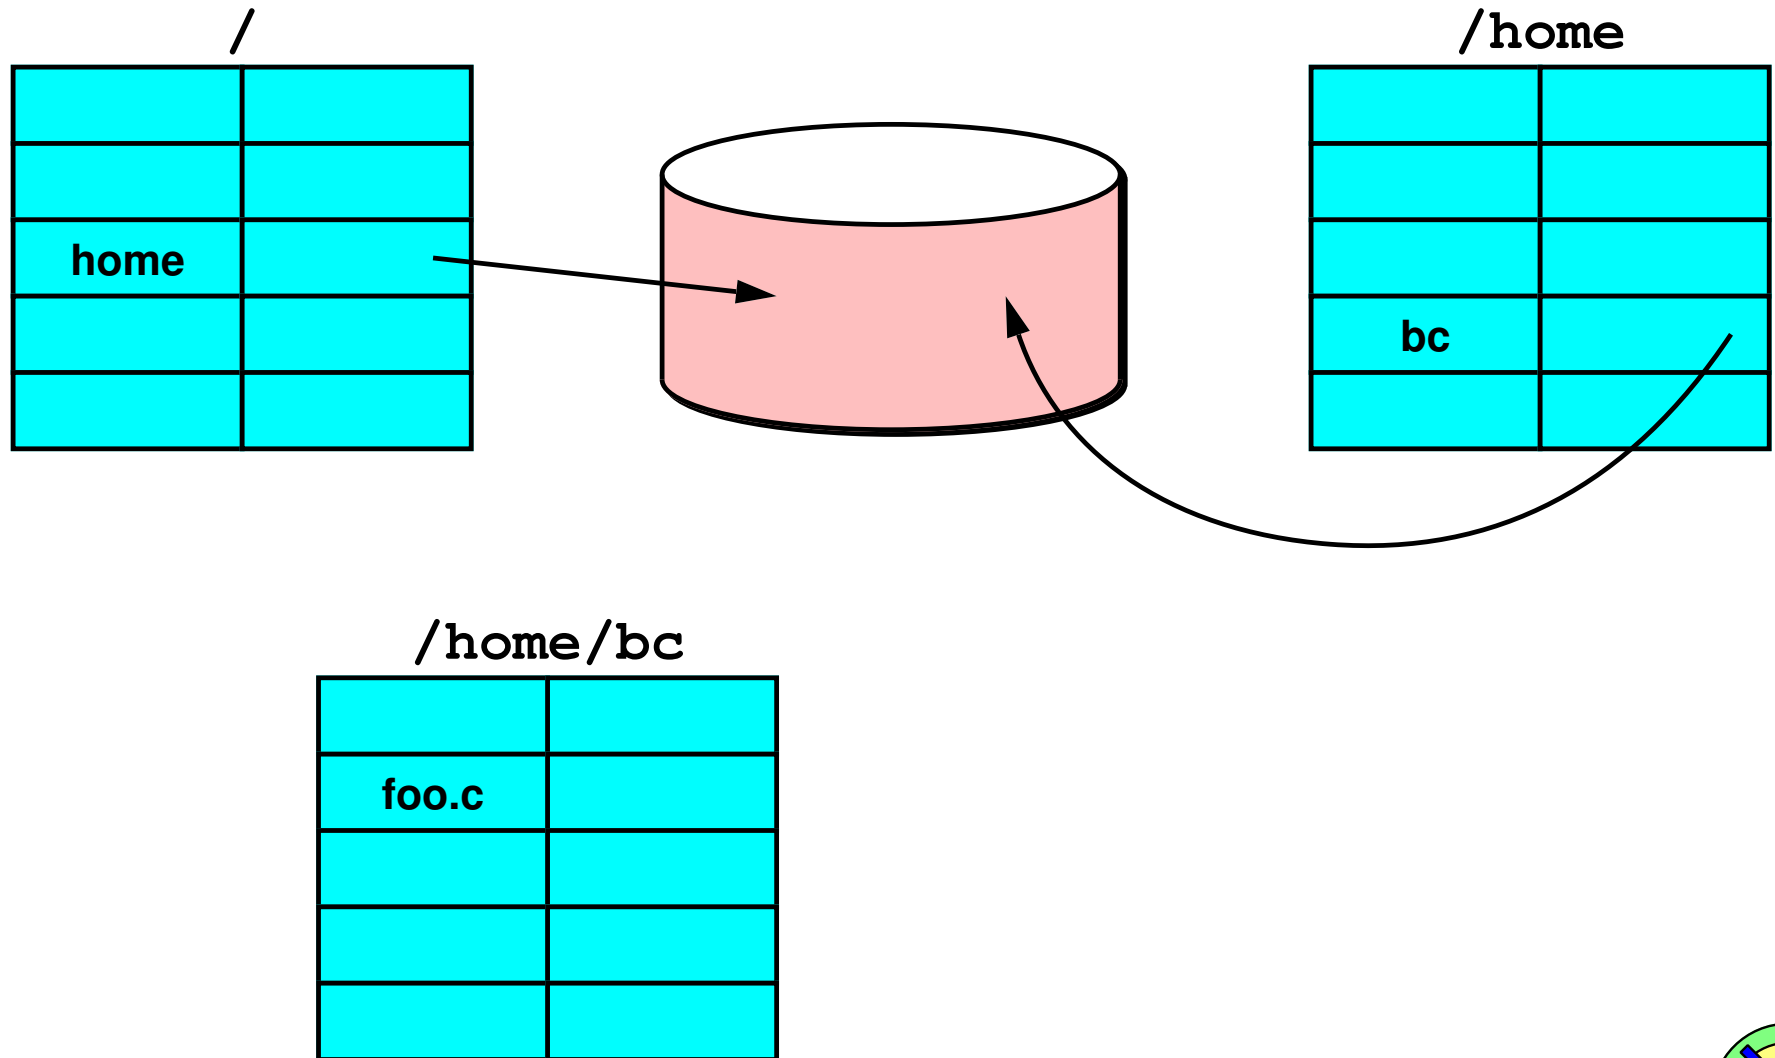
Look Up Inode Number Of A Path

➡ Ex: how do figure out the inode number of `"/home/bc/foo.c"`?



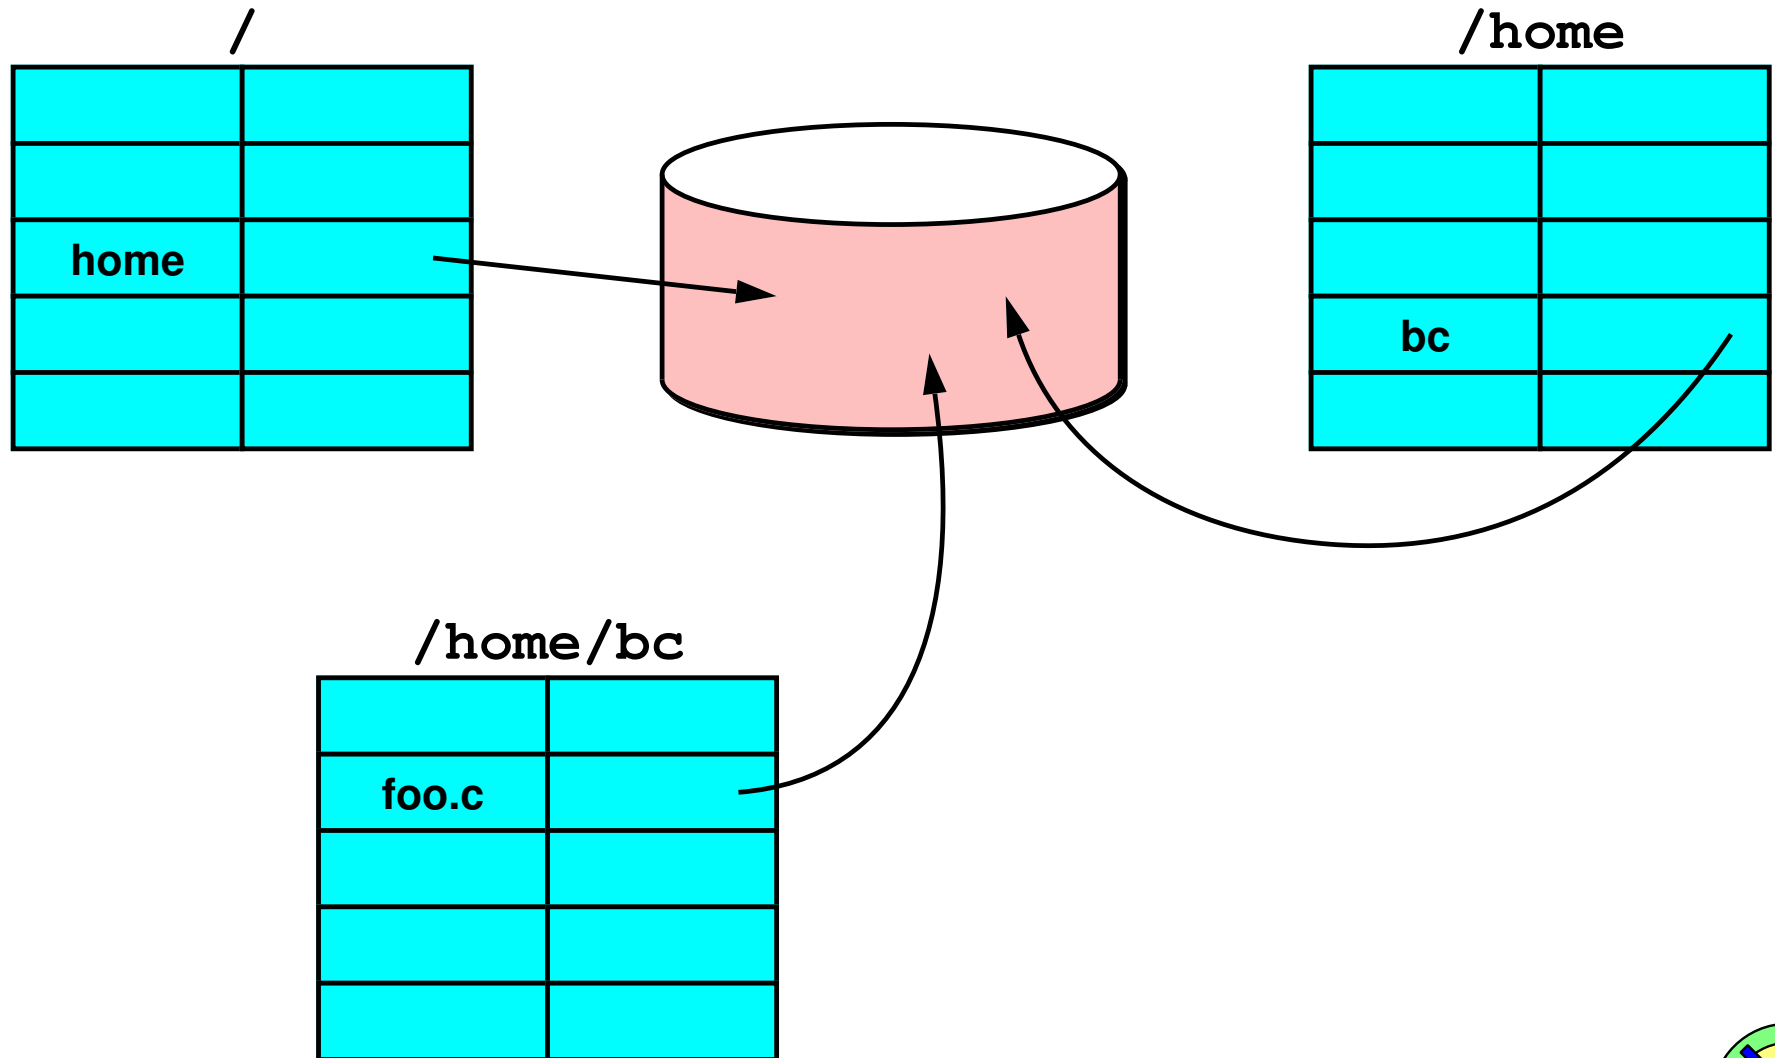
Look Up Inode Number Of A Path

➡ Ex: how do figure out the inode number of `"/home/bc/foo.c"`?




Look Up Inode Number Of A Path

➡ Ex: how do figure out the inode number of `"/home/bc/foo.c"`?



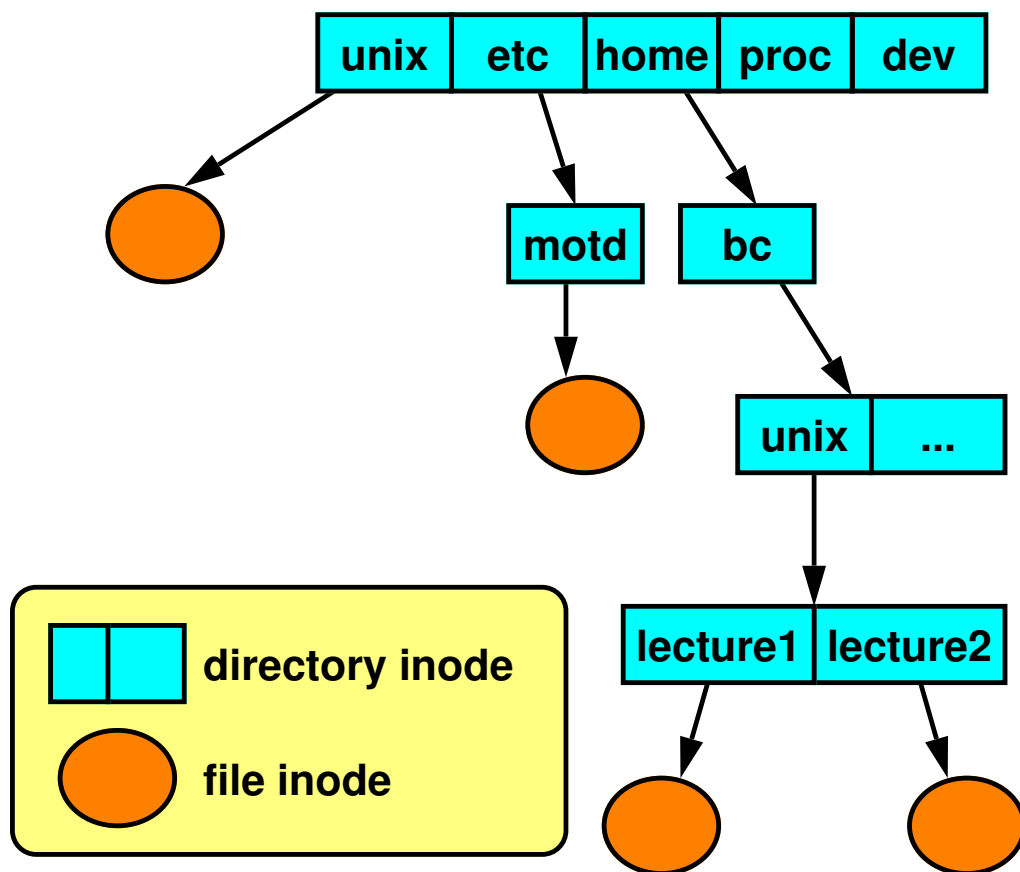
Directory Hierarchy

- ➡ Unix and many other OSes allow limited deviation from trees
 - ➡ *hard links*
 - reference to a *file* (not a directory) in one directory that also appears in another
 - using the `link()` system call or the "ln" shell command
 - ➡ *soft links* or *symbolic links*
 - a special kind of *file* containing the *name* of another file or directory
 - using the `symlink()` system call or the "ln -s" shell command
- ➡ Why hard link cannot be used on a directory?
 - ➡ to avoid cycles
 - ➡ Unix directory hierarchy can be viewed as a *directed acyclic graph* (DAG)
 - can be traversed efficiently




Hard Links

```
% ln /unix /etc/image
```



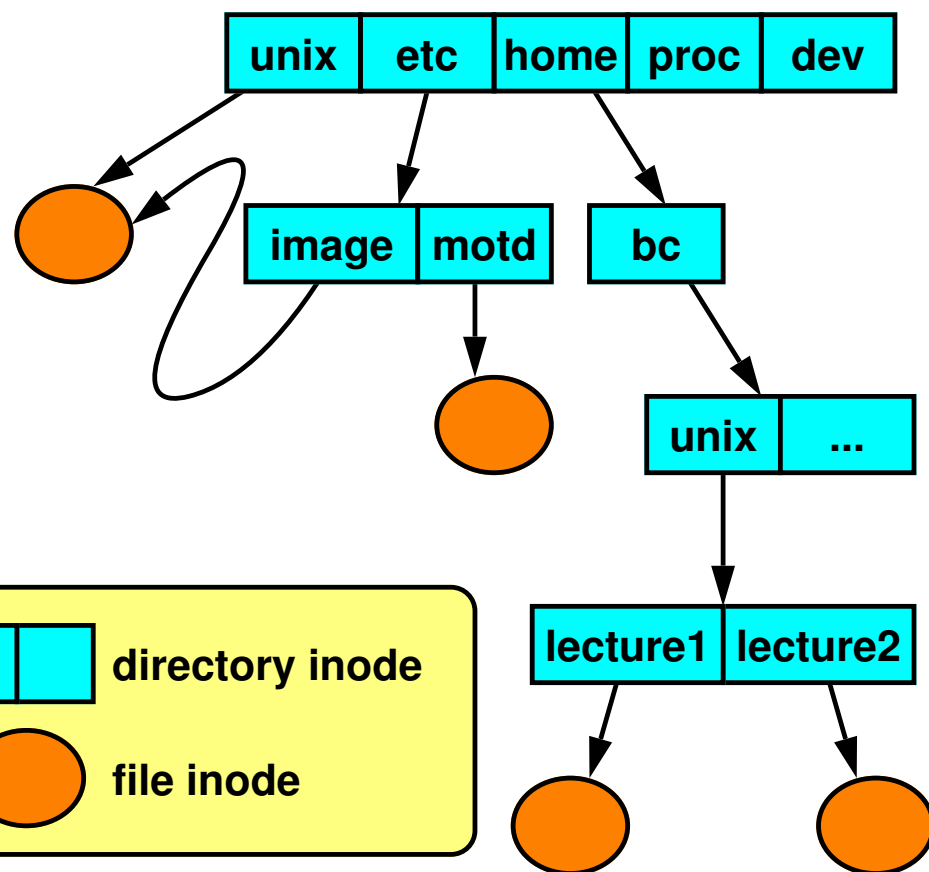
.	1
..	1
unix	117
etc	4
home	18
proc	36
dev	93

.	4
..	1
motd	33




Hard Links

```
% ln /unix /etc/image
```




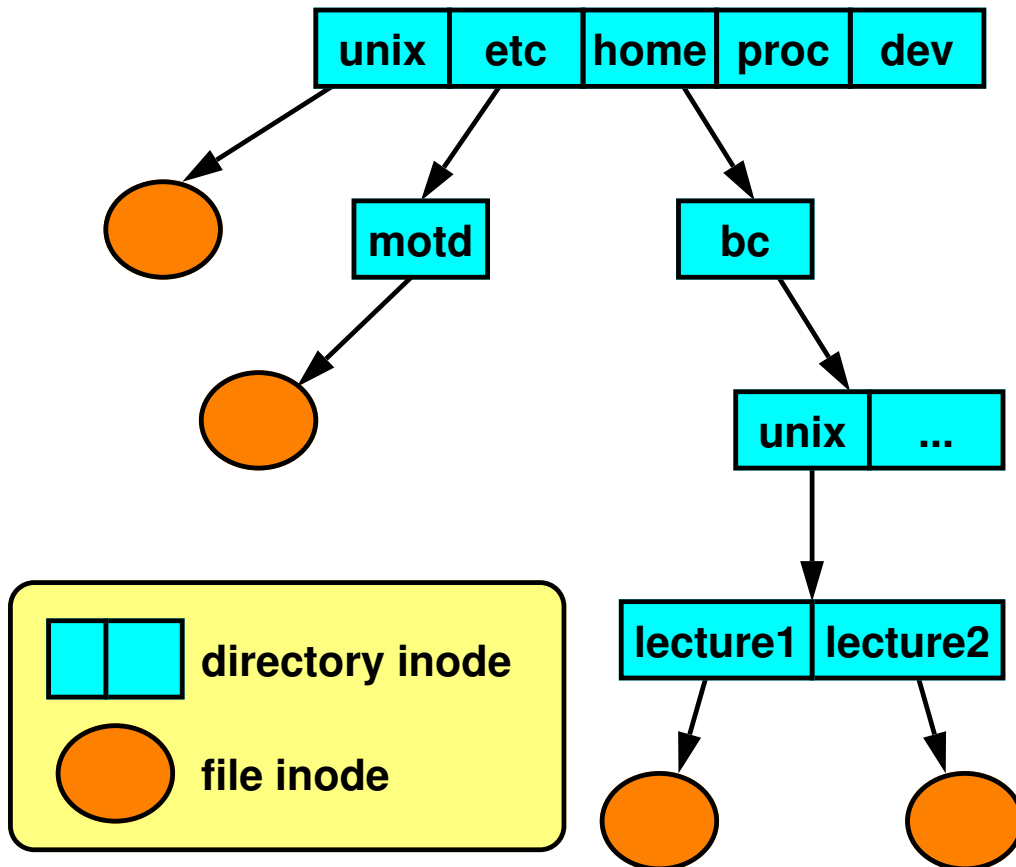
.	1
..	1
unix	117
etc	4
home	18
proc	36
dev	93

.	4
..	1
motd	33
image	117




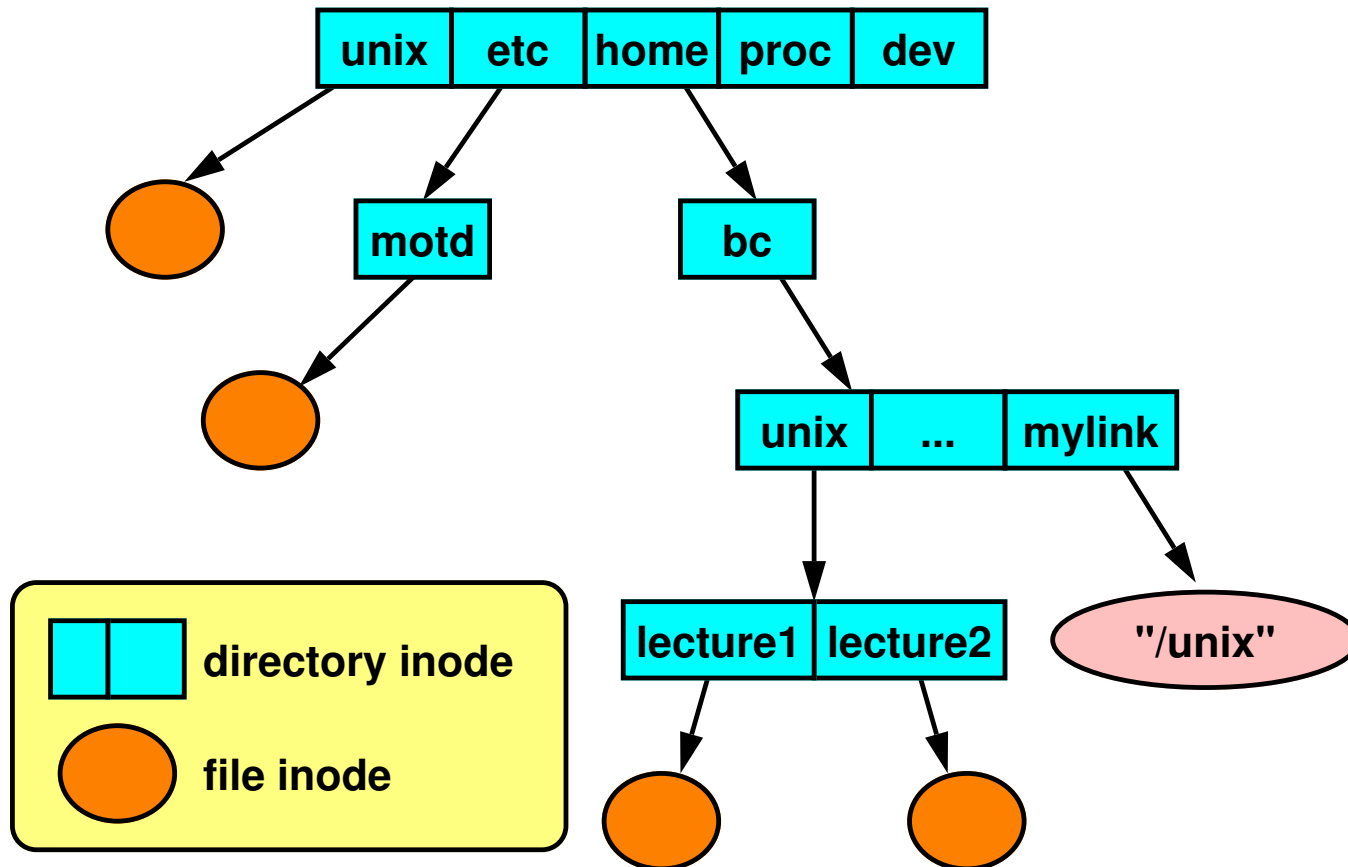
Soft Links

```
% ln -s /unix /home/bc/mylink
```



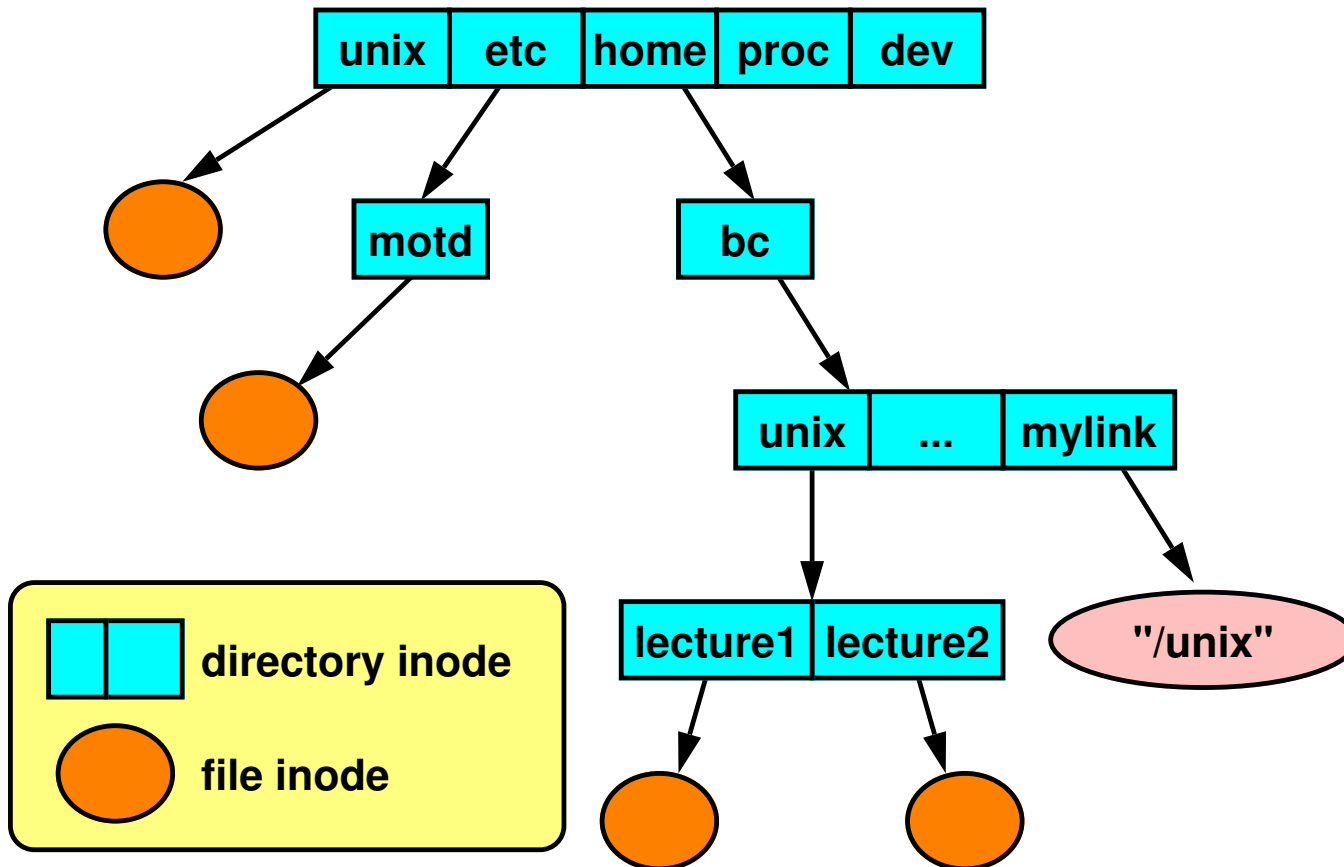
Soft Links

```
% ln -s /unix /home/bc/mylink
```



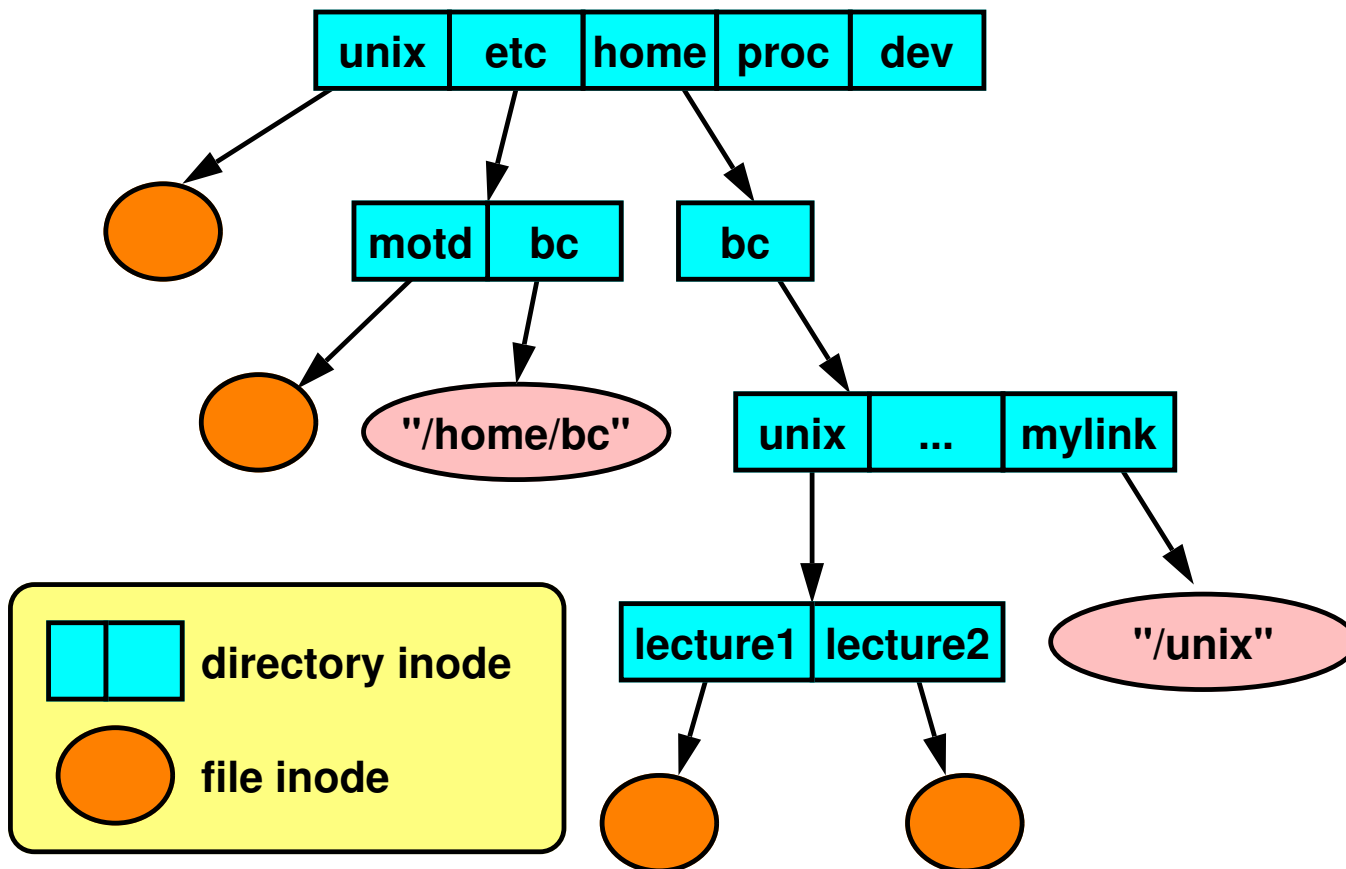
Soft Links

```
% ln -s /unix /home/bc/mylink  
% ln -s /home/bc /etc/bc
```



Soft Links

```
% ln -s /unix /home/bc/mylink  
% ln -s /home/bc /etc/bc
```




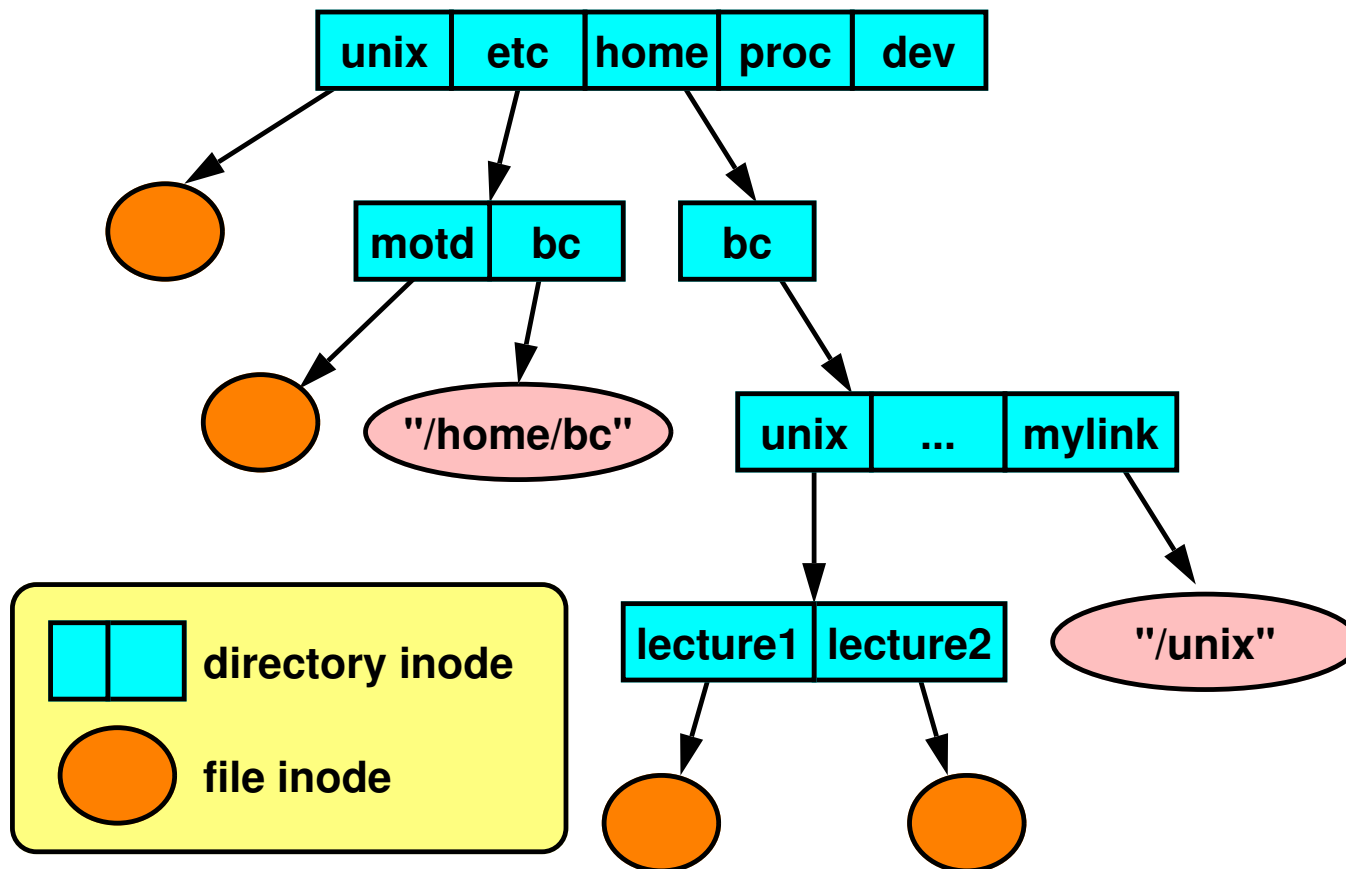
Soft Links

% `ls -l /etc/bc/unix/lecture1`

▢ same as "`ls -l /home/bc/unix/lecture1`", or is it?


○ yes for the "root" account, may be no for the "bc" account

◆ see "access protection"




Working Directory

- ➡ Maintained *in kernel for each process*
- paths not starting from "/" start with the working directory
 - get by using the `getcwd()` system call
 - set by using the `chdir()` system call
 - displayed (via shell) using "pwd"




Access Protection

- ➡ OS needs to make sure that only authorized processes are allowed access to system resources
 - various ways to provide this
- ➡ Unix (and many other systems, such as Windows) associates with files some indication of which *security principals* are allowed access
 - along with what sort of access is allowed
- ➡ A *security principal* is normally a user or group of users
 - a "user" can be an identity used by processes performing system functions
 - each running process can have several security principals associated with it
 - all processes have a user identification and a set of group identifications
 - for Sixth-Edition Unix, only one user ID and one group ID



Access Protection

- ➡ Each file has associated with it a set of access permissions
 - there are 3 classes of security principals:
 - *user*: owner of the file
 - *group*: group owner of the file
 - *others*: everyone else
 - for each of the 3 classes of principals, specify what sorts of operations on the file are allowed
 - the operations are grouped into 3 classes:
 - *read*: can read a file or directory
 - *write*: can write a file or directory
 - *execute*: one must have *execute permission* for a *directory* in order to *follow a path through it*
- ➡ *Rules for checking permissions*
 - 1) determines the *smallest class of principals the requester belongs to* (user being smallest and others being largest)
 - 2) then it checks for appropriate permissions with that class



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm      593 Dec 17 13:34 x
```

```
./B:
```


```
total 2
```

```
-r--rw-rw-  1 bill  adm      446 Dec 17 13:34 x
```

```
-rw----rw-  1 trina  adm      446 Dec 17 13:45 y
```

➡ Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

1) Q: May `andy` list the contents of directory `A`?



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm      593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```


```
-r--rw-rw-  1 bill  adm      446 Dec 17 13:34 x
```

```
-rw----rw-  1 trina  adm      446 Dec 17 13:45 y
```

➡ Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

1) Q: May `andy` list the contents of directory `A`?

A: No



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```


```
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
```

```
-rw----rw-  1 trina  adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

2) Q: May `andy` read `A/x`?



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm      593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```

```
-r--rw-rw-  1 bill  adm      446 Dec 17 13:34 x
```


```
-rw----rw-  1 trina  adm      446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

2) Q: May `andy` read `A/x`?

A: Yes



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm      593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```


```
-r--rw-rw-  1 bill  adm      446 Dec 17 13:34 x
```

```
-rw----rw-  1 trina  adm      446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

3) Q: May `trina` list the contents of directory `B`?



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr---  2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```

```
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
```


```
-rw---rw-  1 trina  adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

3) Q: May `trina` list the contents of directory `B`?

A: Yes



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```


```
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
```

```
-rw----rw-  1 trina adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

4) Q: May `trina` modify `B/y`?



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm      593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```

```
-r--rw-rw-  1 bill  adm      446 Dec 17 13:34 x
```


```
-rw----rw-  1 trina adm      446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

4) Q: May `trina` modify `B/y`?

A: No



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```


```
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
```

```
-rw----rw-  1 trina  adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

5) Q: May `bill` modify `B/x`?



Permissions Example

```
% ls -lR
.:
total 2
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B


./A:
total 1
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x

./B:
total 2
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
-rw----rw-  1 trina adm     446 Dec 17 13:45 y
```

➡ Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

5) Q: May `bill` modify `B/x`?

A: No



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```


```
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
```

```
-rw----rw-  1 trina  adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

6) Q: May `bill` read `B/y`?



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```

```
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
```


```
-rw-----rw- 1 trina adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

6) Q: May `bill` read `B/y`?

A: No




Open

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
int open(const char *path, int options [, mode_t mode])
```



options


- **O_RDONLY** open for reading only
- **O_WRONLY** open for writing only
- **O_RDWR** open for reading and writing
- **O_APPEND** set the file offset to end of file prior to each write
- **O_CREAT** if the file does not exist, then create it, setting its mode to mode adjusted by umask
- **O_EXCL**: if **O_EXCL** and **O_CREAT** are set, then open fails if the file exists
- **O_TRUNC** delete any previous contents of the file
- **O_NONBLOCK** don't wait if I/O cannot be done immediately



Setting File Permissions


```
#include <sys/types.h>
#include <sys/stat.h>
int chmod(const char *path, mode_t mode)
```

- sets the file permissions of the given file to those specified in mode
- only the owner of a file and the superuser may change its permissions
- nine combinable possibilities for mode (read/write/execute for user, group, and others)
- S_IRUSR (0400), S_IWUSR (0200), S_IXUSR (0100)
- S_IRGRP (040), S_IWGRP (020), S_IXGRP (010)
- S_IROTH (04), S_IWOTH (02), S_IXOTH (01)
- note: numeric prefix of 0 means the number is in octal format




Creating a File

- ➡ Use either open or creat
 - ▬ `open(const char *pathname, int flags, mode_t mode)`
 - flags must include `O_CREAT`
 - ▬ `creat(const char *pathname, mode_t mode)`
 - ▬ open is preferred
- ➡ The mode parameter helps specify the permissions of the newly created file
 - ▬ `permissions = mode & ~umask`



Umask


- ➡ Standard programs create files with "maximum needed permissions" as *mode*
 - ▬ compilers: 0777
 - ▬ editors: 0666
- ➡ Per-process parameter, *umask*, used to *turn off* undesired permission bits
 - ▬ e.g., turn off all permissions for others, write permission for group: set umask to 027
 - ▬ compilers: $\text{permissions} = 0777 \ \& \ \sim(027) = 0750$
 - ▬ editors: $\text{permissions} = 0666 \ \& \ \sim(027) = 0640$
 - ▬ set with `umask ()` system call or (usually) `umask` shell command



1.4 Beyond A Simple OS

➡ Extensions

➡ New Functionality




What Else?



Beyond Sixth-Edition Unix (1975)

- multiple threads per process
 - how is the process model affected?
- virtual memory
 - in Sixth-Edition Unix, all currently running process had to fit into the computer's memory at once, along with the OS
 - virtual memory separates the address space from physical resources
- name everything using directory-system path names
 - e.g., /proc
- security
 - Unix solution is pretty elegant
 - new types of requirement such as permissions to add new software, perform backups, etc.



What Else?



New functionalities

- networking
 - much is beyond the scope of this class
- interactive, multimedia user interface
 - make sure interactive user receives excellent response
- software complexity
 - in Sixth-Edition Unix, to add a new device you need to:
 - ◆ write a "device driver" to handle the device
 - ◆ modify OS source code by adding references to the driver to a few tables
 - ◆ recompile the OS and reboot your computer
 - plug-and-play is desirable
 - ◆ need to support dynamic linking of modules into a running system
 - microkernel
 - ◆ how much of the OS functionality can be moved out of the kernel?

