

Ch 2: Multithreaded Programming

Bill Cheng

<http://merlot.usc.edu/cs402-s16>



Copyright © William C. Cheng

Overview

- Why threads?
- How to program with threads?
- what is the API?
- Synchronization
 - mutual exclusion
 - semaphores
 - condition variables
- Pitfall of thread programmings



Copyright © William C. Cheng

Concurrency

- Many things occur simultaneously in the OS
 - e.g., data coming from a disk, data coming from the network, data coming from the keyboard, mouse got clicked, jobs need to get executed
- If you have multiple processors, you may be able to handle things in parallel
 - that's real concurrency
- If you only have one processor, you may want to make it look like things are running in parallel
 - do multiplexing to create the illusion
 - as it turns out, it's a good idea to do this even if you have only have one processor
- The down side is that if you want concurrency, you have to have **concurrency control** or bad things can happen



Copyright © William C. Cheng

Why Threads?




- Many things are easier to do with threads
 - **multithreading** is a **powerful paradigm**
 - makes your design **cleaner**, and therefore, less buggy
- Many things run faster with threads
 - if you are just waiting, don't waste CPU cycles, give the CPU to someone else, **without explicitly** giving up the CPU
- **kernel threads** vs. **user threads**
 - basic concepts are the same
 - can easily do programming assignments for user-level threads
 - for kernel programming assignments, you need to fill out missing parts of various kernel threads




Copyright © William C. Cheng

A Simple Example: rlogind




Copyright © William C. Cheng

A Simple Example: rlogind



Copyright © William C. Cheng

- will be very difficult to implement this without using threads
- if you want to handle a large number of requests simultaneously



- Threads Creation & Termination
- Threads & C++
- Synchronization
- Thread Safety
- Deviations


2.2 Programming With Threads

- don't have to call select()

```

int incoming(int x_in, int x_out) {
    :0;
    int eof = 0;
    char buf[BSIZE];
    int size;
    while(!eof) {
        size = read(x_in, buf, BSIZE);
        if (size == 0)
            eof = 1;
        if (write(x_out, buf, size) <= 0)
            eof = 1;
    }
}

```



Life With Threads

Diagram illustrating a single-threaded database system. Four horizontal arrows labeled "Requests" point towards a vertical pink rectangle labeled "Database". A curved pink arrow loops back from the bottom of the database to the top, indicating a single queue of requests.

Single-Threaded Database Server


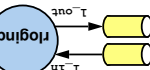
```

LogOutput(int r_in, int r_out, int l_in, int l_out) {
    fd_set in = 0; out;
    int want_l_write = 0, want_r_write = 1;
    int want_l_read = 1, want_r_read = 1;
    int eof = 0, tszize, fszize, wret;
    char tbuf[BSIZE], tbuf[BSIZE];


   fcntl(r_in, F_SETFL, O_NONBLOCK);
    fcntl(r_out, F_SETFL, O_NONBLOCK);
    fcntl(l_in, F_SETFL, O_NONBLOCK);
    fcntl(l_out, F_SETFL, O_NONBLOCK);

    while (1) {
        FD_ZERO(&in);
        FD_ZERO(&out);
        FD_SET(r_out, &out);
        FD_SET(r_read, &in);
        FD_SET(l_out, &out);
        FD_SET(l_read, &in);
        if (want_r_write) FD_SET(r_in, &in);
        if (want_l_write) FD_SET(l_in, &in);
        if (want_r_read) FD_SET(r_out, &out);
        if (want_l_read) FD_SET(l_out, &out);
        select(MAXFD, &in, &out, 0, 0);
        if (FD_ISSET(r_in, &in)) {
            if ((tszize = read(r_in, tbuf, BSIZE)) < 0) {
                if (want_l_read == 1)
                    want_l_read = 0;
            } else {
                if (eof == 1)
                    want_r_write = 1;
            }
        }
    }
}

```



Life Without Threads

[illegible]

Life Without Threads

Copyright © William C. Cheng

17

✕

```

start_servers ( ) {
    DWORD id;
    HANDLE thread;
    for (i=0; i<nr_of_server_threads; i++)
        thread = CreateThread(
            0, // security attributes
            0, // default # of stack pages allocated
            server, // first procedure
            arg, // argument
            0, // creation flags
            0, // thread ID
        );
    // perform service
    return (0);
}

DWORD WINAPI server(void *arg) {
    // perform service
    return (0);
}

```

→ We won't talk about Win32 much

Creating a Win32 Thread

Operating Systems - CSCI 402

Copyright © William C. Cheng

15

```

start_servers ( ) {
    pthread_t thread;
    for (i=0; i<nr_of_server_threads; i++)
        pthread_create(&thread,
            0, // argument;
            server,
            arg,
        );
    // perform service
    return (0);
}

```

→ every thread needs a separate stack

→ first stack frame in every child thread corresponds to server()

→ one arg in each of these stack frames

Creating a POSIX Thread

Operating Systems - CSCI 402

Copyright © William C. Cheng

13

```

#include <pthread.h>

SYNOPSIS
    man pthread_create

Compile and link with -pthread.

→ the start routine is also known as the "first procedure" or "thread function" of the child thread
→ it's like main() for the child thread
→ the "thread ID" of the newly created thread will be returned
→ in the first argument of pthread_create()
→ may not be a Thread Control Block

```

Creating a POSIX Thread

Operating Systems - CSCI 402

Copyright © William C. Cheng

16

```

pthread_create(&in_thread,
    0, // incoming,
    pthread_create(&out_thread,
        0, // outgoing,
        l_in, r_out); // Cannot do this ...
    );
pthread_create(&in_thread,
    0, // outgoing,
    l_in, r_out); // Cannot do this ...
    );
}

```

→ How do we wait till they are done? *

Complications

Operating Systems - CSCI 402

Copyright © William C. Cheng

16

→ keep thread handle in the stack

→ keep thread handle in the heap

→ need to make sure that eventually you will call the following to not leak memory

→ keep thread handle in the stack

→ keep thread handle in the heap

→ need to make sure that eventually you will call the following to not leak memory

Creating a POSIX Thread

Operating Systems - CSCI 402

Copyright © William C. Cheng

14

```

start_servers ( ) {
    pthread_t thread;
    for (i=0; i<nr_of_server_threads; i++)
        pthread_create(&thread,
            0, // thread ID
            // default attributes
            server, // first procedure
            argument); // argument
    // perform service
    return (0);
}

```

→ POSIX 1003.1c standard

→ pthread_create() returns 0 if successful

→ child thread starts executing here

→ arg = argument (from caller)

→ child thread ends when return from its start routine / first procedure

Creating a POSIX Thread

Operating Systems - CSCI 402

Copyright © William C. Cheng

Thread Termination

Thread return values

- which threads receive these values
- how do they do it?
- clearly, receiving thread must wait until the producer thread produced it, i.e., producer thread has terminated
- so we must have a way for one thread to wait for another thread to terminate
- must have a way to say which thread you are waiting for
- need a unique identifier
- tricky if it can be reused

To wait for another thread to terminate

```
int pthread_join(pthread_t thread,
                 void **ret_value);
```

Operating Systems - CSCI 402

Copyright © William C. Cheng

Multiple Arguments

Need to be careful how to pass argument to new thread when you call `pthread_create()`

- there is no way to pass multiple arguments in either POSIX or Win32
- passing address of a *local* variable (like the previous example) only works if we are certain the this storage doesn't go out of scope until the thread is done with it
- passing address of a *static* or a *global* variable only works if we are certain that only one thread at a time is using the storage
- passing address of a *dynamically* allocated storage only works if we can free the storage when, and only when, the thread is finished with it
- this would not be a problem if the language supports garbage collection

Ask yourself, "How can you be sure?"

- if the answer is, "I hope it works", then you need a different solution

Operating Systems - CSCI 402

Copyright © William C. Cheng

Multiple Arguments

```
typedef struct {
    int first, second;
} two_ints_t;

pthread_t in_thread, out_thread;
two_ints_t in={1_in, 1_out}, out={1_in, 1_out};

pthread_create(&in_thread,
               0,
               incoming,
               &in);

/* How do we wait till they are done? */
pthread_join(in_thread,
             void **p=&first);
return NULL;
```

Operating Systems - CSCI 402

Copyright © William C. Cheng

Thread Termination

How does a thread *self-terminate*?

- return from its "first procedure"
- return a value of type `(void*)`
- call `pthread_exit(ret_value)`
- `ret_value` is of type `(void*)`

```
void *child(void *arg) {
    ...
    if (terminate_now) {
        pthread_exit((void*)1);
    }
    return (void*)2;
}
```

Thread Control Block

Exit/Return Code
TID

Operating Systems - CSCI 402

Copyright © William C. Cheng

When Is The Child Thread Done?

```
pthread_t in_thread, out_thread;
two_ints_t in={1_in, 1_out}, out={1_in, 1_out};

pthread_create(&in_thread, 0, incoming, &in);
pthread_create(&out_thread, 0, outgoing, &out);

pthread_join(in_thread, 0);
pthread_join(out_thread, 0);
```

Operating Systems - CSCI 402

Copyright © William C. Cheng

Multiple Arguments

```
typedef struct {
    int first, second;
} two_ints_t;

pthread_t in_thread, out_thread;
two_ints_t in={1_in, 1_out}, out={1_in, 1_out};

pthread_create(&in_thread,
               0,
               incoming,
               &in);

/* How do we wait till they are done? */
pthread_join(in_thread,
             void **p=&first);
return NULL;
```

Operating Systems - CSCI 402



— a function is just an address (of something in the text/code segment)

```
pthread_create(&t1d,
               0,
               (void *) (* ) func,
               (void *) (* ) 1);
...
int func = 4; // func definition 1
...
void func(int i) { // func definition 2
    ...
    return(0);
}
```

Types



```
pthread_t thread;
pthread_attr_t thr_attr;
...
pthread_attr_init(&thr_attr);
/* establish some attributes */
...
pthread_create(&thread, &thr_attr, start_routine, arg);
pthread_attr_destroy(&thr_attr);
...
— thread attribute only needs to be valid when a thread is created
— therefore, it can be destroyed as soon as the thread is
```

Thread Attributes



- Any thread can join with any other thread
- there's *no parent/child relationships* among threads
- unlike process termination and wait ()
- What happens if a thread terminates and no other thread wants to join with this thread?
- it also goes into a *zombie* state
- all the thread related information is freed up, except for the thread ID and return code
- What if two threads want to join with the same thread?
- after the first thread joins, the thread ID and return code are freed up and the thread ID may get reused
- so don't do this!

Thread Termination




```
start_servers() {
    pthread_t thread;
    int i;
    for (i=0; i<nr_of_servers; i++) {
        pthread_create(&thread, 0, server, 0);
        pthread_detach(thread);
    }
    ...
    server() {
        ...
    }
}
```

Detached Threads



- Difference between pthread_exit () and exit ()
- pthread_exit () terminates only the calling thread
- exit () terminates the process, including all threads running in it
- it will not wait for any thread to terminate
- what will this code do?
- int main(int argc, char *argv[]) {
 // create all the threads
 return(0);
}
- when main () returns, exit () will be called
- as a result, none of the created child threads may get a chance to run



Thread Termination



- Difference between pthread_exit () and exit ()
- pthread_exit () terminates only the calling thread
- exit () terminates the process, including all threads running in it
- it will not wait for any thread to terminate
- what about this code?
- int main(int argc, char *argv[]) {
 // create all the threads
 pthread_exit(0); // exit the main thread
 return(0);
}
- here, pthread_exit () will terminate the main thread, so exit () is never called
- as it turns out, this special case is taken care of in the pthread library implementation
- You should use pthread_join () unless you are absolutely sure

Thread Termination



35



Mutual Exclusion

Copyright © William C. Cheng

Operating Systems - CSCI 402


36

Thread 1:
 $x = x + 1;$

Thread 2:
 $x = x + 1;$

looks like it doesn't matter how you execute, x will be incremented by 2 in the end


choices are

- thread 1 executes x = x+1 then thread 2 executes x = x+1
- thread 2 executes x = x+1 then thread 1 executes x = x+1
- are there other choices?

Threads and Mutual Exclusion

Copyright © William C. Cheng

Operating Systems - CSCI 402



33

Compiling It

```
% gcc -o mat mat.c -lpthread
```

Copyright © William C. Cheng

Operating Systems - CSCI 402


34


2.2.3 Synchronization

In real life, "synchronization" means that you want to do things at the same time

In computer science, "synchronization" could mean the above, **OR**, it means that you want to **prevent** do things at the same time

Copyright © William C. Cheng

Operating Systems - CSCI 402


31

Stack Size

pthread_t thread;
pthread_attr_t thr_attr;
pthread_attr_init(&thr_attr);
pthread_attr_setsize(&thr_attr, 20*1024*1024);
pthread_create(&thr_thread, &thr_attr, start routine, arg);
pthread_attr_destroy(&thr_attr);

the above code set the stack size to 20MB


the default stack size is very large

if you need to create a lot of threads, you need to control the stack size

default stack size is probably around 1MB in Solaris and 8MB in some Linux implementations

Copyright © William C. Cheng

Operating Systems - CSCI 402


32

Example

```

#include <stdio.h>
#include <pthread.h>
#include <string.h>

#define M 3
#define N 4
#define P 5

int A[M][N];
int B[N][P];
int C[M][P];

void *matrix_mult(void *arg) {
    int row = (int) arg, col;
    int i, t;
    for (col=0; col < P; col++) {
        t = 0;
        for (i=0; i < N; i++)
            t += A[row][i] * B[i][col];
        C[row][col] = t;
    }
    return(0);
}

/* initialize the matrices ... */
// create the worker threads
for (i=0; i < M; i++) {
    if (error = pthread_create(
        &thr[i],
        0,
        matrix_mult,
        (void *) i))
        fprintf(stderr,
            "pthread_create: %s",
            strerror(error));
    exit(1);
}
// wait for workers to finish
for (i=0; i < M; i++)
    pthread_join(thr[i], 0);
/* print the results ... */

```

Copyright © William C. Cheng

Operating Systems - CSCI 402

Copyright © William C. Cheng

Taking Multiple Locks

Graph representation ("wait-for" graph)

```

proc1 ( ) {
    pthread_mutex_lock(&m1);
    /* use object 1 */
    pthread_mutex_lock(&m2);
    /* use objects 1 and 2 */
    pthread_mutex_unlock(&m1);
    pthread_mutex_unlock(&m2);
}

proc2 ( ) {
    pthread_mutex_lock(&m2);
    /* use object 2 */
    pthread_mutex_lock(&m1);
    /* use objects 1 and 2 */
    pthread_mutex_unlock(&m1);
    pthread_mutex_unlock(&m2);
}

```

Operating Systems - CSCI 402

Copyright © William C. Cheng

Set Up

if a mutex cannot be initialized statically, do:

```

pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;

int pthread_mutex_init(
    pthread_mutex_t *mutexp,
    pthread_mutexattr_t *attrp)

int pthread_mutex_destroy(
    pthread_mutex_t *mutexp)

```

Usually, mutex attributes are not used

Operating Systems - CSCI 402

Copyright © William C. Cheng

Threads and Mutual Exclusion

Unfortunately, machines do not execute high-level language statements

- they execute machine instructions
- now if thread 1 executes the first (or two) machine instructions
- context switch!
- how can this happen?
- then thread 2 executes all 3 machine instructions
- then later thread 1 executes the remaining machine instructions
- x would have only increased by 1

Thread 1: `x = x+1;`

Thread 2: `x = x+1;`

Operating Systems - CSCI 402

Copyright © William C. Cheng

Necessary Conditions For Deadlocks

All 4 conditions below must be met in order for a deadlock to be possible (no guarantee that a deadlock may occur)

- 1) Bounded resources
 - only a finite number of threads can have concurrent access to a resource
- 2) Wait for resources
 - threads wait for resources to be freed up, without releasing resources that they hold
- 3) No preemption
 - resources cannot be *revoked* from a thread
- 4) Circular wait
 - there exists a set of waiting threads, such that each thread is waiting for a resource held by another

Operating Systems - CSCI 402

Copyright © William C. Cheng

Taking Multiple Locks

Mutex is not a cure-all

- when you have more than one locks, you may get into trouble

```

proc1 ( ) {
    pthread_mutex_lock(&m1);
    /* use object 1 */
    pthread_mutex_lock(&m2);
    /* use objects 1 and 2 */
    pthread_mutex_unlock(&m1);
    pthread_mutex_unlock(&m2);
}

proc2 ( ) {
    pthread_mutex_lock(&m2);
    /* use object 2 */
    pthread_mutex_lock(&m1);
    /* use objects 1 and 2 */
    pthread_mutex_unlock(&m1);
    pthread_mutex_unlock(&m2);
}

```

Operating Systems - CSCI 402

Copyright © William C. Cheng

Threads and Synchronization

code between `pthread_mutex_lock()` and `pthread_mutex_unlock()` is called a **critical section**

- all the critical sections *with respect to a particular mutex* are **"mutually exclusive"**
- the system (not necessarily the OS) guarantees that only *one* critical section can be executing at any point in time
- how it's really done will be covered in Ch 5

Locking a mutex is like getting the key to a safe-deposit box

```

// shared by both threads
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;

int x;

```

Operating Systems - CSCI 402

Copyright © William C. Cheng 47

Deadlock Prevention: Conditional Locking

```

proc1 ( ) {
    pthread_mutex_lock(&m1);
    /* use object 1 */
    pthread_mutex_lock(&m2);
    /* use objects 1 and 2 */
    pthread_mutex_unlock(&m2);
    pthread_mutex_unlock(&m1);
}

proc2 ( ) {
    while (1) {
        pthread_mutex_lock(&m2);
        pthread_mutex_unlock(&m1);
        break;
    }
    pthread_mutex_unlock(&m1);
}

```

Operating Systems - CSCI 402

Copyright © William C. Cheng 45

Deadlock Prevention: Lock Hierarchies

= organize mutexes into levels
 = must not try locking a mutex at level *i* if already holding a mutex at level *j* where *j* < *i*
 = e.g., if holding mutexes at levels 2 and 3, can only wait for a mutex at levels 4 or higher

Operating Systems - CSCI 402

Copyright © William C. Cheng 43

Dealing with Deadlock

= Deadlock is a programming bug
 = one of the oldest bugs
 = it's a tricky one because it only deadlocks *sometimes*

Hard

- = is the system deadlocked?
- = will this move lead to deadlock?
- = this is *detection*
- o if you can detect deadlocks, what do you do after you have detected them?

Easy

- = restrict use of mutexes so that deadlock cannot happen
- = this is *prevention*

= Deadlock is a complicated subject
 = some textbooks spend an entire chapter on deadlocks
 = we will only look at a couple of cases

Operating Systems - CSCI 402

Copyright © William C. Cheng 46

One Mutex, Multiple Critical Sections

= I use "Synchronization Box" to mean executing one of many critical section code with respect to one particular mutex

```

f1 ( ) {
    pthread_mutex_lock(&m);
    /* critical section */
    x++;
    pthread_mutex_unlock(&m);
}

f2 ( ) {
    pthread_mutex_lock(&m);
    /* critical section */
    x--;
    pthread_mutex_unlock(&m);
}

```

Operating Systems - CSCI 402

Copyright © William C. Cheng 46

Deadlock Prevention: Lock Hierarchies

= What if you cannot organize your mutexes in such strict order for deadlock detection?

Operating Systems - CSCI 402

Copyright © William C. Cheng 44

Deadlock Prevention: Lock Hierarchies

= organize mutexes into levels
 = must not try locking a mutex at level *i* if already holding a mutex at level *j* where *j* < *i*
 = e.g., if hold mutexes at levels 2 and 3, can only wait for a mutex at levels 4 or higher

Operating Systems - CSCI 402

Copyright © William C. Cheng

53

Guarded Commands

For exams, you need to know how to write simple pseudo-code in the language of *Guarded Commands*

- you cannot evaluate the guard if your thread is not running
- altogether is an *atomic operation* if the *guard is true*
- evaluating the *guard* and *executing the command sequence* *atomically* mean that it's executed without interruption
- a *guard* is a *boolean expression* (evaluates to true or false)
- (*atomically*) at any time the *guard* is evaluated to be *true*
- this means that the command sequence *can be* executed

```

when (guard) [
    ...
    /*
    execute this code atomically
    once the guard is true,
    */
]
command sequence
  
```

Operating Systems - CSCI 402

Copyright © William C. Cheng

51

Producer-Consumer Problem

When does it require synchronization?

- an overkill (i.e., too *inefficient*)
- if you use a *single mutex* to lock the entire array of buffers, it's
- Most of the time, no interference
- A circular buffer is used

Operating Systems - CSCI 402

Copyright © William C. Cheng

49

One Mutex, Multiple Critical Sections

By calling `pthread_mutex_lock (&m)`, a thread can be placed into a *queue* and *wait* there *indefinitely* for mutex `m` to become available

- multiple threads would join this queue
- queue is served one at a time, like a supermarket checkout
- when it's your thread's turn, `pthread_mutex_lock ()` returns with the mutex locked, your thread can execute critical section code, and then release the mutex

```

f1 () {
    pthread_mutex_lock (&m);
    /* critical section */
    pthread_mutex_unlock (&m);
}

f2 () {
    pthread_mutex_lock (&m);
    /* critical section */
    pthread_mutex_unlock (&m);
}
  
```

I use "Synchronization Box" to mean executing one of many critical section code with respect to one particular mutex

Operating Systems - CSCI 402

Copyright © William C. Cheng

54

Guarded Commands

when (guard) [

```

    ...
    /*
    execute this code atomically
    once the guard is true,
    */
]
command sequence
  
```

Operating Systems - CSCI 402

Copyright © William C. Cheng

52

Producer-Consumer Problem

When does it require synchronization?

- producer needs to be blocked when all slots are full
- consumer needs to be blocked when all slots are empty
- an overkill (i.e., too *inefficient*)
- if you use a *single mutex* to lock the entire array of buffers, it's
- Most of the time, no interference
- A circular buffer is used

Operating Systems - CSCI 402

Copyright © William C. Cheng

50

Beyond Mutexes

Mutex is necessary when shared data is being modified

- although there are cases where using a mutex is an overkill (i.e., too restrictive and inefficient and would lock threads out when it's not necessary)
- we would like to have better concurrency (i.e., "fine-grained parallelism") when complete mutual exclusion is not required
- two major categories to illustrate this
- what if threads don't interfere one another most of the time and synchronization is only required occasionally?
- e.g., *Producer-Consumer problem* (a.k.a., bounded-buffer problem)
- what if some threads just want to *look at* (i.e., *read*) a piece of data?
- e.g., *Readers-Writers problem*

Barrier Synchronization

Operating Systems - CSCI 402

Copyright © William C. Cheng

59

```

when (guard) [
  /*
  once the guard is true,
  execute this code atomically
  */
  ...
]
command sequence

```

Guarded Commands

Copyright © William C. Cheng

60

```

when (guard) [
  /*
  once the guard is true,
  execute this code atomically
  */
  ...
]
command sequence

```

Guarded Commands

Copyright © William C. Cheng

57

```

when (guard) [
  /*
  once the guard is true,
  execute this code atomically
  */
  ...
]
command sequence

```

Guarded Commands

Copyright © William C. Cheng

58

```

when (guard) [
  /*
  once the guard is true,
  execute this code atomically
  */
  ...
]
command sequence

```

Guarded Commands

Copyright © William C. Cheng

55

```

when (guard) [
  /*
  once the guard is true,
  execute this code atomically
  */
  ...
]
command sequence

```

Guarded Commands

Copyright © William C. Cheng

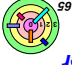
56

```

when (guard) [
  /*
  once the guard is true,
  execute this code atomically
  */
  ...
]
command sequence

```

Guarded Commands


65

Copyright © William C. Cheng

◯ this is often why you would use a semaphore
 ◯ **thread** performs a V operation on the same semaphore
 ◯ one thread performs a P operation on a semaphore, **another**
 ◯ therefore, it must be that thread that unlocks that mutex
 ◯ if a thread locks a mutex, it's holding the lock
 Main difference between a semaphore and a mutex

◯ can be used to solve the producer-consumer problem
 ◯ this is known as a **counting semaphore**

```

semaphore S = N;


void WaitTime ( ) {
    P (S);
    ...
    /* no more than N threads
    here at once */
    V (S);
}

P (S) operation:
    when (S < 0) [
        S = S - 1;
    ]
V (S) operation:
    [
        S = S + 1;
    ]

```

Mutexes with Semaphores

Operating Systems - CSCI 402


66

Copyright © William C. Cheng

```


void Produce (char item) {
    char item;
    P (occupied);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V (empty);
    return (item);
}

char Consume ( ) {
    char item;
    P (occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V (empty);
}

```

Producer/Consumer with Semaphores

Operating Systems - CSCI 402


63

Copyright © William C. Cheng

◯ this is known as a **binary semaphore**

```

semaphore S = 1;


void OneAtATime ( ) {
    P (S);
    ...
    /* code executed mutually
    exclusively */
    V (S);
}

P (S) operation:
    when (S < 0) [
        S = S - 1;
    ]
V (S) operation:
    [
        S = S + 1;
    ]

```

Mutexes with Semaphores

Operating Systems - CSCI 402


64

Copyright © William C. Cheng

◯ if you use it this way, nothing
 ◯ this looks just like mutex, what have we really gained?
 So, you can lock a data structure using a binary semaphore

```


do:
    pthread_mutex_lock (&m);
    x = x + 1;
    pthread_mutex_unlock (&m);

```

Instead of doing

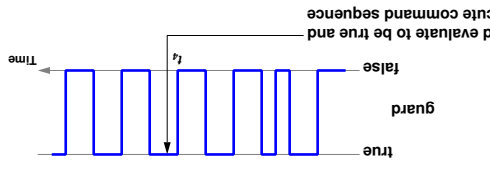
Implement A Mutex With A Binary Semaphore

Operating Systems - CSCI 402


61

Copyright © William C. Cheng

◯ **atomic**: as if it's executed in an instance of time (duration = 0)
 ◯ this is okay because it's just pseudo-code



guard evaluate to be true and execute command sequence


once the guard is true, execute this code atomically

when (guard) [

command sequence

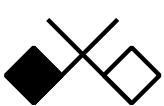
Guarded Commands

Operating Systems - CSCI 402


62

Copyright © William C. Cheng

◯ **A semaphore, S, is a nonnegative integer** on which there are exactly two operations defined by two guarded commands
 ◯ P (S) operation (implemented as a guarded command):
 ◯ when (S > 0) [
 S = S - 1;
]
 ◯ V (S) operation (implemented as a guarded command):
 ◯ [S = S + 1;
 there are no other means for manipulating the value of S
 ◯ other than initializing it



Semaphores

Operating Systems - CSCI 402

Copyright © William C. Cheng

71

```

void Produce(char item) {
    int nextin = 0;
    Semaphore occupied = 0;
    Semaphore empty = B;

    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextout = nextout + 1;
    V(occupied);
    return(item);
}

void Consume(char item) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextin = nextin + 1;
    V(empty);
}

```

Producer/Consumer with Semaphores

Operating Systems - CSCI 402

Copyright © William C. Cheng

69

```

void Produce(char item) {
    int nextin = 0;
    Semaphore occupied = 0;
    Semaphore empty = B;

    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextout = nextout + 1;
    V(occupied);
    return(item);
}

void Consume(char item) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextin = nextin + 1;
    V(empty);
}

```

Producer/Consumer with Semaphores

Operating Systems - CSCI 402

Copyright © William C. Cheng

67

```

void Produce(char item) {
    int nextin = 0;
    Semaphore occupied = 0;
    Semaphore empty = B;

    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextout = nextout + 1;
    V(occupied);
    return(item);
}

void Consume(char item) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextin = nextin + 1;
    V(empty);
}

```

Producer/Consumer with Semaphores

Operating Systems - CSCI 402

Copyright © William C. Cheng

72

```

void Produce(char item) {
    char item;
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextout = nextout + 1;
    V(occupied);
    return(item);
}

void Consume(char item) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextin = nextin + 1;
    V(empty);
}

```

Producer/Consumer with Semaphores

Operating Systems - CSCI 402

Copyright © William C. Cheng

70

```

void Produce(char item) {
    char item;
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextout = nextout + 1;
    V(occupied);
    return(item);
}

void Consume(char item) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextin = nextin + 1;
    V(empty);
}

```

Producer/Consumer with Semaphores

Operating Systems - CSCI 402

Copyright © William C. Cheng

68

```

void Produce(char item) {
    char item;
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextout = nextout + 1;
    V(occupied);
    return(item);
}

void Consume(char item) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextin = nextin + 1;
    V(empty);
}

```

Producer/Consumer with Semaphores

Operating Systems - CSCI 402

Copyright © William C. Cheng

77

nextout 1
nextin 1
occupied 0
empty 7

Consumer → Producer

note: producer continue to produce

```

void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    if (nextin == B)
        nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(occupied);
    return(item);
}

char Consume() {
    P(occupied);
    char item;
    item = buf[nextout];
    if (nextout == B)
        nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(empty);
    return item;
}
    
```

Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;

Producer/Consumer with Semaphores

Operating Systems - CSCI 402

Copyright © William C. Cheng

75

nextout 0
nextin 1
occupied 0
empty 7

Consumer → Producer

note: producer continue to produce

```

void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    if (nextin == B)
        nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(occupied);
    return(item);
}

char Consume() {
    P(occupied);
    char item;
    item = buf[nextout];
    if (nextout == B)
        nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(empty);
    return item;
}
    
```

Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;

Producer/Consumer with Semaphores

Operating Systems - CSCI 402

Copyright © William C. Cheng

73

nextout 0
nextin 1
occupied 1
empty 7

Consumer → Producer

note: producer continue to produce

```

void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    if (nextin == B)
        nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(occupied);
    return(item);
}

char Consume() {
    P(occupied);
    char item;
    item = buf[nextout];
    if (nextout == B)
        nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(empty);
    return item;
}
    
```

Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;

Producer/Consumer with Semaphores

Operating Systems - CSCI 402

Copyright © William C. Cheng

76

nextout 1
nextin 1
occupied 0
empty 8

Consumer → Producer

note: producer continue to produce

```

void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    if (nextin == B)
        nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(occupied);
    return(item);
}

char Consume() {
    P(occupied);
    char item;
    item = buf[nextout];
    if (nextout == B)
        nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(empty);
    return item;
}
    
```

Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;

Producer/Consumer with Semaphores

Operating Systems - CSCI 402

Copyright © William C. Cheng

76

nextout 1
nextin 1
occupied 0
empty 7

Consumer → Producer

note: producer continue to produce

```

void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    if (nextin == B)
        nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(occupied);
    return(item);
}

char Consume() {
    P(occupied);
    char item;
    item = buf[nextout];
    if (nextout == B)
        nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(empty);
    return item;
}
    
```

Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;

Producer/Consumer with Semaphores

Operating Systems - CSCI 402

Copyright © William C. Cheng

74

nextout 0
nextin 1
occupied 0
empty 7

Consumer → Producer

note: producer continue to produce

```

void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    if (nextin == B)
        nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(occupied);
    return(item);
}

char Consume() {
    P(occupied);
    char item;
    item = buf[nextout];
    if (nextout == B)
        nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(empty);
    return item;
}
    
```

Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;

Producer/Consumer with Semaphores

Operating Systems - CSCI 402

Copyright © William C. Cheng

Producer-Consumer with POSIX Semaphores

```

void produce(char item) {
    sem_wait(&empty);
    buf[nextin++] = item;
    if (nextin == B)
        nextin = 0;
    sem_post(&occupied);
}

char consume() {
    char item;
    P(occupied);
    item = buf[nextout];
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}

```

```

void produce(char item) {
    char consume() {
        char item;
        P(occupied);
        item = buf[nextout++];
        sem_wait(&occupied);
        if (nextout >= B)
            nextout = 0;
        sem_post(&empty);
        return(item);
    }
}

```

83

Copyright © William C. Cheng

Producer/Consumer with Semaphores

```

Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;

void Produce(char item) {
    char consume() {
        char item;
        P(occupied);
        item = buf[nextout];
        if (nextout == B)
            nextout = 0;
        V(empty);
        return(item);
    }
}

```

81

Copyright © William C. Cheng

Producer/Consumer with Semaphores

empty	8
occupied	0
nextin	1
nextout	1

```

void Produce(char item) {
    char consume() {
        char item;
        P(occupied);
        item = buf[nextout];
        if (nextout == B)
            nextout = 0;
        V(empty);
        return(item);
    }
}

```

note: producer continue to produce

87

Copyright © William C. Cheng

Implementation Of Guarded Commands

In general, the **guard** can be **complicated** and involving the evaluation of several variables (e.g., $a < 3 \ \& \ (b) \leq c$)

- the guard (which evaluates to either true or false) **keeps** **changing its value, continuously** and by multiple threads **simultaneously**
- how can we "capture" the instance of time when it evaluates to true so we can execute the command sequence atomically?
- we have to "sample" it, i.e., take snap shot of all the variables that are involved, but how?
- a mutex is involved, but how?
- need to be efficient
- need something else (**known as condition variables**)
- need a bunch of **rules** to follow

84

Copyright © William C. Cheng

POSIX Semaphores

```

#include <semaphore.h>
sem_t semaphore;
int err;
int pshared = 0; // not shared among processes
int init_value = B; // initial value

err = sem_init(&semaphore, pshared, init_value);
err = sem_destroy(&semaphore);
err = sem_wait(&semaphore); // conditional P operation
err = sem_trywait(&semaphore); // V operation
err = sem_post(&semaphore);

```

82

Copyright © William C. Cheng

Producer/Consumer with Semaphores

```

void Produce(char item) {
    char consume() {
        char item;
        P(occupied);
        item = buf[nextout];
        if (nextout == B)
            nextout = 0;
        V(empty);
        return(item);
    }
}

```

if producer is fast and consumer slow, no one waits

if producer is fast and producer slow, consumer may wait

if consumer is fast and producer slow, consumer may wait

80

Copyright © William C. Cheng

Implementation Of Guarded Commands

POSIX provides *condition variables* for programmers to implement guarded commands

— a *condition variable* is a *queue of threads* waiting for some sort of notification (an "event" or "condition")

- threads, waiting for a guard to become true, join such a queue
- they wait for a specific *condition* to be *signaled*
- they wait for the *right time* to *evaluate the guard*
- (cont...)

Copyright © William C. Cheng

Implementation Of Guarded Commands

POSIX provides *condition variables* for programmers to implement guarded commands

— conceptually, an "event" happens in an instance of a condition (signaling/broadcasting of a condition) if this "event" is zero

- if you are not waiting for it, you'll miss it
- how do you make sure you won't miss an event?
- you have to *follow the protocol* (for multiple interacting threads to follow) described here

Copyright © William C. Cheng

Implementation Of Guarded Commands

POSIX provides *condition variables* for programmers to implement guarded commands

— *atomically* unlocks mutex and wait for the "event"

- with respect to the *operation of the mutex*

Copyright © William C. Cheng

Implementation Of Guarded Commands

POSIX provides *condition variables* for programmers to implement guarded commands

— a *condition variable* is a *queue of threads* waiting for some sort of notification (an "event" or "condition")

- threads that do something to *potentially* change the truth value of the guard can then wake up the threads that were waiting in the queue
- they can *signal* or *broadcast* the *condition*
- no guarantee* that the guard will be true when it's time for *another thread* to evaluate the guard

Copyright © William C. Cheng

Implementation Of Guarded Commands

POSIX provides *condition variables* for programmers to implement guarded commands

1) pthread_cond_wait (*cv, pthread_mutex_t *mutex)

- should only call pthread_cond_wait () if you have the mutex *locked*
- atomically* unlocks mutex and wait for the "event"
- when the event is signaled/broadcasted, pthread_cond_wait () returns with the mutex *locked*

Copyright © William C. Cheng

Implementation Of Guarded Commands

POSIX provides *condition variables* for programmers to implement guarded commands

— *atomically* unlocks mutex and wait for the "event"

- with respect to the *operation of the mutex*

Copyright © William C. Cheng 95

→ if you don't follow these rules, your code will have **race conditions** (i.e., timing-dependent behavior)

Guarded command	POSIX implementation
<pre> when (guard) [statement 1; ... statement n;] </pre>	<pre> /* code * modifying * the guard */ pthread_mutex_lock(&mutex); ... pthread_cond_broadcast(&cv); pthread_mutex_unlock(&mutex); </pre>

Operating Systems - CSCI 402

Copyright © William C. Cheng 96

→ don't believe that `pthread_cond_signal/broadcast()` can be called without locking the mutex

Guarded command	POSIX implementation
<pre> when (guard) [statement 1; ... statement n;] </pre>	<pre> /* code * modifying * the guard */ pthread_mutex_lock(&mutex); ... pthread_cond_broadcast(&cv); pthread_mutex_unlock(&mutex); </pre>

Operating Systems - CSCI 402

Copyright © William C. Cheng 93

Implementation Of Guarded Commands

critical sections

→ with respect to a mutex, a thread can be

- waiting in the mutex queue
- got the lock and inside the "synchronization box"
- only one thread can be inside the "synchronization box"
- waiting in the CV queue
- or outside
- with respect to a mutex, a, b, c are variables that can affect the value of the guard
- can only access (i.e., read/write) them if a thread is inside the "synchronization box" (i.e., has the mutex locked)

→ Synchronization: mutex, condition variables, guards,

Diagram: A box labeled "Synchronization Box" contains a CV (condition variable) and a mutex. A thread can be in the box, waiting in the CV queue, or outside. A thread can only access the CV or mutex if it has the mutex locked.

Copyright © William C. Cheng 94

Implementation Of Guarded Commands

critical sections

→ when you **signal CV**

- one thread in the CV queue gets moved to the mutex queue
- when you **broadcast CV**
- all threads in the CV queue get moved to the mutex queue
- you can only get **added** to the CV queue if you have the mutex locked
- you can only **modify** the variables in the guard if you have the mutex locked
- you can only **read** the variables in the guard (i.e., evaluate the guard) if you have the mutex locked
- you can only **execute** critical section code if you have the mutex locked

→ Synchronization: mutex, condition variables, guards,

Diagram: A box labeled "Synchronization Box" contains a CV (condition variable) and a mutex. A thread can be in the box, waiting in the CV queue, or outside. A thread can only access the CV or mutex if it has the mutex locked.

Copyright © William C. Cheng 91

Implementation Of Guarded Commands

→ POSIX provides **condition variables** for programmers to implement guarded commands

- atomically unlocks mutex and wait for the "event"
- with respect to the **operation of the mutex**

Diagram: A timeline showing a thread waiting for an event (may wait forever) and then unlocking the mutex. The thread is shown waiting for the event, then the event occurs, and the thread unlocks the mutex.

Copyright © William C. Cheng 92

Implementation Of Guarded Commands

→ POSIX provides **condition variables** for programmers to implement guarded commands

- should only call `pthread_cond_signal()` or `pthread_cond_broadcast()` if you have the corresponding **mutex locked**

2) `pthread_cond_broadcast(pthread_cond_t *cv)`

Diagram: A box labeled "Synchronization Box" contains a CV (condition variable) and a mutex. A thread can be in the box, waiting in the CV queue, or outside. A thread can only access the CV or mutex if it has the mutex locked.



- since readers is part of the guard in a guarded command, in the implementation of `[readers--]`, you must signal/broadcast the corresponding condition used to implement that guard
- in this case, only have to signal if readers becomes 0

```

{
    [!--writers]
    /* write */
    [
        ++sriter
    ] ((0 == sriter))
    when (0 == sriter)
    {
        repeat
    }
}

{
    [!--readers]
    /* read */
    [
        ++siter
    ] ((0 == siter))
    when
    {
        repeat
    }
}

```

Readers-Writers Pseudocode



- also, since `writers` is part of the guards in guarded commands (and these two guards are not identical), in the implementation [1] you must signal/broadcast the corresponding conditions used to implement these guards

```

{
    [':--sreliim]
    /* wrii' /*
    [
        '++sreliim
    ] ((0 == srepeear) ←
?? (0 == sreliim) uehm
    } ( )reliim
}

```

```

{
    [':--srepeear]
    /* paa' /*
    [
        '++srepeear
    ] (0 == sreliim) uehm
    } ( )repeear
}

```

Readers-Writers Pseudocode



```

{
    [!--srəɬʰim]
        /* əɬʰim */
        [
            '++səɬʰim
        ] ((0 == srəɬʰim) uəhm
33 (0 == sɛɬʰim) ) ) sɛɬʰim
}

{
    [!--srəpəər]
        /* pəər */
        [
            '++srəpəər
        ] ((0 == srəɬʰim) uəhm
      ] (0 == sɛɬʰim) ) ) rəpəər
}

```

Readers-Writers Pseudocode



- the sanity checks are really not necessary

```

    {
        [writem]
        /* write */
        ((1 == writers))
        assert(readers == 0) // sanity check
        [
            writers++;
        ] ((0 == readers))
        assert(writers == 0) when (writers == 0)
        } ( ) readm
    }

    {
        [readm]
        /* read */
        ((0 < readers))
        assert(writers == 0) // sanity check
        [
            readers++;
        ] ((0 == writers)) when (writers == 0)
        } ( ) readm
    }
}

```

Pseudocode with Assertions




```
pthread_cond_t cv = PTHREAD_COND_INITIALIZER;

// If a condition variable cannot be initialized statically, do:
int pthread_cond_init(
    pthread_cond_t *cvp,
    pthread_condattr_t *attrp)

int pthread_cond_destroy(
    pthread_cond_t *cvp)
```

Set Up



Readers-Writers Problem

Copyright © William C. Cheng 107

← This is an example of how to give threads priority *without* assigning priorities to threads!

← now it's unfair to the readers

← isn't writing more important than reading anyway?

Solving The Starvation Problem

```

reader ( ) {
    while (writers == 0) {
        readers++;
        [read * /]
        [readers--]
    }
}

writer ( ) {
    while (writers == 0) {
        [writers++]
        [active_writers++]
        [writers--]
        active_writers--;
    }
}

```

Operating Systems - CSCI 402

Copyright © William C. Cheng 105

← one mutex (m) and two condition variables (readersQ and writersQ)

```

reader ( ) {
    pthread_mutex_lock(&m);
    while (writers == 0) {
        pthread_cond_wait(
            &readersQ, &m);
        readers++;
        pthread_mutex_unlock(&m);
        /* read */
        pthread_mutex_lock(&m);
        if (--readers == 0)
            pthread_cond_signal(
                &writersQ);
        pthread_mutex_unlock(&m);
    }
}

writer ( ) {
    pthread_mutex_lock(&m);
    while (readers == 0) {
        pthread_cond_wait(
            &writersQ, &m);
        writers++;
        pthread_mutex_unlock(&m);
        /* write */
        pthread_mutex_lock(&m);
        pthread_cond_broadcast(
            &readersQ);
        writers--;
    }
}

```

Operating Systems - CSCI 402

Copyright © William C. Cheng 103

Readers-Writers Pseudocode

← don't have to worry about this readers

← don't have to worry about this writers

◊ it's *not wrong* to signal/broadcast here, it's just *wasteful/inefficient*

◊ you need to look at your program logic and figure when signal/broadcast conditions won't be useful

```

reader ( ) {
    while (writers == 0) {
        readers++;
        [read * /]
        [readers--]
    }
}

writer ( ) {
    while (readers == 0) {
        [writers++]
        [writers--]
    }
}

```

Operating Systems - CSCI 402

Copyright © William C. Cheng 106

Improved Reader

```

reader ( ) {
    pthread_mutex_lock(&m);
    while (!writers) {
        pthread_cond_wait(
            &readersQ, &m);
        readers++;
        pthread_mutex_unlock(&m);
        /* read */
        pthread_mutex_lock(&m);
        if (--readers == 0)
            pthread_cond_signal(
                &writersQ);
        pthread_mutex_unlock(&m);
    }
}

```

← exactly the same as before!

Operating Systems - CSCI 402

Copyright © William C. Cheng 106

The Starvation Problem

Can the writer never get a chance to write?

← yes, if there are always readers

← so, this implementation can be unfair to writers

← once a writer arrives, shut the door on new readers

← writers now means the number of writers *wanting* to write

← use active_writers to make sure that only one writer can do the actual writing at a time

```

reader ( ) {
    while (writers == 0) {
        readers++;
        [read * /]
        [readers--]
    }
}

writer ( ) {
    while (writers == 0) {
        [writers++]
        [writers--]
    }
}

```

Operating Systems - CSCI 402

Copyright © William C. Cheng 104

Solution with POSIX Threads

to be even more "efficient", can use multiple CVs

← so you don't have to wake up a thread unnecessarily

◊ here we use one CV for reader's guard and one CV for writer's guard (since their expressions are different)

"Synchronization Box"

```

reader ( ) {
    while (writers == 0) {
        readers++;
        [read * /]
        [readers--]
    }
}

writer ( ) {
    while (readers == 0) {
        [writers++]
        [writers--]
    }
}

```

Operating Systems - CSCI 402



- if the n^{th} thread wakes up all the other blocked threads, most likely, **none** of these threads will see `count == n`

```
int count = 0;
pthread_mutex_t m;
pthread_cond_t BarrierQueue;

void barrier_sync()
{
    pthread_mutex_lock(&m);
    while (count > n)
    {
        pthread_cond_wait(&BarrierQueue, &m);
    }
    count = 0;
    pthread_mutex_unlock(&m);
}
```



- if the n^{th} thread wakes up all the other blocked threads, most likely, **none** of these threads will see `count == n`
- `moving count = 0` around won't help
- cannot guarantee all n threads will exit the barrier

```
int count = 0;
pthread_mutex_t m;
pthread_cond_t BarrierQueue;
void barrier_sync()
{
    pthread_mutex_lock(&m);
    if (++count > n)
    {
        while (count > n)
        {
            pthread_cond_wait(&BarrierQueue, &m);
        }
        pthread_cond_broadcast(&BarrierQueue);
    }
    pthread_mutex_unlock(&m);
    count = 0;
}
```




Ex: fork/join (fork to create parallel execution)

- ▶ When a thread reaches a barrier, it must stop (do nothing) and simply wait for other threads to arrive at the same barrier
 - ▬ when all the threads that were suppose to arrive at the barrier have all arrived at the barrier, they are all given the signal to proceed forward
 - the barrier is then reset

- barrier have all arrived at the barrier, they are all given the signal to proceed forward
- the barrier is then reset

signal to proceed forward

when all the threads that were suppose to arrive at the barrier have all arrived at the barrier, they are all given the



http://pubs.opengroup.org/onlinepubs/00969699/functions/pthread_cond_signal.html

- the idea here is to have the last thread broadcast the condition while all the other threads are blocked at waiting for the condition to be signaled
- as it turns out, pthread_cond_wait () might return

- for the condition to be signaled
- as it turns out, pthread_cond_wait () might return

```
int count = 0;
pthread_mutex_t m;
pthread_cond_t BarrierQueue;
void barrier_sync()
{
    pthread_mutex_lock(&m);
    if (++count > n)
    {
        pthread_cond_wait(&BarrierQueue, &m);
    }
    pthread_mutex_unlock(&m);
}
```



```

writer
{
    pthread_mutex_lock(&m);
    writers++;
    while (i() == 0)
    {
        (active_writers)++;
        pthread_cond_wait(&writersQ, &m);
    }
    (active_writers) == 0
    {
        pthread_cond_broadcast(&writersQ);
        pthread_mutex_unlock(&m);
    }
}

```

```
int pthread_rwlock_init(&rwlock, attr) {
    pthread_rwlock_t *lock,
    pthread_rwlockattr_t *attr;
    int ret;

    lock = malloc(sizeof(pthread_rwlock_t));
    if (!lock) return -1;

    ret = pthread_rwlockattr_tattr(lock, attr);
    if (ret != 0) return ret;

    return 0;
}
```





- ≡ don't use count in the guard since it's problematic!
- ≡ introduce a new guard

```
int count = 0;
pthread_mutex_t m;
pthread_cond_t BarrierQueue;

void barrier_sync() {
    pthread_mutex_lock(&m);
    if (++count > number) {
        int my_generation = generation;
        while(my_generation == generation)
            pthread_cond_wait(&BarrierQueue, &m);
    } else {
        count = 0;
        generation++;
        pthread_cond_broadcast(&BarrierQueue);
    }
    pthread_mutex_unlock(&m);
}
```

Barrier in POSIX Threads



More From POSIX!

```
int pthread_barrier_init(
    pthread_barrier_t *barrier,
    pthread_barrierattr_t *attr,
    unsigned int count);
int pthread_barrier_destroy(
    pthread_barrier_t *barrier);
int pthread_barrier_wait(
    pthread_barrier_t *barrier);
```