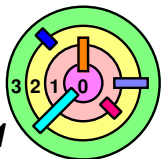


# Ch 2: Multithreaded Programming


Bill Cheng

*<http://merlot.usc.edu/cs402-s16>*




# Overview

- ➡ Why threads?
- ➡ How to program with threads?
  - what is the API?
- ➡ Synchronization
  - mutual exclusion
  - semaphores
  - condition variables
- ➡ Pitfall of thread programmings



# Concurrency


- ➡ Many things occur simultaneously in the OS
  - e.g., data coming from a disk, data coming from the network, data coming from the keyboard, mouse got clicked, jobs need to get executed
- ➡ If you have multiple processors, you may be able to handle things in parallel
  - that's real concurrency
- ➡ If you only have one processor, you may want to make it look like things are running in parallel
  - do multiplexing to create the illusion
  - as it turns out, it's a good idea to do this even if you have only have one processor
- ➡ The down side is that if you want concurrency, you have to have *concurrency control* or bad things can happen






# Why Threads?





- ➡ Many things are easier to do with threads
  - *multithreading* is a *powerful paradigm*
  - makes your design *cleaner*, and therefore, less buggy
- ➡ Many things run faster with threads
  - if you are just waiting, don't waste CPU cycles, give the CPU to someone else, *without explicitly* giving up the CPU
- ➡ *Kernel threads* vs. *user threads*
  - basic concepts are the same
  - can easily do programming assignments for user-level threads
    - that's why we start here (to get your warmed up)!
    - for kernel programming assignments, you need to fill out missing parts of various kernel threads




# A Simple Example: rlogind



# A Simple Example: rlogind



- for a socket, `l_in` = `l_out`, i.e., you read and write using the same file descriptor



# Life Without Threads

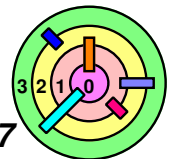

```

logind(int r_in, int r_out, int l_in, int l_out) {
    fd_set in = 0, out;
    int want_l_write = 0, want_r_write = 0;
    int want_l_read = 1, want_r_read = 1;
    int eof = 0, tsize, fsize, wret;
    char fbuf[BSIZE], tbuf[BSIZE];

    fcntl(r_in, F_SETFL, O_NONBLOCK);
    fcntl(r_out, F_SETFL, O_NONBLOCK);
    fcntl(l_in, F_SETFL, O_NONBLOCK);
    fcntl(l_out, F_SETFL, O_NONBLOCK);

    while(!eof) {
        FD_ZERO(&in);
        FD_ZERO(&out);
        if (want_l_read) FD_SET(l_in, &in);
        if (want_r_read) FD_SET(r_in, &in);
        if (want_l_write) FD_SET(l_out, &out);
        if (want_r_write) FD_SET(r_out, &out);
        select(MAXFD, &in, &out, 0, 0);
        if (FD_ISSET(l_in, &in)) {
            if ((tsize = read(l_in, tbuf, BSIZE)) > 0) {
                want_l_read = 0;
                want_r_write = 1;
            } else { eof = 1; }
        }
    }
}

```





# Life Without Threads

```


if (FD_ISSET(r_in, &in)) {
    if ((fsize = read(r_in, fbuf, BSIZE)) > 0) {
        want_r_read = 0;
        want_l_write = 1;
    } else { eof = 1; }
}
if (FD_ISSET(l_out, &out)) {
    if ((wret = write(l_out, fbuf, fsize)) == fsize) {
        want_r_read = 1;
        want_l_write = 0;
    } else if (wret >= 0) {
        tsize -= wret;
    } else { eof = 1; }
}
if (FD_ISSET(r_out, &out)) {
    if ((wret = write(r_out, tbuf, tsize)) == tsize) {
        want_l_read = 1;
        want_r_write = 0;
    } else if (wret >= 0) {
        tsize -= wret;
    } else { eof = 1; }
}
}
}

```





# Life With Threads



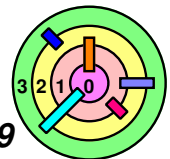
```
incoming(int r_in, int l_out) {
    int eof = 0;
    char buf[BSIZE];
    int size;

    while (!eof) {
        size = read(r_in, buf, BSIZE);
        if (size <= 0)
            eof = 1;
        if (write(l_out, buf, size) <= 0)
            eof = 1;
    }
}
```



```
outgoing(int l_in, int r_out) {
    int eof = 0;
    char buf[BSIZE];
    int size;

    while (!eof) {
        size = read(l_in, buf, BSIZE);
        if (size <= 0)
            eof = 1;
        if (write(r_out, buf, size) <= 0)
            eof = 1;
    }
}
```


— don't have to call `select()`




# Single-Threaded Database Server



# Multithreaded Database Server



- will be very difficult to implement this without using threads if you want to handle a large number of requests simultaneously



# 2.2 Programming With Threads


➡ *Threads Creation & Termination*

➡ **Threads & C++**

➡ **Synchronization**

➡ **Thread Safety**

➡ **Deviations**



# Creating a POSIX Thread

➡ `man pthread_create`


## SYNOPSIS

```
#include <pthread.h>
```

```
int pthread_create(  
    pthread_t *thread,  
    const pthread_attr_t *attr,  
    void *(*start_routine)(void *),  
    void *arg);
```

Compile and link with `-pthread`.

- ➡ the `start_routine` is also known as the "***first procedure***" or "***thread function***" of the child thread
  - it's like `main()` for the child thread
- ➡ the "thread ID" of the newly created thread will be returned in the first argument of `pthread_create()`
  - may ***not*** be a ***Thread Control Block***



# Creating a POSIX Thread


```
start_servers( ) {
    pthread_t thread;
    int i;
    for (i=0; i<nr_of_server_threads; i++)
        pthread_create(&thread,    // thread ID
                      0,           // default attributes
                      server,      // first procedure
                      argument);   // argument
}
```

```
void *server(void *arg) {
    // perform service
    return(0);
}
```

← child thread **starts** executing here  
arg = argument (from caller)

← child thread **ends** when **return**  
from its **start routine / first procedure**

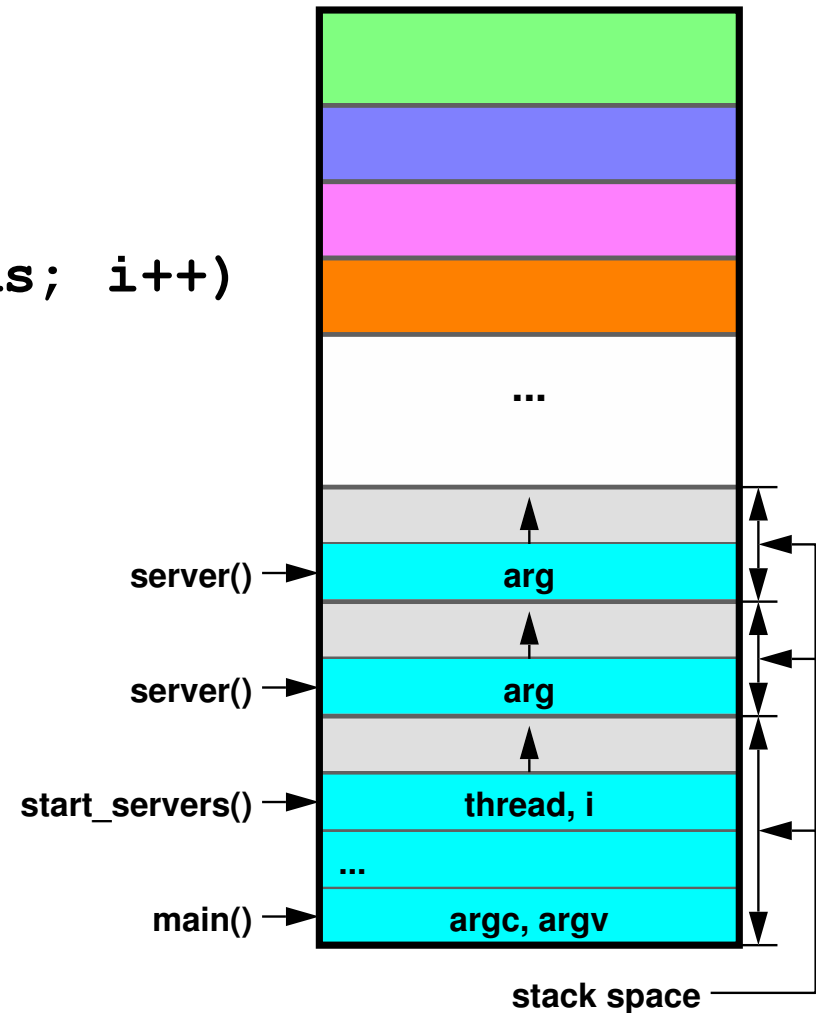
- ▢ pthread\_create() returns 0 if successful
- ▢ POSIX 1003.1c standard



# Creating a POSIX Thread

```
start_servers( ) {
    pthread_t thread;
    int i;
    for (i=0; i<nr_of_server_threads; i++)
        pthread_create(&thread,
                      0,
                      server,
                      argument);
}

void *server(void *arg) {
    // perform service
    return(0);
}
```



- ▢ every thread needs a separate stack
  - first stack frame in every child thread corresponds to `server()`
  - ◆ one `arg` in each of these stack frames


# Creating a POSIX Thread

- ➡ These are the same:
- ▮ keep thread handle in the stack

```
pthread_t thread;  
pthread_create(&thread, ...);
```
  - ▮ keep thread handle in the heap

```
pthread_t *thread_ptr =  
    (pthread_t*)malloc(sizeof(pthread_t));  
pthread_create(thread_ptr, ...);
```
  - need to make sure that eventually you will call the following to not leak memory

```
free(thread_ptr);
```






# Creating a Win32 Thread

```
start_servers( ) {  
    HANDLE thread;  
    DWORD id;  
    int i;  
    for (i=0; i<nr_of_server_threads; i++)  
        thread = CreateThread(  
            0,           // security attributes  
            0,           // default # of stack pages allocated  
            server,       // first procedure  
            arg,          // argument  
            0,           // default attributes  
            0,           // creation flags  
            &id);         // thread ID  
}
```


```
DWORD WINAPI server(void *arg) {  
    // perform service  
    return(0);  
}
```

— We won't talk about Win32 much



# Complications

```
rlogind(int r_in, int r_out, int l_in, int l_out) {  
    pthread_t in_thread, out_thread;  
  
    pthread_create(&in_thread,  
                  0,  
                  incoming,  
                  r_in, l_out); // Cannot do this ...  
    pthread_create(&out_thread,  
                  0,  
                  outgoing,  
                  l_in, r_out); // Cannot do this ...  
    /* How do we wait till they are done? */  
}
```



# Multiple Arguments


```
typedef struct {
    int first, second;
} two_ints_t;

rlogind(int r_in, int r_out, int l_in, int l_out) {
    pthread_t in_thread, out_thread;

    two_ints_t in={r_in, l_out}, out={l_in, r_out};
    pthread_create(&in_thread,
                  0,
                  incoming,
                  &in);


    ...
    /* How do we wait till they are done? */
}

void *incoming(void *arg) {
    two_ints_t *p=(two_ints_t*)arg;
    ... p->first ...
    return NULL;
}
```



# Multiple Arguments

```
typedef struct {
    int first, second;
} two_ints_t;
```



```
rlogind(int r_in, int r_out, int l_in, int l_out) {
    pthread_t in_thread, out_thread;
```

```
    two_ints_t in={r_in, l_out}, out={l_in, r_out};
    pthread_create(&in_thread,
```

```
        0,
        incoming,
        &in);
```




```
    ...
```

```
    /* How do we wait till they are done? */
```


```
}
```

```
void *incoming(void *arg) {
    two_ints_t *p=(two_ints_t*)arg;
    p->first ...
    ...
}
```




# Multiple Arguments

- ➡ Need to be careful how to pass argument to new thread when you call `pthread_create()`
- there is no way to pass multiple arguments in either POSIX or Win32
  - passing address of a *local* variable (like the previous example) only works if we are certain the this storage doesn't go out of scope until the thread is done with it
  - passing address of a *static* or a *global* variable only works if we are certain that only one thread at a time is using the storage
  - passing address of a *dynamically* allocated storage only works if we can free the storage when, and only when, the thread is finished with it
    - this would not be a problem if the language supports garbage collection
- ➡ Ask yourself, "How can you be sure?"
- if the answer is, "I hope it works", then you need a different solution



# When Is The Child Thread Done?

```
rlogind(int r_in, int r_out, int l_in, int l_out) {  
    pthread_t in_thread, out_thread;  
    two_ints_t in={r_in, l_out}, out={l_in, r_out};  
  
    pthread_create(&in_thread, 0, incoming, &in);  
    pthread_create(&out_thread, 0, outgoing, &out);  
  
    pthread_join(in_thread, 0);  
    pthread_join(out_thread, 0);  
}
```



# Thread Termination




## Thread return values

- which threads receive these values
- how do they do it?
  - clearly, receiving thread must wait until the producer thread produced it, i.e., producer thread has terminated
  - so we must have a way for one thread to wait for another thread to terminate
- must have a way to say which thread you are waiting for
  - need a unique identifier
  - tricky if it can be reused



## To wait for another thread to terminate

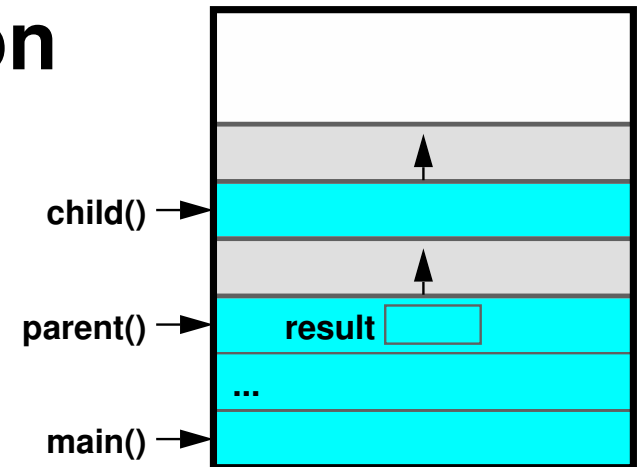
```
int pthread_join(thread_t thread,  
                 (void **)ret_value);
```



# Thread Termination

➡ How does a thread *self-terminate*?

- 1) return from its "first procedure"
  - ◆ return a value of type (void\*)
- 2) call `pthread_exit (ret_value)`
  - ◆ `ret_value` is of type (void\*)




```
parent () {
    pthread_t thread;
    void *result;
    pthread_create(&thread,
        0, child, 0);
    pthread_join(thread,
        (void**)&result);
    switch ((int)result) {
    case 1: ...
    case 2: ...
    }
    ...
}
```

```
void *child(void *arg) {
    ...
    if (terminate_now) {
        pthread_exit((void*)1);
    }
    return((void*)2);
}
```

Thread Control Block

TID
Exit/Return Code
...







# Thread Termination

## ➡ Difference between `pthread_exit()` and `exit()`

- ▬ `pthread_exit()` terminates only the calling thread
- ▬ `exit()` terminates the process, including all threads running in it
  - it will not wait for any thread to terminate
  - what will this code do?

```
int main(int argc, char *argv[]) {  
    // create all the threads  
    return(0);  
}
```

- when `main()` returns, `exit()` will be called
  - ◆ as a result, none of the created child threads may get a chance to run




# Thread Termination

- ➡ Difference between `pthread_exit()` and `exit()`
- ▬ `pthread_exit()` terminates only the calling thread
  - ▬ `exit()` terminates the process, including all threads running in it
  - it will not wait for any thread to terminate
  - what about this code?

```
int main(int argc, char *argv[]) {  
    // create all the threads  
    pthread_exit(0); // exit the main thread  
    return(0);  
}
```

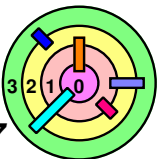
- here, `pthread_exit()` will terminate the main thread, so `exit()` is never called
  - ◆ as it turns out, this special case is taken care of in the pthread library implementation

- ➡ You should use `pthread_join()` unless you are absolutely sure




# Thread Termination

- ➡ Any thread can join with any other thread
  - there's *no parent/child relationships* among threads
    - unlike process termination and `wait()`
- ➡ What happens if a thread terminates and no other thread wants to join with this thread?
  - it also goes into a *zombie* state
    - all the thread related information is freed up, except for the thread ID and return code
- ➡ What if two threads want to join with the same thread?
  - after the first thread joins, the thread ID and return code are freed up and the thread ID may get reused
  - so don't do this!



# Detached Threads

```
start_servers( ) {  
    pthread_t thread;  
    int i;  
    for (i=0; i<nr_of_server_threads; i++) {  
        pthread_create(&thread, 0, server, 0);  
        pthread_detach(thread);  
    }  
    ...  
}  
  
server( ) {  
    ...  
}
```



# Types


```
pthread_create(&tid,
               0,
               (void *(*)(void *))func,
               (void *)1);
```

```
int func = 4; // func definition 1
```

```
void func(int i) { // func definition 2
    ...
}
```

```
void *func(void *arg) { // func definition 3
    int i = (int)arg;
    ...
    return(0);
}
```


⇒ a function is just an address (of something in the text/code segment)



# Thread Attributes

```
pthread_t thread;  
pthread_attr_t thr_attr;  
  
pthread_attr_init(&thr_attr);  
/* establish some attributes */  
...  
pthread_create(&thread, &thr_attr, startroutine, arg);  
pthread_attr_destroy(&thr_attr);  
...
```


- thread attribute only needs to be valid when a thread is created
  - therefore, it can be destroyed as soon as the thread is created



# Stack Size

```
pthread_t thread;  
pthread_attr_t thr_attr;  
  
pthread_attr_init(&thr_attr);  
pthread_attr_setstacksize(&thr_attr, 20*1024*1024);  
...  
pthread_create(&thread, &thr_attr, startroutine, arg);  
pthread_attr_destroy(&thr_attr);
```

- ▢ the above code set the stack size to 20MB
- ▢ the default stack size is very large
  - if you need to create a lot of threads, you need to control the stack size
  - default stack size is probably around 1MB in Solaris and 8MB in some Linux implementations



# Example

```

#include <stdio.h>
#include <pthread.h>
#include <string.h>

#define M 3
#define N 4
#define P 5

int A[M][N];
int B[N][P];
int C[M][P];

void *matmult(void *arg) {
    int row = (int)arg, col;
    int i, t;

    for (col=0; col < P; col++) {
        t = 0;
        for (i=0; i<N; i++)
            t += A[row][i] * B[i][col];
        C[row][col] = t;
    }
    return(0);
}


```

```

main( ) {
    int i;
    pthread_t thr[M];
    int error;

    /* initialize the matrices ... */
    ...
    // create the worker threads
    for (i=0; i<M; i++) {
        if (error = pthread_create(
            &thr[i],
            0,
            matmult,
            (void *)i)) {
            fprintf(stderr,
                "pthread_create: %s",
                strerror(error));
            exit(1);
        }
    }
    // wait for workers to finish
    for (i=0; i<M; i++)
        pthread_join(thr[i], 0)
    /* print the results ... */
}


```





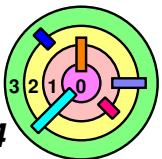
# Compiling It

```
% gcc -o mat mat.c -lpthread
```




## 2.2.3 Synchronization

- ➡ In real life, "synchronization" means that you want to do things at the same time
- ➡ In computer science, "synchronization" could mean the above, **OR**, it means that you want to **prevent** do things at the same time



# Mutual Exclusion



# Threads and Mutual Exclusion


Thread 1:

$x = x + 1;$

Thread 2:

$x = x + 1;$

- ▢ looks like it doesn't matter how you execute,  $x$  will be incremented by 2 in the end
  - choices are
    - ◆ thread 1 executes  $x = x + 1$  then thread 2 executes  $x = x + 1$
    - ◆ thread 2 executes  $x = x + 1$  then thread 1 executes  $x = x + 1$
  - are there other choices?




# Threads and Mutual Exclusion

Thread 1:

```
x = x+1;
/*
ld    r1, x
add   r1, 1
st    r1, x
*/
```

Thread 2:

```
x = x+1;
/*
ld    r1, x
add   r1, 1
st    r1, x
*/
```



➡ Unfortunately, machines do not execute high-level language statements

- they execute machine instructions
- now if thread 1 executes the first (or two) machine instructions
- context switch!
  - how can this happen?
- then thread 2 executes all 3 machine instructions
- then later thread 1 executes the remaining machine instructions
- x would have only increased by 1




# Threads and Synchronization

```
// shared by both threads
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
int x;
...
pthread_mutex_lock(&m);
x = x+1;
pthread_mutex_unlock(&m);
```

➡ Locking a mutex is like getting the key to a safe-deposit box

*critical section*



- code between `pthread_mutex_lock()` and `pthread_mutex_unlock()` for a particular *mutex* is called a *critical section with respect to that mutex*
  - all the critical sections *with respect to a particular mutex* are "*mutually exclusive*"
    - ◆ the system (not necessarily the OS) guarantees that only *one* critical section can be executing at any point in time
  - how it's really done will be covered in Ch 5



# Set Up

```
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
```




If a mutex cannot be initialized statically, do:

```
int pthread_mutex_init(  
    pthread_mutex_t *mutexp,  
    pthread_mutexattr_t *attrp)
```

```
int pthread_mutex_destroy(  
    pthread_mutex_t *mutexp)
```



Usually, mutex attributes are not used



# Taking Multiple Locks




**Mutex is not a cure-all**

— when you have more than one locks, you may get into trouble

```
proc1( ) {  
    pthread_mutex_lock(&m1);  
    /* use object 1 */  
    pthread_mutex_lock(&m2);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m2);  
    pthread_mutex_unlock(&m1);  
}
```

```
proc2( ) {  
    pthread_mutex_lock(&m2);  
    /* use object 2 */  
    pthread_mutex_lock(&m1);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m1);  
    pthread_mutex_unlock(&m2);  
}
```






# Taking Multiple Locks

```
proc1( ) {  
    pthread_mutex_lock(&m1);  
    /* use object 1 */  
    pthread_mutex_lock(&m2);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m2);  
    pthread_mutex_unlock(&m1);  
}
```

```
proc2( ) {  
    pthread_mutex_lock(&m2);  
    /* use object 2 */  
    pthread_mutex_lock(&m1);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m1);  
    pthread_mutex_unlock(&m2);  
}
```

➡ Graph representation ("wait-for" graph)



# Necessary Conditions For Deadlocks

➡ All 4 conditions below must be met in order for a deadlock to be possible (no guarantee that a deadlock may occur)

## 1) Bounded resources

➡ only a finite number of threads can have concurrent access to a resource

## 2) Wait for resources


➡ threads wait for resources to be freed up, without releasing resources that they hold

## 3) No preemption

➡ resources cannot be *revoked* from a thread


## 4) Circular wait

➡ there exists a set of waiting threads, such that each thread is waiting for a resource held by another




# Dealing with Deadlock


- ➡ Deadlock is a programming bug
  - one of the oldest bug
  - it's a tricky one because it only deadlocks *sometimes*
- ➡ Hard
  - is the system deadlocked?
  - will this move lead to deadlock?
  - this is *detection*
    - if you can detect deadlocks, what do you do after you have detected them?
- ➡ Easy
  - restrict use of mutexes so that deadlock cannot happen
  - this is *prevention*
- ➡ Deadlock is a complicated subject
  - some textbooks spend an entire chapter on deadlocks
  - we will only look at a couple of cases




# Deadlock Prevention: Lock Hierarchies




- organize mutexes into levels
- must not try locking a mutex at level  $i$  if already holding a mutex at level  $j$  if  $i \leq j$ , otherwise it's okay
  - e.g., if hold mutexes at levels 2 and 3, can only wait for a mutex at levels 4 or higher




# Deadlock Prevention: Lock Hierarchies




- organize mutexes into levels
- must not try locking a mutex at level  $i$  if already holding a mutex at level  $j$  if  $i \leq j$ , otherwise it's okay
  - e.g., if holding mutexes at levels 2 and 3, can only wait for a mutex at levels 4 or higher



# Deadlock Prevention: Lock Hierarchies




What if you cannot organize your mutexes in such strict order for deadlock detection?



# Deadlock Prevention: Conditional Locking



```
proc1( ) {  
    pthread_mutex_lock(&m1);  
    /* use object 1 */  
    pthread_mutex_lock(&m2);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m2);  
    pthread_mutex_unlock(&m1);  
}  
  
proc2( ) {  
    while (1) {  
        pthread_mutex_lock(&m2);  
  
        if (!pthread_mutex_trylock(&m1))  
            break;  
        pthread_mutex_unlock(&m2);  
    }  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m1);  
    pthread_mutex_unlock(&m2);  
}
```



# One Mutex, Multiple Critical Sections

```
f1() {
    pthread_mutex_lock(&m);
    x++;
    pthread_mutex_unlock(&m);
}
```

```
f2() {
    pthread_mutex_lock(&m);
    x--;
    pthread_mutex_unlock(&m);
}
```



- I use "Synchronization Box" to mean executing one of many critical section code with respect to one particular mutex




# One Mutex, Multiple Critical Sections

```
f1 () {
    pthread_mutex_lock (&m);
    x++; } critical section
    pthread_mutex_unlock (&m);
}
```

```
f2 () {
    pthread_mutex_lock (&m);
    x--; } critical section
    pthread_mutex_unlock (&m);
}
```

"Synchronization Box"



— I use "Synchronization Box" to mean executing one of many critical section code with respect to one particular mutex


- ➡ By calling `pthread_mutex_lock (&m)`, a thread can be placed into a *queue* and *wait* there *indefinitely* for mutex `m` to become available
- multiple threads would join this queue
  - queue is served one at a time, like a supermarket checkout
  - when it's your thread's turn, `pthread_mutex_lock ()` returns with the mutex locked, your thread can execute critical section code, and then release the mutex




# Beyond Mutexes

- ➡ **Mutex is necessary when shared data is being modified**
  - ▬ although there are cases where using a mutex is an overkill (i.e., too restrictive and inefficient and would lock threads out when it's not necessary)
    - we would like to have better concurrency (i.e., *"fine-grained parallelism"*) when complete mutual exclusion is not required
  - ▬ two major categories to illustrate this
    - 1) what if threads don't interfere one another most of the time and synchronization is only required occasionally?
      - ◆ e.g., *Producer-Consumer problem* (a.k.a., bounded-buffer problem)
    - 2) what if some threads just want to *look at* (i.e., *read*) a piece of data?
      - ◆ e.g., *Readers-Writers problem*


## ➡ Barrier Synchronization




# Producer-Consumer Problem




- ➡ A circular buffer is used
- ➡ Most of the time, no interference
  - if you use a *single mutex* to lock the entire array of buffers, it's an overkill (*i.e., too inefficient*)
- ➡ When does it require synchronization?



# Producer-Consumer Problem



- ➡ A circular buffer is used
- ➡ Most of the time, no interference
  - if you use a *single mutex* to lock the entire array of buffers, it's an overkill (*i.e., too inefficient*)
- ➡ When does it require synchronization?
  - producer needs to be blocked when all slots are full
  - consumer needs to be blocked when all slots are empty



# Guarded Commands

```


when (guard) [
  /*
    once the guard is true,
    execute this code atomically
  */
  ...
]

```

} *command sequence*

- ⇒ this means that the command sequence *can be* executed (*atomically*) *at any time* the *guard* is evaluated to be *true*
  - a *guard* is a *boolean expression* (evaluates to true or false)
  - *atomically* mean that it's executed without interruption
    - ◆ *evaluting the guard* and *executing the command sequence* altogether is an *atomic operation* *if* the *guard is true*
    - ◆ you cannot evaluate the guard if your thread is not running

➡ For exams, you need to know how to write simple pseudo-code in the language of *Guarded Commands*



# Guarded Commands

```
when (guard) [
```

```
  /*
```



```
    once the guard is true,  
    execute this code atomically
```

```
  */
```

```
  ...
```

```
]
```

*command sequence*



# Guarded Commands

```
when (guard) [
```

```
  /*
```


```
    once the guard is true,  
    execute this code atomically
```

```
  */
```

```
  ...
```

```
]
```

*command sequence*



# Guarded Commands

```
when (guard) [
```

```
  /*
```


```
    once the guard is true,  
    execute this code atomically
```

```
  */
```

```
  ...
```

```
]
```

*command sequence*





# Guarded Commands

```
when (guard) [
```

```
  /*
```


```
    once the guard is true,  
    execute this code atomically
```

```
  */
```

```
  ...
```

```
]
```

*command sequence*



# Guarded Commands

```
when (guard) [
```

```
  /*
```


```
    once the guard is true,  
    execute this code atomically
```

```
  */
```

```
  ...
```

```
]
```

*command sequence*



# Guarded Commands

```
when (guard) [
```

```
  /*
```


```
    once the guard is true,  
    execute this code atomically
```

```
  */
```

```
  ...
```

```
]
```

*command sequence*




# Guarded Commands

```


when (guard) [
  /*
    once the guard is true,
    execute this code atomically
  */
  ...
]

```

*command sequence*



- please understand that ***command sequence***  $\neq$  ***critical section***
  - evaluate the guard to be true ***and*** execute command sequence ***together*** is done inside one critical section



# Guarded Commands

```
when (guard) [
```

```
  /*
```


```
    once the guard is true,  
    execute this code atomically
```

```
  */
```

```
  ...
```


```
]
```

*command sequence*




guard evaluate to be true and  
execute command sequence


- **atomic**: as if it's executed in an instance of time (duration = 0)
  - this is okay because it's just pseudo-code



# Semaphores



- ➡ A **semaphore**,  $S$ , is a **nonnegative integer** on which there are exactly two operations defined by two guarded commands
- ➡  $P(S)$  operation (implemented as a guarded command):
    - **when**  $(S > 0)$  [
      - $S = S - 1;$
  - ➡  $V(S)$  operation (implemented as a guarded command):
    - $[S = S + 1;]$
  - ➡ there are no other means for manipulating the value of  $S$ 
    - other than initializing it



# Mutexes with Semaphores

```
semaphore S = 1;
```

```
void OneAtATime( ) {
    P(S);
    ...
    /* code executed mutually
       exclusively */
    ...
    V(S);
}
```


⇒ P(S) operation:

○ when (S > 0) [  
     S = S - 1;  
   ]

⇒ V(S) operation:

○ [S = S + 1;]

⇒ this is known as a *binary semaphore*



# Implement A Mutex With A Binary Semaphore



Instead of doing

```
pthread_mutex_lock (&m) ;  
x = x+1 ;  
pthread_mutex_unlock (&m) ;
```

= do:


```
S = 1 ;  
P (S) ;  
x = x+1 ;  
V (S) ;
```



So, you can lock a data structure using a binary semaphore

= this looks just like mutex, what have we really gained?

- if you use it this way, nothing





# Mutexes with Semaphores

```
semaphore S = N;

void NAtATime( ) {
    P(S);
    ...
    /* no more than N threads
       here at once */
    ...
    V(S);
}
```

= P(S) operation:

○ when (S > 0) [  
     S = S - 1;  
   ]

= V(S) operation:

○ [S = S + 1;]

= this is known as a *counting semaphore*

= can be used to solve the producer-consumer problem


➡ Main difference between a semaphore and a mutex

= if a thread locks a mutex, it's holding the lock

○ therefore, it must be that thread that unlocks that mutex

= one thread performs a P operation on a semaphore, *another thread* performs a V operation on the same semaphore

○ this is often why you would use a semaphore




# Producer/Consumer with Semaphores

```
Semaphore empty = B;  
Semaphore occupied = 0;  
int nextin = 0;  
int nextout = 0;
```

```
void Produce(char item) {  
    P(empty);  
    buf[nextin] = item;  
    nextin = nextin + 1;  
    if (nextin == B)  
        nextin = 0;  
    V(occupied);  
}
```

```
char Consume( ) {  
    char item;  
    P(occupied);  
    item = buf[nextout];  
    nextout = nextout + 1;  
    if (nextout == B)  
        nextout = 0;  
    V(empty);  
    return item;  
}
```



# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```



```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	8
occupied	0
nextin	0
nextout	0

Consumer

Producer



# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```



```
void Produce(char item) {
  → P(empty);
  buf[nextin] = item;
  nextin = nextin + 1;
  if (nextin == B)
    nextin = 0;
  V(occupied);
}
```

```
char Consume( ) {
  char item;
  →✗ P(occupied);
  item = buf[nextout];
  nextout = nextout + 1;
  if (nextout == B)
    nextout = 0;
  V(empty);
  return(item);
}
```

empty	8
occupied	0
nextin	0
nextout	0

Consumer

Producer



# Producer/Consumer with Semaphores



```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    P(empty);
    → buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    →✗ P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	7
occupied	0
nextin	0
nextout	0

Consumer ←      → Producer



# Producer/Consumer with Semaphores



```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    → nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    →✗ P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	7
occupied	0
nextin	0
nextout	0

Consumer ←      → Producer



# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```



```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    → if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    → × P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	7
occupied	0
nextin	1
nextout	0

Consumer

Producer



# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```



```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    → V(occupied);
}
```

```
char Consume( ) {
    char item;
    →✗ P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	7
occupied	0
nextin	1
nextout	0

Consumer

Producer





# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```


```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    → P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	7
occupied	1
nextin	1
nextout	0

Consumer

Producer



≡ note: producer  
continue to produce

# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```


```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    → item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```


empty	7
occupied	0
nextin	1
nextout	0

Consumer

Producer



≡ note: producer  
continue to produce



# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```


```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    → nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	7
occupied	0
nextin	1
nextout	0

Consumer

Producer



≡ note: producer  
continue to produce

# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```


```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```



empty	7
occupied	0
nextin	1
nextout	1

Consumer ←      → Producer



≡ note: producer  
continue to produce

# Producer/Consumer with Semaphores


```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```


```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    → V(empty);
    return(item);
}
```

empty	7
occupied	0
nextin	1
nextout	1

Consumer ←      → Producer



≡ note: producer  
continue to produce



# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```


```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return item;
}
```



empty	8
occupied	0
nextin	1
nextout	1

Consumer ←      → Producer



≡ note: producer  
continue to produce


# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```


```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	8
occupied	0
nextin	1
nextout	1



= note: producer  
continue to produce




# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

- if produce and consume at same rate, no one waits
- if producer is fast and consumer slow, producer may wait
- if consumer is fast and producer slow, consumer may wait





# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```


```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

➡ **Mutex** by itself is more "**coarse grain**"

➡ you may use one mutex to control access to the number of empty and occupied cells, nextin, and nextout

➡ **Semaphore** is more "**fine grain**"


➡ but **not general** enough



# POSIX Semaphores

```
#include <semaphore.h>
sem_t semaphore;
int err;
int pshared = 0;    // not shared among processes
int init_value = B; // initial value

err = sem_init(&semaphore, pshared, init_value);
err = sem_destroy(&semaphore);
err = sem_wait(&semaphore);    /* P operation */
err = sem_trywait(&semaphore); /* conditional P
                                operation
                                */
err = sem_post(&semaphore);    /* V operation */
```



# Producer-Consumer with POSIX Semaphores


```
void produce(char item) {
    sem_wait(&empty);
    buf[nextin++] = item;
    if (nextin >= B)
        nextin = 0;
    sem_post(&occupied);
}
```

```
char consume() {
    char item;
    sem_wait(&occupied);
    item = buf[nextout++];
    if (nextout >= B)
        nextout = 0;
    sem_post(&empty);
    return(item);
}
```

---

```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```




# Implementation Of Guarded Commands

```

when (guard) [
    /* command sequence */
    . . .
]

```

- ➡ In general, the *guard* can be *complicated* and involving the evaluation of several variables (e.g.,  $a > 3 \ \&\& \ f(b) \leq c$ )
- ➡ the guard (which evaluates to either true or false) *keeps changing its value, continuously* and by multiple threads *simultaneously*
  - ➡ how can we "capture" the instance of time when it evaluates to true so we can execute the command sequence atomically?
    - we have to "sample" it, i.e., take snap shot of all the variables that are involved and then evaluate it
    - a mutex is involved, but how?
      - ◆ need to be efficient
    - need something else (*known as condition variables*)
    - need a bunch of *rules* to follow



# Implementation Of Guarded Commands

```


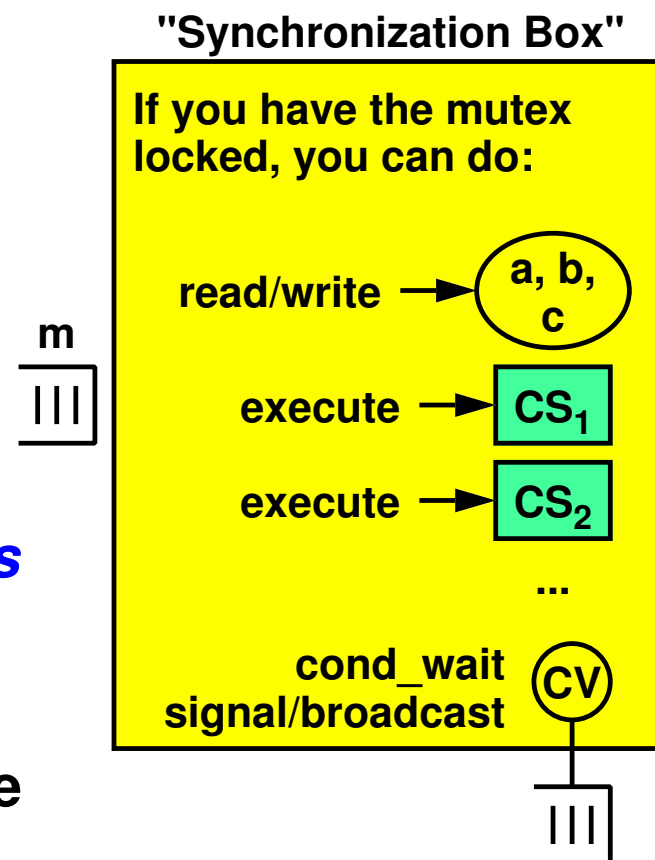
when (guard) [
    /* command sequence */
    ...
]

```

➡ POSIX provides *condition variables* for programmers to implement guarded commands

— a *condition variable* is a *queue of threads* waiting for some sort of notification (an "*event*" or "*condition*")

- threads, waiting for a guard to become true, join such a queue
  - ◆ they wait for a specific *condition* to be *signaled*
  - ◆ they wait for *the right time* to *evaluate the guard*
- (cont...)



# Implementation Of Guarded Commands

```


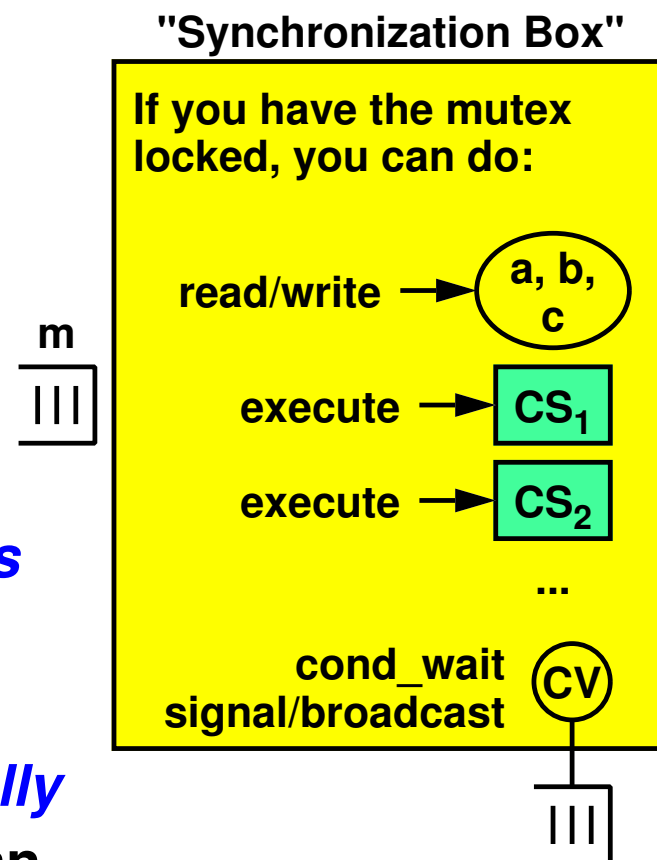
when (guard) [
    /* command sequence */
    . . .
]

```

➡ POSIX provides *condition variables* for programmers to implement guarded commands

— a *condition variable* is a *queue of threads* waiting for some sort of notification (an "*event*" or "*condition*")

- threads that do something to *potentially* change the truth value of the guard can then wake up the threads that were waiting in the queue
  - ◆ they can *signal* or *broadcast* the *condition*
  - ◆ *no guarantee* that the guard will be true when it's time for *another thread* to evaluate the guard



# Implementation Of Guarded Commands

```

when (guard) [
    /* command sequence */
    ...
]

```

➡ POSIX provides *condition variables* for programmers to implement guarded commands


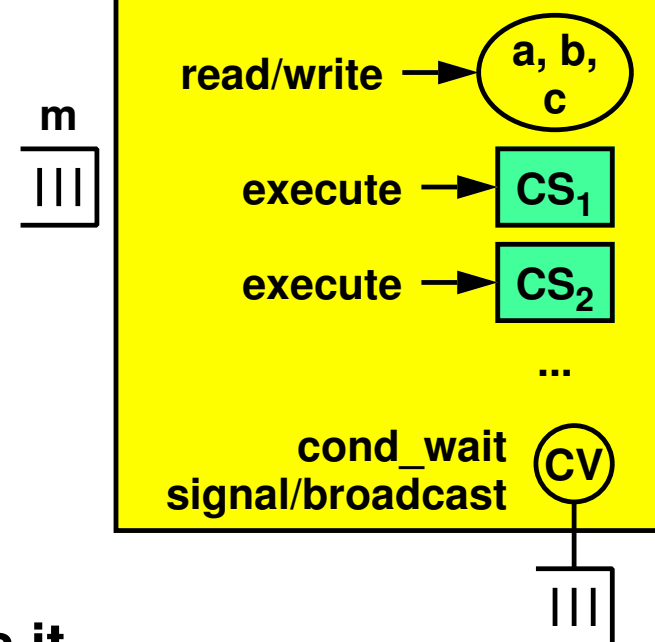
— conceptually, an "event"

(signaling/broadcasting of a condition)  
happens in an instance of time (duration of this "event" is zero)

- if you are not waiting for it, you'll miss it
- how do you make sure you won't miss an event?
  - ◆ you have to *follow the protocol* (for multiple interacting threads to follow) described here

## "Synchronization Box"

If you have the mutex locked, you can do:



# Implementation Of Guarded Commands

```

when (guard) [
    /* command sequence */
    ...
]

```

➡ POSIX provides *condition variables* for programmers to implement guarded commands

```


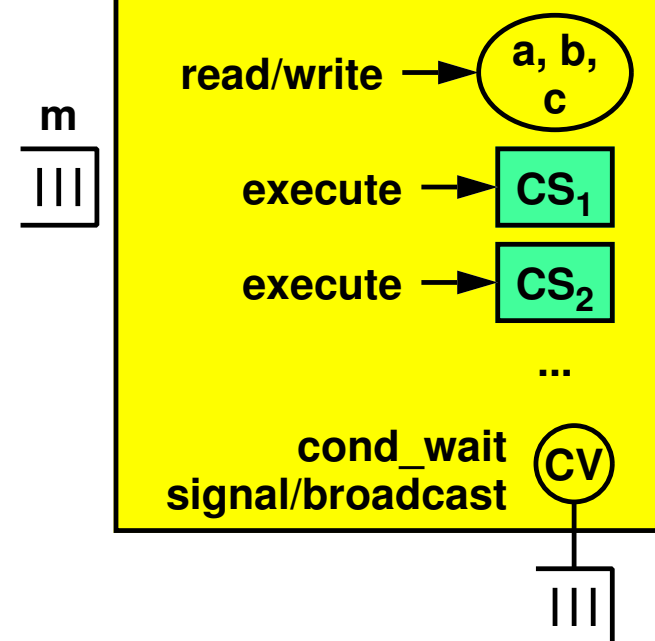
1) pthread_cond_wait (
    pthread_cond_t *cv,
    pthread_mutex_t *mutex)

```

- should only call `pthread_cond_wait()` if you have the mutex *locked*
- *atomically* unlocks mutex and wait for the "event"
- when the event is signaled/broadcasted, `pthread_cond_wait()` returns with the mutex *locked*

## "Synchronization Box"

If you have the mutex locked, you can do:





# Implementation Of Guarded Commands

```

when (guard) [
    /* command sequence */
    ...
]



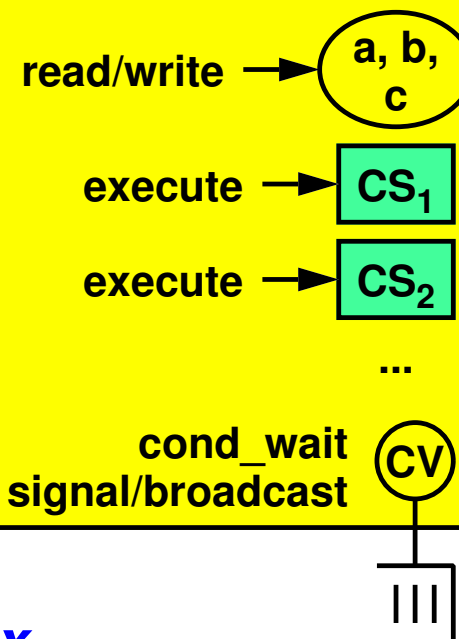
```

➡ POSIX provides *condition variables* for programmers to implement guarded commands

- *atomically* unlocks mutex and wait for the "event"
- ◆ with respect to the *operation of the mutex*

## "Synchronization Box"

If you have the mutex locked, you can do:



# Implementation Of Guarded Commands

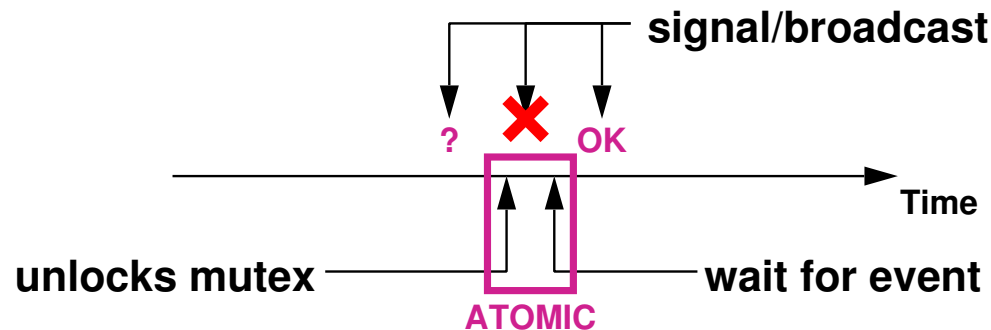
```

when (guard) [
    /* command sequence */
    ...
]

```


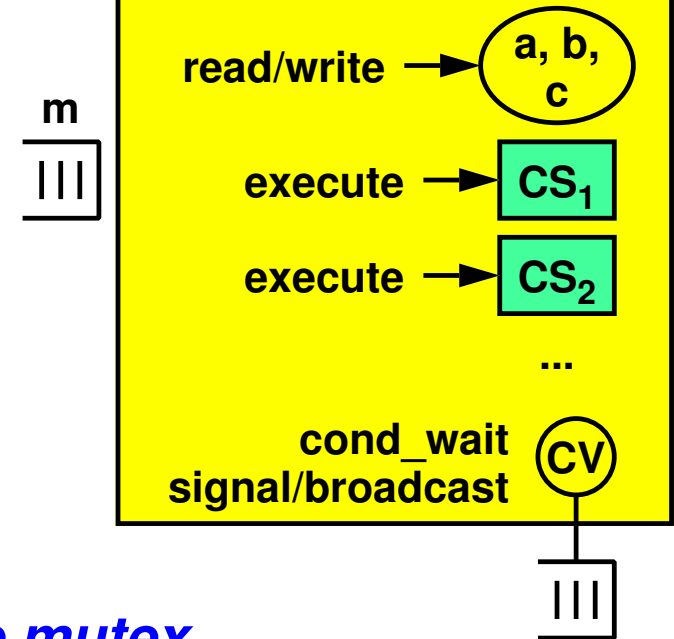
➡ POSIX provides *condition variables* for programmers to implement guarded commands

- *atomically* unlocks mutex and wait for the "event"
- ◆ with respect to the *operation of the mutex*



## "Synchronization Box"

If you have the mutex locked, you can do:



# Implementation Of Guarded Commands

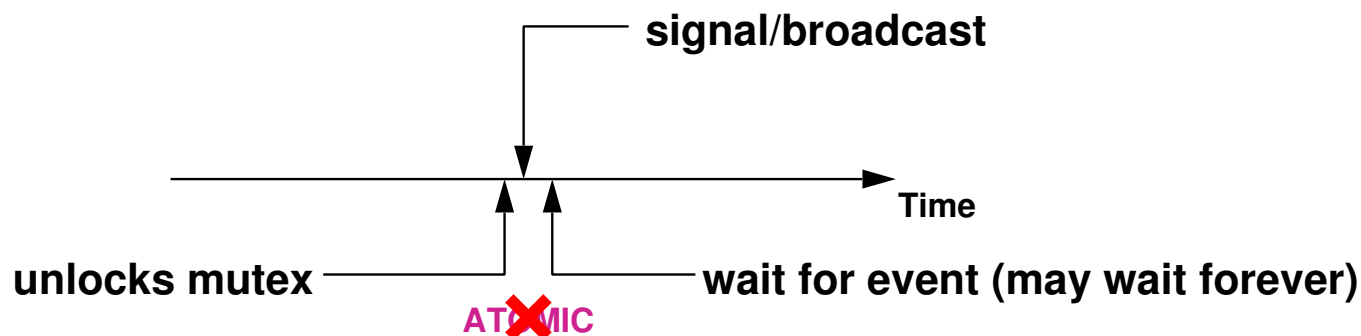
```

when (guard) [
    /* command sequence */
    ...
]

```


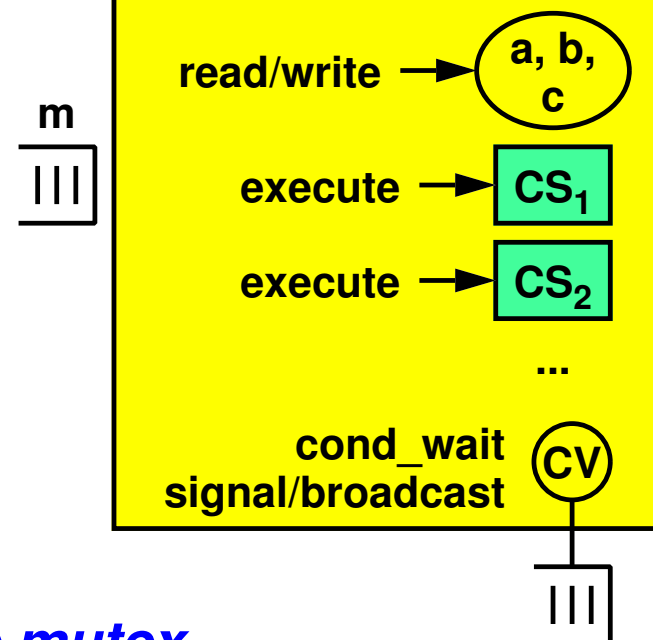
➡ POSIX provides *condition variables* for programmers to implement guarded commands

- *atomically* unlocks mutex and wait for the "event"
- ◆ with respect to the *operation of the mutex*



## "Synchronization Box"

If you have the mutex locked, you can do:



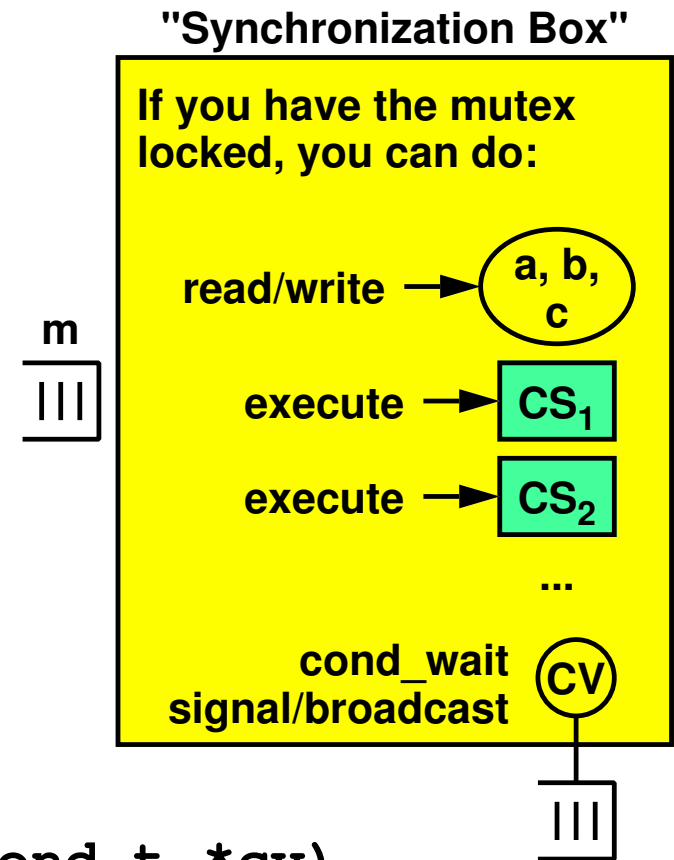
# Implementation Of Guarded Commands

```


when (guard) [
    /* command sequence */
    . . .
]

```

➡ POSIX provides *condition variables* for programmers to implement guarded commands



- 2) `pthread_cond_broadcast(pthread_cond_t *cv)`  
`pthread_cond_signal(pthread_cond_t *cv)`
  - should only call `pthread_cond_broadcast()` or `pthread_cond_signal()` if you have the corresponding mutex *locked*



# Implementation Of Guarded Commands

➡ **Synchronization:** mutex, condition variables, guards, critical sections


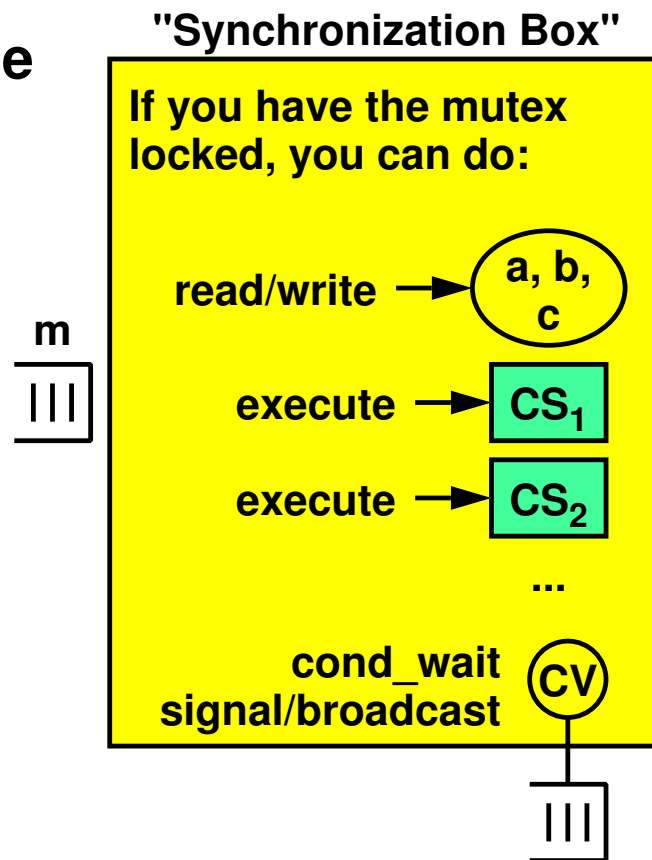
— with respect to a mutex, a thread can be

- waiting in the mutex queue
- got the lock and inside the "synchronization box"
  - ◆ only one thread can be inside the "synchronization box"

- waiting in the CV queue
- or outside

— with respect to a mutex, a, b, c are variables that can affect the value of the guard


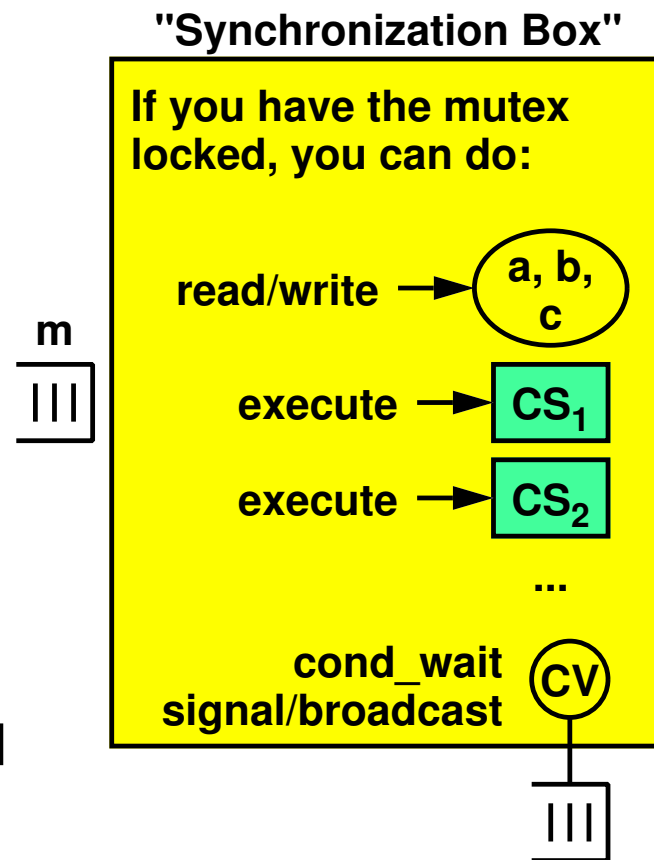
- can only access (i.e., read/write) them if a thread is inside the "synchronization box" (i.e., has the mutex locked)



# Implementation Of Guarded Commands

➡ **Synchronization:** mutex, condition variables, guards, critical sections


- when you **signal** CV
  - **one** thread in the CV queue gets moved to the mutex queue
- when you **broadcast** CV
  - **all** threads in the CV queue get moved to the mutex queue
- you can only get **added** to the CV queue if you have the mutex locked
- you can only **modify** the variables in the guard if you have the mutex locked
- you can only **read** the variables in the guard (i.e., evaluate the guard) if you have the mutex locked
- you can only **execute** critical section code if you have the mutex locked



# POSIX Condition Variables

<i>Guarded command</i>	<i>POSIX implementation</i>
<pre> when (guard) [     statement 1;     ...     statement n; ]</pre>	<pre> pthread_mutex_lock(&amp;mutex); while (!guard)     pthread_cond_wait(         &amp;cv,         &amp;mutex); statement 1; ... statement n; pthread_mutex_unlock(&amp;mutex);</pre>
<pre> [ /* code    * modifying    * the guard    */ ]</pre>	<pre> pthread_mutex_lock(&amp;mutex); /*code modifying the guard:*/ ... pthread_cond_broadcast(&amp;cv); pthread_mutex_unlock(&amp;mutex);</pre>


- if you don't follow these rules, your code will have **race conditions** (i.e., timing-dependent behavior)



# POSIX Condition Variables

<i>Guarded command</i>	<i>POSIX implementation</i>
<pre> when (guard) [     statement 1;     ...     statement n; ] </pre>	<pre> pthread_mutex_lock(&amp;mutex); while (!guard)     pthread_cond_wait(         &amp;cv,         &amp;mutex); statement 1; ... statement n; pthread_mutex_unlock(&amp;mutex); </pre>
<pre> [ /* code    * modifying    * the guard    */ ] </pre>	<pre> pthread_mutex_lock(&amp;mutex); /*code modifying the guard:*/ ... pthread_cond_broadcast(&amp;cv); pthread_mutex_unlock(&amp;mutex); </pre>

- don't believe that `pthread_cond_signal/broadcast()` can be called without locking the mutex





# Set Up


```
pthread_cond_t cv = PTHREAD_COND_INITIALIZER;
```

➡ If a condition variable cannot be initialized statically, do:


```
int pthread_cond_init(  
    pthread_cond_t *cvp,  
    pthread_condattr_t *attrp)
```

```
int pthread_cond_destroy(  
    pthread_cond_t *cvp)
```

➡ Usually, condition variable attributes are not used




# Readers-Writers Problem



# Readers-Writers Pseudocode

```
reader( ) {  
    when (writers == 0) [  
        readers++;  
    ]  
    /* read */  
    [readers--;]  
}
```

```
writer( ) {  
    when ((writers == 0) &&  
        (readers == 0)) [  
        writers++;  
    ]  
    /* write */  
    [writers--;]  
}
```



# Pseudocode with Assertions

```

reader( ) {
  when (writers == 0) [
    readers++;
  ]
  // sanity check
  assert((writers == 0) &&
        (readers > 0));
  /* read */
  [readers--;]
}


```

```

writer( ) {
  when ((writers == 0) &&
        (readers == 0)) [
    writers++;
  ]
  // sanity check
  assert((readers == 0) &&
        (writers == 1));
  /* write */
  [writers--;]
}

```

— the sanity checks are really not necessary



# Readers-Writers Pseudocode

```

reader( ) {
  when (writers == 0) [
    readers++;
  ]
  /* read */
  [readers--;]
}


```

```

writer( ) {
  when ((writers == 0) &&
        (readers == 0)) [
    writers++;
  ]
  /* write */
  [writers--;]
}

```

- = since `readers` is part of the guard in a guarded command, in the implementation of `[readers--;]`, you must signal/broadcast the corresponding condition used to implement that guard
  - in this case, only have to signal if `readers` becomes 0



# Readers-Writers Pseudocode

```

reader( ) {
  when (writers == 0) [
    readers++;
  ]
  /* read */
  [readers--;]
}


```

```

writer( ) {
  when ((writers == 0) &&
        (readers == 0)) [
    writers++;
  ]
  /* write */
  [writers--;]
}

```

- also, since `writers` is part of the guards in guarded commands (and these two guards are not identical), in the implementation of `[writers--;]`, you must signal/broadcast the corresponding conditions used to implement these guards



# Readers-Writers Pseudocode

```

reader( ) {
    when (writers == 0) [
        readers++;
    ]
    /* read */
    [readers--;]
}


```

```

writer( ) {
    when ((writers == 0) &&
        (readers == 0)) [
        writers++;
    ]
    /* write */
    [writers--;]
}

```

- don't have to worry about this readers
- don't have to worry about this writers
- you need to look at your program logic and figure when signal/broadcast conditions won't be useful
  - ◆ it's *not wrong* to signal/broadcast here, it's just *wasteful/inefficient*



# Solution with POSIX Threads

```

reader( ) {
    when (writers == 0) [
        readers++;
    ]
    /* read */
    [readers--;]
}



```

```

writer( ) {
    when ((writers == 0) &&
        (readers == 0)) [
        writers++;
    ]
    /* write */
    [writers--;]
}

```

- ➡ to be even more "efficient", can use multiple CVs
- ▬ so you don't have to wake up a thread unnecessarily
  - here we use one CV for reader's guard and one CV for writer's guard (since their expressions are different)





# Solution with POSIX Threads

```

reader( ) {
    pthread_mutex_lock(&m);
    while (!(writers == 0))
        pthread_cond_wait(
            &readersQ, &m);
    readers++;
    pthread_mutex_unlock(&m);
    /* read */
    pthread_mutex_lock(&m);
    if (--readers == 0)
        pthread_cond_signal(
            &writersQ);
    pthread_mutex_unlock(&m);
}


```

```

writer( ) {
    pthread_mutex_lock(&m);
    while (!(readers == 0) &&
        (writers == 0))
        pthread_cond_wait(
            &writersQ, &m);
    writers++;
    pthread_mutex_unlock(&m);
    /* write */
    pthread_mutex_lock(&m);
    writers--;
    pthread_cond_signal(
        &writersQ);
    pthread_cond_broadcast(
        &readersQ);
    pthread_mutex_unlock(&m);
}

```

— one mutex (m) and two condition variables (readersQ and writersQ)



# The Starvation Problem

```
reader( ) {
    when (writers == 0) [
        readers++;
    ]
    /* read */
    [readers--;]
}
```

```
writer( ) {
    when ((writers == 0) &&
        (readers == 0)) [
        writers++;
    ]
    /* write */
    [writers--;]
}
```




Can the writer never get a chance to write?

- yes, if there are always readers
- so, this implementation can be unfair to writers



Solution

- once a writer arrives, shut the door on new readers
  - `writers` now means the number of writers *wanting* to write
  - use `active_writers` to make sure that only one writer can do the actual writing at a time



# Solving The Starvation Problem

```

reader( ) {
    when (writers == 0) [
        readers++;
    ]
    /* read */
    [readers--;]
}

```


```

writer( ) {
    [writers++;]
    when ((readers == 0) &&
        (active_writers == 0))
    [
        active_writers++;
    ]
    /* write */
    [ writers--;
      active_writers--;
    ]
}

```

- now it's unfair to the readers
- isn't writing more important than reading anyway?


➡ This is an example of how to give threads priority *without* assigning priorities to threads!



# Improved Reader

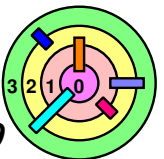
```
reader( ) {  
    pthread_mutex_lock(&m);  
    while (!(writers == 0))  
        pthread_cond_wait(  
            &readersQ, &m);  
    readers++;  
    pthread_mutex_unlock(&m);  
    /* read */  
    pthread_mutex_lock(&m);  
    if (--readers == 0)  
        pthread_cond_signal(  
            &writersQ);  
    pthread_mutex_unlock(&m);  
}
```

⇒ exactly the same as before!




# Improved Writer

```
writer( ) {  
    pthread_mutex_lock(&m);  
    writers++;  
    while (!((readers == 0) &&  
        (active_writers == 0))) {  
        pthread_cond_wait(&writersQ, &m);  
    }  
    active_writers++;  
    pthread_mutex_unlock(&m);  
    /* write */  
    pthread_mutex_lock(&m);  
    writers--;  
    active_writers--;  
    if (writers > 0)  
        pthread_cond_signal(&writersQ);  
    else  
        pthread_cond_broadcast(&readersQ);  
    pthread_mutex_unlock(&m);  
}
```




# New, From POSIX!


```
int pthread_rwlock_init(  
    pthread_rwlock_t *lock,  
    pthread_rwlockattr_t *att);  
int pthread_rwlock_destroy(  
    pthread_rwlock_t *lock);  
int pthread_rwlock_rdlock(  
    pthread_rwlock_t *lock);  
int pthread_rwlock_wrlock(  
    pthread_rwlock_t *lock);  
int pthread_rwlock_tryrdlock(  
    pthread_rwlock_t *lock);  
int pthread_rwlock_trywrlock(  
    pthread_rwlock_t *lock);  
int pthread_timedrwlock_rdlock(  
    pthread_rwlock_t *lock, struct timespec *ts);  
int pthread_timedrwlock_wrlock(  
    pthread_rwlock_t *lock, struct timespec *ts);  
int pthread_rwlock_unlock(  
    pthread_rwlock_t *lock);
```



# Barriers



- ➡ When a thread reaches a barrier, it must stop (do nothing) and simply wait for other threads to arrive at the same barrier
  - when all the threads that were suppose to arrive at the barrier have all arrived at the barrier, they are all given the signal to proceed forward
    - the barrier is then reset
- ➡ Ex: fork/join (fork to create parallel execution)



# A Solution?

```
int count = 0;
pthread_mutex_t m;
pthread_cond_t BarrierQueue;
void barrier_sync() {
    pthread_mutex_lock(&m);
    if (++count < n) {
        pthread_cond_wait(&BarrierQueue, &m);
    } else {
        count = 0;
        pthread_cond_broadcast(&BarrierQueue);
    }
    pthread_mutex_unlock(&m);
}
```

- ➡ the idea here is to have the last thread broadcast the condition while all the other threads are blocked at waiting for the condition to be signaled
- ➡ as it turns out, `pthread_cond_wait()` might return *spontaneously*, so this won't work

○ [http://pubs.opengroup.org/onlinepubs/009604599/functions/pthread\\_cond\\_signal.html](http://pubs.opengroup.org/onlinepubs/009604599/functions/pthread_cond_signal.html)




# A Solution?

```

int count = 0;
pthread_mutex_t m;
pthread_cond_t BarrierQueue;
void barrier_sync() {
    pthread_mutex_lock(&m);
    if (++count < n) {
        while (count < n)
            pthread_cond_wait(&BarrierQueue, &m);
    } else {
        pthread_cond_broadcast(&BarrierQueue);
        count = 0;
    }
    pthread_mutex_unlock(&m);
}

```

- if the  $n^{\text{th}}$  thread wakes up all the other blocked threads, most likely, *none* of these threads will see `count == n`




# A Solution?

```

int count = 0;
pthread_mutex_t m;
pthread_cond_t BarrierQueue;
void barrier_sync() {
    pthread_mutex_lock(&m);
    if (++count < n) {
        while (count < n)
            pthread_cond_wait(&BarrierQueue, &m);
    } else {
        pthread_cond_broadcast(&BarrierQueue);
    }
    pthread_mutex_unlock(&m);
    count = 0;
}

```


- ⇒ if the  $n^{\text{th}}$  thread wakes up all the other blocked threads, most likely, *none* of these threads will see `count == n`
- ⇒ moving `count = 0` around won't help
  - cannot guarantee all  $n$  threads will exit the barrier



# Barrier in POSIX Threads

```
int count = 0;
pthread_mutex_t m;
pthread_cond_t BarrierQueue;
void barrier_sync() {
    pthread_mutex_lock(&m);
    if (++count < number) {
        int my_generation = generation;
        while(my_generation == generation)
            pthread_cond_wait(&BarrierQueue, &m);
    } else {
        count = 0;
        generation++;
        pthread_cond_broadcast(&BarrierQueue);
    }
    pthread_mutex_unlock(&m);
}
```

- ≡ don't use count in the guard since its problematic!
- ≡ introduce a new guard



# More From POSIX!

```
int pthread_barrier_init(  
    pthread_barrier_t *barrier,  
    pthread_barrierattr_t *attr,  
    unsigned int count);  
int pthread_barrier_destroy(  
    pthread_barrier_t *barrier);  
int pthread_barrier_wait(  
    pthread_barrier_t *barrier);
```

