

# Ch 3: Basic Concepts

Bill Cheng

<http://merlot.usc.edu/cs402-s16>



Copyright © William C. Cheng



Copyright © William C. Cheng

- Procedures
- Threads & Coroutines
- Systems Calls
- Interrupts

## 3.1 Context Switching

Operating Systems - CSCI 402



Copyright © William C. Cheng


- What's the execution context of a thread?
  - if we are going to talk about context switching, we need to know what we are switching and how to get back
- The *execution context* of a thread is the *current state* of our thread
  - what does the execution context include?
    - CPU registers, including the *instruction pointer*, *stack pointer*, *base/frame pointer*, etc.
    - stack
    - open files
    - etc.
    - i.e., things that may affect the execution of the thread
  - turns out the stack is complicated
    - in reality, it's just the *current stack frame* of the current thread
    - what's below it (and the rest of the address space) is also part of the thread's state

### Context

Operating Systems - CSCI 402



Copyright © William C. Cheng



- What's Next?
  - So far, we have talked about abstractions
    - processes, files, threads
    - stuff at the user level
  - We are not ready to talk about the OS yet
  - Next step is something in between



Copyright © William C. Cheng



- Context Switching
  - The magic of OS
    - to provide the illusion that applications run concurrently and each application thinks it's the only application running on the processor
  - The OS switches the processor from one application to another
    - switching happens transparently to the applications
  - What is the OS doing when an application is running?

Application1   Application2   Application3

Operating Systems - CSCI 402



Copyright © William C. Cheng

- Procedures
- Threads & Coroutines
- Systems Calls
- Interrupts

## 3.1 Context Switching

Operating Systems - CSCI 402

Copyright © William C. Cheng

### Intel x86 (32-Bit): Subroutine Linkage

Who sets what?

- caller** is explicitly setup by the **callee**
- args** is copied explicitly by the **caller** into the stack frame by a "call" machine instruction
- ebp** is copied explicitly by the **callee**
- registers are saved explicitly by the **callee**
- as it turned out, for x86, some registers are designated to be saved by the callee code
- space for local variables is **created** explicitly by the **callee** code as well as initialization of these variables

What does the stack frame look like for the following function?

```
void func() { printf("I'm here.\n"); }
```

Operating Systems - CSCI 402

Copyright © William C. Cheng

### Intel x86 (32-Bit): Subroutine Linkage

In reality, there can be stuff between stack frames

- e.g., by convention, specific registers are saved and restored by the caller (this can depend on the compiler)

Operating Systems - CSCI 402

Copyright © William C. Cheng

### Intel x86 (32-Bit): Subroutine Linkage

- esp** points to the end of the current stack frame
- it is used to prepare the next stack frame
- ebp** contains the caller's **instruction pointer**
- this is the **return address**!

Please be reminded that our address space is up-side-down (comparing against the textbook)

Operating Systems - CSCI 402

Copyright © William C. Cheng

### Intel x86 (32-Bit): Subroutine Linkage

- ebp** contains the caller's **base (frame) pointer** register
- this is a link to the caller's **stack frame**
- eax** contains the return value of a function
- some fields are not always present, compiler decides

Please be reminded that our address space is up-side-down (comparing against the textbook)

Operating Systems - CSCI 402

Copyright © William C. Cheng

### Subroutines

```
int main() {
    int x, y;
    // computers x^y
    int result = 1;
    for (i=0; i<y; i++)
        result *= x;
    return (0);
}

int sub(int x, int y) {
    // computers x^y
    int result = 1;
    for (i=0; i<y; i++)
        result *= x;
    return (result);
}
```

- You are in **main()** and are ready to call **sub()**
- how do you make sure that **sub()** has the right context to execute the code in **sub()**?
- you need to prepare the context for **sub()**
- how do you make sure that you can return from **sub()** and restore the **main()** context and continue to execute properly?
- you need to first **save** the context of **main()**
- after return from **sub()**, you need to **restore** the context
- of **main()** so **main()** can **resume** execution

Operating Systems - CSCI 402

Copyright © William C. Cheng

### Subroutines

- The context of **main()** includes CPU registers, any global variables (none here) and its local variables, **i** and **a**
- The context of **sub()** includes its arguments, **x** and **y**
- any global variables, none here
- its local variables, **i** and **result**
- its arguments, **x** and **y**

Global variables are in fixed location in the address space

Local variables and arguments are in **current stack frame**

Operating Systems - CSCI 402

Copyright © William C. Cheng

17

main: pushl %ebp  
movl %esp, %ebp  
pushl %esi  
pushl %edi  
addl \$8, %esp  
movl \$0, %eax  
popl %edi  
popl %esi  
movl %ebp, %esp  
popl %ebp  
ret

set up stack frame  
pushl %ebp  
movl %esp, %ebp  
pushl %edi  
pushl %esi  
subl \$8, %esp

push args  
movl -12(%ebp), %eax  
pushl %eax  
call sub  
addl \$8, %esp  
movl %eax, -16(%ebp)

get result  
i = sub(a, 1);

pop args  
int i;  
int a;

set return value and restore frame  
addl \$8, %esp  
movl \$0, %eax  
popl %edi  
popl %esi  
movl %ebp, %esp  
popl %ebp  
ret

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng

15

main: pushl %ebp  
movl %esp, %ebp  
pushl %esi  
pushl %edi  
subl \$8, %esp

set up stack frame  
pushl %ebp  
movl %esp, %ebp  
pushl %edi  
pushl %esi  
subl \$8, %esp

push args  
movl -12(%ebp), %eax  
pushl %eax  
call sub  
addl \$8, %esp  
movl %eax, -16(%ebp)

get result  
i = sub(a, 1);

pop args  
int i;  
int a;

set return value and restore frame  
addl \$8, %esp  
movl \$0, %eax  
popl %edi  
popl %esi  
movl %ebp, %esp  
popl %ebp  
ret

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng

13

main: pushl %ebp  
movl %esp, %ebp  
pushl %esi  
pushl %edi  
subl \$8, %esp

set up stack frame  
pushl %ebp  
movl %esp, %ebp  
pushl %edi  
pushl %esi  
subl \$8, %esp

push args  
movl -12(%ebp), %eax  
pushl %eax  
call sub  
addl \$8, %esp  
movl %eax, -16(%ebp)

get result  
i = sub(a, 1);

pop args  
int i;  
int a;

set return value and restore frame  
addl \$8, %esp  
movl \$0, %eax  
popl %edi  
popl %esi  
movl %ebp, %esp  
popl %ebp  
ret

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng

16

main: pushl %ebp  
movl %esp, %ebp  
pushl %esi  
pushl %edi  
subl \$8, %esp

set up stack frame  
pushl %ebp  
movl %esp, %ebp  
pushl %edi  
pushl %esi  
subl \$8, %esp

push args  
movl -12(%ebp), %eax  
pushl %eax  
call sub  
addl \$8, %esp  
movl %eax, -16(%ebp)

get result  
i = sub(a, 1);

pop args  
int i;  
int a;

set return value and restore frame  
addl \$8, %esp  
movl \$0, %eax  
popl %edi  
popl %esi  
movl %ebp, %esp  
popl %ebp  
ret

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng

16

main: pushl %ebp  
movl %esp, %ebp  
pushl %esi  
pushl %edi  
subl \$8, %esp

set up stack frame  
pushl %ebp  
movl %esp, %ebp  
pushl %edi  
pushl %esi  
subl \$8, %esp

push args  
movl -12(%ebp), %eax  
pushl %eax  
call sub  
addl \$8, %esp  
movl %eax, -16(%ebp)

get result  
i = sub(a, 1);

pop args  
int i;  
int a;

set return value and restore frame  
addl \$8, %esp  
movl \$0, %eax  
popl %edi  
popl %esi  
movl %ebp, %esp  
popl %ebp  
ret

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng

14

main: pushl %ebp  
movl %esp, %ebp  
pushl %esi  
pushl %edi  
subl \$8, %esp

set up stack frame  
pushl %ebp  
movl %esp, %ebp  
pushl %edi  
pushl %esi  
subl \$8, %esp

push args  
movl -12(%ebp), %eax  
pushl %eax  
call sub  
addl \$8, %esp  
movl %eax, -16(%ebp)

get result  
i = sub(a, 1);

pop args  
int i;  
int a;

set return value and restore frame  
addl \$8, %esp  
movl \$0, %eax  
popl %edi  
popl %esi  
movl %ebp, %esp  
popl %ebp  
ret

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng

23

```

main:
    pushl %ebp
    movl %esp, %ebp
    pushl %esi
    pushl %edi
    addl $8, %esp
    movl %eax, -16(%ebp)
    ...
    pop args;
    call sub
    int i;
    int a;
    i = sub(a, 1);
    return(0);
}
    
```

push args:  $\text{movl } -12(\%ebp), \%eax$   
 stack frame:  $\text{subl } \$8, \%esp$   
 set up:  $\text{pushl } \%esp$ ,  $\text{pushl } \%edi$ ,  $\text{pushl } \%ebp$   
 restore frame:  $\text{popl } \%edi$ ,  $\text{popl } \%esi$ ,  $\text{movl } \%ebp, \%esp$   
 set return:  $\text{ret}$

Intel x86: Subroutine Code (1)

Copyright © William C. Cheng

24

```

main:
    pushl %ebp
    movl %esp, %ebp
    pushl %esi
    pushl %edi
    addl $8, %esp
    movl %eax, -16(%ebp)
    ...
    pop args;
    call sub
    int i;
    int a;
    i = sub(a, 1);
    return(0);
}
    
```

push args:  $\text{movl } -12(\%ebp), \%eax$   
 stack frame:  $\text{subl } \$8, \%esp$   
 set up:  $\text{pushl } \%esp$ ,  $\text{pushl } \%edi$ ,  $\text{pushl } \%ebp$   
 restore frame:  $\text{popl } \%edi$ ,  $\text{popl } \%esi$ ,  $\text{movl } \%ebp, \%esp$   
 set return:  $\text{ret}$

Intel x86: Subroutine Code (1)

Copyright © William C. Cheng

21

```

main:
    pushl %ebp
    movl %esp, %ebp
    pushl %esi
    pushl %edi
    addl $8, %esp
    movl %eax, -12(%ebp), %eax
    ...
    pop args;
    call sub
    int i;
    int a;
    i = sub(a, 1);
    return(0);
}
    
```

push args:  $\text{movl } -12(\%ebp), \%eax$   
 stack frame:  $\text{subl } \$8, \%esp$   
 set up:  $\text{pushl } \%esp$ ,  $\text{pushl } \%edi$ ,  $\text{pushl } \%ebp$   
 restore frame:  $\text{popl } \%edi$ ,  $\text{popl } \%esi$ ,  $\text{movl } \%ebp, \%esp$   
 set return:  $\text{ret}$

Intel x86: Subroutine Code (1)

Copyright © William C. Cheng

22

```

main:
    pushl %ebp
    movl %esp, %ebp
    pushl %esi
    pushl %edi
    addl $8, %esp
    movl %eax, -12(%ebp), %eax
    ...
    pop args;
    call sub
    int i;
    int a;
    i = sub(a, 1);
    return(0);
}
    
```

push args:  $\text{movl } -12(\%ebp), \%eax$   
 stack frame:  $\text{subl } \$8, \%esp$   
 set up:  $\text{pushl } \%esp$ ,  $\text{pushl } \%edi$ ,  $\text{pushl } \%ebp$   
 restore frame:  $\text{popl } \%edi$ ,  $\text{popl } \%esi$ ,  $\text{movl } \%ebp, \%esp$   
 set return:  $\text{ret}$

Intel x86: Subroutine Code (1)

Copyright © William C. Cheng

19

```

main:
    pushl %ebp
    movl %esp, %ebp
    pushl %esi
    pushl %edi
    addl $8, %esp
    movl %eax, -16(%ebp)
    ...
    pop args;
    call sub
    int i;
    int a;
    i = sub(a, 1);
    return(0);
}
    
```

push args:  $\text{movl } -12(\%ebp), \%eax$   
 stack frame:  $\text{subl } \$8, \%esp$   
 set up:  $\text{pushl } \%esp$ ,  $\text{pushl } \%edi$ ,  $\text{pushl } \%ebp$   
 restore frame:  $\text{popl } \%edi$ ,  $\text{popl } \%esi$ ,  $\text{movl } \%ebp, \%esp$   
 set return:  $\text{ret}$

Intel x86: Subroutine Code (1)

Copyright © William C. Cheng

20

```

main:
    pushl %ebp
    movl %esp, %ebp
    pushl %esi
    pushl %edi
    addl $8, %esp
    movl %eax, -16(%ebp)
    ...
    pop args;
    call sub
    int i;
    int a;
    i = sub(a, 1);
    return(0);
}
    
```

push args:  $\text{movl } -12(\%ebp), \%eax$   
 stack frame:  $\text{subl } \$8, \%esp$   
 set up:  $\text{pushl } \%esp$ ,  $\text{pushl } \%edi$ ,  $\text{pushl } \%ebp$   
 restore frame:  $\text{popl } \%edi$ ,  $\text{popl } \%esi$ ,  $\text{movl } \%ebp, \%esp$   
 set return:  $\text{ret}$

Intel x86: Subroutine Code (1)

Copyright © William C. Cheng

29

main:

```

pushl %ebp
movl %esp, %ebp
pushl %esi
pushl %edi
pushl %eax
movl -12(%ebp), %eax
pushl %eax
call sub
addl $8, %esp
movl %eax, -16(%ebp)
get result: {
    ...
    i = sub(a, 1);
    return(0);
}

```

push args: {  
pushl %eax  
movl -12(%ebp), %eax  
}

pop args: {  
addl \$8, %esp  
movl %eax, -16(%ebp)  
}

get result: {  
...  
i = sub(a, 1);  
return(0);  
}

set return  
value and  
restore frame

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng

27

main:

```

pushl %ebp
movl %esp, %ebp
pushl %esi
pushl %edi
pushl %eax
movl -12(%ebp), %eax
pushl %eax
call sub
addl $8, %esp
movl %eax, -16(%ebp)
get result: {
    ...
    i = sub(a, 1);
    return(0);
}

```

push args: {  
pushl %eax  
movl -12(%ebp), %eax  
}

pop args: {  
addl \$8, %esp  
movl %eax, -16(%ebp)  
}

get result: {  
...  
i = sub(a, 1);  
return(0);  
}

set return  
value and  
restore frame

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng

25

main:

```

pushl %ebp
movl %esp, %ebp
pushl %esi
pushl %edi
pushl %eax
movl -12(%ebp), %eax
pushl %eax
call sub
addl $8, %esp
movl %eax, -16(%ebp)
get result: {
    ...
    i = sub(a, 1);
    return(0);
}

```

push args: {  
pushl %eax  
movl -12(%ebp), %eax  
}

pop args: {  
addl \$8, %esp  
movl %eax, -16(%ebp)  
}

get result: {  
...  
i = sub(a, 1);  
return(0);  
}

set return  
value and  
restore frame

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng

30

main:

```

pushl %ebp
movl %esp, %ebp
pushl %esp, %ebp
pushl %edi
pushl %esi
subl $8, %esp
pushl $1
movl -12(%ebp), %eax
pushl %eax
call sub
addl $8, %esp
movl %eax, -16(%ebp)
get result: {
    ...
    i = sub(a, 1);
    return(0);
}

```

push args: {  
pushl %eax  
movl -12(%ebp), %eax  
}

pop args: {  
addl \$8, %esp  
movl %eax, -16(%ebp)  
}

get result: {  
...  
i = sub(a, 1);  
return(0);  
}

set return  
value and  
restore frame

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng

28

main:

```

pushl %ebp
movl %esp, %ebp
pushl %esp, %ebp
pushl %edi
pushl %esi
subl $8, %esp
pushl $1
movl -12(%ebp), %eax
pushl %eax
call sub
addl $8, %esp
movl %eax, -16(%ebp)
get result: {
    ...
    i = sub(a, 1);
    return(0);
}

```

push args: {  
pushl %eax  
movl -12(%ebp), %eax  
}

pop args: {  
addl \$8, %esp  
movl %eax, -16(%ebp)  
}

get result: {  
...  
i = sub(a, 1);  
return(0);  
}

set return  
value and  
restore frame

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng

26

main:

```

pushl %ebp
movl %esp, %ebp
pushl %esp, %ebp
pushl %edi
pushl %esi
subl $8, %esp
pushl $1
movl -12(%ebp), %eax
pushl %eax
call sub
addl $8, %esp
movl %eax, -16(%ebp)
get result: {
    ...
    i = sub(a, 1);
    return(0);
}

```

push args: {  
pushl %eax  
movl -12(%ebp), %eax  
}

pop args: {  
addl \$8, %esp  
movl %eax, -16(%ebp)  
}

get result: {  
...  
i = sub(a, 1);  
return(0);  
}

set return  
value and  
restore frame

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng

35

```

sub:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl $1, -4(%ebp)
    movl $0, -8(%ebp)
    movl -4(%ebp), %ecx
    movl -8(%ebp), %eax
    jge endloop
    cmpi 12(%ebp), %eax
    jg endloop
    // computers x^y
    int i;
    for (i=0; i<y; i++)
        result = x;
    return(result);
}
    
```

init locals

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng

36

```

sub:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl $1, -4(%ebp)
    movl $0, -8(%ebp)
    movl -4(%ebp), %ecx
    movl -8(%ebp), %eax
    jge endloop
    cmpi 12(%ebp), %eax
    jg endloop
    // computers x^y
    int i;
    for (i=0; i<y; i++)
        result = x;
    return(result);
}
    
```

init locals

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng

33

```

sub:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl $1, -4(%ebp)
    movl $0, -8(%ebp)
    movl -4(%ebp), %ecx
    movl -8(%ebp), %eax
    jge endloop
    cmpi 12(%ebp), %eax
    jg endloop
    // computers x^y
    int i;
    for (i=0; i<y; i++)
        result = x;
    return(result);
}
    
```

init locals

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng

34

```

sub:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl $1, -4(%ebp)
    movl $0, -8(%ebp)
    movl -4(%ebp), %ecx
    movl -8(%ebp), %eax
    jge endloop
    cmpi 12(%ebp), %eax
    jg endloop
    // computers x^y
    int i;
    for (i=0; i<y; i++)
        result = x;
    return(result);
}
    
```

init locals

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng

31

```

main:
    pushl %ebp
    movl %esp, %ebp
    pushl %edi
    pushl %esi
    subl $8, %esp
    ...
    call sub
    addl $8, %esp
    movl %eax, -16(%ebp)
    pushl %eax
    int i;
    int a;
    get result:
        i = sub(a, 1);
    ...
    addl $8, %esp
    movl $0, %eax
    popl %edi
    popl %esi
    movl %ebp, %esp
    ret
    
```

set up stack frame

push args

get result

set return value and restore frame

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng

32

```

main:
    pushl %ebp
    movl %esp, %ebp
    pushl %edi
    pushl %esi
    subl $8, %esp
    ...
    call sub
    addl $8, %esp
    movl %eax, -16(%ebp)
    pushl %eax
    int i;
    int a;
    get result:
        i = sub(a, 1);
    ...
    addl $8, %esp
    movl $0, %eax
    popl %edi
    popl %esi
    movl %ebp, %esp
    ret
    
```

set up stack frame

push args

get result

set return value and restore frame

Intel x86: Subroutine Code (1)

Operating Systems - CSCI 402

Copyright © William C. Cheng 41

```

ret
popl %ebp
movl %ebp, %esp
movl -4(%ebp), %eax
movl %ecx, -4(%ebp)
endloop:
jmp begloop
addl $1, %eax
jge endloop
imull 8(%ebp), %ecx
cmpl 12(%ebp), %eax
begloop:
movl -8(%ebp), %eax
movl -4(%ebp), %ecx
movl $0, -8(%ebp)
movl $1, -4(%ebp)
subl $8, %esp
movl %esp, %ebp
pushl %ebp

```

init locals

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng 39

```

ret
popl %ebp
movl %ebp, %esp
movl -4(%ebp), %eax
movl %ecx, -4(%ebp)
endloop:
jmp begloop
addl $1, %eax
jge endloop
imull 8(%ebp), %ecx
cmpl 12(%ebp), %eax
begloop:
movl -8(%ebp), %eax
movl -4(%ebp), %ecx
movl $0, -8(%ebp)
movl $1, -4(%ebp)
subl $8, %esp
movl %esp, %ebp
pushl %ebp

```

init locals

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng 37

```

ret
popl %ebp
movl %ebp, %esp
movl -4(%ebp), %eax
movl %ecx, -4(%ebp)
endloop:
jmp begloop
addl $1, %eax
jge endloop
imull 8(%ebp), %ecx
cmpl 12(%ebp), %eax
begloop:
movl -8(%ebp), %eax
movl -4(%ebp), %ecx
movl $0, -8(%ebp)
movl $1, -4(%ebp)
subl $8, %esp
movl %esp, %ebp
pushl %ebp

```

init locals

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng 42

```

ret
popl %ebp
movl %ebp, %esp
movl -4(%ebp), %eax
movl %ecx, -4(%ebp)
endloop:
jmp begloop
addl $1, %eax
jge endloop
imull 8(%ebp), %ecx
cmpl 12(%ebp), %eax
begloop:
movl -8(%ebp), %eax
movl -4(%ebp), %ecx
movl $0, -8(%ebp)
movl $1, -4(%ebp)
subl $8, %esp
movl %esp, %ebp
pushl %ebp

```

init locals

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng 40

```

ret
popl %ebp
movl %ebp, %esp
movl -4(%ebp), %eax
movl %ecx, -4(%ebp)
endloop:
jmp begloop
addl $1, %eax
jge endloop
imull 8(%ebp), %ecx
cmpl 12(%ebp), %eax
begloop:
movl -8(%ebp), %eax
movl -4(%ebp), %ecx
movl $0, -8(%ebp)
movl $1, -4(%ebp)
subl $8, %esp
movl %esp, %ebp
pushl %ebp

```

init locals

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng 38

```

ret
popl %ebp
movl %ebp, %esp
movl -4(%ebp), %eax
movl %ecx, -4(%ebp)
endloop:
jmp begloop
addl $1, %eax
jge endloop
imull 8(%ebp), %ecx
cmpl 12(%ebp), %eax
begloop:
movl -8(%ebp), %eax
movl -4(%ebp), %ecx
movl $0, -8(%ebp)
movl $1, -4(%ebp)
subl $8, %esp
movl %esp, %ebp
pushl %ebp

```

init locals

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng

47

```

sub:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl $1, -4(%ebp)
    movl $0, -8(%ebp)
    movl -4(%ebp), %ecx
    movl -8(%ebp), %eax
    jge endloop
    cmpi 12(%ebp), %eax
    jg endloop
    jmp beginloop
beginloop:
    addl $1, %eax
    imull 8(%ebp), %ecx
    int 1;
    for (i=0; i<y; i++)
        result *= x;
    return(result);
}
    
```

init locals

args
ebp
ebp
saved registers (empty)
local variables (8 bytes)

esp → ebp →

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng

46

```

sub:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl $1, -4(%ebp)
    movl $0, -8(%ebp)
    movl -4(%ebp), %ecx
    movl -8(%ebp), %eax
    jge endloop
    cmpi 12(%ebp), %eax
    jg endloop
    jmp beginloop
beginloop:
    addl $1, %eax
    imull 8(%ebp), %ecx
    int 1;
    for (i=0; i<y; i++)
        result *= x;
    return(result);
}
    
```

init locals

args
ebp
ebp
saved registers (empty)
local variables (8 bytes)

esp → ebp →

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng

45

```

sub:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl $1, -4(%ebp)
    movl $0, -8(%ebp)
    movl -4(%ebp), %ecx
    movl -8(%ebp), %eax
    jge endloop
    cmpi 12(%ebp), %eax
    jg endloop
    jmp beginloop
beginloop:
    addl $1, %eax
    imull 8(%ebp), %ecx
    int 1;
    for (i=0; i<y; i++)
        result *= x;
    return(result);
}
    
```

init locals

args
ebp
ebp
saved registers (empty)
local variables (8 bytes)

esp → ebp →

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng

46

```

sub:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl $1, -4(%ebp)
    movl $0, -8(%ebp)
    movl -4(%ebp), %ecx
    movl -8(%ebp), %eax
    jge endloop
    cmpi 12(%ebp), %eax
    jg endloop
    jmp beginloop
beginloop:
    addl $1, %eax
    imull 8(%ebp), %ecx
    int 1;
    for (i=0; i<y; i++)
        result *= x;
    return(result);
}
    
```

init locals

args
ebp
ebp
saved registers (empty)
local variables (8 bytes)

esp → ebp →

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng

43

```

sub:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl $1, -4(%ebp)
    movl $0, -8(%ebp)
    movl -4(%ebp), %ecx
    movl -8(%ebp), %eax
    jge endloop
    cmpi 12(%ebp), %eax
    jg endloop
    jmp beginloop
beginloop:
    addl $1, %eax
    imull 8(%ebp), %ecx
    int 1;
    for (i=0; i<y; i++)
        result *= x;
    return(result);
}
    
```

init locals

args
ebp
ebp
saved registers (empty)
local variables (8 bytes)

esp → ebp →

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402

Copyright © William C. Cheng

44

```

sub:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl $1, -4(%ebp)
    movl $0, -8(%ebp)
    movl -4(%ebp), %ecx
    movl -8(%ebp), %eax
    jge endloop
    cmpi 12(%ebp), %eax
    jg endloop
    jmp beginloop
beginloop:
    addl $1, %eax
    imull 8(%ebp), %ecx
    int 1;
    for (i=0; i<y; i++)
        result *= x;
    return(result);
}
    
```

init locals

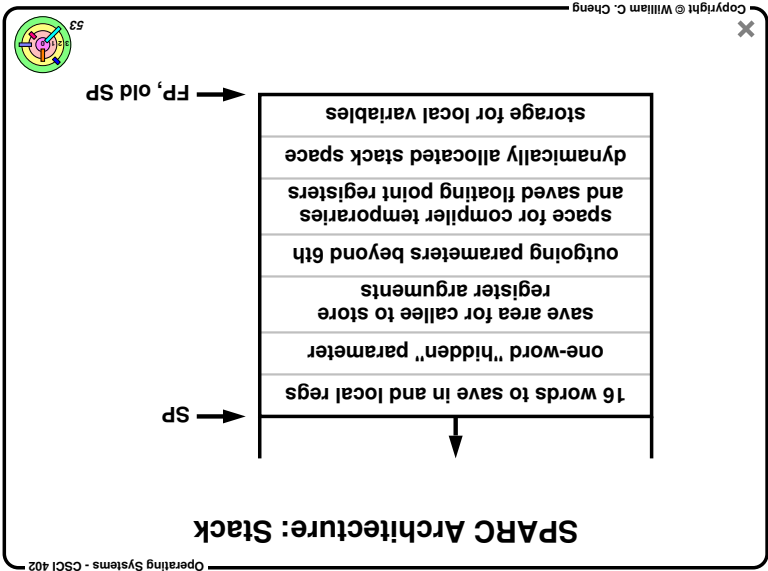
args
ebp
ebp
saved registers (empty)
local variables (8 bytes)

esp → ebp →

Intel x86: Subroutine Code (2)

Operating Systems - CSCI 402





Copyright © William C. Cheng

54

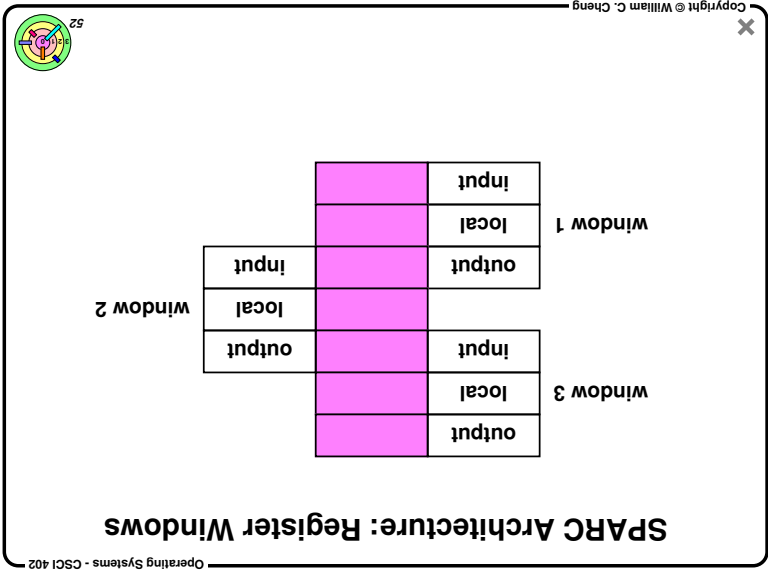
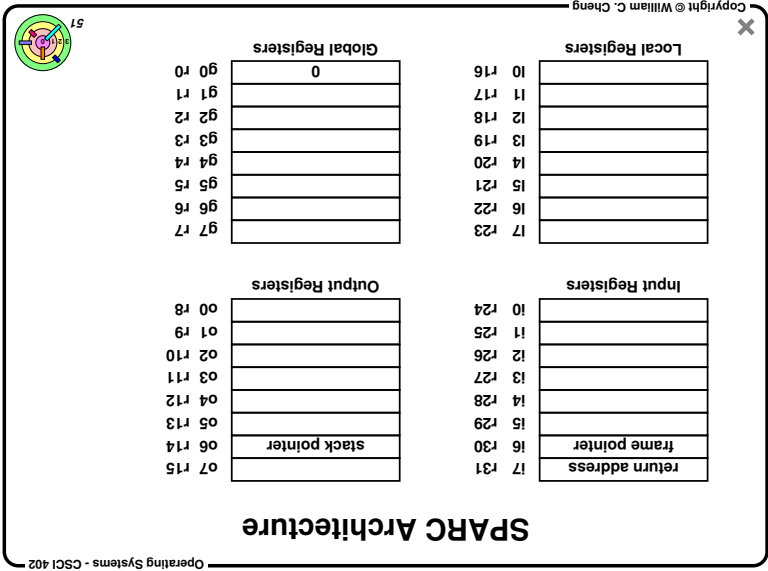
### SPARC Architecture: Subroutine Code

```

ld [%fp-8], %o0
! put local var (a) into out register
mov l, %o1
! deal with 2nd parameter
call sub
nop
st %o0, [%fp-4]
! store result into local var (1)
...
sub:
save %sp, -64, %sp
! push a new stack frame
add %i0, %i1, %i0
! compute sum
ret
! return to caller
restore
! pop frame off stack (in delay slot)

```

Operating Systems - CSCI 402



Copyright © William C. Cheng

49

### Intel x86: Subroutine Code (2)

```

sub:
pushl %ebp
movl %esp, %ebp
subl $8, %esp
movl $1, %eax
jge endloop
cmpl 12(%ebp), %eax
beginloop:
movl -8(%ebp), %eax
movl -4(%ebp), %ecx
movl -8(%ebp), %ecx
movl $0, %eax
movl -4(%ebp), %ecx
movl -8(%ebp), %ecx
movl $1, %eax
jmp beginloop
endloop:
movl %ecx, -4(%ebp)
movl %ebp, %esp
popl %ebp
ret

```

init locals

args

ebp

esp

Operating Systems - CSCI 402

Copyright © William C. Cheng

50

### Intel x86: Subroutine Code (2)

```

sub:
pushl %ebp
movl %esp, %ebp
subl $8, %esp
movl $1, %eax
jge endloop
cmpl 12(%ebp), %eax
beginloop:
movl -8(%ebp), %eax
movl -4(%ebp), %ecx
movl -8(%ebp), %ecx
movl $0, %eax
movl -4(%ebp), %ecx
movl -8(%ebp), %ecx
movl $1, %eax
jmp beginloop
endloop:
movl %ecx, -4(%ebp)
movl %ebp, %esp
popl %ebp
ret

```

init locals

args

ebp

esp

Operating Systems - CSCI 402

# 3.1 Context Switching


- Procedures
- Threads & Coroutines
- Systems Calls
- Interrupts



Copyright © William C. Cheng



- normally, threads are independent of one another and don't directly control one another's execution
- threads can be made aware of each other and be able to transfer control from one thread to another
- this is known as *coroutine linkage*



Representing Threads




Copyright © William C. Cheng



Copyright © William C. Cheng

- A thread's context
  - its stack, its register state
  - can be stored in a *thread control block* (directly or indirectly)
  - To transfer control from one thread to another is equivalent to copying the thread control block of the target thread into the current thread context



Representing Threads



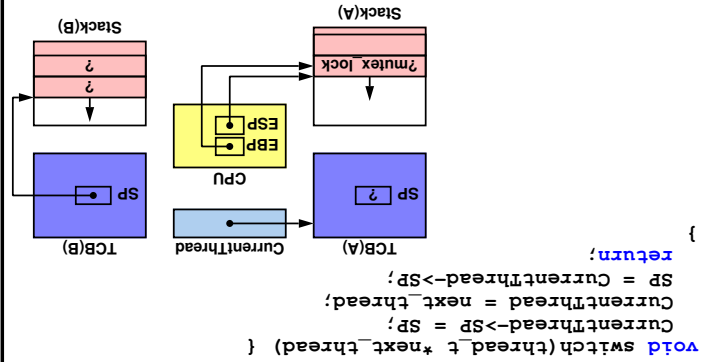
Copyright © William C. Cheng



- thread A calls switch() to switch to thread B
- thread B's TCB has the *exact context* of thread B right before thread B was *last suspended*
- context information in thread A's TCB is *out-dated*



Copyright © William C. Cheng



Switching Between Threads



Copyright © William C. Cheng



Copyright © William C. Cheng



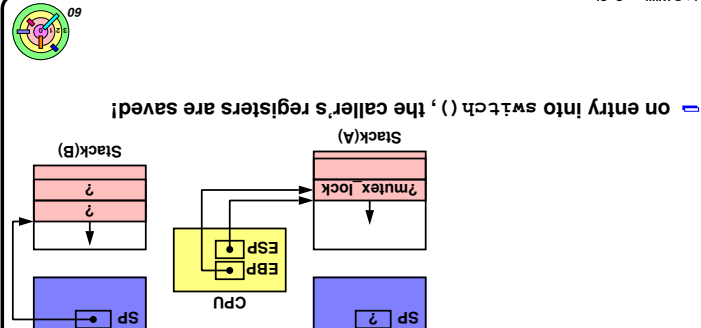
```
switch:
    ;enter switch, creating new stack frame
    pushl %ebp; set fp to point to new frame
    movl %esp, %ebp; save esp register
    pushl %eax; save eax register
    movl CurrentThread, %eax; load address of caller's TCB
    movl CurrentThread, %eax; store target TCB address
    movl 8(%ebp), %eax; put new TCB address into esp
    movl CurrentThread, %eax; restore target thread's SP
    movl SP(%eax), %eax; restore target thread's esp register
    ;we're now in the context of the target thread!
    popl %eax; restore target thread's esp register
    popl %ebp; pop target thread's fp
    ret; return to caller within target thread
```

```
void switch(thread_t *next_thread) {
    CurrentThread->SP = SP;
    CurrentThread = next_thread;
    SP = CurrentThread->SP;
    return;
}
```

Switching Between Threads



Copyright © William C. Cheng



```
void switch(thread_t *next_thread) {
    CurrentThread->SP = SP;
    CurrentThread = next_thread;
    SP = CurrentThread->SP;
    return;
}
```

Switching Between Threads



Copyright © William C. Cheng



Copyright © William C. Cheng

65

the thread control block of the target thread is copied into

the current thread context

```

    }
    SP = CurrentThread->SP;
    CurrentThread = next_thread;
    CurrentThread->SP = SP;
    void switch(thread_t *next_thread) {

```

Switching Between Threads

Operating Systems - CSCI 402

Copyright © William C. Cheng

66

fetch the target thread's stack pointer (esp for x86) from its thread control block and loads it into the actual stack pointer

which thread executes this?

does it matter?

```

    }
    SP = CurrentThread->SP;
    CurrentThread = next_thread;
    CurrentThread->SP = SP;
    void switch(thread_t *next_thread) {

```

Switching Between Threads

Operating Systems - CSCI 402

Copyright © William C. Cheng

63

then the current stack pointer is saved into current thread's thread control block

```

    }
    SP = CurrentThread->SP;
    CurrentThread = next_thread;
    CurrentThread->SP = SP;
    void switch(thread_t *next_thread) {

```

Switching Between Threads

Operating Systems - CSCI 402

Copyright © William C. Cheng

64

the thread control block of the target thread is copied into

the current thread context

```

    }
    SP = CurrentThread->SP;
    CurrentThread = next_thread;
    CurrentThread->SP = SP;
    void switch(thread_t *next_thread) {

```

Switching Between Threads

Operating Systems - CSCI 402

Copyright © William C. Cheng

61

on entry into switch(), the caller's registers are saved!

```

    }
    SP = CurrentThread->SP;
    CurrentThread = next_thread;
    CurrentThread->SP = SP;
    void switch(thread_t *next_thread) {

```

Switching Between Threads

Operating Systems - CSCI 402

Copyright © William C. Cheng

62

then the current stack pointer is saved into current thread's thread control block

```

    }
    SP = CurrentThread->SP;
    CurrentThread = next_thread;
    CurrentThread->SP = SP;
    void switch(thread_t *next_thread) {

```

Switching Between Threads

Operating Systems - CSCI 402

Copyright © William C. Cheng

71

### Switching Between Threads

```

void switch(thread_t *next_thread) {
    CurrentThread->SP = SP;
    CurrentThread = next_thread;
    return;
}

```

Note: one very interesting thing happened in this call

- usually, a single thread executes the entire procedure call
- with `switch()`, at the beginning of the procedure call, one thread is executing
- half way through the procedure call, another thread starts to execute
- so, one thread enters the `switch()` call, and a different thread leaves the `switch()` call!

This is an elegant way of switching threads

- all threads come here to switch to another thread

Copyright © William C. Cheng

Copyright © William C. Cheng

72

### ... in x86 Assembler

```

switch:
    !enter switch, creating new thread
    pushl %ebp; pushl %sp;
    movl %esp, %ebp; setl %sp;
    pushl %esi; save %esi register;
    movl CurrentThread, %esi;
    movl %esp, SP(%esi); save %esp;
    movl 8(%ebp), CurrentThread; store target TCB address
    !into CurrentThread;
    movl CurrentThread, %esi; put new TCB address into %esi
    !we're now in the context of the target thread!
    popl %esi; restore target thread's %esi register
    popl %ebp; pop target thread's %ebp
    ret; return to caller within target thread

```

Copyright © William C. Cheng

Copyright © William C. Cheng

69

### Switching Between Threads

```

void switch(thread_t *next_thread) {
    CurrentThread->SP = SP;
    CurrentThread = next_thread;
    return;
}

```

- on return from `switch()`, the registers (`ebp` and `eip` for `x86`) are restored into the current thread, which is the target thread!
- which thread executes `return`?

Copyright © William C. Cheng

Copyright © William C. Cheng

70

### Switching Between Threads

```

void switch(thread_t *next_thread) {
    CurrentThread->SP = SP;
    CurrentThread = next_thread;
    return;
}

```

- if thread control blocks were user-space data structures, threads were switched *without* getting the kernel involved!
- Note: `SP` field inside `TCB(B)` no longer tracks `ESP` in CPU

Copyright © William C. Cheng

Copyright © William C. Cheng

67

### Switching Between Threads

```

void switch(thread_t *next_thread) {
    CurrentThread->SP = SP;
    CurrentThread = next_thread;
    return;
}

```

- fetch the target thread's stack pointer (`esp` for `x86`) from its thread control block and loads it into the actual stack pointer
- hmm... which thread executes this?
- both? either? `EIP`?

Copyright © William C. Cheng

Copyright © William C. Cheng

68

### Switching Between Threads

```

void switch(thread_t *next_thread) {
    CurrentThread->SP = SP;
    CurrentThread = next_thread;
    return;
}

```

- on return from `switch()`, the registers (`ebp` and `eip` for `x86`) are restored into the current thread, which is the target thread!
- which thread executes `return`?

Copyright © William C. Cheng

Copyright © William C. Cheng

77

**System Calls**

➤ A *trap* is a type of "*software interrupt*"

➤ interrupt handler will invoke trap handler

```

int_handler(int_code) {
    ...
    syscall_handler(trap_code) {
        ...
        if (int_code == SYSCTL)
            syscall_handler();
        ...
        if (trap_code == write_code)
            write_handler();
    }
}

```

**User**

```

prog() {
    ...
    write(fd, buffer, size);
    ...
    trap(write_code);
    ...
}

```

**Kernel**

```

write_handler() frame
syscall_handler() frame
int_handler() frame

```

**User Stack**

```

write() frame
prog() frame

```

**Kernel Stack**

Operating Systems - CSCI 402

Copyright © William C. Cheng

78

**System Calls**

➤ More details on the "*trap*" machine instruction

- 1) Trap into the kernel with all *interrupt disabled* and processor mode set to *kernel mode*
- 2) The *Hardware Abstraction Layer (HAL)* save IP and SP in "temporary locations" in kernel space (e.g., the *interrupt stack*)
- additional registers may be saved
- *HAL* is *hardware-dependent* (outside the scope of this class) for the corresponding user process (information from PCB)
- 4) HAL sets IP to *interrupt handler* (written in C)
- pop user IP and SP from "temporary location" and push them onto kernel stack, then *re-enable interrupt*
- 5) On return from the trap handler, disable interrupt and executes a special "return" instruction to *return to user process*

➤ Similar sequence happens when you get *hardware interrupt*

➤ iret on x86

Operating Systems - CSCI 402

Copyright © William C. Cheng

75

**System Calls**

➤ A system call involves the transfer of control from user code to system/kernel code and back

➤ there is no thread switching!

➤ depending on the OS implementation, this can view this as a user thread *change status* and becomes a kernel thread

➤ and executes in privileged mode

➤ and executing operating-system code

➤ effectively, it's part of the OS

➤ in reality, more complex than just changing status

➤ then it changed back to a user thread

Operating Systems - CSCI 402

Copyright © William C. Cheng

76

**System Calls**

➤ Most systems provide threads with two stacks

➤ one for use in user mode

➤ and one for use in kernel mode

➤ in some systems, one kernel stack is shared by all threads in the same user process

➤ therefore, when a thread performs a system call and switches from user mode to kernel mode

➤ it switches to use its kernel-mode stack

➤ the kernel cannot use the user-space stack because it cannot trust the user process

Operating Systems - CSCI 402

Copyright © William C. Cheng

73

**... in SPARC Assembler**

```

switch:
    save %sp, -64, %sp
    t 3
    ! Trap into the OS to force
    ! window overflow.
    ! Save CurrentThread's SP in
    ! control block.
    mov %i0, %g0
    ! Set CurrentThread to be
    ! target thread.
    ! Set SP to that of target thread
    ! return to caller (in target
    ! thread's context).
    ! Pop frame off stack (in delay
    ! slot).
    restore

```

Operating Systems - CSCI 402

Copyright © William C. Cheng

74

**3.1 Context Switching**

➤ Procedures

➤ Threads & Coroutines

➤ Systems Calls

➤ Interrupts

Operating Systems - CSCI 402

Copyright © William C. Cheng

## Interrupts

- Interrupt context needs a stack
- which stack should it use?
- there are several possibilities
- 1) allocate a new stack each time an interrupt occurs
- too slow
- 2) have one stack shared by all interrupt handlers
- not often done
- 3) interrupt handler could borrow a stack from the thread it is interrupting
- most common

83

Operating Systems - CSCI 402

Copyright © William C. Cheng

## Interrupts

- Do not confuse **interrupts** with **signals** (even though the terminologies related to them are similar)
- signals** are generated by the kernel
- they are delivered to the **user process**
- signal** ≠ **software interrupt**
- interrupts** are generated by the hardware
- they are delivered to the **kernel**
- they are delivered to the HAL and then the kernel
- When an **interrupt** occurs, the processor puts aside the current context and switch to an **interrupt context**
- the current context can be a **thread context** or **another interrupt context**
- when the interrupt handler is finishes, the processor generally resumes the original context

81

Operating Systems - CSCI 402

Copyright © William C. Cheng

## Context Switch

- The big idea here is that in order to perform a context switch, you must first save your context
- therefore, you must know what constitutes the context
- then you save all of it
- what's the **minimum** amount of context to save?
- context can be stored in several places
- stack
- thread control block (e.g., in a system call, the TCB contains pointers to **both** the corresponding user stack frame and the kernel stack frame)
- etc.
- when switching back, you must restore the context
- In general, it's difficult to make a "clean" context switch
- when you switch from context A to context B
- there may be time you are in the context of both A and B
- there may be time you are in neither contexts

79

Operating Systems - CSCI 402

Copyright © William C. Cheng

## Currently Executing User Thread

84

Operating Systems - CSCI 402

Copyright © William C. Cheng

## Interrupting A User Thread

- If interrupt occurs when a **user thread** is executing in the CPU
- Disable interrupt** and set processor mode to **kernel mode**
- The **Hardware Abstraction Layer (HAL)** save IP and SP in "temporary locations" in kernel space (e.g., the **interrupt stack**)
- additional registers may be saved
- HAL** is **hardware-dependent** (outside the scope of this class)
- HAL sets the SP to point to the **kernel stack** designated for the corresponding user process (information from PCB)
- HAL sets IP to **interrupt handler** (written in C)
- pop user IP and SP from "temporary location" and push them onto kernel stack, then **re-enable interrupt**
- On return from the trap handler, disable interrupt and executes a special "return" instruction to **return to user process**
- let on x86

What about interrupting a **kernel thread** or an **interrupt service routine**?

82

Operating Systems - CSCI 402

Copyright © William C. Cheng

## 3.1 Context Switching

- Procedures
- Threads & Coroutines
- Systems Calls
- Interrupts

80

Operating Systems - CSCI 402

Copyright © William C. Cheng

89

## Interrupt Mask

- Interrupt can be **masked**, i.e., temporarily blocked
- if an interrupt occurs while it is masked, the interrupt indication remains **pending**
- once it is unmasked, the processor is interrupted
- How interrupts are masked is architecture-dependent
- common approaches
  - 1) hardware register implements a **bit vector / mask**
  - if a particular bit is set, the corresponding interrupt class is enable (or disabled)
  - the kernel masks interrupts by setting bits in the register
  - when an interrupt does occur, the corresponding mask bit is set in the register (block other interrupts of the same class)
  - cleared when the handler returns
  - 2) hierarchical interrupt levels (more common)

Copyright © William C. Cheng

87

## Interrupts

- For approaches (2) and (3), there is no way to suspend one interrupt handler and **resume** the execution of another
- since there is only **one stack** for all the interrupt handlers
- therefore, the handler of the most recent interrupt must **run to completion**
- when it's done, the stack frame is removed, and the next-most-recent interrupt now must run to completion
- this is a **big deal!**
- once you have interrupt handlers running, a normal thread (no matter how important it is) cannot run until **all** interrupt handlers complete
- this is why an interrupt service routine should do as little as possible (and figure out a way to do the rest later)
- if we have approach (1), then we won't have this problem

Copyright © William C. Cheng

85

## Currently Executing Kernel Thread

Copyright © William C. Cheng

90

## Interrupt Mask

- Interrupt can be **masked**, i.e., temporarily blocked
- if an interrupt occurs while it is masked, the interrupt indication remains **pending**
- once it is unmasked, the processor is interrupted
- How interrupts are masked is architecture-dependent
- common approaches
  - 1) hardware register implements a **bit vector / mask**
  - 2) hierarchical interrupt levels (more common)
  - the processor masks interrupts by setting an **Interrupt Priority Level (IPL)** in a hardware register
  - all interrupts with the current or lower levels are masked
  - the kernel masks a class of interrupts by setting the IPL to a particular value
  - when an interrupt does occur, the current IPL is set to that of the level the interrupt belongs
  - restores to previous value on handler return

Copyright © William C. Cheng

88

## Interrupts

- What if an interrupt service routine takes too long to run?
  - be done on a queue of some sort, then arranges for it to be done in some other context at a later time
  - still need to do *something* in the interrupt handler
  - 1) **unblock a kernel thread** that's sleeping in the corresponding I/O queue
  - 2) **start the next I/O operation** on the same device
  - this approach is used in many systems, including Windows and Linux
  - will discuss further in Ch 5

Copyright © William C. Cheng

86

## Currently Executing Another Interrupt Service Routine

## 3.2 Input/Output Architectures

- Architectural concerns
  - memory-mapped I/O
  - programmed I/O (PIO)
  - direct memory access (DMA)
- Software concerns
  - device drivers
  - concurrency of I/O and computation

### Input/Output





91

Copyright © William C. Cheng

Operating Systems - CSCI 402

### What Does A Computer Look Like?



- LSI-11
- processor for PDP-11
- Boards are connected over a "bus"
- various standards for PDP-11
- Unibus, Q-Bus, etc.
- connect to backplane bus

<http://hampage.hu/pdp-11/lsi11.html>




93

Copyright © William C. Cheng

Operating Systems - CSCI 402

### Simple I/O Architecture

- memory-mapped I/O
- PIO (programmed I/O)
- perform I/O operations by reading or writing data in the controller registers one byte or word at a time over the bus




95

Copyright © William C. Cheng

### Simple I/O Architecture

- memory-mapped I/O
- two categories of devices
- PIO (programmed I/O)
- perform I/O operations by reading or writing data in the controller registers one byte or word at a time over the bus




96

Copyright © William C. Cheng

Operating Systems - CSCI 402

### Simple I/O Architecture

- memory-mapped I/O
- all controllers listen on the bus to determine if a request is for itself or not
- memory controller behaves differently from other controllers, i.e., it passes the bus request to primary memory others "process" the bus request
- and respond to relatively few addresses
- memory is not really a "device"



94

Copyright © William C. Cheng

Operating Systems - CSCI 402



Copyright © William C. Cheng 101

Simple I/O Architecture

- memory-mapped I/O
- two categories of devices
- PIO (programmed I/O)
- DMA (direct memory access)
- the controller performs the I/O itself
- the processor writes to the controller to tell it where to transfer the results to
- the controller takes over and transfers data between itself and primary memory

Operating Systems - CSCI 402

Copyright © William C. Cheng 99

Simple I/O Architecture

- memory-mapped I/O
- two categories of devices
- PIO (programmed I/O)
- perform I/O operations by reading or writing data in the bus
- controller registers one byte or word at a time over the bus

Operating Systems - CSCI 402

Copyright © William C. Cheng 97

Simple I/O Architecture

- memory-mapped I/O
- two categories of devices
- PIO (programmed I/O)
- perform I/O operations by reading or writing data in the bus
- controller registers one byte or word at a time over the bus

Operating Systems - CSCI 402

Copyright © William C. Cheng 102

Simple I/O Architecture

- memory-mapped I/O
- two categories of devices
- PIO (programmed I/O)
- DMA (direct memory access)
- the controller performs the I/O itself
- the processor writes to the controller to tell it where to transfer the results to
- the controller takes over and transfers data between itself and primary memory

Operating Systems - CSCI 402

Copyright © William C. Cheng 100

Simple I/O Architecture

- memory-mapped I/O
- two categories of devices
- PIO (programmed I/O)
- DMA (direct memory access)
- the controller performs the I/O itself
- the processor writes to the controller to tell it where to transfer the results to
- the controller takes over and transfers data between itself and primary memory

Operating Systems - CSCI 402

Copyright © William C. Cheng 96

Simple I/O Architecture

- memory-mapped I/O
- two categories of devices
- PIO (programmed I/O)
- perform I/O operations by reading or writing data in the bus
- controller registers one byte or word at a time over the bus

Operating Systems - CSCI 402

- ## Direct Memory Access

- ## Device Drivers



- ## Programmed I/O

- ## DMA Registers

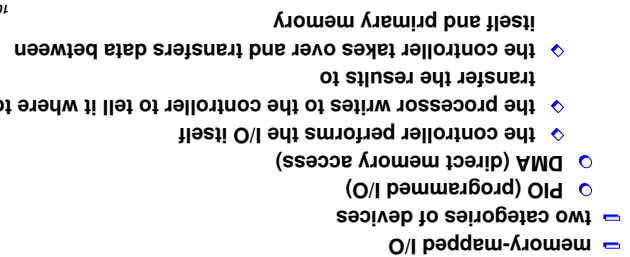
■ Operating Systems - CSCI 402



Operating Systems - CSCI 402

- ## Simple I/O Architecture

Operating Systems - CSCI 402



- memory-mapped I/O
- two categories of devices
  - PIO (programmed I/O)
  - DMA (direct memory access)
- ◆ the controller performs the I/O itself
- ◆ the processor writes to the controller to tell it where to transfer the results to
- ◆ the controller takes over and transfers data between itself and primary memory

Copyright © William C. Cheng

111

Even in Sixth-Edition Unix, the internal driver interface is often asynchronous

- start\_read/start\_write () returns a handle identifying the operation that has started
- a thread can call the wait () method to synchronously wait for I/O completion
- it's possible for multiple threads to invoke wait () with the same handle, if they all want the same block from a file

```

class disk {
public:
    virtual handle_t start_read(request_t) = 0;
    virtual handle_t start_write(request_t) = 0;
    virtual status_t wait(handle_t) = 0;
    virtual status_t interrupt() = 0;
};

```

A Bit More Realistic

Operating Systems - CSCI 402

Copyright © William C. Cheng

109

Each type of disk driver is a **subclass** of the disk class and has its own implementation of these functions

- each disk driver looks like a generic disk to the OS
- this gets compiled into an array of function pointers (which is what C++ code gets compiled into)
- in reality, there are no object classes and no polymorphism
- the CPU doesn't even know about data structures
- the CPU only knows about memory addresses and how to execute machine instructions

```

class disk {
public:
    virtual status_t read(request_t) = 0;
    virtual status_t write(request_t) = 0;
    virtual status_t interrupt() = 0;
};

```

... in C++

Operating Systems - CSCI 402

Copyright © William C. Cheng

112

When I/O costs dominate computation costs

- use I/O processors (a.k.a. channels) to handle much of the I/O work
- important in large data-processing applications
- can even download program into a channel

```

graph LR
    Mem[Memory] <--> C1[Channel]
    Mem <--> C2[Channel]
    Mem <--> C3[Channel]
    C1 <--> Cont1[Controller]
    C2 <--> Cont2[Controller]
    C3 <--> Cont3[Controller]
    Mem <--> HW[Hardware Chip]

```

I/O Processors: Channels

Operating Systems - CSCI 402

Copyright © William C. Cheng

110

This is a synchronous interface

- a user thread would call the read/write () method
- this starts the device and the user thread would block
- the device driver's interrupt method is called in the interrupt context
- if I/O is completed, the thread is unblocked and return from the read/write () method

```

class disk {
public:
    virtual status_t read(request_t) = 0;
    virtual status_t write(request_t) = 0;
    virtual status_t interrupt() = 0;
};

```

... in C++

Operating Systems - CSCI 402