


Dynamic Storage Allocation


➡ Where in the kernel do you need to do memory allocation?

- ▬ stack space
- ▬ `malloc()`
- ▬ `fork()`
- ▬ various OS data structures
 - process control block
 - thread control block
 - mutex (it's a queue)
- ▬ etc.




Dynamic Storage Allocation

- ➡ Goal: allow dynamic creation and destruction of data structures
- ➡ Concerns:
 - efficient use of storage
 - efficient use of processor time
- ➡ Example:
 - *first-fit* vs. *best-fit* allocation



Allocation Example



1300 bytes free (first)

1200 bytes free (last)


Allocate 1000 bytes:

First Fit


Best Fit

Allocate 1100 bytes:

Allocate 250 bytes:



Allocation Example




1300 bytes free (first)

1200 bytes free (last)

Allocate 1000 bytes:


First Fit

Best Fit




Allocate 1100 bytes:

Allocate 250 bytes:



Allocation Example




1300 bytes free (first)

1200 bytes free (last)


Allocate 1000 bytes:

First Fit


Best Fit




Allocate 1100 bytes:



Allocate 250 bytes:



Allocation Example




1300 bytes free (first)

1200 bytes free (last)


Allocate 1000 bytes:

First Fit



Best Fit




Allocate 1100 bytes:



Allocate 250 bytes:



Allocation Example




1300 bytes free (first)

1200 bytes free (last)

Allocate 1000 bytes:


First Fit



300

1200

Best Fit



1300

200


Allocate 1100 bytes:



300


100

Allocate 250 bytes:




50

100



Allocation Example




1300 bytes free (first)

1200 bytes free (last)

Allocate 1000 bytes:


First Fit



300

1200


Best Fit



1300

200

Allocate 1100 bytes:



300


100



200


200

Allocate 250 bytes:




50

100



Allocation Example




1300 bytes free (first)

1200 bytes free (last)

Allocate 1000 bytes:


First Fit



300

1200

Best Fit



1300

200

Allocate 1100 bytes:



300


100



200

200

Allocate 250 bytes:



50


100



Stuck!

200

200




Fragmentation




First-fit vs. ***best-fit*** allocation

- studies have shown that first-fit works better
- best-fit tends to leave behind a large number of regions of memory that are too small to be useful
 - best-fit tends to create smallest left-over blocks!
- this is the general problem of ***fragmentation***
 - ***internal fragmentation***: unusable memory is contained within an allocated region (e.g., buddy system)
 - ***external fragmentation***: unusable memory is separated into small blocks and is interspersed by allocated memory (e.g., best-fit)




Implementing First Fit: Data Structures




- a linked list of *free blocks*
 - don't need to manage allocated blocks
- use a doubly-linked list
 - insertion and deletion are fast, i.e., $O(1)$, once you know where to insert or delete

Liberation of Storage



This is known as *coalescing*


- in order to make coalescing possible, you need to know that *size* of the blocks above and below the block being freed
- you also need to know if they are *allocated* or *free*



Boundary Tags



Allocated Block




Free Block



This is known as *coalescing*

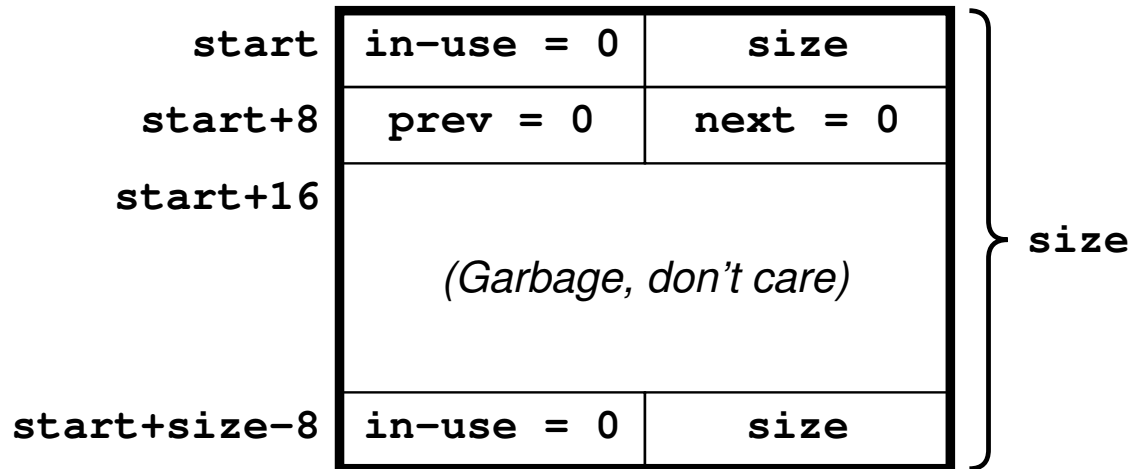
- in order to make coalescing possible, you need to know that *size* of the blocks above and below the block being freed
- you also need to know if they are *allocated* or *free*



Detailed Examples

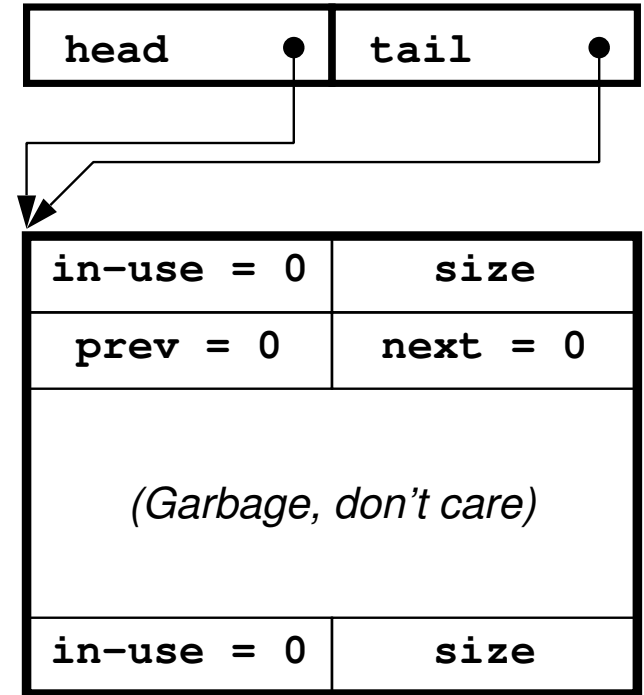


Free block


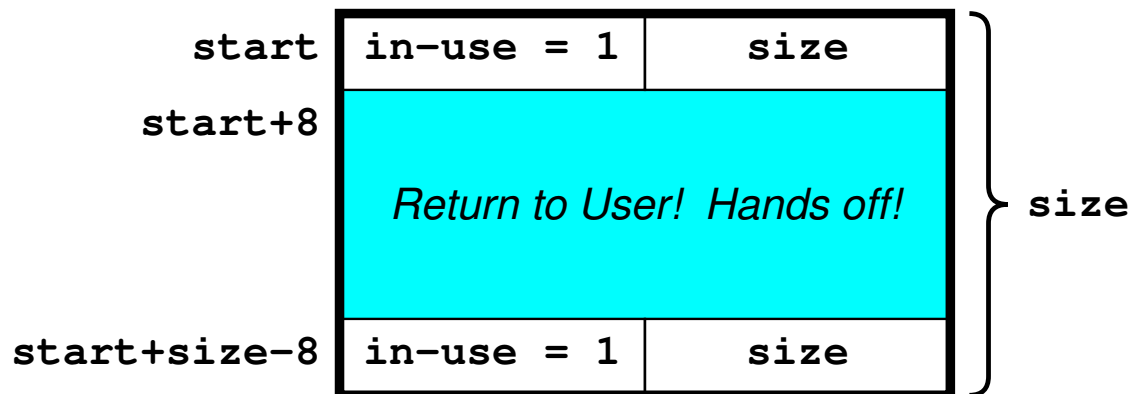


Free list

Free List

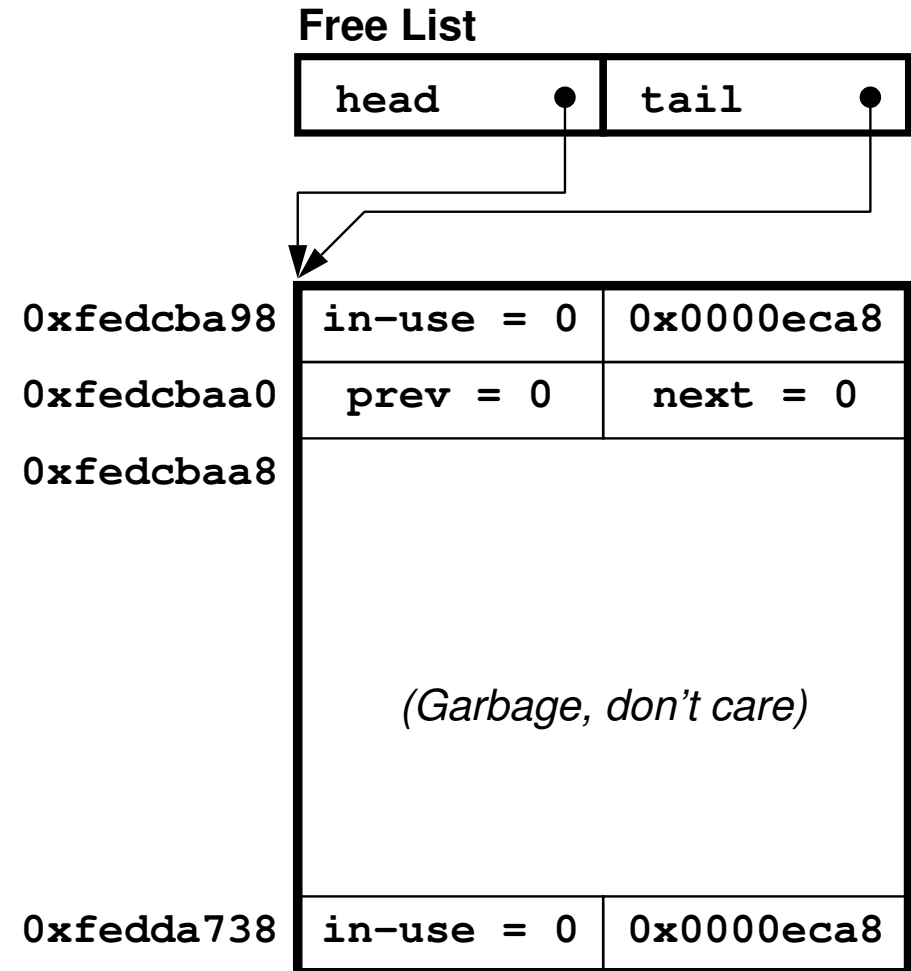


In-use block



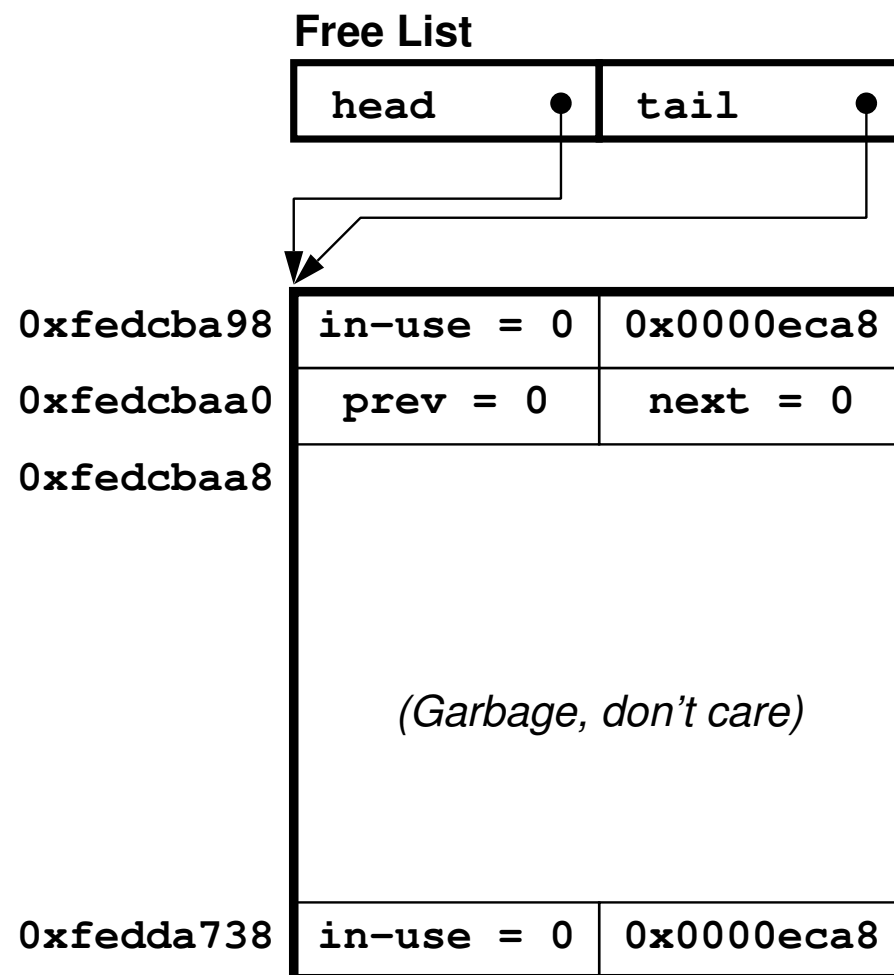
malloc() Example

- ➡ Ex: Heap starts at 0xfedcba98 and size of the heap is 0x0000eca8 (60,584) bytes
- the Free List contains one free block and it looks like this:




malloc() Example

- ➡ Ex: Heap starts at 0xfedcba98 and size of the heap is 0x0000eca8 (60,584) bytes
- the Free List contains one free block and it looks like this:




- ➡ Ex: Request block size is 100
- split the block into two
 - busy block size is 116
 - remaining free block size is $60584 - 116 = 60468 = 0xec34$



malloc() Example

- ➡ Ex: Heap starts at 0xfedcba98 and size of the heap is 0x0000eca8 (60,584) bytes
- the Free List contains one free block and it looks like this:

Free List



return

0xfedcba98

0xfedcbba0

...

0xfedcbb04

0xfedcbb0c

0xfedcbb14


0xfedcbb1c

...

0xfedda738



in-use = 1	0x00000074
Return to user! Hands off!	
in-use = 1	0x00000074
in-use = 0	0x0000ec34
prev = 0	next = 0
(Garbage, don't care)	
in-use = 0	0x0000ec34

- ➡ Ex: Request block size is 100
- split the block into two
 - busy block size is 116
 - remaining free block size is $60584 - 116 = 60468 = 0xec34$





free () Example

- ➡ After K blocks of memory have been allocated (and assume that none of them have been deallocated)
- in the memory layout, the first K blocks are used block, followed by one free block



free () Example

- ➡ After **K** blocks of memory have been allocated (and assume that none of them have been deallocated)
 - ▬ in the memory layout, the first **K** blocks are used block, followed by one free block
- ➡ Memory blocks can be freed in any order
 - ▬ when a memory block is freed, we need to check if the blocks before it and after it are also free
- ➡ If neither of them are free, we just need to insert the newly freed block into the Free List (at the right place)
 - ▬ need to **search** the Free List to find insertion point
 - ▬ searching through a linear list is "slow", **$O(n)$**
- ➡ Otherwise, we can **merge/coalesce** the block in question with neighboring free block(s)



free () Example



Ex: free (Y)

- Y-16 tells you if the *previous* block is free or not
- Y-8+Z tells you if the *next* block is free or not
 - where Z is what's in Y-4

Y-8 - (* (Y-12))

Y-16

Y-8

Y


Y-8+Z

in-use=?	size
?	
in-use=?	size
in-use=1	size=Z
Return to user! Hands off!	
in-use=1	size=Z
in-use=?	size
?	
in-use=?	size




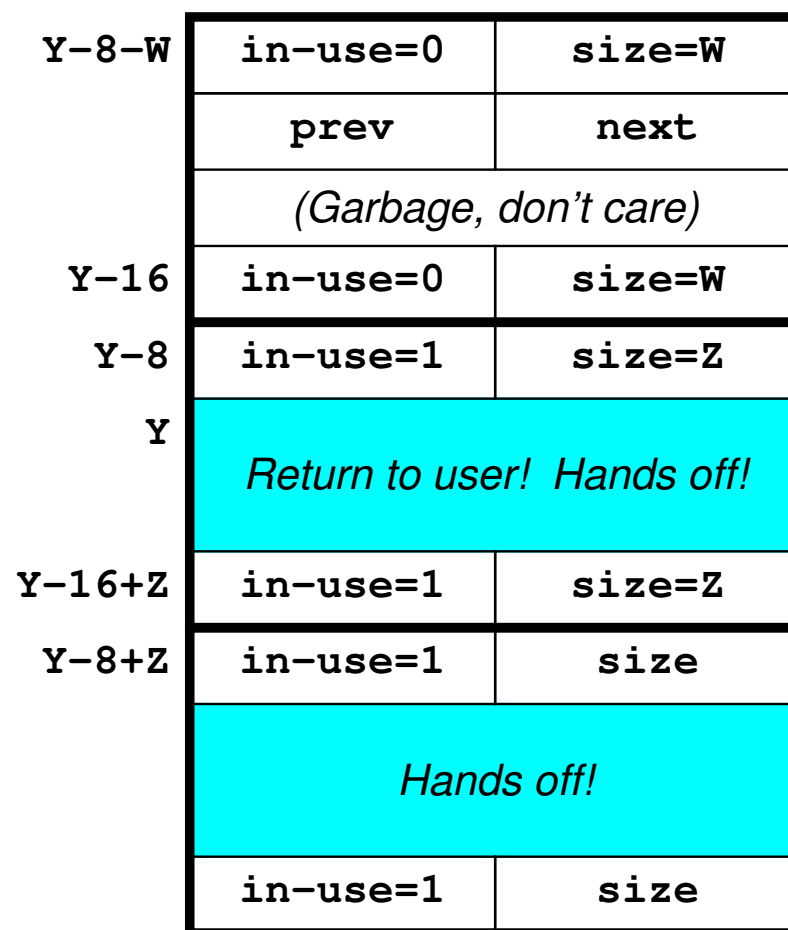
Coalescing:

- need to make sure that everything is consistent



free () Example


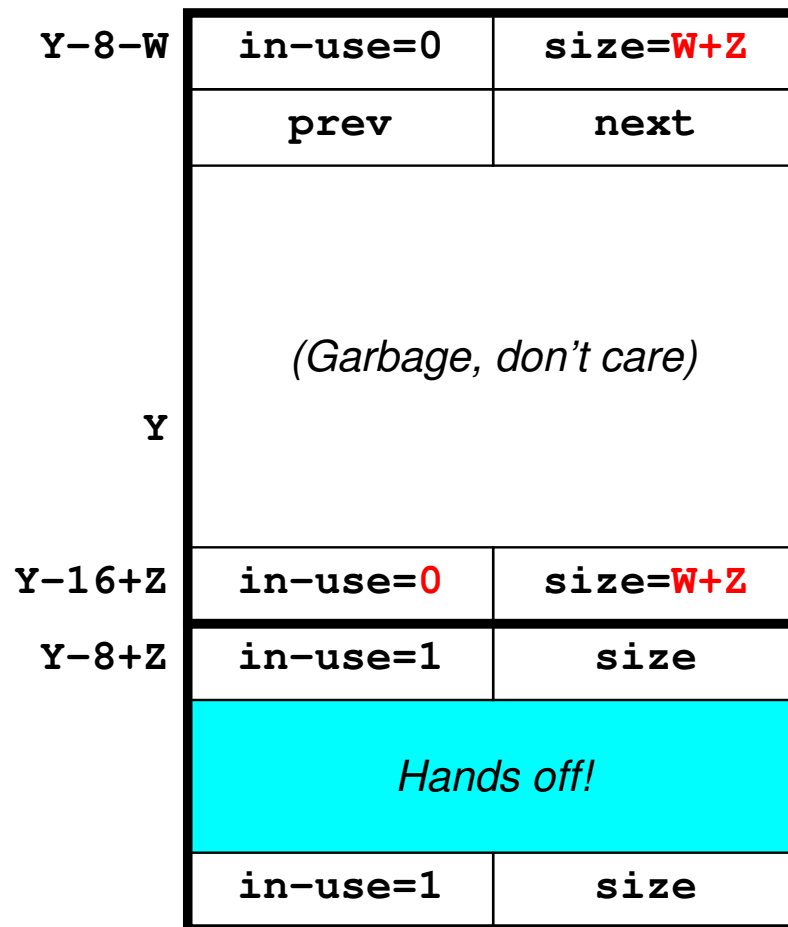
- ➡ Ex: free (Y) and *previous block is free* and *next block is busy*
- ▢ i.e., $Y-16$ is 0 and $Y-8+Z$ is 1
 - where Z is what's in $Y-4$ and W is what's in $Y-12$
 - ▢ furthermore, $Y-8-W$ is on the Free List
 - ▢ coalesce this block and the *previous* block



free () Example

➡ Ex: free (Y) and *previous block is free* and *next block is busy*


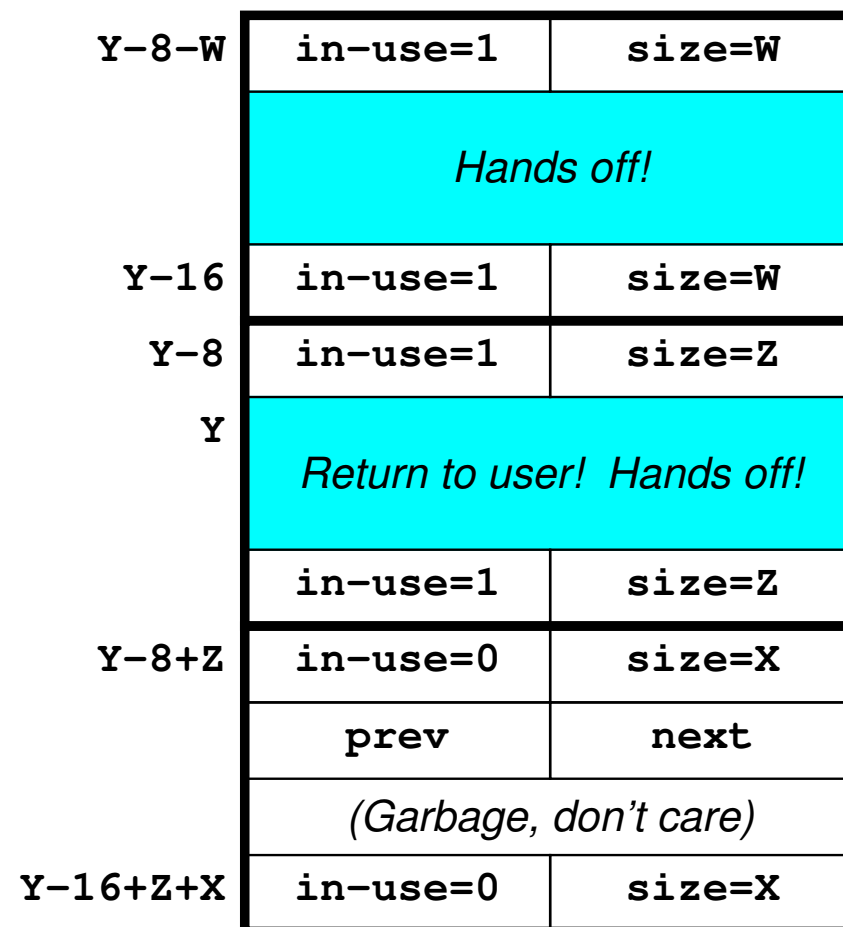
- ➡ i.e., $Y-16$ is 0 and $Y-8+Z$ is 1
 - where Z is what's in $Y-4$ and W is what's in $Y-12$
- ➡ furthermore, $Y-8-W$ is on the Free List
- ➡ coalesce this block and the *previous* block
 - easy!
 - just change $Y-12+Z$ and $Y-4-W$ to $W+Z$ and $Y-16+Z$ to 0
 - don't even need to change prev and next!



free () Example

➡ Ex: free (Y) and *previous block is busy* and *next block is free*

- ➡ i.e., Y-16 is 1 and Y-8+Z is 0
 - where Z is what's in Y-4 and X is what's in Y-4+Z
- ➡ furthermore, Y-8+Z is on the Free List
- ➡ coalesce this block and the *next* block



free () Example

➡ Ex: free (Y) and *previous block is busy* and *next block is free*

— i.e., Y-16 is 1 and Y-8+Z is 0

- where Z is what's in Y-4 and X is what's in Y-4+Z

— furthermore, Y-8+Z is on the Free List

— coalesce this block and the *next* block


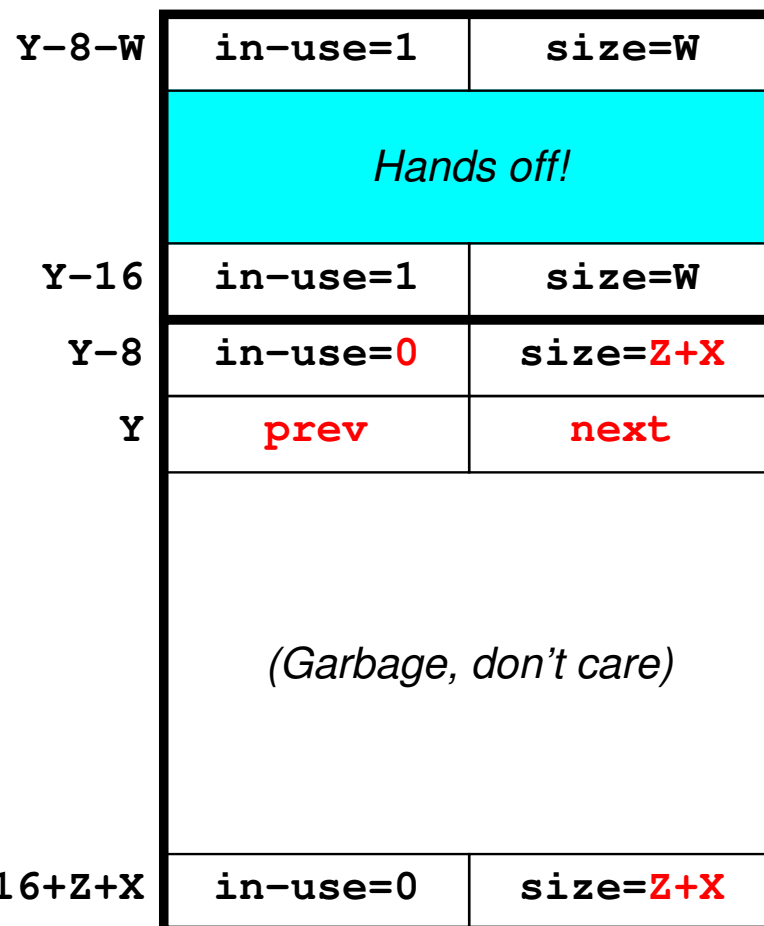
- just change Y-4 and Y-12+Z+X to Z+X and Y-8 to 0

- move prev and next pointers

- adjust next field in previous block in Free List

- adjust prev field in next block in Free List


- may need to update where Free List points



free () Example

- ➡ Ex: free (Y) and *previous block is free* and *next block is also free*
- ▬ i.e., $Y-16$ is 0 and $Y-8+Z$ is 0
 - where Z is what's in $Y-4$, X is what's in $Y-4+Z$, and W is what's in $Y-12$
 - ▬ blocks starting at $Y-8-W$ and $Y-8+Z$ are both on the Free List and next to and point at each other
 - ▬ coalesce all 3 blocks

$Y-8-W$	in-use=0	size=W
$Y-W$	prev	$Y-8+Z$
	(Garbage, don't care)	
$Y-16$	in-use=0	size=W
$Y-8$	in-use=1	size=Z
Y	Return to user! Hands off!	
	in-use=1	size=Z
$Y-8+Z$	in-use=0	size=X
$Y+Z$	$Y-8-W$	next
	(Garbage, don't care)	
	in-use=0	size=X



free () Example

➡ Ex: free (Y) and *previous block is free* and *next block is also free*


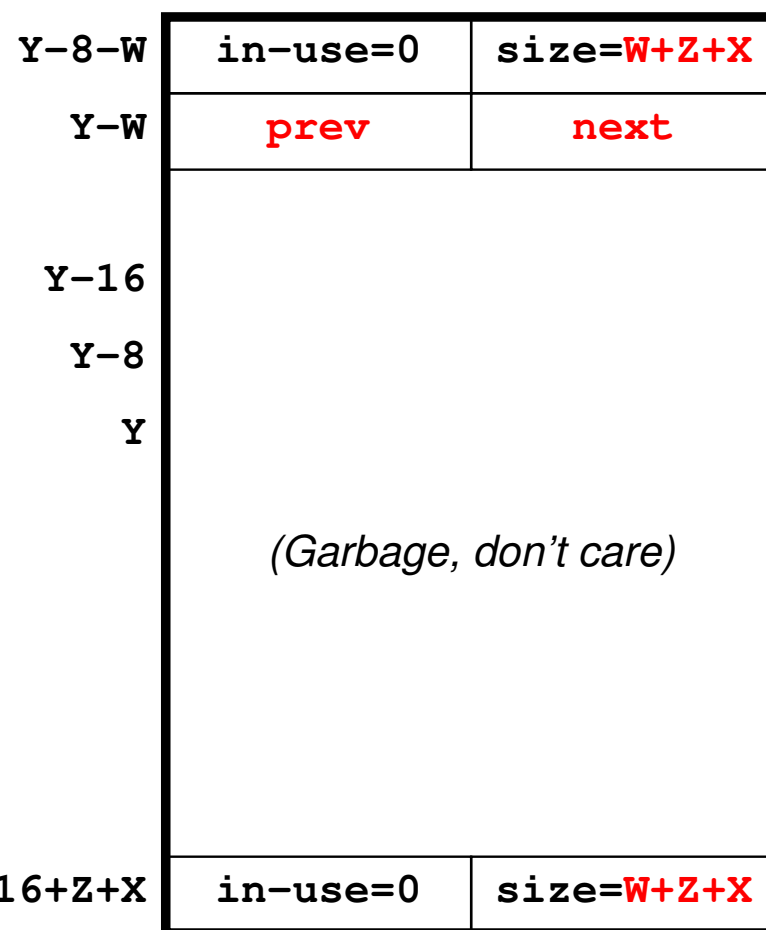
— i.e., $Y-16$ is 0 and $Y-8+Z$ is 0

- where Z is what's in $Y-4$,
 X is what's in $Y-4+Z$, and
 W is what's in $Y-12$

— blocks starting at $Y-8-W$ and
 $Y-8+Z$ are both on the Free List
and next to and point at each other

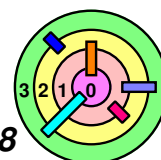
— coalesce all 3 blocks

- just change $Y-4-W$ and
 $Y-12+Z+X$ to $W+Z+X$
- copy next from $Y+Z+4$ to $Y-W+4$
- adjust prev field in the new next block in Free List to
point to $Y-8-W$
- may need to update where Free List points




First-fit & Best-fit Algorithms

- ➡ Memory allocator must run fast
 - it does not check if the free list is in a consistent state
 - just like our warmup 1 assignment
- ➡ One bad bit in the memory allocator data structure and it can break the memory allocator code
 - if you write into a *boundary tag*, your program may die in `malloc()` or `free()`
 - what would happen if you call `free()` twice on the same address?
 - user/application code can *corrupt the memory allocation chain* easily
 - the result can lead to *segmentation faults*
 - unfortunately, the corruption can *stay hidden* for a long time and *eventually* lead to a segmentation fault
 - ◆ memory corruption bugs are very difficult to squash




First-fit Algorithm

- ➡ Let n be the number of free blocks on the free list
 - ▬ `malloc()` is $O(n)$
 - ▬ `free(ptr)` is $O(n)$
 - occurs when the blocks around the block containing `ptr` are both in-use
- ➡ Such performance is unacceptable in the kernel




3.3 Dynamic Storage Allocation

- ➡ Best-fit & First-fit Algorithms
- ➡ *Buddy System*
- ➡ Slab Allocation




Buddy Lists

Ex: malloc(4000)



- blocks get evenly divided into two blocks that are buddies with each other
 - can only merge with your buddy if your buddy is also free
- internal fragmentation*
 - Ex: malloc(4000)
 - return a 4K block




Buddy Systems

➡ Faster memory allocation system (at the cost of more fragmentation, internal fragmentation)

— restrict block size to be a power of 2

- 1) all blocks of size 2^k start at location x where $x \bmod 2^k = 0$
- 2) given a block starting at location x such that $x \bmod 2^k = 0$
 - ◆ $BUDDY_k(x) = x + 2^k$ if $x \bmod 2^{k+1} = 0$
 - ◆ $BUDDY_k(x) = x - 2^k$ if $x \bmod 2^{k+1} = 2^k$
 - ◆ Ex: $BUDDY_2(1010100) = 1010000$
- 3) only buddies can be merged
- 4) try to coalesce buddies when storage is deallocated
- k different available block lists, one for each block size
- When request a block of size 2^k and none is available:
 - 1) split smallest block $2^j > 2^k$ into a pair of blocks of size 2^{j-1}
 - 2) place block on appropriate free list and try again



Buddy Systems



Data Structure

1) doubly-linked list (not circular) FREE list indexed by k

- ◆ links stored in actual blocks
- ◆ $\text{FREE}[k]$ points to first available block of size 2^k

2) each block contains

- ◆ in-use bit
- ◆ size
- ◆ NEXT and PREV links for FREE list


— lots of details

- read `weenix` source code for its "page allocator"




Can get greater variety in block sizes using Fibonacci sequence of block sizes so $b_j = b_{j-1} + b_{j-2}$

— ratio of successive block sizes is $2/3$ instead of $1/2$




High-level Example of Buddy Algorithm

Ex: 16 "pages" (minimum allocation is 1 page)




k	free[k]
0	Ω
1	Ω
2	Ω
3	Ω
4	0

1) allocate a block of size 2




k	free[k]
0	Ω
1	Ω 2
2	Ω 4
3	Ω 8
4	0 Ω

2) allocate a block of size 4




k	free[k]
0	Ω
1	Ω 2
2	Ω 4 Ω
3	Ω 8
4	0 Ω




High-level Example of Buddy Algorithm

Ex: 16 "pages" (minimum allocation is 1 page)

3) allocate a block of size 2




4) allocate a block of size 2




k	free[k]
0	Ω
1	0 2 Ω
2	0 4 Ω
3	0 8
4	0 Ω

k	free[k]
0	Ω
1	0 2 10
2	0 4 12
3	0 8 Ω
4	0 Ω




3.3 Dynamic Storage Allocation


- ➡ Best-fit & First-fit Algorithms
- ➡ Buddy System
- ➡ *Slab Allocation*



Slab Allocation




— see weenix kernel code!



Slab Allocation

- ➡ **Objects are allocated and freed frequently**
 - **allocation involves**
 - finding an appropriate-sized storage
 - initialize it
 - ◆ pointers need to point at the right places
 - ◆ may even need to initialize synchronization data structures
 - **deallocation involves**
 - tearing down the data structures
 - freeing the storage
 - **lots of "overhead"**
- ➡ **Difficulties with dynamic storage allocation**
 - you cannot predict what an application will ask for
 - but it's not true for the kernel
 - e.g., can allocate a slab of process control blocks at a time
 - ◆ return one of them from a slab



Slab Allocation

- ➡ **Slab Allocation**
 - sets up a separate cache for each type of object to be managed
 - contiguous sets of pages called **slabs**, allocated to hold objects
 - we will cover "pages" later, won't get into too much detail now
- ➡ Whenever a **slab** is allocated, a constructor is called to initialize all the objects it hold
 - this is where you pay for initialization, but it's done in a **batch**
- ➡ As **objects** are being allocated, they are taken from the set of existing slabs in the cache
 - objects are considered "preallocated" since they have all been initialized already
- ➡ As **objects** are being freed, they are simply marked as free
 - don't have to free up storage
 - when appropriate can free up an entire slab

