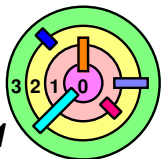


3.4 Linking & Loading

➡ *Static Linking & Loading*

➡ **Shared Libraries**



Remember This?

```

main:
→ pushl %ebp
  movl %esp, %ebp
  pushl %esi
  pushl %edi
  subl $8, %esp
  ...
  pushl $1
  movl -12(%ebp), %eax
  pushl %eax
  call sub
  addl $8, %esp
  movl %eax, -16(%ebp)
  ...
  addl $8, %esp
  movl $0, %eax
  popl %edi
  popl %esi
  movl %ebp, %esp
  popl %ebp
  ret

```

set up stack frame

push args

pop args; get result

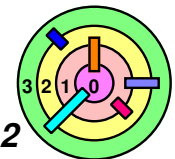
set return value and restore frame



```

int main() {
    int i;
    int a;
    ...
    i = sub(a, 1);
    ...
    return(0);
}

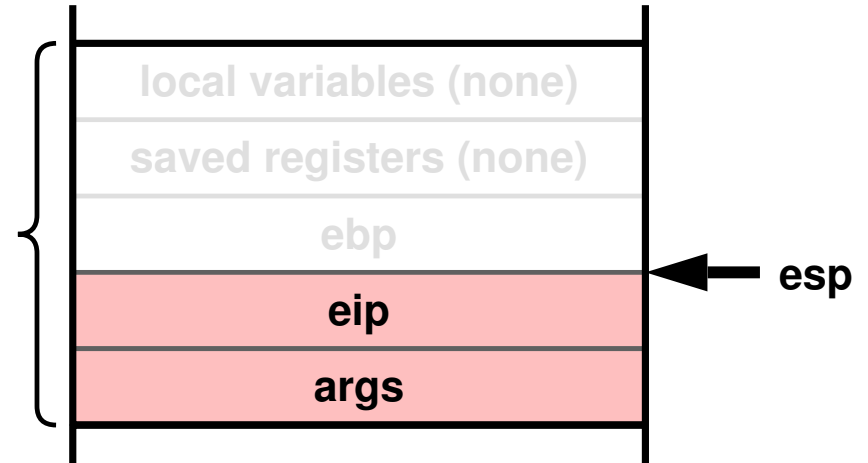
```



Something Simpler

```
int main(int argc,
        char *[]) {
    return(argc);
}
```

stack frame
of main()



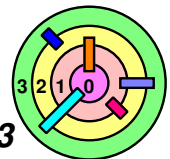
main:

```
→ pushl %ebp
   movl %esp, %ebp
   movl 8(%ebp), %eax
   movl %ebp, %esp
   popl %ebp
   ret
```

set up
stack frame

set return
value and
restore frame

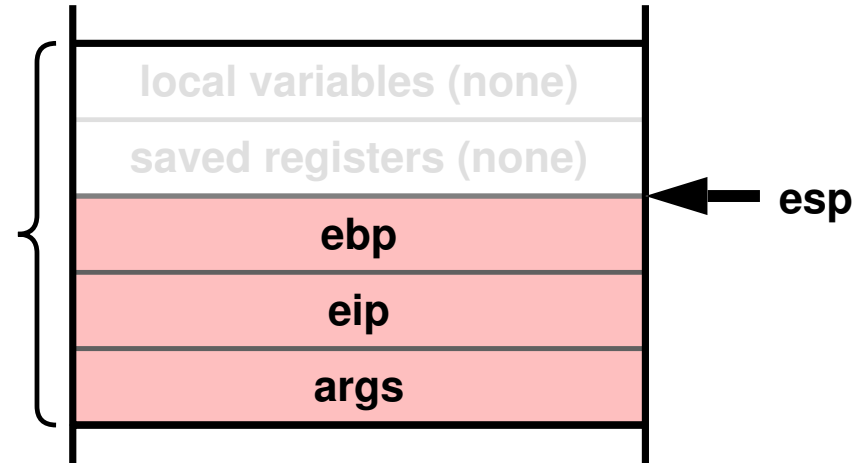
➡ Does location matter?



Something Simpler

```
int main(int argc,
         char *[]) {
    return(argc);
}
```

stack frame
of main()



main:

```

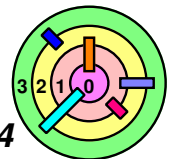
    pushl %ebp
    → movl %esp, %ebp
    movl 8(%ebp), %eax
    movl %ebp, %esp
    popl %ebp
    ret

```

set up
stack frame

set return
value and
restore frame

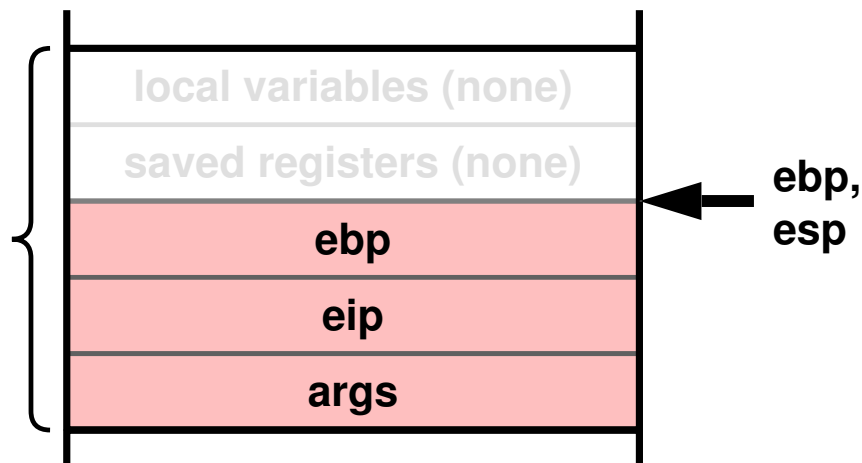
➡ Does location matter?



Something Simpler

```
int main(int argc,
         char *[]) {
    return(argc);
}
```

stack frame
of main()

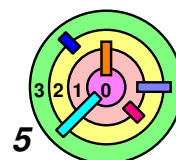


```
main:
    pushl %ebp
    movl %esp, %ebp
    → movl 8(%ebp), %eax
    movl %ebp, %esp
    popl %ebp
    ret
```

set up
stack frame

set return
value and
restore frame

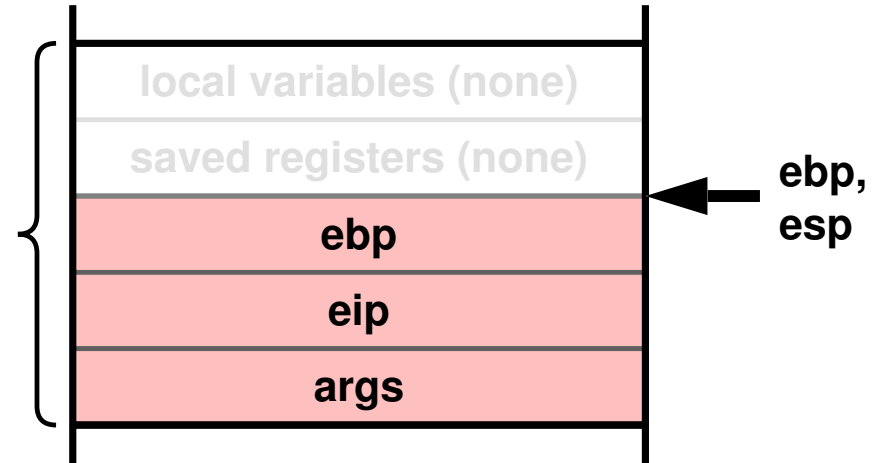
➡ Does location matter?



Something Simpler

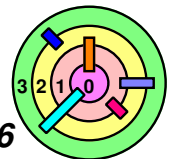
```
int main(int argc,
        char *[]) {
    return(argc);
}
```

stack frame
of main()



```
main:
    pushl %ebp          } set up
    movl %esp, %ebp     } stack frame
    movl 8(%ebp), %eax
    → movl %ebp, %esp    } set return
    popl %ebp           } value and
    ret                 } restore frame
```

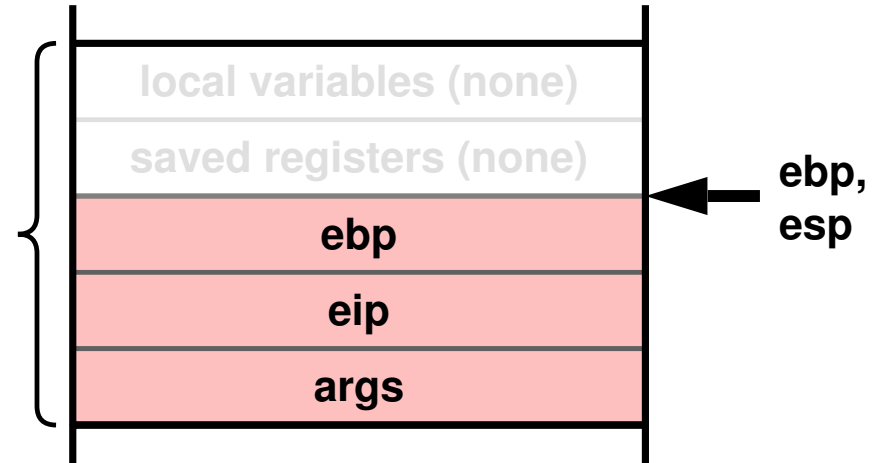
➡ Does location matter?



Something Simpler

```
int main(int argc,
         char *[]) {
    return(argc);
}
```

stack frame
of main()

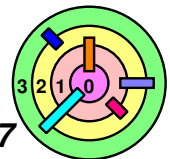


```
main:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    movl %ebp, %esp
    → popl %ebp
    ret
```

set up
stack frame

set return
value and
restore frame

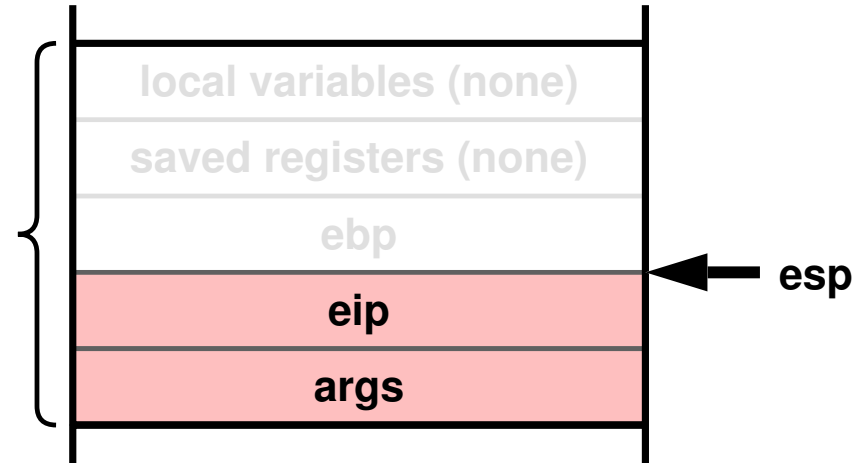
➡ Does location matter?



Something Simpler

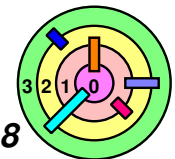
```
int main(int argc,
        char *[]) {
    return(argc);
}
```

stack frame
of main()



```
main:
    pushl %ebp           } set up
    movl %esp, %ebp      } stack frame
    movl 8(%ebp), %eax
    movl %ebp, %esp      } set return
    popl %ebp            } value and
    ret                  } restore frame
```

➡ Does location matter?

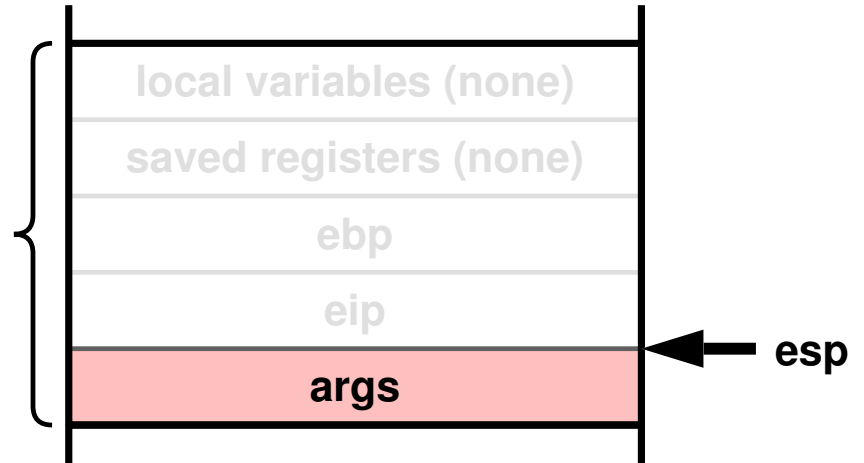


Something Simpler

```
int main(int argc,
        char *[]) {
    return(argc);
}
```

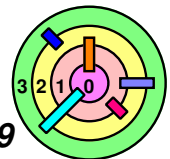
```
main:
    pushl %ebp          } set up
    movl %esp, %ebp     } stack frame
    movl 8(%ebp), %eax
    movl %ebp, %esp      } set return
    popl %ebp            } value and
    ret                  } restore frame
```

stack frame
of main()



Does location matter?

- ➡ if everything can be accessed relative to the *frame pointer*, then you don't need to know the actual address of an object
 - just use *relative-addresses* to *access variables*
 - the *code* is also *location-independent*



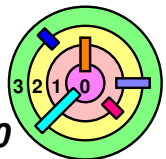
Location Matters ...

```
int X = 6;  
int *aX = &X;
```

```
void subr(int i) {  
    printf("i = %d\n", i);  
}
```

```
int main( ) {  
    void subr(int);  
    int y = X;  
    subr(y);  
    return(0);  
}
```

- ➡ Why does it matter here?
- need to put the *address* of x into aX
 - *what* is the address of x?
 - ◆ *when* do you know?
 - remember, both x and aX are in the *data segment*
 - *who* would put the actual value into aX?
 - also need to put the *address* of subr into main()



Coping

➡ Relocation

- ➡ modify internal references in memory depending on where module is expected to be *loaded*
 - one of the *exec* system calls loads a program into memory
 - ◆ everything is laid out carefully in memory
- ➡ modules requiring relocation are said to be *relocatable*
- ➡ the act of modifying such a module to *resolve these references* is called *relocation*
- ➡ the program that performs relocation is called a *linker*

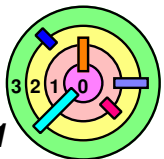
➡ Two main functions of a *linker*

- 1) *relocation*
- 2) *symbol resolution*

textbook
is wrong

➡ A *loader* loads a program into memory

- ➡ a "relocating loader" may perform additional relocation



A Slight Revision

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```

main.c

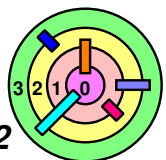
```
#include <stdio.h>
int X;

void subr(int i) {
    printf("i = %d\n", i);
}
```

subr.c

```
% gcc -o prog main.c subr.c
```

- ⇒ main.c is compiled into main.o
 - ⇒ subr.c is compiled into subr.o
- } relocatable modules
- ⇒ ld is then invoked to combine them into prog
 - ld knows where to find printf()
 - prog can be loaded into memory through one of the **exec** system calls



A Slight Revision

```
extern int X;  
int *aX = &X;  
  
int main( ) {  
    void subr(int);  
    int y = *aX;  
    subr(y);  
    return(0);  
}
```

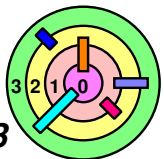
main.c

```
#include <stdio.h>  
int X;  
  
void subr(int i) {  
    printf("i = %d\n", i);  
}
```

subr.c

```
% gcc -o prog main.c subr.c
```

- how does ld decides what needs to be done?
- main.c contains undefined references to X and subr()
 - instructions for doing this are provided in main.o
- later on, when the actual locations for these are determined, ld will modify them when main.o is *copied* into prog



A Slight Revision

```
extern int X;  
int *aX = &X;  
  
int main( ) {  
    void subr(int);  
    int y = *aX;  
    subr(y);  
    return(0);  
}
```

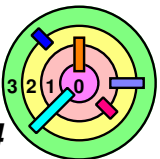
main.c

```
#include <stdio.h>  
int X;  
  
void subr(int i) {  
    printf("i = %d\n", i);  
}
```

subr.c

```
% gcc -o prog main.c subr.c
```

- main.o must contain a list of external symbols, along with their types, and instructions for updating this code

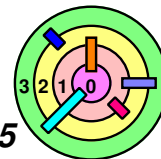


main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0: aX:		
0:	.long	X
4:		
0:	.text	; offset restarts; what follows is ; text (read-only code)
0:	.globl	main
0: main:		
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



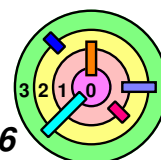
main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

What follows goes into the *data* segment

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



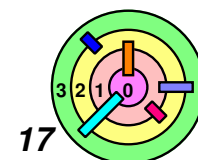
main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

What follows goes into the **text** segment

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



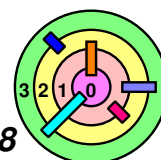
main.s

Offset	Op	Arg
→ 0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
→ 0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

offset got restarted because
segments are relocatable

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



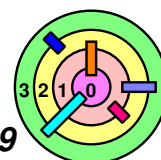
main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
→ 0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
→ 0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

.global directive means that the symbol mentioned is defined *here* and is *exported*
 ➡ i.e., can be referenced by other modules
 ➡ aX and main are global

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



main.s

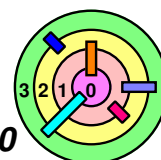
Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

aX is 4 bytes long and put
the value of x here

→ x will remain *unresolved*

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



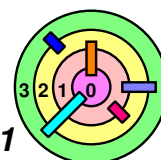
main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

these 3 places require
relocation

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



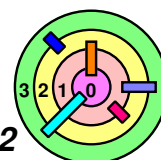
main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

this call is a PC-relative call
 ➡ what's stored at offset 20 is not the absolute address of subr, but a relative address
 ➡ this is just how x86 works

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```

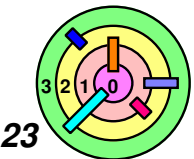


subr.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS is required ; for global X
4:		
0:	.text	; offset restarts; what follows is ; text (read-only code)
0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp, %ebp ; point to new frame
3:	pushl	8(%ebp) ; push i onto stack
6:	pushl	\$printfarg ; push address of printf arg ; onto stack
11:	call	printf
16:	addl	\$8, %esp ; pop arguments
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

```
#include <stdio.h>
int X;

void subr(int i) {
    printf("i = %d\n", i);
}
```



subr.s

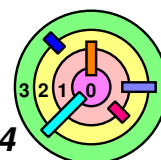
Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS ; for global X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the fra
1:	movl	%esp, %ebp ; point t
3:	pushl	8(%ebp) ; push i on
6:	pushl	\$printfarg ; push ad ; onto st
11:	call	printf
16:	addl	\$8, %esp ; pop argum
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

this is how you create a string constant

- ▢ this one is 8 bytes long
- ▢ and local to this module (since it's not global)

```
#include <stdio.h>
int X;
```

```
void subr(int i) {
    printf("i = %d\n", i);
}
```



subr.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS ; for global X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the fra
1:	movl	%esp, %ebp ; point t
3:	pushl	8(%ebp) ; push i on
6:	pushl	\$printfarg ; push ad ; onto st
11:	call	printf
16:	addl	\$8, %esp ; pop argum
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

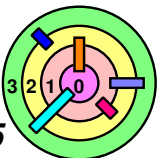
this is how you create a string constant

- ▢ this one is 8 bytes long
- ▢ and local to this module (since it's not global)
- ▢ it is used here

```
#include <stdio.h>
```

```
int X;
```

```
void subr(int i) {  
    printf("i = %d\n", i);  
}
```



subr.s

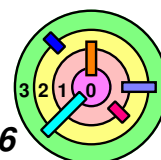
Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS ; for global X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the fra
1:	movl	%esp, %ebp ; point t
3:	pushl	8(%ebp) ; push i on
6:	pushl	\$printfarg ; push ad ; onto st
11:	call	printf
16:	addl	\$8, %esp ; pop argum
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

4 bytes is required in the **bss** segment for this global variable

☞ .comm directive also means that the symbol mentioned is defined **here** and is **exported**

```
#include <stdio.h>
int X;
```

```
void subr(int i) {
    printf("i = %d\n", i);
}
```



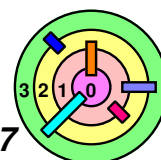
subr.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS ; for global X
4:		
0:	.text	; offset restarts; w ; text (read-only co
→ 0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the fra
1:	movl	%esp, %ebp ; point t
3:	pushl	8(%ebp) ; push i on
6:	pushl	\$printfarg ; push ad ; onto st
11:	call	printf
16:	addl	\$8, %esp ; pop argum
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

subr is a global symbol
exported from here

```
#include <stdio.h>
int X;

void subr(int i) {
    printf("i = %d\n", i);
}
```



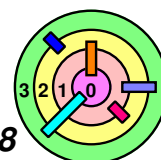
subr.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS ; for global X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the fra
1:	movl	%esp, %ebp ; point t
3:	pushl	8(%ebp) ; push i on
→ 6:	pushl	\$printfarg ; push ad ; onto st
→ 11:	call	printf
16:	addl	\$8, %esp ; pop argum
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

relocation is required for
printf and printfarg

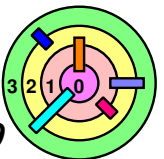
```
#include <stdio.h>
int X;

void subr(int i) {
    printf("i = %d\n", i);
}
```



Object Files

- ➡ An object file describes what's in the data, bss, and text segments in separate sections
- ➡ Along with each section is a list of:
 - ▬ global symbols
 - ▬ undefined symbols
 - ▬ instructions for relocation
 - these instructions indicate
 - ◆ which locations within the section must be modified
 - ◆ which symbol's value is used to modify the location
 - a symbol's value is the address that is ultimately determined for it
 - typically, this address is added to the location being modified
- ➡ To inspect an object file on Unix
 - ▬ nm - list symbols from object files
 - ▬ objdump - display information from object files



subr.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS ; for global X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the fra
1:	movl	%esp, %ebp ; point t
3:	pushl	8(%ebp) ; push i onto stack
→ 6:	pushl	\$printfarg ; push address of string ; onto stack
→ 11:	call	printf
16:	addl	\$8, %esp ; pop arguments from stack
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

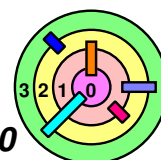
relocation is required for:

▢ printfarg at offset 7

▢ printf at offset 12

```
#include <stdio.h>
int X;
```

```
void subr(int i) {
    printf("i = %d\n", i);
}
```



subr.o

Data:

Size: 8

Contents: "i = %d\n"

bss:

Size: 4

Global: X, offset 0

Text:

Size: 24

Global: subr, offset 0

Undefined: printf

Relocation:

offset 7, size 4, value: addr of printfarg

offset 12, size 4, value: PC-relative addr of
printf

Contents: [machine instructions]

relocation is required for:

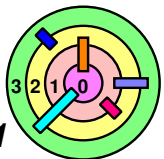
▢ printfarg at offset 7

▢ printf at offset 12

```
#include <stdio.h>
```

```
int X;
```

```
void subr(int i) {  
    printf("i = %d\n", i);  
}
```



subr.o

Data:

Size: 8

Contents: "i = %d\n"

bss:

Size: 4

→ Global: X, offset 0

Text:

Size: 24

→ Global: subr, offset 0

Undefined: printf

Relocation:

offset 7, size 4, value: addr of printfarg

offset 12, size 4, value: PC-relative addr of
printf

Contents: [machine instructions]

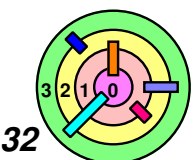
X and subr are *exported*

→ needed in main.o

```
#include <stdio.h>
```

```
int X;
```

```
void subr(int i) {  
    printf("i = %d\n", i);  
}
```



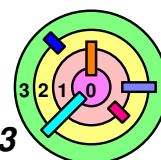
main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the fram
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y onto stack
19:	call	subr
24:	addl	\$4,%esp ; remove y from stack
27:	movl	\$0,%eax ; set return value to 0
31:	movl	%ebp, %esp ; restore stack pointer
33:	popl	%ebp ; pop frame pointer
35:	ret	

these 2 places remained
unresolved

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



main.o

Data:

Size: 4

Global: aX, offset 0

Undefined: X

Relocation: offset 0, size 4, value

Contents: 0x00000000

these 2 places remained

unresolved

→ they are noted in main.o

bss:

Size: 0

Text:

Size: 36

Global: main, offset 0

Undefined: subr

Relocation:

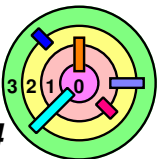
offset 7, size 4, value: addr of aX

offset 20, size 4, value: PC-relative
addr of subr

Contents: [machine instructions]

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



subr.o

Data:

Size: 8

Contents: "i = %d\n"

bss:

Size: 4

Global: X, offset 0

Text:

Size: 24

Global: subr, offset 0

Undefined: printf

Relocation:

offset 7, size 4, value: addr of printfarg

offset 12, size 4, value: PC-relative addr of
printf

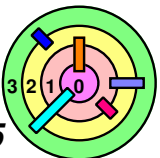
Contents: [machine instructions]

printf remained *unresolved*

```
#include <stdio.h>
```

```
int X;
```

```
void subr(int i) {  
    printf("i = %d\n", i);  
}
```



printf.o

Data:

Size: 1024

Global: StandardFiles

Contents: ...

assume that printf.o looks like this

→ write is *unresolved*

bss:

Size: 256

Text:

Size: 12000

Global: printf, offset 100

...



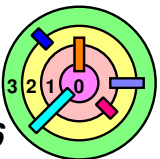
Undefined: write

Relocation:

offset 211, value: addr of StandardFiles

offset 723, value: PC-relative addr of write

Contents: [machine instructions]



write.o

Data:

Size: 0

bss:

Size: 4

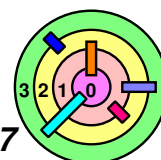
Global: errno, offset 0

Text:

Size: 16

Contents: [machine
instructions]

and write.o looks like this



startup function

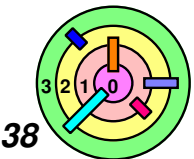
```
Data:
  Size:  0

bss:
  Size:  0

Text:
  Size:  36
  Undefined: main
  Relocation:
    offset 21, value: main
  Contents: [machine
             instructions]
```

every C program contains a *startup* routine that is called first

- ▢ it calls `main()`
- ▢ if `main()` returns, it calls `exit()`
- ▢ our example is incomplete



prog

Text

main	4096
subr	4132
printf	4156
write	16156
startup	16172

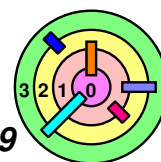
Data

aX	16384
printfargs	16388
StandardFiles	16396

BSS

X	17420
errno	17424

- ⇒ this is how `ld` might set things up
 - `ld` *lays out* the *address space*
 - `ld` allocates memory in *pages* (typically 4KB each)
- ⇒ `main` does not start at location 0
 - first "page" is typically made inaccessible so that references to null pointers will fail (get SIGSEG)



prog

Text

main	4096
subr	4132
printf	4156
write	16156
startup	16172

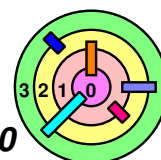
Data

aX	16384
printfargs	16388
StandardFiles	16396

BSS

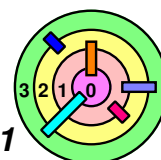
X	17420
errno	17424

- ➡ due to the use of "*pages*", the data segment needs to start at a page boundary (i.e., multiple of page size)
 - this way, the *text segment* can be made *read-only* while the *data* and *bss segments* made *read-write*
 - here we assume 4KB pages (therefore, pages start at 4096, 8192, 12288, 16384, etc.)



Virtual Memory Basics

- ➡ A process has, say, a 32-bit address space
 - that's 4GB of memory
- ➡ Our prog process, when it starts, only needs about 16KB for text+data+bss
 - plus more for stack
- ➡ Allocating 4GB of physical memory will be a *huge waste*
- ➡ Solution: *page table* to *map virtual to physical addresses*
 - OS allocate *pages* of *physical memory*
 - a page is 4KB in many systems
 - a page corresponds to physical memory that can be located (or "*mapped*") anywhere in virtual memory
 - one level of *indirection* to get to the physical memory
 - the hardware makes this transparent
- ➡ We will spend a lot of time talking about virtual memory (Ch 7)



Virtual Memory Basics

Text

main 4096
subr 4132
printf 4156
write 16156
startup 16172

Data

aX 16384
printfargs 16388
StandardFiles 16396

BSS

X 17420
errno 17424

Page Table

Start	Access	Physical Addr
0	-	-
4096	R	•
8192	R	•
12288	R	•
16384	R/W	•

ask buddy system to
allocate these pages

