

3.4 Linking & Loading

- Static Linking & Loading
- Shared Libraries

- A library is just a collection of .o files
- the linker is used to create libraries
- Two types of libraries
 - static library
 - dynamic (or shared) library

Libraries



Copyright © William C. Cheng



Copyright © William C. Cheng

Creating a Static Library

```
% cat sub1.c
void sub1() { puts("sub1"); }
% cat sub2.c
void sub2() { puts("sub2"); }
% cat sub3.c
void sub3() { puts("sub3"); }
% gcc -c sub1.c sub2.c sub3.c
% ls
sub1.o sub2.o sub3.o
% ar cr libpriv1.a sub1.o sub2.o sub3.o
% ar t libpriv1.a
sub1.o
sub2.o
sub3.o
% gcc -o prog prog.c -L. -lpriv1
```

- puts () is unresolved in libpriv1.a

Operating Systems - CSCI 402



Copyright © William C. Cheng

- will try to resolve puts () first in the priv1 library, then in the myputs library, then in the c library

```
% gcc -c myputs.c
% ar cr libmyputs.a myputs.o
% gcc -o prog prog.c -L. -lpriv1 -lmyputs

int puts(char *s)
{
    write(1, "My puts: ", 6);
    write(1, s, strlen(s));
    write(1, "\n", 1);
    return 1;
}

library
```

- Want to use my version of puts () instead of what's in the C library


Substitution

Operating Systems - CSCI 402

Copyright © William C. Cheng

Shared Libraries

- prog must contain *everything* needed for execution
 - duplicate code, may be lots of duplicate code, e.g., printf ()
 - take up *disk space*
 - take up *memory*
- Need a way to share things like printf ()
 - on disk, and
 - in memory
- if printf () required no relocation, then it's easy
- if printf () required relocation, then it's more complicated
 - problem: processes want a shared function to be at different addresses



Operating Systems - CSCI 402

Copyright © William C. Cheng

Using a Static Library

- ```
% cat prog.c
int main()
{
 sub1();
 sub2();
 sub3();
}
% gcc -o prog prog.c -L. -lpriv1

Where does puts () come from?
% gcc -o prog prog.c -L. -lpriv1 -L/lib -lc

The order of the libraries matter
% will try to resolve references first in the priv1 library (either libpriv1.a or libpriv1.so) and then in the c library (either libc.a or libc.so)
```



Operating Systems - CSCI 402

Copyright © William C. Cheng



Operating Systems - CSCI 402

Copyright © William C. Cheng

11

Please note that `ld` is not the same as `mov` in x86 CPU

- each process maintains a private table, pointed to by register `r1`
- table contains addresses of shared routines
- don't call functions directly
- make a position-independent call (i.e., an indirect call)
- i.e., call the function located at a *fixed index* into the table
- implemented as two instructions in the above example

Position-Independent Code

Operating Systems - CSCI 402

Copyright © William C. Cheng

9

What about this?

Sharing

Operating Systems - CSCI 402

Copyright © William C. Cheng

7

Sharing

Operating Systems - CSCI 402

Copyright © William C. Cheng

12

Position-Independent Code Details

Processor-dependent; x86 32-bit version:

ELF requires 3 data structures for each dynamic executable and shared object

- the *procedure linkage table (PLT)*
- read-only executable code, *shared* by all processes
- essentially stubs for calling subroutines
- the *global offset table (GOT)*
- read-write data, private (to each process)
- relocated dynamically for each process
- the *dynamic structure*
- read-only data, *shared* by all processes
- contains relocation info and symbol table

Operating Systems - CSCI 402

Copyright © William C. Cheng

10

Relocation and Shared Libraries

Approaches

- Limited sharing: relocate separately for each process
- have a single copy of `printf()` on disk
- as `printf()` gets copied into memory, perform relocation
- this would work, but still end up with too many copies of `printf()` in memory
- Prelocation*: relocate libraries ahead of time
- difficult to prelocate all shared functions
- may need to perform relocation

*Position-Independent Code*: no need for relocation

- producing code that can be placed anywhere in memory without requiring modification
- need *indirection*

Operating Systems - CSCI 402

Copyright © William C. Cheng

8

Sharing

Looks like it can work

Operating Systems - CSCI 402



- ```
% gcc -o prog prog.c -L. -lpravl -lmyputs -w1,-xpath .  
%  
ldd prog  
libmyputs.so => ./../libmyputs.so  
libc.so.6 => /lib/tls/i686/cmov/libc.so.6  
/lib/ld-linux.so.2  
%.progr  
My puts : sub1  
My puts : sub2  
My puts : sub3
```

Creating a Shared Library (2)



- When a program is invoked via the **exec** system call
 - the code that is first given control is **ld**, so, the **run-time linker**
 - the job of **ld**, so is to complete the linking and relocation steps, if necessary
- it does some initial set up of linkages
- then calls the actual program code
- it may be called upon later to do some further loading and linking

Shared Library Details



- ➡ **Shared libraries** are used extensively in many modern systems
 - ➡ often implemented with either preallocation or position-independent code
 - ➡ in Windows, they are known as **Dynamic-Link Libraries (DLLs)** (so files)
 - ➡ in Unix, they are known as **shared objects** (so files)
 - ➡ vs. **static libraries** (a files)
 - ➡ they need not be loaded when a program starts up
 - can be loaded when needed, i.e., **on-demand**
 - this way, the startup time of a program may be reduced
- ➡ **Disadvantages of DLLs and shared objects**
 - ➡ they can have **dependencies**
 - ➡ different **versions** of the same library

Shared Libraries in Practice



- [illegible]

Versioning



- ```
% gcc -fPIC -c myputs.c
% ld -o libmyputs.so myputs.o
% gcc -o prog prog.c -L. -lmyputs
./prog
./prog: error while loading shared libraries:
libmyputs.so: cannot open shared object file: No
such file or directory
% ldd prog
libmyputs.so => not found
libmyputs.so => not found
11b9c.so.6 => /lib64/ld-linux.so.2
11b9c.so.6 => /lib64/cmov/libc.so.6
11b9c.so.2 => /lib64/ld-linux.so.2
```

## Creating a Shared Library (1)



- x86 ELF (Executable and Linking Format)
- used in Unix/Linux systems
- not used in either MacOS X or Windows
- Creating and using a shared library
- Substitution
- Shared library details
- Versioning
- Dynamic linking
- Interpositioning

## Linking and Loading on Linux with ELF



- glibc
- compiles code
- does *static linking*
- searches list of libraries
- adds *references to shared objects*
- runtime
  - program invokes `ld`, so (or `ld-linux`, so) to *finish linking*
  - maps in shared objects
  - does relocation and procedure linking as required
  - `dlsym()` invokes `ld` to do *more linking*

## What's Going On ...



```
% gcc -o tputs tputs.c
% ./tputs
This is a boring message.
% setenv _LTPRELOAD ./ltpbmytputs.so.1
% ./tputs
calling mytputs: This is a boring message.
```

- LD\_PRELOAD
  - environment variable checked by ld, so
  - specifies additional shared objects to search (first) when program is started

## Delayed Wrapping



- ```
% cat myputs.c
#include <stdio.h>

int puts(const char *s)
{
    int pptr = (const char *)RTLD_NEXT;
    while(2,"calling myputs:", 16);
    return (*pptr)(s);
}
```

How To ...




- D_GNU_SOURCE is needed or won't recognize RLD_NEXT
- ldconfig may be in /sbin

```
% gcc -fPIC -c myputs.c -D-GNU_SOURCE
% ld -shared -soname libmyputs.so.1 \
-o libmyputs.so.1.0 myputs.o -ldl
% ldconfig -v -n
% cat tputs.c
% int main()
% {
puts("This is a boring message.");
return 0;
}
% gcc -o tputs tputs.c ./libmyputs.so.1 -Wl,-rpath
calling myputs: This is a boring message.
```

Compiling/Linking it




- prog thinks it's calling puts ()



Interpositioning



- prog thinks it's calling puts ()
- interpose your puts ()
- you can call the original puts () that prog thought it was calling
- security problem? "DLL injection" if you pick up a DLL unknowingly



Interpositioning

Extra Slides

- Processor-dependent; x86 32-bit version:
- ELF requires 3 data structures for each dynamic executable and shared object
 - the *procedure linkage table (PLT)*
 - read-only executable code, *shared* by all processes
 - essentially stubs for calling subroutines
 - the *global offset table (GOT)*
 - read-write data, private (to each process)
 - relocated dynamically for each process
 - the *dynamic structure*
 - read-only data, *shared* by all processes
 - contains relocation info and symbol table

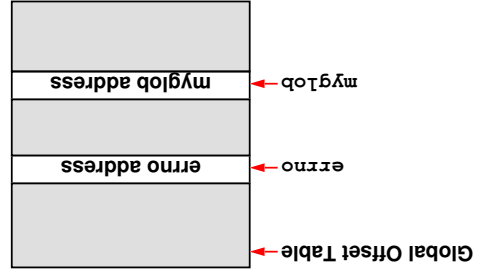
Position-Independent Code Details



Copyright © William C. Cheng



Global-Offset Table: Data References



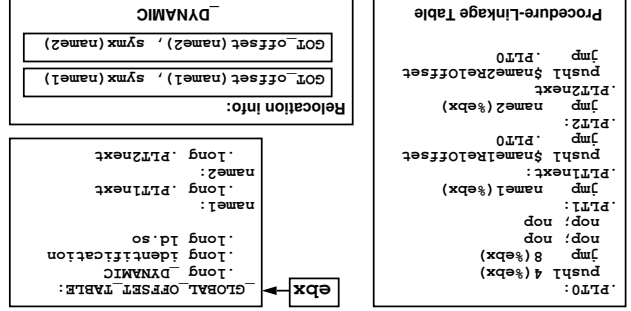
- Modules refer to *global variables* indirectly by looking up their addresses in this table
- a *register* contains the address of the table
- modules refer to entries via their *offsets*
- When a module is loaded into memory
- ld.so puts the actual addresses into the GOT



Copyright © William C. Cheng



Before Calling name1



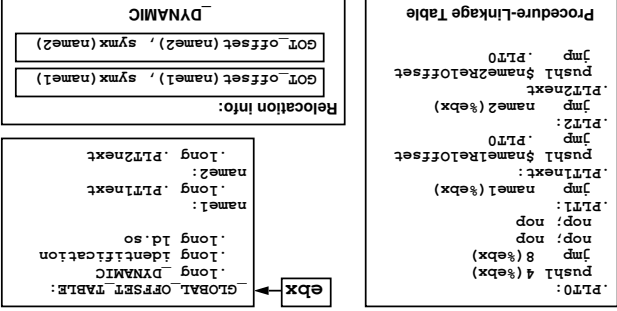
- Before the first call to name1
- the actual address of the name1 procedure is *not* in the GOT
- first call to name1 invokes ld.so with indication of the above fact
- ld.so finds name1 and update the GOT



Copyright © William C. Cheng



Before Calling name1



- How did this work?
- references from module to name1 are statically linked to entry .PLT1 in the procedure-linkage table
- the index into the symbol table in _DYNAMIC tells use where to start (in this example, it's 0 for name1)



Copyright © William C. Cheng



Procedure References

- More complicated than data references
- Lots of them
- Many are never used
- Fix them up on demand



Copyright © William C. Cheng



39

Copyright © William C. Cheng

After Calling name1

Procedure-Linkage Table

pushl 4(%ebx)

jmp 8(%ebx)

nop

nop

pushl name1(%ebx)

pushl name1ReelOffset

jmp .PLT1

pushl name2(%ebx)

pushl name2ReelOffset

jmp .PLT2

pushl .PLT0

jmp

Relocation info:

GOT_offset(name1), symx(name1)

GOT_offset(name2), symx(name2)

DYNAMIC

GOT_offset(name1), symx(name1)

GOT_offset(name2), symx(name2)

ebx

GLOBAL_OFFSET_TABLE:

pushl _DYNAMIC

long identification

long Id.so

name1:

name2:

pushl .PLT2next

Subsequent call to name1 invokes more directly

37

Copyright © William C. Cheng

Before Calling name1

Procedure-Linkage Table

pushl 4(%ebx)

jmp 8(%ebx)

nop

nop

pushl name1(%ebx)

pushl name1ReelOffset

jmp .PLT1

pushl name2(%ebx)

pushl name2ReelOffset

jmp .PLT2

pushl .PLT0

jmp

Relocation info:

GOT_offset(name1), symx(name1)

GOT_offset(name2), symx(name2)

DYNAMIC

GOT_offset(name1), symx(name1)

GOT_offset(name2), symx(name2)

offset 8

ebx

GLOBAL_OFFSET_TABLE:

pushl _DYNAMIC

long identification

long Id.so

name1:

name2:

pushl .PLT2next

ld, so can figure out who was making the request (TOS)

TOS-4 contains the GOT offset (12 in our example) of where to write the result into

TOS-8 contains index into symbol table in _DYNAMIC

38

Copyright © William C. Cheng

Before Calling name1

Procedure-Linkage Table

pushl 4(%ebx)

jmp 8(%ebx)

nop

nop

pushl name1(%ebx)

pushl name1ReelOffset

jmp .PLT1

pushl name2(%ebx)

pushl name2ReelOffset

jmp .PLT2

pushl .PLT0

jmp

Relocation info:

GOT_offset(name1), symx(name1)

GOT_offset(name2), symx(name2)

DYNAMIC

GOT_offset(name1), symx(name1)

GOT_offset(name2), symx(name2)

ebx

GLOBAL_OFFSET_TABLE:

pushl _DYNAMIC

long identification

long Id.so

name1:

name2:

pushl .PLT2next

ld, so writes the actual address of the name1 procedure into the name1 entry of the GOT

unwinds the stack a bit and then passes control to name1