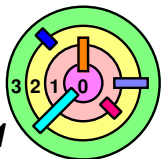


3.4 Linking & Loading


➡ Static Linking & Loading

➡ *Shared Libraries*



Libraries


- ➡ A library is just a collection of `.o` files
 - the linker is used to create libraries
- ➡ Two types of libraries
 - *static* library
 - *dynamic* (or *shared*) library



Creating a Static Library

```
% cat sub1.c
void sub1() { puts("sub1"); }
% cat sub2.c
void sub2() { puts("sub2"); }
% cat sub3.c
void sub3() { puts("sub3"); }
% gcc -c sub1.c sub2.c sub3.c
% ls
sub1.c  sub2.c  sub3.c
sub1.o  sub2.o  sub3.o
% ar cr libpriv1.a sub1.o sub2.o sub3.o
% ar t libpriv1.a
sub1.o
sub2.o
sub3.o
%
```

➡ `puts()` is unresolved in `libpriv1.a`



Using a Static Library


```
% cat prog.c
int main() {
    sub1();
    sub2();
    sub3();
}
% gcc -o prog prog.c -L. -lpriv1
```

➡ Where does `puts()` come from?

```
% gcc -o prog prog.c -L. -lpriv1 -L/lib -lc
```

➡ The order of the libraries matter

- will try to resolve references first in the `priv1` library (either `libpriv1.a` or `libpriv1.so`) and then in the `c` library (either `libc.a` or `libc.so`)




Substitution

➡ Want to use my version of `puts()` instead of what's in the C library

```
% cat myputs.c
int puts(char *s) {
    write(1, "My puts: ", 9);
    write(1, s, strlen(s));
    write(1, "\n", 1);
    return 1;
}
% gcc -c myputs.c
% ar cr libmyputs.a myputs.o
% gcc -o prog prog.c -L. -lpriv1 -lmyputs
```

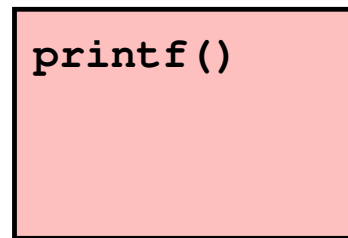
— will try to resolve `puts()` first in the `priv1` library, then in the `myputs` library, then in the `c` library



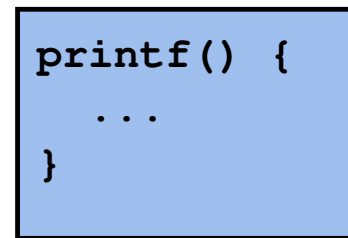
Shared Libraries

- ➡ prog must contain *everything* needed for execution
 - ▬ duplicate code, may be lots of duplicate code, e.g., `printf()`
 - take up *disk space*
 - take up *memory*

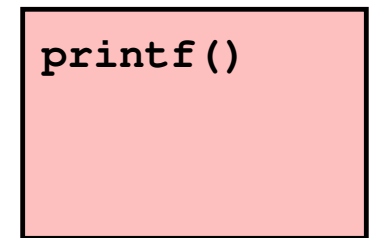
- ➡ Need a way to share things like `printf()`
 - ▬ on disk, and
 - ▬ in memory



Process A




C Library





Process B


- ➡ If `printf()` required no relocation, then it's easy
 - ▬ just make sure `ld` use the right address consistently
- ➡ If `printf()` required relocation, then it's more complicated
 - ▬ problem: processes want a shared function to be at different addresses





Sharing




Sharing





➡ Looks like it can work



Sharing




➡ What about this?




Relocation and Shared Libraries

➡ Approaches

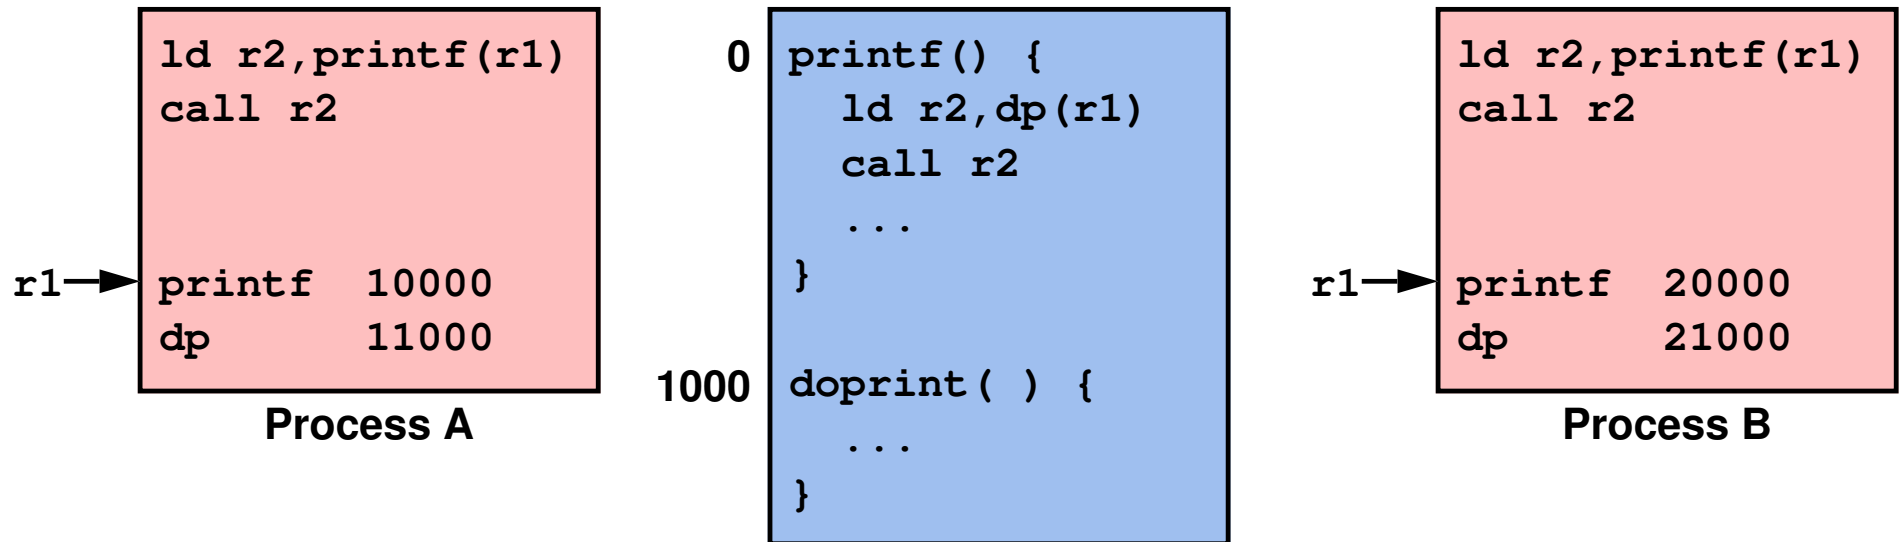
- Limited sharing: relocate separately for each process
 - have a single copy of `printf()` on disk
 - as `printf()` gets copied into memory, perform relocation
 - this would work, but still end up with too many copies of `printf()` in memory
- **Prerelocation:** relocate libraries ahead of time
 - difficult to prerelocate all shared functions
 - ◆ may need to preform rerelocation



- **Position-Independent Code:** no need for relocation
 - producing code that can be placed anywhere in memory without requiring modification
 - need **indirection**




Position-Independent Code




- ▢ each process maintains a private table, pointed to by register `r1`
 - table contains addresses of shared routines
- ▢ don't call functions directly
 - make a position-independent call (i.e., an indirect call)
 - i.e., call the function located at a *fixed index* into the table pointed by `r1`
 - ◆ implemented as two instructions in the above example

➡ Please note that `ld` is not the same as `mov` in x86 CPU




Position-Independent Code Details

- ➡ Processor-dependent; x86 32-bit version:
- ➡ ELF requires 3 data structures for each dynamic executable and shared object
 - ⇒ the *procedure linkage table (PLT)*
 - read-only executable code, *shared* by all processes
 - essentially stubs for calling subroutines
 - ⇒ the *global offset table (GOT)*
 - read-write data, private (to each process)
 - relocated dynamically for each process
 - ⇒ the *dynamic structure*
 - read-only data, *shared* by all processes
 - contains relocation info and symbol table



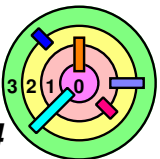
Shared Libraries In Practice

- ➡ **Shared libraries** are used extensively in many modern systems
 - often implemented with either prereslocation or position-independent code
 - in Windows, they are known as **Dynamic-Link Libraries (DLLs)**
 - in Unix, they are known as **shared objects** (.so files)
 - vs. **static libraries** (.a files)
 - they need not be loaded when a program starts up
 - can be loaded when needed, i.e., **on-demand**
 - this way, the startup time of a program may be reduced
- ➡ Disadvantages of DLLs and shared objects
 - they can have **dependencies**
 - different **versions** of the same library




Linking and Loading on Linux with ELF

- ➡ **x86 ELF (Executable and Linking Format)**
 - used in Unix/Linux systems
 - not used in either MacOS X or Windows
- ➡ **Creating and using a shared library**
- ➡ **Substitution**
- ➡ **Shared library details**
- ➡ **Versioning**
- ➡ **Dynamic linking**
- ➡ **Interpositioning**



Creating a Shared Library (1)


```
% gcc -fPIC -c myputs.c
% ld -shared -o libmyputs.so myputs.o
% gcc -o prog prog.c -L. -lpriv1 -lmyputs
% ./prog
./prog: error while loading shared libraries:
libmyputs.so: cannot open shared object file: No
such file or directory
% ldd prog
libmyputs.so => not found
libc.so.6 => /lib/tls/i686/cmov/libc.so.6
/lib/ld-linux.so.2 => /lib/ld-linux.so.2
```



Creating a Shared Library (2)

```
% gcc -o prog prog.c -L. -lpriv1 -lmyputs -Wl,-rpath .
% ldd prog
libmyputs.so => ./libmyputs.so
libc.so.6 => /lib/tls/i686/cmov/libc.so.6
/lib/ld-linux.so.2 => /lib/ld-linux.so.2
% ./prog
My puts: sub1
My puts: sub2
My puts: sub3
```


- `"-Wl,-rpath ."` means that what comes after `-Wl` are linker options (i.e., pass them to the linker)
 - in this example, the linker will be invoked with `"-rpath ."`
- also try `"-Wl,-rpath, ."` if the space character is giving you trouble




Versioning

```
% gcc -fPIC -c myputs.c
% ld -shared -soname libmyputs.so.1 \
    -o libmyputs.so.1.0 myputs.o
% ldconfig -v -n .
% ln -s libmyputs.so.1 libmyputs.so
% gcc -o prog1 prog1.c -L. -lpriv1 -lmyputs \
    -Wl,-rpath .
% vi myputs.c
% gcc -fPIC -c myputs.c
% ld -shared -soname libmyputs.so.2 \
    -o libmyputs.so.2 myputs.o
% rm -f libmyputs.so
% ldconfig -v -n .
% ln -s libmyputs.so.2 libmyputs.so
% gcc -o prog2 prog2.c -L. -lpriv1 -lmyputs \
    -Wl,-rpath .
```


- ▢ "libmyputs.so.1" is the *soname*
- ▢ "libmyputs.so.1.0" is the *real name*
- ▢ "libmyputs.so" is the *linker name*




Interpositioning



— prog thinks it's calling `puts ()`



Interpositioning



- ▬ prog thinks it's calling `puts()`
- ▬ interpose your `puts()`
 - you can call the original `puts()` that prog thought it was calling
 - security problem?!
 - ◆ "DLL injection" if you pick up a DLL unknowingly

How To ...


```
% cat myputs.c
#include <dlfcn.h>

int puts(const char *s) {
    int (*pptr)(const char *);

    pptr = (int (*)(const char*)) dlsym(RTLD_NEXT, "puts");

    write(2, "calling myputs: ", 16);
    return (*pptr)(s);
}
```


- ⇒ `dlsym()` returns a *function pointer* for the named function
- ⇒ `RTLD_NEXT` asks for the next occurrence of the named function
 - `RTLD_DEFAULT` will get you the first occurrence of the named function using the default library search order



Compiling/Linking It


```
% gcc -fPIC -c myputs.c -D_GNU_SOURCE
% ld -shared -soname libmyputs.so.1 \
    -o libmyputs.so.1.0 myputs.o -ldl
% ldconfig -v -n .
% cat tputs.c
int main() {
    puts("This is a boring message.");
    return 0;
}
% gcc -o tputs tputs.c ./libmyputs.so.1 -Wl,-rpath .
% ./tputs
calling myputs: This is a boring message.
%
```

- ▢ `-D_GNU_SOURCE` is needed or won't recognize `RTLD_NEXT`
- ▢ `ldconfig` may be in `/sbin`



What's Going On ...


- ➡ gcc/ld
 - ➡ compiles code
 - ➡ does *static linking*
 - ➡ searches list of libraries
 - ➡ adds *references* to *shared objects*
- ➡ *runtime*
 - ➡ program invokes ld.so (or ld-linux.so) to *finish linking*
 - ➡ maps in shared objects
 - ➡ does relocation and procedure linking as required
 - ➡ dlsym() invokes ld.so to do *more linking*




Delayed Wrapping

- ➡ **LD_PRELOAD**
- environment variable checked by `ld.so`
 - specifies additional shared objects to search (first) when program is started

```
% gcc -o tputs tputs.c
% ./tputs
This is a boring message.
% setenv LD_PRELOAD ./libmyputs.so.1
% ./tputs
calling myputs: This is a boring message.
%
```




Extra Slides




Position-Independent Code Details

- ➡ Processor-dependent; x86 32-bit version:
- ➡ ELF requires 3 data structures for each dynamic executable and shared object
 - ⇒ the *procedure linkage table (PLT)*
 - read-only executable code, *shared* by all processes
 - essentially stubs for calling subroutines
 - ⇒ the *global offset table (GOT)*
 - read-write data, private (to each process)
 - relocated dynamically for each process
 - ⇒ the *dynamic structure*
 - read-only data, *shared* by all processes
 - contains relocation info and symbol table

Global-Offset Table: Data References




- Modules refer to *global variables* indirectly by looking up their addresses in this table
 - a *register* contains the address of the table
 - modules refer to entries via their *offsets*
- When a module is loaded into memory
 - ld.so puts the actual addresses into the GOT



Procedure References

- ➡ More complicated than data references
- ➡ Lots of them
- ➡ Many are never used
- ➡ Fix them up on demand



Before Calling name1

```
.PLT0:
    pushl 4(%ebx)
    jmp    8(%ebx)
    nop; nop
    nop; nop
.PLT1:
    jmp    name1(%ebx)
.PLT1next:
    pushl $name1RelOffset
    jmp    .PLT0
.PLT2:
    jmp    name2(%ebx)
.PLT2next:
    pushl $name2RelOffset
    jmp    .PLT0
```

Procedure-Linkage Table

ebx →

```
_GLOBAL_OFFSET_TABLE:
    .long _DYNAMIC
    .long identification
    .long ld.so

name1:
    .long .PLT1next
name2:
    .long .PLT2next
```

Relocation info:

GOT_offset(name1), symx(name1)


GOT_offset(name2), symx(name2)

_DYNAMIC



Before the first call to name1

- the actual address of the name1 procedure is *not* in the GOT
- first call to name1 invokes ld.so with indication of the above fact
- ld.so finds name1 and update the GOT



Before Calling name1

```
.PLT0:
    pushl 4(%ebx)
    jmp    8(%ebx)
    nop; nop
    nop; nop
.PLT1:
    jmp    name1(%ebx)
.PLT1next:
    pushl $name1RelOffset
    jmp    .PLT0
.PLT2:
    jmp    name2(%ebx)
.PLT2next:
    pushl $name2RelOffset
    jmp    .PLT0
```

Procedure-Linkage Table

ebx →

```
_GLOBAL_OFFSET_TABLE:
    .long _DYNAMIC
    .long identification
    .long ld.so

name1:
    .long .PLT1next
name2:
    .long .PLT2next
```

Relocation info:

GOT_offset(name1), symx(name1)


GOT_offset(name2), symx(name2)

_DYNAMIC


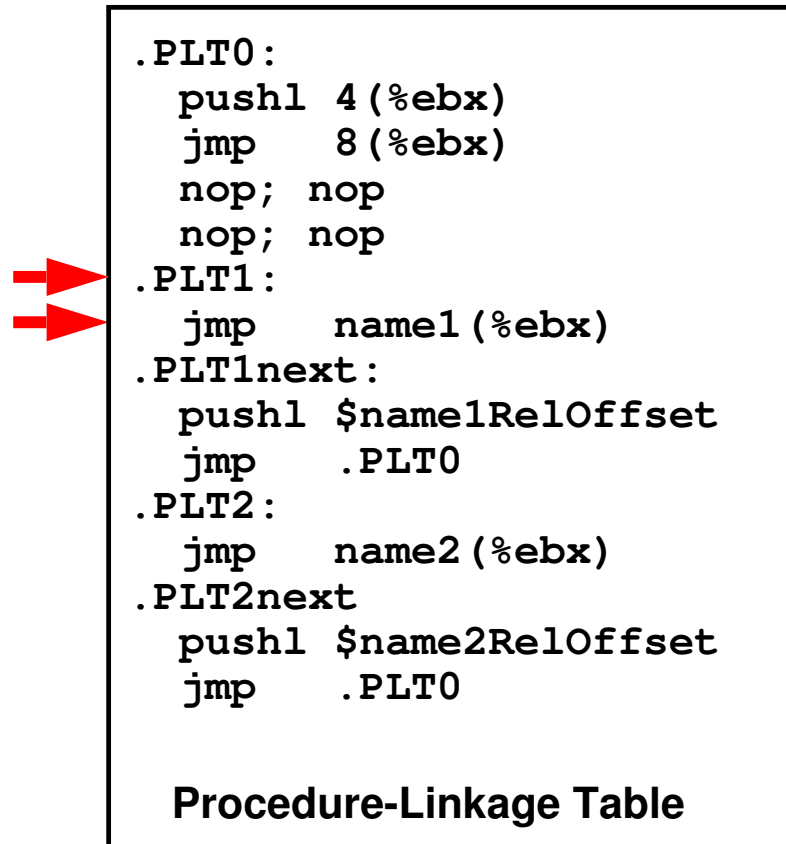


How did this work?

- references from module to name1 are statically linked to entry .PLT1 in the procedure-linkage table
- the index into the symbol table in _DYNAMIC tells us where to start (in this example, it's 0 for name1)



Before Calling name1



Relocation info:

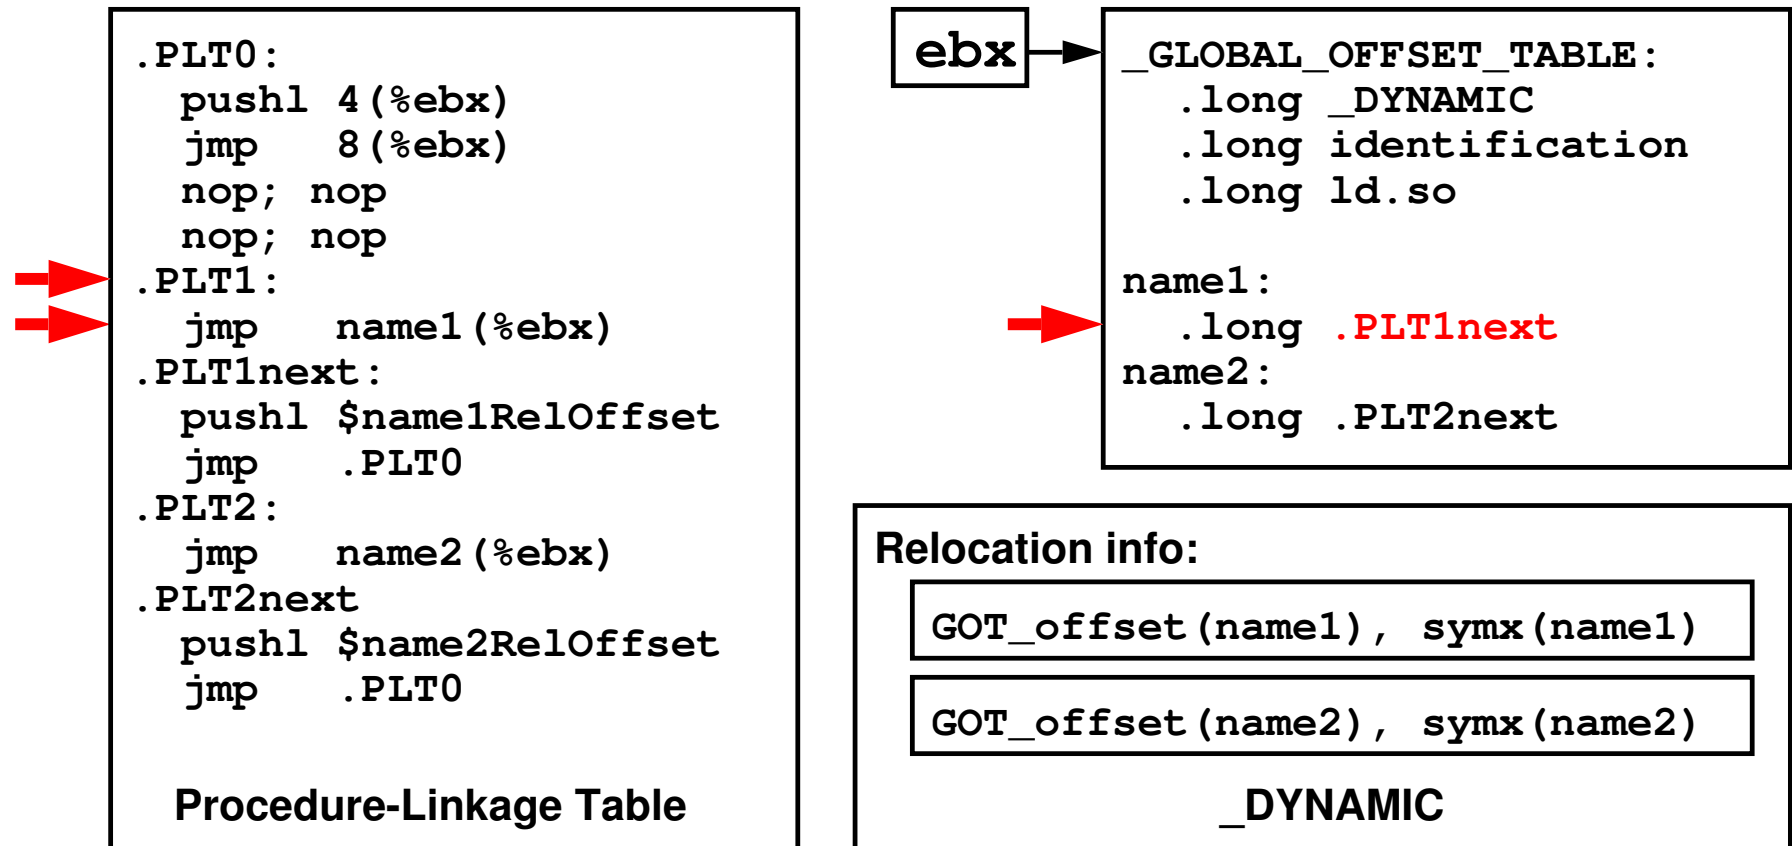
GOT_offset(name1), symx(name1)

GOT_offset(name2), symx(name2)

_DYNAMIC

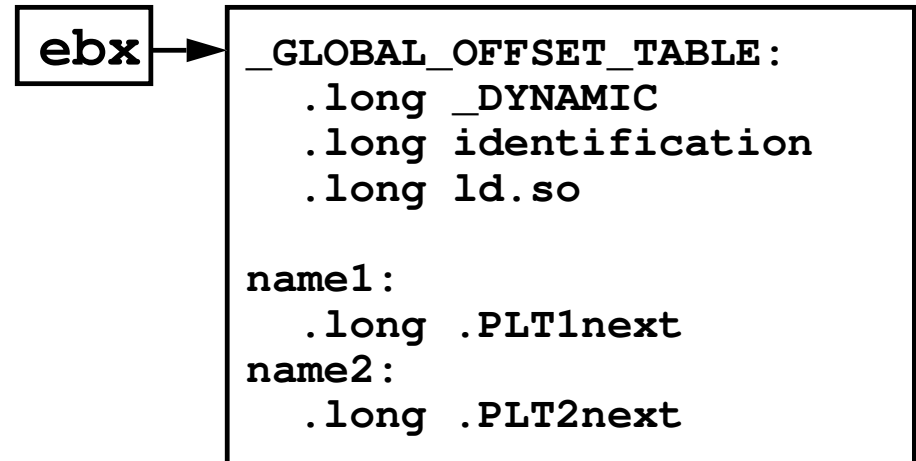
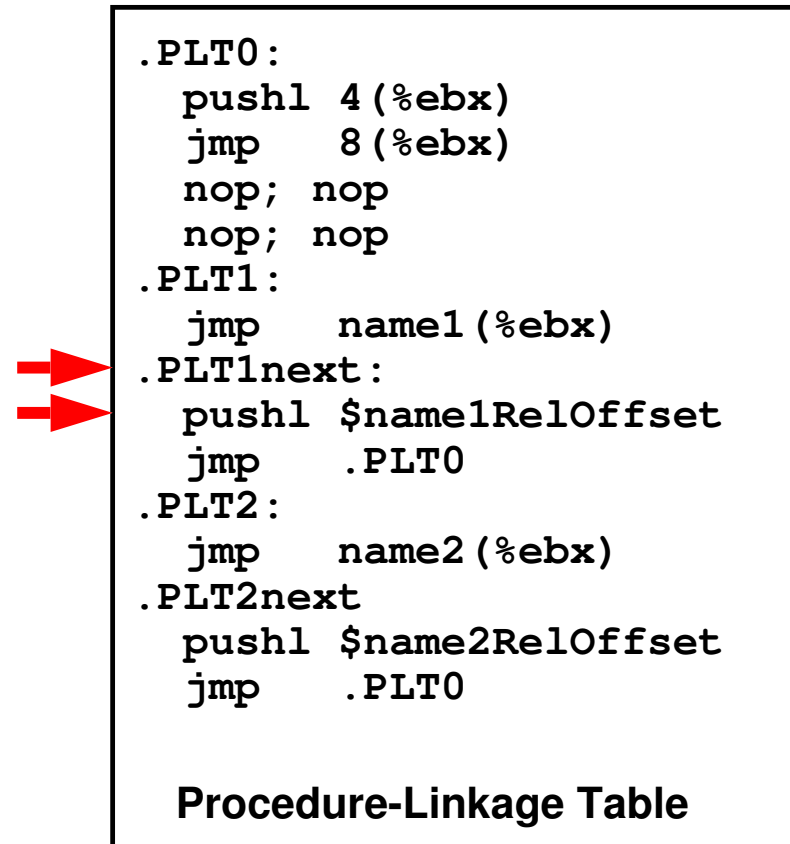
- unconditional jump to the address contained in the name1 offset of the GOT (pointed to by the register ebx)

Before Calling name1



- unconditional jump to the address contained in the `name1` offset of the GOT (pointed to by the register `ebx`)
 - initially, this address is of the instruction following the jump instruction

Before Calling name1




Relocation info:

GOT_offset(name1), symx(name1)

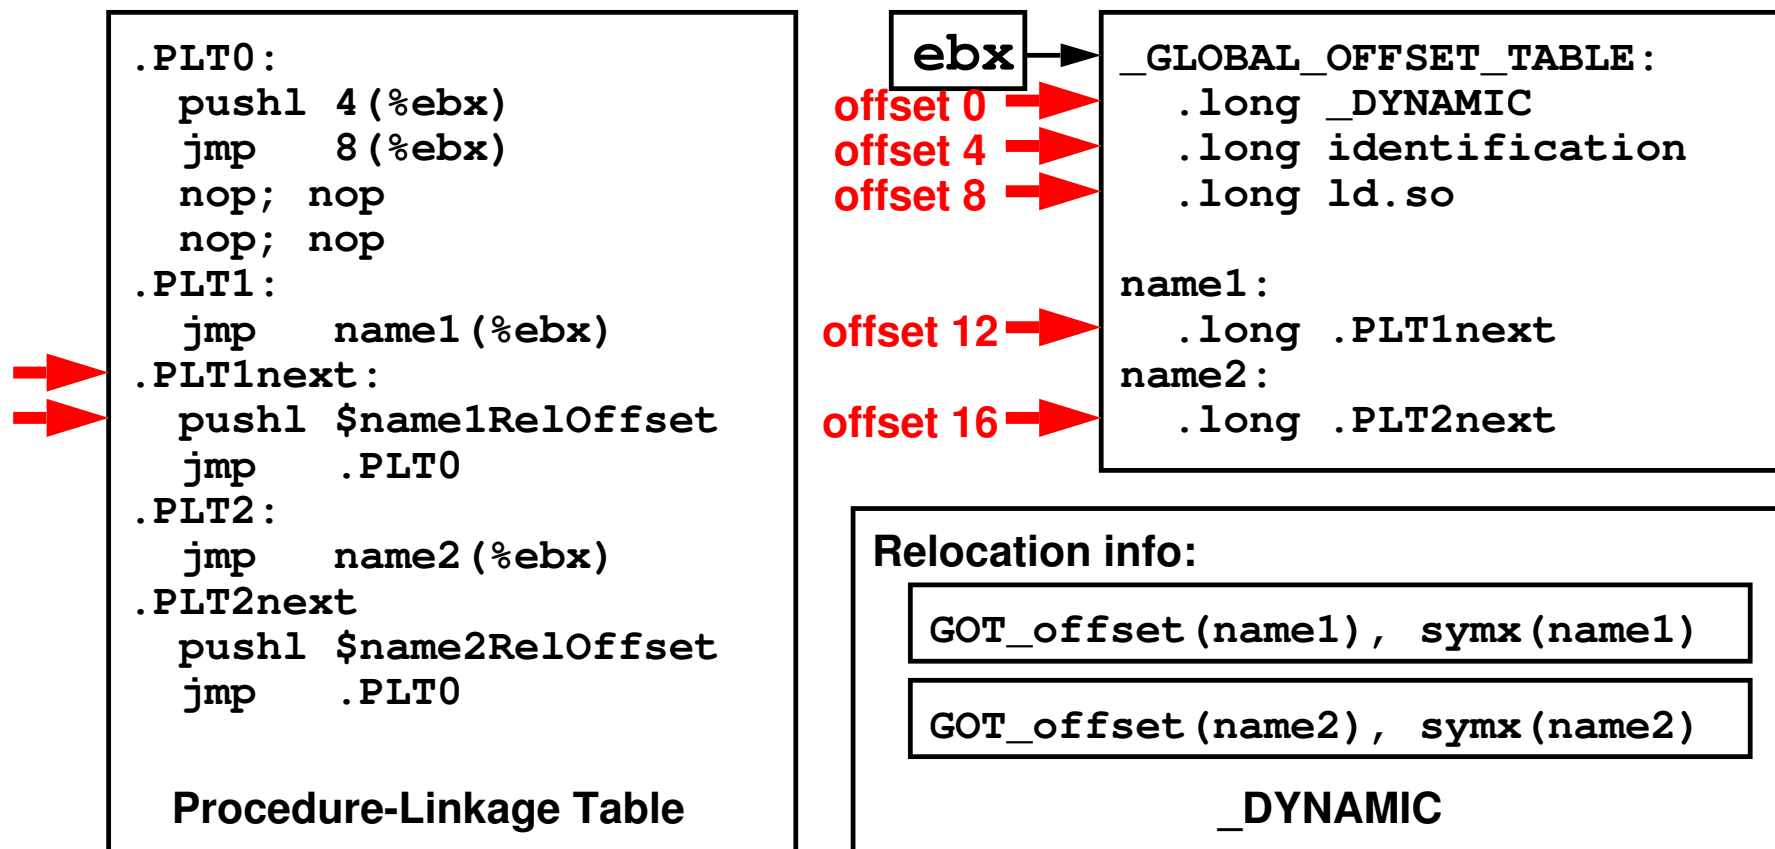
GOT_offset(name2), symx(name2)

_DYNAMIC

- pushd onto the stack the offset of the name1 entry in the relocation table
 - in the above example, this is 12 (if name1 is the 4th entry in the GOT)



Before Calling name1



- pushd onto the stack the offset of the `name1` entry in the relocation table
 - in the above example, this would be 12 (if `name1` is the 4th entry in the GOT)

Before Calling name1

```

.PLT0:
    pushl 4(%ebx)
    jmp 8(%ebx)
    nop; nop
    nop; nop
.PLT1:
    jmp name1(%ebx)
.PLT1next:
    pushl $name1RelOffset
    jmp .PLT0
.PLT2:
    jmp name2(%ebx)
.PLT2next:
    pushl $name2RelOffset
    jmp .PLT0

```

Procedure-Linkage Table

ebx →

```

_GLOBAL_OFFSET_TABLE:
    .long _DYNAMIC
    .long identification
    .long ld.so

name1:
    .long .PLT1next
name2:
    .long .PLT2next

```

Relocation info:

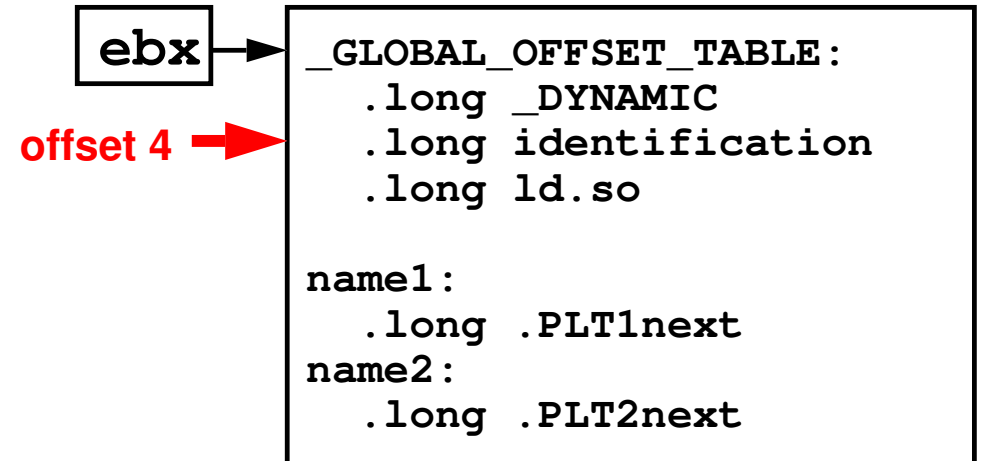
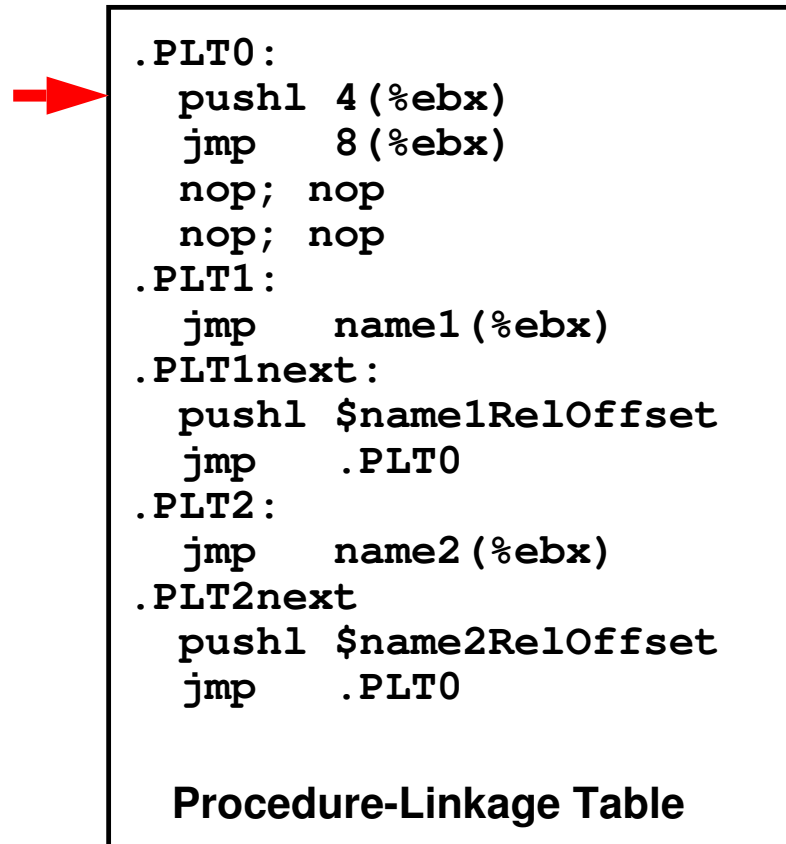
GOT_offset(name1), symx(name1)

GOT_offset(name2), symx(name2)

_DYNAMIC

- unconditional jump to the beginning of the procedure-linkage table, .PLT0

Before Calling name1



Relocation info:

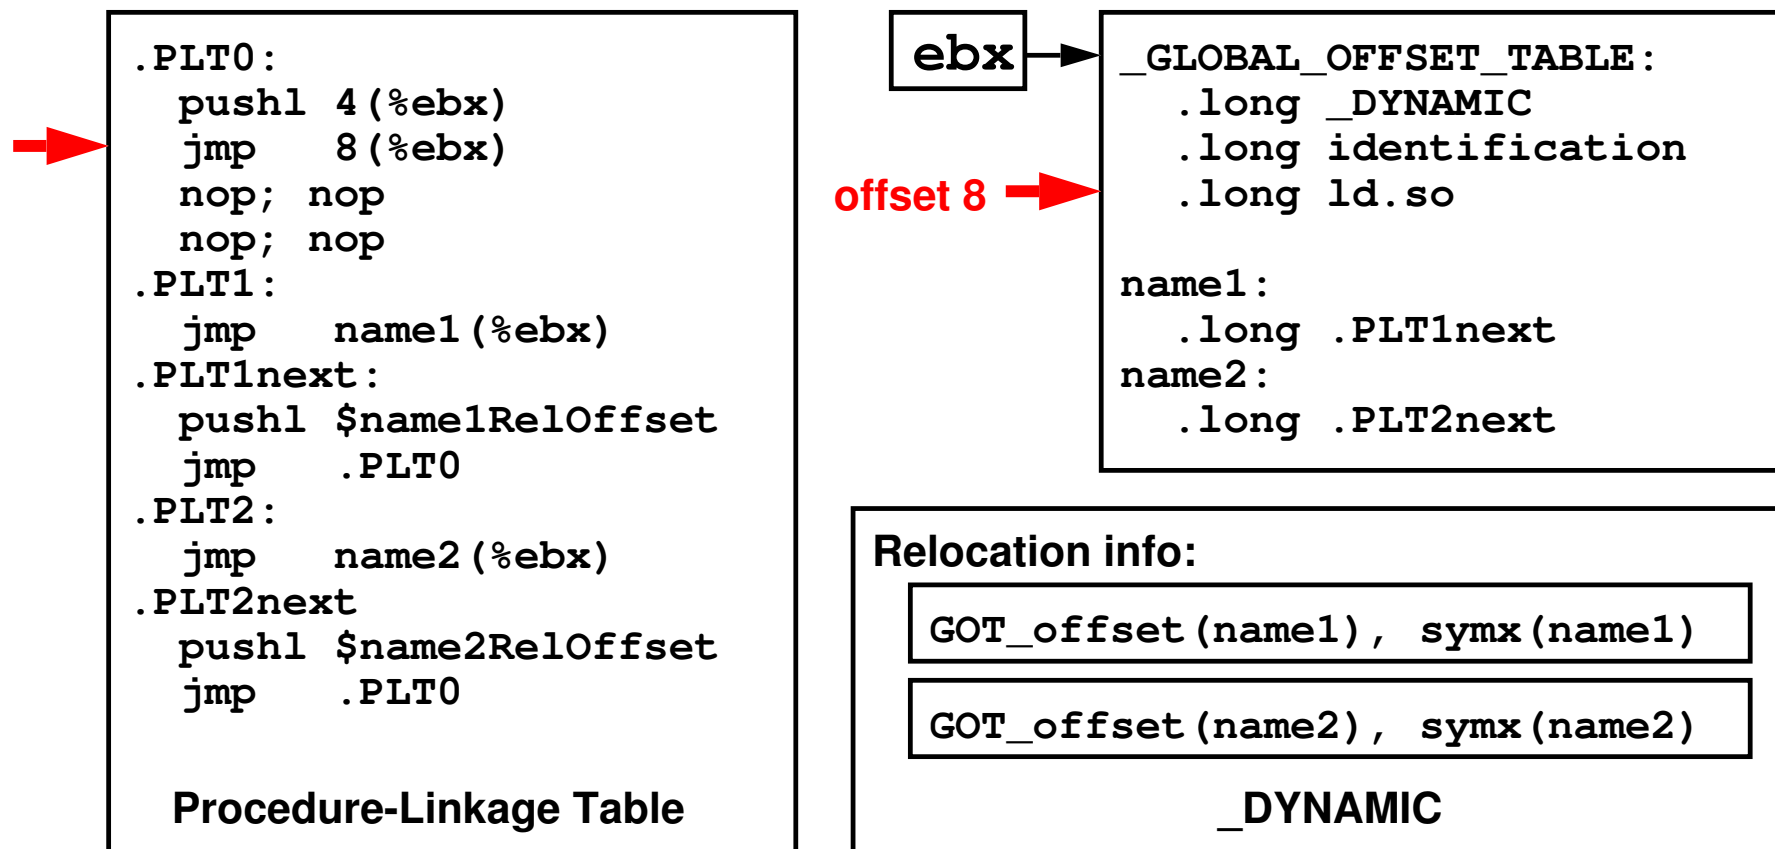
GOT_offset(name1), symx(name1)

GOT_offset(name2), symx(name2)


_DYNAMIC

- pushd the "identification" of the current executing module on to the stack

Before Calling name1



- unconditional jump to `ld.so`
 - `ld.so` can figure out who was making the request (TOS)
 - TOS-4 contains the GOT offset (12 in our example) of where to write the result into
 - TOS-8 contains index into symbol table in `_DYNAMIC`



Before Calling name1

```
.PLT0:
    pushl 4(%ebx)
    jmp    8(%ebx)
    nop; nop
    nop; nop
.PLT1:
    jmp    name1(%ebx)
.PLT1next:
    pushl $name1RelOffset
    jmp    .PLT0
.PLT2:
    jmp    name2(%ebx)
.PLT2next:
    pushl $name2RelOffset
    jmp    .PLT0
```

Procedure-Linkage Table

ebx →

```
_GLOBAL_OFFSET_TABLE:
    .long _DYNAMIC
    .long identification
    .long ld.so

name1:
    .long name1
name2:
    .long .PLT2next
```

Relocation info:


GOT_offset(name1), symx(name1)

GOT_offset(name2), symx(name2)

_DYNAMIC

➡ ld.so writes the actual address of the name1 procedure into the name1 entry of the GOT

— unwinds the stack a bit and then passes control to name1



After Calling name1

```
.PLT0:
    pushl 4(%ebx)
    jmp    8(%ebx)
    nop; nop
    nop; nop
.PLT1:
    jmp    name1(%ebx)
.PLT1next:
    pushl $name1RelOffset
    jmp    .PLT0
.PLT2:
    jmp    name2(%ebx)
.PLT2next:
    pushl $name2RelOffset
    jmp    .PLT0
```

Procedure-Linkage Table

ebx →

```
_GLOBAL_OFFSET_TABLE:
    .long _DYNAMIC
    .long identification
    .long ld.so

name1:
    .long name1
name2:
    .long .PLT2next
```

Relocation info:

GOT_offset(name1), symx(name1)

GOT_offset(name2), symx(name2)

_DYNAMIC

➡ Subsequent call to name1 invokes more directly

