

Design

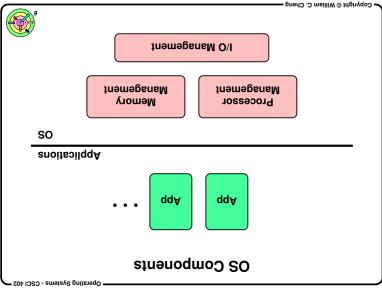
Ch 4: Operating-System

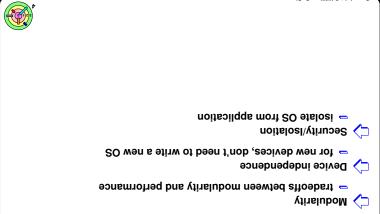
Copyright © William C. Cheng

- generally less efficient design

no support for bit-mapped displays and mice

Early 1980s OS, so we can focus on the basic OS issues



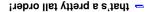




OS Design Issues

Copyright © William C. Cheng 🗕

Performance

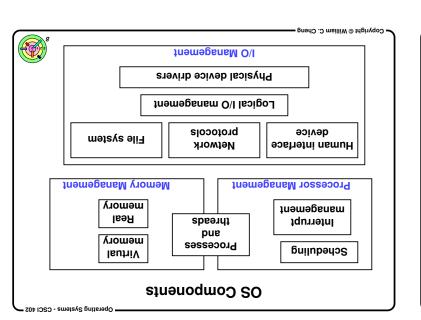


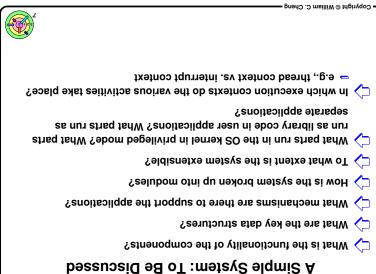
- efficiency of application

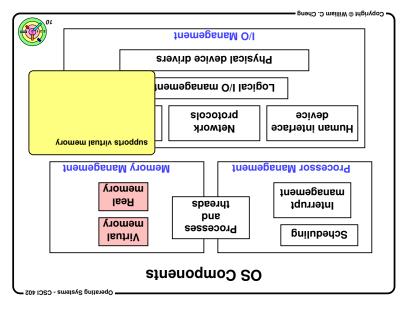
- → while being secure and efficient
- = the OS needs to provide a consistent and usable interface
 - "computer" Applications views the OS as the "computer"
 - eint froqque of bebeen si 20 fahw = We will start with a simple hardware configuration
 - - virtual memory (Ch 7)
 - file systems (Ch 6)
 - Scheduling (Ch 5) We will introduce new components in this chapter

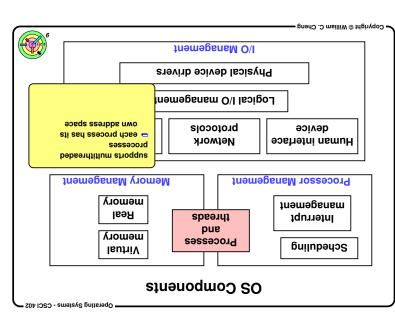
 - how performance concerns are factoered in
 - how is the software structured
 - how they interact with each other
 - SO ns ofni seog tshw = We will now look at how OSes are constructed

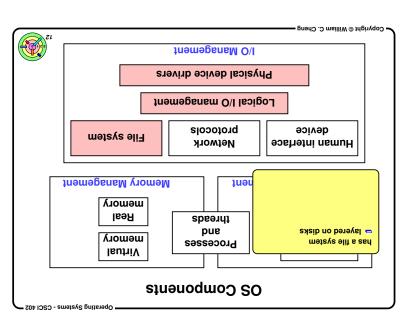


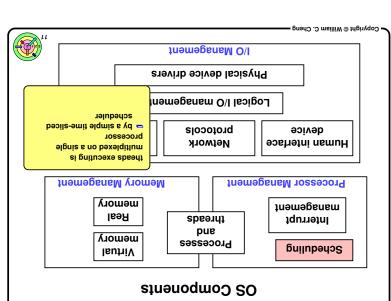


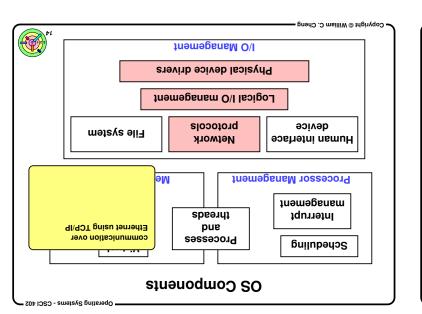


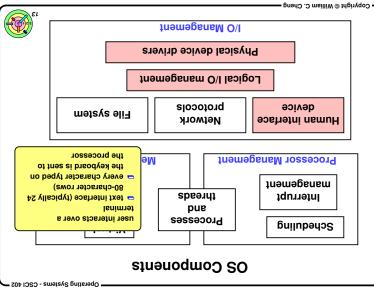


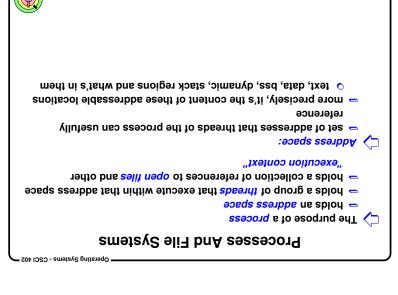


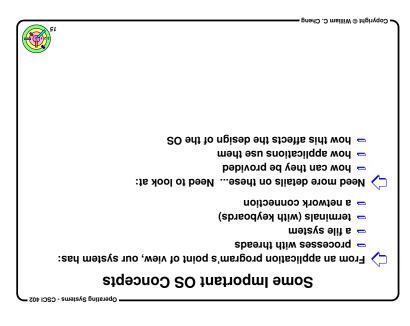


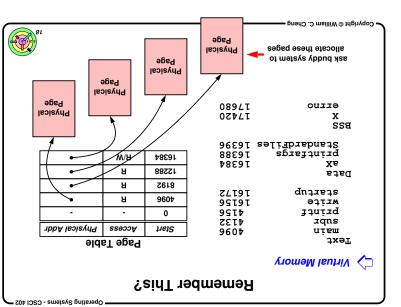


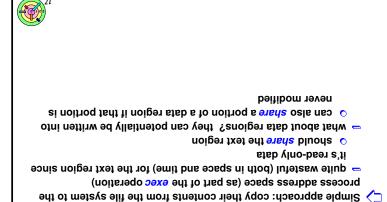










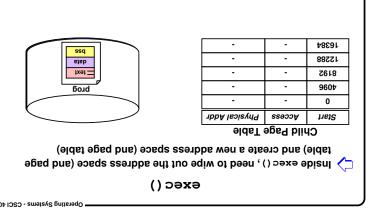


- how should the OS initialize these address space regions?

Processes And File Systems

Copyright © William C. Cheng

:ənssi ngisəd 🔷



Page Physical Page Physical Page Physical Page 16384 16384 12288 12288 26 L R 8195 9601 9601 Physical Addr ssəɔɔ\ Child Page Table Physical Addr esess A Start Parent Page Table Inside fork (), can simply copy parent's page table to child

Processes Can Share Memory Pages

Memory Map

address space in the first place? For the text region, why bother copying the executable file into the

- can just map the file into the address space (Ch 7)
- SO əht ni tqəonoo tnstroqmi ns si **g**niqqsm \, 🔾
- o mapping let the OS tie the regions of the address space to
- the file system

- o if several processes are executing the same program, then address space and files are divided into pieces, called pages

- at most one copy of that program's text page is in memory
- setup, using hardware address translation facilities, to = text regions of all processes running this program are
- this type of mapping is known as shared mapping share these pages



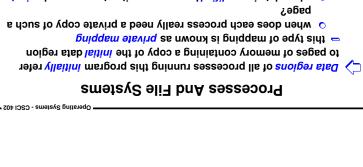
16384

76 L8

9601

SSACOSS

Child Page Table



the kernel also uses memory map to keep track of the mapping

atab

- when data is modified by a process, it gets a new and private



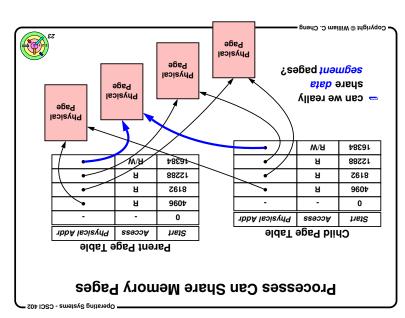
from virtual pages to physical pages

Physical Addr

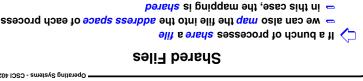
The kernel uses a memory map to keep track of the mapping

Memory Map

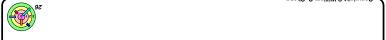
from virtual pages to file pages

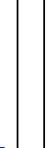






- we can also map the file into the address space of each process
- when one process modifies a page, no private copy is made
- o instead, the original page itself is modified
- and changes are written back to the file everyone gets the changes
- O more on issues in Ch 6





- their pages are initialized, with zeros; copy-on-write mapping 6 The dynamic/heap and stack regions use a special form of private modified are copied the basic idea is that only those pages of memory that are

- a process gets a private copy of the page after a thread in the

Copy-On-Write

o managed by anonymous objects in weenix objects in weenix these are known as anonymous pages

process performs a write for the first time

Copy-on-write (COW):



block/page is the basic unit = block/page is the blo Mapping files into address space is one way to perform I/O on files

Block I/O vs. Sequential I/O

this is referred to as block I/O

Some devices cannot be mapped into the address space

a message via a network connection - e.g., receiving characters typed into the keyboard, sending

snch as read() and write() need a more traditional approach using explicit system calls

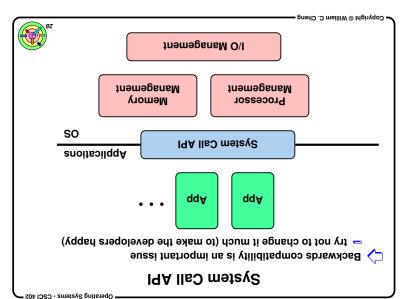
O/l leitneyers as sequential I/O

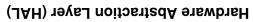
the keyboard It also makes sense to be able to read a file like reading from

or write it out to a network connection should be able to use the same code to write output to a file = similarly, a program that produces lines of text as output

Portability

makes life easier! (and make code more robust)





different code to configure interrupts, hardware timers, etc. e.g., different manufacturers for x86 machines will require Portability across machine configuration

- e.g., may need additional code for context switching, system Portability across processor families

With a well-defined Hardware Abstraction Layer, most of the OS is calls, interrupting handler, virtual memmory management, etc.

relink with the kernel eaniting new HAL routines porting an OS to a new computer is done by machine and processor independent

For a monolithic OS, it is achieved through the use of a portable across various hardware platforms Ti is desirable to have a portable operating system

processor-specific operations within the kernel a portable interface to machine configuration and Hardware Abstraction Layer (HAL)

Hardware Hardware Abstraction Layer System Call API Applications



o independent name space (i.e., named independently from

Devices

Oberating Systems - CSCI 402



A Framework for Devices

Low-level Kernel (will come back to talk about this after Ch 7)

Processes & Threads

Storage Management (will come back to talk about this after Ch 5)

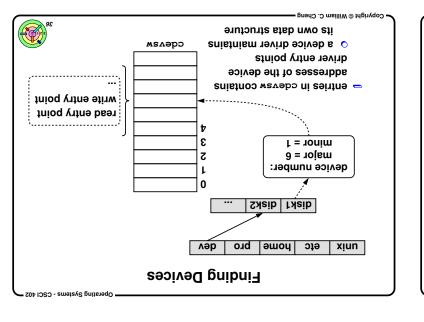
e two choices Device naming

device discovery achice independence Challenges in supporting devices

devices are named as files other things in the system)

"tty" A Oberating Systems - CSCI 402





A Framework for Devices

Device driver:

a major device number - identifies the device driver actually a pair of numbers - every device is identified by a device "number", which is

managed by the same device driver a minor device number - device index for all devices

Special entries were created in the file system to refer to devices

- usually in the /dev directory

• e.g., /dev/diskl, /dev/disk2 each marked as a special file

a special file does not contain data

it refers to devices by their major and minor device

o if you do "1s -1", you can see the device numbers ♦

(character device switch) - statically allocated array in the kernel called cdevsw

Data structure in the early Unix systems

Device Drivers in Early Unix Systems

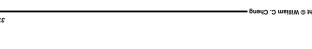
information such as: The kernel was statically configured to contain device-specific

interrupt-vector locations

them all)

- locations of device-control registeres on whatever bus the
- device was attached to

- a kernel must be custom configured for each installation
- Static approach was simple, but cannot be easily extended



members of such a class and choose among them (or use an application can enumerate all currently connected

a device can register itself as members of one or more such

Discovering Devices

- how should the choice be presented to applications?

What about the case where different devices acted similarly?

what about applications? how can they reference

select the appropriate device driver and load into the kernel

once a connected device is identified, system software would

must know how to interact with the USB meta-driver via the

any device that goes onto a USB (Universal Serial Bus)

Step 2: find the needed device drivers and dynamically link

Step 1: discover the device without the benefit of having the

Device Probing

a USB meta-driver is installed into the kernel

a meta-drive handles a particular kind of bus

them into the kernel

Windows has the notion of interface classes

e.g., touchpad on a laptop and USB mouse

dynamically discovered devices?

USB protocol

Solution: use meta-drivers

Sob of gnint their state of the control of the cont

- e.g., USB (Universal Serial Bus)

but how do you achieve this?

relevant device driver in the kernel

Discovering Devices

- eoiton bluow 20 ==
- find a device driver

discovers them

Processes & Threads

A Framework for Devices

- what kind of device is it?
- Where is the driver?

user-level application assigns names based on rules

o downside of this approach is that device naming

lookup the names from a database of names known as devis

In some Linux systems, entries are added into /dev as the kernel

multiple similar devices, but how does application choose?

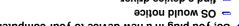
Storage Management (will come back to talk about this after Ch 5)

(Monolithic Kernel)

4.1 A Simple System

Low-level Kernel (will come back to talk about this after Ch 7)





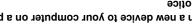
assign a name, but how is it chosen?

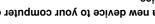
provided by an administrator

some current Linux systems use udev

conventions not universally accepted

o what's an application to do?





So, you plug in a new device to your computer on a particular bus

impractical as the number of supported devices gets big

for a number of similar but not identical installations

This allowed one kernel image to be built that could be useful device-control-register locations

boot time is kind of long

invoked at boot time

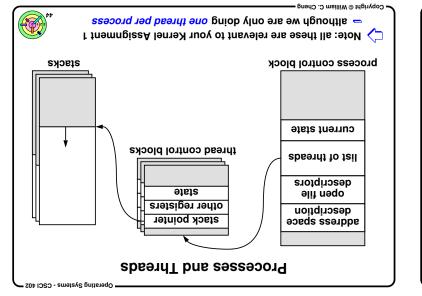
o including identifying and recording interrupt-vector and probe the relevant buses for devices and configure them

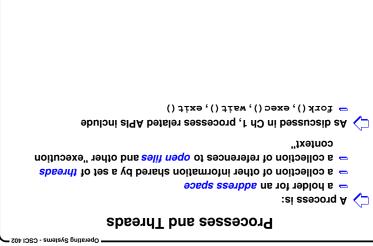
Each device driver includes a probe routine

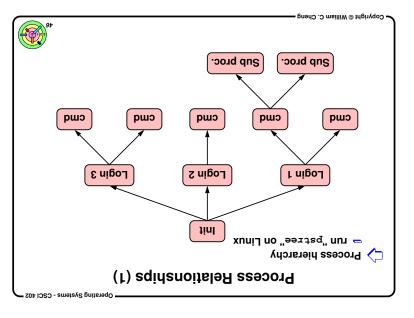
 (still require that a kernel contain all necessary device configured when the system booted allow the devices to to be found and automatically

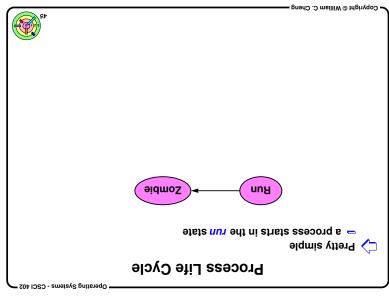
First step to improve the old way

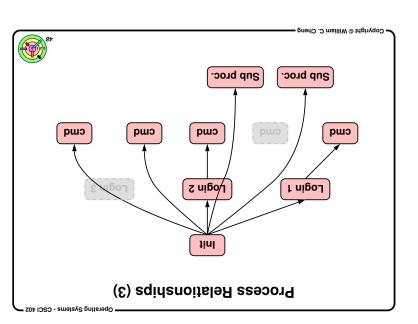
Device Probing

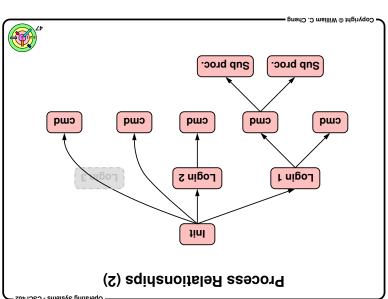


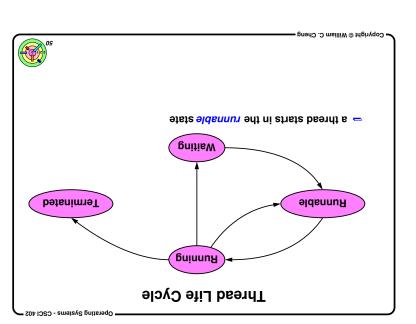


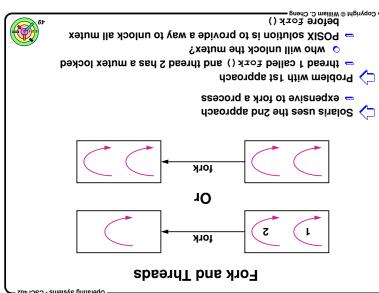


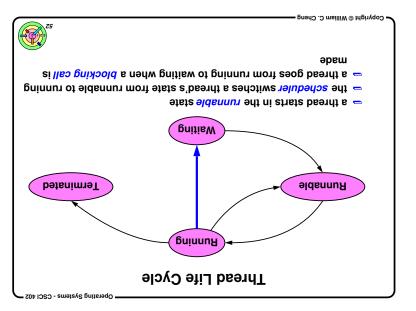


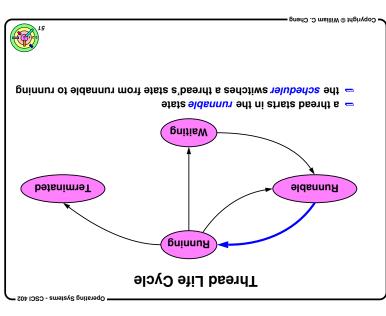


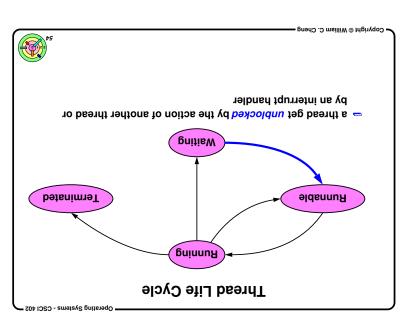


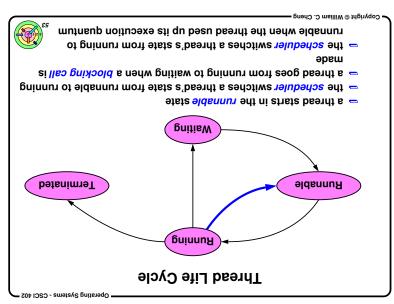








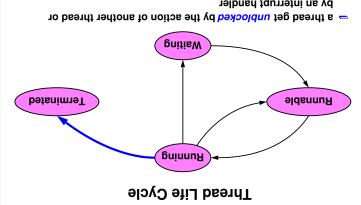




Thread Life Cycle

Does pthread_exit () delete the thread (completely) that calls it?

- What's left in the thread after it calls pthread_exit()?
- o needs to keep thread ID and return code around = its thread control block
- its stack
- o how can a thread delete its own stack? no way!



by an interrupt handler

- be in the running state just before that = in order for a thread to enter the terminated state, it has to
- → what if pthread_cancel() is invoked when the thread

Satate gninnur and ni ton ei

Copyright © William C. Cheng

Thread Life Cycle

the thread's stack space? Who is deleting the thread control block and freeing up

Thread is not detached

- ${\color{blue}\circ}$ the thread that calls ${\tt pthread_join}$ () does the clean up = it can be taken care of in the $pthread_join()$ code
- can do this is one of two ways If a thread is detached (our simple OS does not support this)
- basically doing pthread_join() 1) use a special reaper thread
- scheduler schedule a thread to run) free them when it's convenient (e.g., when the 2) queue these threads on a list and have other threads

