


Ch 4: Operating-System Design

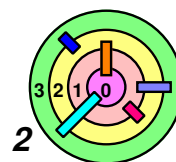
Bill Cheng

<http://merlot.usc.edu/cs402-s16>




OS Design

- ➡ We will now look at how OSes are constructed
 - what goes into an OS
 - how they interact with each other
 - how is the software structured
 - how performance concerns are factored in
- ➡ We will introduce new components in this chapter
 - scheduling (Ch 5)
 - file systems (Ch 6)
 - virtual memory (Ch 7)
- ➡ We will start with a simple hardware configuration
 - what OS is needed to support this
- ➡ Applications views the OS as the "computer"
 - the OS needs to provide a *consistent* and *usable interface*
 - while being *secure* and *efficient*
 - that's a pretty tall order!




OS Design

- ➡ Our goal is to *build a general-purpose OS*
- can run a variety of applications
 - some are interactive
 - many use network communication
 - all read/write to a file system
 - it's like most general-purpose OSes
 - Linux
 - Solaris
 - FreeBSD
 - Mac OS X
 - Chromium OS (has a Linux kernel)
 - Windows (the only one that's not directly based on Unix)
 - all these OSes are quite similar, functionally!
they all provide:
 - processes
 - threads
 - file systems
 - network protocols with similar APIs
 - user interface with display, mouse, keyboard
 - access control based on file ownership and that file owners can control

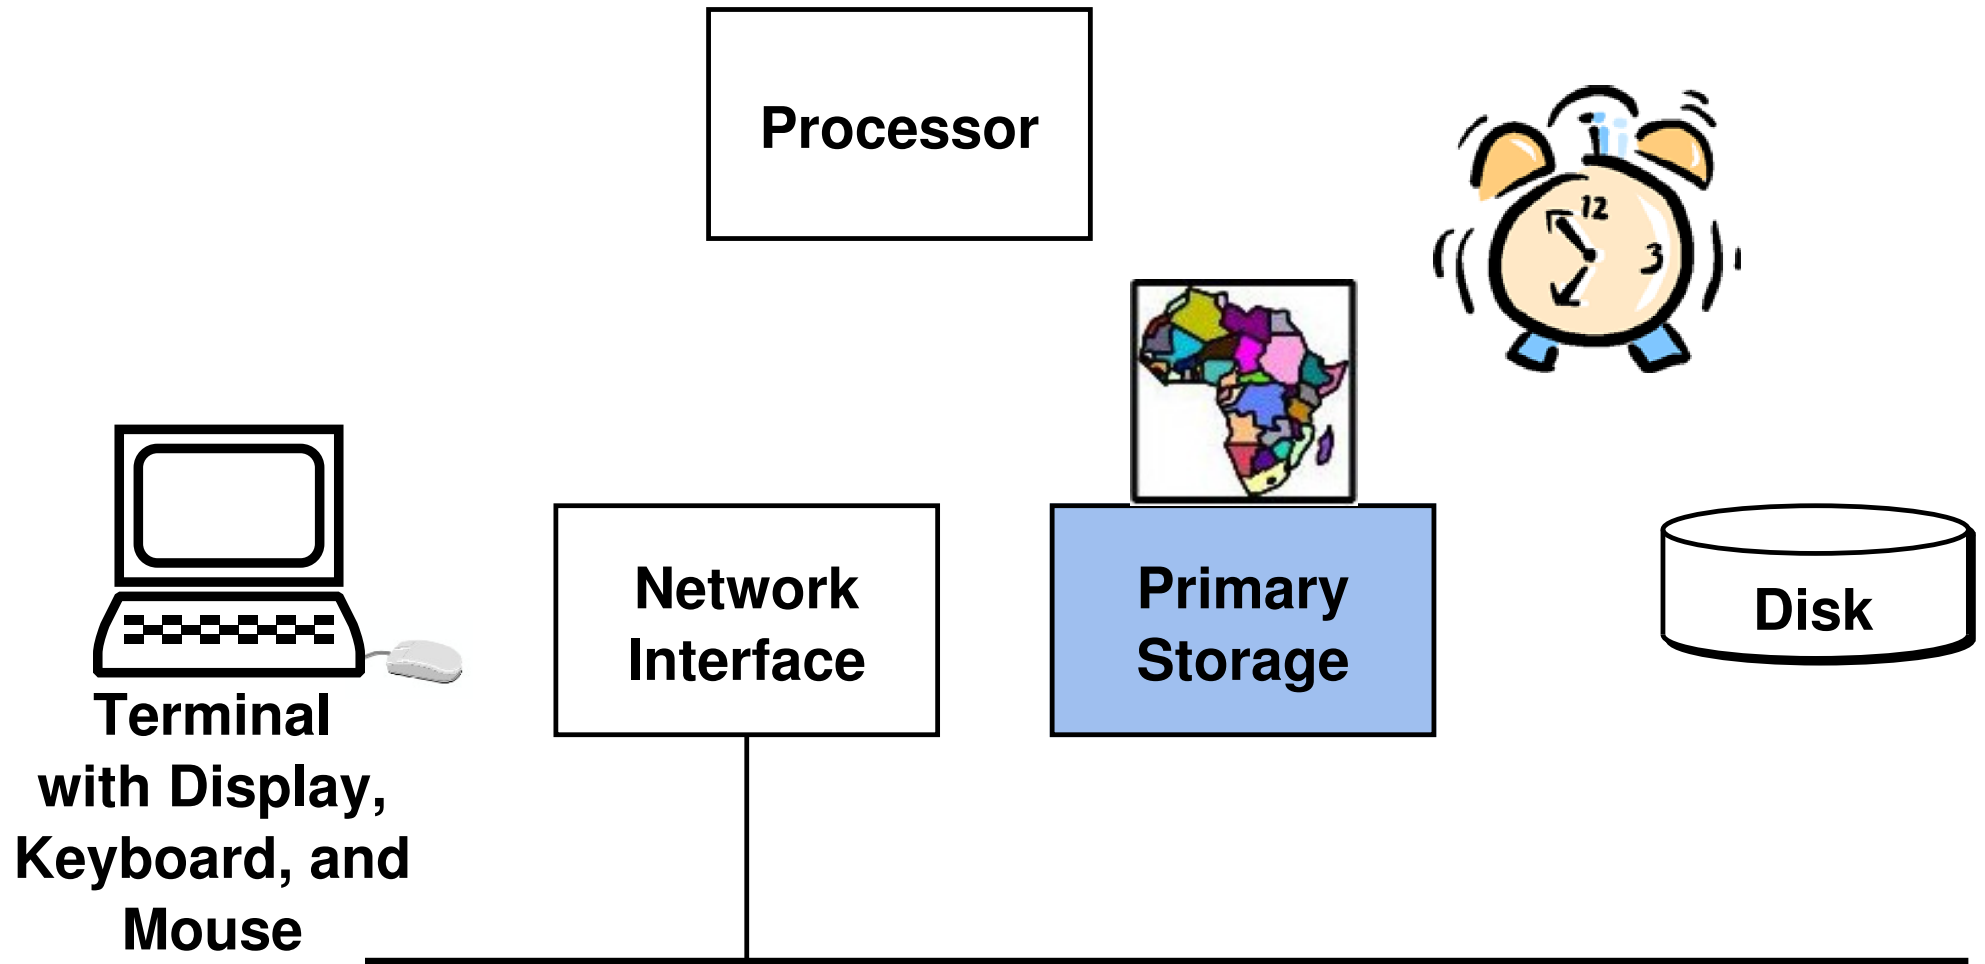


OS Design Issues


- ➡ **Performance**
 - efficiency of application
- ➡ **Modularity**
 - tradeoffs between modularity and performance
- ➡ **Device independence**
 - for new devices, don't need to write a new OS
- ➡ **Security/Isolation**
 - isolate OS from application





Simple Configuration




- ➡ Early 1980s OS, so we can focus on the basic OS issues
- no support for bit-mapped displays and mice
 - generally *less efficient* design



OS Components




...


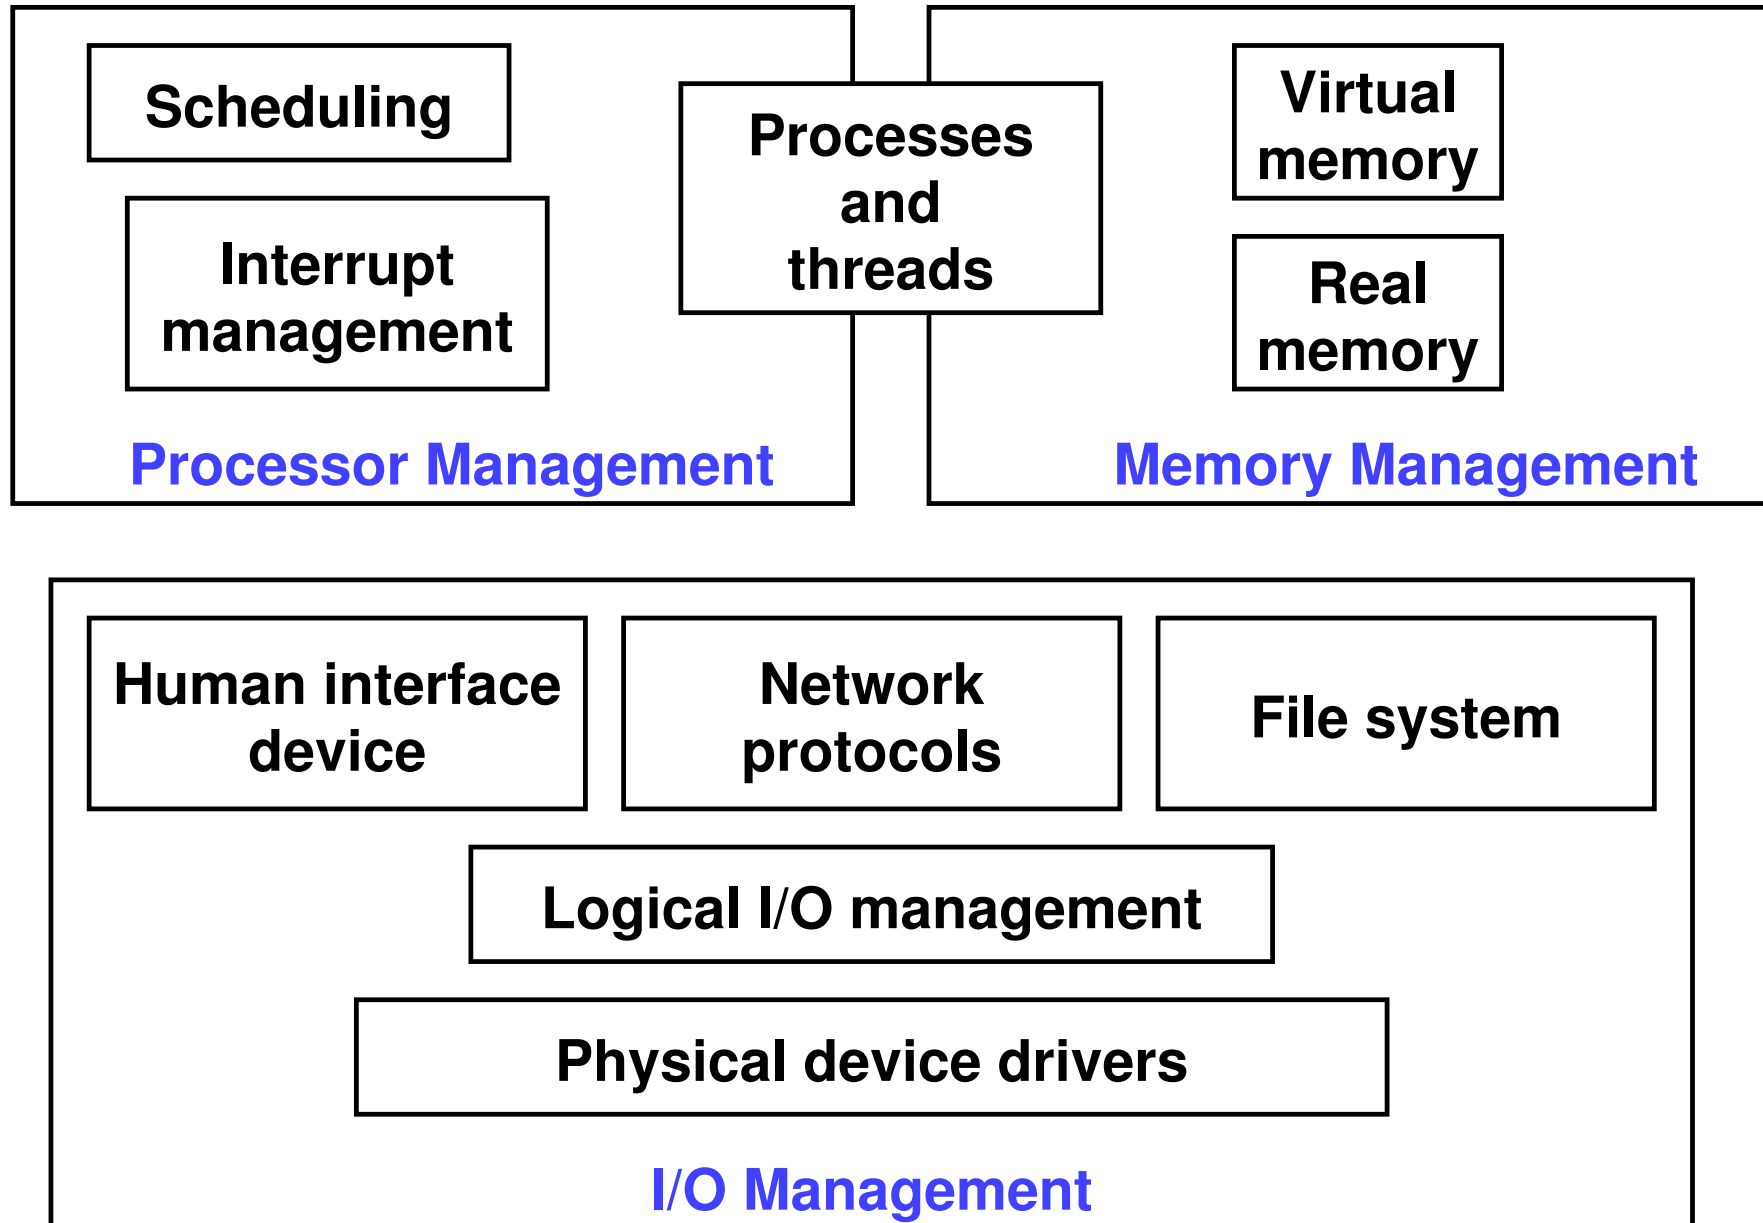
Applications**OS****Processor
Management****Memory
Management****I/O Management**

A Simple System: To Be Discussed


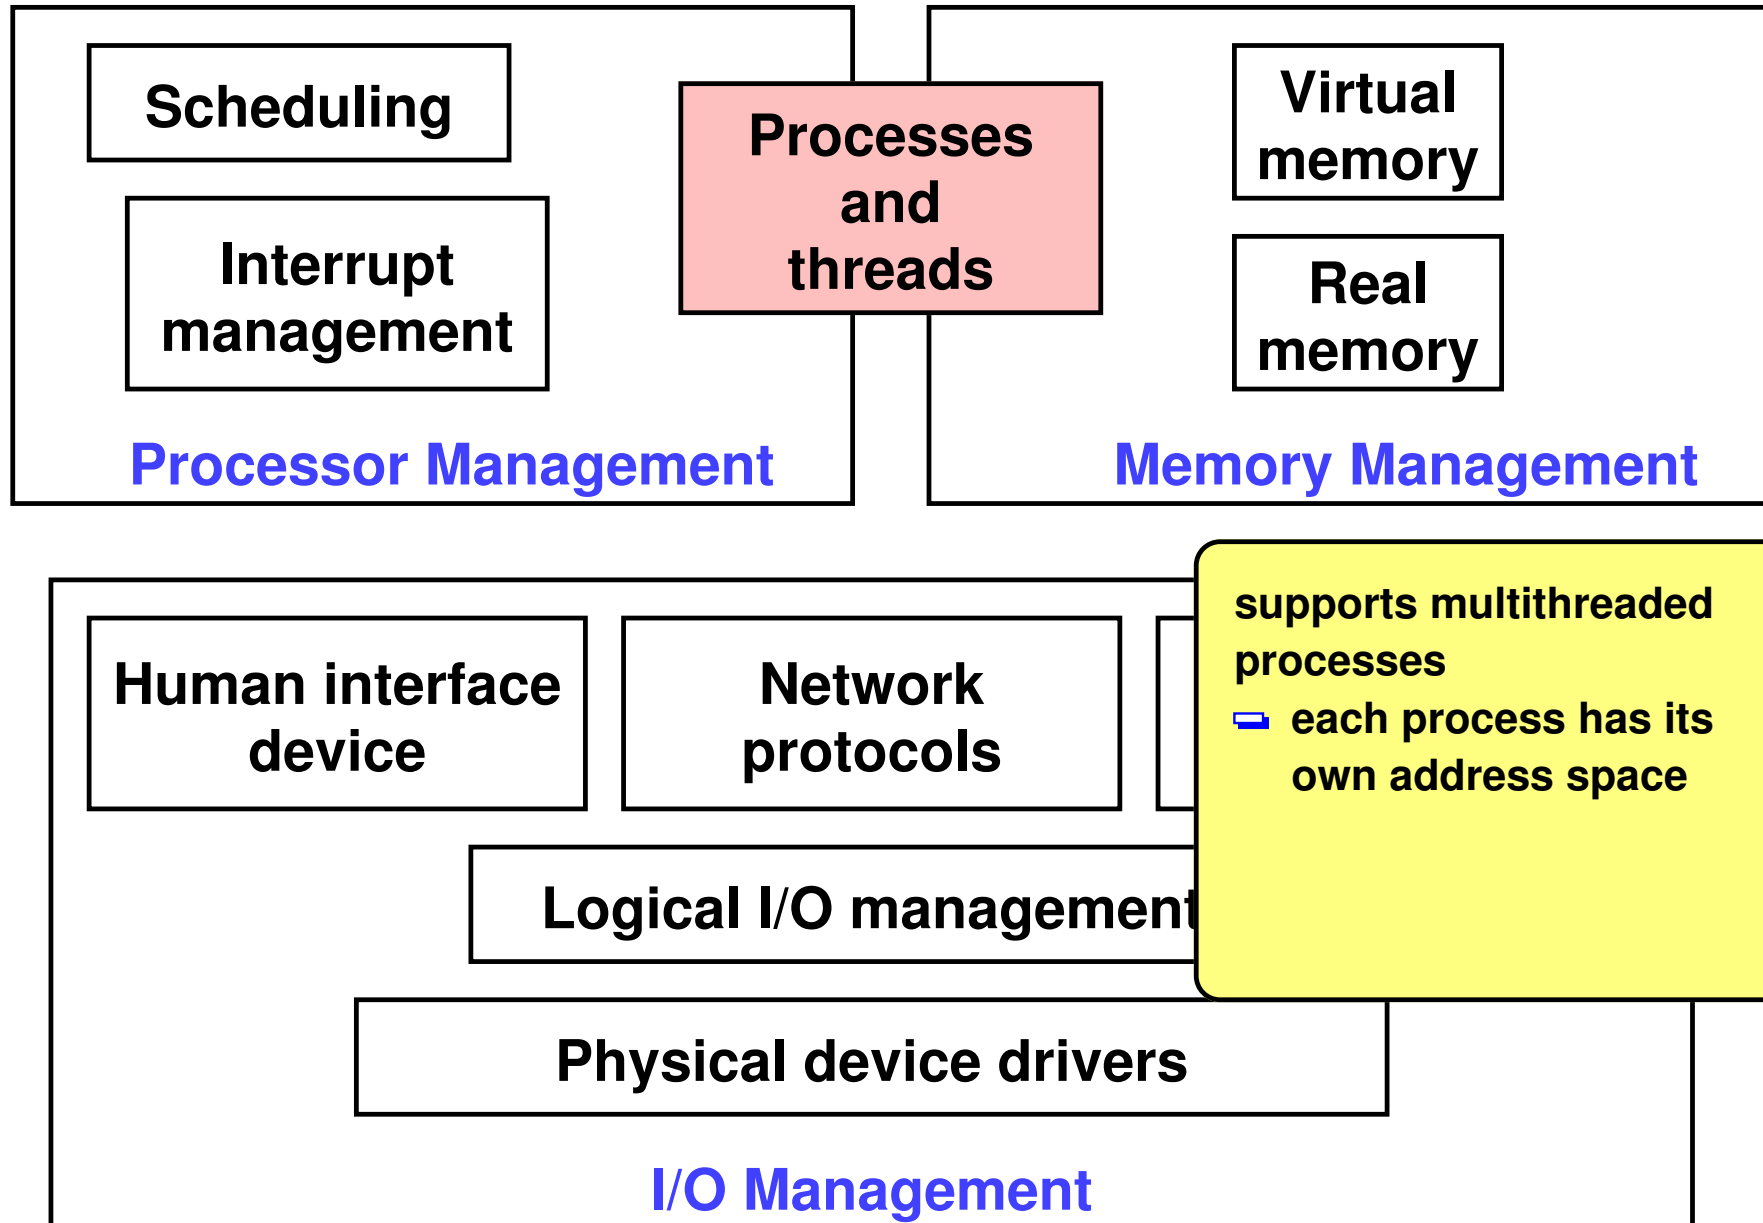
- ➡ What is the functionality of the components?
- ➡ What are the key data structures?
- ➡ What mechanisms are there to support the applications?
- ➡ How is the system broken up into modules?
- ➡ To what extent is the system extensible?
- ➡ What parts run in the OS kernel in privileged mode? What parts run as library code in user applications? What parts run as separate applications?
- ➡ In which execution contexts do the various activities take place?
 - e.g., thread context vs. interrupt context



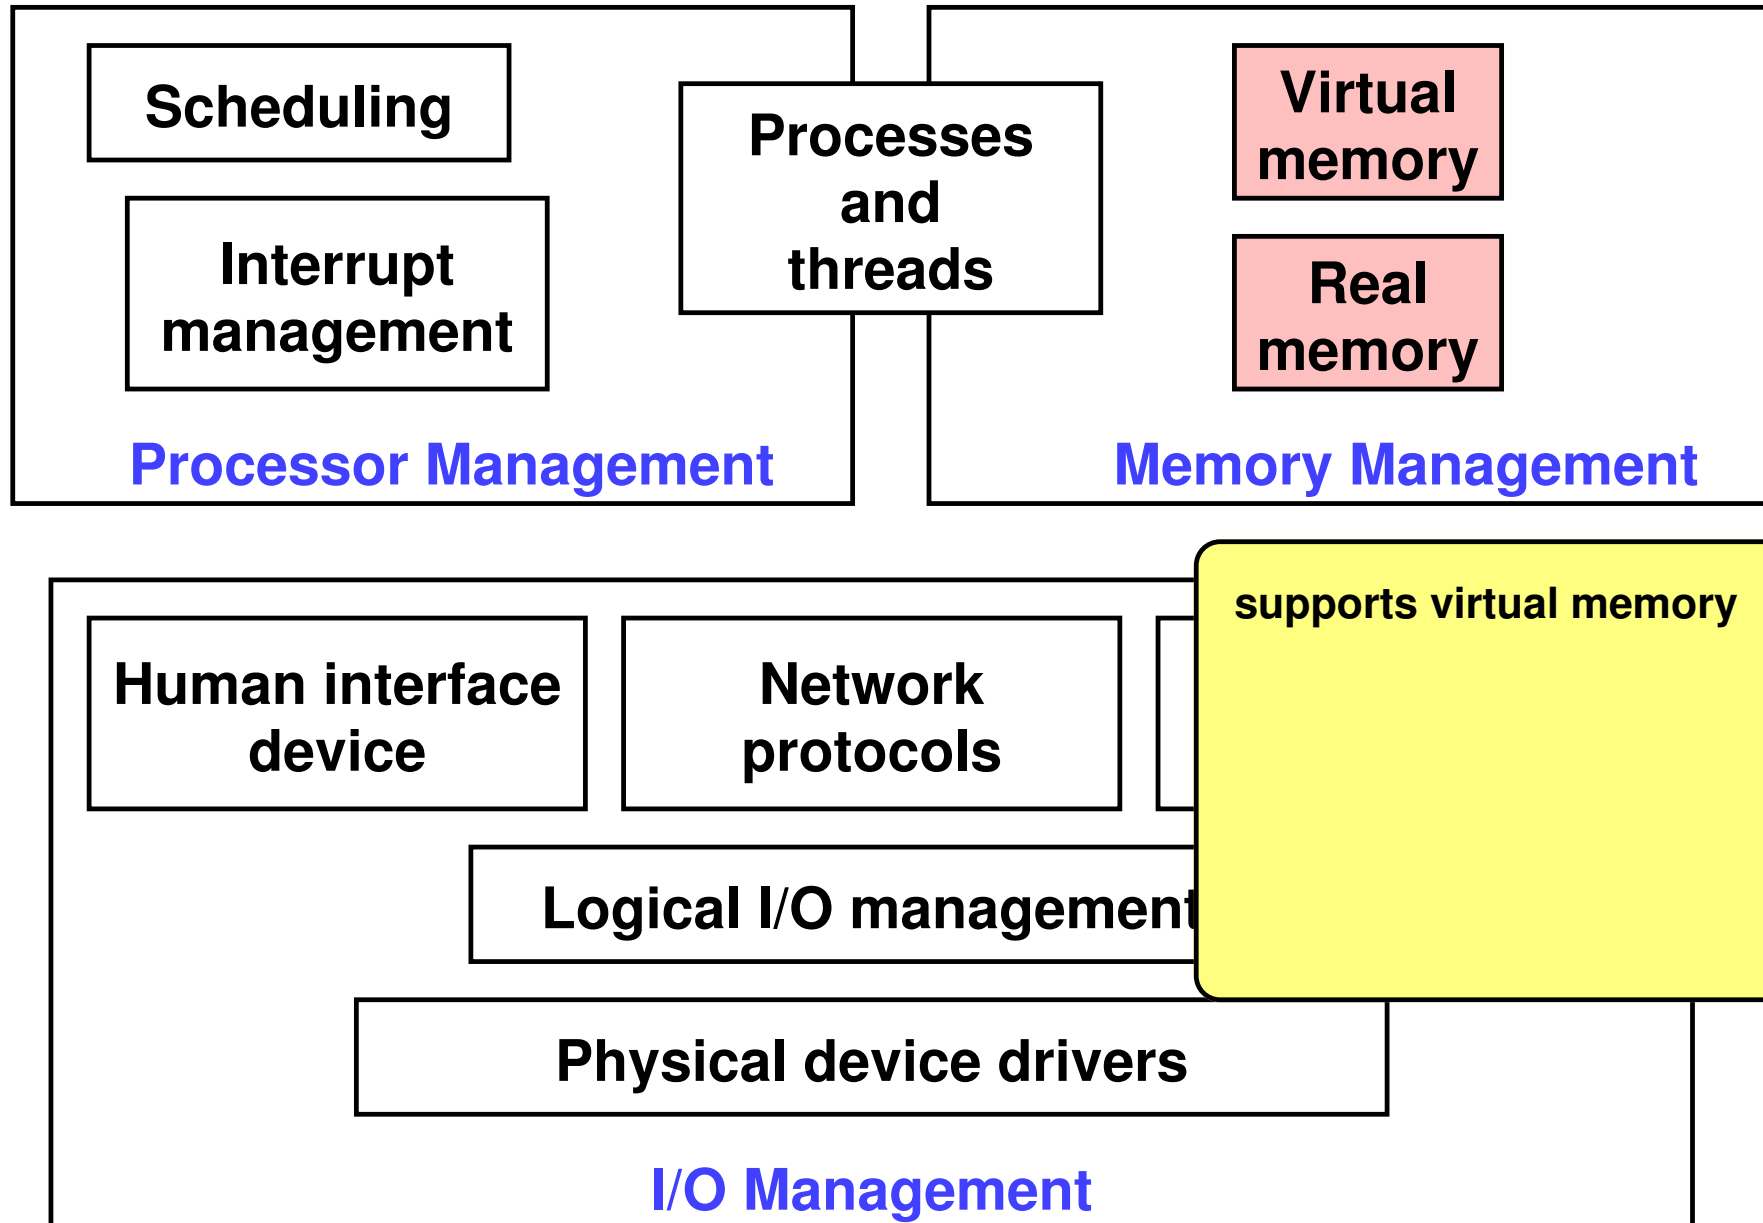
OS Components



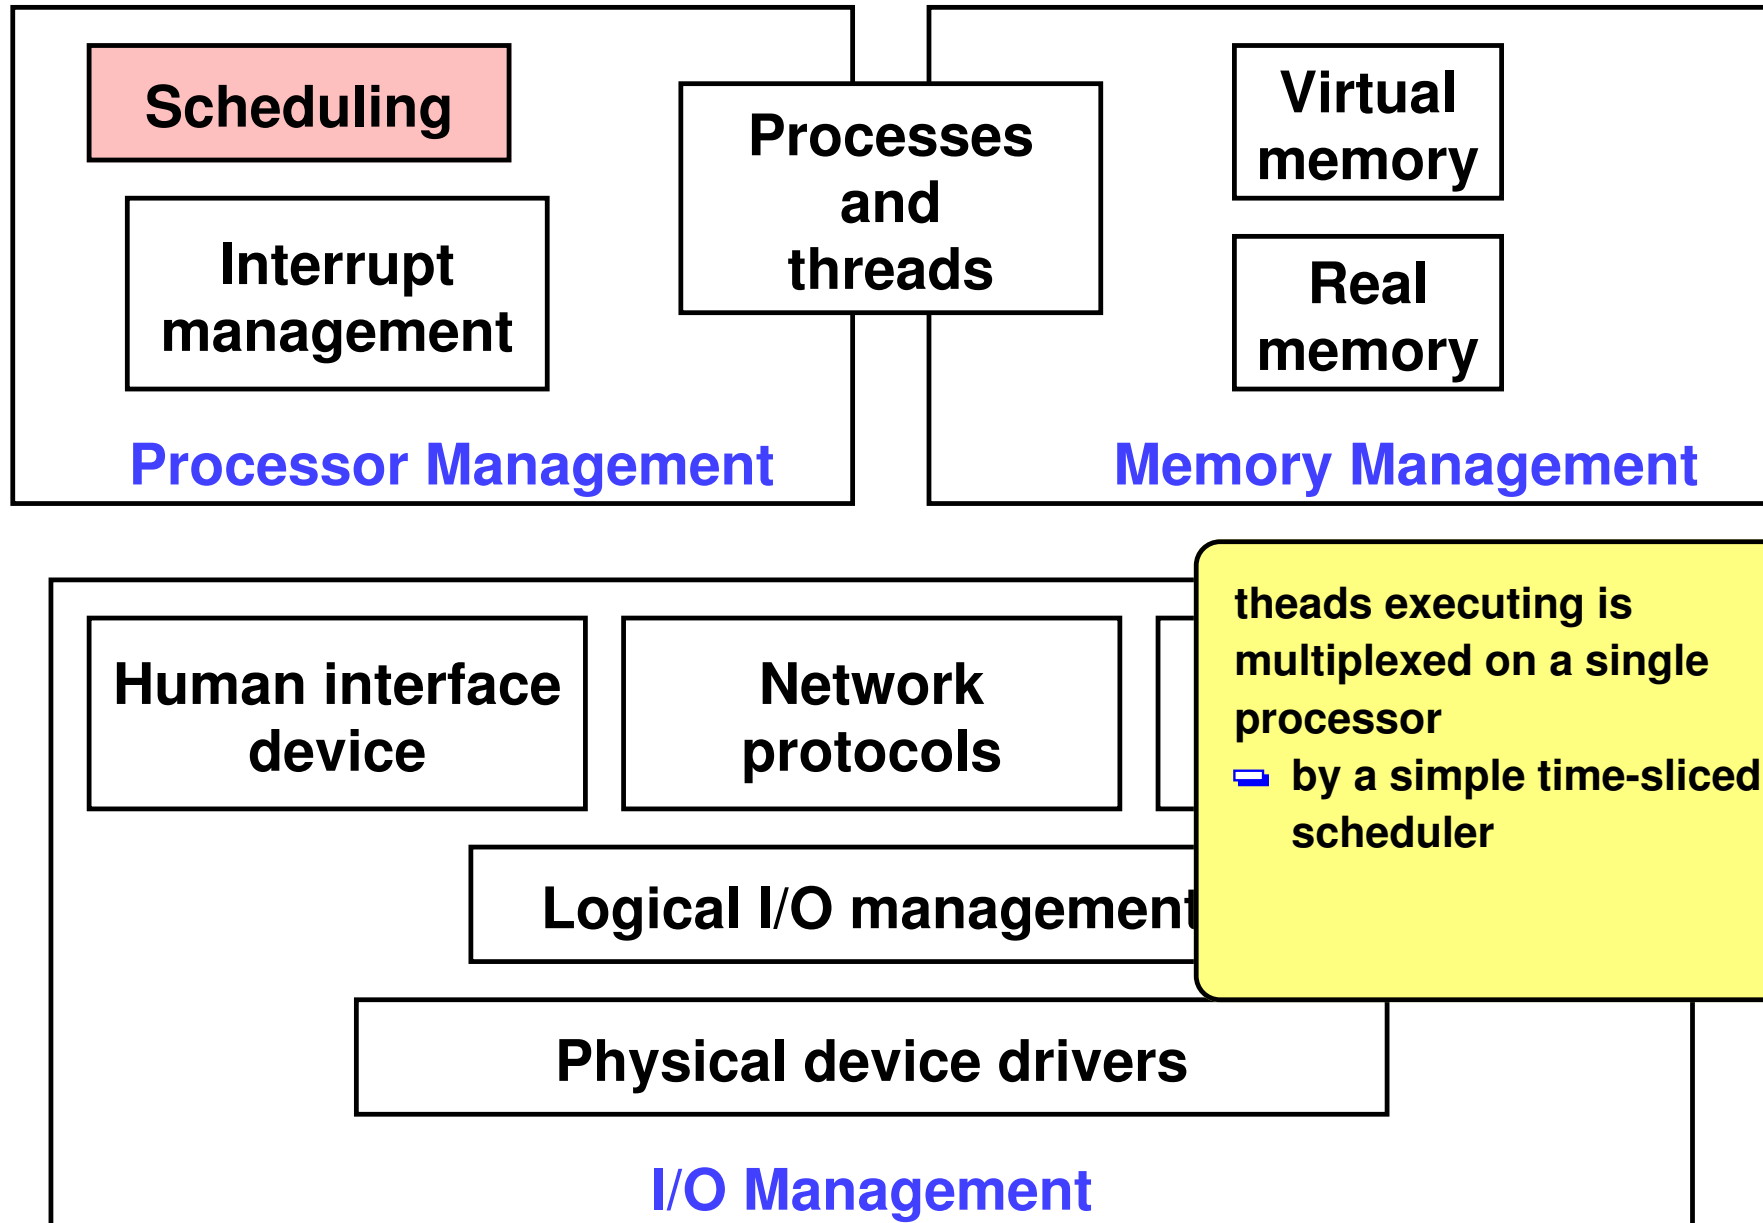
OS Components



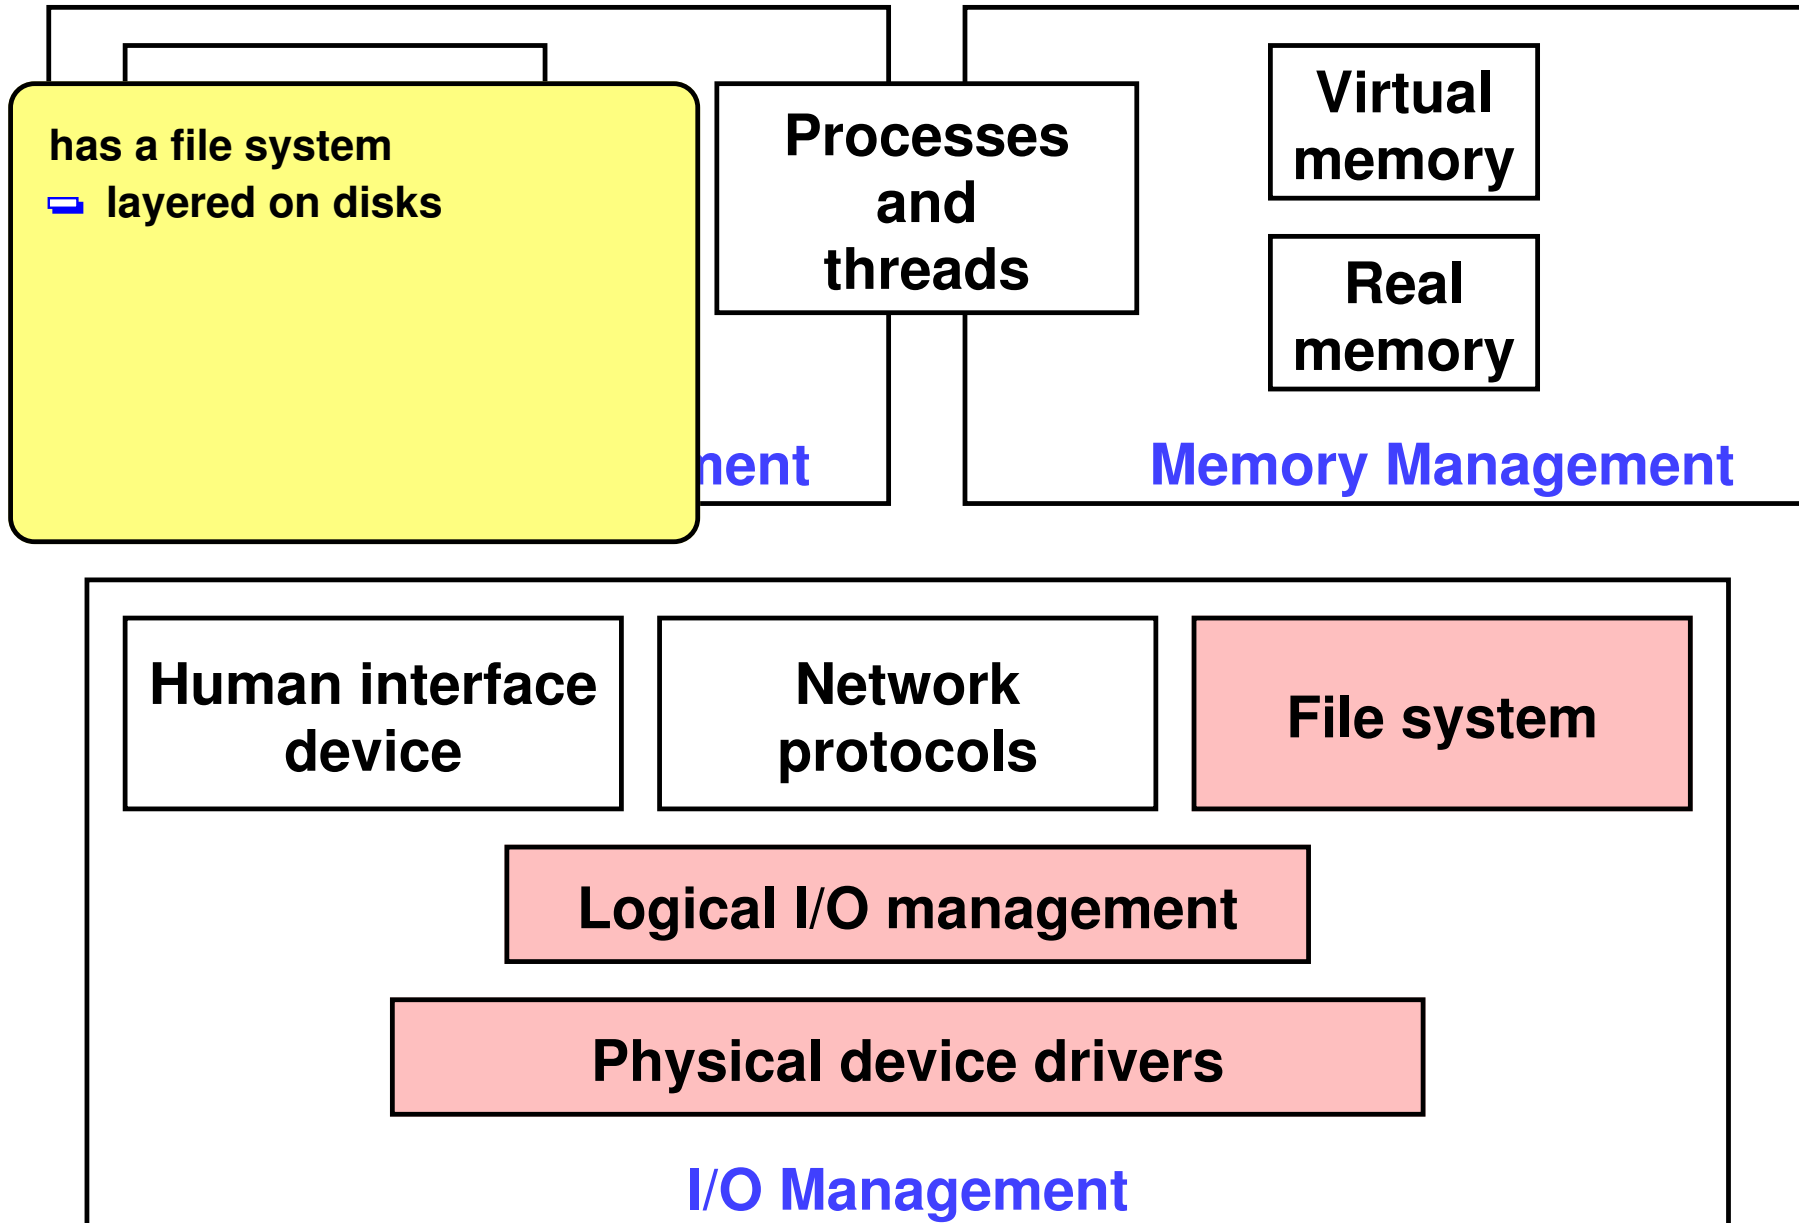
OS Components



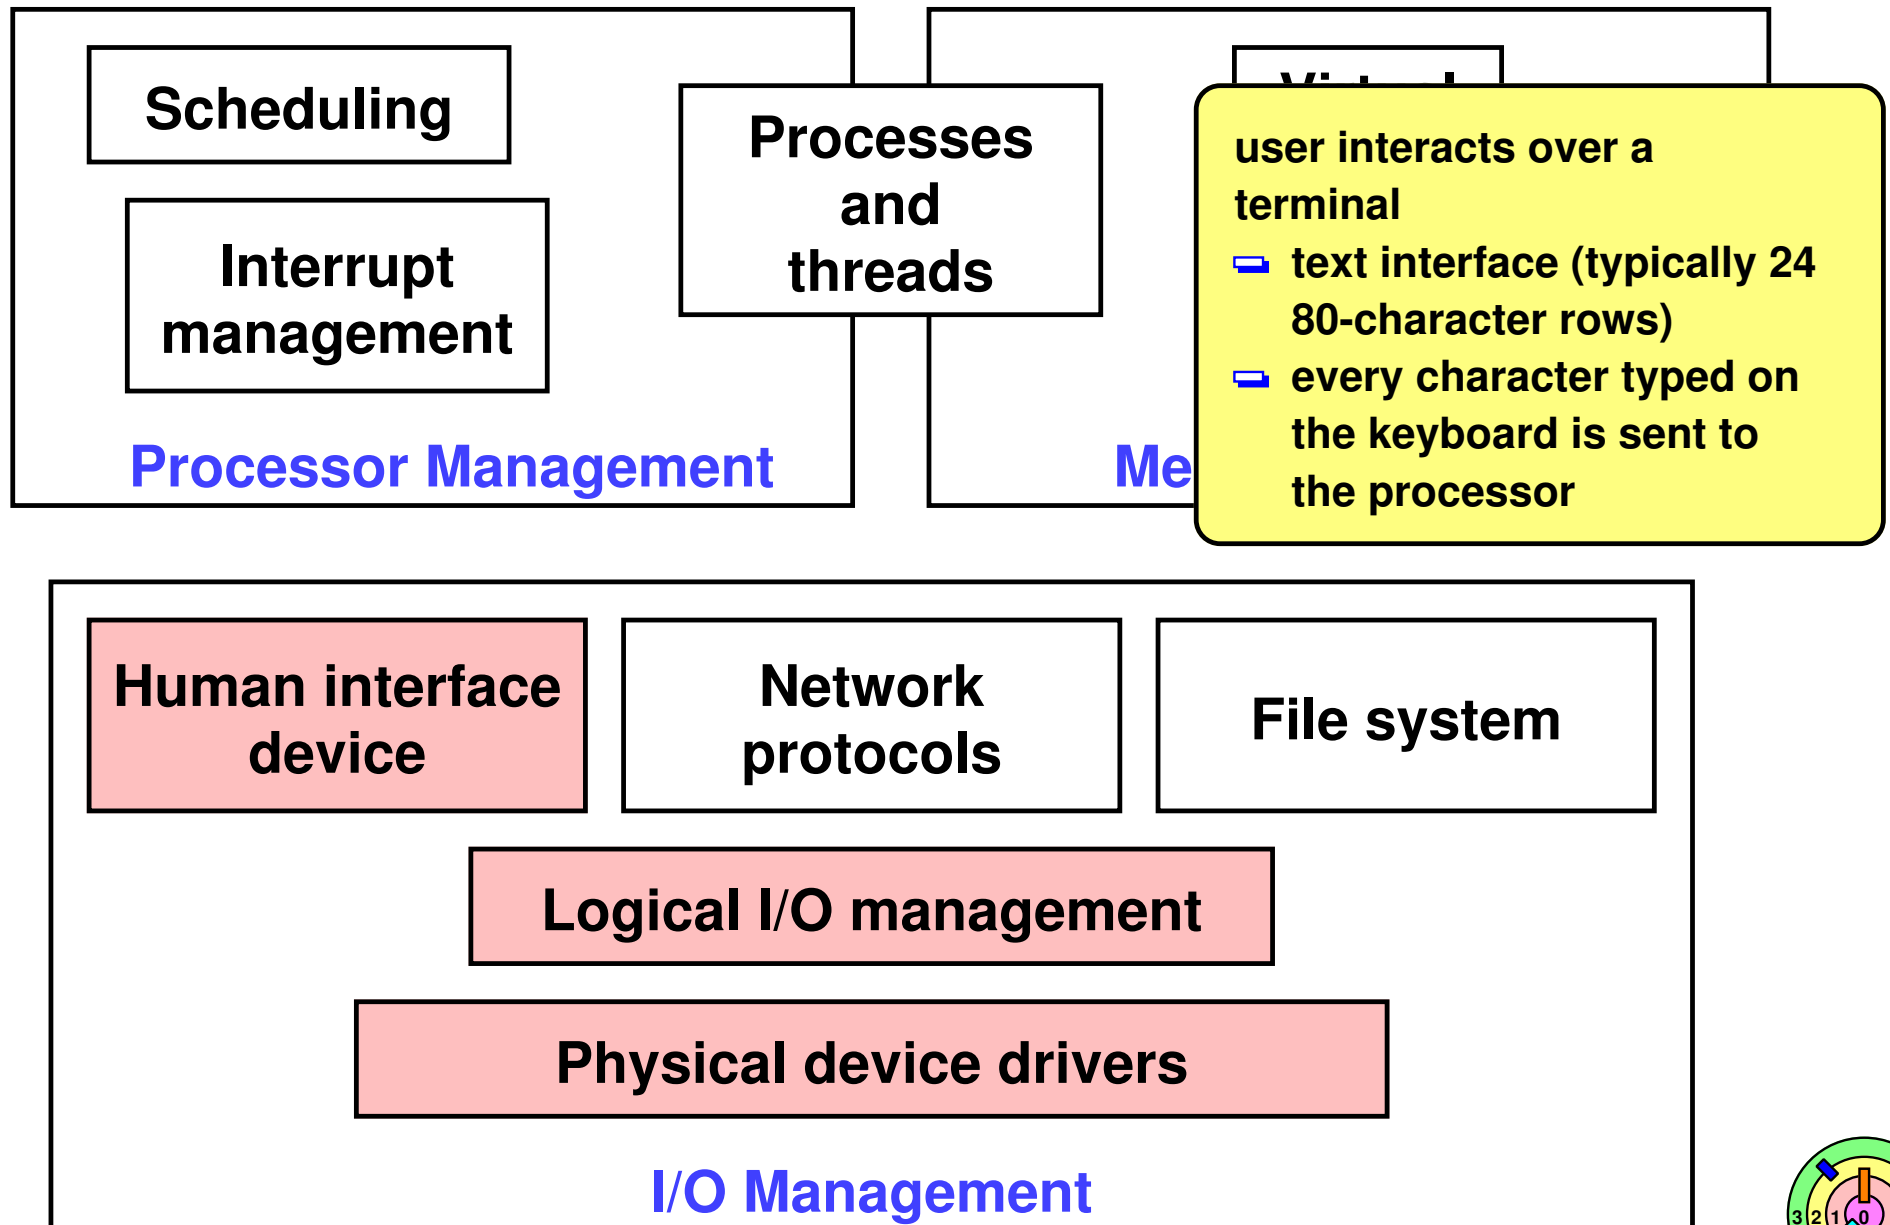
OS Components



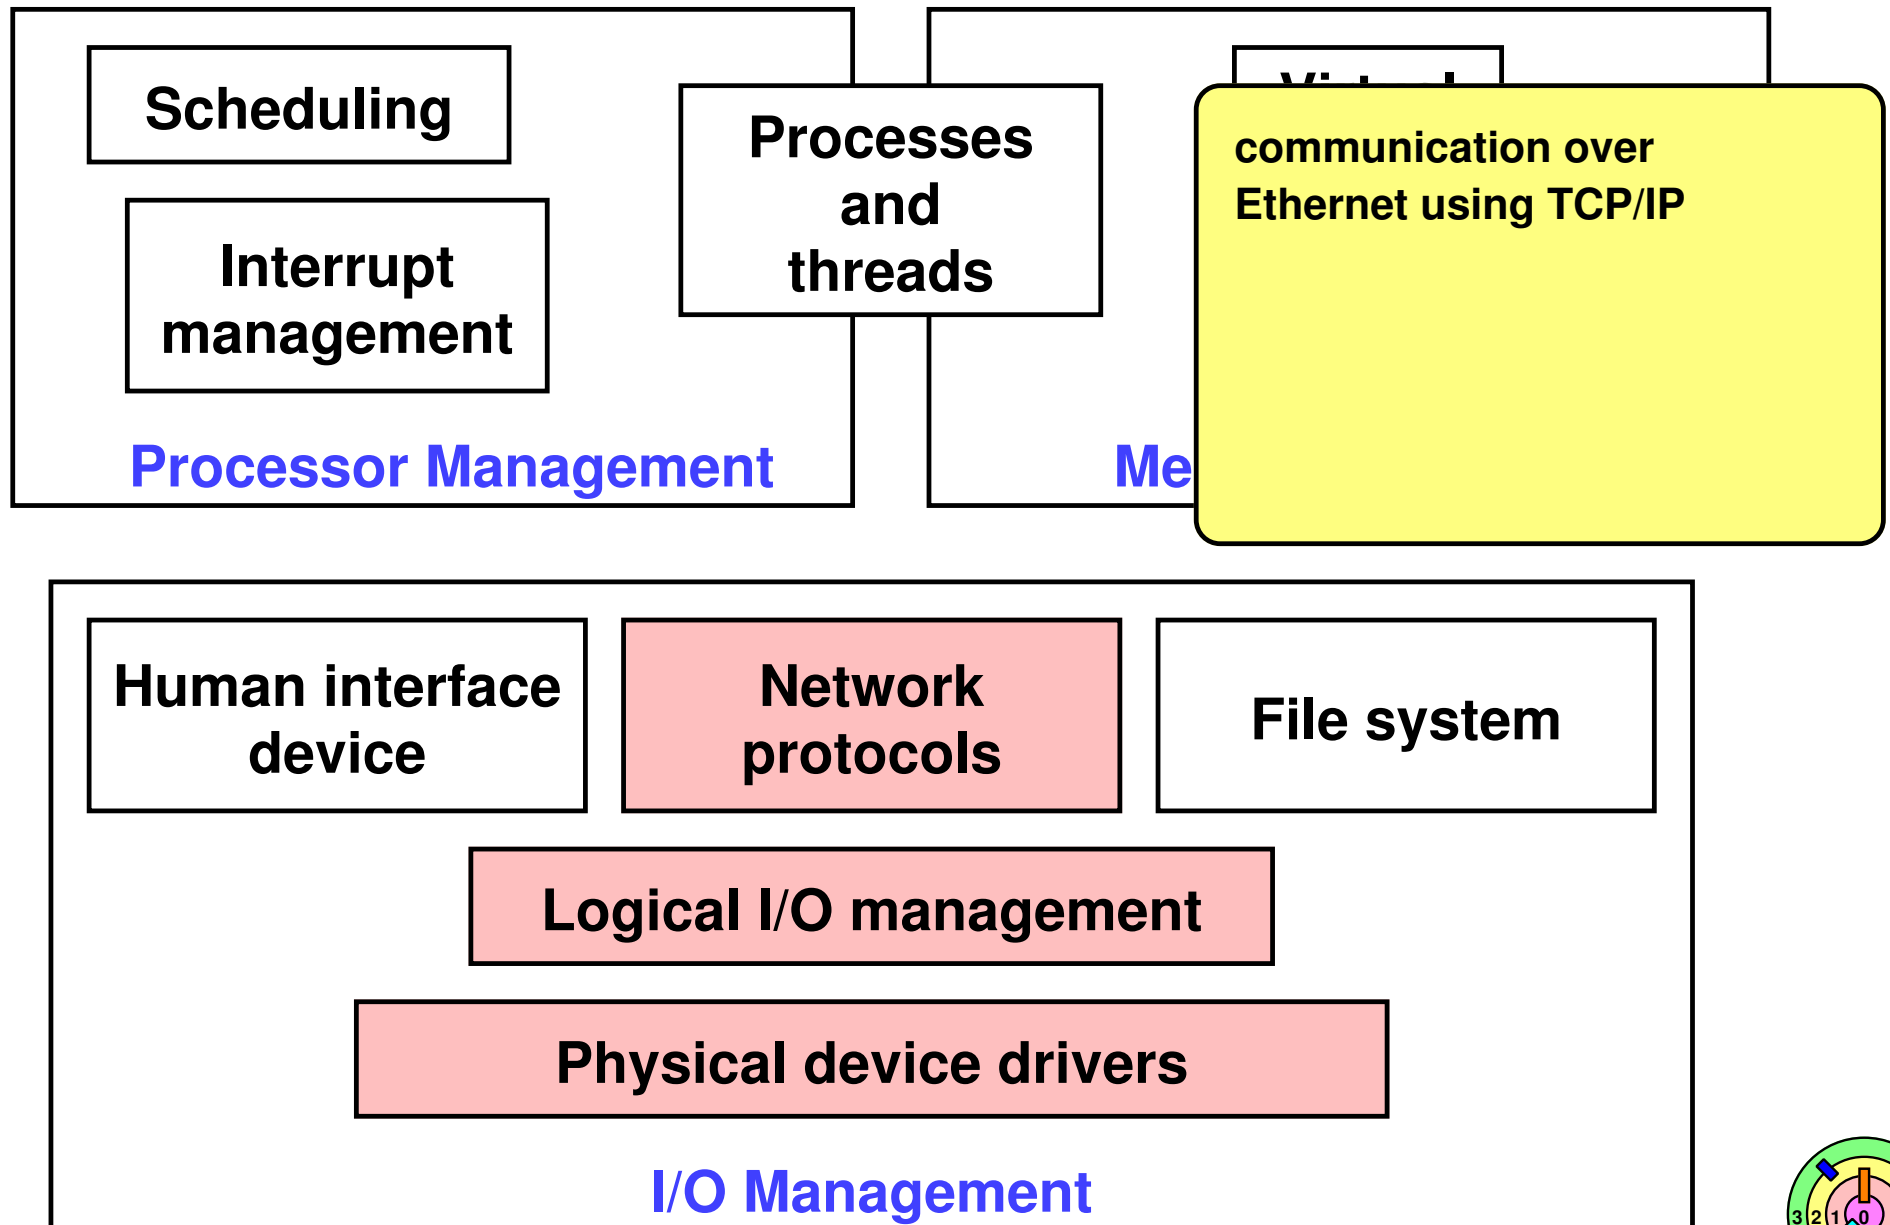
OS Components



OS Components

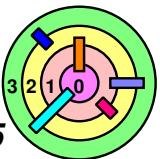


OS Components



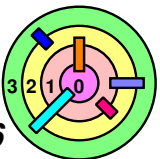
Some Important OS Concepts

- ➡ From an application program's point of view, our system has:
- ▢ processes with threads
 - ▢ a file system
 - ▢ terminals (with keyboards)
 - ▢ a network connection
- ➡ Need more details on these... Need to look at:
- ▢ how can they be provided
 - ▢ how applications use them
 - ▢ how this affects the design of the OS



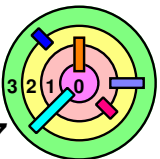
Processes And File Systems

- ➡ The purpose of a *process*
 - ➡ holds an *address space*
 - ➡ holds a group of *threads* that execute within that address space
 - ➡ holds a collection of references to *open files* and other *"execution context"*
- ➡ *Address space:*
 - ➡ set of addresses that threads of the process can usefully reference
 - ➡ more precisely, it's the content of these addressable locations
 - text, data, bss, dynamic, stack regions and what's in them



Processes And File Systems

- ➡ Design issue:
 - ➡ how should the OS *initialize* these address space regions?
- ➡ Simple approach: copy their contents from the file system to the process address space (as part of the *exec* operation)
 - ➡ quite wasteful (both in space and time) for the text region since it's read-only data
 - should *share* the text region
 - ➡ what about data regions? they can potentially be written into
 - can also *share* a portion of a data region if that portion is never modified



Remember This?

➡ Virtual Memory

Text

```

main      4096
subr      4132
printf    4156
write     16156
startup   16172

```

Data

```

aX        16384
printfargs 16388
StandardFiles 16396

```

BSS

```

X         17420
errno     17680

```

Page Table

Start	Access	Physical Addr
0	-	-
4096	R	•
8192	R	•
12288	R	•
16384	R/W	•


ask buddy system to
allocate these pages ➡

Physical
Page

Physical
Page

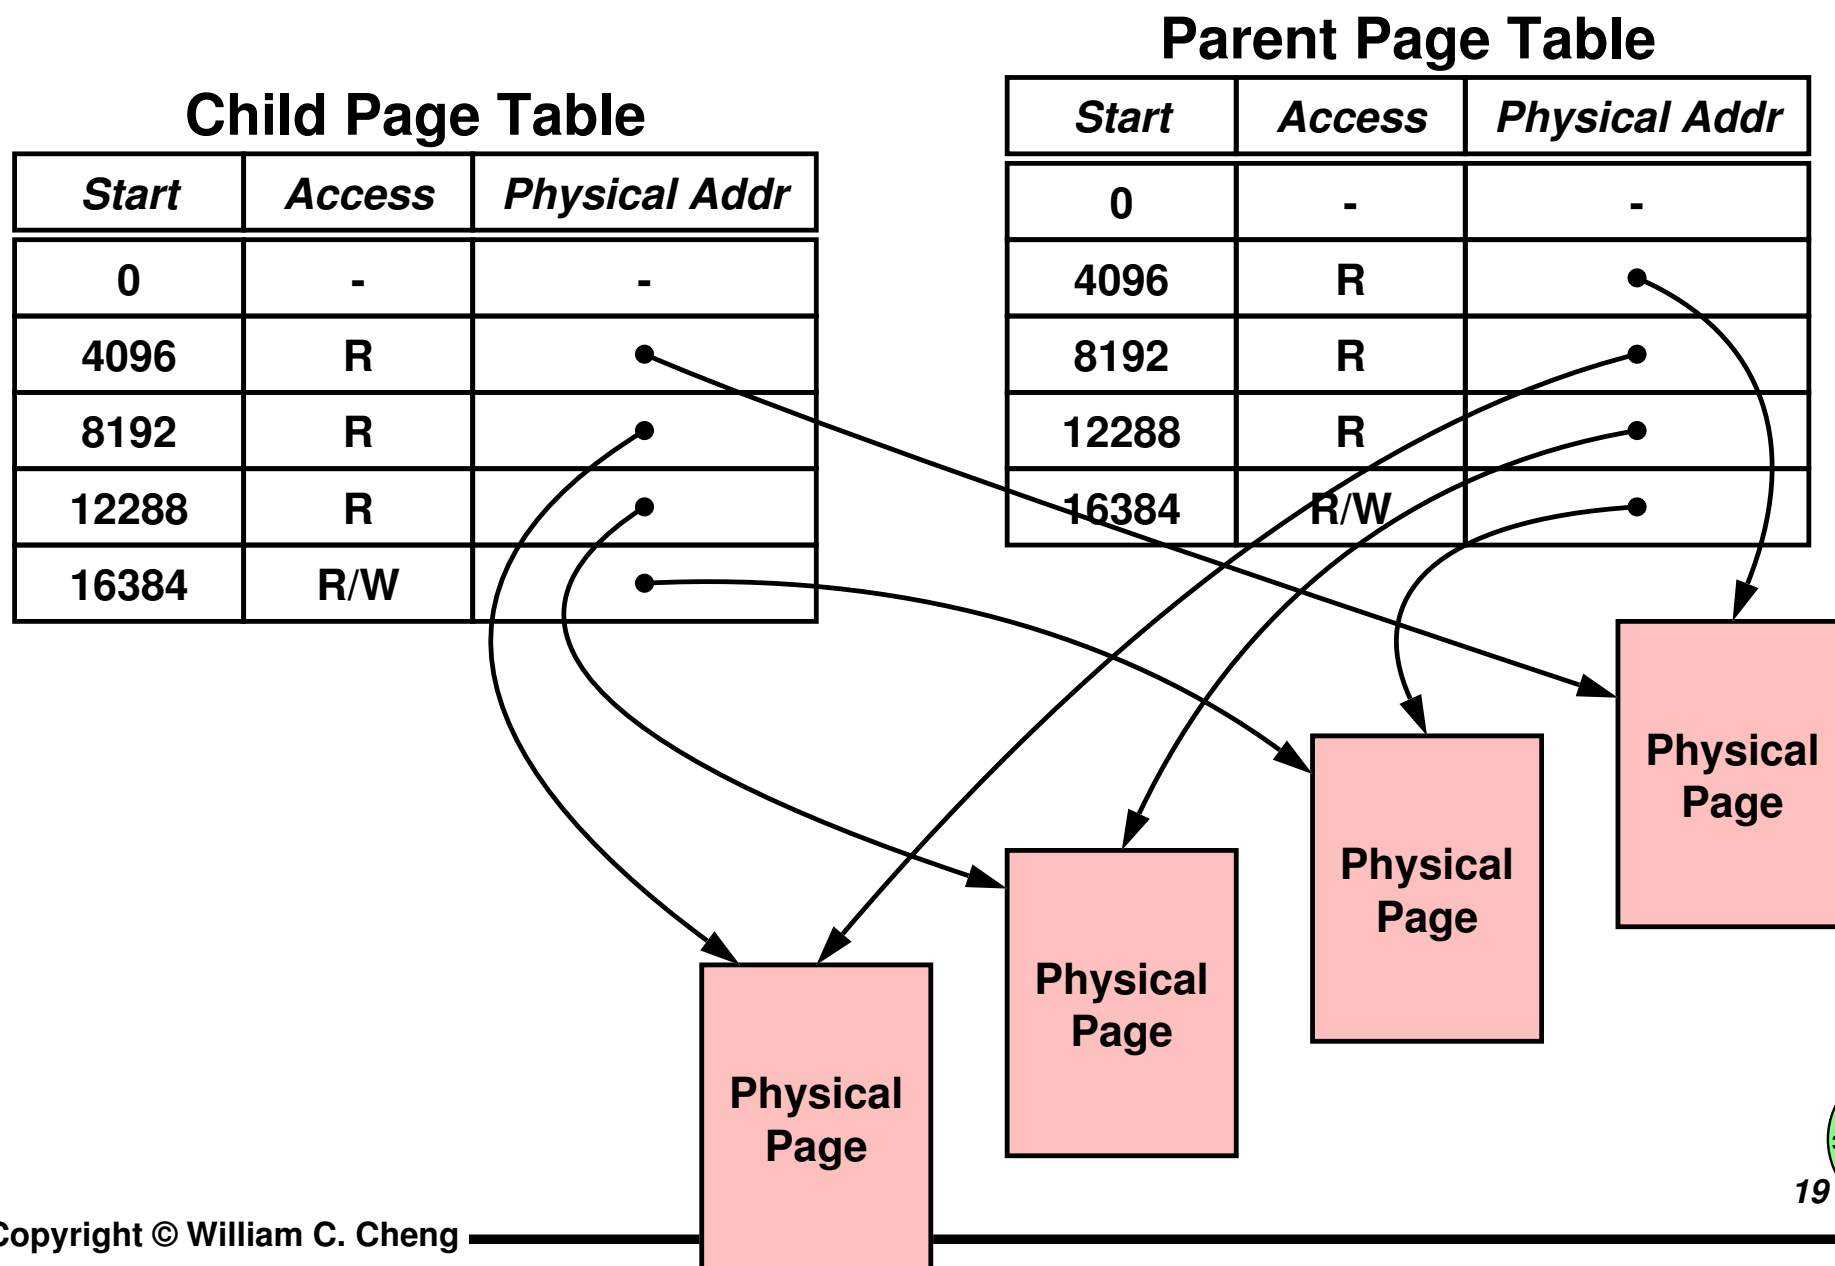
Physical
Page

Physical
Page



Processes Can Share Memory Pages

➡ Inside `fork()`, can simply copy parent's page table to child




exec ()

➡ Inside `exec ()`, need to wipe out the address space (and page table) and create a new address space (and page table)


Child Page Table

<i>Start</i>	<i>Access</i>	<i>Physical Addr</i>
0	-	-
4096	-	-
8192	-	-
12288	-	-
16384	-	-



Memory Map

- ➡ For the text region, why bother copying the executable file into the address space in the first place?
- ➡ can just *map* the *file* into the *address space* (Ch 7)
 - *mapping* is an important concept in the OS
 - *mapping* let the OS *tie* the regions of the address space to the file system
 - address space and files are divided into pieces, called *pages*
 - if several processes are executing the same program, then at most one copy of that program's text page is in memory at once
 - ➡ *text regions* of all processes running this program are setup, using hardware address translation facilities, to share these pages
 - this type of mapping is known as *shared mapping*




Memory Map

➡ The kernel uses a *memory map* to keep track of the mapping from *virtual pages* to *file pages*

Child Page Table

Start	Access	Physical Addr
0	-	-
4096	←	-
8192	←	-
12288	←	-
16384	-	-



OS

➡ the kernel also uses *memory map* to keep track of the mapping from *virtual pages* to *physical pages*

Processes Can Share Memory Pages

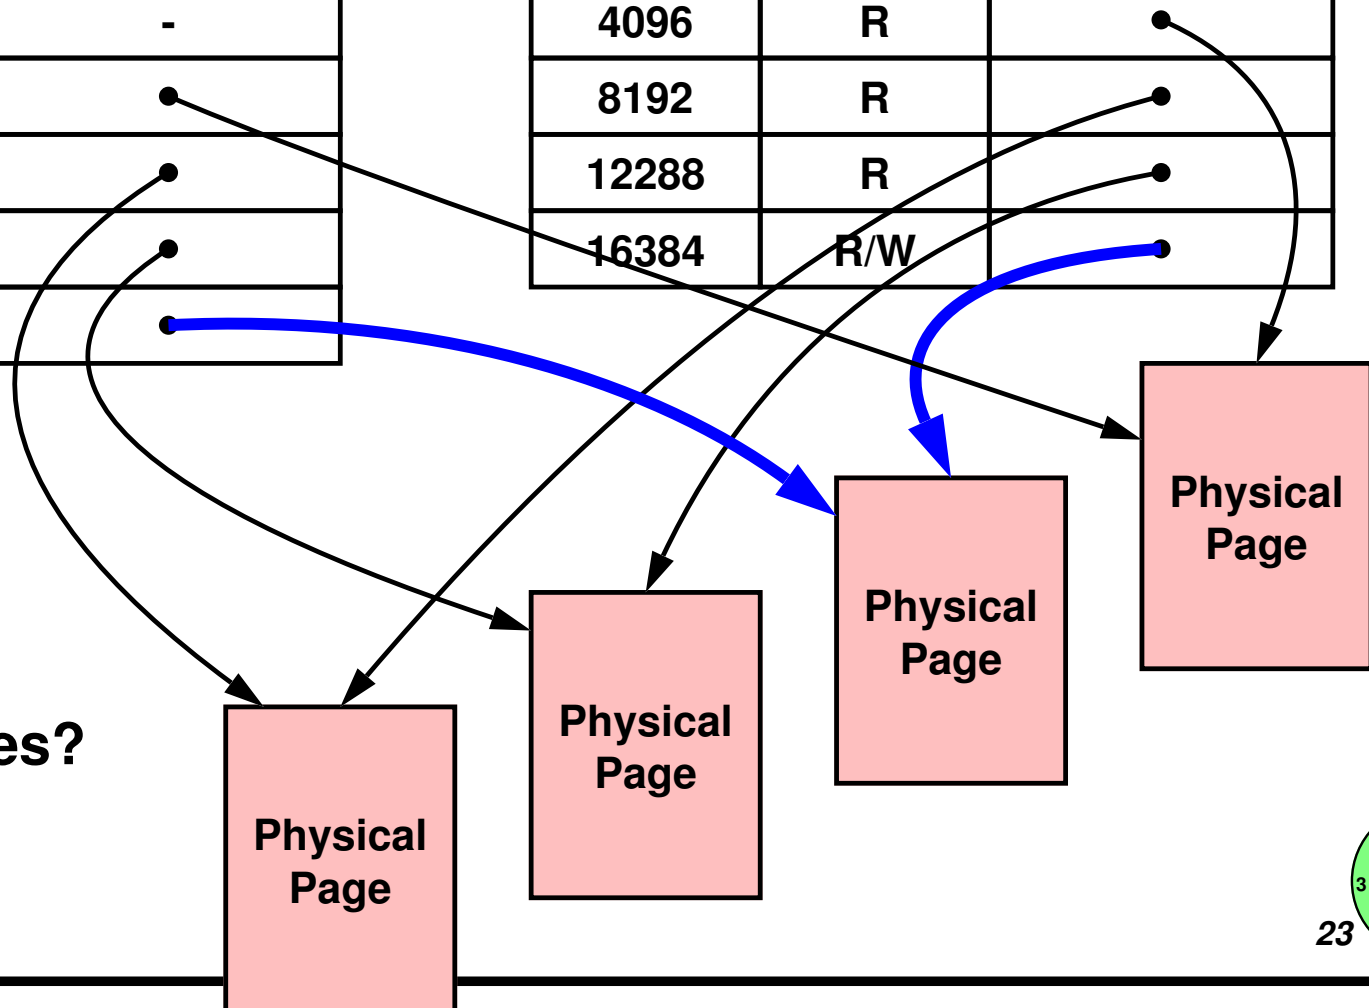
Child Page Table

<i>Start</i>	<i>Access</i>	<i>Physical Addr</i>
0	-	-
4096	R	•
8192	R	•
12288	R	•
16384	R/W	•

Parent Page Table


<i>Start</i>	<i>Access</i>	<i>Physical Addr</i>
0	-	-
4096	R	•
8192	R	•
12288	R	•
16384	R/W	•

— can we really share *data segment* pages?



Processes And File Systems

- ➡ **Data regions** of all processes running this program **initially** refer to pages of memory containing a copy of the **initial** data region
- this type of mapping is known as **private mapping**
 - when does each process really need a private copy of such a page?
 - when data is **modified** by a process, it gets a **new** and **private** copy of the initial page




Copy-On-Write

➡ *Copy-on-write (COW):*

- ➡ a process gets a *private* copy of the page after a thread in the process performs a *write* for the *first time*
 - the basic idea is that only those pages of memory that are modified are copied


➡ The dynamic/heap and stack regions use a special form of private mapping

- ➡ their pages are initialized, with zeros; *copy-on-write*
 - these are known as *anonymous pages*
 - managed by *anonymous objects* in `weenix`




Shared Files

- ➡ If a bunch of processes *share* a *file*
- we can also *map* the file into the *address space* of each process
 - in this case, the mapping is *shared*
 - when one process modifies a page, no private copy is made
 - instead, the original page itself is modified
 - everyone gets the changes
 - and changes are written back to the file
 - ◆ more on issues in Ch 6



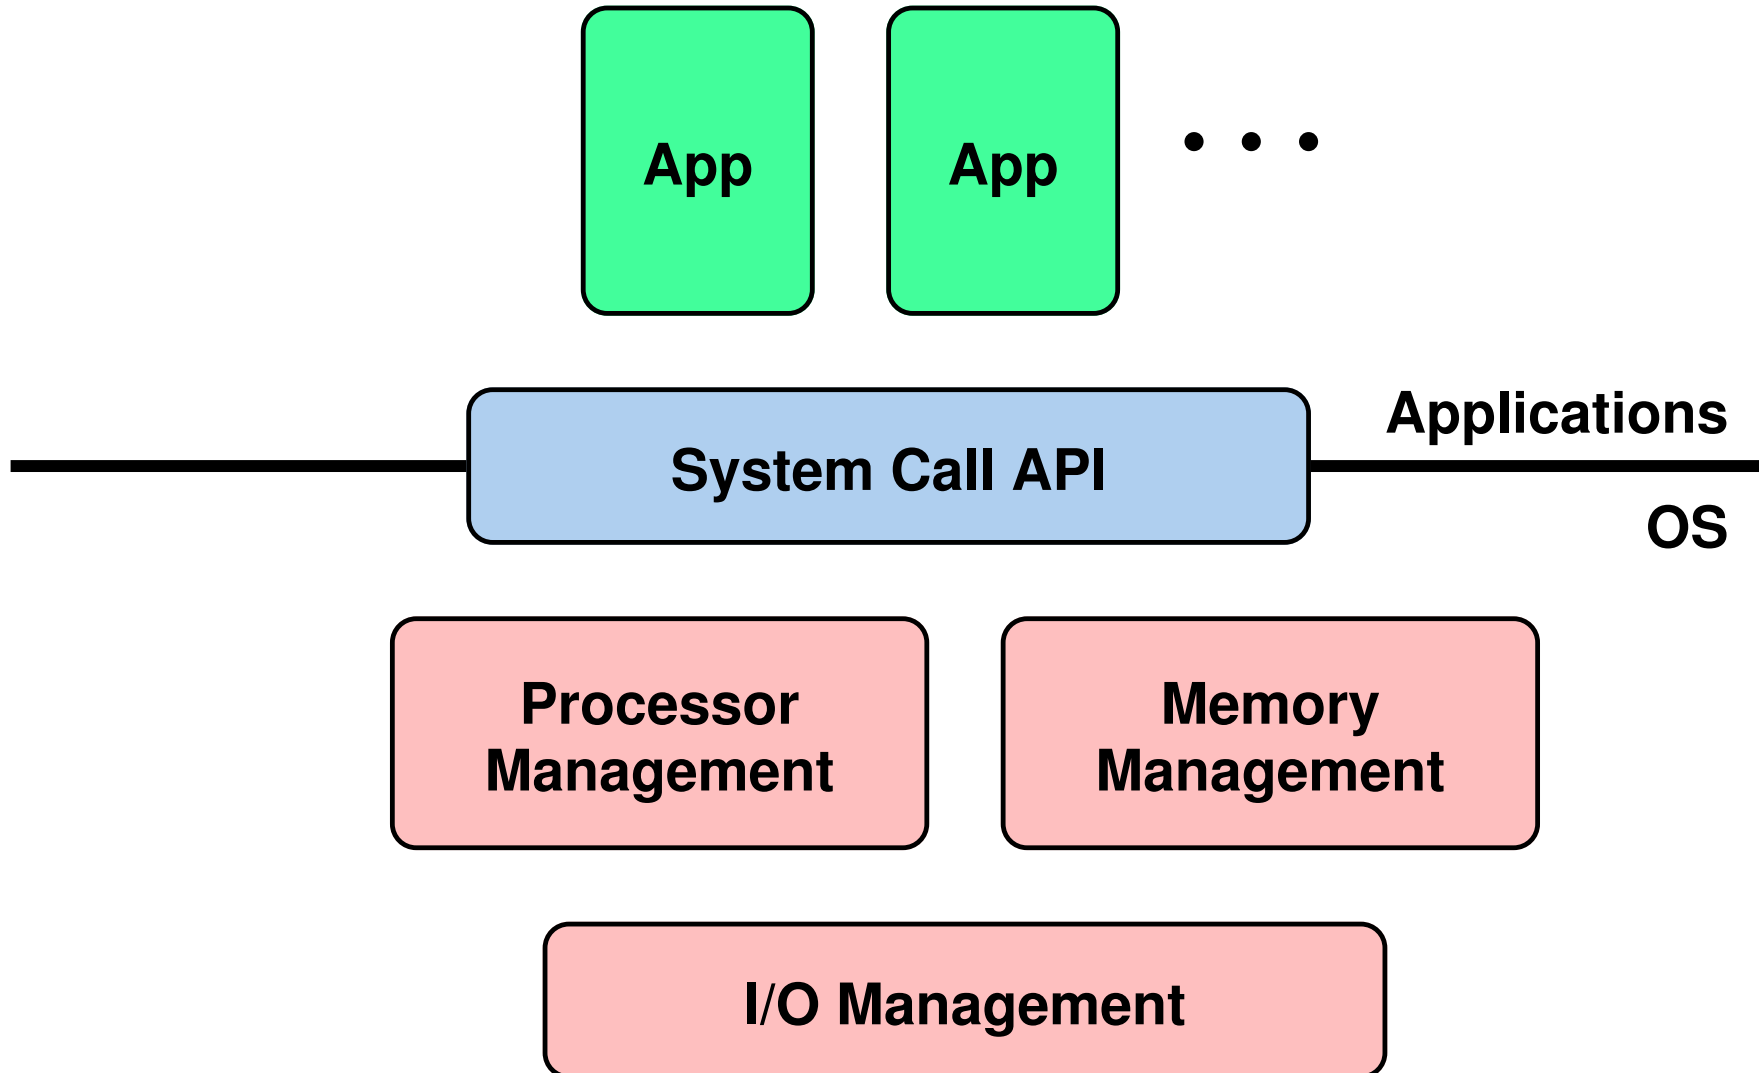
Block I/O vs. Sequential I/O

- ➡ Mapping files into address space is one way to perform I/O on files
 - block/page is the basic unit
 - this is referred to as *block I/O*
- ➡ Some devices cannot be mapped into the address space
 - e.g., receiving characters typed into the keyboard, sending a message via a network connection
 - need a more traditional approach using explicit system calls such as `read()` and `write()`
 - this is referred to as *sequential I/O*
- ➡ It also makes sense to be able to read a file like reading from the keyboard
 - similarly, a program that produces lines of text as output should be able to use the same code to write output to a file or write it out to a network connection
 - makes life easier! (and make code more robust)





System Call API

- ➡ Backwards compatibility is an important issue
- try not to change it much (to make the developers happy)




Portability

- ➡ It is desirable to have a portable operating system
 - ▬ portable across various hardware platforms
- ➡ For a monolithic OS, it is achieved through the use of a *Hardware Abstraction Layer (HAL)*
 - ▬ a *portable interface* to *machine configuration* and *processor-specific operations* within the kernel



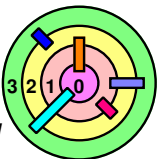
Hardware Abstraction Layer (HAL)

- ➡ Portability across machine configuration
 - e.g., different manufacturers for x86 machines will require different code to configure interrupts, hardware timers, etc.
- ➡ Portability across processor families
 - e.g., may need additional code for context switching, system calls, interrupting handler, virtual memory management, etc.
- ➡ With a well-defined Hardware Abstraction Layer, most of the OS is *machine* and *processor independent*
 - porting an OS to a new computer is done by
 - writing new HAL routines
 - relink with the kernel



4.1 A Simple System (Monolithic Kernel)

- ➡ *A Framework for Devices*
- ➡ Low-level Kernel (will come back to talk about this after Ch 7)
- ➡ Processes & Threads
- ➡ Storage Management (will come back to talk about this after Ch 5)



Computer Terminal



➡ VT100

A "tty"




Devices

➡ Challenges in supporting devices

- device independence
- device discovery

➡ Device naming

- two choices
 - independent name space (i.e., named independently from other things in the system)
 - devices are named as files



A Framework for Devices



Device driver:

- every device is identified by a device "number", which is actually a pair of numbers
 - a *major device number* - identifies the device driver
 - a *minor device number* - device index for all devices managed by the same device driver




Special entries were created in the file system to refer to devices

- usually in the `/dev` directory
 - e.g., `/dev/disk1`, `/dev/disk2` each marked as a *special file*
 - ◆ a *special file* does not contain data
 - ◆ it refers to devices by their major and minor device numbers
 - ◆ if you do `"ls -l"`, you can see the device numbers


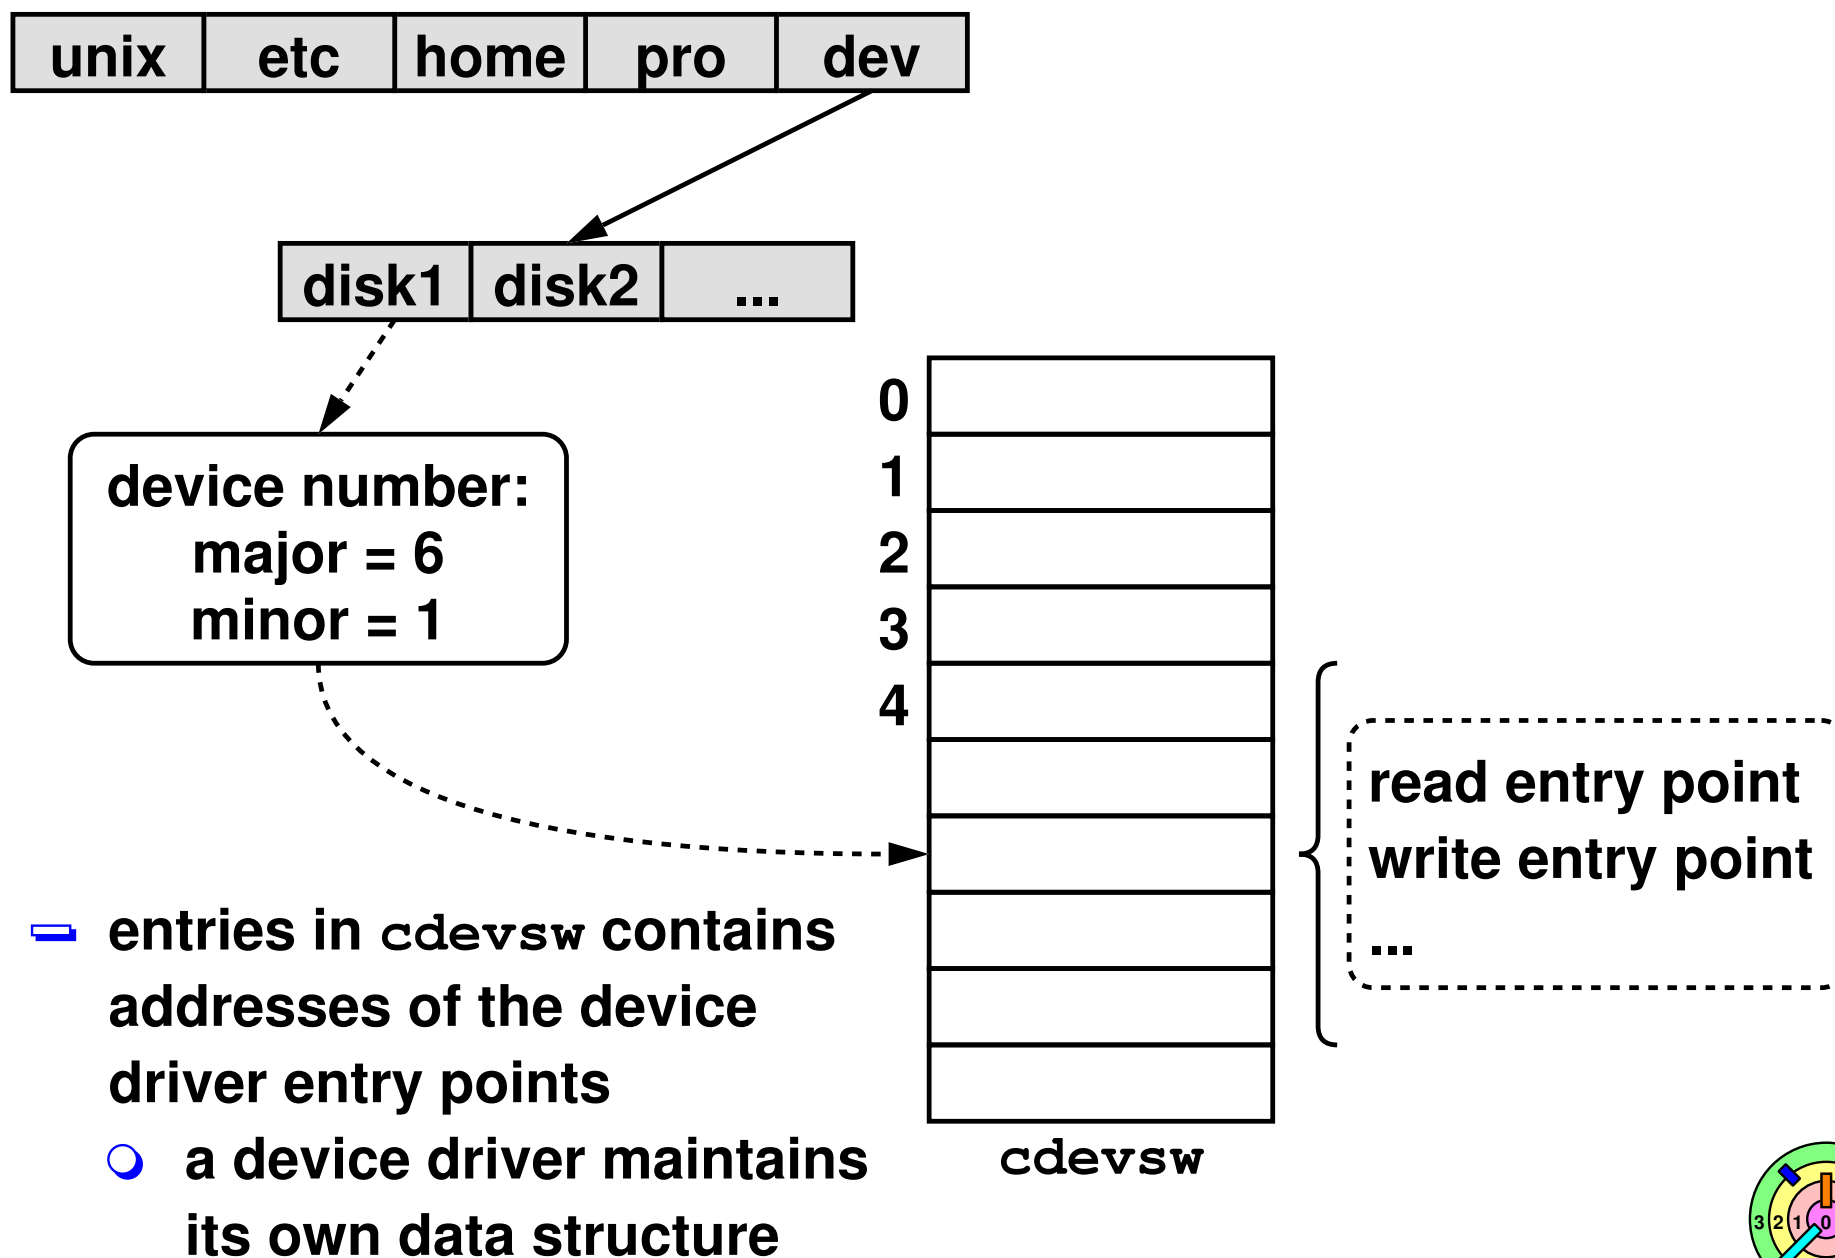


Data structure in the early Unix systems

- statically allocated array in the kernel called `cdevsw` (character device switch)




Finding Devices




Device Drivers in Early Unix Systems

- ➡ The kernel was statically configured to contain device-specific information such as:
 - ▬ interrupt-vector locations
 - ▬ locations of device-control registers on whatever bus the device was attached to
- ➡ Static approach was simple, but cannot be easily extended
 - ▬ a kernel must be custom configured for each installation



Device Probing

- ➡ First step to improve the old way
 - allow the devices to be found and automatically configured when the system booted
 - (still require that a kernel contain all necessary device drivers)
- ➡ Each device driver includes a *probe routine*
 - invoked at boot time
 - probe the relevant buses for devices and configure them
 - including identifying and recording interrupt-vector and device-control-register locations
- ➡ This allowed one kernel image to be built that could be useful for a number of similar but not identical installations
 - boot time is kind of long
 - impractical as the number of supported devices gets big



Device Probing

➡ What's the right thing to do?

Step 1: discover the device without the benefit of having the relevant device driver in the kernel

Step 2: find the needed device drivers and dynamically link them into the kernel

— but how do you achieve this?

➡ Solution: use meta-drivers

— a meta-driver handles a particular kind of bus


— e.g., USB (Universal Serial Bus)

- a USB meta-driver is installed into the kernel

- any device that goes onto a USB (Universal Serial Bus) must know how to interact with the USB meta-driver via the *USB protocol*


- once a connected device is identified, system software would select the appropriate device driver and load into the kernel

- what about applications? how can they reference dynamically discovered devices?




Discovering Devices

- ➡ So, you plug in a new device to your computer on a particular bus
 - OS would notice
 - find a device driver
 - what kind of device is it?
 - where is the driver?
 - assign a name, but how is it chosen?
 - multiple similar devices, but how does application choose?
- ➡ In some Linux systems, entries are added into `/dev` as the kernel discovers them
 - lookup the names from a database of names known as `devfs`
 - downside of this approach is that device naming conventions not universally accepted
 - what's an application to do?
 - some current Linux systems use `udev`
 - user-level application assigns names based on rules provided by an administrator




Discovering Devices

- ➡ What about the case where different devices acted similarly?
 - e.g., touchpad on a laptop and USB mouse
 - how should the choice be presented to applications?
- ➡ Windows has the notion of *interface classes*
 - a device can register itself as members of one or more such classes
 - an application can *enumerate* all currently connected members of such a class and choose among them (or use them all)




4.1 A Simple System (Monolithic Kernel)

- ➡ A Framework for Devices
- ➡ Low-level Kernel (will come back to talk about this after Ch 7)
- ➡ *Processes & Threads*
- ➡ Storage Management (will come back to talk about this after Ch 5)

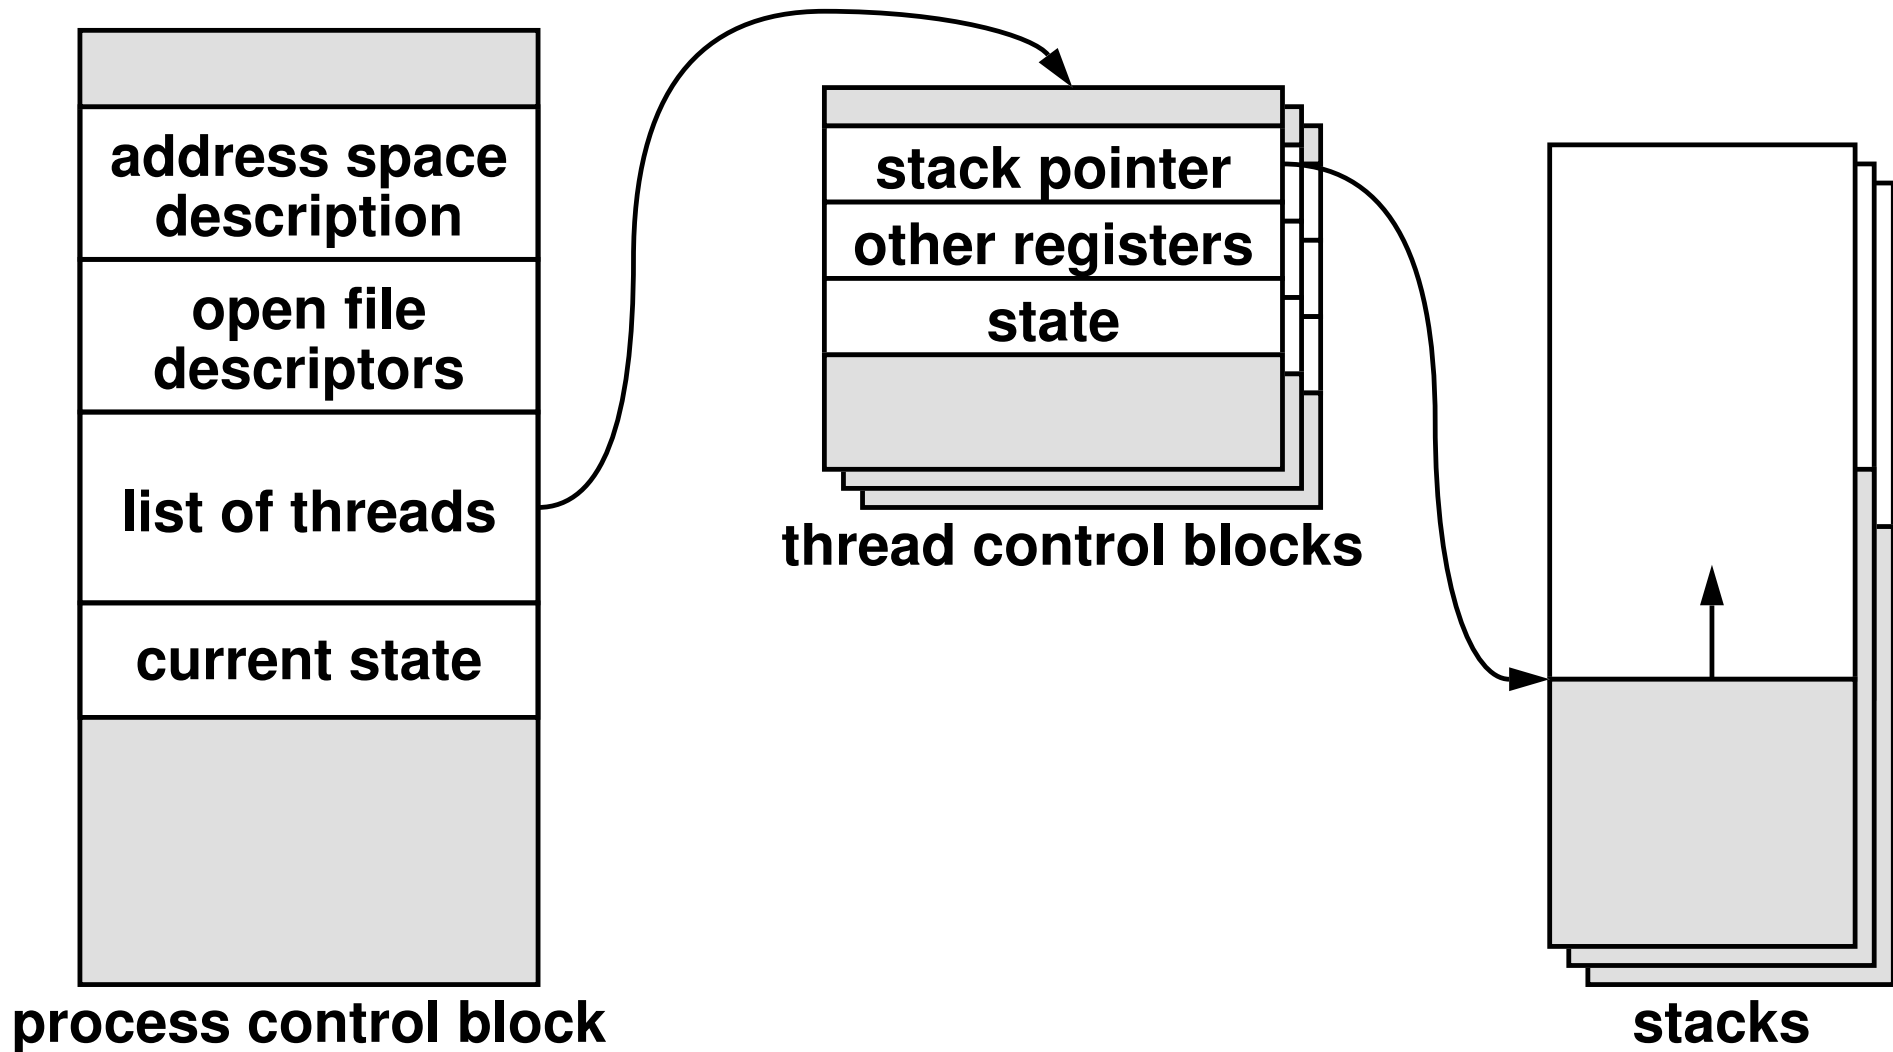


Processes and Threads

- ➡ A process is:
- ➡ a holder for an *address space*
 - ➡ a collection of other information shared by a set of *threads*
 - ➡ a collection of references to *open files* and other "execution context"
- ➡ As discussed in Ch 1, processes related APIs include
- ➡ `fork()`, `exec()`, `wait()`, `exit()`



Processes and Threads





➡ Note: all these are relevant to your Kernel Assignment 1
— although we are only doing *one thread per process*

Process Life Cycle



Pretty simple

— a process starts in the *run* state


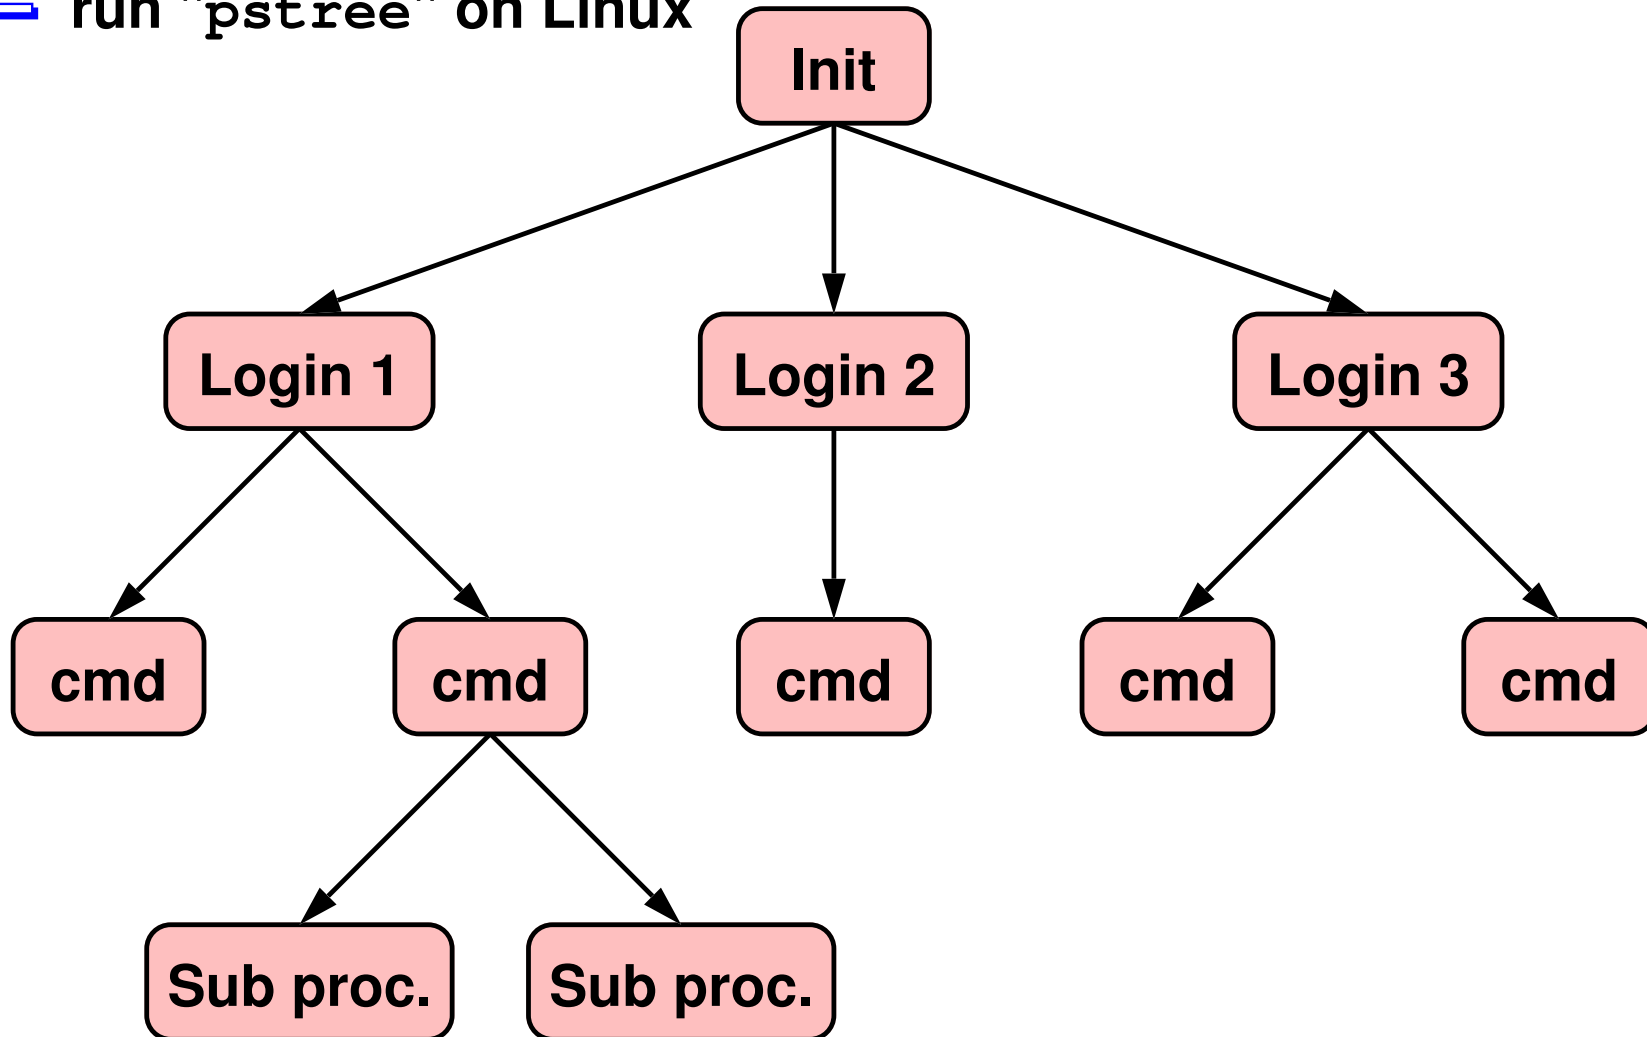


Process Relationships (1)


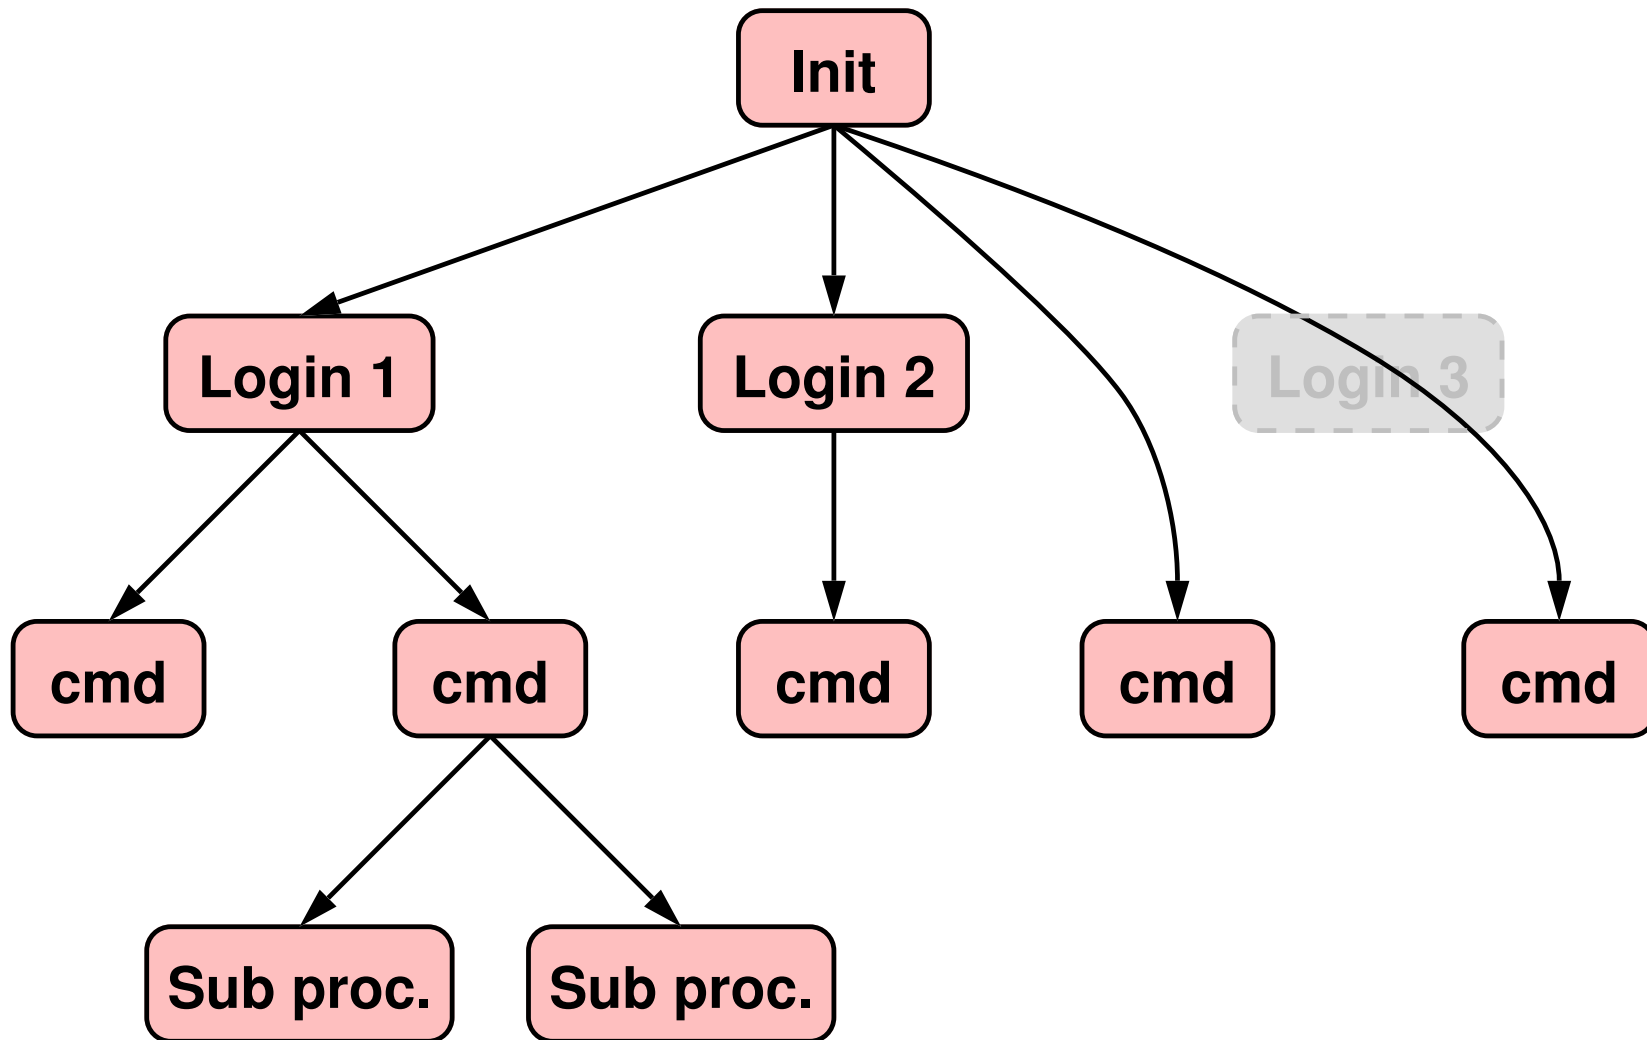


Process hierarchy

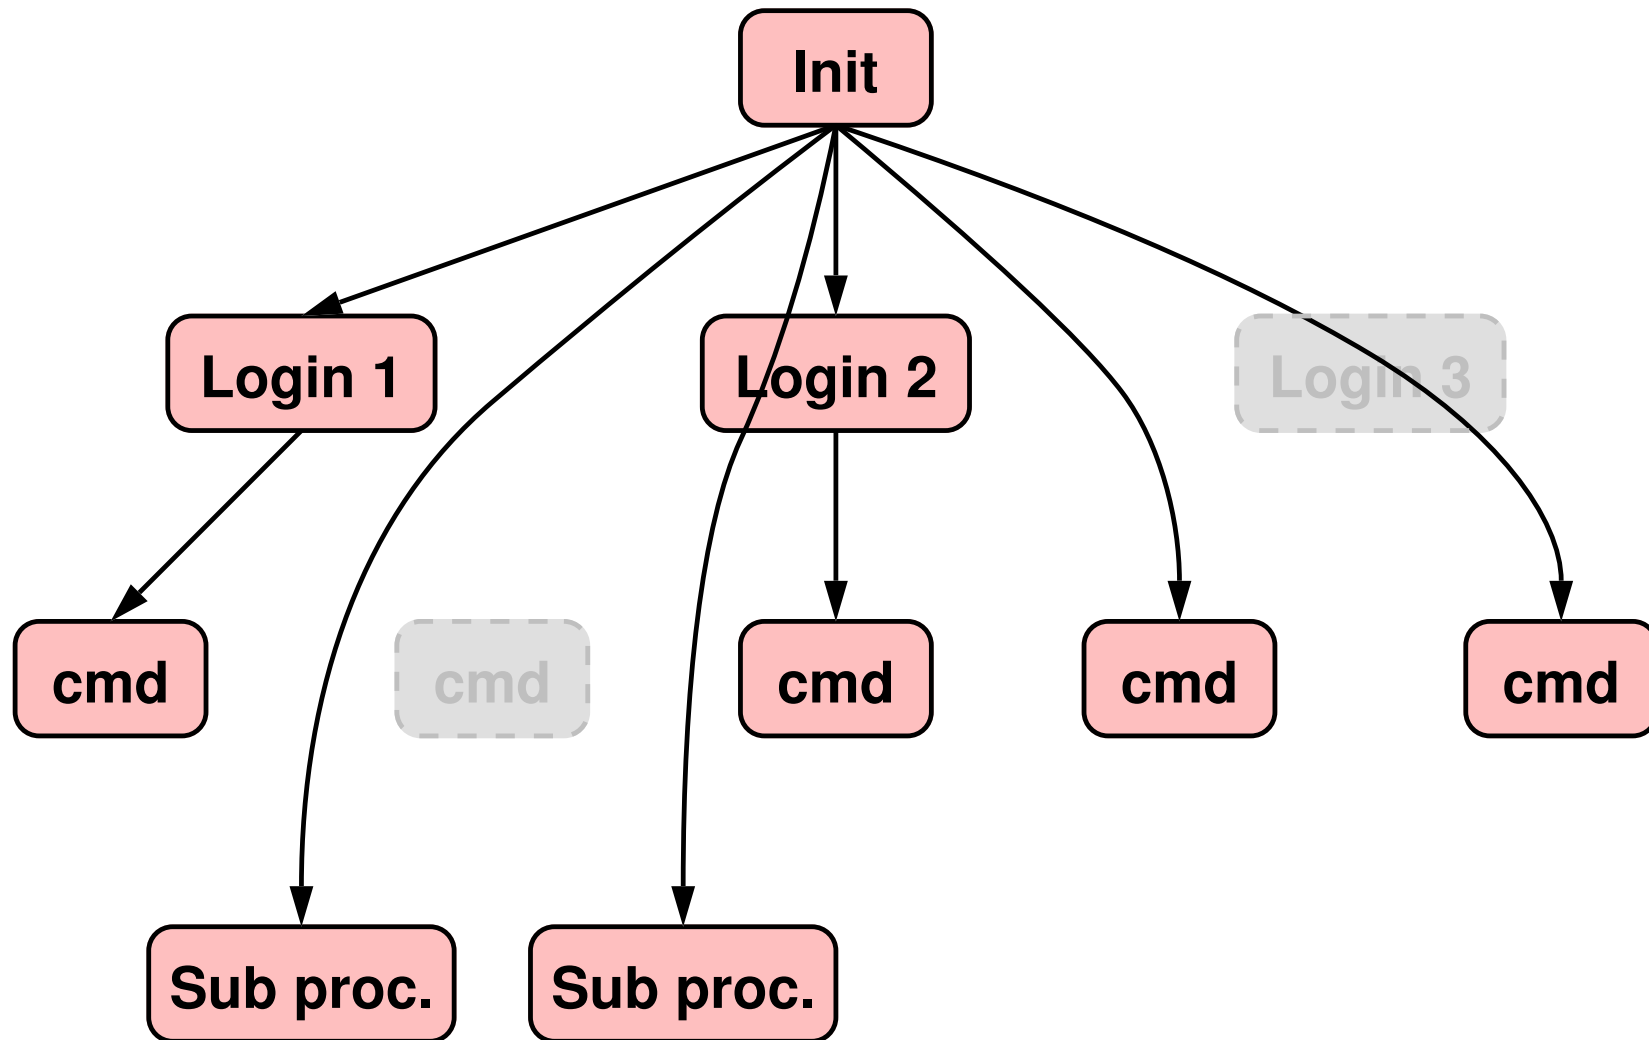
— run "pstree" on Linux




Process Relationships (2)



Process Relationships (3)

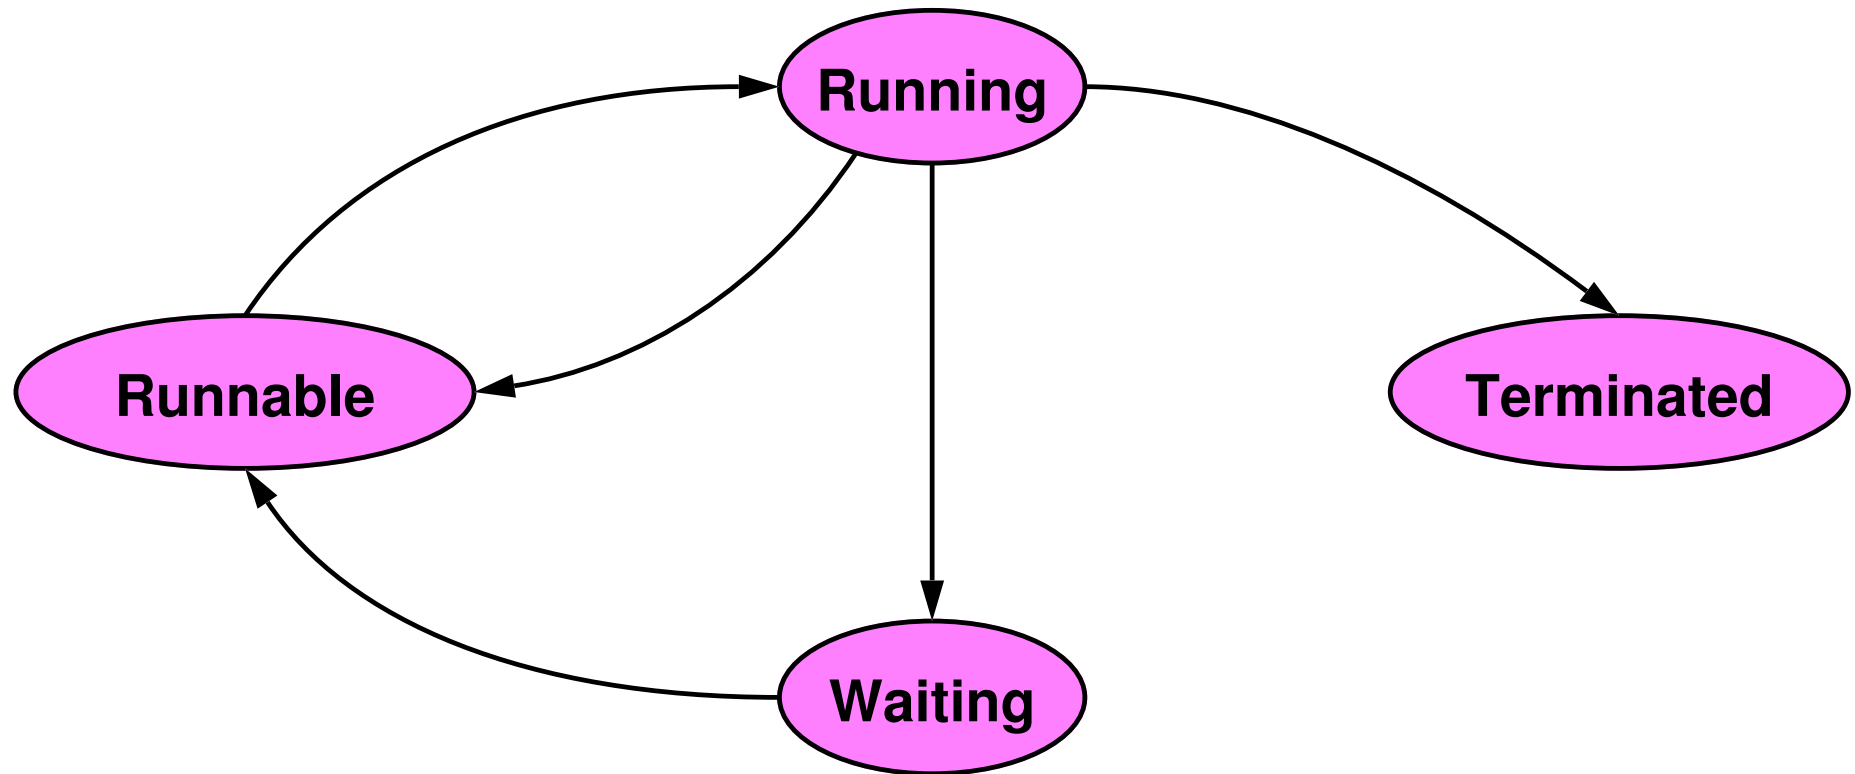


Fork and Threads



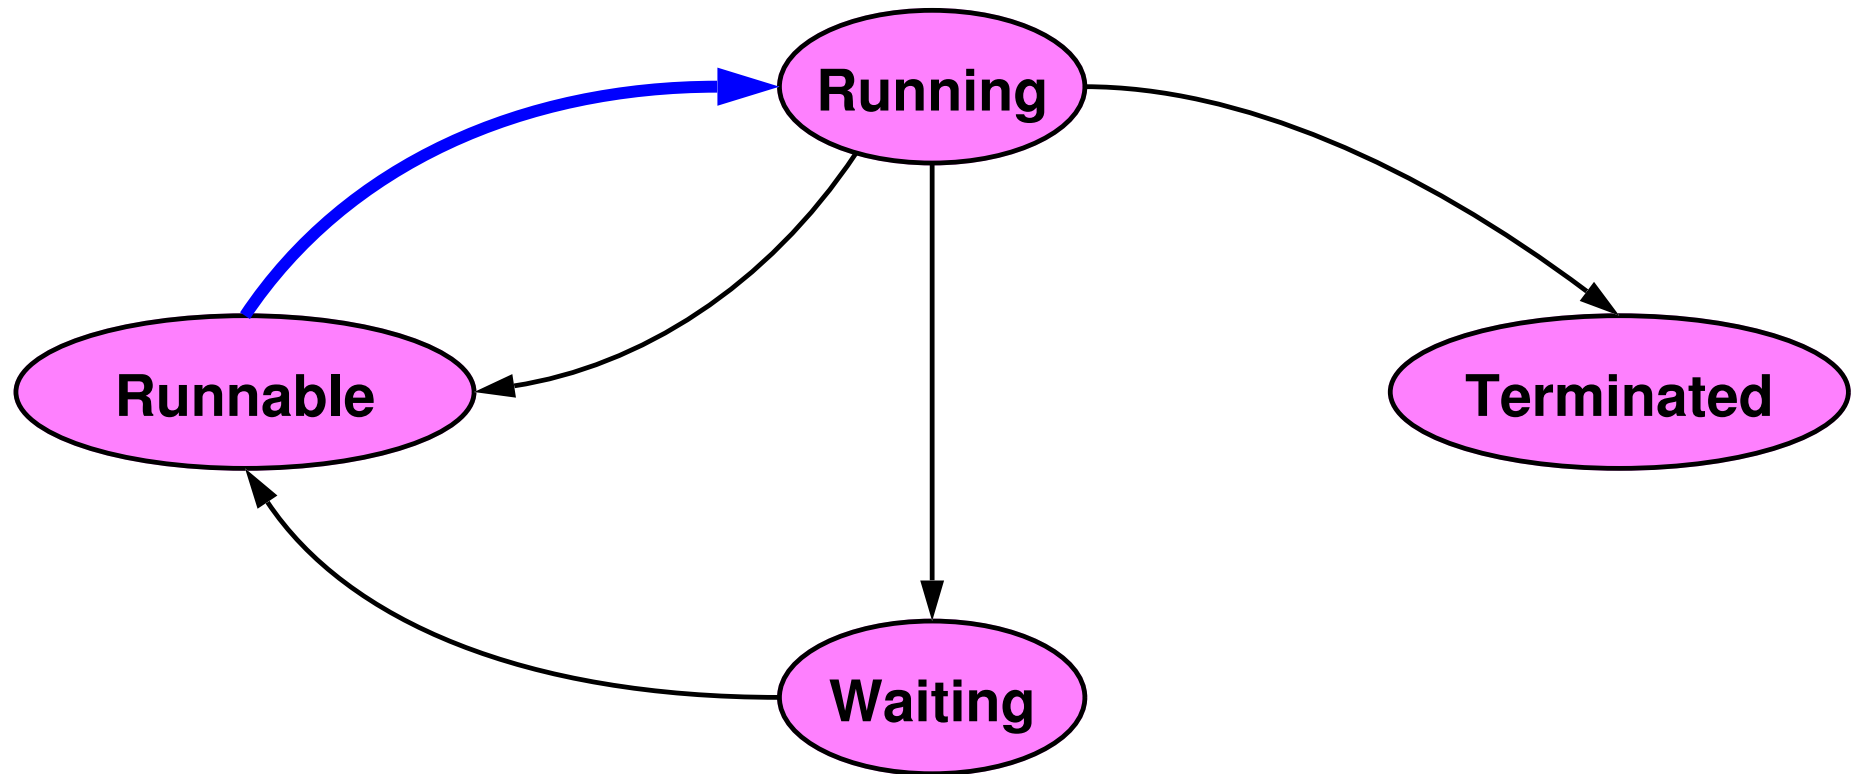
- ➡ Solaris uses the 2nd approach
 - expensive to fork a process
- ➡ Problem with 1st approach
 - thread 1 called `fork()` and thread 2 has a mutex locked
 - who will unlock the mutex?
 - POSIX solution is to provide a way to unlock all mutex before `fork()`

Thread Life Cycle




— a thread starts in the *runnable* state

Thread Life Cycle




- = a thread starts in the *runnable* state
- = the *scheduler* switches a thread's state from runnable to running

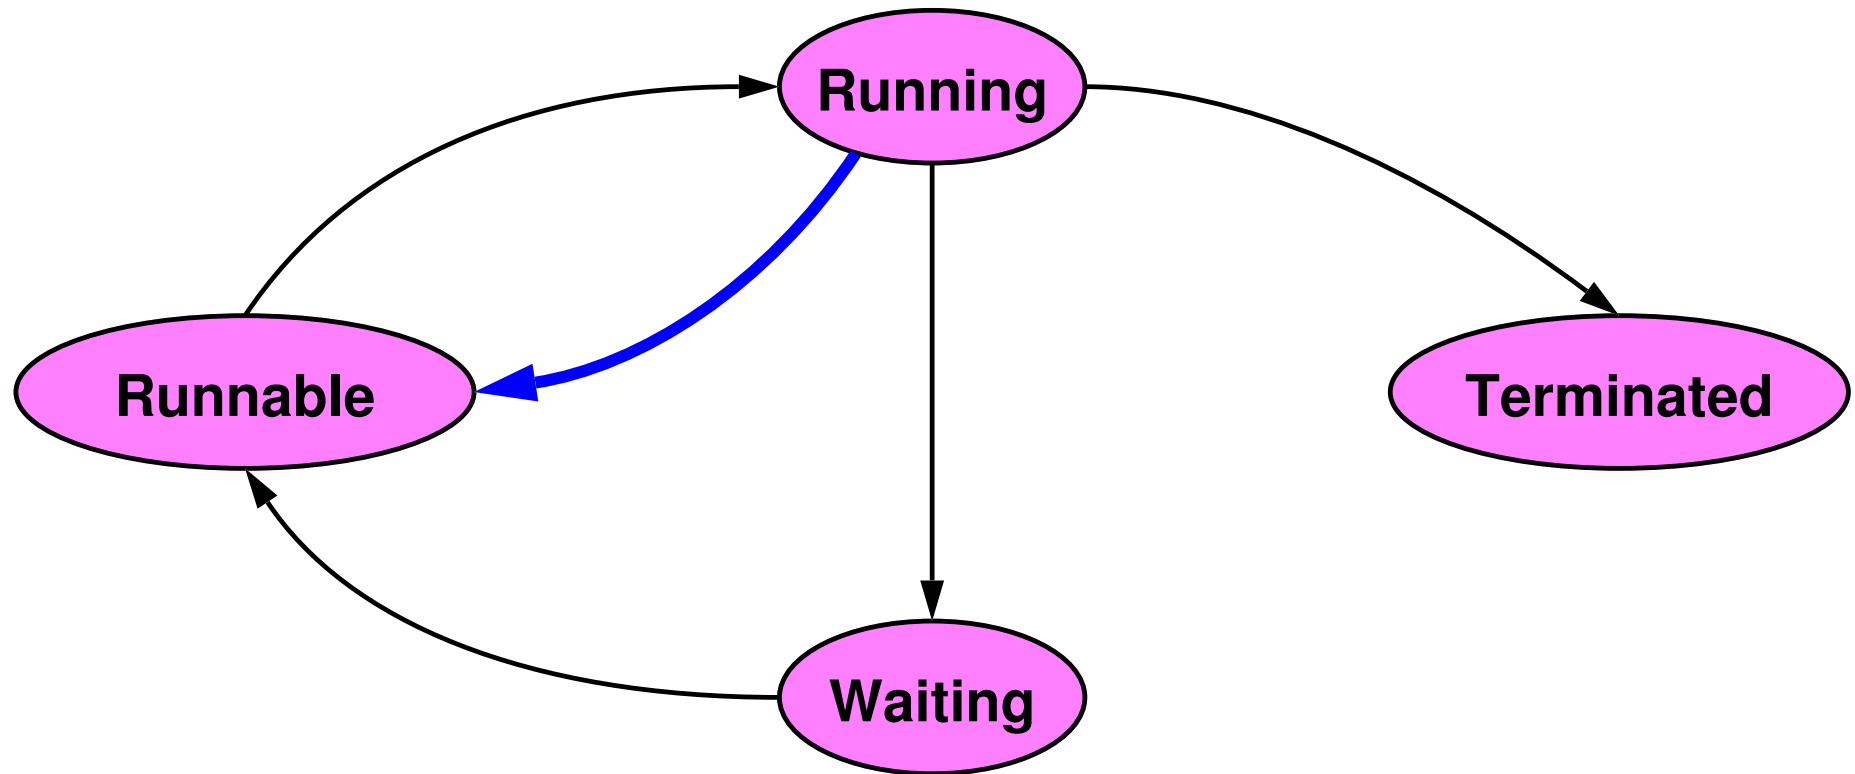
Thread Life Cycle



- ⇒ a thread starts in the *runnable* state
- ⇒ the *scheduler* switches a thread's state from runnable to running
- ⇒ a thread goes from running to waiting when a *blocking call* is made

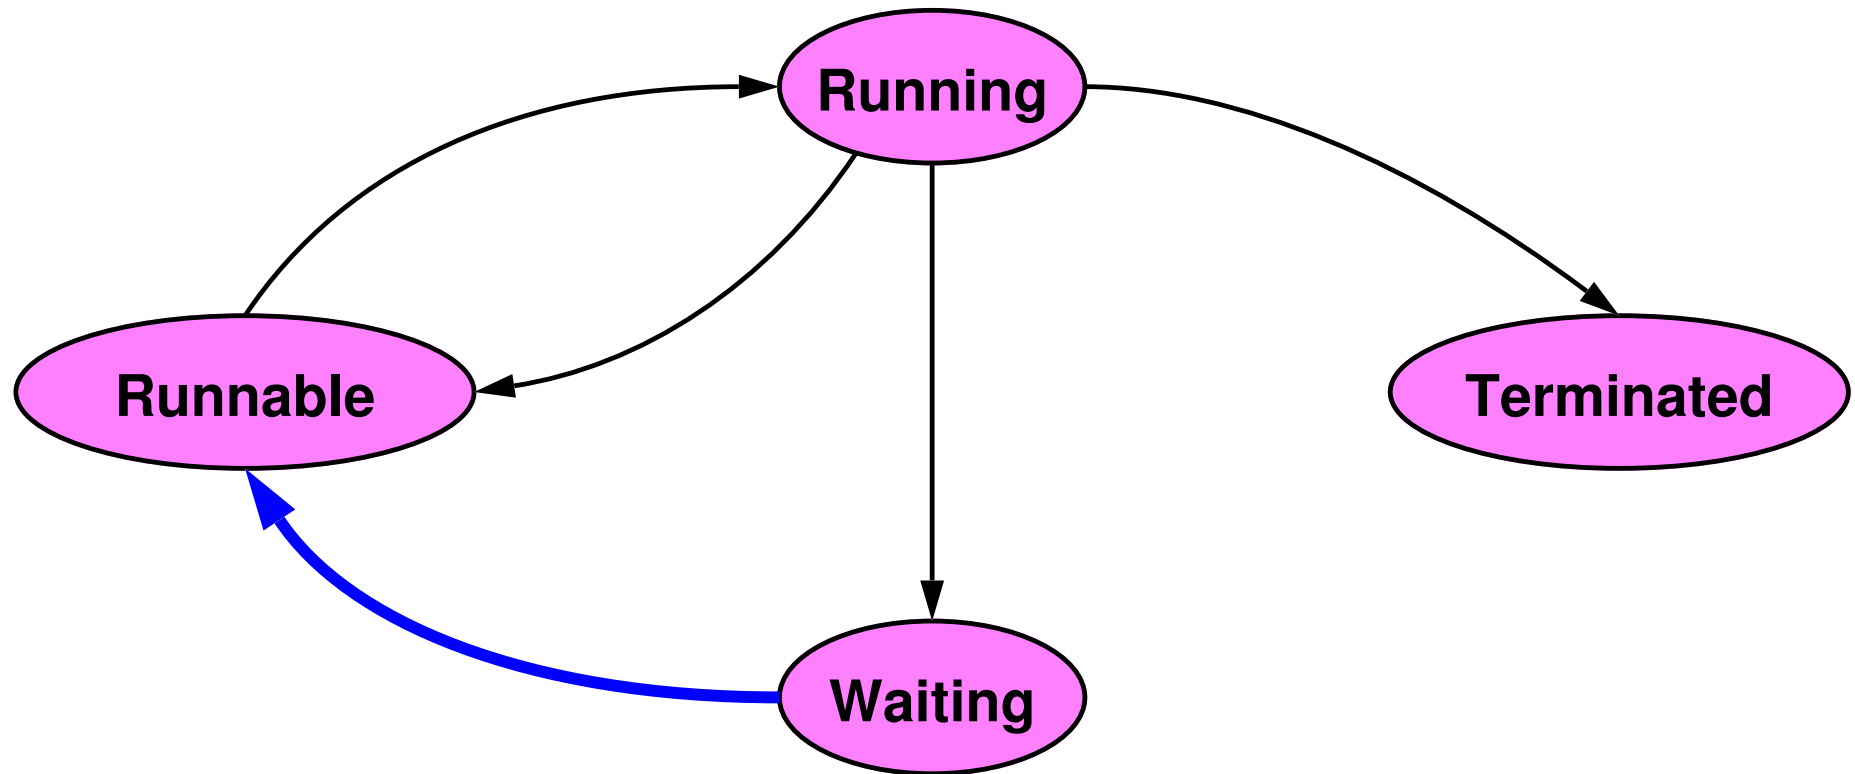


Thread Life Cycle




- a thread starts in the *runnable* state
- the *scheduler* switches a thread's state from runnable to running
- a thread goes from running to waiting when a *blocking call* is made
- the *scheduler* switches a thread's state from running to runnable when the thread used up its execution quantum

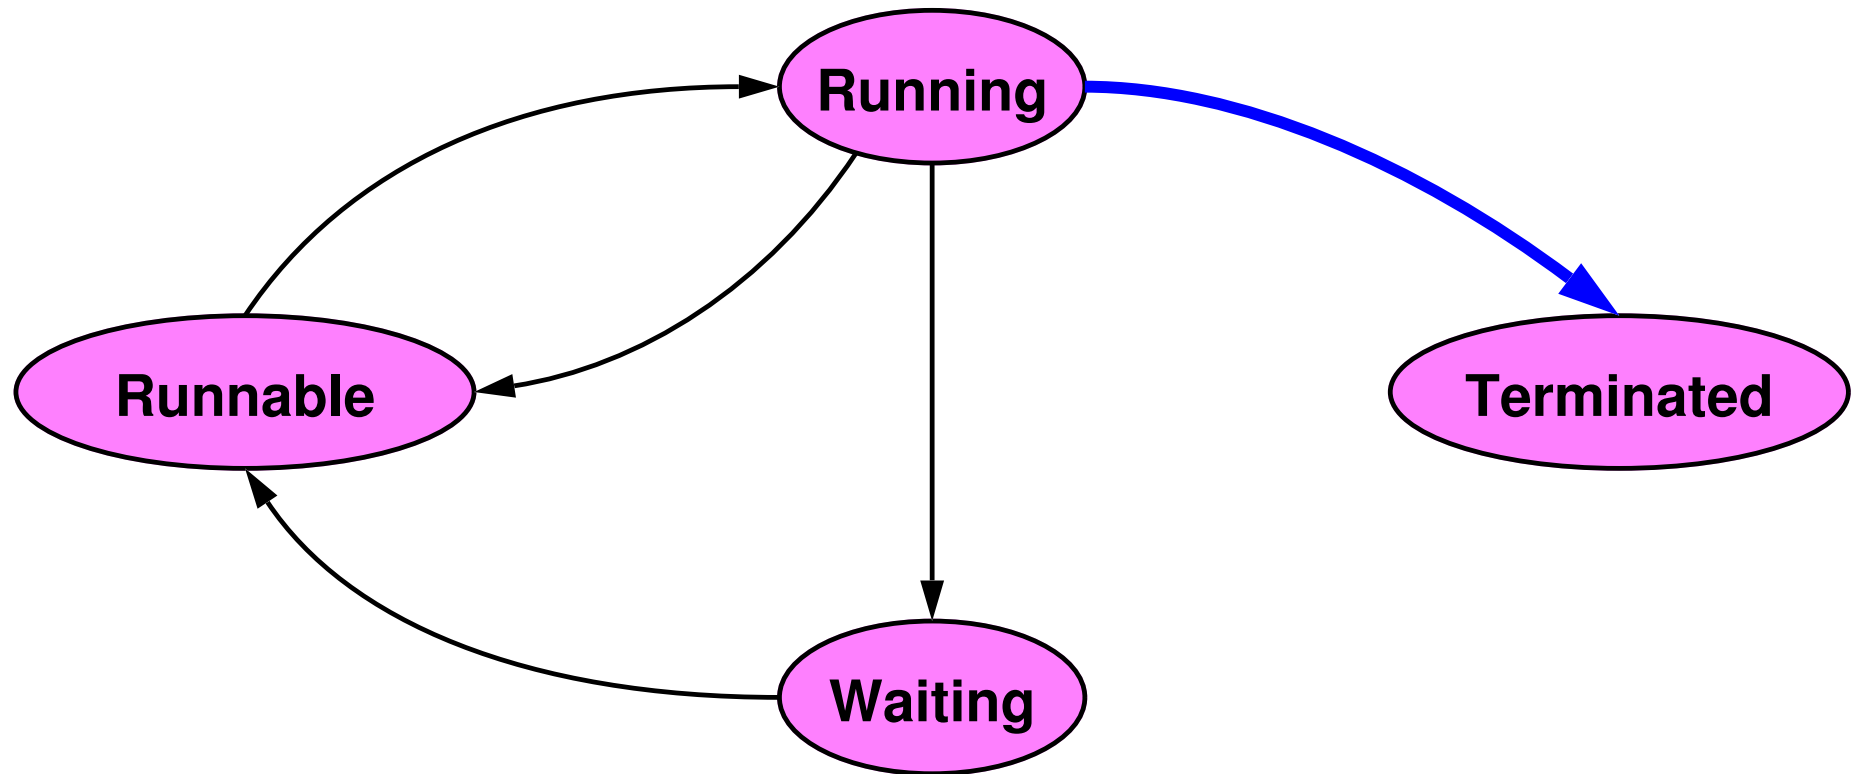
Thread Life Cycle




- a thread get *unblocked* by the action of another thread or by an interrupt handler



Thread Life Cycle




- a thread get **unblocked** by the action of another thread or by an interrupt handler
- in order for a thread to enter the terminated state, it has to be in the running state just before that
 - what if `pthread_cancel()` is invoked when the thread is not in the running state?



Thread Life Cycle

- ➡ Does `pthread_exit()` delete the thread (completely) that calls it?
 - no
- ➡ What's left in the thread after it calls `pthread_exit()`?
 - its thread control block
 - needs to keep thread ID and return code around
 - its stack
 - how can a thread *delete its own stack*? no way!



Thread Life Cycle

- ➡ Who is deleting the *thread control block* and freeing up the thread's *stack* space?
- ➡ If a thread is not detached
 - it can be taken care of in the `pthread_join()` code
 - the thread that calls `pthread_join()` does the clean up
- ➡ If a thread is detached (our simple OS does not support this)
 - can do this is one of two ways
 - 1) use a special *reaper thread*
 - ◆ basically doing `pthread_join()`
 - 2) queue these threads on a list and have other threads free them when it's convenient (e.g., when the scheduler schedule a thread to run)

