

4.1 A Simple System (Monolithic Kernel)

- A Framework for Devices
- Low-level Kernel (will come back to talk about this after Ch 7)
- Processes & Threads
- Storage Management



Copyright © William C. Cheng

Storage Space

- Where to store data?
 - **primary storage**, i.e., physical memory
 - directly addressable
 - **secondary storage**, i.e., disk-based storage
- What would it take to support the idea of virtual memory, i.e., application's "view" of memory?
- An application only works with "virtual memory" (as far as an application is concerned, "virtual memory" is "real memory")
 - e.g., map a 1GB file into memory
 - this memory is **virtual memory**
 - can **allocate** 1GB of **virtual memory** while there's only 256MB of **physical memory**
 - the OS makes sure that real primary storage is available when necessary
- **Virtual Memory** ties everything together!



Copyright © William C. Cheng

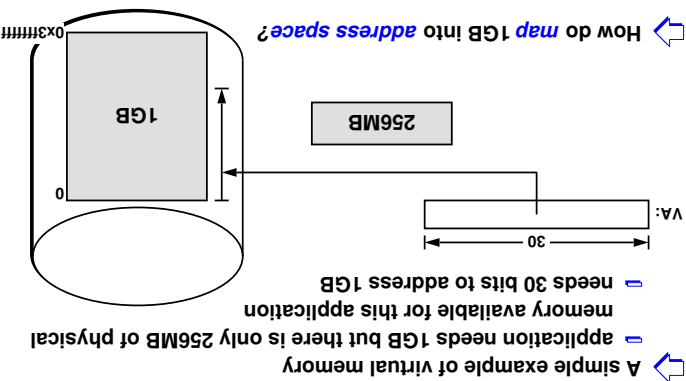
Memory Management Concerns

- **Mapping** virtual addresses to real ones
- Determining which addresses are **valid**, i.e., refer to allocated memory, and which are not
- Keeping track of which real objects, if any, are mapped into each range of virtual addresses
- Deciding what should to keep in primary storage (RAM) and what to fetch from elsewhere



Copyright © William C. Cheng

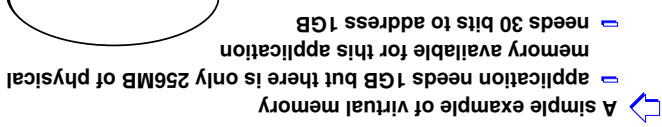
Storage Space

- A simple example of virtual memory
 - application needs 1GB but there is only 256MB of physical memory available for this application
 - needs 30 bits to address 1GB
- How do **map** 1GB into **address space**?
 



Copyright © William C. Cheng

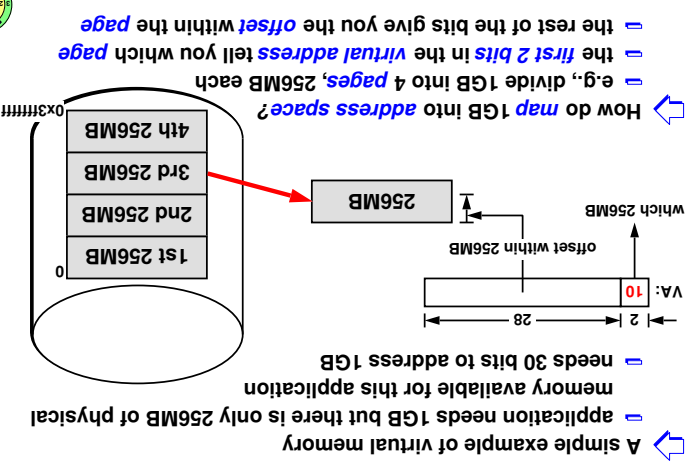
Storage Space

- A simple example of virtual memory
 - application needs 1GB but there is only 256MB of physical memory available for this application
 - needs 30 bits to address 1GB
- How do **map** 1GB into **address space**?
 



Copyright © William C. Cheng

Storage Space

- A simple example of virtual memory
 - application needs 1GB but there is only 256MB of physical memory available for this application
 - needs 30 bits to address 1GB
- How do **map** 1GB into **address space**?
 



Copyright © William C. Cheng

[illegible]

Storage Space

Hardware Memory Map

Start	Access	Physical Addr
0	-	-
4096	R	-
8192	R	-
12288	R	-
16384	R/W	-

Page Table

Segmentation Fault

If it's the wrong 256MB that's in primary memory

Storage Space

Storage Space

The diagram illustrates a 4-way set-associative cache. On the left, a vertical stack of four 256MB sets is shown, indexed from 0 to 3. The 4th set is highlighted. An arrow points from this set to a larger 256MB block. This block contains a 10-bit tag and a 26-bit data field. The 10-bit tag is compared with the value 11 in the VA register. The 28-bit address is split into a 2-bit tag and a 26-bit data field.

Storage Space

Copyright © William C. Cheng

In Reality, Have To Deal With Mapped Files

- An example to demonstrate a dilemma
 - one process is using all of its primary storage allocation
 - it then maps a file into its address space and starts accessing that file
 - should the memory that's needed to buffer this file be charged against the files subsystem or charged against the process?
- if charged against the files subsystem
 - if the newly mapped file takes up all the buffer space in the files subsystem, it's unfair to other processes
- if charged against the process
 - if other processes are sharing the same file, other processes are getting a free ride (in terms of memory usage)
 - even worse, another process may increase the memory usage of this process (double unfair!)

Operating Systems - CSCI 402

Copyright © William C. Cheng

How OS Makes Virtual Memory Work?

- if a thread access a virtual memory location that's both in primary memory and mapped by the hardware's map
 - no action by the OS
- if a thread access a virtual memory location that's *not in primary memory* or if the *translation is not in the memory map*
 - a *page fault* is occurred and the OS is invoked
 - OS checks the `as_region` address space data structures to make sure the reference is valid
 - if it's valid, the OS does whatever that's necessary to locate or create the object of the reference
 - find, or if necessary, make room for it in primary storage
 - if it's not already there, and put it there
- details in Ch 7

Two issues need further discussion

- how is the *primary storage* managed?
- how are these objects managed in *secondary storage*?

Operating Systems - CSCI 402

Copyright © William C. Cheng

Address Space Representation

recall that there is something called "address space description" in a PCB

This is related to Kernel Assignment 3 where you need to create and manage *address spaces* / *memory maps*

- `as_region` (address space region data structure) contains:
 - start address, length, access permissions, shared or private*
 - if mapped to a file, pointer to the corresponding *file object*
- This is related to Kernel Assignment 3 where you need to create and manage *address spaces* / *memory maps*

Operating Systems - CSCI 402

Copyright © William C. Cheng

In Reality, Have To Deal With Mapped Files

- it's difficult to be *fair*
 - it's difficult to even define what *fair* means
- We will discuss some solutions in Ch 7
 - for now, we use the following solution
 - give each participant (processes, file subsystem, etc.) a minimum amount of storage
 - leave some additional storage available for all to compete

Operating Systems - CSCI 402

Copyright © William C. Cheng

How Is The Primary Storage Managed?

- Who needs primary memory?
 - application processes
 - terminal-handling subsystem
 - communication subsystem
 - I/O subsystem
- They *compete* for available memory
 - it's difficult to be "fair" (what does it even mean?)
- if primary memory is managed poorly
 - one subsystem can use up all the available memory
 - then other subsystem won't get to run
 - this many lead to OS crash when a subsystem runs out of memory
- if there are no mapped files, the solution can be simple
 - equally divide the primary memory among the participants
 - this way, they won't compete

Operating Systems - CSCI 402

Copyright © William C. Cheng

Address Space Representation

In this example, text and data map portions of the same file

- text* is marked read-execute and *shared*
- data* is marked read-write and *private* to mean that changes will be private, i.e., will not affect other processes executed from the same file

Operating Systems - CSCI 402





File System Cache


- ➡ **Recently used blocks** in a file are kept in a *file system cache*
- ➡ the primary storage holding these blocks might be mapped into one or more address spaces of processes that have this file mapped
 - blocks are available for immediate access by read and write system calls
- ➡ A simple *hash function* is used to locate file blocks in the cache
 - ➡ keyed by *inode number*
- ➡ More details in Ch 6



File Object

-  The file object is like an *abstract class* in C++
-  = subclasses of file object are the *actual* file objects
- ```

class FileObject
{
 unsigned short recount;
 unsigned short access;
 unsigned int file_pos;
 ...
 virtual int create(const char * , int , FileO
 virtual int read(int , void * , int) ;
 virtual int write(int , const void * , int) ;
 ...
 }

```
-  But wait ...
- what's this about C++?
  - real operating systems are written in C ...
  - check out the DRIVERS kernel documentation (we skipped this week's assignment)



## How Are Objects Managed in Secondary Storage?

- ➡ The *file system* is used to manage objects in secondary storage
- ➡ The file system is usually divided into two parts
  - ➡ *file system independent*
    - supports the "file abstraction"
    - on Windows, this is called the "*/O manager*"
    - on Unix, this is called the "*virtual file system (VFS)*"
  - ➡ *file system dependent*
    - ◇ Kernel Assignment 2
    - on Windows, this is called the "file system"
    - on Unix, this is called the "actual file system"



## File Object in C

```
typedef struct
{ unsigned short
 !recount;
```

```
struct file_ops *file_op; /* function pointers (can use indirection) */
} fileObject;
```

- A file object uses an *array of function pointers*
- this is how C implements *C++ polymorphism*
- one for each operation on a file
- where they point to is (actual) file system dependent
- but the (virtual) interface is the same to higher level of the OS
- Loose coupling between the actual file system and storage devices
- The actual file system is written to talk to the devices in a device-independent manner
- i.e., using major and minor device numbers to reference the device and using standard interface provided by the device driver



## Open-File Data Structures

