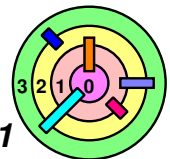


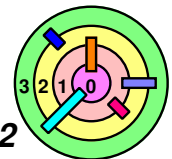
# 4.1 A Simple System (Monolithic Kernel)

- ➡ A Framework for Devices
- ➡ Low-level Kernel (will come back to talk about this after Ch 7)
- ➡ Processes & Threads
- ➡ *Storage Management*




# Storage Space

- ➡ Where to store data?
  - ➡ *primary storage*, i.e., physical memory
    - directly addressable
  - ➡ *secondary storage*, i.e., disk-based storage
- ➡ What would it take to support the idea of virtual memory, i.e., application's "view" of memory?
- ➡ An application only works with "virtual memory" (as far as an application is concerned, "virtual memory" is "real memory")
  - ➡ e.g., map a 1GB file into memory
    - this memory is *virtual memory*
  - ➡ can *allocate* 1GB of *virtual memory* while there's only 256MB of *physical memory*
  - ➡ the OS makes sure that real primary storage is available when necessary
- ➡ *Virtual Memory* ties everything together!




# Memory Management Concerns

- ➡ **Mapping** virtual addresses to real ones
- ➡ Determining which addresses are **valid**, i.e., refer to allocated memory, and which are not
- ➡ Keeping track of which real objects, if any, are mapped into each range of virtual addresses
- ➡ Deciding what should to keep in primary storage (RAM) and what to fetch from elsewhere




# Storage Space

- ➡ A simple example of virtual memory
- application needs 1GB but there is only 256MB of physical memory available for this application
  - needs 30 bits to address 1GB

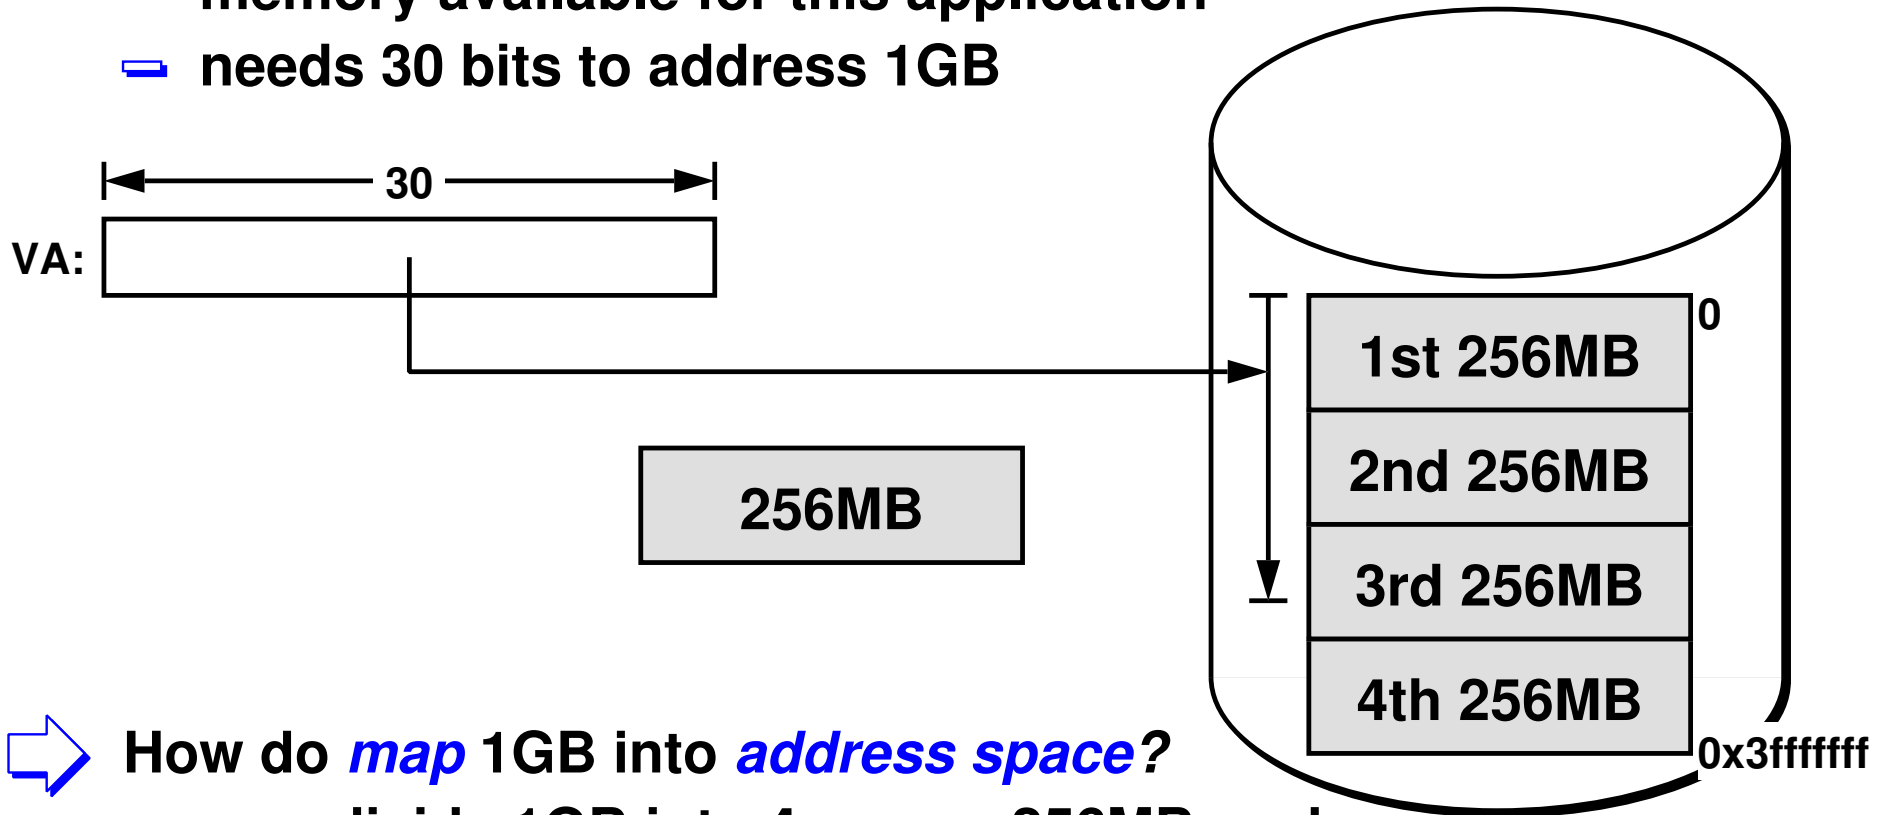


➡ How do *map* 1GB into *address space*?




# Storage Space

- ➡ A simple example of virtual memory
- application needs 1GB but there is only 256MB of physical memory available for this application
  - needs 30 bits to address 1GB




- ➡ How do *map* 1GB into *address space*?
- e.g., divide 1GB into 4 *pages*, 256MB each




# Storage Space

- ➡ A simple example of virtual memory
  - application needs 1GB but there is only 256MB of physical memory available for this application
  - needs 30 bits to address 1GB

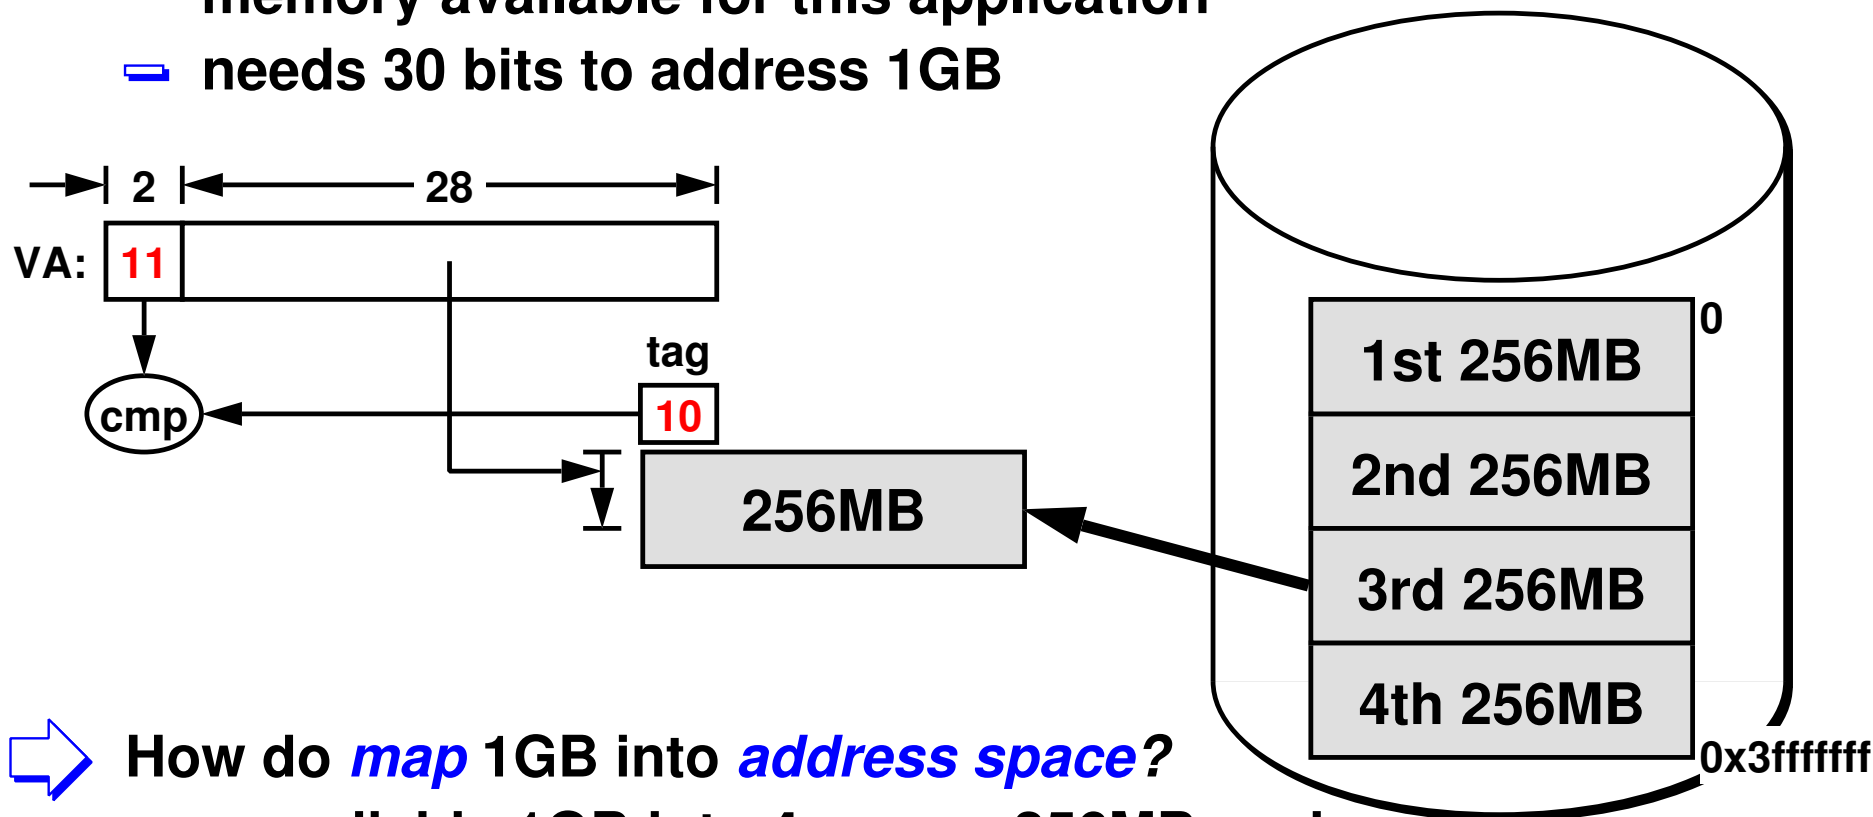


- ➡ How do *map* 1GB into *address space*?
  - e.g., divide 1GB into 4 *pages*, 256MB each
  - the *first 2 bits* in the *virtual address* tell you which *page*
  - the rest of the bits give you the *offset* within the *page*




# Storage Space

- ➡ A simple example of virtual memory
  - ▢ application needs 1GB but there is only 256MB of physical memory available for this application
  - ▢ needs 30 bits to address 1GB




- ➡ How do *map* 1GB into *address space*?
  - ▢ e.g., divide 1GB into 4 *pages*, 256MB each
  - ▢ the *first 2 bits* in the *virtual address* tell you which *page*
  - ▢ the rest of the bits give you the *offset* within the *page*
  - ▢ check to see if the right page is in *physical* memory




# Storage Space

- ➡ A simple example of virtual memory
- application needs 1GB but there is only 256MB of physical memory available for this application
  - needs 30 bits to address 1GB



- ➡ If it's the wrong page that's in primary memory
- accessing it will cause a *page fault*
  - during a page fault, OS brings the right page into real memory
  - then the thread is allow to proceed with accessing the memory





# Segmentation Fault



➡ A valid virtual address must be ultimately *resolvable* by the OS to a location in the physical memory

- if it cannot be resolved, the virtual address is considered an *invalid* virtual address
- referencing an invalid virtual address will cause a *segmentation fault* (the OS will deliver SIGSEG to the process)
  - the default action would be to terminate the process
- e.g., virtual address 0

➡ A *page fault* is a *segmentation fault*


Page Table

<i>Start</i>	<i>Access</i>	<i>Physical Addr</i>
0	-	-
4096	R	•
8192	R	•
12288	R	•
16384	R/W	•




# Storage Space

- ➡ A simple example of virtual memory
  - ▬ application needs 1GB but there is only 256MB of physical memory available for this application
  - ▬ needs 30 bits to address 1GB

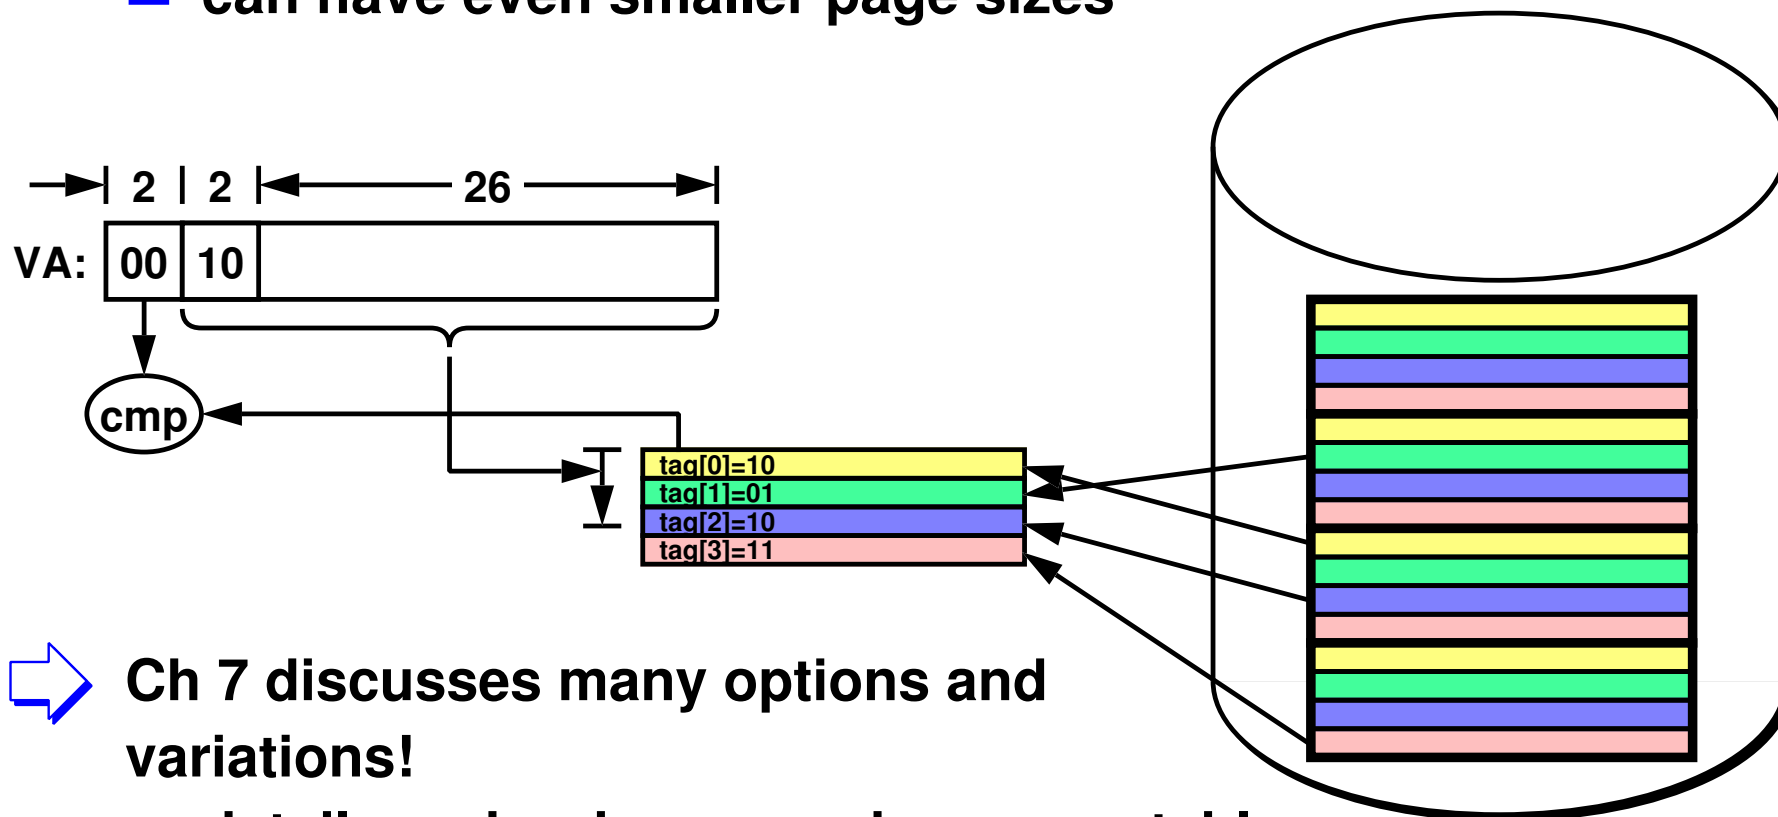


- ➡ If it's the wrong 256MB that's in primary memory
  - ▬ accessing it will cause a *page fault*
  - ▬ this "simple" approach has *really poor performance*
    - takes too long to copy 256MB
    - why just use 2 leading bits? different organizations?




# Storage Space

- ➡ A more complicated scheme with a smaller page size
  - compare to determine if there is a *hit* or not
  - can have even smaller page sizes

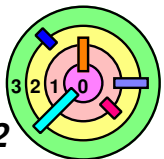


- ➡ Ch 7 discusses many options and variations!
  - details on hardware, such as page tables, translation look-aside buffers, etc.
  - details on OS software, such as how to implement *memory map*, *copy-on-write*, etc.

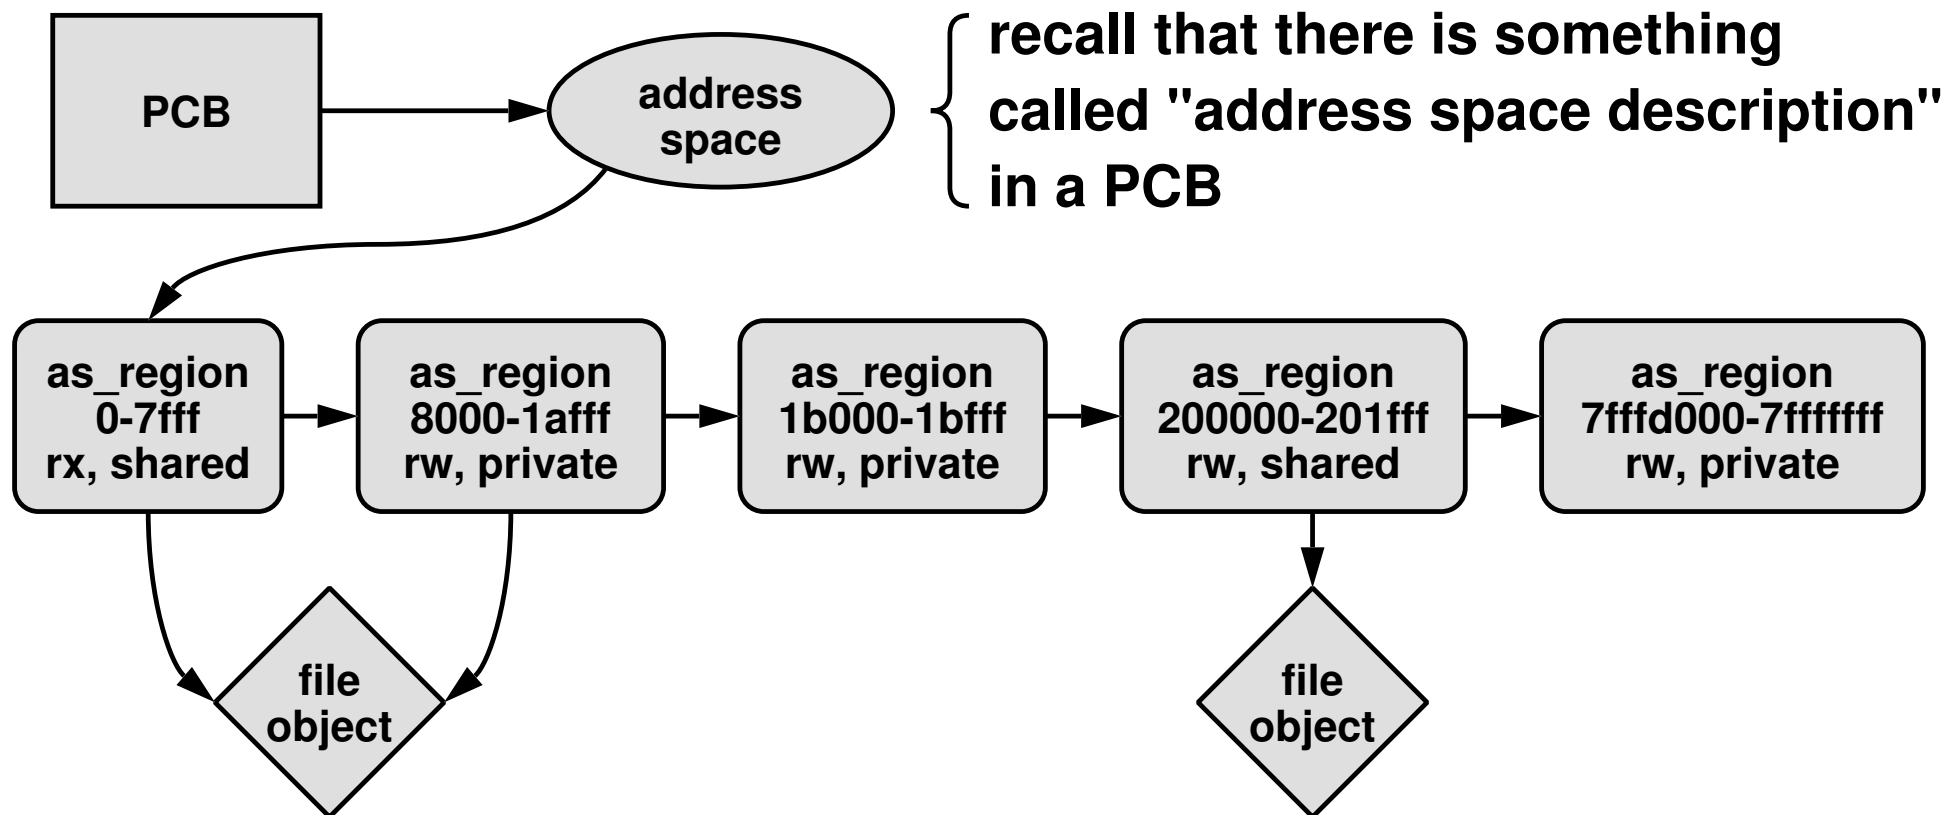


# Hardware Memory Map


- ➡ In reality, the OS is too slow since *every* virtual address needs to be resolved
- some of the virtual memory mechanisms must be built into the *hardware*
    - in some cases, the hardware is given the complete *"map"* (i.e., mapping from virtual to physical address)
    - in some other cases, only a partial map is given to the hardware
    - in either case, OS needs to provide some map to the hardware and needs a *data structure* for the map
      - ◆ often referred as the *memory map*, or *mmap*



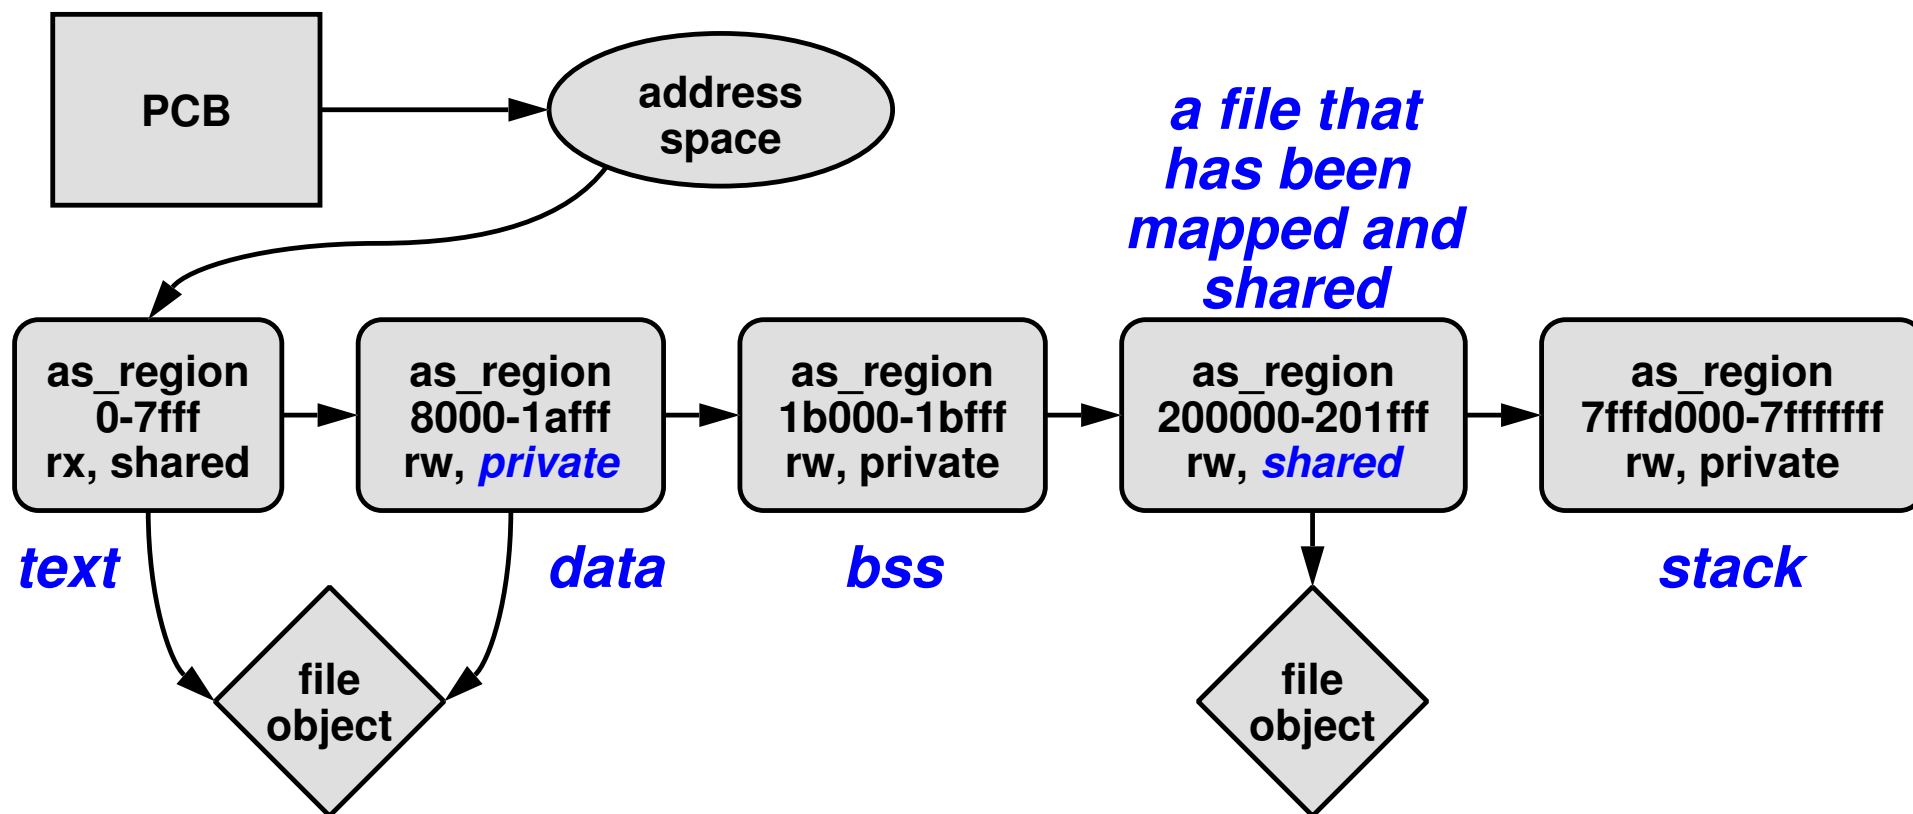
# Address Space Representation




- ➡ as\_region (address space region data structure) contains:
  - ▬ *start address*, *length*, *access permissions*, *shared* or *private*
  - ▬ if mapped to a file, pointer to the corresponding *file object*
- ➡ This is related to Kernel Assignment 3 where you need to create and manage *address spaces* / *memory maps*



# Address Space Representation




- ➡ In this example, text and data map portions of the same file
- **text** is marked read-execute and **shared**
  - data is marked read-write and **private** to mean that changes will be private, i.e., will not affect other processes exec'ed from the same file




# How OS Makes Virtual Memory Work?

- ➡ If a thread access a virtual memory location that's both in primary memory and mapped by the hardware's map
  - no action by the OS
- ➡ If a thread access a virtual memory location that's *not in primary memory* or if the *translation is not in the memory map*
  - a **page fault** is occurred and the OS is invoked
    - OS checks the `as_region` address space data structures to make sure the reference is valid
      - ◆ if it's valid, the OS does whatever that's necessary to locate or create the object of the reference
      - ◆ find, or if necessary, make room for it in primary storage if it's not already there, and put it there
      - ◆ details in Ch 7
- ➡ Two issues need further discussion
  - how is the *primary storage* managed?
  - how are these objects managed in *secondary storage*?



# How Is The Primary Storage Managed?

- ➡ Who needs primary memory?
  - application processes
  - terminal-handling subsystem
  - communication subsystem
  - I/O subsystem
- ➡ They *compete* for available memory
  - it's difficult to be "fair" (what does it even mean?)
- ➡ If primary memory is managed poorly
  - one subsystem can use up all the available memory
    - then other subsystem won't get to run
    - this many lead to OS crash when a subsystem runs out of memory
- ➡ If there are no mapped files, the solution can be simple
  - equally divide the primary memory among the participants
    - this way, they won't compete






# In Reality, Have To Deal With Mapped Files

- ➡ **An example to demonstrate a dilemma**
  - one process is using all of its primary storage allocation
  - it then maps a file into its address space and starts accessing that file
  - should the memory that's needed to buffer this file be charged against the files subsystem or charged against the process?
- ➡ **If charged against the files subsystem**
  - if the newly mapped file takes up all the buffer space in the files subsystem, it's unfair to other processes
- ➡ **If charged against the process**
  - if other processes are sharing the same file, other processes are getting a free ride (in terms of memory usage)
  - even worse, another process may increase the memory usage of this process (double unfair!)


# In Reality, Have To Deal With Mapped Files

- ➡ It's difficult to be *fair*
  - it's difficult to even define what *fair* means
- ➡ We will discuss some solutions in Ch 7
  - for now, we use the following solution
    - give each participant (processes, file subsystem, etc.) a minimum amount of storage
    - leave some additional storage available for all to compete

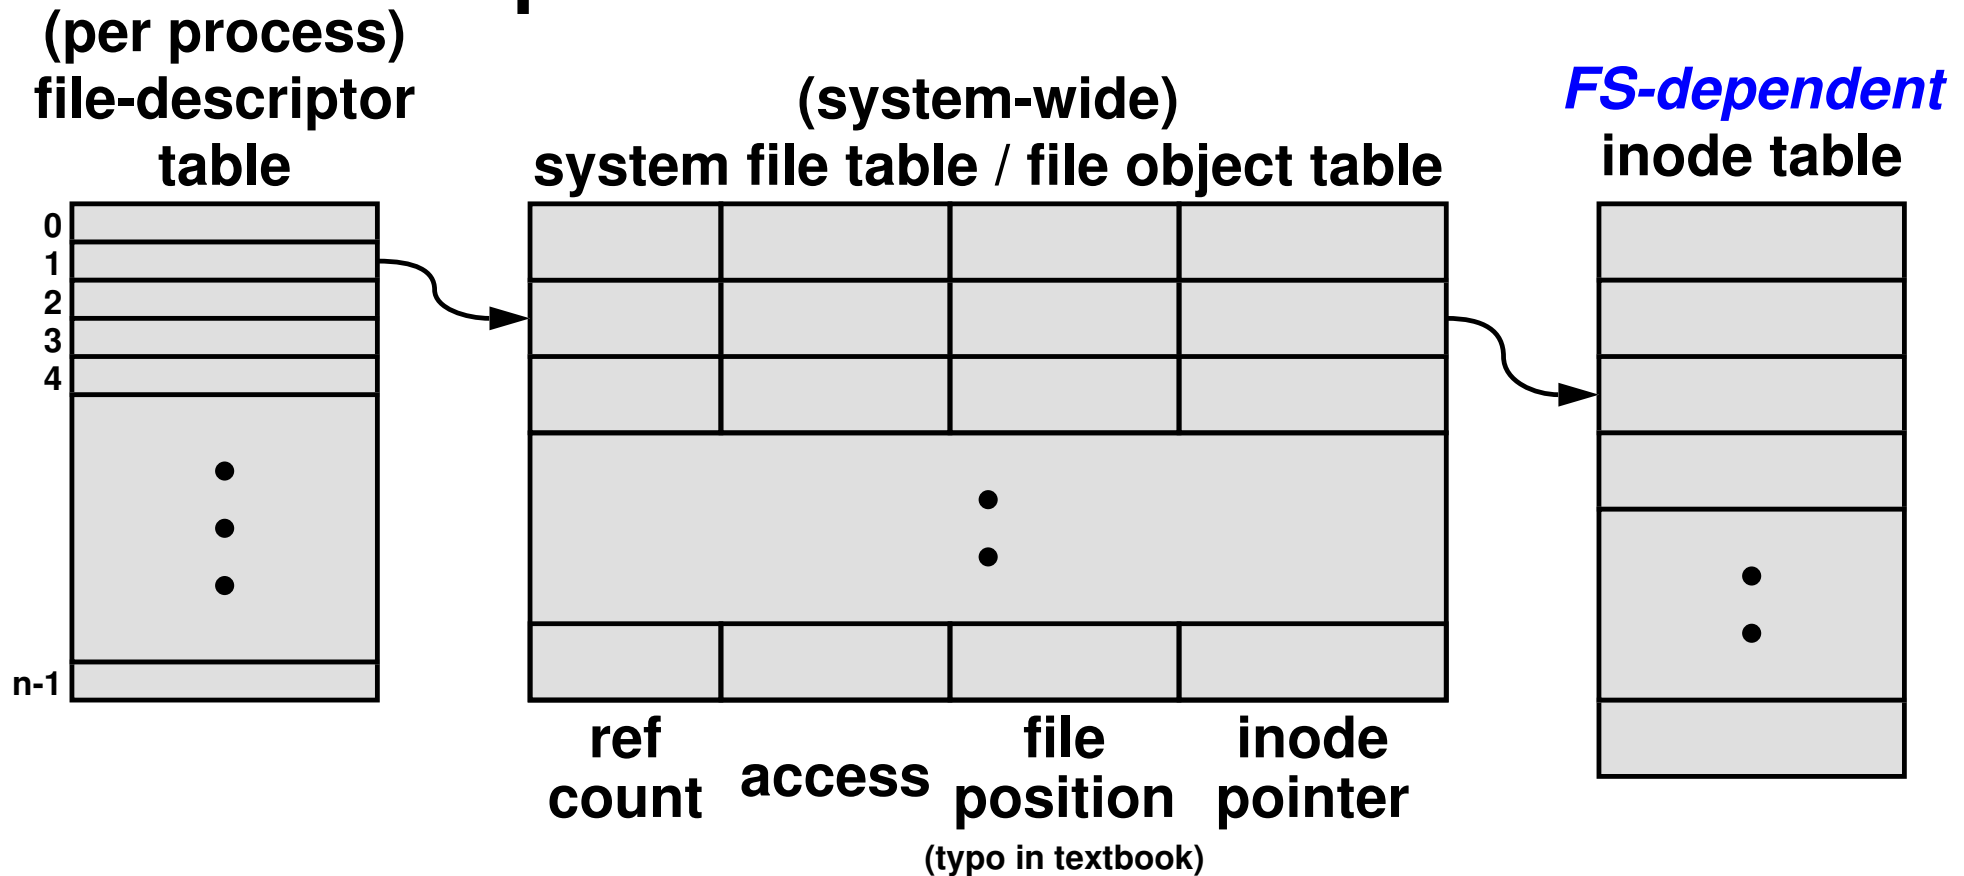


# How Are Objects Managed In Secondary Storage?


- ➡ The *file system* is used to manage objects in secondary storage
- ➡ The file system is usually divided into two parts
  - ⇒ *file system independent*
    - supports the "file abstraction"
    - on Windows, this is called the "*I/O manager*"
    - on Unix, this is called the "*virtual file system (VFS)*"
      - ◆ Kernel Assignment 2
  - ⇒ *file system dependent*
    - on Windows, this is called the "file system"
    - on Unix, this is called the "actual file system"



# Open-File Data Structures



- ➡ In the kernel, *each process* has its own *file-descriptor table*
  - ▢ the kernel also maintains *system file table* (or *file object table*)
- ➡ The *file object / inode* forms the *boundary* between *VFS* and the actual file system (i.e., points to *file-system-dependent* stuff)
  - ▢ how can this be done?




# File Object

- ➡ The file object is like an *abstract class* in C++  
— subclasses of file object are the *actual* file objects

```
class FileObject {  
    unsigned short refcount;  
    unsigned short access;  
    unsigned int file_pos;  
    ...  
    virtual int create(const char *, int, FileObject **);  
    virtual int read(int, void *, int);  
    virtual int write(int, const void *, int);  
    ...  
};
```


- ➡ But wait ...
- what's this about C++?
    - real operating systems are written in C ...
    - checkout the DRIVERS kernel documentation (we skipped this weenix assignment)



# File Object in C

```
typedef struct {  
    unsigned short refcount;  
    ...  
    struct file_ops *file_op;  
    /* function pointers (can use indirection) */  
} FileObject;
```

- ➡ A file object uses an *array of function pointers*
  - this is how C implements *C++ polymorphism*
  - one for each operation on a file
  - where they point to is (actual) file system dependent
  - but the (virtual) interface is the same to higher level of the OS
- ➡ Loose coupling between the actual file system and storage devices
  - the actual file system is written to talk to the devices in a device-independent manner
    - i.e., using major and minor device numbers to reference the device and using standard interface provided by the device driver



# File System Cache

- ➡ **Recently used blocks** in a file are kept in a **file system cache**
  - ▬ the primary storage holding these blocks might be mapped into one or more address spaces of processes that have this file mapped
    - blocks are available for immediate access by read and write system calls
- ➡ A simple **hash function** is used to locate file blocks in the cache
  - ▬ keyed by **inode number**
- ➡ More details in Ch 6

