

## 4.1 A Simple System (Monolithic Kernel)

- A Framework for Devices
- Low-level Kernel
- Processes & Threads
- Storage Management



Copyright © William C. Cheng

Operating Systems - CSCI 402



➤ VT100



Terminals

Copyright © William C. Cheng

Operating Systems - CSCI 402

- Issues
  - 1) terminals are slow and characters generation are too fast
    - need to tell the application to slow down and wait for the terminal to catch up
    - so, we need an *output buffer* to buffer the output and send characters to the terminal from the buffer
    - we have an instance of the producer-consumer problem! characters arrive from the keyboard even though there isn't a waiting read request from an application
    - so, we need an *input buffer* to buffer incoming characters and wait for an application to issue a read request
    - we have another instance of the producer-consumer problem!



Copyright © William C. Cheng

Operating Systems - CSCI 402



Copyright © William C. Cheng

- Let's talk about how devices are handled, starting at the lowest levels of the kernel
  - (although *bottom-up* is not a good way to *design* an OS) but it may be a reasonable way to implement OS components
- We will start by looking at two such devices
  - terminals
  - network communication

### Low-Level Kernel

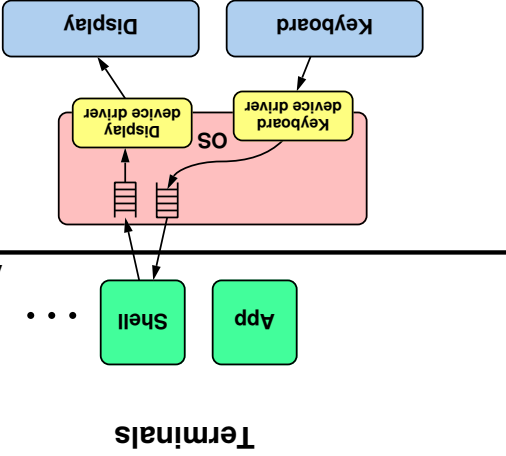


- Long obsolete, but still relevant
  - on Linux, you would probably use a "terminal" program (such as `xterm` or `gnome-terminal`) to interact with the system
  - on Windows, `putty` or `xwin-32` brings up a "terminal" for you to interact with a remote system
  - `ssh` client program interact with `sshd` on a server
  - once authenticated, `sshd` forks to exec `tcsh`/`bash`
  - you login session is *on the target machine*
    - i.e., if you login as root and type "halt", you would halt the machine!
- How to interact with a *terminal device*?
  - characters to be displayed are simply sent to the output routine of the *serial-line driver*
  - to fetch characters that have been typed at the keyboard, a call can be made to its input routine
  - as it turns out, not to straight-forward

### Terminals

Copyright © William C. Cheng

Operating Systems - CSCI 402



Copyright © William C. Cheng

Operating Systems - CSCI 402

Copyright © William C. Cheng

## Terminal

3) input characters may need to be processed in some way before they reach the application

- ◇ e.g., characters typed at the keyboard are *echoed* back to the display
- ◇ characters may be grouped into lines of text and subject to simple *editing* (such as backspace)
- ◇ some applications prefer to process all characters themselves, including their editing

Additional issue

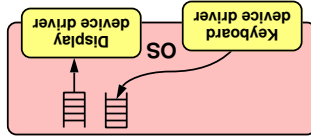
Copyright © William C. Cheng

Copyright © William C. Cheng

## Terminals

To deal with concerns (1) and (2)

- ◇ how about just using an *interrupt handler*?
- in the *read-completion interrupt*, the handler moves one character from the device to the input queue and issue another read request to the device and blocks
- ◇ if the queue is full, the character is thrown away
- ◇ the application thread must *mask interrupts* when it's taking a character from the queue
- ◇ can do the same for the output queue ...



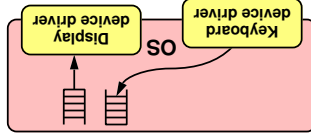
Copyright © William C. Cheng

Copyright © William C. Cheng

## Terminals

To deal with concerns (1) and (2)

- ◇ use two queues, one for input and one for output
- ◇ characters are placed on the output queue and taken from the input queue
- i.e., application write to the output queue and read from the input queue
- a thread producing output would block if output queue is full
- a thread consuming input would block if input queue is empty
- ◇ what about the other ends of these queues? who are handling them?



Copyright © William C. Cheng

Copyright © William C. Cheng

## Terminals

To deal with concern (3)

- ◇ remember, once you allow an application to take a character, you cannot ask for it back
- can only give it to the application when there is no chance that you will want it back
- ◇ this happens when a line is completed
- ◇ therefore, we need *two input queues*
- one for the *partial-line*
- ◇ subject to editing
- the other contain characters from *completed lines*
- in the read-completion interrupt, the handler moves one character from the device to the *partial-line queue*
- ◇ if the input character is a carriage-return, the entire content of the partial-line queue is moved to the completed-line queue

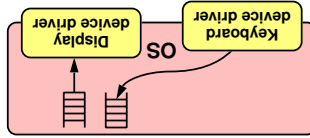
Copyright © William C. Cheng

Copyright © William C. Cheng

## Terminals

To deal with concerns (1) and (2)

- ◇ how about just using an *interrupt handler*?
- ◇ can do the same for the output queue
- in the *write-completion interrupt*, the handler moves one character from the output queue to the device and issue another write request to the device and blocks
- ◇ if the application writes to an *empty queue*, it would *setup the write-completion interrupt handler* and issue a write request to the device



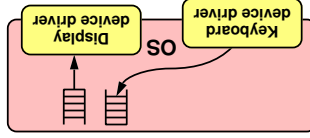
Copyright © William C. Cheng

Copyright © William C. Cheng

## Terminals

To deal with concerns (1) and (2)

- ◇ how about using a keyboard reading thread (that would do the following)?
- 1) issue a read to the device
- 2) block itself and wait for interrupt from the device
- 3) when interrupt occurs, the thread is woken up
- 4) the thread reads from the device and move one character from the device to the input queue
- 5) goto step 1
- ◇ this approach of using *thread context* seems to be an overkill and may be inefficient



Copyright © William C. Cheng

Copyright © William C. Cheng

## Where to Put the Modules?

Where to put the terminal driver and the line-discipline module?

- 1) kernel
- 2) separate user process
- 3) library routines that are linked into application processes

➡ **driver** should be in the kernel since device registers access needs to be protected from arbitrary manipulation by application programs

➡ **line-discipline** may be shared by multiple applications putting it in **library routines** will make it **difficult to share**

➡ one terminal with many user applications can it go into a separate **user process**?

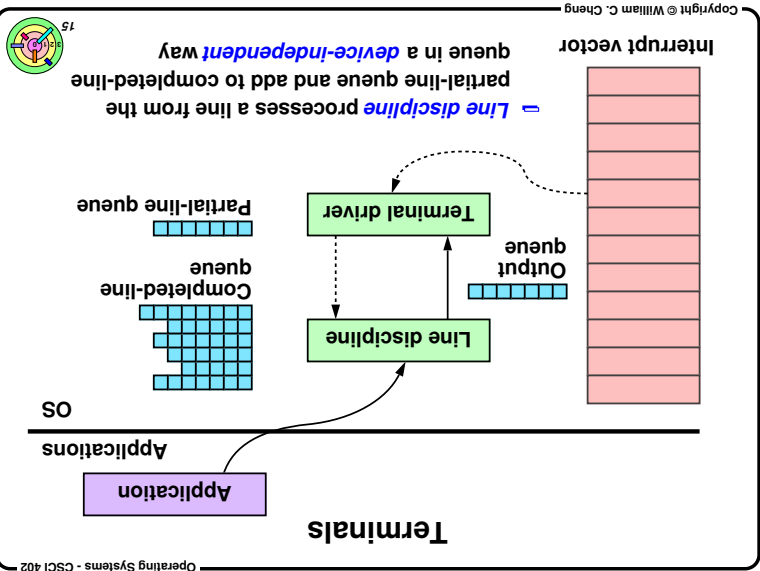
➡ but can have serious **performance problems** would need to transfer characters into the line-discipline process, then transfer to another process

➡ putting it in the **kernel** seems to be the best choice

➡ although kernel code is **hard to modify**, replace, and debug

17

Operating Systems - CSCI 402



Copyright © William C. Cheng

## Terminals

➡ To deal with concern (3)

➡ when a character is typed, it should go to the display immediately (due to the **echoing** requirement)

➡ it may be competing with the output thread, but it's okay (and that's how it's done in Unix)

➡ Windows handle this differently

➡ typed characters are only echoed when an application consumes them

➡ therefore, echoing is not done in the interrupt context

➡ echoing is done in the context of the thread consuming the characters (i.e., the one that calls read())

13

Operating Systems - CSCI 402

Copyright © William C. Cheng

## Terminals and Pseudo Terminals

➡ Modern systems do not have terminals

➡ they often have bit-mapped displays, keyboards and mice connected via USB

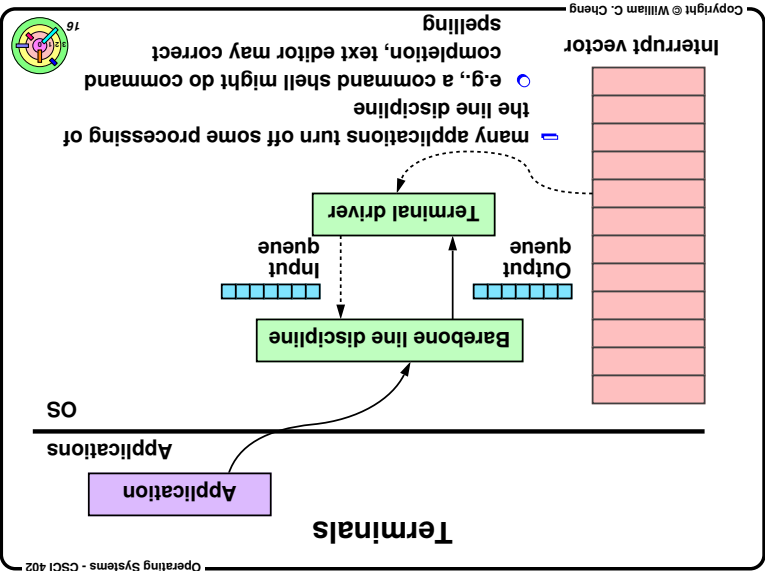
➡ a **window manager** implements windows on the display and determines which applications receive typed input (**input focus**)

➡ a server might support remote sessions where applications receive input and send output over a network

➡ they use **pseudoterminals** which implements a line discipline whose input comes from and output goes to a **controlling application** (and not a physical device)

16

Operating Systems - CSCI 402



Copyright © William C. Cheng

## Modularization

➡ **Device independence** consideration (figure out the **common part**)

➡ is common

➡ actually, character processing is performed in situations where the source and sink of characters aren't even a serial line

➡ e.g., bit-mapped display, network connection

➡ therefore, it makes sense to separate the device dependent part from the **common, device independent** part

➡ it can interact with any device driver capable of handling terminals

➡ can even use a different line-discipline module to deal with an alternative character set

14

Operating Systems - CSCI 402

