



4.1 A Simple System (Monolithic Kernel)

- ➡ A Framework for Devices
- ➡ *Low-level Kernel*
- ➡ Processes & Threads
- ➡ Storage Management



Low-Level Kernel


- ➡ Let's talk about how devices are handled, starting at the lowest levels of the kernel
 - (although *bottom-up* is not a good way to *design* an OS)
 - but it may be a reasonable way to implement OS components
- ➡ We will start by looking at two such devices
 - terminals
 - network communication



Terminals




➡ VT100


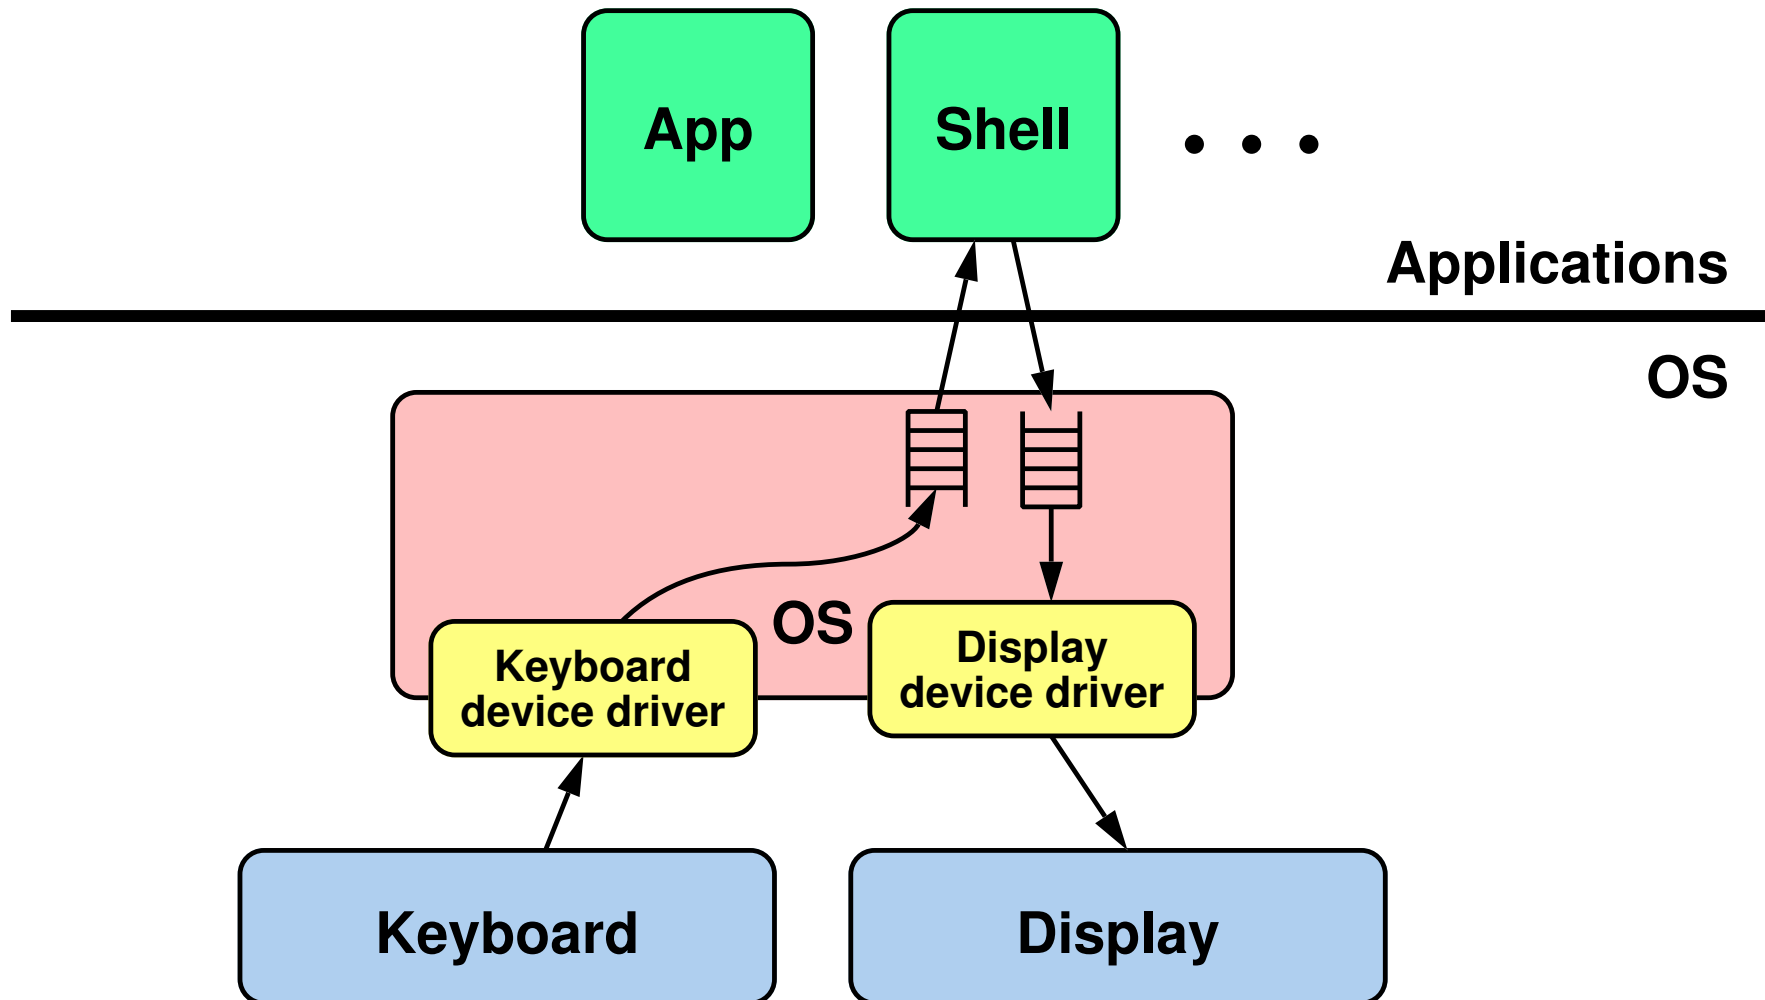


Terminals

- ➡ Long obsolete, but still relevant
 - on Linux, you would probably use a "terminal" program (such as `xterm` or `gnome-terminal`) to interact with the system
 - on Windows, `putty` or `xwin-32` brings up a "terminal" for you to interact with a remote system
 - `ssh` client program interact with `sshd` on a server
 - ◆ once authenticated, `sshd` forks to `exec tcsh/bash`
 - you login session is *on the target machine*
 - ◆ i.e., if you login as root and type "`halt`", you would halt the machine!
- ➡ How to interact with a *terminal device*?
 - characters to be displayed are simply sent to the output routine of the *serial-line driver*
 - to fetch characters that have been typed at the keyboard, a call can be made to its input routine
 - as it turns out, not to straight-forward


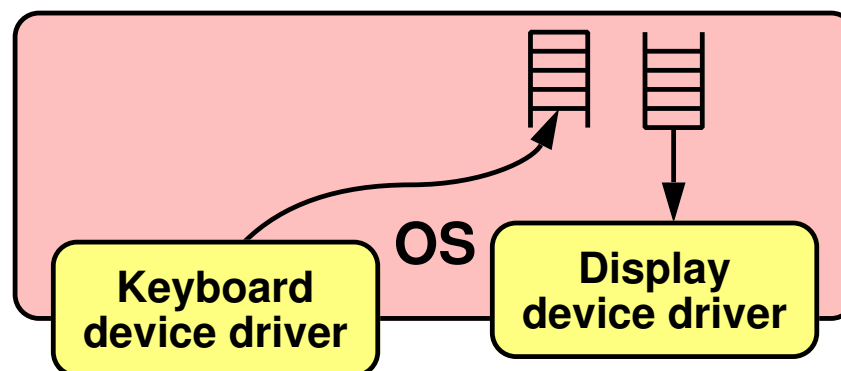


Terminals





Terminals

- ➡ To deal with concerns (1) and (2)
- use two queues, one for input and one for output
 - characters are placed on the output queue and taken from the input queue in the context of the *application thread*
 - i.e., application write to the output queue and read from the input queue
 - a thread producing output would block if output queue is full
 - a thread consuming input would block if input queue is empty
 - what about the other ends of these queues? who are handling them?




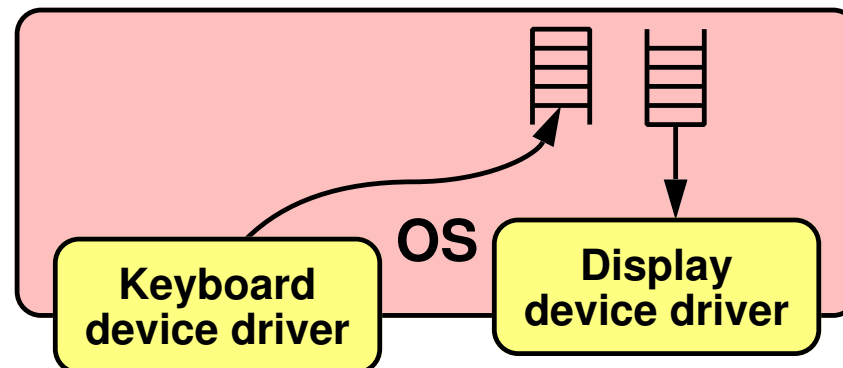
Terminals

- ➡ To deal with concerns (1) and (2)
- ➡ how about using a keyboard reading thread (that would do the following)?
 - 1) issue a read to the device
 - 2) block itself and wait for interrupt from the device
 - 3) when interrupt occurs, the thread is woken up
 - 4) the thread reads from the device and move one character from the device to the input queue
 - 5) goto step 1
 - ➡ this approach of using *thread context* seems to be an overkill and may be inefficient





Terminals

- ➡ To deal with concerns (1) and (2)
- ▢ how about just using an *interrupt handler*?
 - in the *read-completion interrupt*, the handler moves one character from the device to the input queue and issue another read request to the device and blocks
 - ◆ if the queue is full, the character is thrown away
 - ◆ the application thread must *mask interrupts* when it's taking a character from the queue
 - ▢ can do the same for the output queue ...



Terminals

- ➡ To deal with concerns (1) and (2)
- ▬ how about just using an *interrupt handler*?
 - ▬ can do the same for the output queue
 - in the *write-completion interrupt*, the handler moves one character from the output queue to the device and issue another write request to the device and blocks
 - ◇ if the application writes to an *empty queue*, it would *setup the write-completion interrupt handler* and issue a write request to the device




Terminal




Additional issue

- 3) input characters may need to be processed in some way before they reach the application
 - ◆ e.g., characters typed at the keyboard are *echoed* back to the display
 - ◆ characters may be grouped into lines of text and subject to simple *editing* (such as backspace)
 - ◆ some applications prefer to process all characters themselves, including their editing




Terminals

- ➡ To deal with concern (3)
- remember, once you allow an application to take a character, you cannot ask for it back
 - can only give it to the application when there is no chance that you will want it back
 - ◆ this happens when a line is completed
 - therefore, we need *two input queues*
 - one for the *partial-line*
 - ◆ subject to editing
 - the other contain characters from *completed lines*
 - in the read-completion interrupt, the handler moves one character from the device to the *partial-line queue*
 - ◆ if the input character is a carriage-return, the entire content of the partial-line queue is moved to the completed-line queue




Terminals

- ➡ To deal with concern (3)
- when a character is typed, it should go to the display immediately (due to the *echoing* requirement)
 - it may be competing with the output thread, but it's okay (and that's how it's done in Unix)
 - Windows handle this differently
 - ◆ typed characters are only echoed when an application consumes them
 - ◆ therefore, echoing is not done in the interrupt context
 - ◆ echoing is done in the context of the thread consuming the characters (i.e., the one that calls `read()`)



Modularization

- ➡ ***Device independence*** consideration (figure out the ***common*** part)
 - for many different serial-line devices, character processing is common
 - actually, character processing is performed in situations where the source and sink of characters aren't even a serial line
 - ◆ e.g., bit-mapped display, network connection
 - therefore, it makes sense to separate the device dependent part from the ***common, device independent*** part
 - promotes ***reusability***
- ➡ A separate module, known as the ***line-discipline*** module in some systems, provides the ***common character-handling code***
 - it can interact with any device driver capable of handling terminals
 - can even use a different line-discipline module to deal with an alternative character set


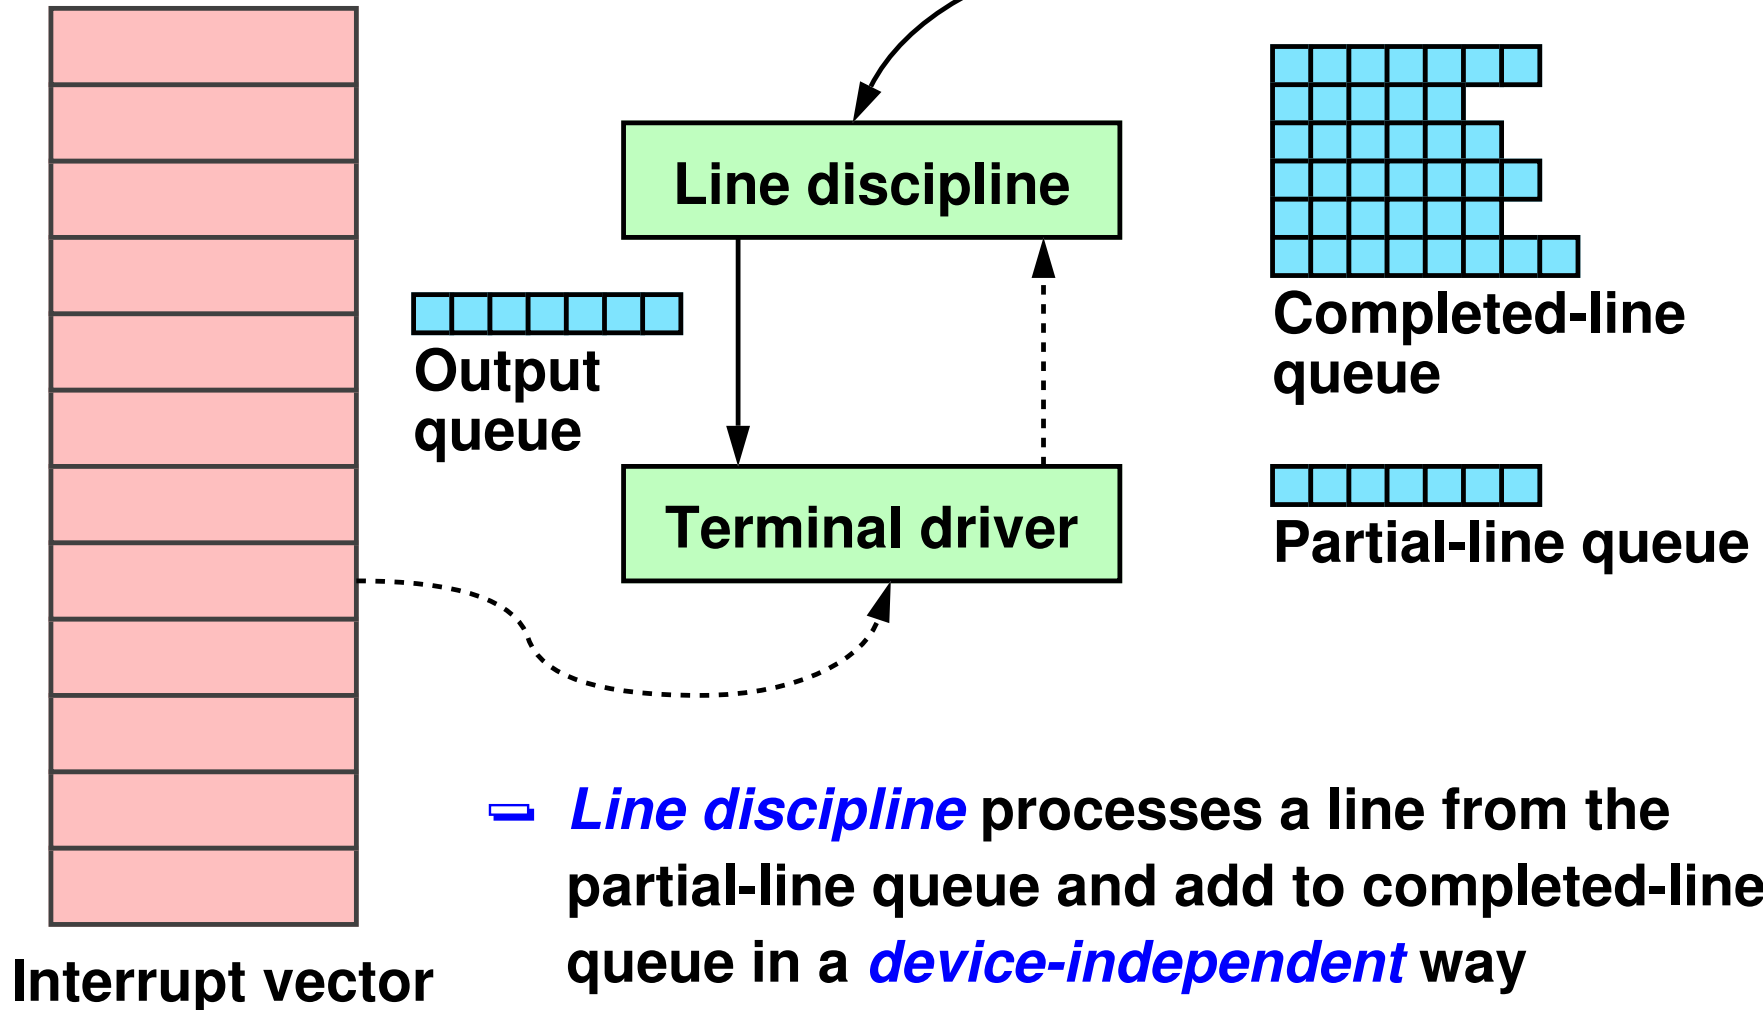


Terminals

Application

Applications

OS

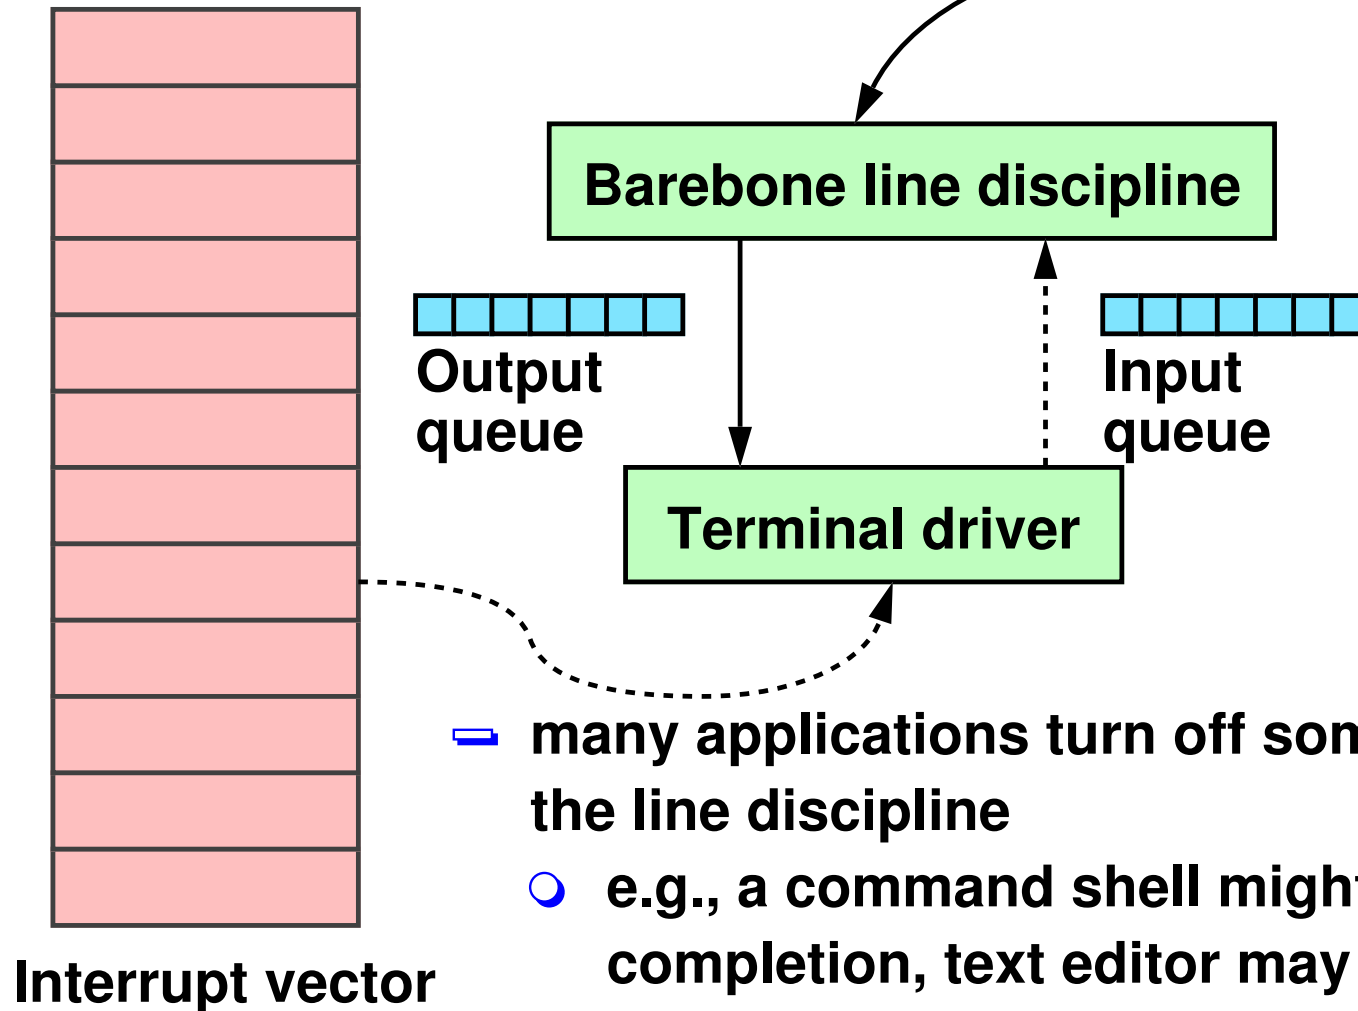


Terminals

Application

Applications

OS



- many applications turn off some processing of the line discipline
 - e.g., a command shell might do command completion, text editor may correct spelling

Interrupt vector

Where to Put the Modules?



Where to put the terminal driver and the line-discipline module?

1) kernel

2) separate user process

3) library routines that are linked into application processes

⇒ **driver** should be in the kernel since device registers access needs to be protected from arbitrary manipulation by application programs

⇒ **line-discipline** may be shared by multiple applications

○ putting it in **library routines** will make it **difficult to share** one terminal with many user applications


○ can it go into a separate **user process**?

◆ but can have serious **performance problems**

◆ would need to transfer characters into the line-discipline process, then transfer to another process

○ putting it in the **kernel** seems to be the best choice

◆ although kernel code is **hard to modify**, replace, and debug




Terminals and Pseudo Terminals

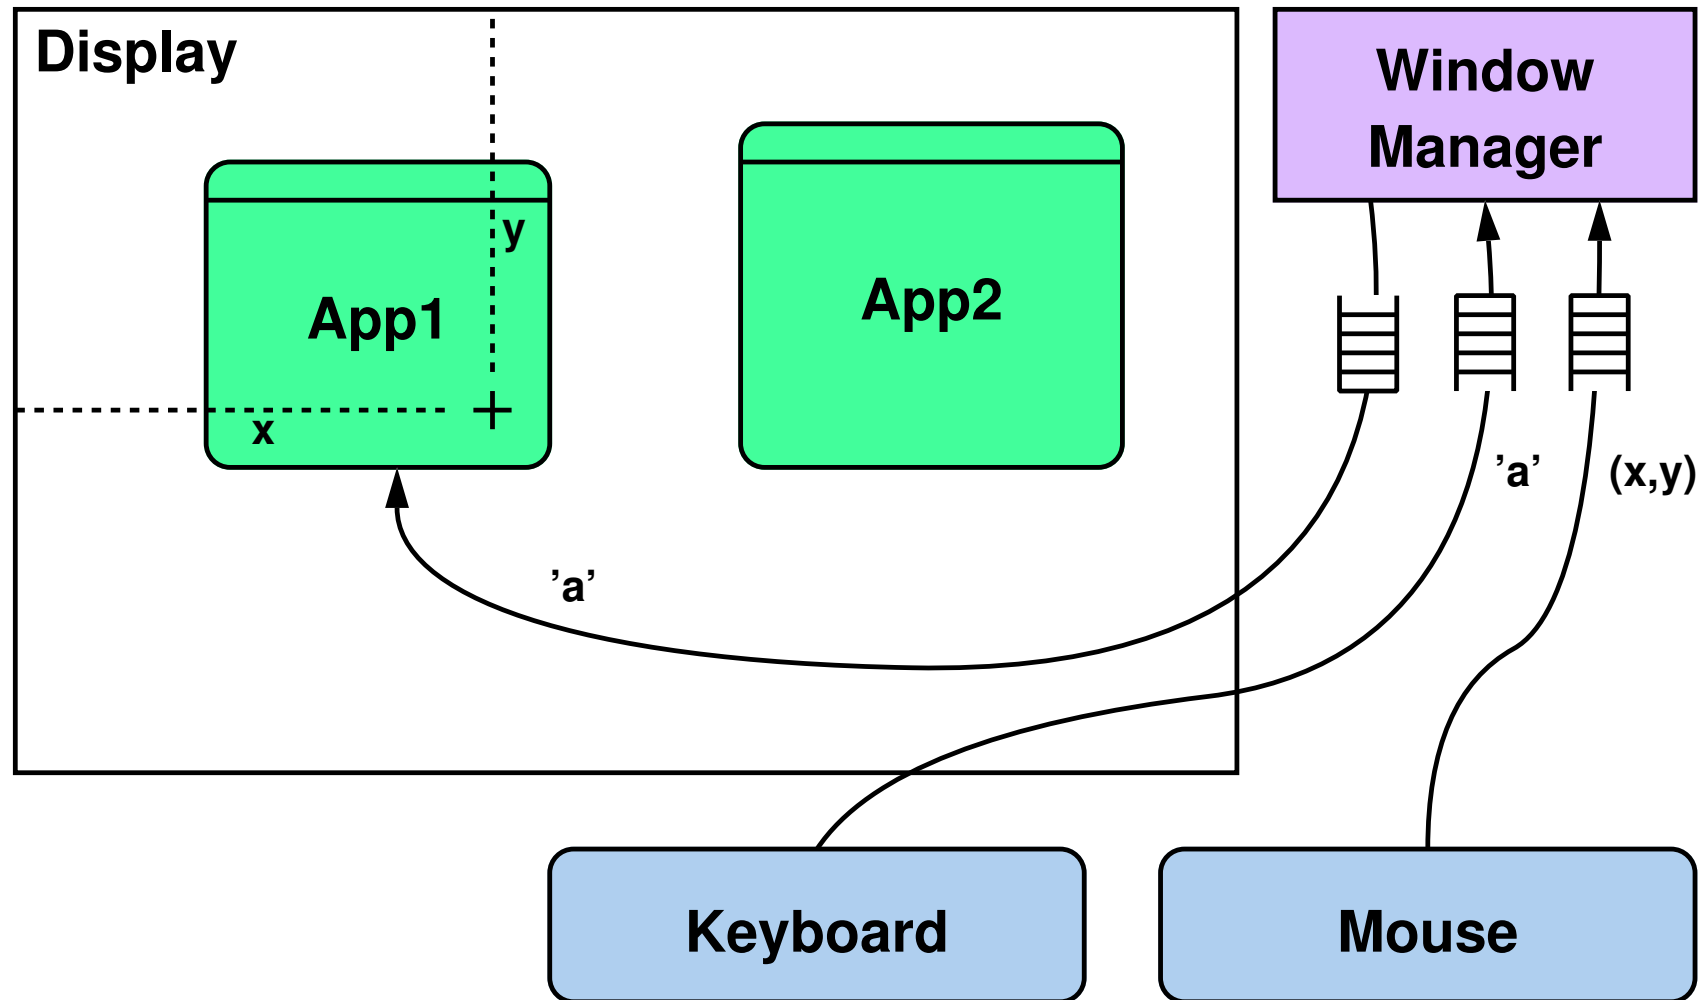


Modern systems do not have terminals

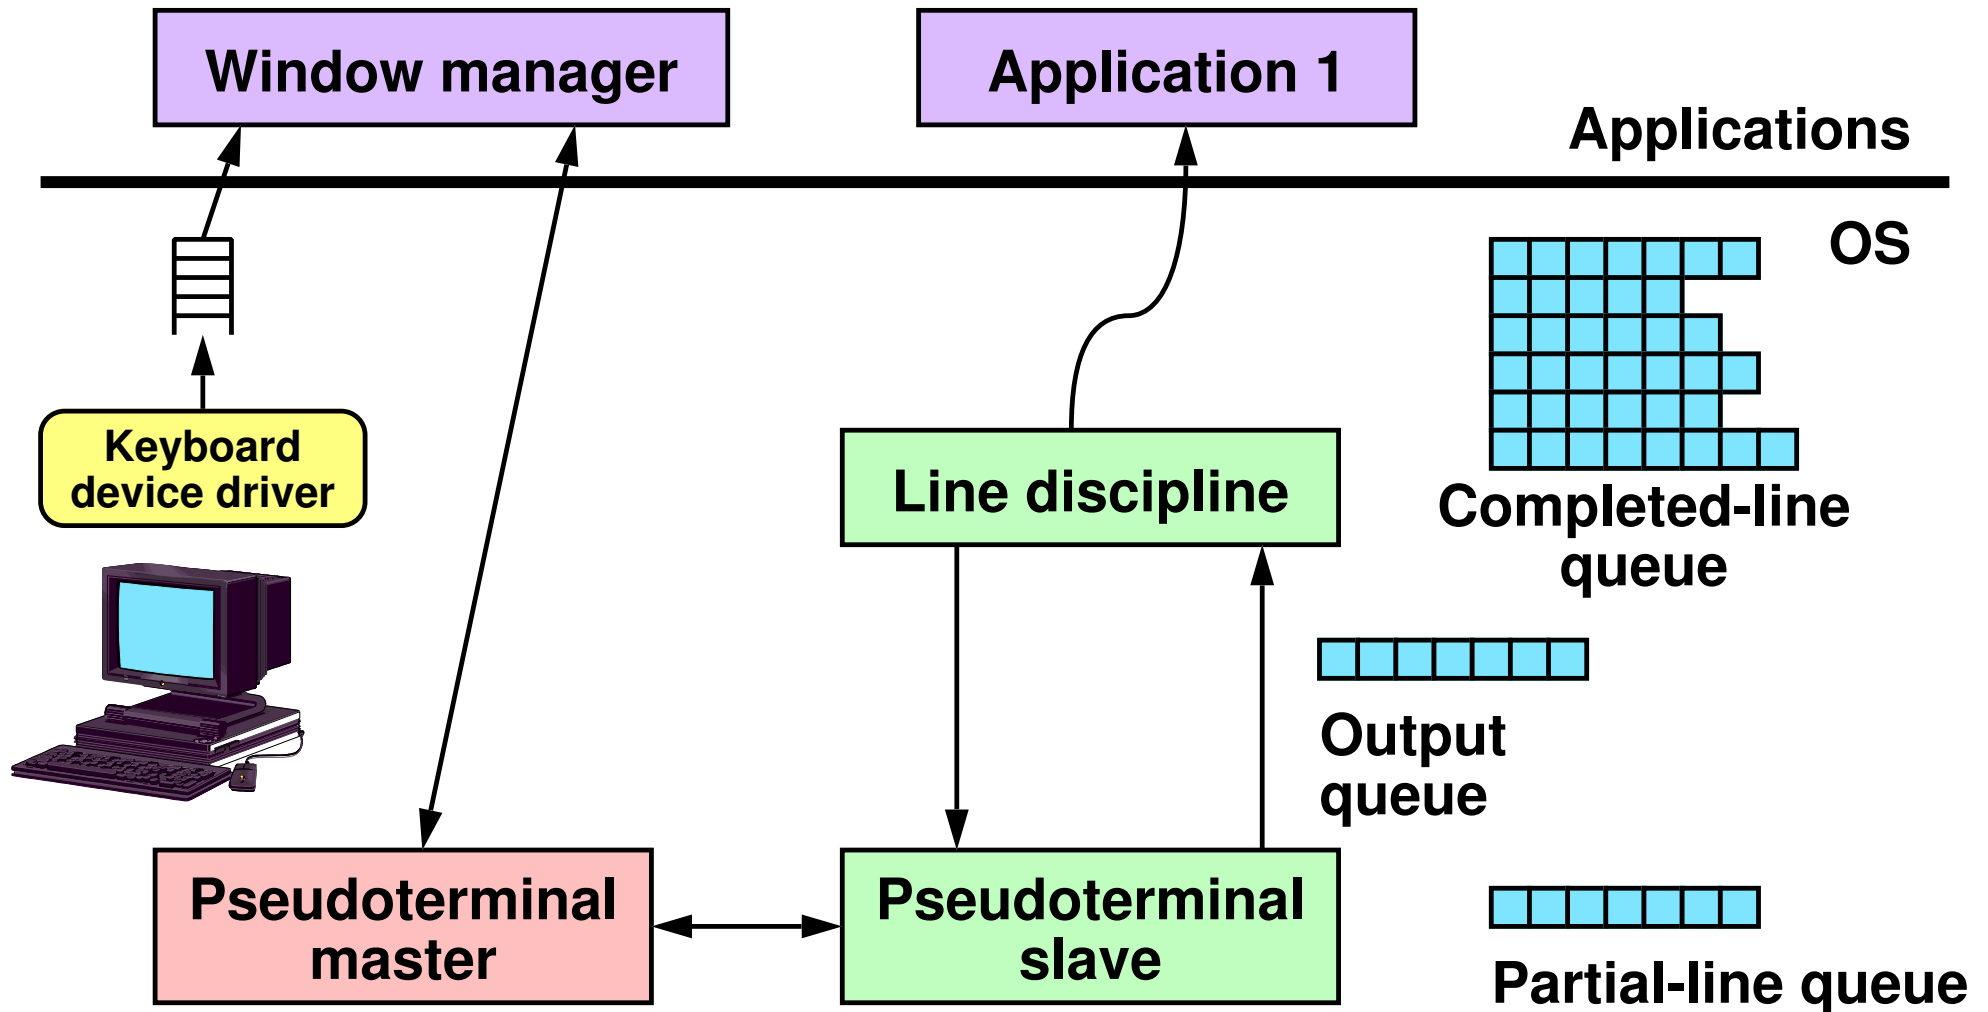
- they often have bit-mapped displays, keyboards and mice connected via USB
- a *window manager* implements windows on the display and determines which applications receive typed input (*input focus*)
- a server might support remote sessions where applications receive input and send output over a network
- they use *pseudoterminals*
 - which implements a line discipline whose input comes from and output goes to a *controlling application* (and not a physical device)




Bitmapped Display



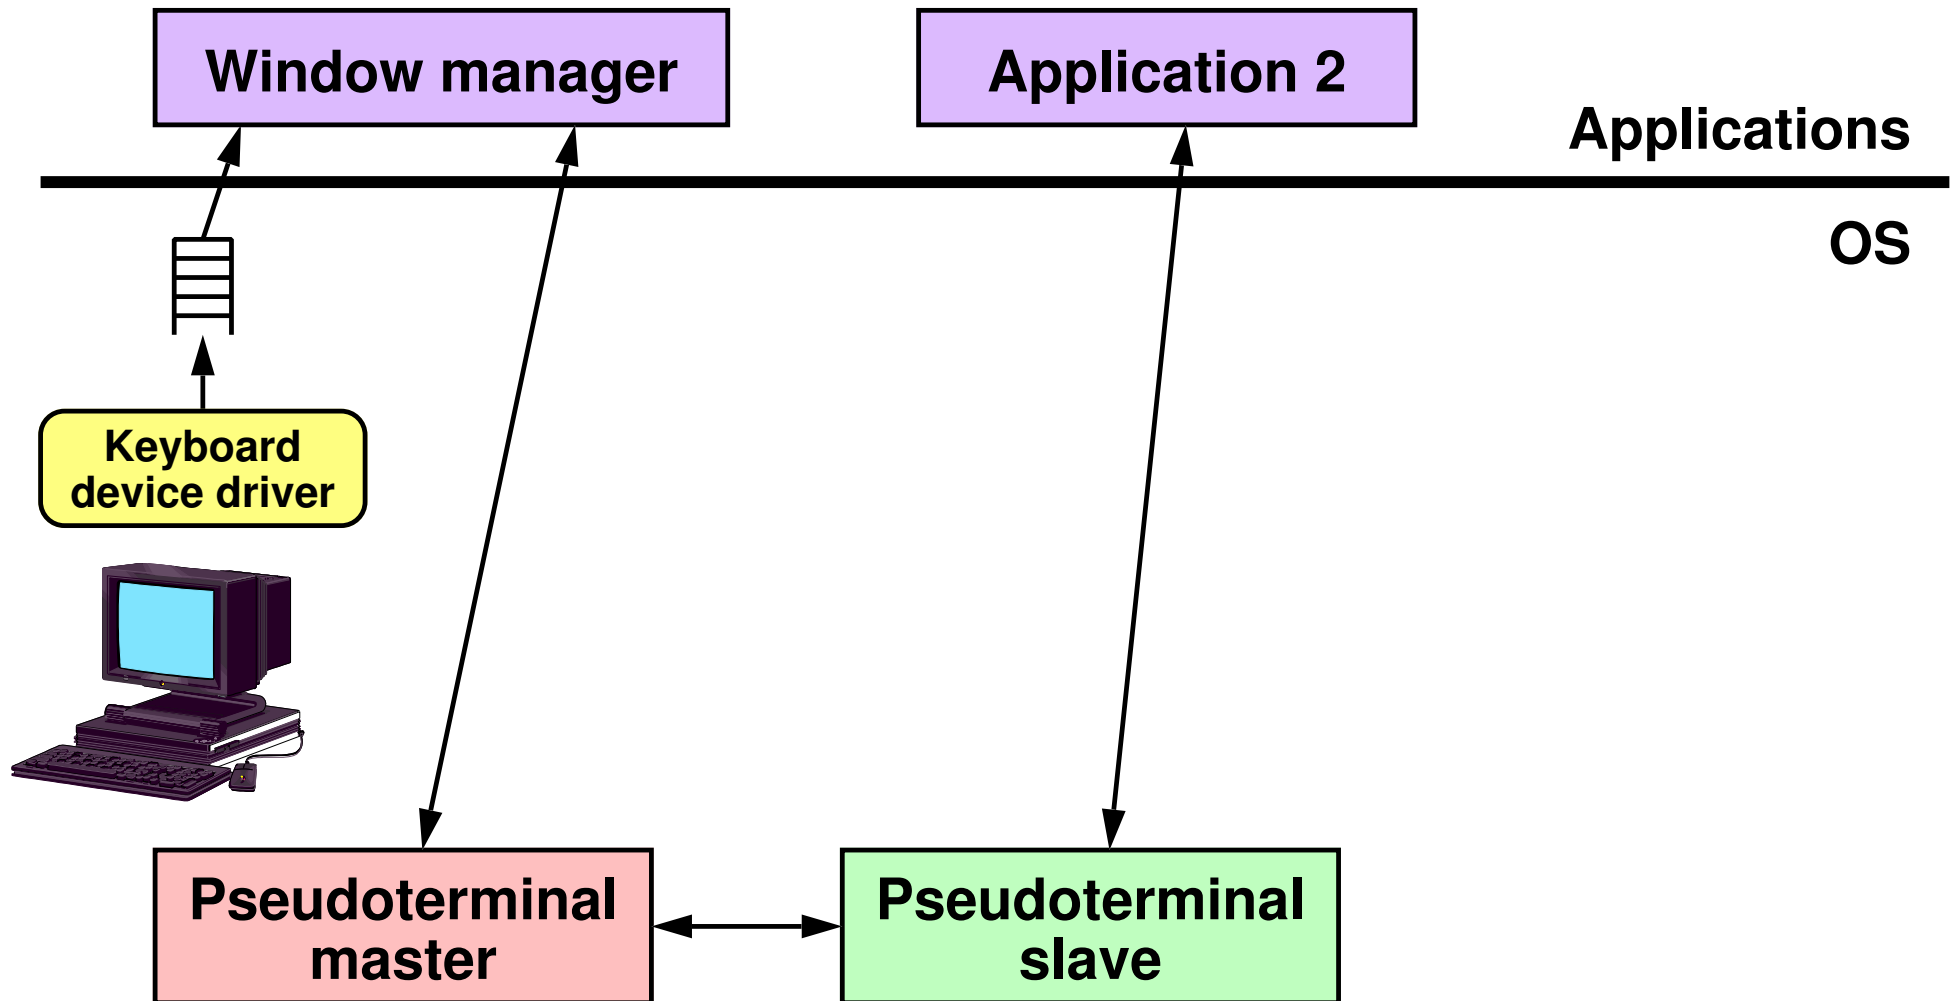
Pseudo Terminals




- the OS provides a pair of entities (*pseudoterminal master* and *pseudoterminal slave*) that appear to applications as *devices*



Pseudo Terminals




- pseudoterminal slave acts like a *terminal device driver* to the rest of the OS
- can by-pass the line discipline module if desired




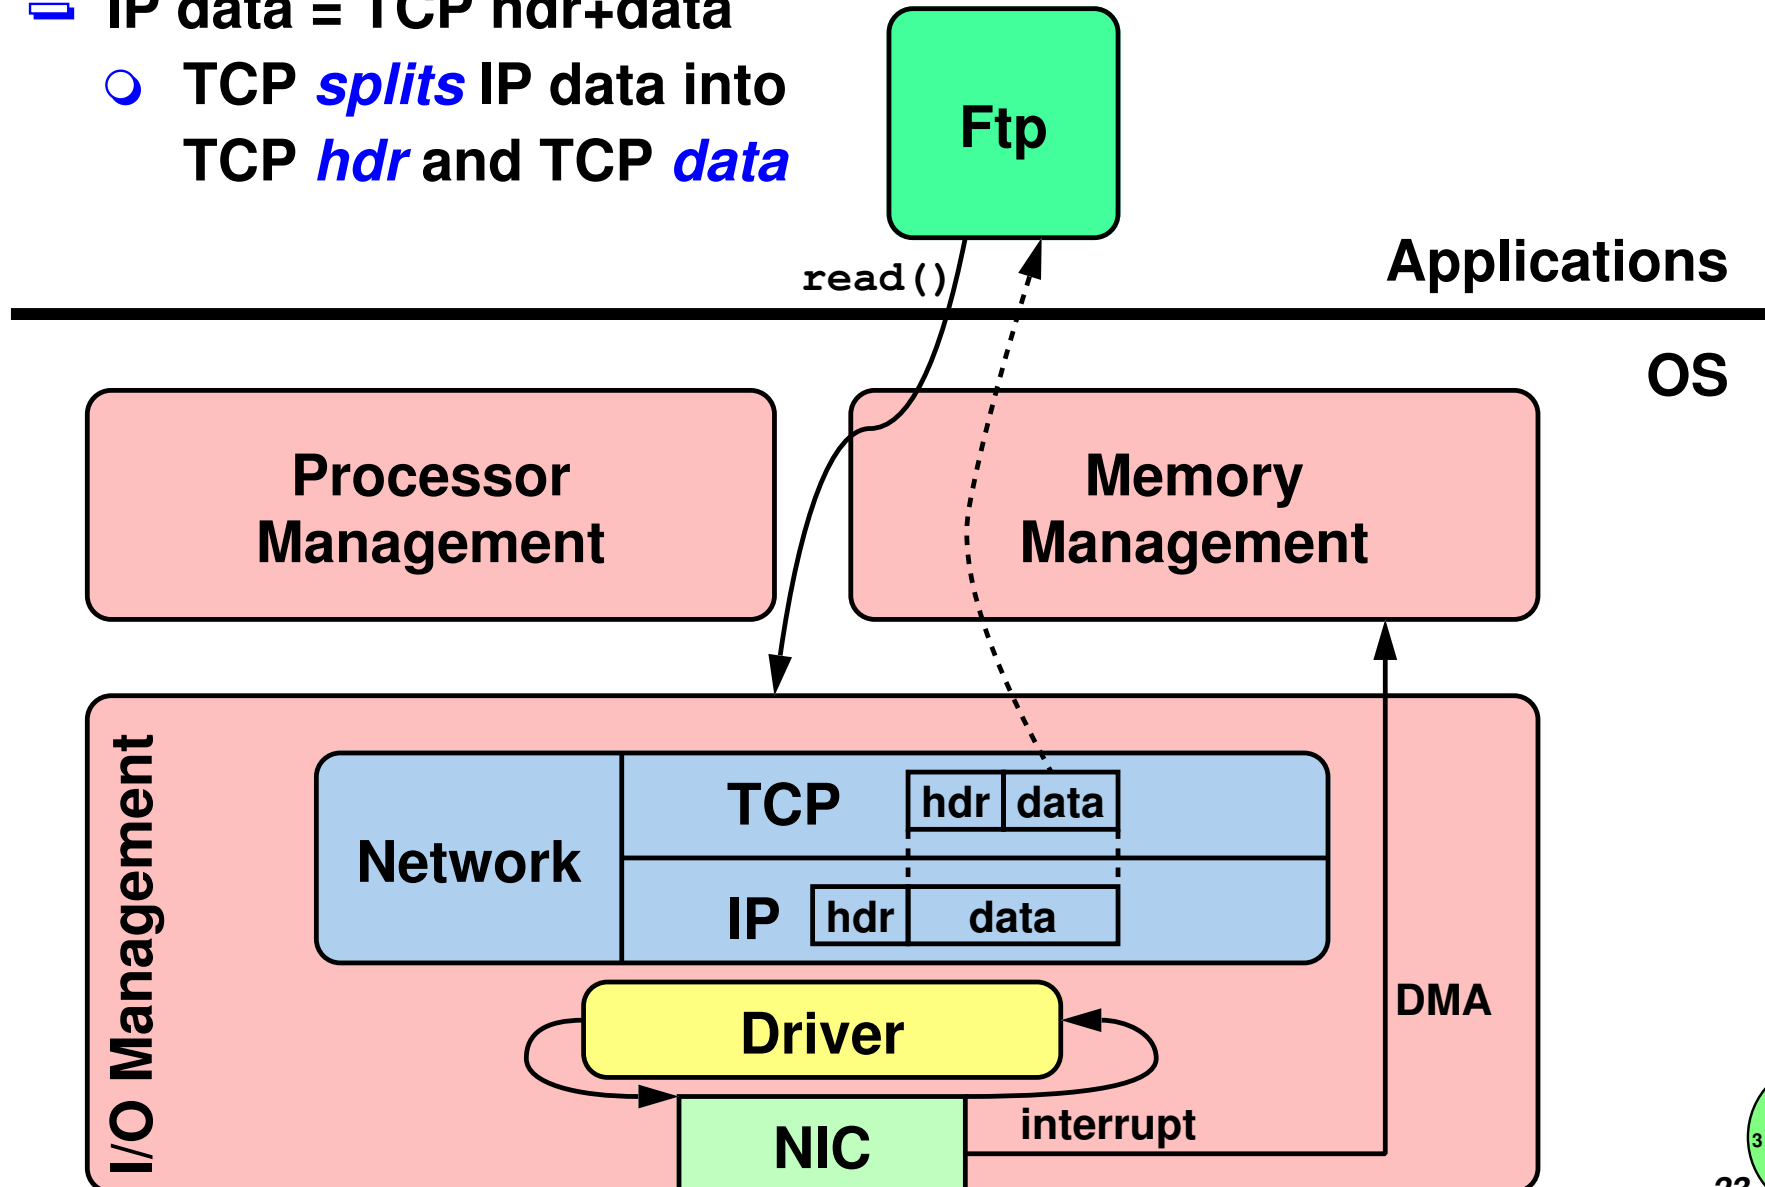
Network Communication

- ➡ Network communication and terminal handling are very similar, in architecture and implementation
- device is called a *Network Interface Card (NIC)*
 - data arrives in a packet (instead of a character)
 - incoming data must be processed via *network-protocol modules* similar to line-discipline module
- ➡ Main difference
- performance is crucial in network communication
 - data from keyboard can go at tens of characters per second
 - data going to display can go at a few thousand characters per second (for character-based display)
 - protocols are *layered* on top of one another
 - *data* in lower layer is views as *header + data* in higher layer
 - cannot afford to copy network data from queue to queue!
 - no copying allowed!
 - must pass by memory address!




Network Communication

- ➡ Ex: TCP (details in Ch 9)
- ▢ IP data = TCP hdr+data
 - TCP *splits* IP data into TCP *hdr* and TCP *data*





Network Communication


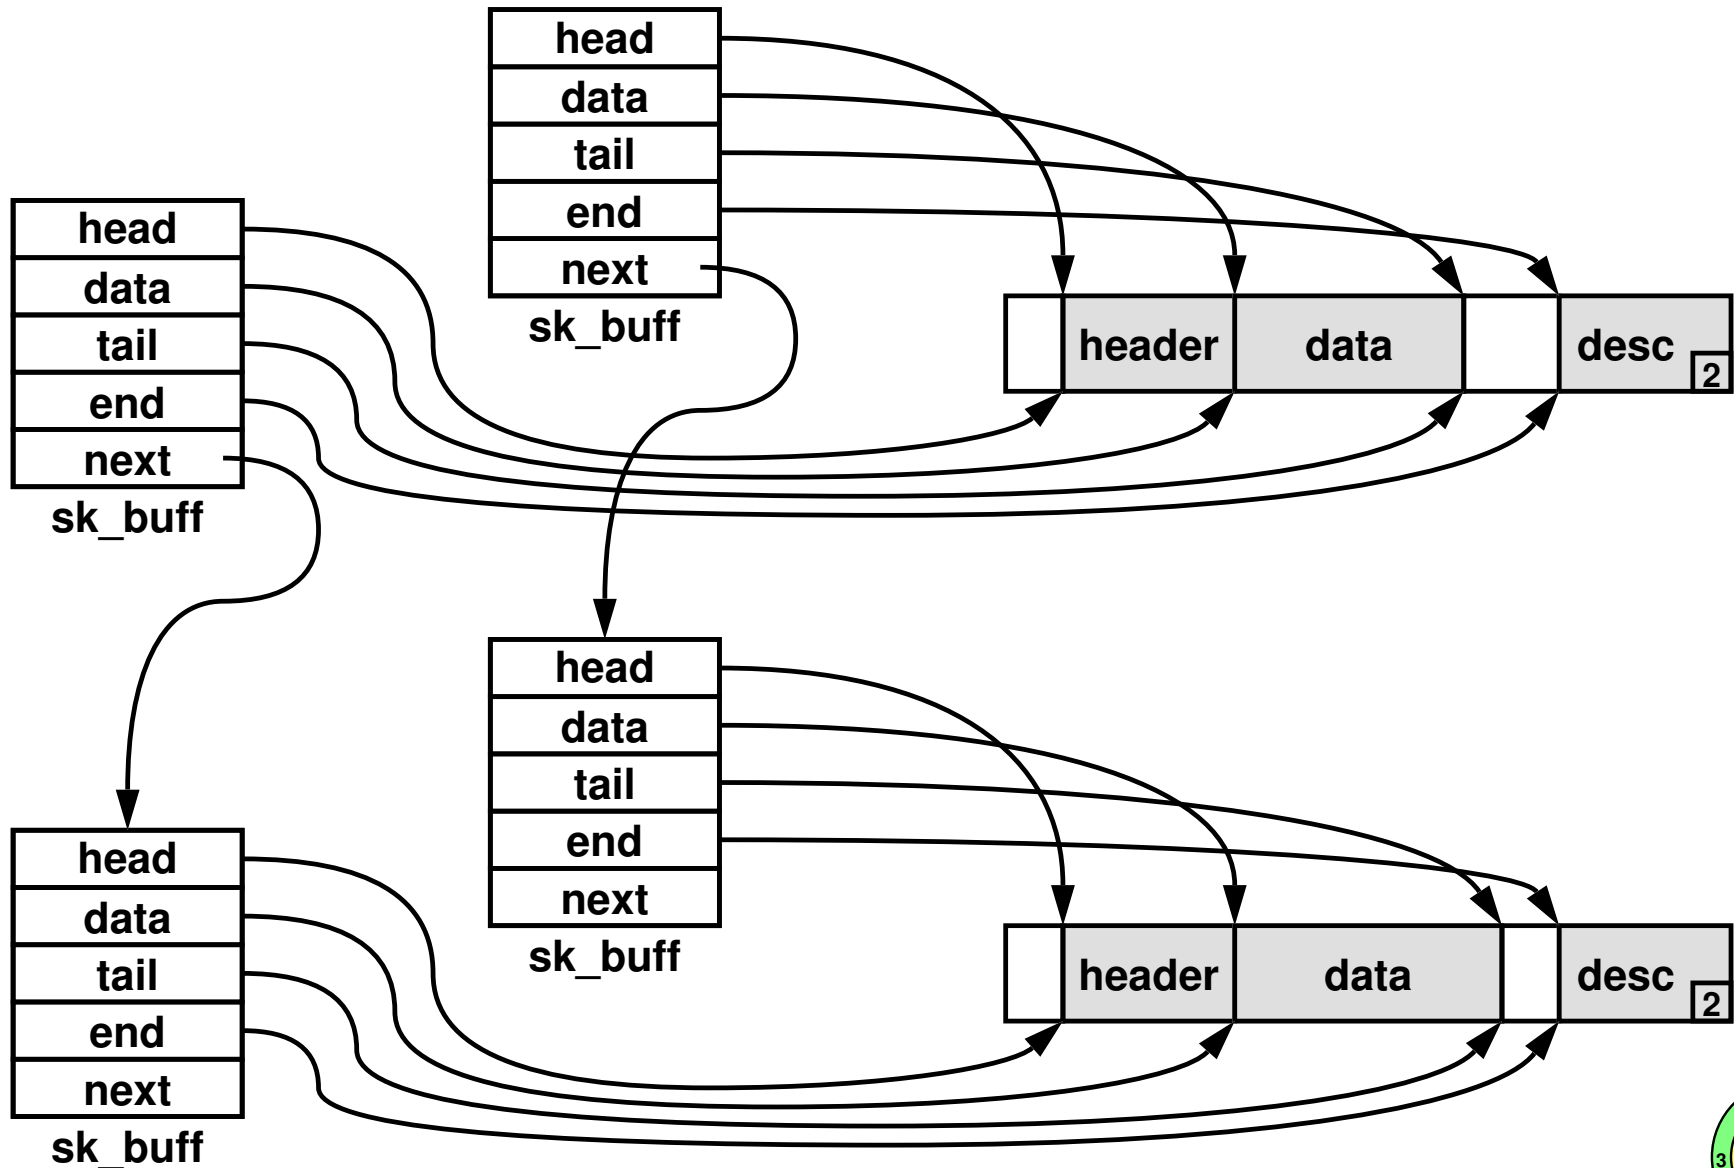
- ➡ Ex: TCP (details in Ch 9)
- ➡ Performance challenge:
 - need to be able to pass blocks from one *module* to the next *without copying* data
 - copying came from
 - ◆ splitting "header" and "data"
 - ◆ copying data into application-provided buffer
 - append headers to the beginning of outgoing packets; remove headers from incoming packets (*known as layering*)
 - *hold on to packets* for possible retransmission
 - request and *respond to time-out* notifications
- ➡ To accomplish all this in our simple OS, we use a data structure adapted from Linux
 - `sk_buff` (socket buffer)



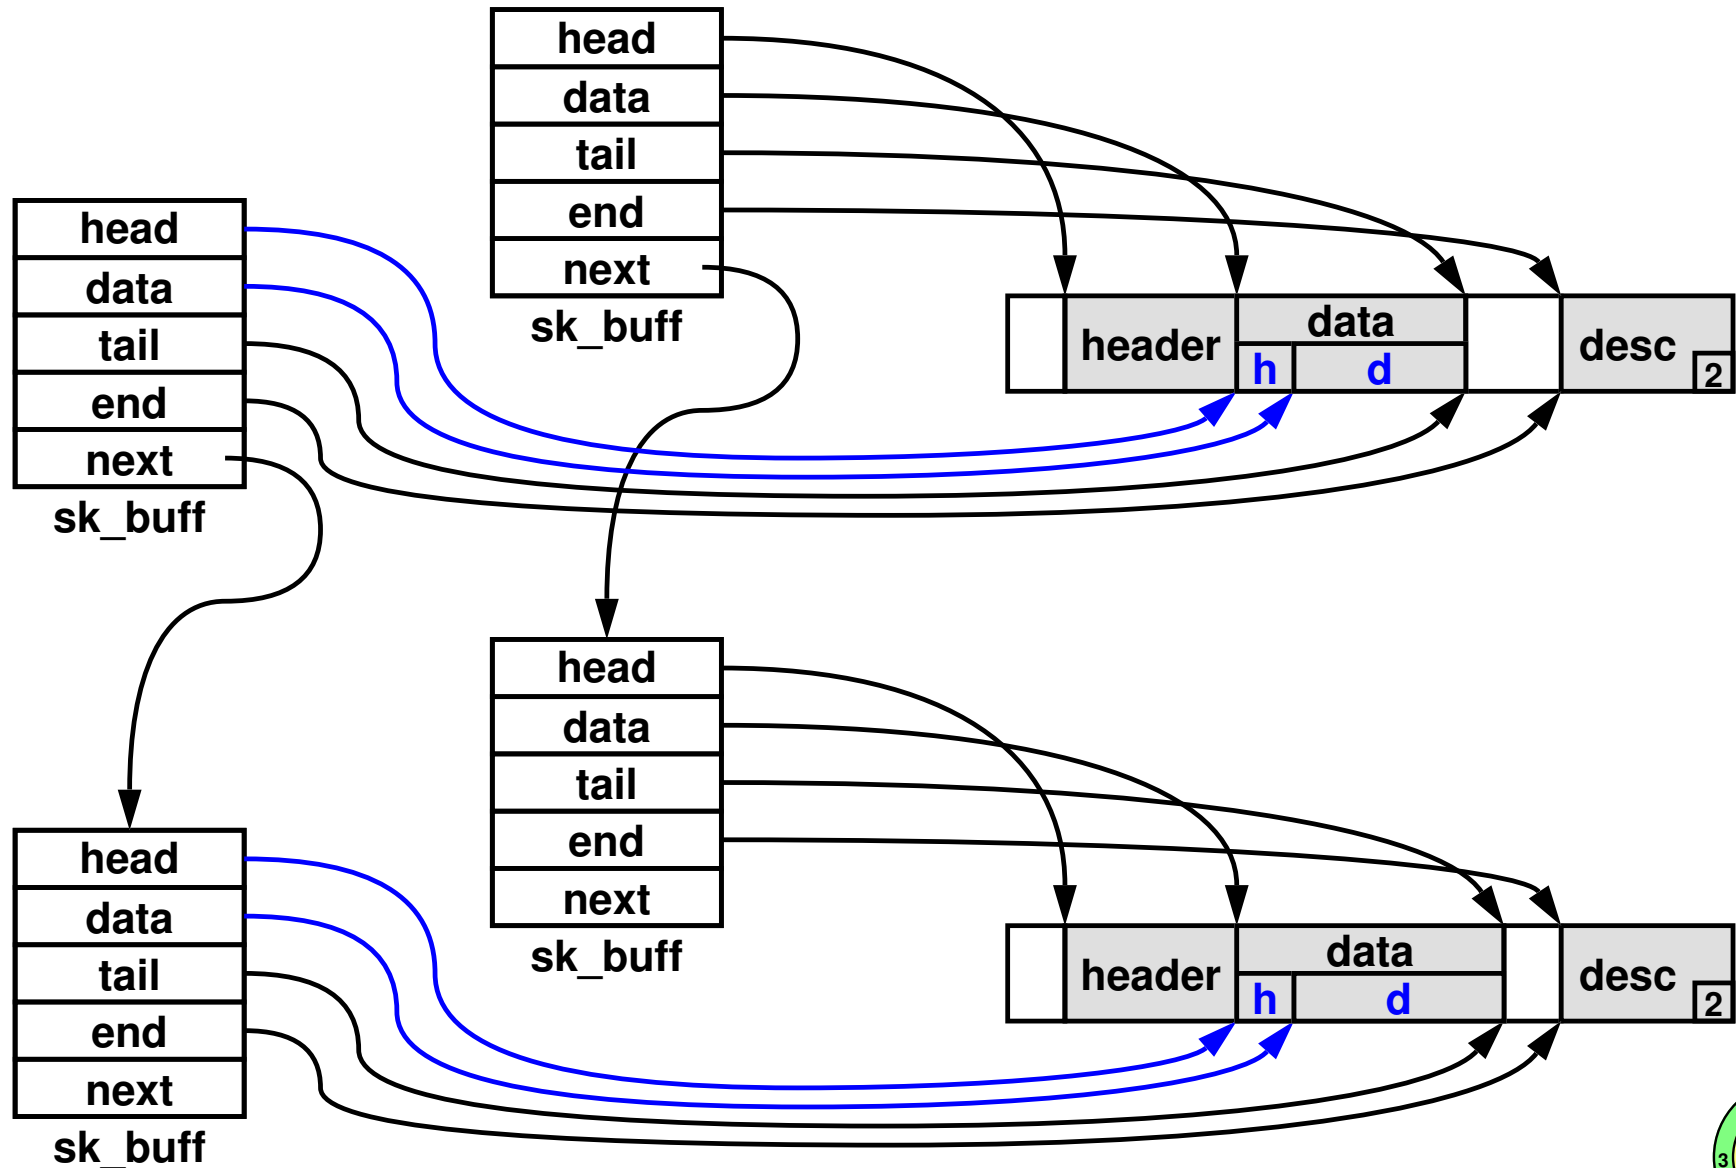
Two Queued Segments



Passed to the Another Module...



Copyright © William C. Cheng



Support Timeout

- ➡ Lots of timers in network programming!
 - if you send a message/packet that needs a response or an acknowledgement (such as in TCP internal), and
 - if it's possible for the message/packet to be lost
 - you need to set a "reasonable" timeout
- ➡ To implement *timeout*, can use a *callback* mechanism
 - use a function pointer and pass it to the interval timer
 - when timeout occurs, call the callback function
 - if the acknowledgement was received before timeout occurred
 - need to cancel the timer
 - can also specify a cancel routine

