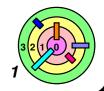
Ch 5: Processor Management

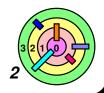
Bill Cheng

http://merlot.usc.edu/cs402-s16



Processor Management

- Threads *Implementation*
 - lock/mutex implementation on multiprocessors
- Interrupts
- Scheduling
- Linux/Windows Scheduler

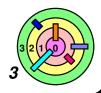


5.1 Threads Implementations



A Simple Thread Implementation

Multiple Processors



Threads Implementation



The ultimate goal of the OS is to support user-level applications

we will discuss various strategies for supporting threads



Where are operations on threads implemented?

- in the kernel?
- or in user-level library?



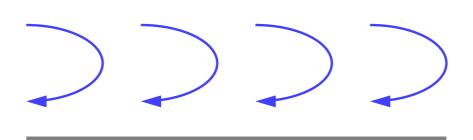
Approaches

- one-level model (threads are implemented in the kernel)
 - variable-weight processes
- two-level model (threads are implemented in user library)
 - \circ N×1
 - \circ M \times N
 - scheduler activations model



One-Level Model

User

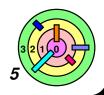


Kernel





Processors



One-Level Model



The simplest and most direct approach is the one-level model

- all aspects of the thread implementation are in the kernel
 - i.e., all thread routines (e.g., pthread_mutex_lock) called by user code are all system calls
- each user thread is mapped one-to-one to a kernel thread



If a thread calls pthread_create()

- it's a system call, so it traps into the kernel
- the kernel creates a thread control block
 - associate it with the process control block
- the kernel creates a kernel and a user stack for this thread



What about pthread_mutex_lock()

- why does it have to be done in the kernel?
- it's not necessary to protect the threads from each other!
 - you definitely don't need the kernel to protect threads from each other



One-Level Model



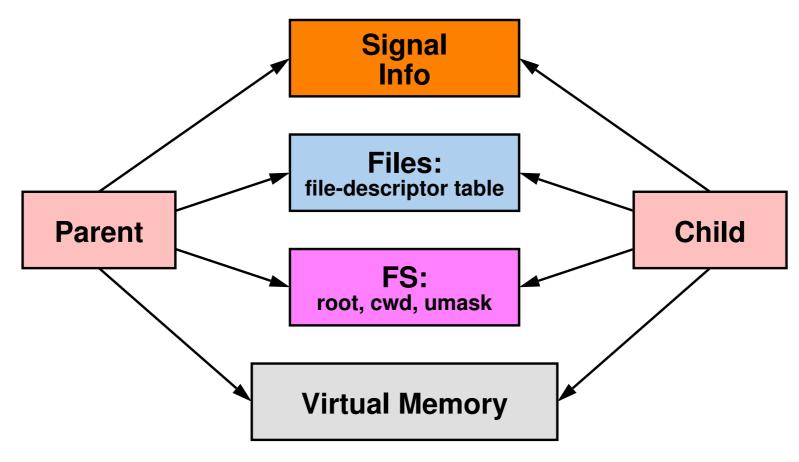
Problem: system calls are expensive

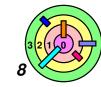
- if pthread_mutex_lock finds the mutex available, it should return quickly (and lock the mutex)
 - if this can be done in user code, it can be 20 times faster (for the case where the mutex is available)
 - in Win32 threads, an equivalent of a mutex is represented in a user-level data structure
 - if such an object is not locked, it returns quickly
 - if such an object is locked, it makes a system call and blocks in the kernel



Variable-Weight Processes

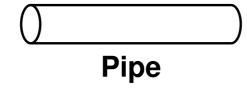
- Variant of one-level model
- Portions of parent process selectively copied into or shared with child process
- Children created using clone() system call



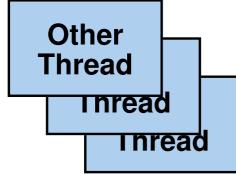


Linux Threads (pre 2.6)

Initial Thread



Manager Thread



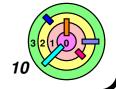


NPTL in Linux 2.6



Native POSIX-Threads Library

- full POSIX-threads semantics on improved variable-weight processes
- threads of a "process" form a thread group
 - getpid() returns process ID of first thread in group
 - any thread in group can wait for any other to terminate
 - signals to process delivered by kernel to any thread in group
- new kernel-supported synchronization construct: futex (fast user-space mutex)
 - used to implement mutexes, semaphores, and condition variables



Two-Level Model



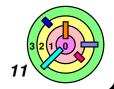
In the two-level model, a user-level library plays a major role

what an user-level application perceives as a thread is implemented within user-level library code

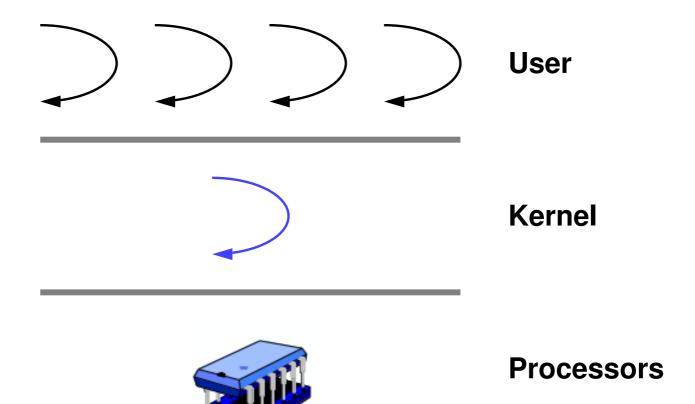


Two versions

- single kernel thread (per user process)
- multiple kernel threads (per user process)



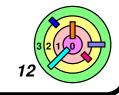
Two-Level Model - One Kernel Thread





This is one of the earliest ways of implementing threads

- threads are implemented entirely in the user level
 - thread control block, mutex in user space
 - thread stack allocated by user library code
- mostly done on uniprocessors



Two-Level Model - One Kernel Thread



- Within a process, user threads are multiplexed not on the processor, but on a kernel-supported thread
- the OS multiplexes kernel threads (or equivalently, processes)
 on the processor
- kernel does not know about the existance of user threads



User thread creation

- a stack and a thread control block is allocated
- thread is put on a queue of runnable threads
 - wait for its turn to become the running thread



Synchronization implementation

- relative straightforward
- e.g., mutex (one queue per mutex)
 - if a thread must block, it simply queues itself on a wait queue and calls context-switch routine to pass control to the first thread on the runnable queue

Two-Level Model - One Kernel Thread

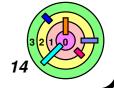




fast, because no system calls for thread-related APIs

Major disadvantage

- what if a thread makes a system call (for a non-thread-related API)?
 - it gets blocked in the kernel
 - no other user thread in the process can run



Coping ...

```
ssize_t read(int fd, void *buf, size_t count)
{
    ssize_t ret;
    while (1) {
        if ((ret = real_read(fd, buf, count)) == -1) {
            if (errno == EWOULDBLOCK) {
                 sem_wait(&FileSemaphore[fd]);
                  continue;
            }
            break;
    }
    return(ret);
}
```



Solution is to have a non-blocking read() called real_read()

- real_read() either returns immediately with data in buf
- or returns immediately with an error code in errno
 - EWOULDBLOCK means that a real read() would block, i.e., data is not ready to be read

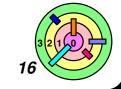
Coping ...



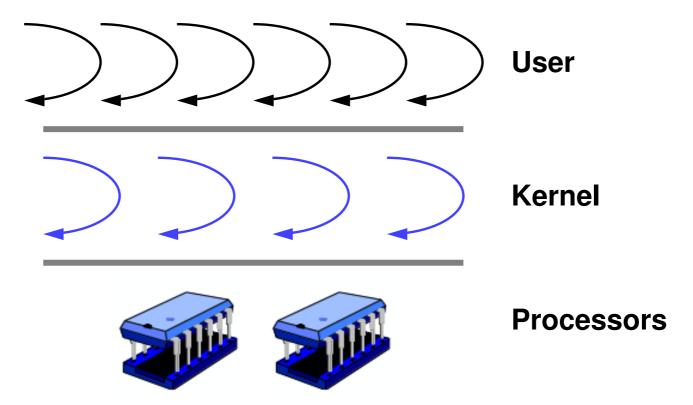
perhaps a signal handler will invoke sem_post() to when data is ready to be read



only works for some I/O objects - not a general solution



Two-Level Model: Multiple Kernel Threads

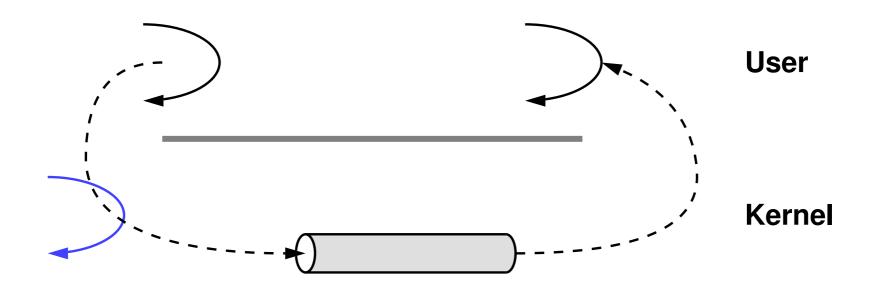




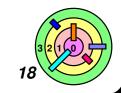
- Implementation is similar to the two-level model with a single kernel thread
 - no system calls (for thread-related APIs)
 - if we don't have enough kernel threads per user process, we end up having the same problem with the N-to-1 model



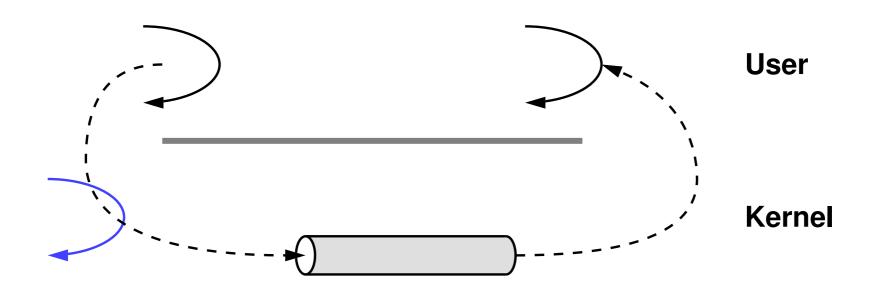
Deadlock



- Ex: two threads are communicating using a pipe (this is essentially a kernel implementation of the producer-consmer problem)
 - first user thread writes to a full pipe and get blocked in the kernel
 - first thread just happened to use the last kernel thread
 - 2nd thread wants to read the pipe to unblock the first thread, but cannot because no kernel thread left



Deadlock





Solaris solution: automatically create a new kernel thread

an obvious solution



Recap - Problems



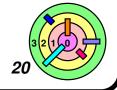
Two-level model does not solve the I/O blocking problem

- if there are N kernel threads and if N user threads are blocked in I/O
 - no other user threads can make progress



Another problem: Priority Inversion

- user-level thread schedulers are not aware of the kernel-level thread scheduler
 - it may know the number of kernel threads
- how can the user-level scheduler talk to the kernel-level scheduler?
 - people have tried this, but it's complicated
- it's possible to have a higher priority user thread scheduled on a lower priority kernel thread and vice versa

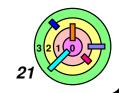


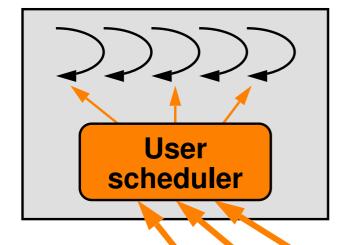
Scheduler Activations Model

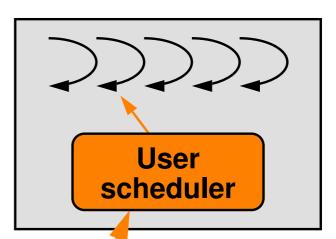


The scheduler activations model is radically different from the other models

- in other models, we think of the kernel as providing some kernel thread contexts
 - then multiplexing these contexts on processors using the kernel's scheduler
- in scheduler activations model, we divvy up processors to processes, and processes determine which threads get to use these processors
 - the kernel should supply however many kernel contexts it finds necessary



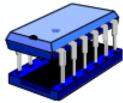


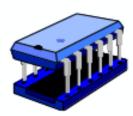


User Kernel

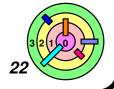
Kernel scheduler











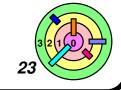


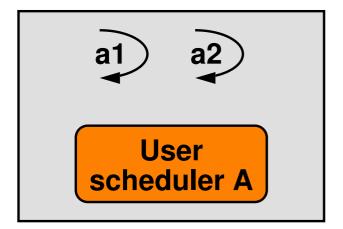
Let's say a process starts up running a single thread

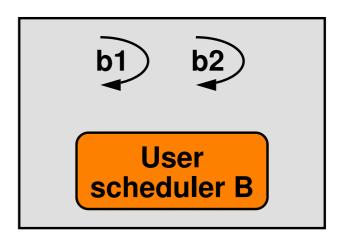
- kernel scheduler assigns a processor to the process
- if the thread blocks, the process gives up the processor to the kernel scheduler



- Suppose the user program creates a new thread and parallelism is desired
- code in user-level library notifies the kernel that it needs two processors
- when a processor becomes available, the kernel creates a new kernel context
 - the kernel places an upcall to the user-level library, effectively giving it the processor
 - the user-level library code assigns this processor to the new thread



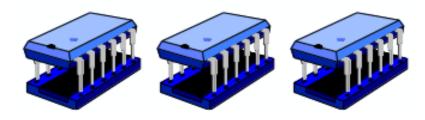




User

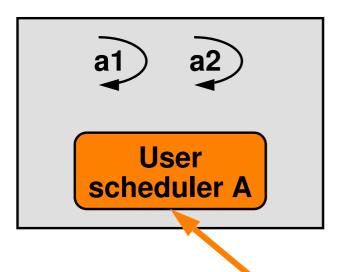
Kernel

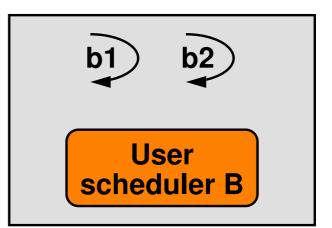
Kernel scheduler







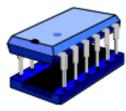


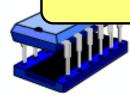


kernel scheduler does an upcall to offer processor 1 to user scheduler A

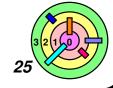
Kernel scheduler

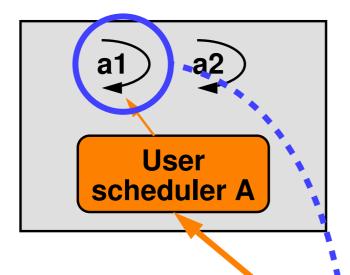


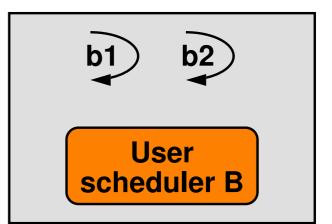




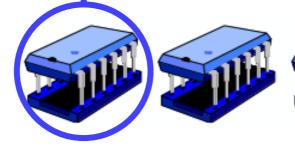








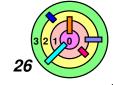
Kernel scheduler

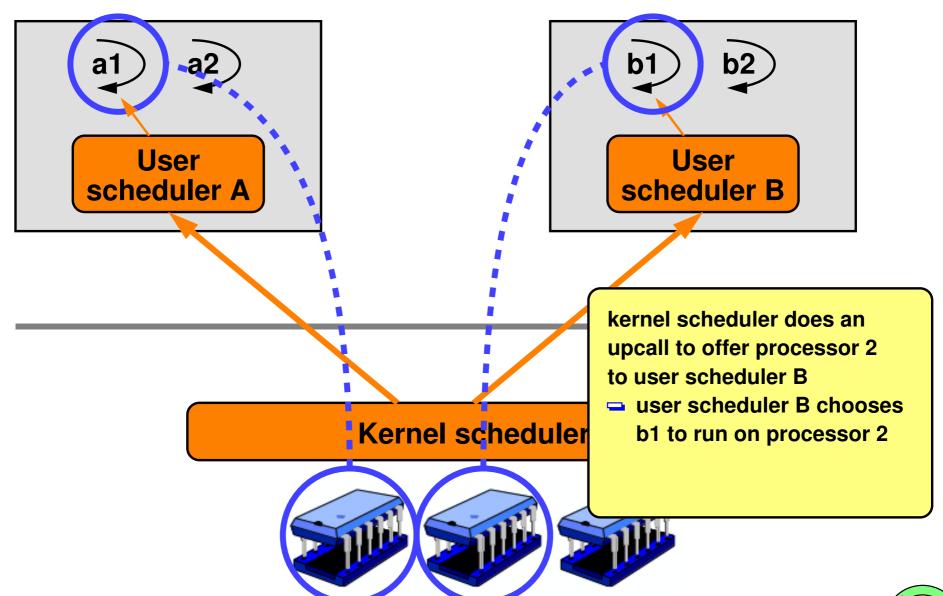


kernel scheduler does an upcall to offer processor 1 to user scheduler A

- user scheduler A choosesa1 to run on processor 1
- kernel does not choose threads, just processes

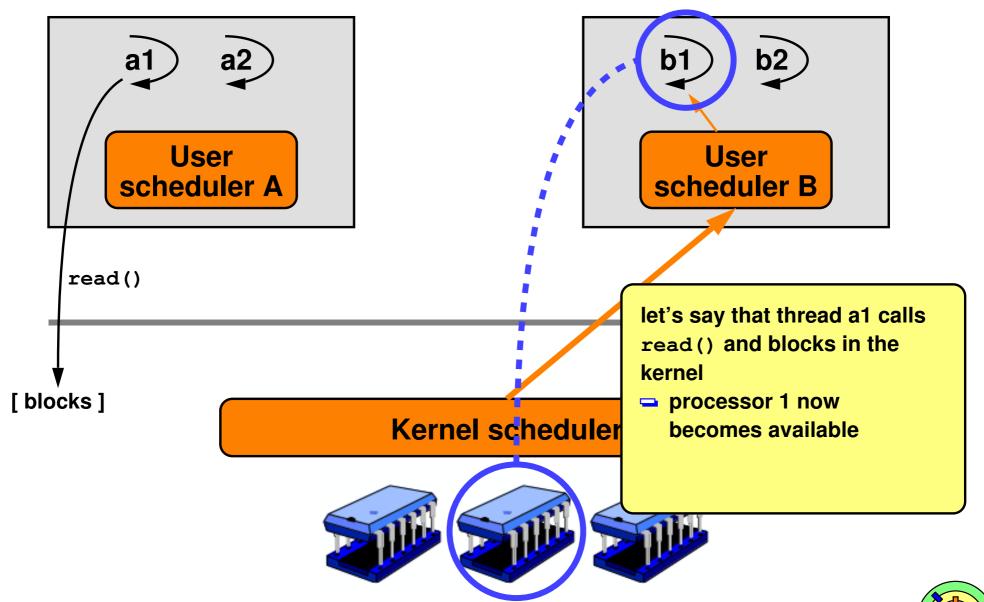








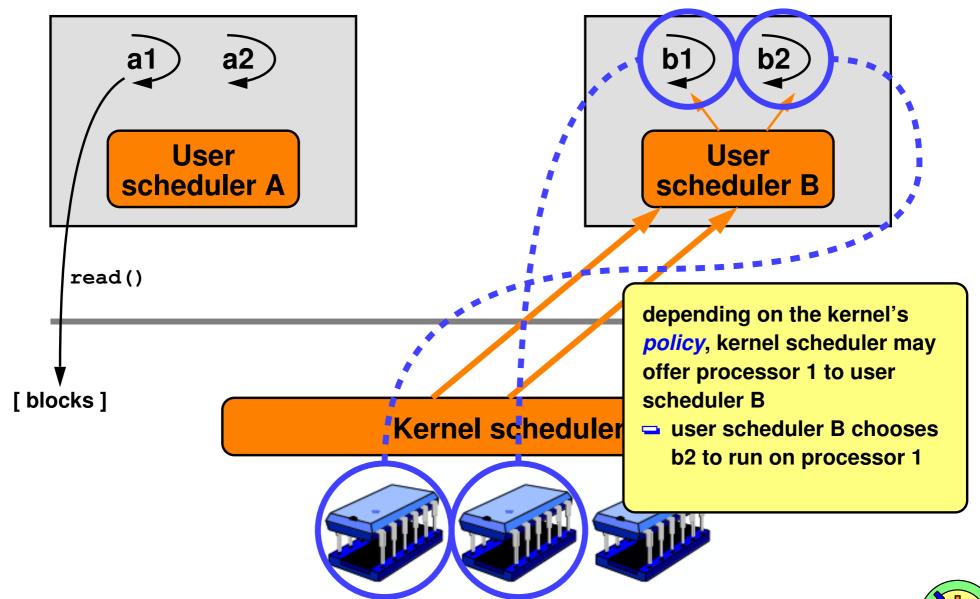






Kernel scheduler can have various scheduling policies

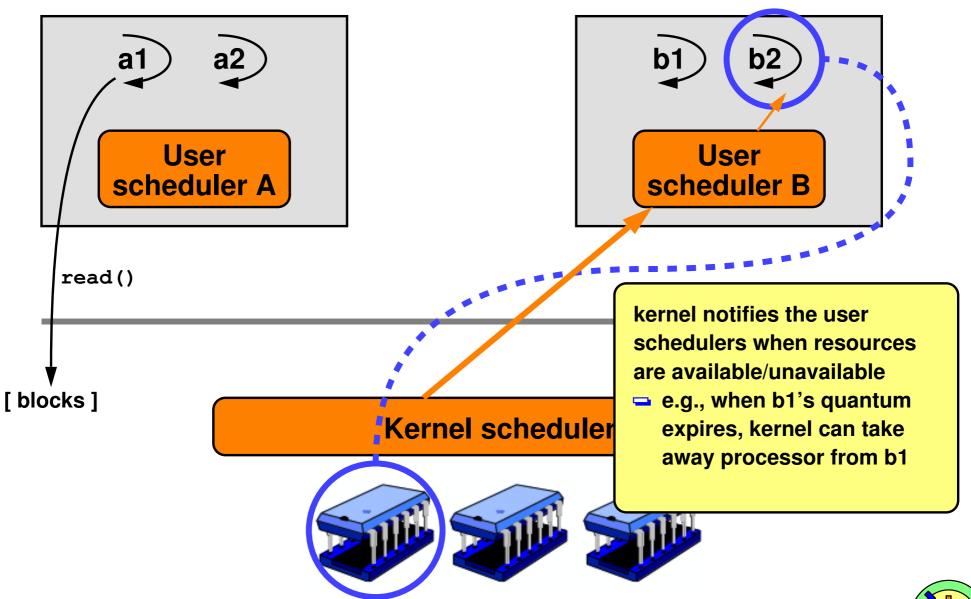






Kernel scheduler can have various scheduling policies







Kernel scheduler can have various scheduling policies



5.1 Threads Implementations

Strategies

A Simple Thread Implementation

Multiple Processors



A Simple Threads Implementation



Threads implementation considerations

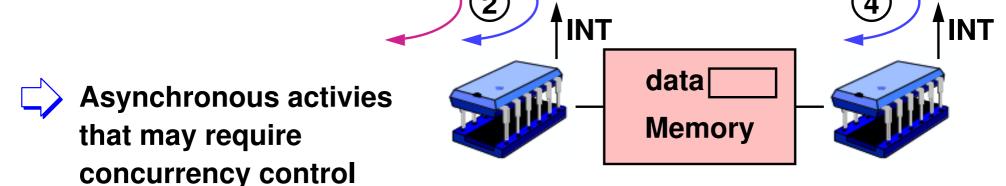
- data structures
- thread switching
- synchronization
 - how to implmement mutexes?
 - spin locks
 - sleep/blocking locks
 - futexes
 - please keep in mind that a mutex can be implemented in the kernel and in the user space



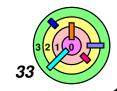
A Simple Threads Implementation

The challenge with implementing mutexes is that you have to ensure that they perform correctly under different kinds of





- 1) an *interrupt handler* running on the *same processor* that accesses the same data structure
- 2) another thread running on the same processor may preempt this thread and accesses the same data structure
- 3) an *interrupt handler* running on *another processor* might access the same data structure
- 4) another thread running on another processor might access the same data structure



A Simple Threads Implementation



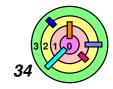
This implementation is the basis for user-level threads package

- "thread" can mean kernel thread or user thread
- mutex does not need to be a kernel data structure

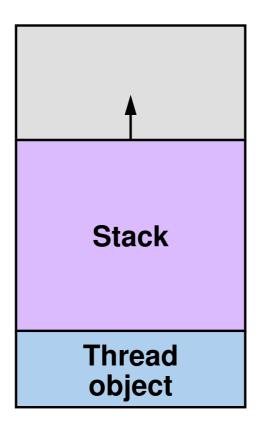


Straight-threads implementation

- everything happens in thread contexts
 - o no interrupt
 - therefore, no preemption
- one processor
- this is like your kernel 1 with DRIVERS=0 in Config.mk



Basic Representation



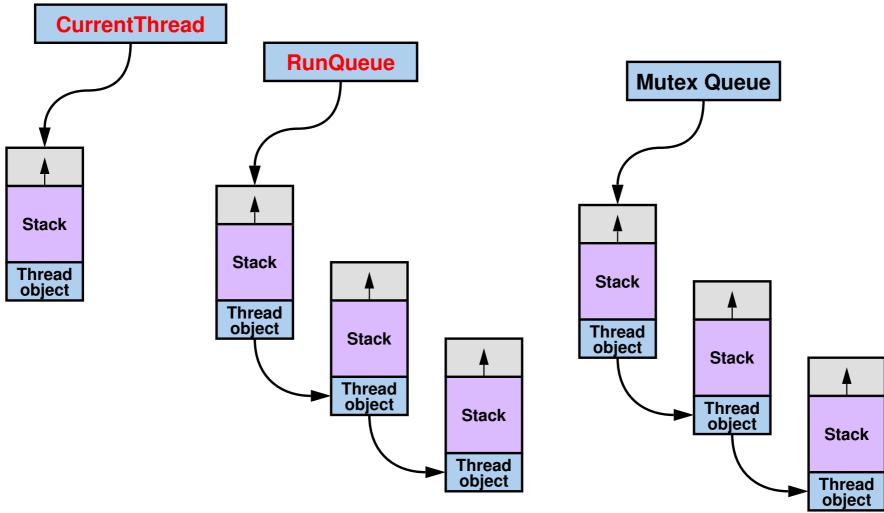


We will depict a thread like this (to be more compact)

 although we know that a thread control block is separated from a thread's stack

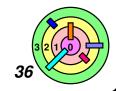


A Collection of Threads





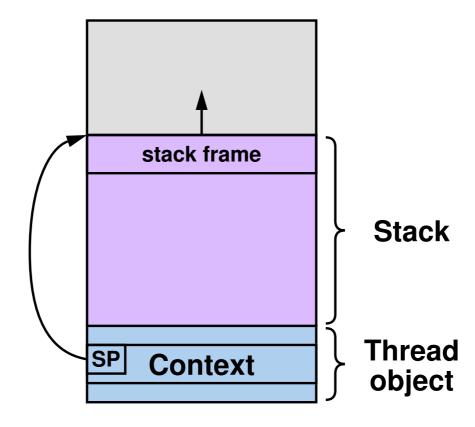
- your kernel assignment looks like this
 - o at any time, you should know where your threads are



Context Pointer



Recall from Ch 3



if this thread is not currently running, "stack frame" corresponds to switch()

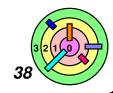


Straight-threads - Thread Switch



Need a thread_switch() function to yield the processor

- switch() in Ch 3 has a target thread argument
- swapcontext (old, new) saves the caller's context into the old context and restores from the new context
- note that the RunQueue may be empty, so this code is incomplete
- before you get here, the current thread is queued onto somewhere else already (e.g., a mutex queue)



Straight-threads - Synchronization



According to the textbook

```
void mutex_lock(mutex_t *m) {
  if (m->locked) {
    enqueue (m->queue, CurrentThread);
    thread_switch();
  } else
    m->locked = 1;
void mutex unlock(mutex t *m) {
  if (queue_empty(m->queue))
    m->locked = 0;
  else
    enqueue (runqueue, dequeue (m->queue));
```

- mutex_unlock() does not seem to work because when it returns, the mutex can be locked and the new mutex holder is not holding the mutex
- after further analysis, it actually does work!

Straight-threads - Synchronization

```
void mutex_lock(mutex_t *m) {
  if (m->locked) {
    enqueue (m->queue, CurrentThread);
    thread_switch();
  } else
    m->locked = 1;
void mutex_unlock(mutex_t *m) {
  if (queue_empty(m->queue))
    m->locked = 0;
  else
    enqueue (runqueue, dequeue (m->queue));
```



Why is the code atomic?



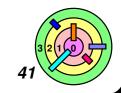
Straight-threads - Synchronization

```
void mutex_lock(mutex_t *m) {
  if (m->locked) {
    enqueue (m->queue, CurrentThread);
    thread_switch();
  } else
    m->locked = 1;
void mutex_unlock(mutex_t *m) {
  if (queue_empty(m->queue))
    m->locked = 0;
  else
    enqueue (runqueue, dequeue (m->queue));
```



Why is the code atomic?

- single process and no interrupts
- no way to preempt a thread's execution
 - a thread holds on to the processor as long as it wants, until it reliquishes processor all by itself



5.1 Threads Implementations



A Simple Thread Implementation

Multiple Processors



Straight-threads - Multiple Processors



thread_switch() is no longer sufficient

it's meant for uniprocessor

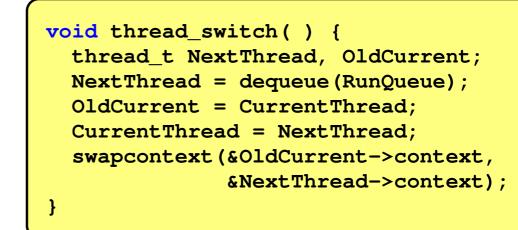


Simple approach

run on each processor an idle thread

```
void idle_thread() {
   while(1) {
     enqueue(runqueue, CurrentThread)
     thread_switch()
   }
}
```

- code is incomplete (because thread_switch() is incomplete, the way it was presented here)
- this thread never blocks, so there is always something to run to avoid boundary condition
- normal threads join the RunQueue when ready

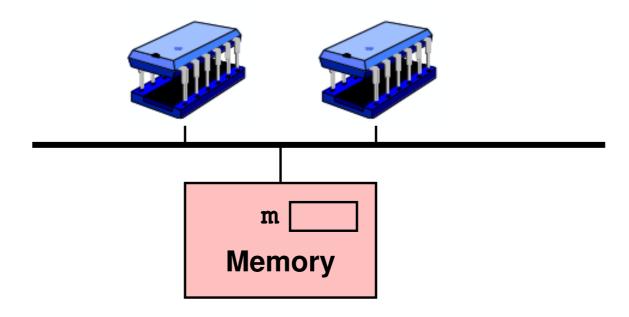


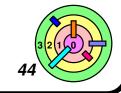
Straight-threads - Multiple Processors

When there are multiple processors, the difficulty lies in locking

```
if (!m->locked) {
  m->locked = 1;
}
```

if both threads execute the above code concurrently, in different processors, both threads think they got the lock

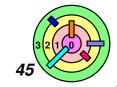






```
int CAS(int *ptr, int old, int new) {
   int tmp = *ptr; // get the value of mutex
   if (tmp == old) // if it equals to old
      *ptr = new; // set it to new
   return tmp; // return old
}
```

- often implemented as a machine-level instruction
 - must execute atomically
 - how?





```
int CAS(int *ptr, int old, int new) {
   int tmp = *ptr; // get the value of mutex
   if (tmp == old) // if it equals to old
      *ptr = new; // set it to new
   return tmp; // return old
}
```

- e.g., assume mutex is *unlocked*, call CAS (&lock, 0, 1)
 - mutex is represented as a bit, 0 if unlocked, 1 if locked

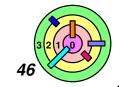
```
A[0..31] —

D[0..31] —

RD _

WR _

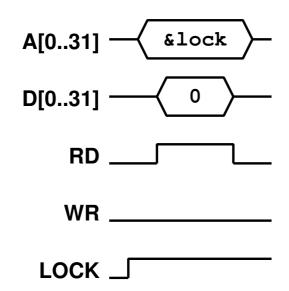
LOCK _
```





```
int CAS(int *ptr, int old, int new) {
  int tmp = *ptr; // get the value of mutex
  if (tmp == old) // if it equals to old
     *ptr = new; // set it to new
  return tmp; // return old
}
```

- e.g., assume mutex is *unlocked*, call CAS (&lock, 0, 1)
 - mutex is represented as a bit, 0 if unlocked, 1 if locked

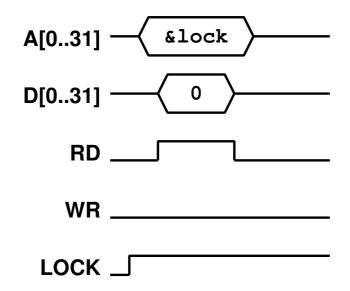


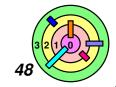




```
int CAS(int *ptr, int old, int new) {
   int tmp = *ptr; // get the value of mutex
   if (tmp == old) // if it equals to old
      *ptr = new; // set it to new
   return tmp; // return old
}
```

- e.g., assume mutex is *unlocked*, call CAS (&lock, 0, 1)
 - mutex is represented as a bit, 0 if unlocked, 1 if locked

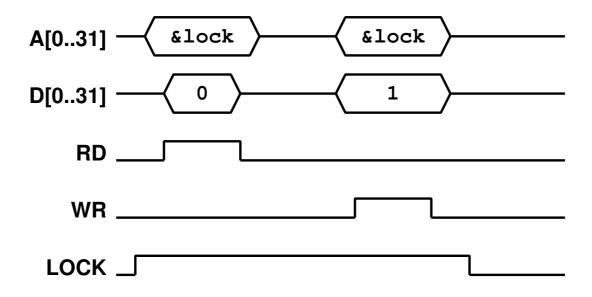






```
int CAS(int *ptr, int old, int new) {
   int tmp = *ptr; // get the value of mutex
   if (tmp == old) // if it equals to old
     *ptr = new; // set it to new
   return tmp; // return old
}
```

- e.g., assume mutex is *unlocked*, call CAS (&lock, 0, 1)
 - mutex is represented as a bit, 0 if unlocked, 1 if locked

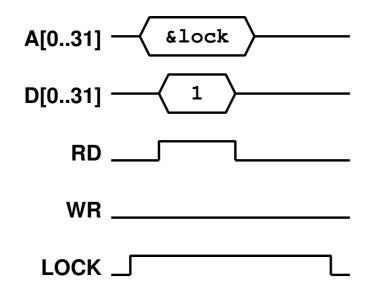


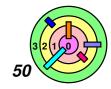




```
int CAS(int *ptr, int old, int new) {
   int tmp = *ptr; // get the value of mutex
   if (tmp == old) // if it equals to old
      *ptr = new; // set it to new
   return tmp; // return old
}
```

- e.g., assume mutex is *locked*, call CAS (&lock, 0, 1)
 - mutex is represented as a bit, 0 if unlocked, 1 if locked







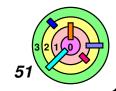
Compare and swap *machine instruction*

```
int CAS(int *ptr, int old, int new) {
   int tmp = *ptr; // get the value of mutex
   if (tmp == old) // if it equals to old
     *ptr = new; // set it to new
   return tmp; // return old
}
```

- often implemented as a machine-level instruction
 - must execute atomically



mutex is represented as a bit, 0 if unlocked, 1 if locked



Spin Lock



Naive spin lock

```
void spin_lock(int *mutex) {
    while(CAS(mutex, 0, 1)) // textbook is wrong
    ;
}

void spin_unlock(int *mutex) {
    *mutex = 0;
}
```

Better spin lock

```
void spin_lock(int *mutex) {
  while (1) {
    if (*mutex == 0) {
        // the mutex was at least momentarily unlocked
        if (!CAS(mutex, 0, 1))
            break; // we have locked the mutex
        // some other thread beat us to it, try again
    }
}
```



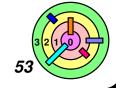
Spin locks are wasteful

- processor time wasted waiting for the lock to be released
- barely acceptable if locks are held only briefly



A better approach is to have a blocking lock

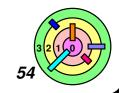
- threads wait by having their execution suspended
- a thread much yield the processor and join a queue of waiting threads
 - later on, get resumed explicitly



```
void blocking_lock(mutex_t *m) {
  if (m->holder != 0) {
    enqueue (m->wait_queue, CurrentThread);
    thread_switch();
  } else
    m->holder = CurrentThread;
void blocking_unlock(mutex_t *m) {
  if (queue_empty(m->wait_queue))
    m->holder = 0;
  else {
    m->holder = dequeue(m->wait_queue);
    enqueue (RunQueue, m->holder);
```



This code only works on a uniprocessor



```
void blocking_lock(mutex_t *m) {
  if (m->holder != 0) {
    enqueue (m->wait_queue, CurrentThread);
    thread_switch();
  } else
    m->holder = CurrentThread;
void blocking_unlock(mutex_t *m) {
  if (queue_empty(m->wait_queue))
    m->holder = 0;
  else {
    m->holder = dequeue(m->wait_queue);
    enqueue (RunQueue, m->holder);
```



On a multiprocessor, it may not work

threads 1 and 2 can both think they've got the lock



```
void blocking_lock(mutex_t *m) {
  if (m->holder != 0) {
    enqueue (m->wait_queue, CurrentThread);
    thread_switch();
   else
    m->holder = CurrentThread;
void blocking_unlock(mutex_t *m) {
  if (queue_empty(m->wait_queue))
    m->holder = 0;
  else {
    m->holder = dequeue(m->wait_queue);
    enqueue (RunQueue, m->holder);
```



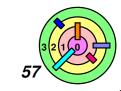
- thread 2 holds the mutex and wait queue is empty and thread 1 tries to lock the mutex at the same time thread 2 is releasing the mutex
- thread 1 may wait forever



```
void blocking_lock(mutex_t *m) {
  spin_lock(m->spinlock); // okay to spin here
  if (m->holder != 0) {
    enqueue (m->wait_queue, CurrentThread);
    thread_switch();
  } else {
    m->holder = CurrentThread;
  spin_unlock (m->spinlock);
void blocking_unlock(mutex_t *m) {
  spin_lock(m->spinlock); // okay to spin here
  if (queue_empty(m->wait_queue)) {
    m->holder = 0;
  } else {
    m->holder = dequeue(m->wait_queue);
    enqueue (RunQueue, m->holder);
  spin_unlock(m->spinlock);
```

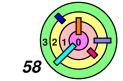


Will deadlock because of thread_switch()



```
void blocking_lock(mutex_t *m) {
  spin_lock (m->spinlock);
  if (m->holder != 0) {
    enqueue (m->wait_queue, CurrentThread);
    spin_unlock(m->spinlock);
    thread_switch();
  } else {
    m->holder = CurrentThread;
    spin_unlock (m->spinlock);
void blocking unlock(mutex t *m) {
  spin_lock (m->spinlock);
  if (queue_empty(m->wait_queue)) {
    m->holder = 0;
  } else {
    m->holder = dequeue(m->wait_queue);
    enqueue (RunQueue, m->holder);
  spin_unlock(m->spinlock);
```



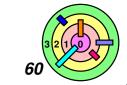


```
void blocking_lock(mutex_t *m) {
  spin_lock (m->spinlock);
  if (m->holder != 0) {
    enqueue (m->wait_queue, CurrentThread);
    spin_unlock(m->spinlock);
    thread_switch();
  } else {
    m->holder = CurrentThread;
    spin_unlock (m->spinlock);
void blocking unlock(mutex t *m) {
  spin_lock (m->spinlock);
  if (queue_empty(m->wait_queue)) {
    m->holder = 0;
  else {
    m->holder = dequeue(m->wait_queue);
    enqueue (RunQueue, m->holder);
  spin_unlock (m->spinlock);
```

59

```
void blocking_lock(mutex_t *m) {
  spin_lock (m->spinlock);
  if (m->holder != 0) {
    enqueue (m->wait_queue, CurrentThread);
    spin_unlock(m->spinlock);
    thread_switch();
  } else {
    m->holder = CurrentThread;
    spin_unlock (m->spinlock);
void blocking unlock(mutex t *m) {
  spin_lock (m->spinlock);
  if (queue_empty(m->wait_queue)) {
    m->holder = 0;
  else {
    m->holder = dequeue(m->wait_queue);
    enqueue (RunQueue, m->holder);
  spin_unlock (m->spinlock);
Can you do spin_unlock() inside thread_switch()?
```





Futexes



Futex: fast user-space mutex

- safe, efficient kernel conditional queueing in Linux
 - most of the time when you try to lock a mutex, it's unlocked; so just go ahead and lock it (no system call)
 - if it's locked (by another thread), then a system call is required for this thread to obtain the lock
- contained in it is an unsigned integer state called value and a queue of waiting threads

Two system calls are provided to support futexes

```
futex_wait(futex_t *futex, int val) {
  if (futex->val == val)
    sleep();
}

futex_wake(futex_t *futex) {
  // wake up one thread from wait queue if
  // there is any
  ...
}
```

Ancillary Functions

```
Add 1 to *val, return its original value
unsigned int atomic_inc(unsigned int *val) {
```

```
// performed atomically
return((*val)++); // textbook is wrong
```

Subtract 1 to *val, return its original value

```
unsigned int atomic_dec(unsigned int *val) {
   // performed atomically
   return((*val)--); // textbook is wrong
}
```

Just like CAS(), both functions return the previous lock value



Attempt 1

```
futex->val
```

0 means unlocked; otherwise, locked

```
void lock(futex_t *futex) {
   unsigned int c;
   while ((c = atomic_inc(&futex->val)) != 0)
     futex_wait(futex, c+1);
}

void unlock(futex_t *futex) {
   futex->val = 0;
   futex_wake(futex);
}
```



Problem with unlock()

slow because futex_wake() is a system call



Problem with lock()

- threads run in lock steps in a multiprocessor environment!
- futex->val may wrap-around



Attempt 2

- futex->val can only take on values of 0, 1, and 2
- 0 means unlocked
- 1 means locked but no waiting thread
- **2** means locked with the possibility of waiting threads

```
void lock(futex_t *futex) {
                                                   textbook
  unsigned int c;
                                                   is wrong
  if ((c = CAS(\&futex->val, 0, 1) != 0)
    do {
      if (c == 2 || (CAS(&futex->val, 1, 2) == 1))
        futex wait(futex, 2);
    } while ((c = CAS(&futex->val, 0, 2)) != 0));
void unlock(futex_t *futex) {
  if (atomic_dec(&futex->val) != 1) {
    futex->val = 0;
    futex wake(futex);
```

Thread Synchronization Summary



- used if the duration of waiting is expected to be small
 - as in the case at the beginning of blocking_lock()
- Sleep (or blocking) locks
 - used if the duration of waiting is expected to be long
- Futexes
 - optimized version of blocking locks
- In your kernel assignmen #1, you need to implement kernel threads
 - very different from user threads
 - keep in mind that the weenix kernel is non-preemptive
 - the kernel is very powerful (and therefore, must be bug free)
 - in kernel assignmen #3, you need to implement user threads/processes (well, still one thread per process)

