


5.2 Interrupts

Interrupt Handling - Overview



Thread Synchronization

Recall asynchronous activities that may require concurrency control

- 1) an *interrupt handler* running on the *same processor* that accesses the same data structure
- 2) *another thread* running on the *same processor* may *preempt* this thread and accesses the same data structure
- 3) an *interrupt handler* running on *another processor* might access the same data structure
- 4) *another thread* running on *another processor* might access the same data structure

let's look at (1) and (3) now



Interrupt Handling

What to do if you have *preemption kernels*?

- threads running in privileged mode may be forced to yield the processor
- so you can use *disable preemption*
- then you can use interrupt masking
- spin locks*



Interrupt Masking

Unmasked interrupts interrupt current processing

What causes interrupts to be masked?

- the occurrence of a particular class of interrupts masks further occurrences
- explicit programmable action

Some architectures impose a hierarchy of interrupt levels

- Intel architectures use APIC
- advanced programmable interrupt controller



Interrupt Handling

We are focusing on dealing with synchronization/concurrency issues

What to do if you have *non-preemption kernels*?

- in these systems, a kernel thread can *never* be preempted *by another thread*
- threads running in privileged mode yield the processor only voluntarily
- this makes the kernel a lot easier to implement! because don't have to implement *locking inside the kernel* for every shared data structure (although sometimes, mutex is still needed to synchronize kernel threads)
- done in early Unix systems
- done in weenix
- this is like your kernel 1 with `DRIVERS=1` in `Config.mk`
- use *interrupt masking*



Copyright © William C. Cheng

Non-Preemptive Kernel Synchronization

```

int x = 0;

void AccessXThread() {
    ...
    x = x+1;
    ...
}

```

Sharing a variable between a thread and an interrupt handler

since we have a non-preemptive kernel, the only thing that can prevent a kernel thread from executing till completion is an interrupt

The above code does not work

cannot use locks to fix it

7

Copyright © William C. Cheng

Example: Disk I/O

```

int disk_write(...) {
    ...
    startIO(); // start disk operation
    enqueue(disk_waitq, CurrentThread);
    thread_switch();
    // wait for disk operation to complete
    ...
    thread_t *thread;
    void disk_intr(...) {
        // handle disk interrupt
        if (thread != 0) {
            thread = deque(disk_waitq);
            enqueue(RunQueue, thread);
            // wakeup waiting thread
        }
    }
    disk_write();
}

```

App

File System

9

Copyright © William C. Cheng

Improved Disk I/O

```

int disk_write(...) {
    ...
    oldIPL = setIPL(diskIPL);
    startIO(); // start disk operation
    enqueue(disk_waitq, CurrentThread);
    thread_switch();
    // wait for disk operation to complete
    ...
    setIPL(oldIPL);
}

```

Solution

mask disk interrupt

11

Copyright © William C. Cheng

Non-Preemptive Kernel Synchronization

```

int x = 0;

void AccessXThread() {
    ...
    setIPL(oldIPL);
    x = x+1;
    ...
}

void AccessXInterrupt() {
    ...
    int x = 0;
}

```

Solution is to mask the interrupt

works well in a non-preemptive kernel

8

Copyright © William C. Cheng

Example: Disk I/O

```

int disk_write(...) {
    ...
    startIO(); // start disk operation
    enqueue(disk_waitq, CurrentThread);
    thread_switch();
    // wait for disk operation to complete
    ...
    thread_t *thread;
    void disk_intr(...) {
        // handle disk interrupt
        if (thread != 0) {
            thread = deque(disk_waitq);
            enqueue(RunQueue, thread);
            // wakeup waiting thread
        }
    }
    disk_write();
}

```

Problem

disk may be too fast

disk_intr() gets called before enqueue()

this is a synchronization problem / race condition

10

Copyright © William C. Cheng

Improved Disk I/O

```

int disk_write(...) {
    ...
    oldIPL = setIPL(diskIPL);
    startIO(); // start disk operation
    enqueue(disk_waitq, CurrentThread);
    thread_switch();
    // wait for disk operation to complete
    ...
    setIPL(oldIPL);
}

```

Solution

mask disk interrupt

Doesn't quite work!

thread_switch() will switch to another thread and won't return back here any time soon to unmask interrupt

who will enable the disk interrupt?

complication caused by the fact that thread_switch() does not function like a normal procedure call

moving setIPL(oldIPL) to before thread_switch() may have race condition in accessing the RunQueue

12

Copyright © William C. Cheng

17

Does it work?

= no, can deadlock in AccessXInterrupt () in case (1)

```

void AccessXThread () {
    SpinLock (&L);
    X = X+1;
    SpinUnlock (&L);
    ...
}

void AccessXInterrupt () {
    SpinLock (&L);
    X = X+1;
    SpinUnlock (&L);
    ...
}
    
```

Solution?

Operating Systems - CSCI 402

Copyright © William C. Cheng

18

Does it work?

Solution ...

```

void AccessXThread () {
    DisableInterrupt ();
    SpinLock (&L);
    X = X+1;
    SpinUnlock (&L);
    ...
}

void AccessXInterrupt () {
    SpinLock_t L = UNLOCKED;
    int x = 0;
    ...
}
    
```

Operating Systems - CSCI 402

Copyright © William C. Cheng

15

Modified thread_switch

```

void thread_switch () {
    thread_t *OldThread;
    int oldIPL;
    oldIPL = setIPL (HIGH_IPL);
    // protect access to RunQueue by
    // masking all interrupts
    while (queue_empty (RunQueue)) {
        // repeatedly allow interrupts, then
        // RunQueue
        setIPL (0); // 0 means no interrupts
        // should halt the CPU
        setIPL (HIGH_IPL);
    }
    // We found a runnable thread
    OldThread = dequeue (RunQueue);
    swapContext (OldThread->context,
        CurrentThread->context);
    setIPL (oldIPL);
}
    
```

If you decide to halt the CPU
in weatix, need to watch out
for a race condition
= it does not "wait" properly
= the correct way to wait for
an asynchronous event is:
1) disable/block it
2) enable/unblock and
wait for it in one atomic
operation

Operating Systems - CSCI 402

Copyright © William C. Cheng

16

Preemptive kernels

What's different?

Recall asynchronous activities that may require concurrency control

- 1) an *interrupt handler* running on the *same processor* that accesses the same data structure
- 2) *another thread* running on the *same processor* may *preempt* this thread and accesses the same data structure
- 3) an *interrupt handler* running on *another processor* might access the same data structure
- 4) *another thread* running on *another processor* might access the same data structure

access the same data structure

Operating Systems - CSCI 402

Copyright © William C. Cheng

13

Modified thread_switch

```

void thread_switch () {
    thread_t *OldThread;
    int oldIPL;
    oldIPL = setIPL (HIGH_IPL);
    // protect access to RunQueue by
    // masking all interrupts
    while (queue_empty (RunQueue)) {
        // repeatedly allow interrupts, then check
        // RunQueue
        setIPL (0); // 0 means no interrupts are masked
        setIPL (HIGH_IPL);
    }
    // We found a runnable thread
    OldThread = CurrentThread;
    CurrentThread = dequeue (RunQueue);
    swapContext (OldThread->context,
        CurrentThread->context);
    setIPL (oldIPL);
}
    
```

Operating Systems - CSCI 402

Copyright © William C. Cheng

14

Modified thread_switch

```

void thread_switch () {
    thread_t *OldThread;
    int oldIPL;
    oldIPL = setIPL (HIGH_IPL);
    // protect access to RunQueue by
    // masking all interrupts
    while (queue_empty (RunQueue)) {
        // repeatedly allow interrupts, then check
        // RunQueue
        setIPL (0); // 0 means no
        setIPL (HIGH_IPL);
    }
    // We found a runnable thread
    OldThread = CurrentThread;
    CurrentThread = dequeue (RunQueue);
    swapContext (OldThread->context,
        CurrentThread->context);
    setIPL (oldIPL);
}
    
```

This code is actually much more tricky that it looks
= it can be invoked by a thread that's not doing I/O
= oldIPL is the oldIPL of a different thread!

Let's say that another thread calls thread_switch ()
= it's not doing I/O
= its oldIPL is set to 0

Now we call thread_switch ()
= our oldIPL set to diskIPL
= then we switch to this other thread and set IPL to 0
= when *all interrupts blocked*
= RunQueue only accessed

Operating Systems - CSCI 402

Copyright © William C. Cheng

23

Interrupt Threads?

Solaris allows interrupts to be handled as threads

Does it make sense to handle interrupts with threads?

- perhaps similar to using `sigwait` for handling signals with threads
- what would be the advantages?
- what would be the disadvantages?

Operating Systems - CSCI 402

Copyright © William C. Cheng

24

Interrupt Threads

```

void InterruptHandler ( ) {
    // deal with interrupt
    ...
    if (!MoreWork)
        return;
    else
        BecomeThread ( ) ;
    P (Semaphore) ; // sleep!
    ...
}

```

Operating Systems - CSCI 402

Copyright © William C. Cheng

21

Solution ...

```

void AccessXThread () {
    DisableInterrupts () ;
    ...
    SpinLock (&L) ;
    X = X+1 ;
    SpinUnlock (&L) ;
    ...
    UnmaskInterrupts () ;
    EnableInterrupt () ;
}

```

Does it work?

Operating Systems - CSCI 402

Copyright © William C. Cheng

22

Solution ...

```

void AccessXThread () {
    DisableInterrupts () ;
    ...
    SpinLock (&L) ;
    X = X+1 ;
    SpinUnlock (&L) ;
    ...
    UnmaskInterrupts () ;
    EnableInterrupt () ;
}

```

Does it work?

Operating Systems - CSCI 402

Copyright © William C. Cheng

19

Solution ...

```

void AccessXThread () {
    DisableInterrupts () ;
    ...
    SpinLock (&L) ;
    X = X+1 ;
    SpinUnlock (&L) ;
    ...
    UnmaskInterrupts () ;
    EnableInterrupt () ;
}

```

Does it work?

Operating Systems - CSCI 402

Copyright © William C. Cheng

20

Solution ...

```

void AccessXThread () {
    DisableInterrupts () ;
    ...
    SpinLock (&L) ;
    X = X+1 ;
    SpinUnlock (&L) ;
    ...
    UnmaskInterrupts () ;
    EnableInterrupt () ;
}

```

Does it work?

Operating Systems - CSCI 402

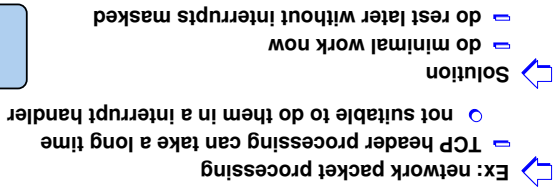
	Interrupt #23's handler frame	Interrupt #15's handler frame	Interrupt #7's handler frame	Kernel stack frames	Current thread's kernel stack
--	-------------------------------	-------------------------------	------------------------------	---------------------	-------------------------------

Operating Systems - CSCI 402

- Windows handles deferred work in a special interrupt context
- DPC (deferred procedure call) is a *software interrupt*

Operating Systems - CSCI 402


Operating Systems - CSCI 402 -



Operating Systems - CSCI 402 -



Operating Systems - CSCI 402 -



Copyright © William C. Cheng

35

```

void TopLevelInterruptHandler(int dev) {
    InterruptVector[dev]();
    if (PreviousMode == UserMode) {
        // the clock interrupted user-mode code
        if (ShouldReschedule)
            Reschedule();
    }
}

// If interrupted a user thread
void TopLevelTrapHandler(...) {
    SpecialTrapHandler();
    if (ShouldReschedule) {
        // the time slice expired
        while the thread was
        in kernel mode */
        Reschedule();
    }
}

```

➡ If interrupted a kernel thread

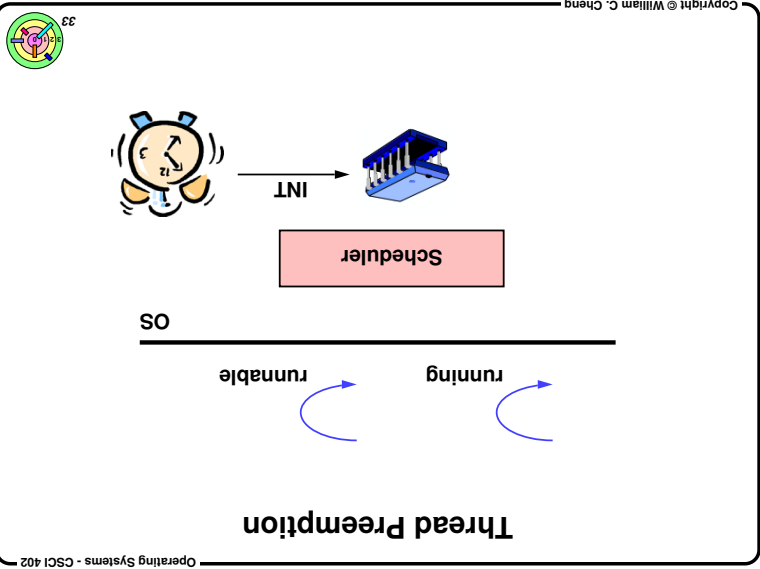
➡ If interrupted a user thread

Preemption: User-Level Only

The work of **rescheduling** is **deferred**

Reschedule() puts the calling thread on the run queue

➡ then call thread_switch() to give up the processor



Copyright © William C. Cheng

31

```

void InterruptHandler() {
    // deal with interrupt
    QueueDPC(MoreWork);
    /* requests a DPC interrupt here */
}

void DPCHandler(...) {
    while (!Empty(DPCQueue)) {
        Work = Dequeue(DPCQueue);
        Work();
    }
}

```

Deferred Procedure Calls

Copyright © William C. Cheng

36

```

void ClockInterruptHandler() {
    // deal with clock interrupt
    if (TimesliceOver()) {
        QueueDPC(Reschedule);
        /* requests a DPC interrupt here */
    }
}

```

➡ Preemptive kernel

➡ preemption can happen for a kernel thread

➡ if clock-interrupt happens, setup the kernel thread to give up the processor when the processor is about to return to the thread's context

➡ how?

➡ e.g., add the Reschedule() function to the DPC queue

Preemption: Full (i.e., Preemptive Kernel)

Copyright © William C. Cheng

34

```

void ClockHandler() {
    // deal with clock interrupt
    if (TimesliceOver())
        ShouldReschedule = 1;
}

```

➡ Non-preemptive kernel

➡ preempt only threads running in user mode

➡ if clock-interrupt happens, just set a global flag

Preemption: User-Level Only

Copyright © William C. Cheng

32

```

void SoftwareInterruptThread() {
    while (TRUE) {
        WaitEvent(Work);
        while (!Empty(WorkQueue)) {
            Work = Dequeue(WorkQueue);
            Work();
        }
    }
}

void InterruptHandler() {
    // deal with interrupt
    Enqueue(WorkQueue, MoreWork);
    SetEvent(Work);
}

```

➡ Linux handles deferred work in a special kernel thread

➡ this kernel thread is scheduled like any other kernel thread

Software Interrupt Threads

Copyright © William C. Cheng

41

Invoking the Signal Handler (2)

Diagram illustrating the state of the system before an interrupt. The Main Line is executing a function `func(int a1, int a2)` with `i=1, j=2`. The Handler is ready to execute. The User Stack contains `Previous frames` and a `func()` frame. The Kernel Stack contains `Registers` and `User` IP, which is set to return to the user context.

Copyright © William C. Cheng

42

Invoking the Signal Handler (3)

Diagram illustrating the state of the system after an interrupt. The Main Line has saved its context (`i=1, j=2`) into the Handler's stack frame. The User Stack now contains a new frame for the signal handler, with its IP set to return to the kernel stack. The Kernel Stack's User IP is now set to the address of the signal handler frame.

Copyright © William C. Cheng

39

Invoking the Signal Handler

Basic idea is to set up the user stack so that the handler is called as a subroutine and so that when it returns, normal execution of the thread may continue

Complications:

- must first save **all** registers and later on restore all of them
- signal mask
- must block the signal and later on unblock the signal
- therefore, when the signal handler returns, it needs to return to some code that restores all the registers, unblocks the signal, then return to the interrupted code

Copyright © William C. Cheng

40

Invoking the Signal Handler (1)

Diagram illustrating the state of the system after the signal handler has been invoked. The Main Line's context (`i=1, j=2`) is saved in the Handler's stack frame. The User Stack contains the signal handler's frame, and the Kernel Stack's User IP points to the signal handler's frame.

Copyright © William C. Cheng

37

Interrupt Handling - Overview

Diagram illustrating the flow of interrupt handling. Interrupts are processed by the Interrupt Handling module, which then directs processing to the Threads, APCs, or Signals. Deferred Work is also shown as a path from Interrupt Handling to Threads.

Copyright © William C. Cheng

38

Directed Processing

Signals: Unix

- perform given action in context of a particular thread in user mode
- e.g., seg fault
- generated by hardware and needs to be delivered to the user process to invoke a signal handler

APC: Windows Asynchronous Procedure Calls

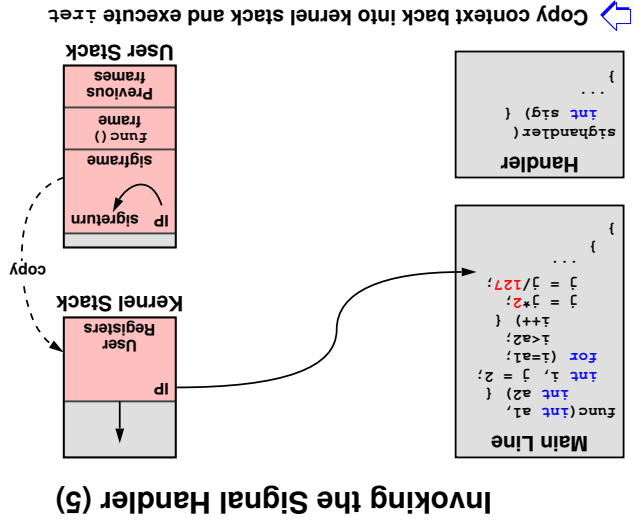
- roughly same thing, but also may be done in **kernel mode**
- thus, the APC mechanism is **more general** than Unix signals

software	Thread	APC	DPC
0	Thread		
1		APC	
2			DPC

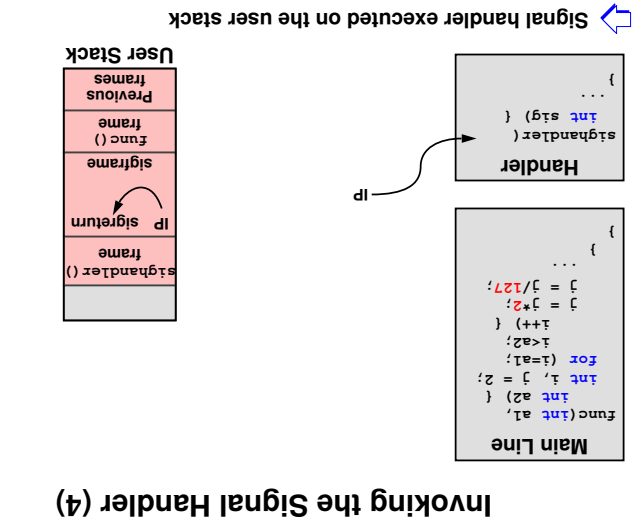
Windows Interrupt Priority Levels

Extra Slides

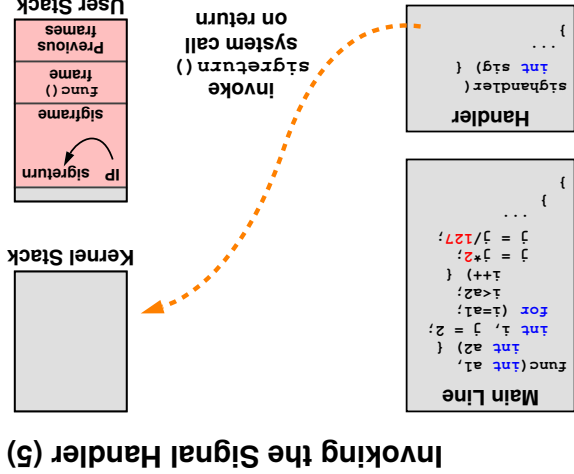
- Two uses
- kernel APC: release of kernel resources
 - user APC: notifying a thread of an external event
- ## Asynchronous Procedure Calls



Invoking the Signal Handler (5)



Invoking the Signal Handler (4)



Invoking the Signal Handler (5)

Copyright © William C. Cheng

51

APC Implementation

- Per-thread list of pending APCs
 - on notification, thread executes them
- User APC
 - thread in alertable state is woken up and executes pending APCs when it returns to user mode
- Kernel APC
 - running thread interrupted by APC interrupt (lowest priority interrupt)
 - waiting thread is "unwaited"
 - execute pending kernel APCs

Copyright © William C. Cheng

49

Kernel APC

- Release of kernel resources
 - interrupt handler cannot free storage for buffer and control blocks until info passed to process
 - can't be done unless in context of process
 - otherwise address space not mapped in
 - interrupt handler requests kernel APC to have thread, running in kernel mode, absorb info in buffer and control blocks and then free them

Copyright © William C. Cheng

50

User APC

- Notifying thread of external event
 - example: asynchronous I/O
 - thread supplies *completion routine* when starting asynchronous I/O request
 - called in thread's context when I/O completes
 - similar to a Unix signal
 - called only when thread is in *alertable wait state*
 - an option in certain blocking system calls