

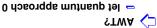
Max-min Fairness criterion:

1) no user receives more than its request
2) no other allocation scheme satisfying condition 1 has a high
minimum allocation
3) condition 2 remains recursively true as we remove the minimal
3) condition 2 remains recursively true as we remove the minimal
user and reduce total resource accordingly



FO

Round Robin + FIFO



- 169 jobs sharing the processor

= run at 1/169th speed for first week

short jobs receive one hour of processor time in 169 hours
 different from SJF since all short jobs finish at about the

same time
same time

long job completes in 336 hoursTWA = 169.99 hours

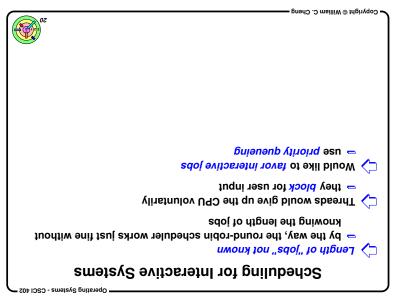
recall that AMT is 252 hours for FIFO in our example and $\,$ 162 for SIFO in our example and

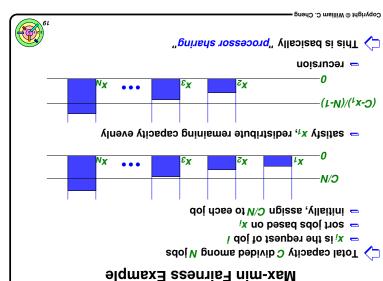
average deviation = 1.96 hours

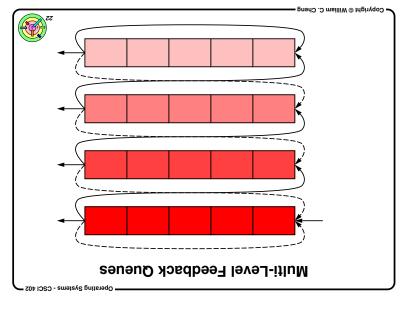
 recall that average deviation = 42.25 hours for FIFO in our example and 52 hours for SJF

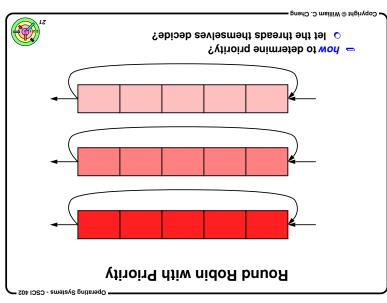


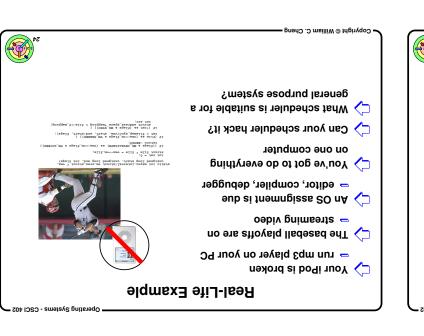
Copyright © William C. Cheng

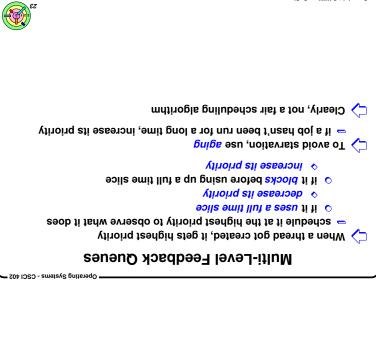












Proportional-Share Scheduling

Ctride scheduling

- 1995 paper by Waldspurger and Weihl

7002 ni xuni ot bebbs -Completely fair scheduling (CFS)

Stride Scheduling Algorithm

Time-sliced, priority-based, preemptive

- every thread is assigned a priority, called a pass value

o single queue, sorted based on pass values, smallest first

every thread is assigned a stride value

tickets in a lottery scheduling scheme stride values are computed according to distribution of

Ly In every iteration / time-slice

1) shedule the thread with the smallest pass value (at the head

of the queue)

3) increment the thread's pass value by its stride value 2) set the global pass value to be the pass value of this thread

as far as exam goes, stick to our presentation - the textbook presented Stride Scheduling differently

7*77777*77

Metered Processors

your thread is 5 times more likely to win than the others

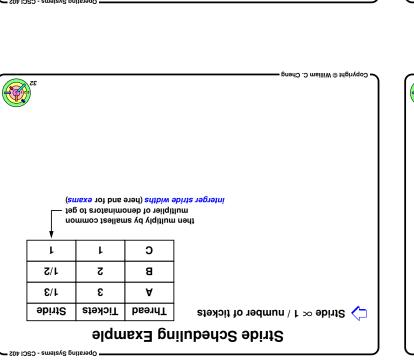
25 lottery tickets are distributed equally to you and your four

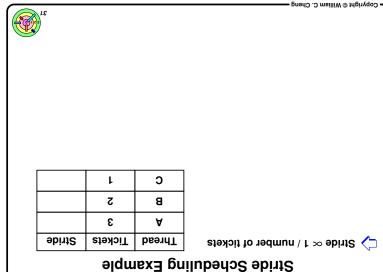
Lottery Scheduling

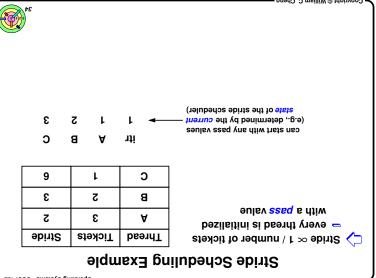
A lottery is held for every scheduling decision

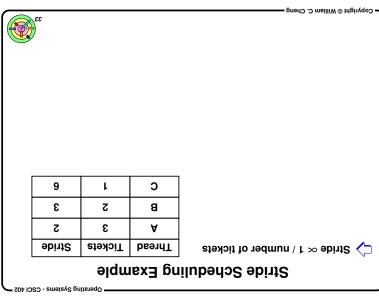
you give 5 tickets to your one thread

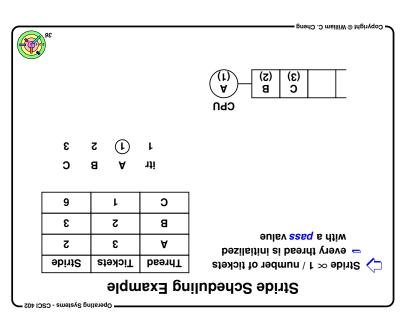
- they give one ticket each to their threads

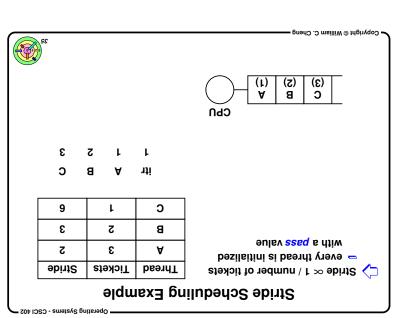




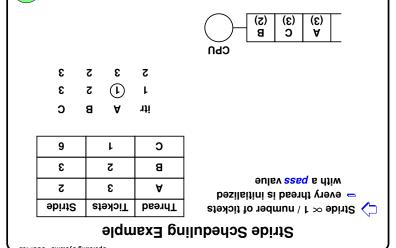


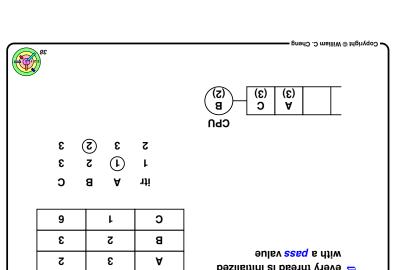


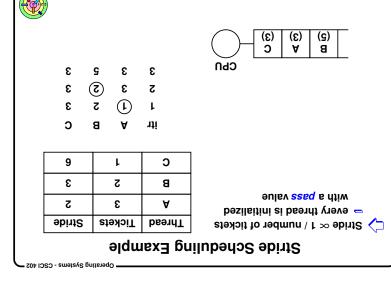


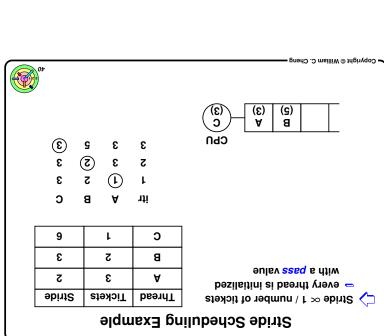


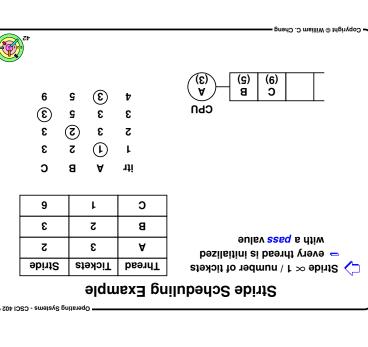


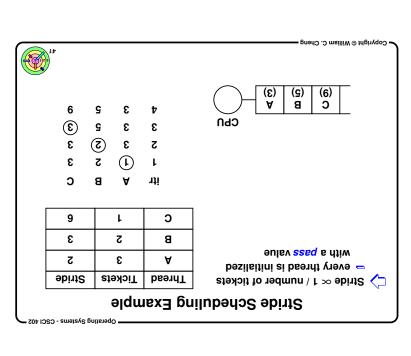


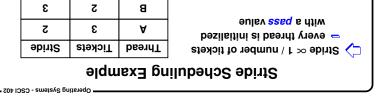


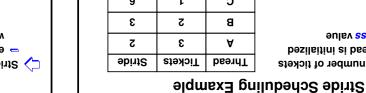












eulsv seed a djiw
- every thread is initialized
Stride \propto 1 $$ number of tickets

anlay <mark>ssi</mark>
bazilaitini zi baa
number of tickets
ataylait to undering

L	ဂ
2	8
3	A
	2

3

\bigcirc	(g) B	(G)	(6) O	Γ
СРИ		-		_

(6) O (2) B

Stride Scheduling Example

9	Ļ	၁
3	2	8
2	3	A
Stride	Tickets	Thread

(9)

3

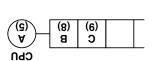
(I)

with a pass value - every thread is initialized Stride ∞ 1 / number of tickets



(9)

9



В

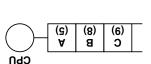
Stride Scheduling Example

3 (**!**)

9 3

3

Thread Tickets

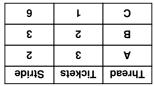


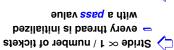
with a pass value

- every thread is initialized

Stride ∞ 1 / number of tickets

Stride Scheduling Example





to the distribution	Conforms of tickets!	\ -

9	Ļ	0
3	2	В
2	3	A
Stride	Tickets	Thread

9	L	ဂ
3	2	В
2	3	A
Stride	Tickets	Thread

3	a v	~*:
9	Ļ	Э
3	2	В
2	з	A
Stride	Tickets	Тһгеаd

J	a v	**!
	_	_
9	ŀ	၁
ε	2	æ
2	3	A
Stride	IICKGIS	ı ukesq

3	a A	Ήi
9	Ļ	0
3	2	В
2	3	A
Stride	SIEKEIS	ı uread



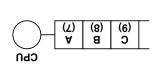
(9)

L

(9)

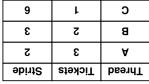
3

L



	-	_	 -	-	
Operating Systems - CSCI 402					

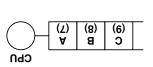
Stride Scheduling Example







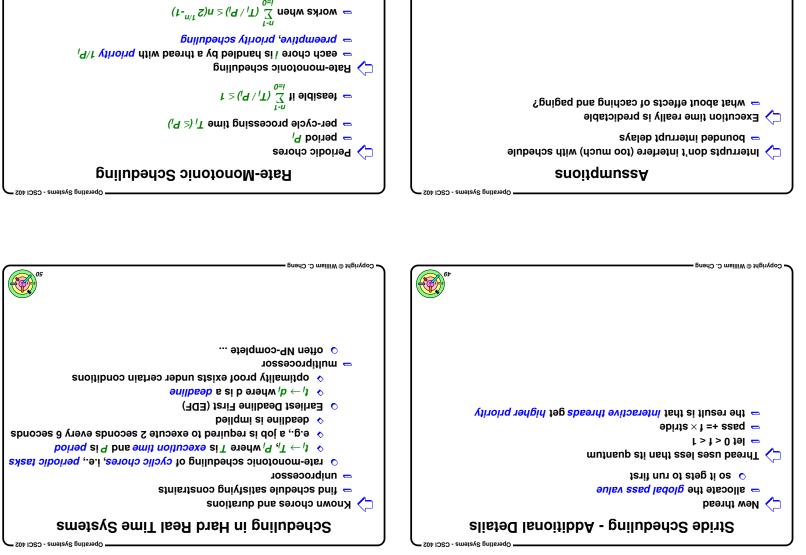


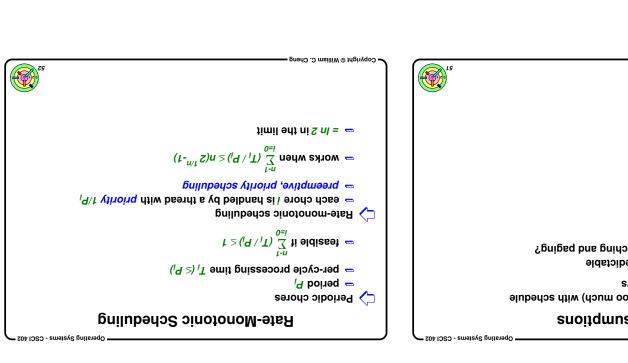


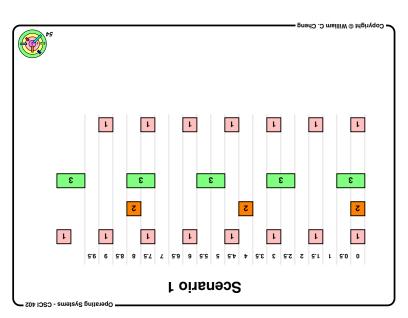
with a pass value

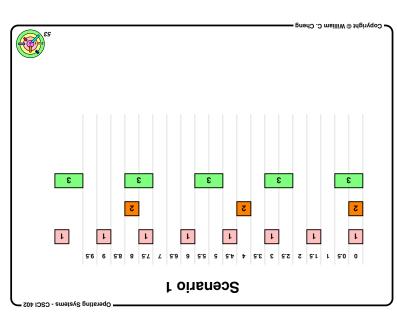
- every thread is initialized

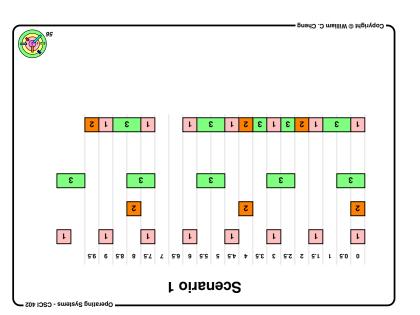
Stride ∞ 1 / number of tickets

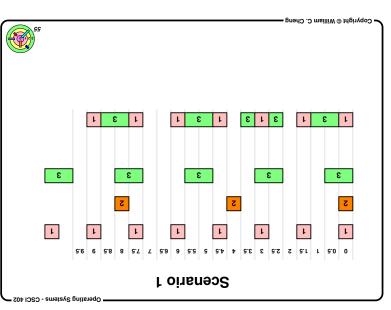


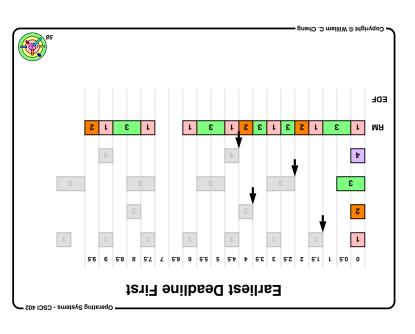


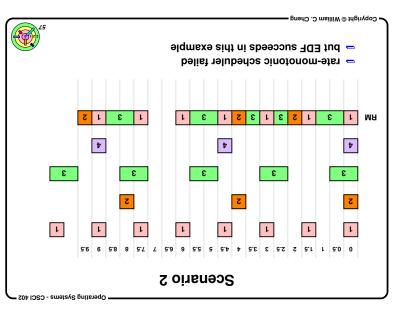


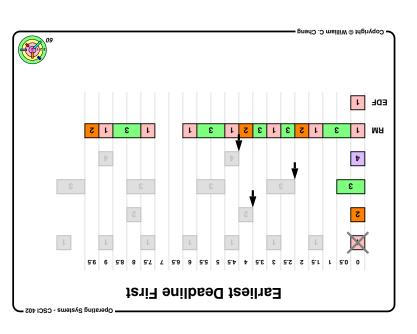


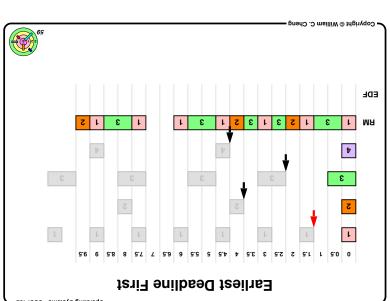


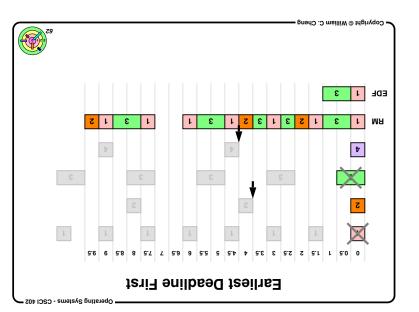


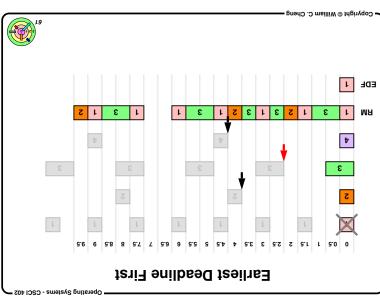


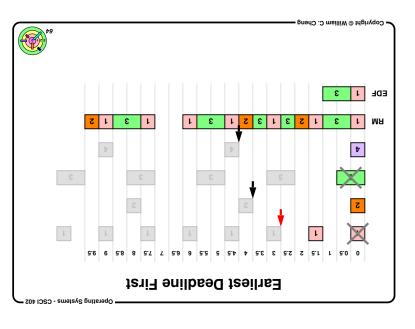


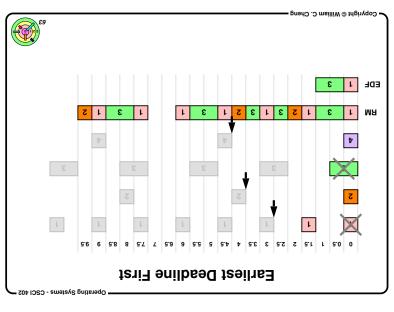


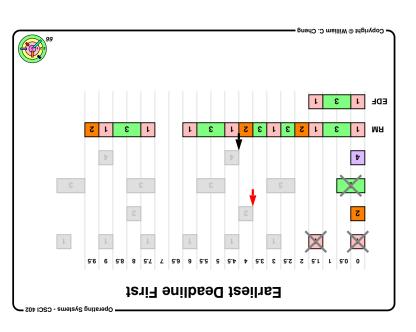


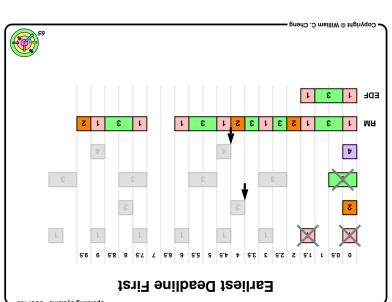


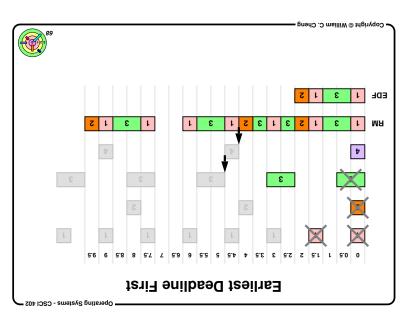


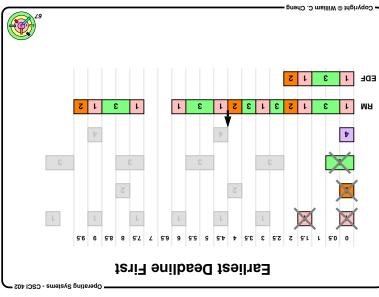


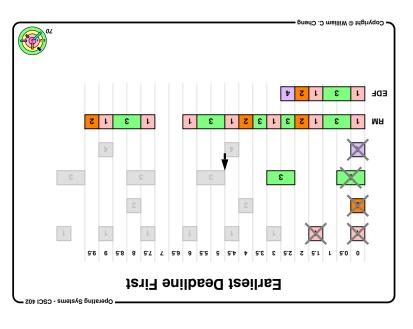


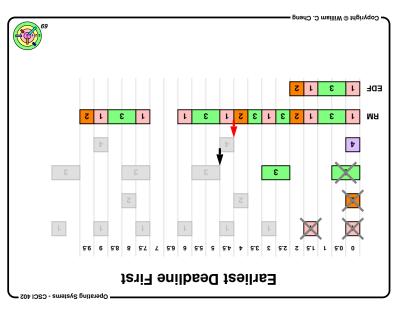


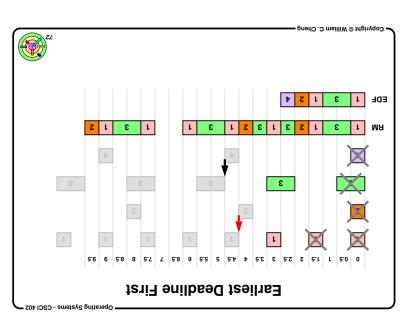


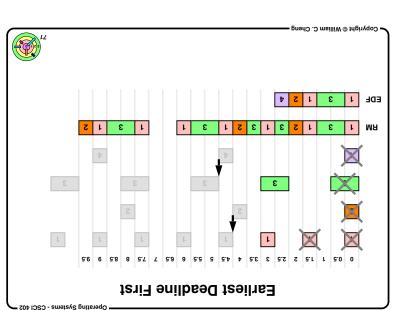


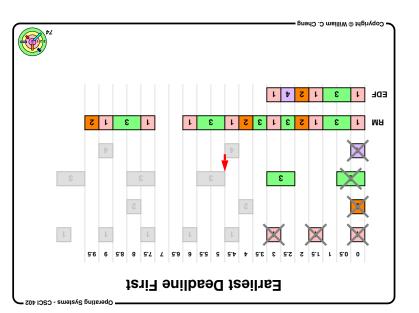


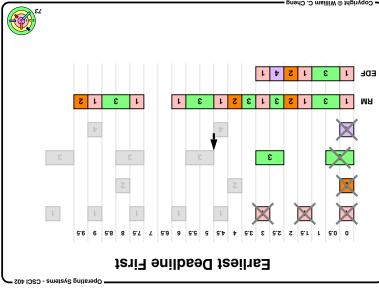


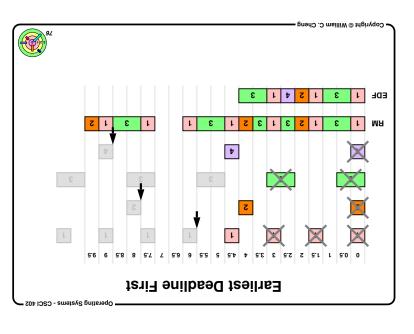


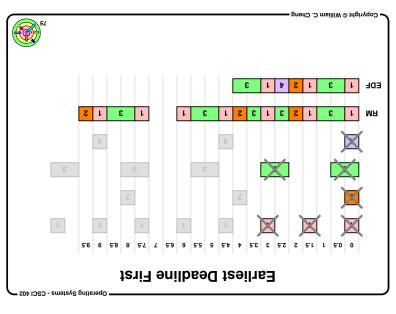


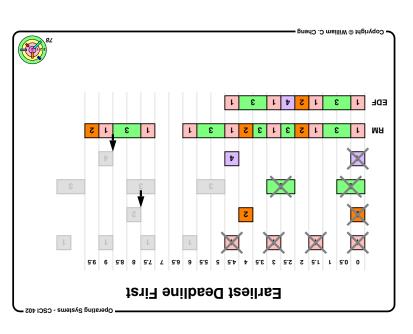


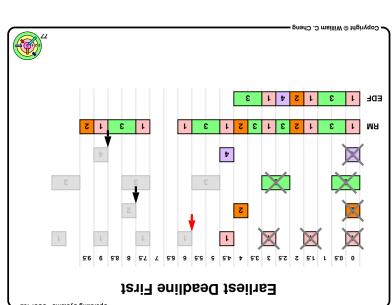


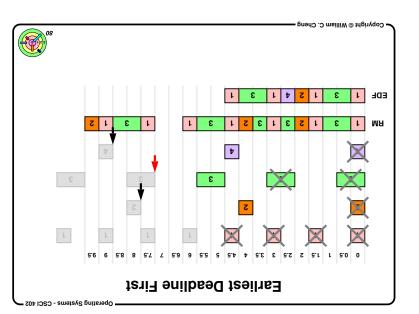


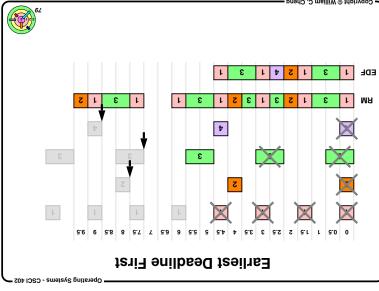


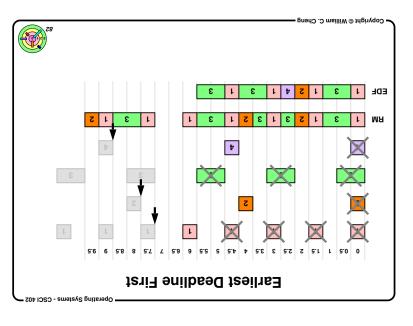


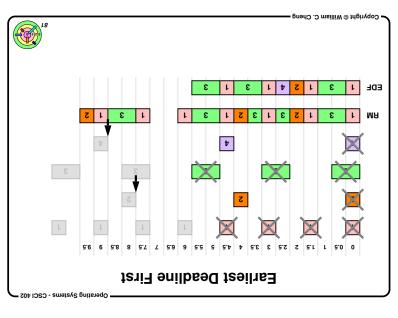


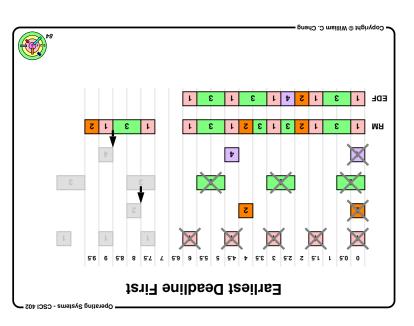


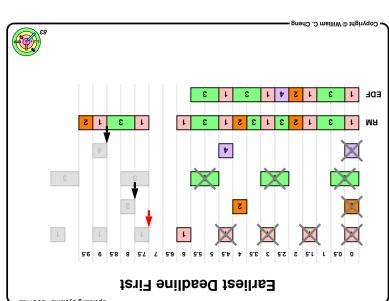


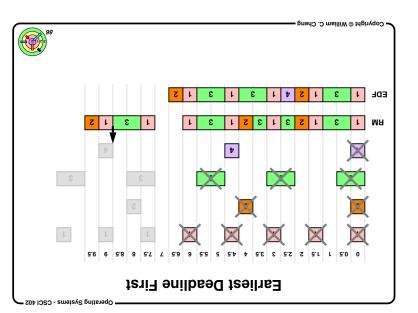


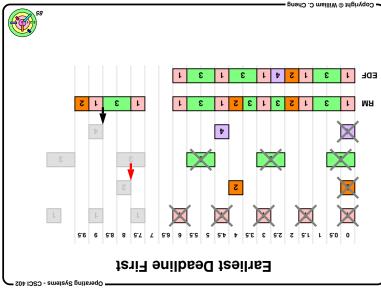


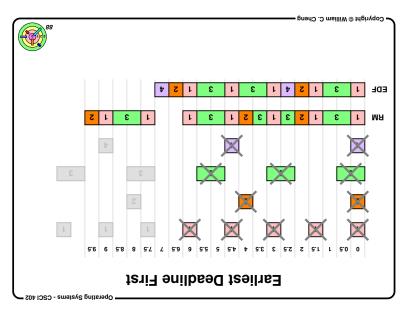


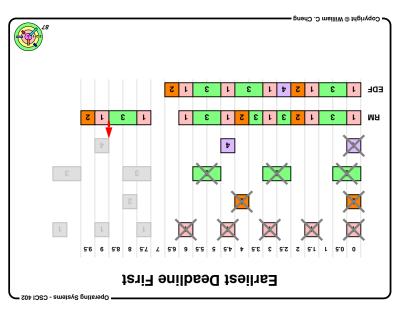


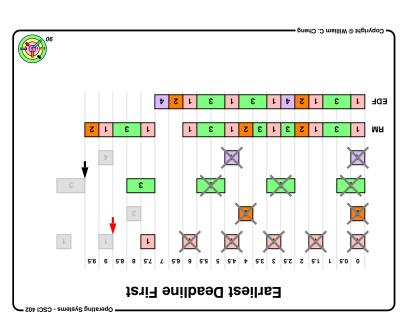


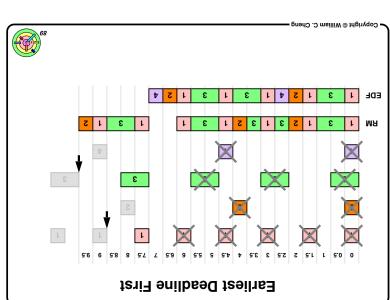


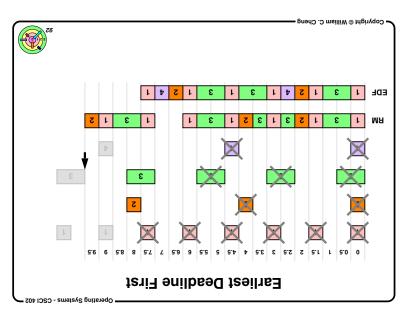


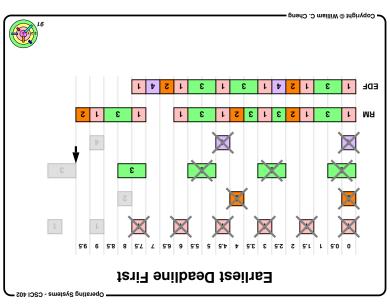


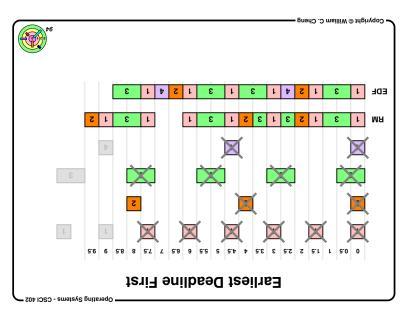


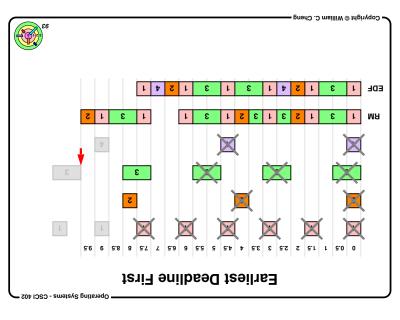


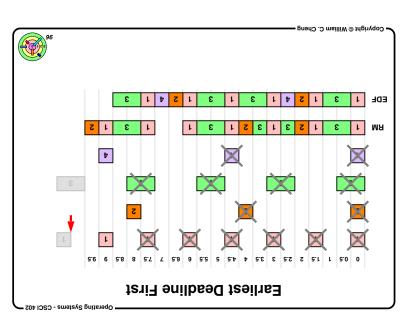


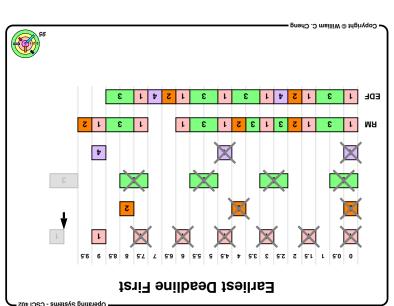


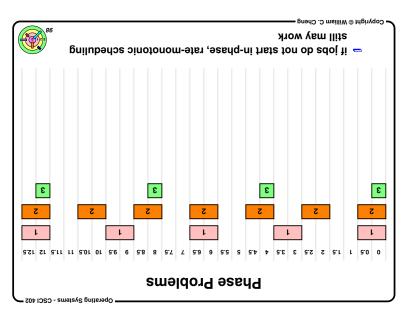


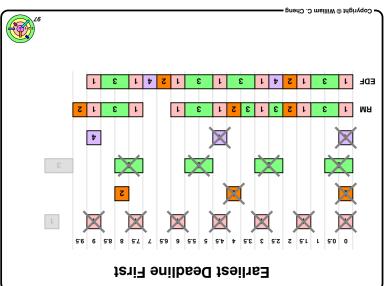


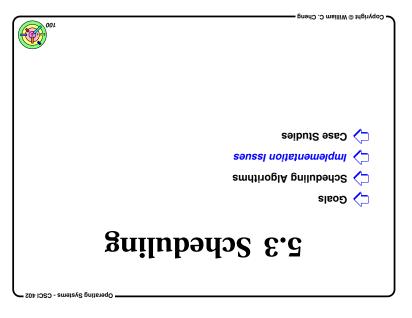


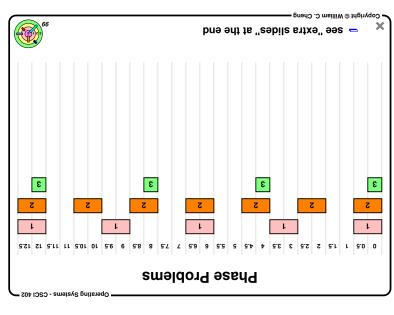


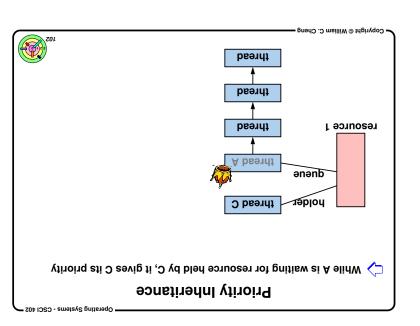


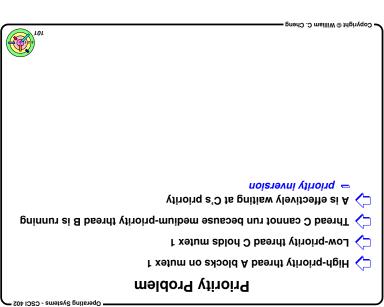


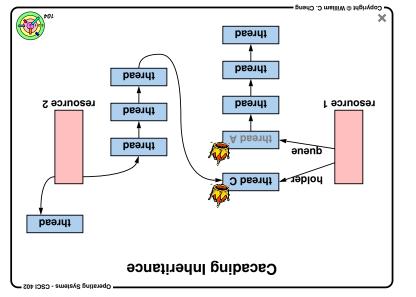


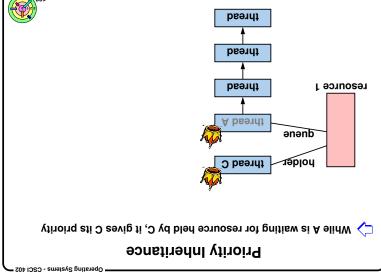




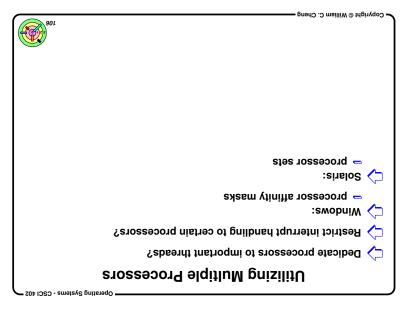


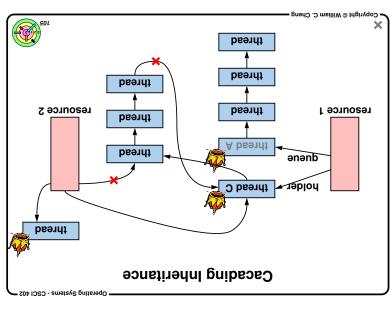


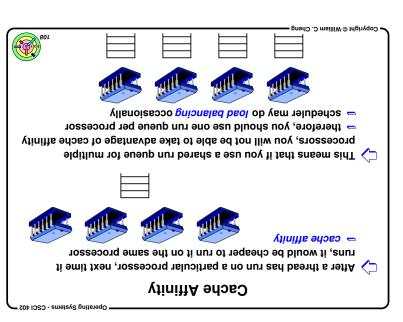


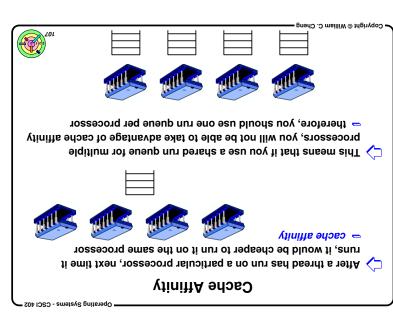


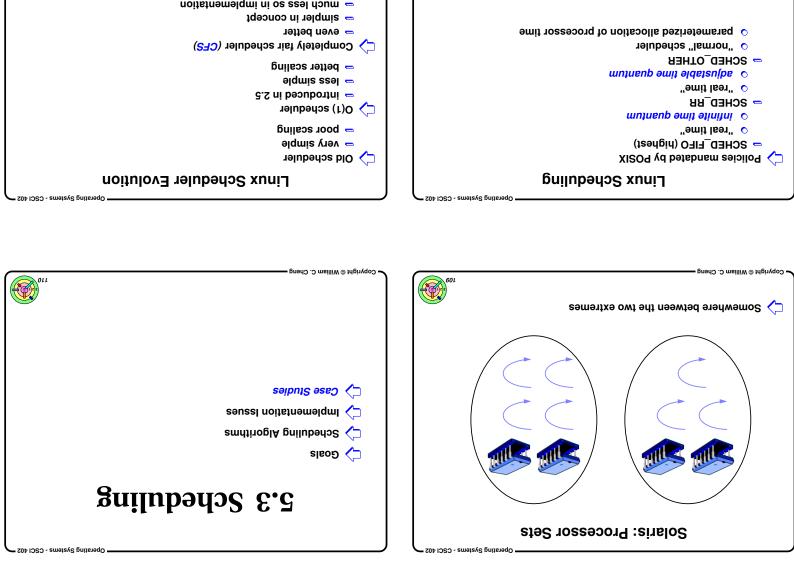
Copyright © William C. Cheng

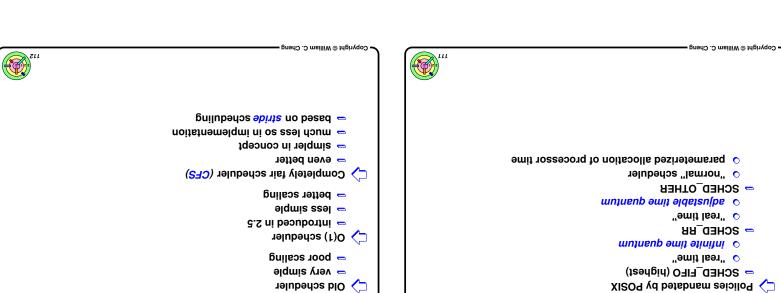


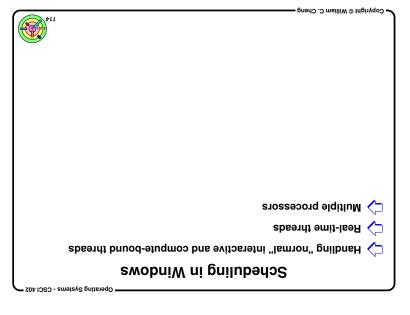


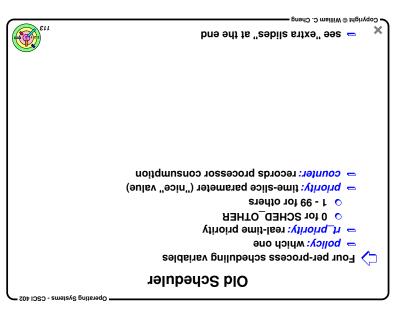


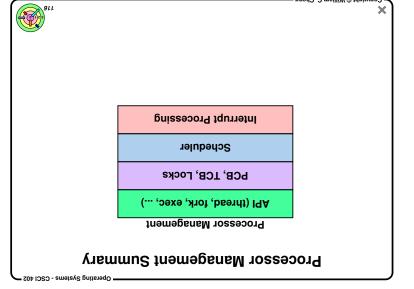


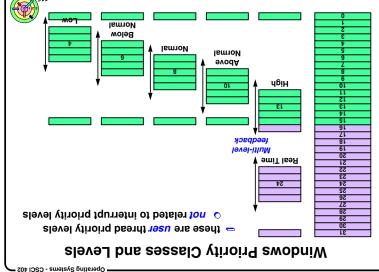


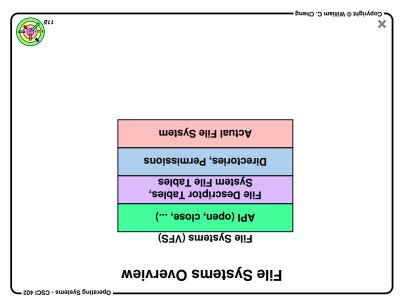


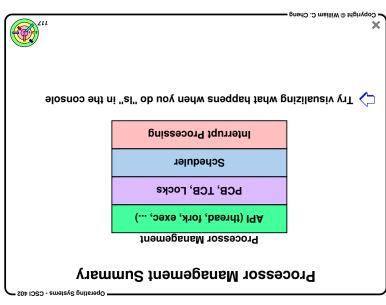


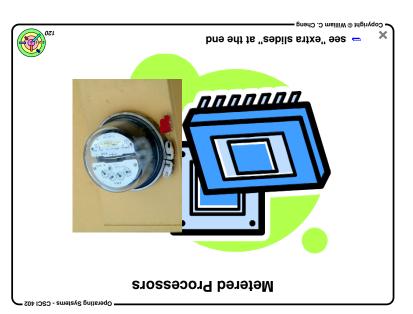


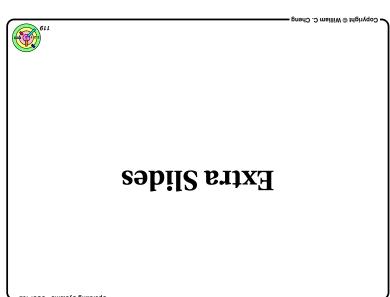


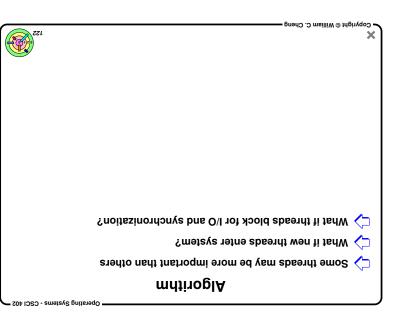


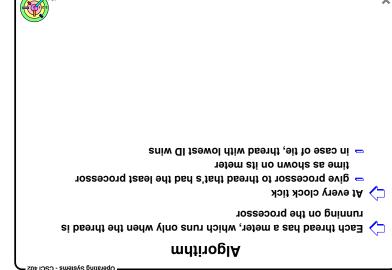


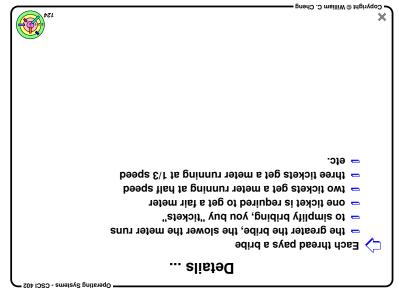




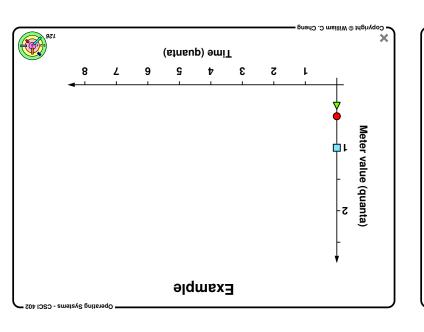


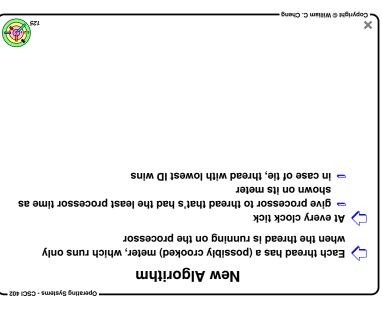


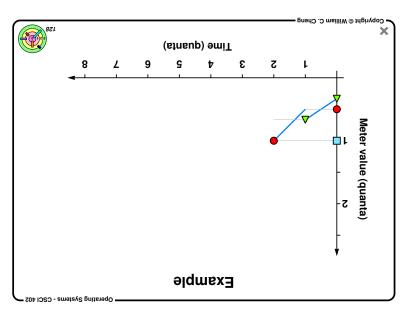


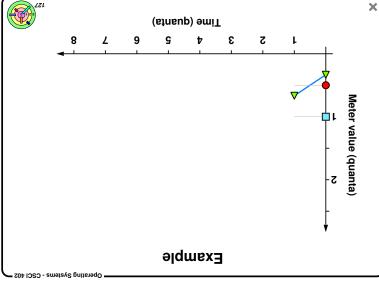


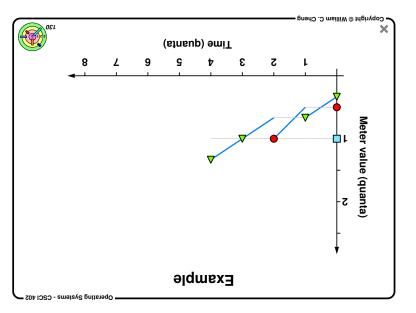


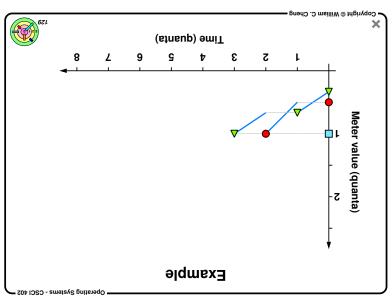


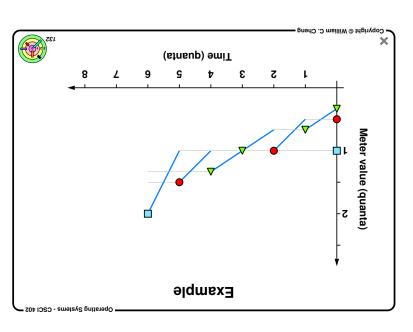


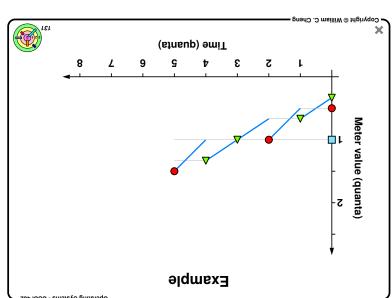


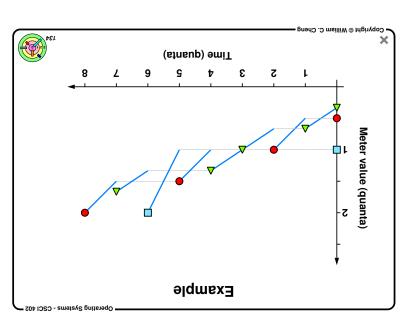


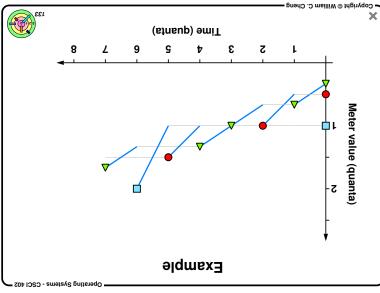


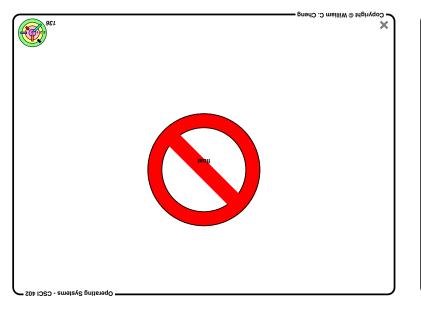


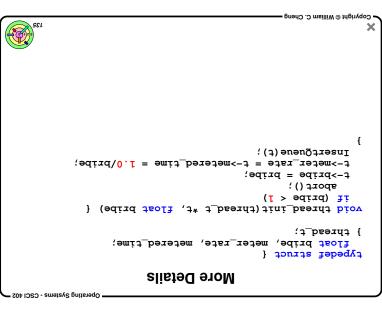












```
More Details (continued)

void OnClockTick() {

thread_t *NextThread;

CurrentThread->metered_time +=

CurrentThread->metered_time +=

CurrentThread->metered_time +=

InsertQueue (CurrentThread);

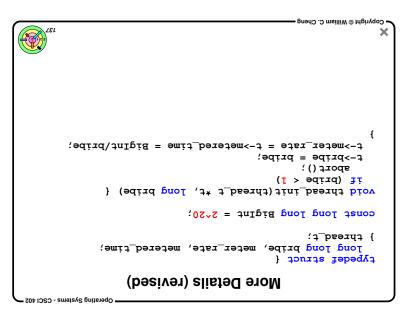
if (NextThread = PullSmallestThreadFromQueue();

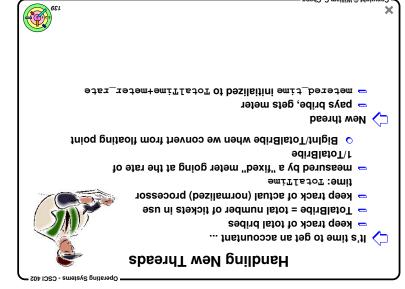
bit (NextThread != CurrentThread);

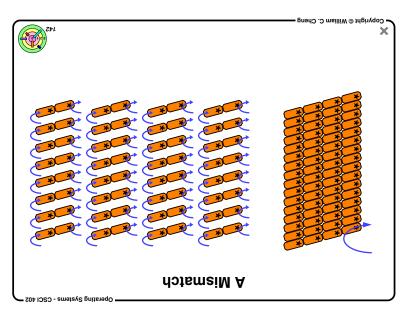
SwitchTo (NextThread);

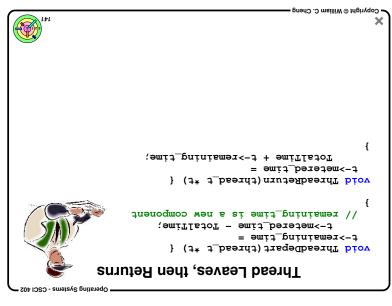
SwitchTo (NextThread);

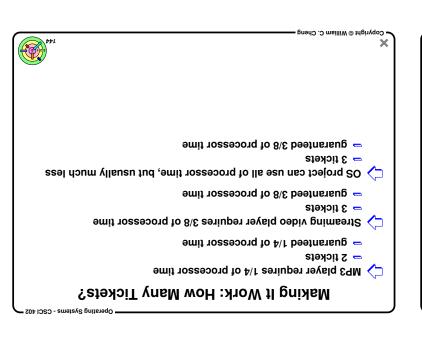
}
```

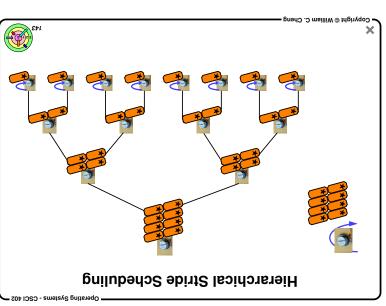


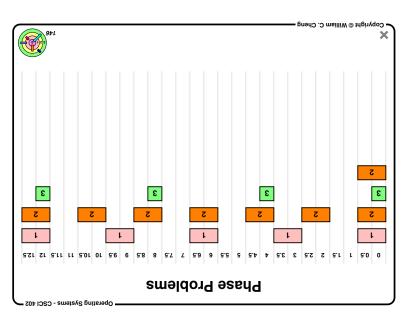


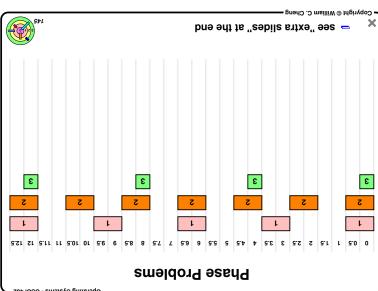


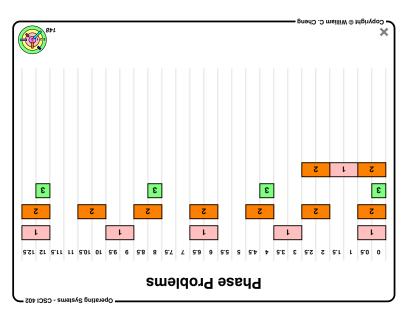


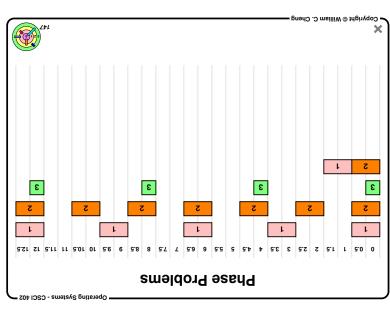


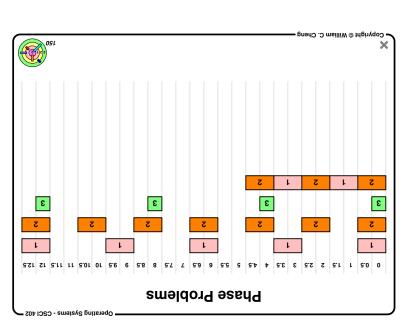


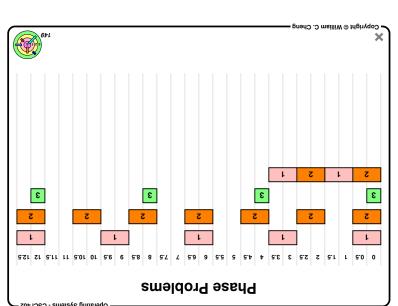


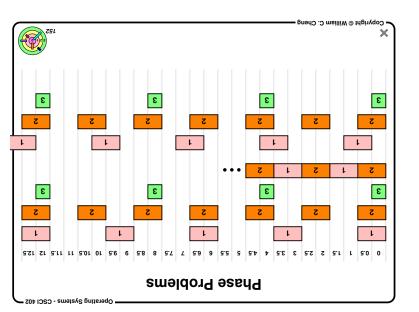


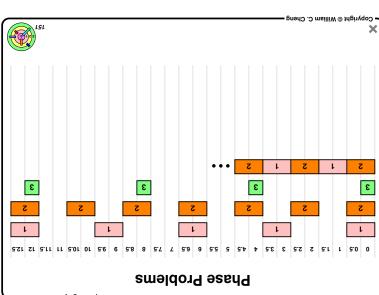


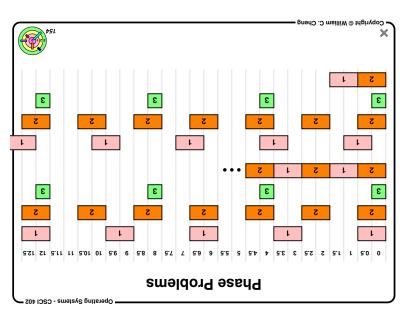


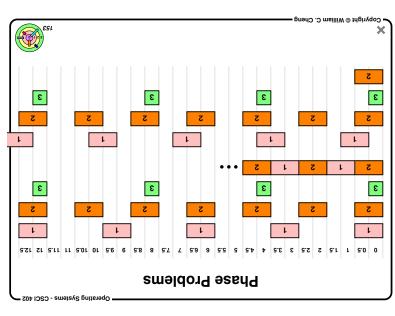


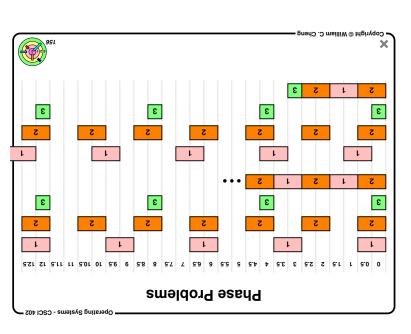


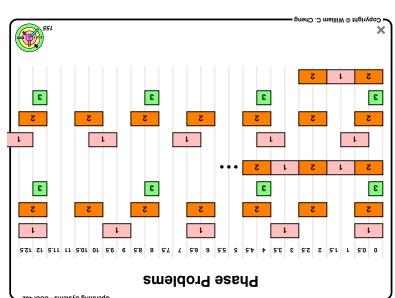


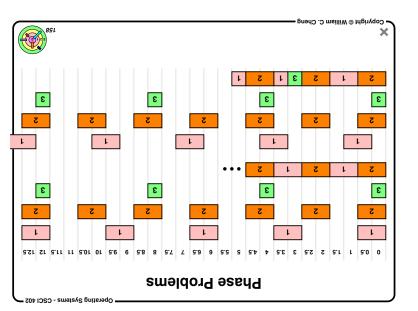


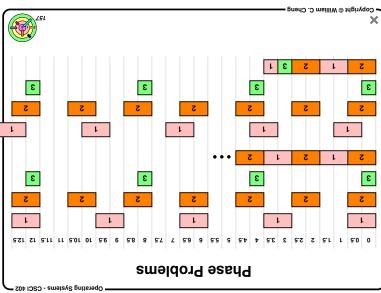


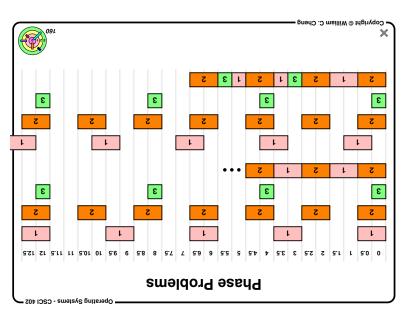


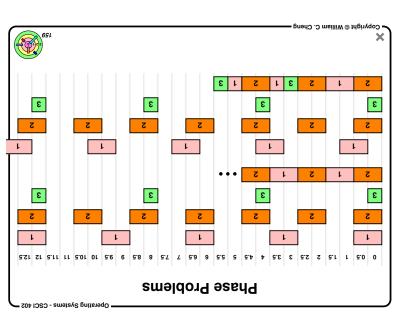


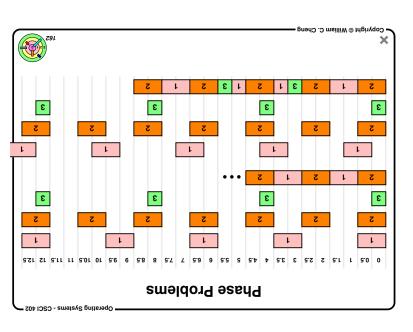


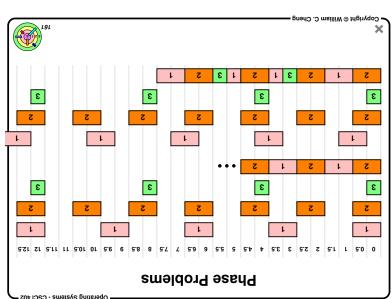


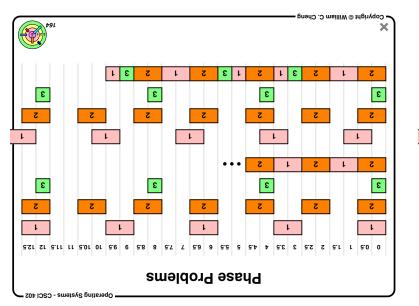


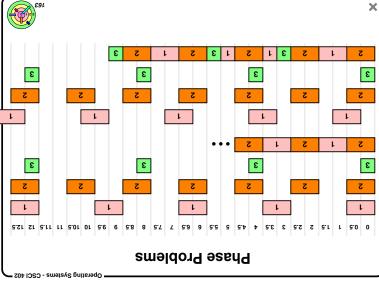


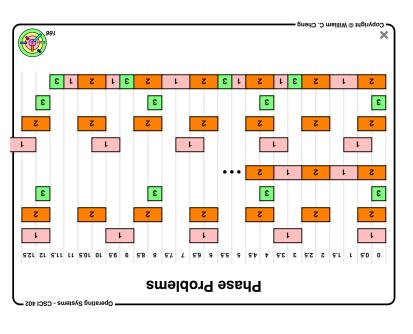


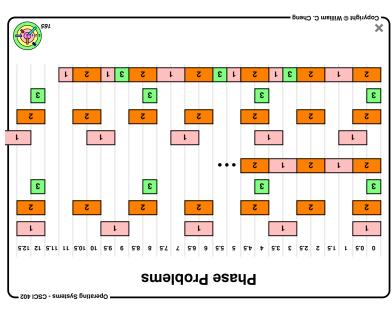


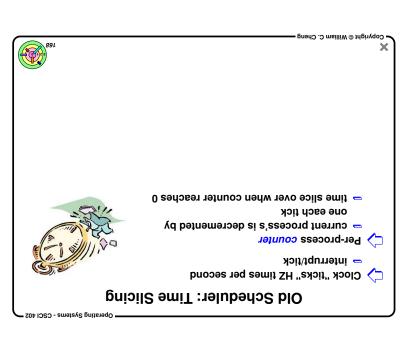


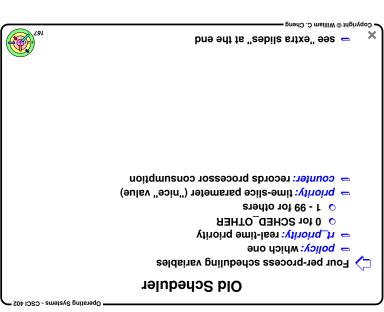


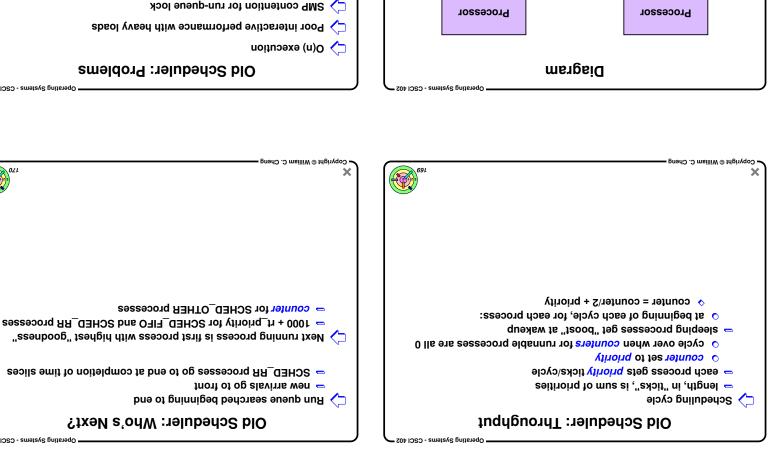


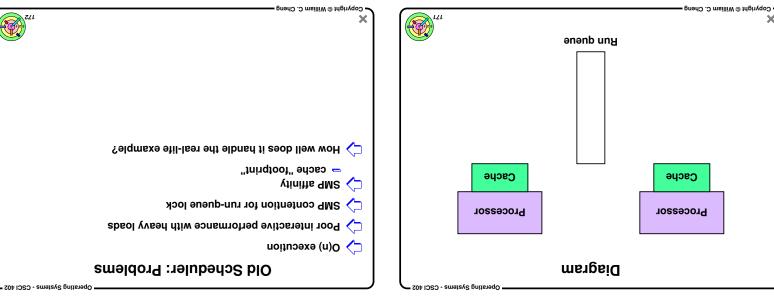


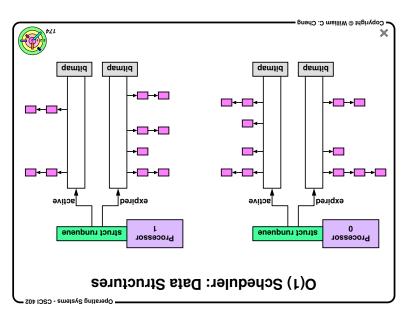


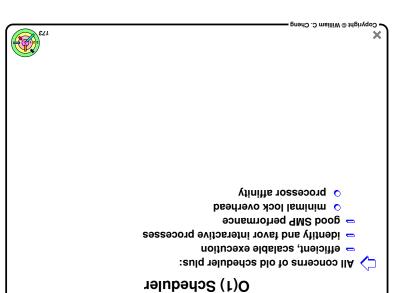


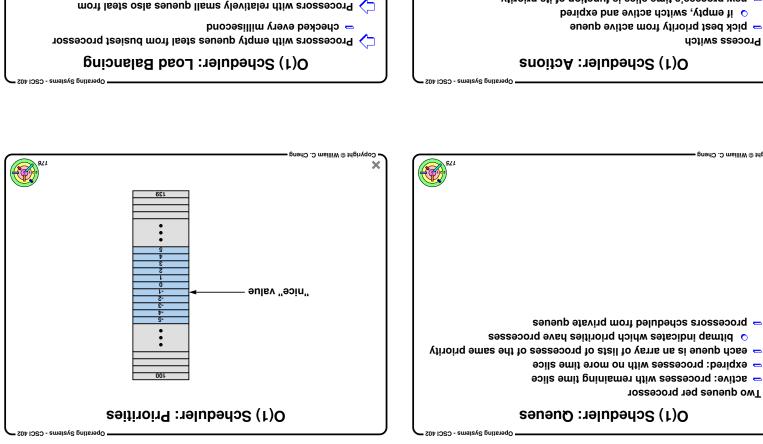






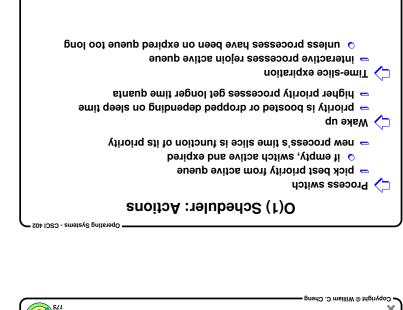






checked every 250 milliseconds

busiest processor



processors scheduled from private queues

expired: processes with no more time slice active: processes with remaining time slice

Two dueues per processor

bitmap indicates which priorities have processes

O(1) Scheduler: Queues

