






Ch 6: File Systems

Bill Cheng

<http://merlot.usc.edu/cs402-s16>




Files




Requirements

- ➡ **Permanent storage**
 - ▬ resides on disk (or alternatives)
 - ▬ survives software and hardware crashes
 - (including loss of disk?)
- ➡ **Quick, easy, and efficient**
 - ▬ satisfies needs of most applications
 - how do applications use permanent storage?



Applications

- ➡ **Software development**
 - ▬ text editors
 - ▬ linkers and loaders
 - ▬ source-code control
- ➡ **Document processing**
 - ▬ editing
 - ▬ browsing
- ➡ **Web stuff**
 - ▬ serving
 - ▬ browsing
- ➡ **Program execution**
 - ▬ paging



Needs




Directories

- convenient naming
- fast lookup



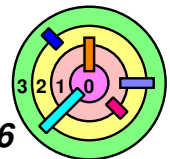
File access

- sequential is very common!
- "random access" is relatively rare



6.1 The Basics of File Systems

- ➡ *UNIX's S5FS*
- ➡ Disk Architecture
- ➡ Problems with S5FS
- ➡ Improving Performance



S5FS




A simple file system

- slow
- not terribly tolerant to crashes
- reasonably efficient in space
 - no compression




Concerns


- on-disk data structures
 - file representation
 - free space



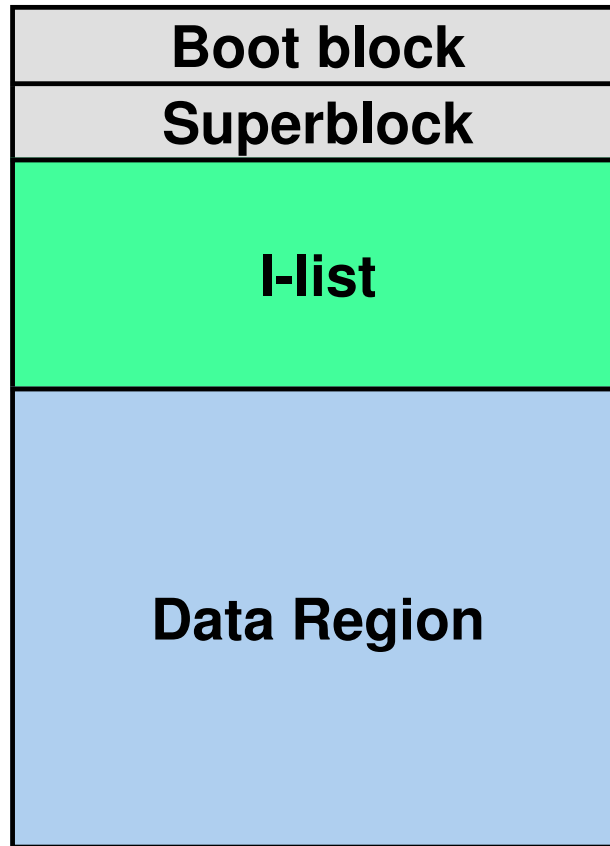
S5FS Layout




- ➡ A disk is simply an array of blocks of 1KB each (old Unix: 512B)
- ➡ A "linear view" (1-D array of blocks) of the disk




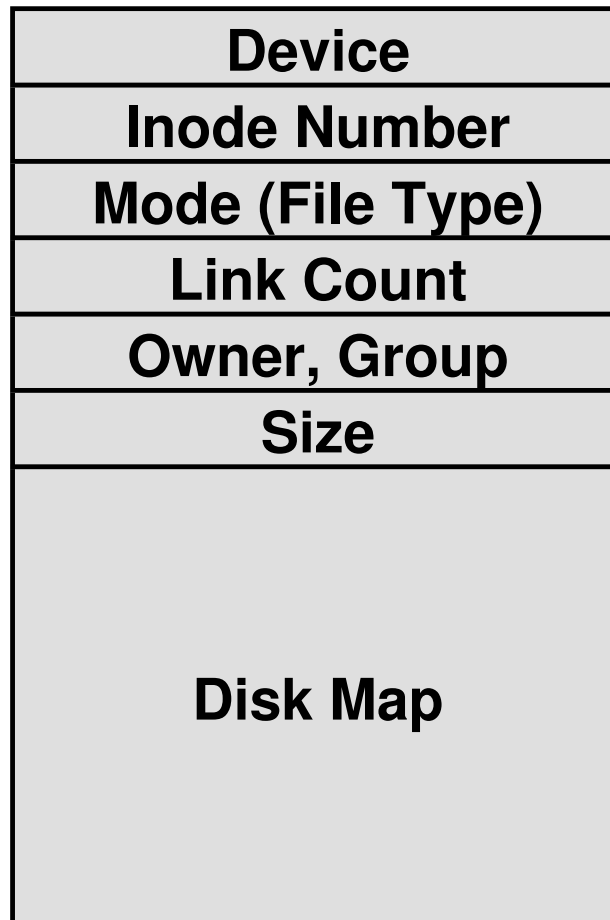
S5FS Layout



- ➡ The *superblock*
 - ➡ describes the layout of the rest of the file system
 - ➡ contains the *head* of the *free list*
- ➡ The *i-list* is an *array* of *index nodes (inodes)*
 - ➡ each representing a *file*




S5FS: Inode




Disk Map

— assuming blocksize = 1KB


0
1
2
3
4
5
6
7
8
9
10
11
12





Disk Map




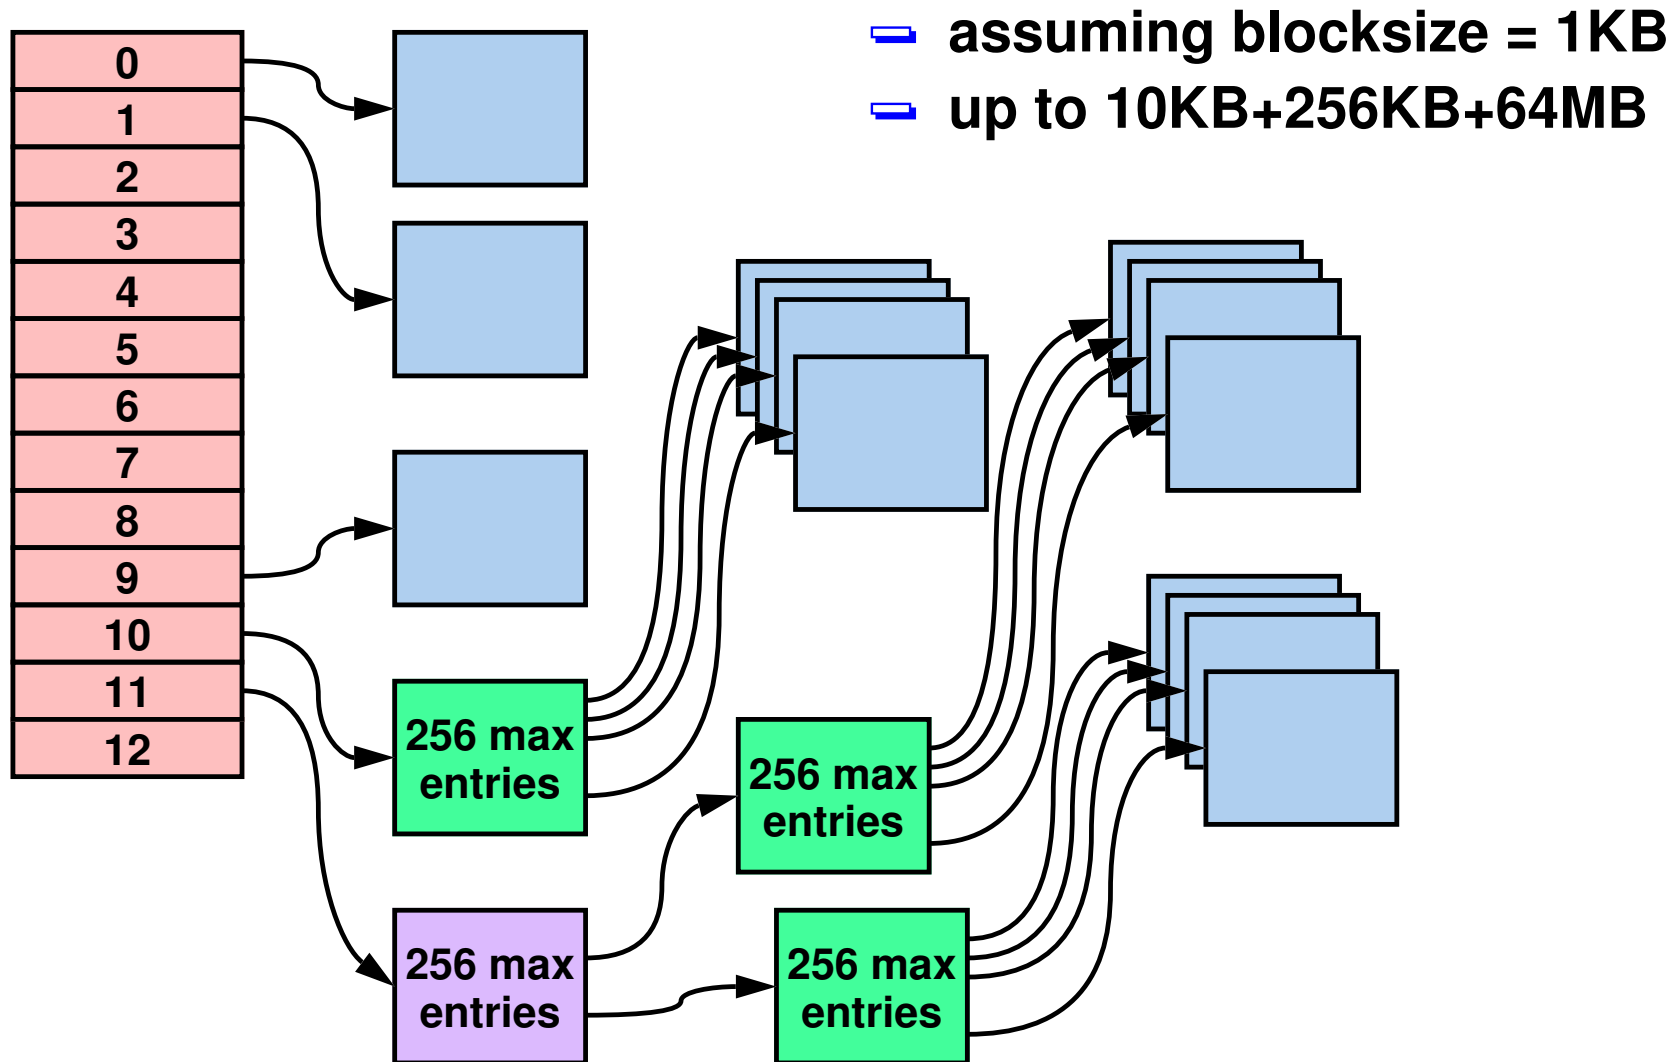
- ▢ assuming blocksize = 1KB
- ▢ up to 10KB




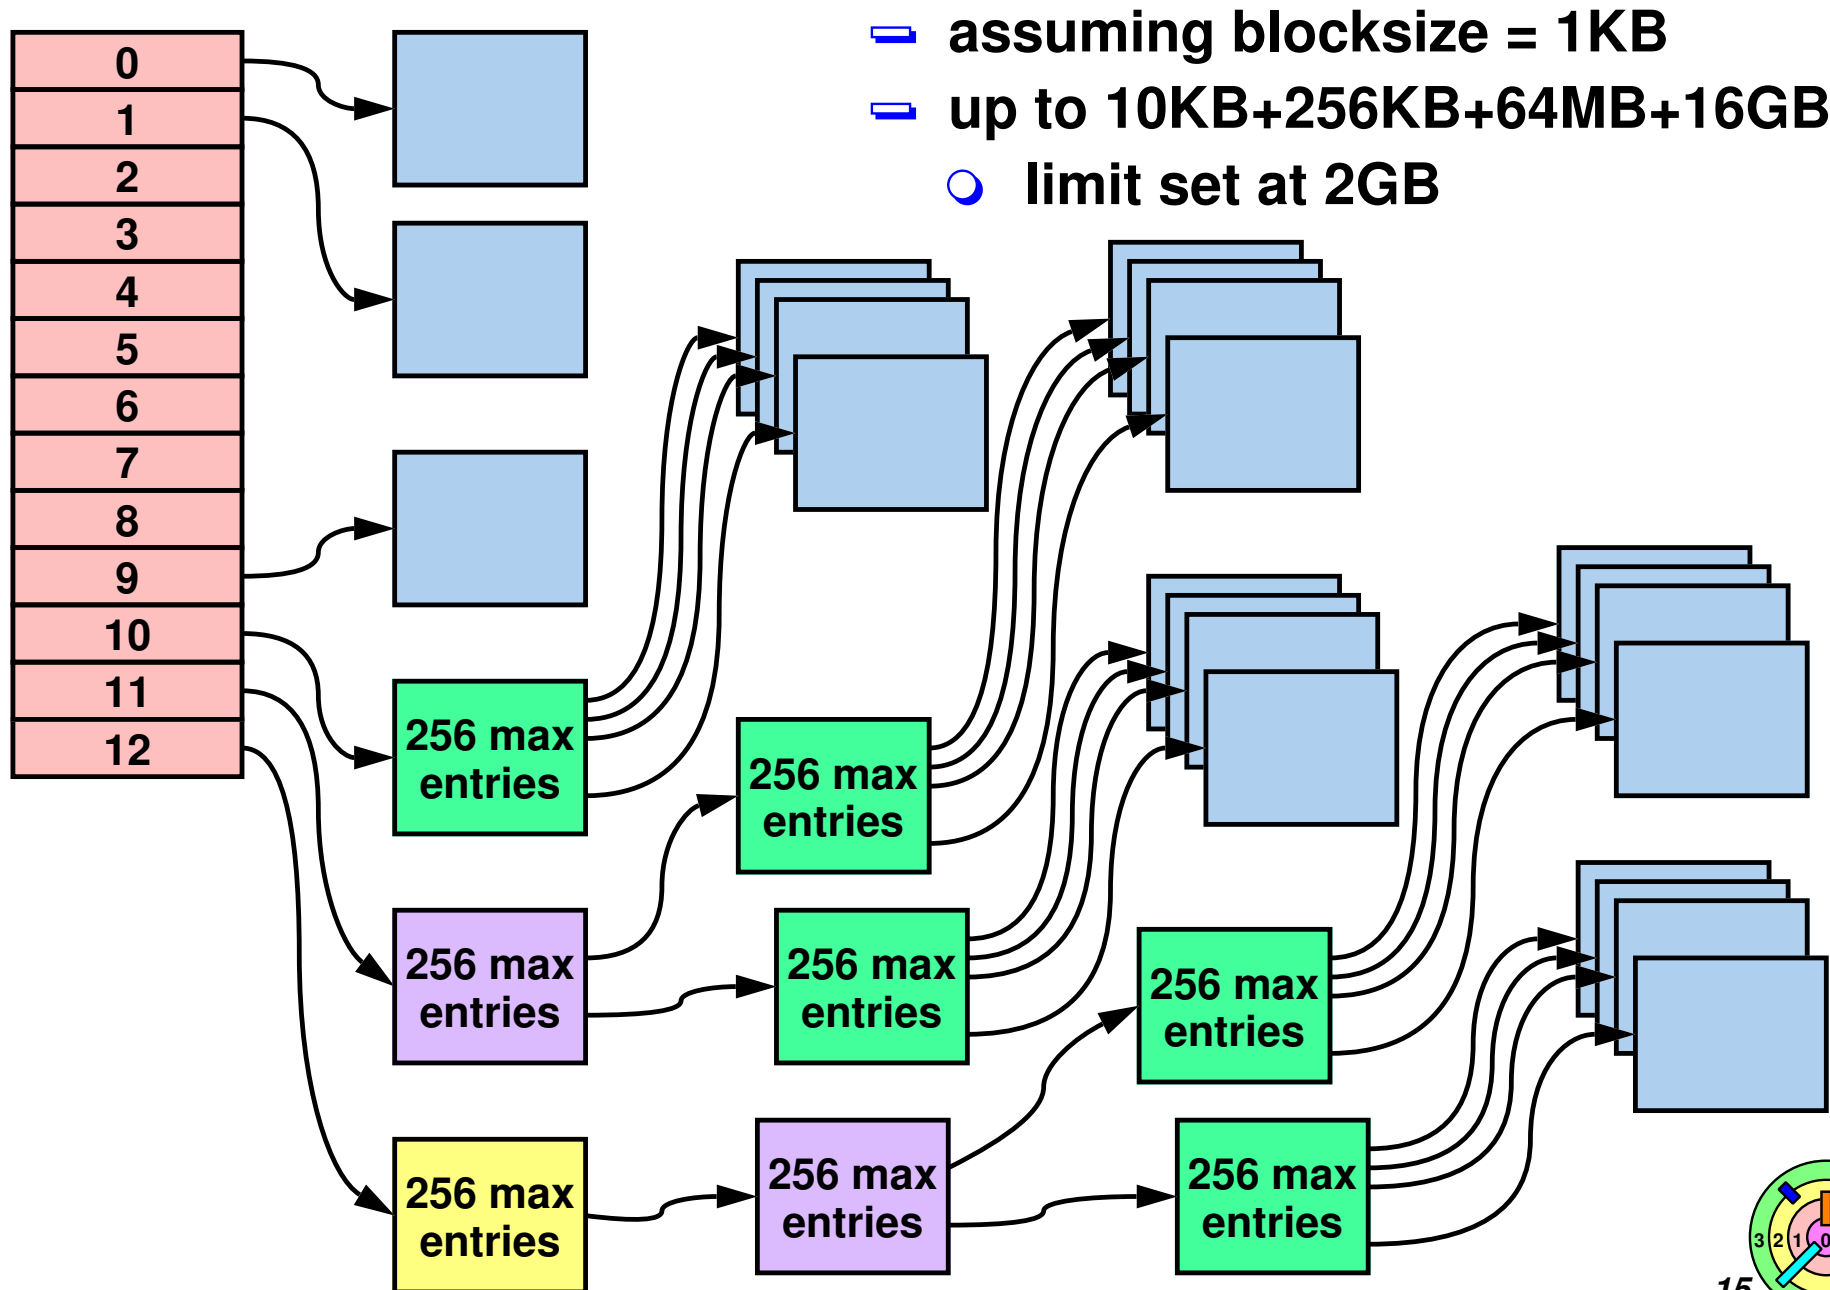
Disk Map



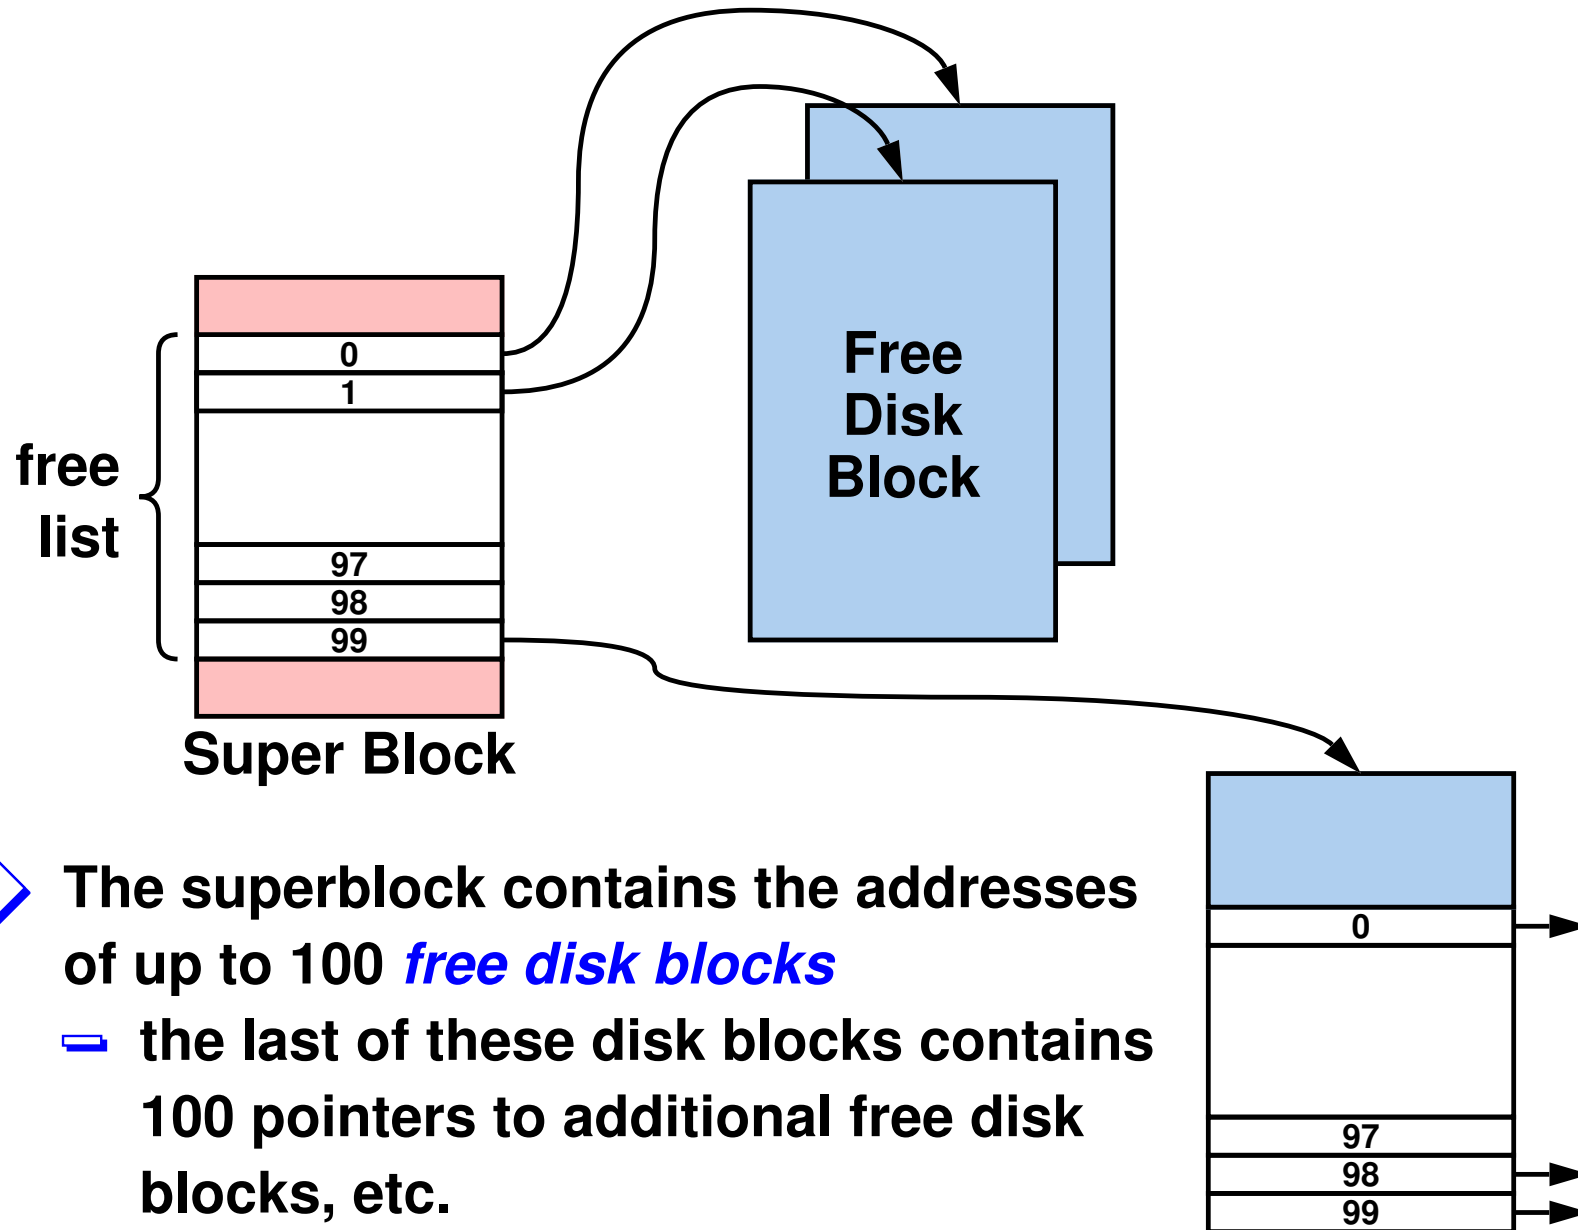
Disk Map



Disk Map




S5FS Free List



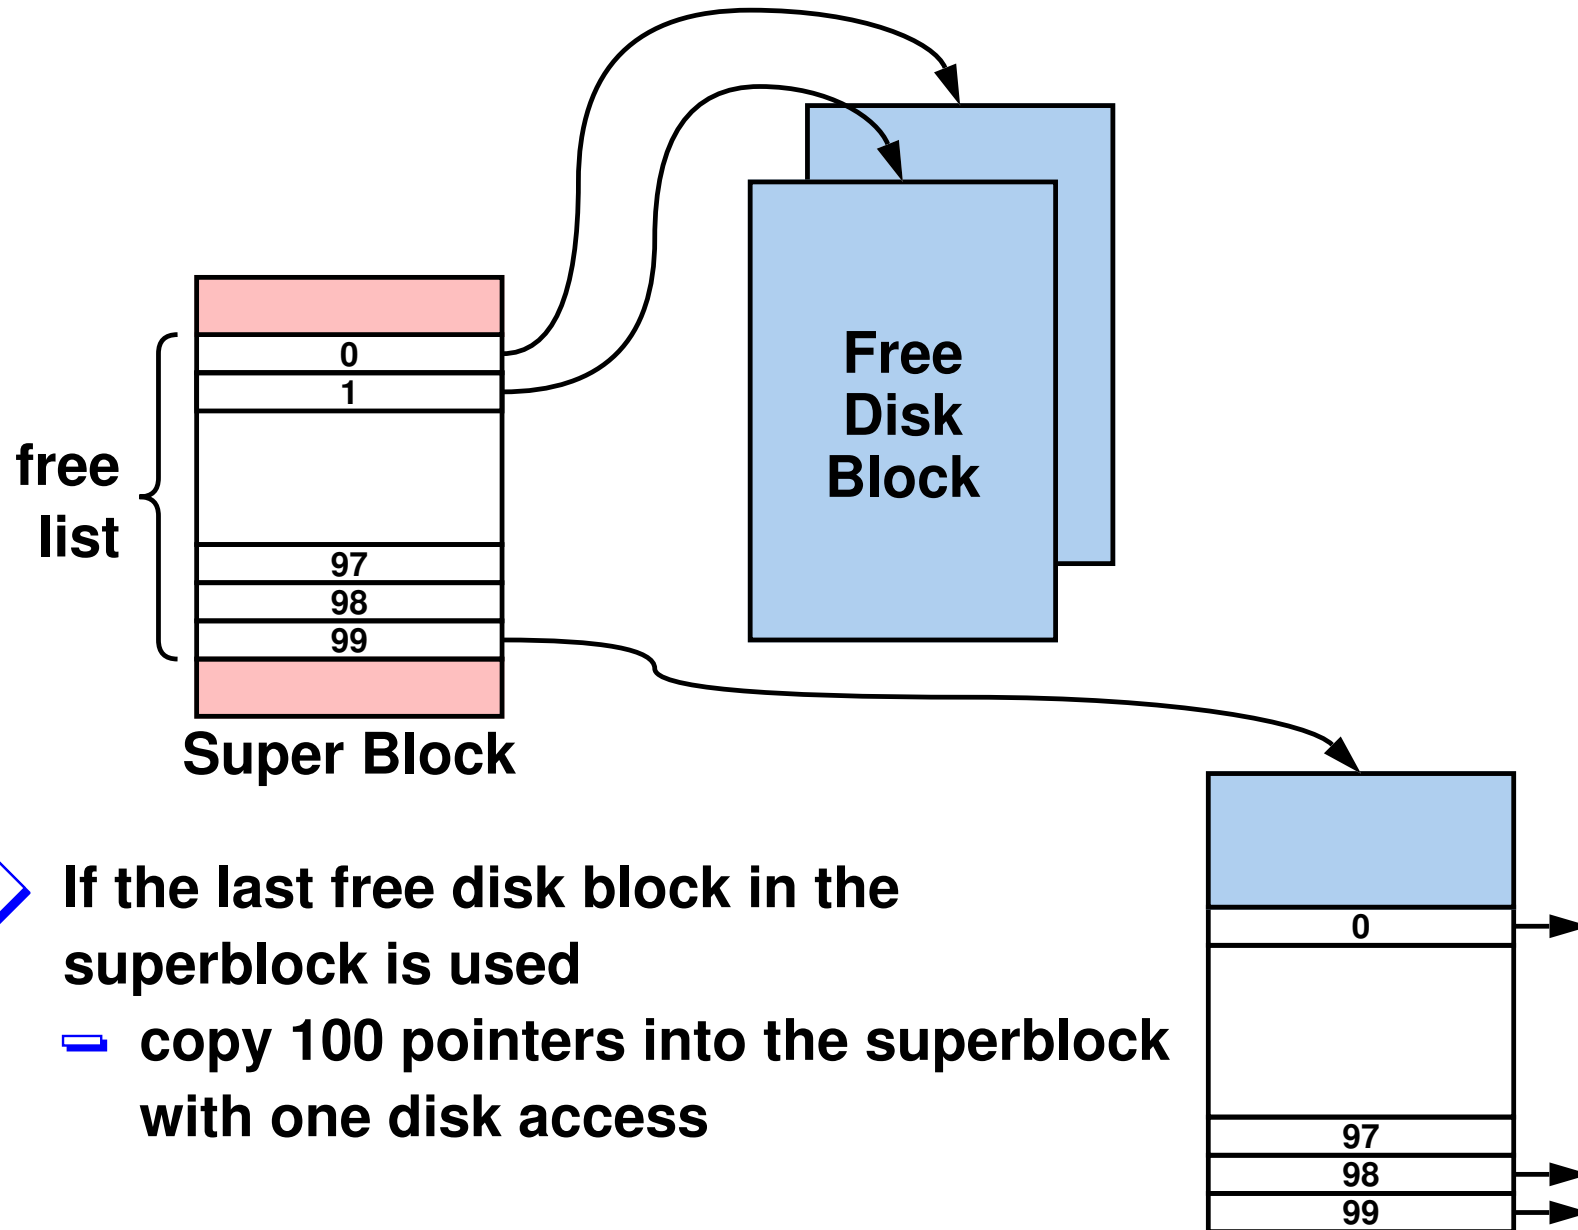
The superblock contains the addresses of up to 100 *free disk blocks*

— the last of these disk blocks contains 100 pointers to additional free disk blocks, etc.


— can find *all* the free disk blocks this way



S5FS Free List




S5FS Free List




When a disk block is freed

- that block's address is added to the list of free blocks in the superblock
- if the list is full
 - write it out and update superblock



S5FS Free Inode List




1		
2		
3		
4	0	
5		
6	0	
7		
8		
9	0	
10		
11	0	
12	0	
13	0	
14		
15		
16	0	

I-list

- ➡ Inodes (in the i-list) are marked free or not free
- ➡ no additional organization in the i-list
 - ➡ the superblock *cache*s free inodes (i.e., in the *inode cache*)


S5FS Free Inode List




1		
2		
3		
4	0	
5		
6	0	
7		
8		
9	0	
10		
11	0	
12	0	
13	0	
14		
15		
16	0	

I-list

- ➡ The *inode cache*
- to allocate an inode, simply mark it not free and remove it from the inode cache
 - to free an inode, simply mark it free and add to the inode cache if there is room




S5FS Free Inode List




1		
2		
3		
4	0	
5		
6	0	
7		
8		
9	0	
10		
11	0	
12	0	
13	0	
14		
15		
16	0	

I-list

- ➡ If the inode cache is empty
 - ▬ scan the i-list to refill it
 - to help out with the scan, the inode cache contains the index of the first free inode in the i-list
 - ◆ need to maintain this entry when necessary




S5FS Free Inode List




1		
2		
3		
4	0	
5		
6	0	
7		
8		
9	0	
10		
11	0	
12	0	
13	0	
14		
15		
16	0	

I-list

- ➡ To *create a file*
- ➡ get a *free block*
 - update *free list*
 - ➡ get a *free inode*
 - update *i-list* and *inode cache*



S5FS Free Inode List




1		
2		
3		
4	0	
5		
6	0	
7		
8		
9	0	
10		
11	0	
12	0	
13	0	
14		
15		
16	0	

I-list

- ➡ To *delete a file*
- ➡ add disk block(s) to *free list*
 - ➡ mark inode free in *i-list* and may be update *inode cache*


S5FS Summary

- ➡ In designing a file system, one tries to minimize the number of disk operations
- ▬ read vs. write
 - ▬ sequential access vs. random access
 - S5FS gives $O(1)$ number of disk operations for random access




6.1 The Basics of File Systems

- ➡ UNIX's S5FS
- ➡ *Disk Architecture*
- ➡ Problems with S5FS
- ➡ Improving Performance




Disk Architecture


- ➡ How do you get to a particular byte of data on the disk?
- what does a disk look like?




- ➡ Clearly, this looks nothing like the S5FS layout
- or any logical file system layout



Disk Architecture



➡ Smallest addressable unit is a **sector**
= disk address = (**head/surface#**, **cylinder/track#**, **sector#**)



Rhinopias Disk Drive

Rotation speed	10,000 RPM
Number of surfaces	8
Sector size	512 bytes
Sectors/track	500-1000; 750 average
Tracks/surface	100,000
Storage capacity	307.2 billion bytes
Average seek time	4 milliseconds
One-track seek time	.2 milliseconds
Maximum seek time	10 milliseconds

