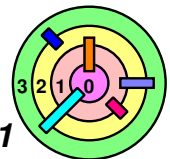



6.1 The Basics of File Systems


- ➡ UNIX's S5FS
- ➡ Disk Architecture
- ➡ *Problems with S5FS*
- ➡ Improving Performance



Disk Architecture




➡ Smallest addressable unit is a **sector**
= disk address = (**head/surface#**, **cylinder/track#**, **sector#**)



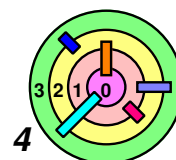
Rhinopias Disk Drive

Rotation speed	10,000 RPM
Number of surfaces	8
Sector size	512 bytes
Sectors/track	500-1000; 750 average
Tracks/surface	100,000
Storage capacity	307.2 billion bytes
Average seek time	4 milliseconds
One-track seek time	.2 milliseconds
Maximum seek time	10 milliseconds




S5FS on Rhinopias (A Marketing Disaster ...)

- ➡ Rhinopias's maximum transfer speed?
 - 63.9 MB/sec
- ➡ S5FS's average speed on Rhinopias?
 - average seek time:
 - < 4 milliseconds (say 2)
 - average rotational latency:
 - ~3 milliseconds
 - per-sector transfer time:
 - negligible
 - time/sector: 5 milliseconds
 - effective transfer speed: 102.4 KB/sec (.16% of maximum)
- ➡ In general, we have:
 - ***access time = seek time + rotational latency + data transfer time***
 - some people would use the term "response time" to mean "access time"



6.1 The Basics of File Systems

- ➡ UNIX's S5FS
- ➡ Disk Architecture
- ➡ Problems with S5FS
- ➡ *Improving Performance*



What to Do About It?




Hardware

- employ pre-fetch buffer
 - filled by hardware with what's underneath head
 - helps reads a bit; doesn't help writes




Software

- better on-disk data structures
 - increase block size
 - minimize seek time
 - reduce rotational latency





FFS



- ➡ Better on-disk organization
- ➡ Longer component names in directories
- ➡ Retains disk map of S5FS



Larger Block Size




Larger Block Size





The Down Side ...

**Smaller
Block Size**





The Down Side ...

**Smaller
Block Size**




The Down Side ...


**Smaller
Block Size**



The Down Side ...




— internal fragmentation





The Down Side ...

**Smaller
Block Size**





Wasted Space

**Larger
Block Size**



The Down Side ...


**Smaller
Block Size**




Wasted Space

= even worse internal fragmentation

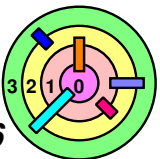

**Larger
Block Size**




Wasted Space




Two Block Sizes ...



Two Block Sizes ...




- e.g., 16KB blocks and 1KB fragments
- best of both worlds





Rules

- ➡ **File-system blocks may be split into fragments that can be independently assigned to files**
 - **fragments assigned to a file must be contiguous and in order**
- ➡ **The number of fragments per block (1, 2, 4, or 8) is fixed for each file system**
- ➡ **Allocation in fragments may only be done on what would be the last block of a file, and only for small files**





Use of Fragments (1)



 File A

 File B

Use of Fragments (2)




File A






File B

— A can grow by 2 segments




Use of Fragments (3)




 File A

 File B




Minimizing Seek Time

- ➡ Keep related things close to one another
- ➡ Separate unrelated things




Cylinder Groups




⇒ recall that *seeking* to the *next cylinder/track* is much *faster*

Minimizing Seek Time




The practice (heuristics):


- attempt to put new inodes in the same cylinder group as their directories
- put inodes for new directories in cylinder groups with "lots" of free space
- put the beginning of a file (first 10KB, i.e., direct blocks) in the inode's cylinder group
- put additional portions of the file (each 2MB) in cylinder groups with "lots" of free space




Locality Of File Access




— if access "d.c", likely to access "e.c"



Locality Of File Access

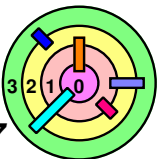


— if access "d.c", likely to access "e.c"





How Are We Doing? (Part 1)

- ➡ Configure Rhinopias with 20 cylinders per group
- 2-MB file fits entirely within one cylinder group
 - average seek time within cylinder group is $\sim .3$ milliseconds
 - average rotational delay still 3 milliseconds
 - .12 milliseconds required for disk head to pass over 8KB block
 - 3.42 milliseconds for each block
 - 2.4 million bytes/second average effective transfer speed
 - *factor of 20 improvement*
 - 3.7% of maximum possible




Minimizing Latency




Numbers


- ➡ **Rhinopias spins at 10,000 RPM**
 - **6 milliseconds/revolution**
- ➡ **100 microseconds required to service disk-completion interrupt and start next operation**
 - **typical of early 1980s**
- ➡ **Each block takes 120 microseconds to traverse disk head**
- ➡ **Reading successive blocks is expensive!**



Minimizing Latency



Block interleaving




How're We Doing Now? (Part 2)




Time to read successive blocks (two-way interleaving):

- after request for second block is issued, must wait 20 microseconds for the beginning of the block to rotate under disk head
- *factor of 15 improvement*
 - together with other improvements, overall, a factor of 300 improvement




How're We Doing Now? (Altogether)

- ➡ Same setup as before
- 2-MB file within one cylinder group
 - actually fits in one cylinder
 - block interleaving employed: every other block is skipped
 - .3-millisecond seek to that cylinder
 - 3-millisecond rotational delay for first block
 - 50 blocks/track, but 25 read in each revolution
 - 10.24 revolutions required to read all of file
 - 32.4 MB/second (50% of maximum possible)




Further Improvements?

- ➡ **S5FS: 0.16% of capacity**
- ➡ **FFS without block interleaving**
 - ▬ **factor of 20 improvement**
 - ▬ **reached 3.8% of capacity**
- ➡ **FFS with block interleaving**
 - ▬ **another factor of 15 improvement**
 - ▬ **reached 50% of capacity**
- ➡ **What next?**




Larger Transfer Units

- ➡ **Allocate in whole tracks or cylinders**
 - too much wasted space
- ➡ **Allocate in blocks, but group them together**
 - transfer many at once



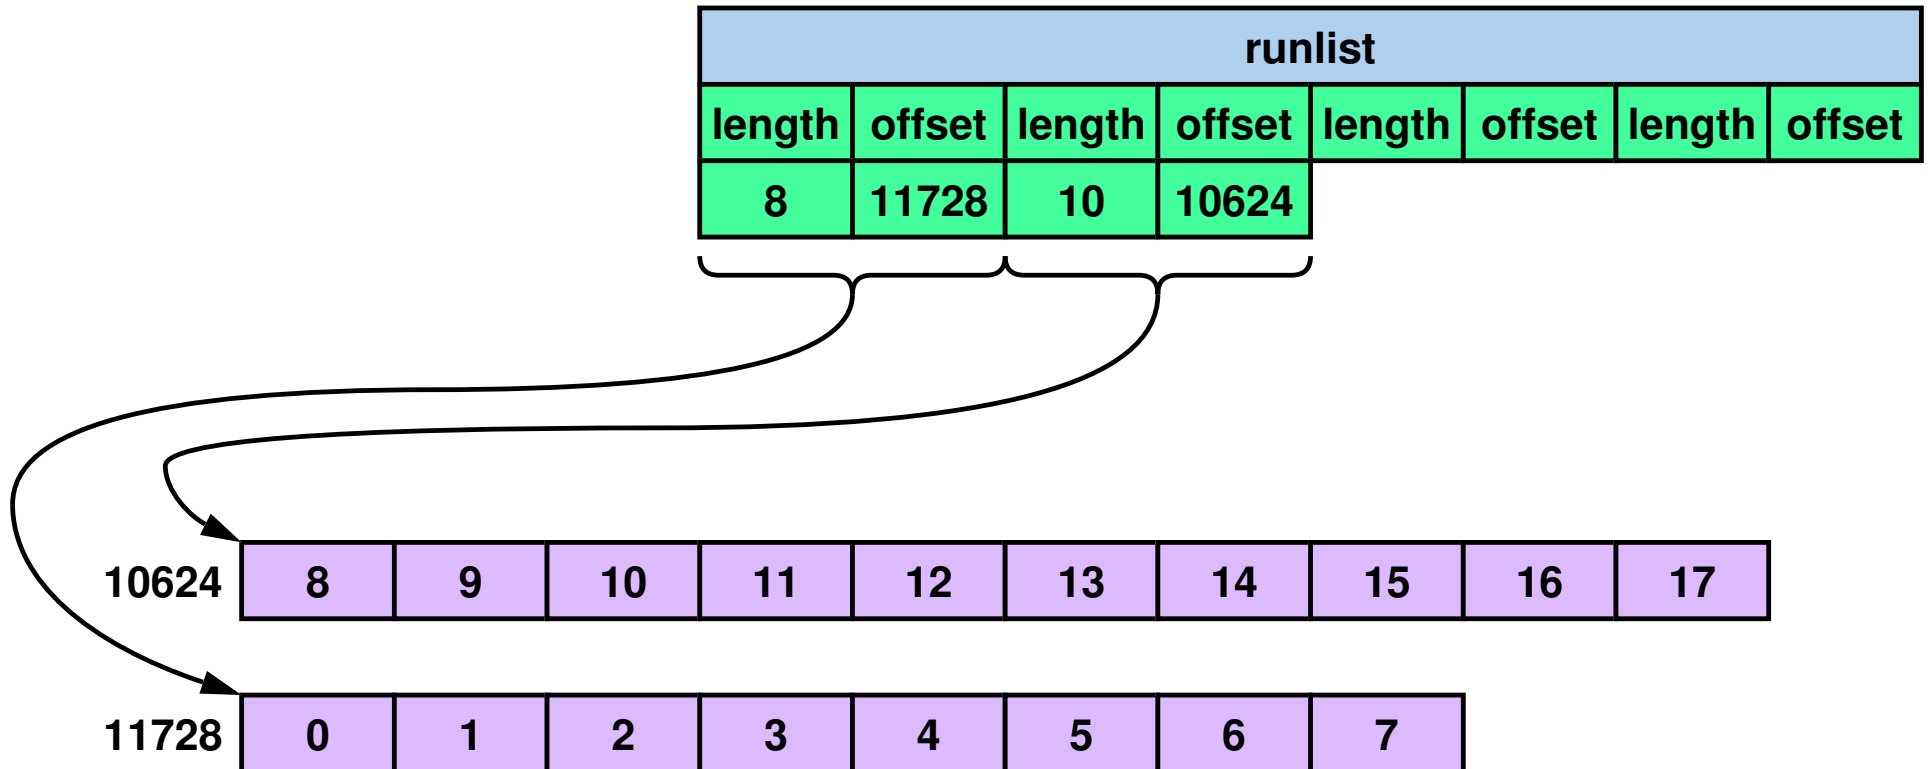
Block Clustering

- ➡ **Allocate space in blocks, eight at a time**
- ➡ **Linux's Ext2 (an FFS clone):**
 - ▬ **allocate eight blocks at a time**
 - ▬ **extra space is available to other files if there is a shortage of space**
- ➡ **FFS on Solaris (~1990)**
 - ▬ **delay disk-space allocation until:**
 - **8 blocks are ready to be written**
 - **or the file is closed**




Extents

➡ Windows



Problems with Extents

- ➡ Could result in highly fragmented disk space
 - lots of small areas of free space
 - external fragmentation
 - solution: use a *defragmenter* to *coalesce* free space
- ➡ Random access
 - *linear search* through a long list of extents
 - solution: *multiple levels*
 - usually two levels



Extents in NTFS

Top-level run list							
length	offset	length	offset	length	offset	length	offset
50000	1076	10000	9738	36000	5192	2200	14024

9738


Runlist							
length	offset	length	offset	length	offset	length	offset
8	11728	10	10624				

10624




50008	50009	50010	50011	50012	50013	50014	50015	50016	50017
-------	-------	-------	-------	-------	-------	-------	-------	-------	-------

11728

50000	50001	50002	50003	50004	50005	50006	50007
-------	-------	-------	-------	-------	-------	-------	-------




Are We There Yet?




Recall What A File Look Like in S5FS & FFS

File Data:
(e.g., PDF)





File On-disk Representation:



S5FS:





FFS:





CPU, Memory, Disk Speeds Over Time

Capacity/Speed




CPU, Memory, Disk Speeds Over Time

Capacity/Speed

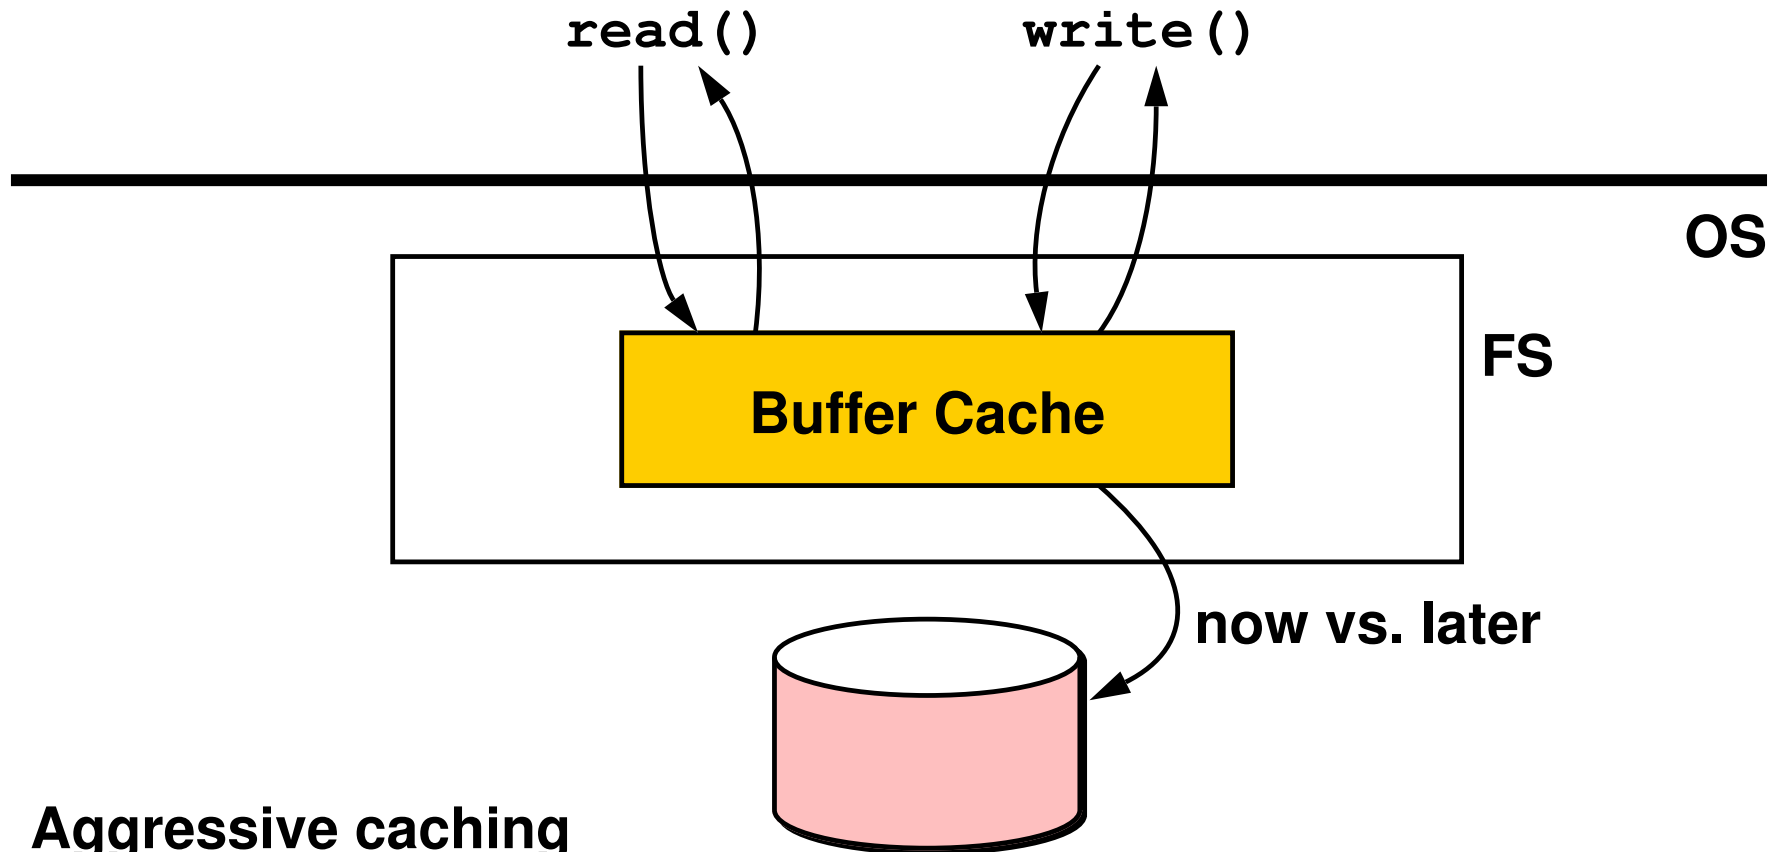


A Different Approach

- ➡ We have lots of primary memory
 - enough to cache all commonly used files
- ➡ Read time from disk doesn't matter
- ➡ Time for writes does matter




The Buffer Cache

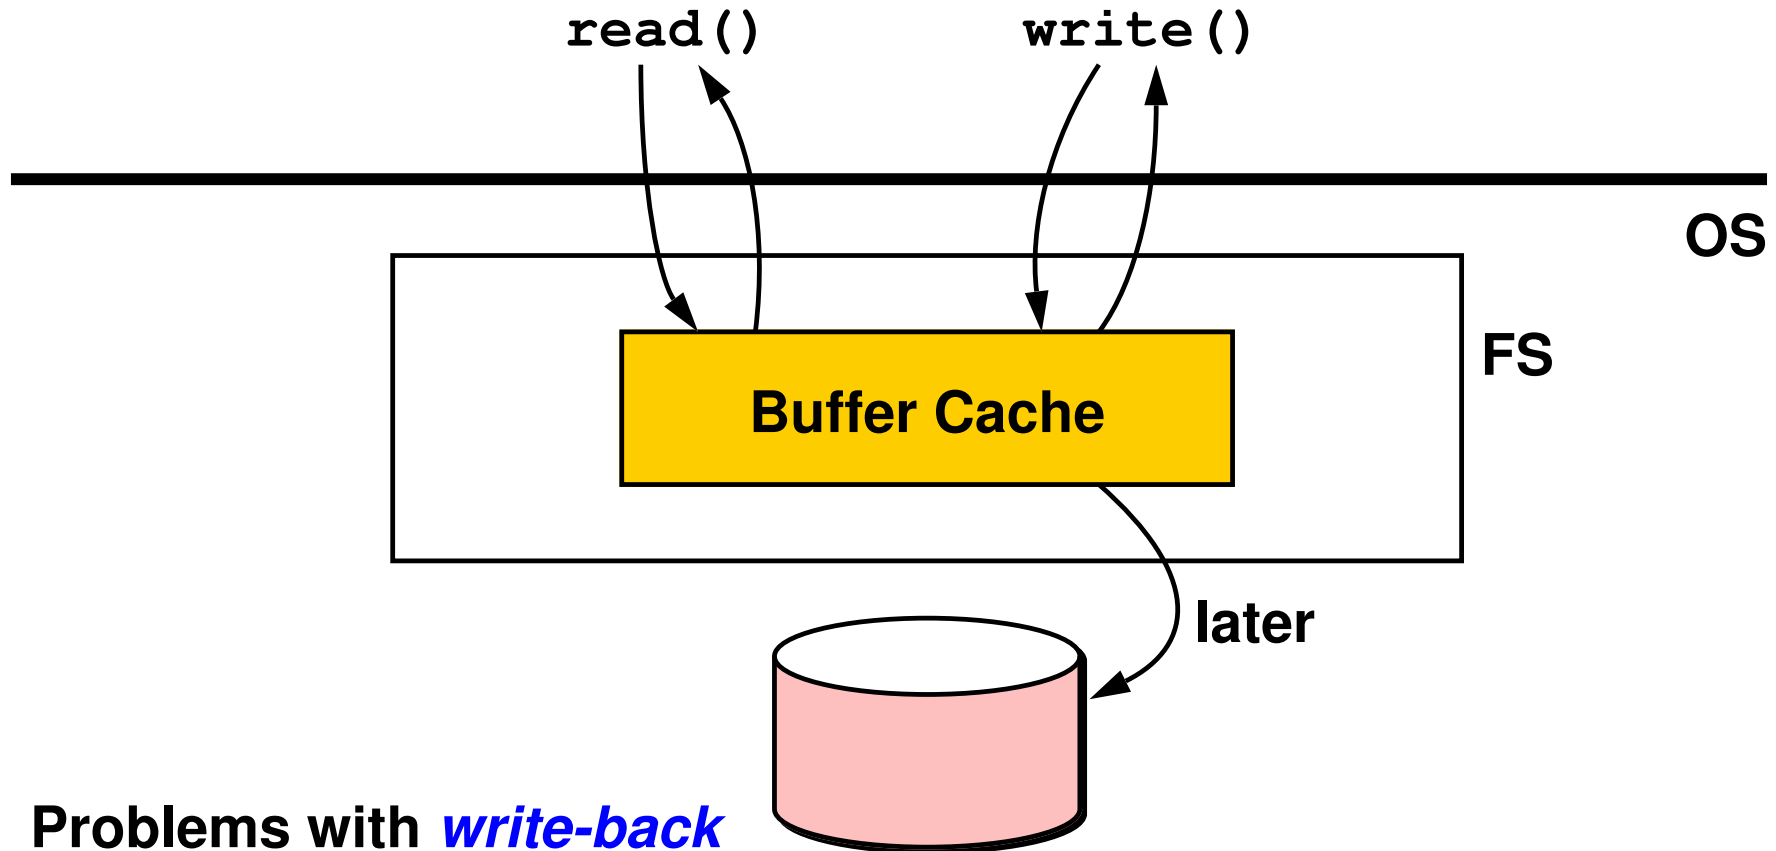


Aggressive caching


- = most read and write will have a *cache hit*
- = for writes, need to update the disk
 - *write through* vs. *write back*



The Buffer Cache




- ➡ Problems with *write-back*
 - writes to the disk can wait, may be for quite a while
 - longer the wait, higher the *risk*
- ➡ Need a file system optimized for *writing*!
 - how?
 - you organize the disk as a very long *log*



Log-Structured File Systems


- ➡ Main principles
- *never delete*
 - *append only*




0 1 2 ...

Log-Structured File Systems

- ➡ Main principles
- *never delete*
 - *append only*




0 1 2 ...




Log-Structured File Systems

- ➡ How does "never delete" and "always append" help with performance?
 - ▬ minimize seek latency
 - ▬ minimize rotational latency
 - write a cylinder at a time
- ➡ **Sprite FS** (a log-structured file system)
 - ▬ through batching, a single, long write can write out everything




LFS Data Placement Example

File On-disk
Representation:




LFS:



LFS Data Placement Example

➡ What happens if you want to modify the file?
 ➡ how does "append-only" really work?

➡ Ex: you create file A and then file B




- ➡ you modify file A, e.g., append to the last block of file A
- ➡ the new file will be referred as A'

LFS Data Placement Example

- ➡ What happens if you want to modify the file?
- ➡ how does "append-only" really work?

- ➡ Ex: you create file A and then file B

LFS:



Inode Map:

A

B


- ➡ you modify file A, e.g., append to the last block of file A
- ➡ the new file will be referred as A'

LFS Data Placement Example

- ➡ What happens if you want to modify the file?
- ➡ how does "append-only" really work?

- ➡ Ex: you create file A and then file B

LFS:



Inode Map:

A ————— B

- ➡ you modify file A, e.g., append to the last block of file A
- ➡ the updated file is still file A
 - but the inode has changed


LFS Data Placement Example

➡ What happens if you want to modify the file?

— how does "append-only" really work?


➡ Ex: you create file A and then file B

LFS:



Inode Map:

- you modify file A, e.g., append to the last block of file A
- the updated file is still file A
 - but the inode has changed
- a *piece* of the *inode map* is appended to the log
 - fixed regions (previous version and current version) on the disk keeps track of *all* the *inode map pieces*
 - ◆ known as *checkpoint file*



LFS Summary




Advantages

- good performance for writes
- can recover from crashes easily through the use of checkpoint files




Disadvantages

- can waste a lot of disk space




Extra Slides



Example

- ➡ We create two single-block files
 - ▬ dir1/file1
 - ▬ dir2/file2
- ➡ FFS
 - ▬ allocate and initialize inode for file1 and write it to disk
 - ▬ update dir1 to refer to it (and update dir1 inode)
 - ▬ write data to file1
 - allocate disk block
 - fill it with data and write to disk
 - update inode
 - ▬ six writes, plus six more for the other file
 - seek and rotational delays

FFS Picture




Example (Continued)





Sprite (a log-structured file system)


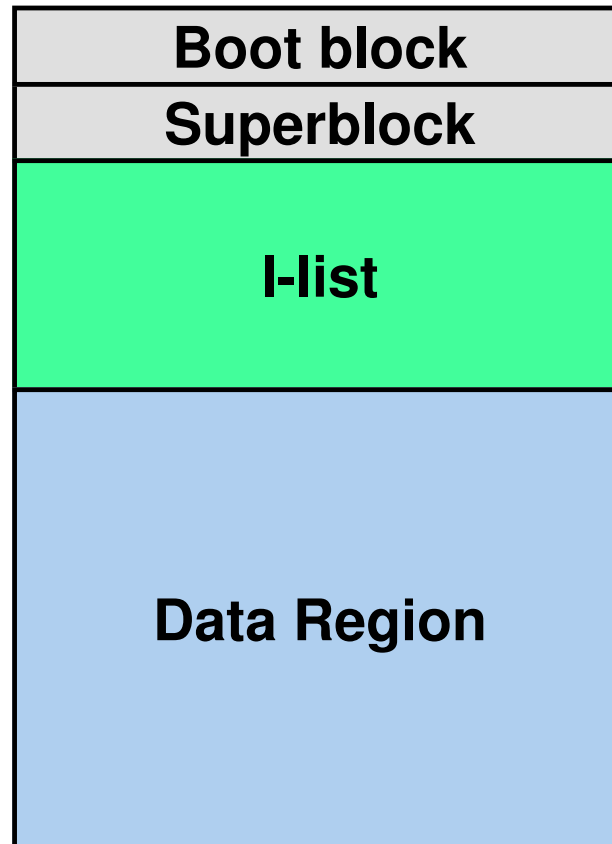
— one single, long write does everything



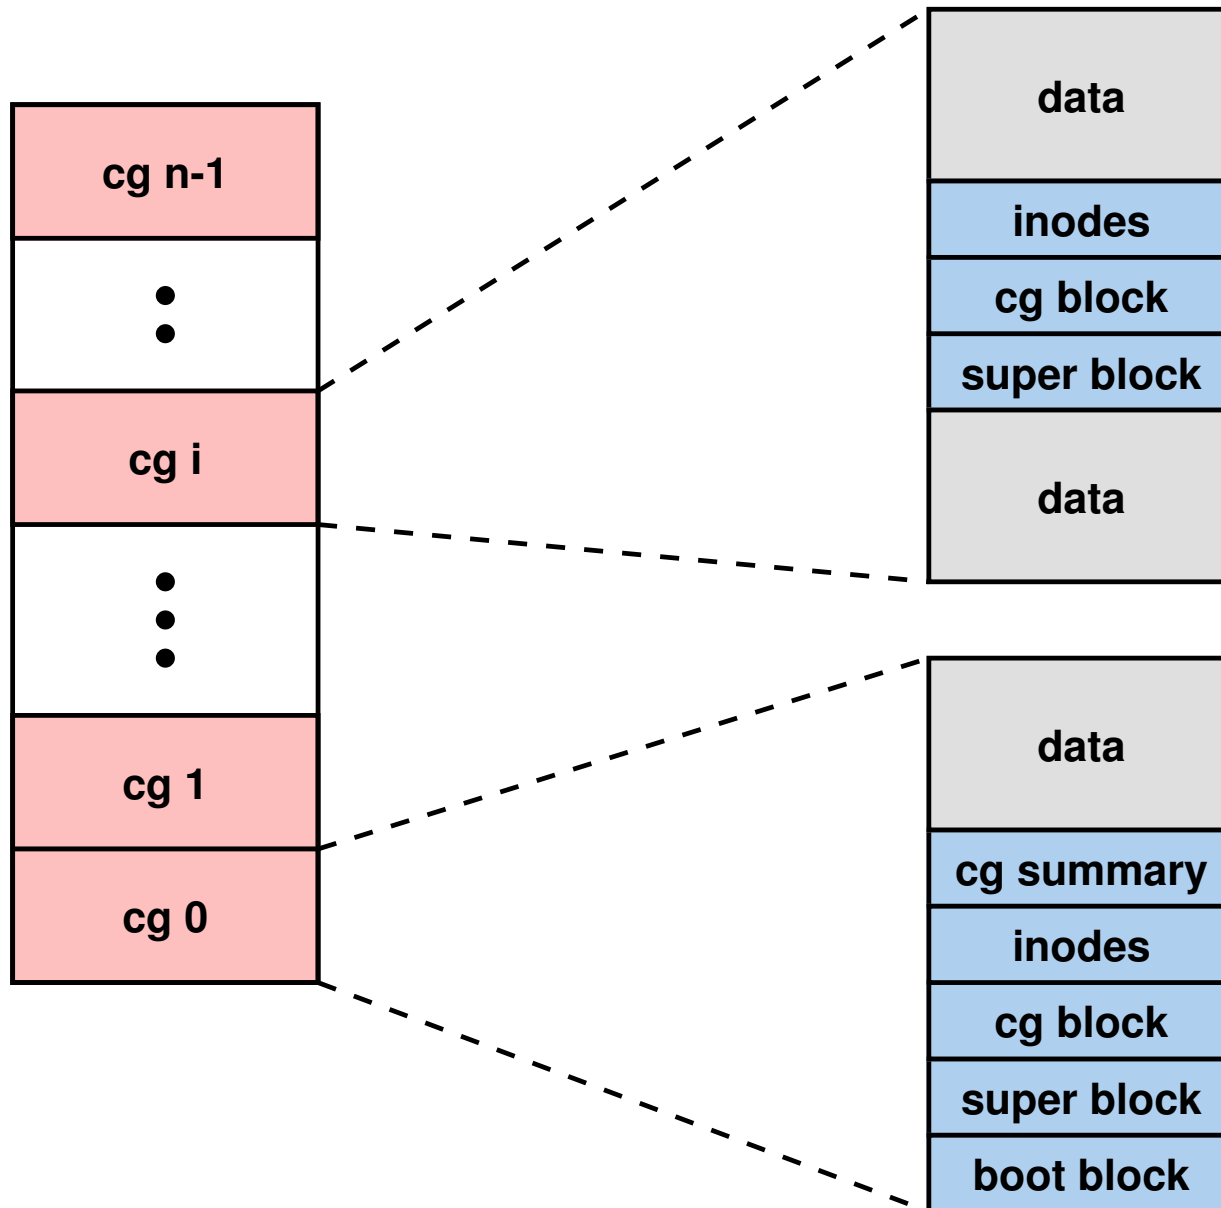
Sprite Picture



S5FS Layouts




FFS Layout




6.1 The Basics of File Systems

- ➡ UNIX's S5FS
- ➡ Disk Architecture
- ➡ Problems with S5FS
- ➡ Improving Performance
- ➡ *Dynamic Inodes*






NTFS Master File Table

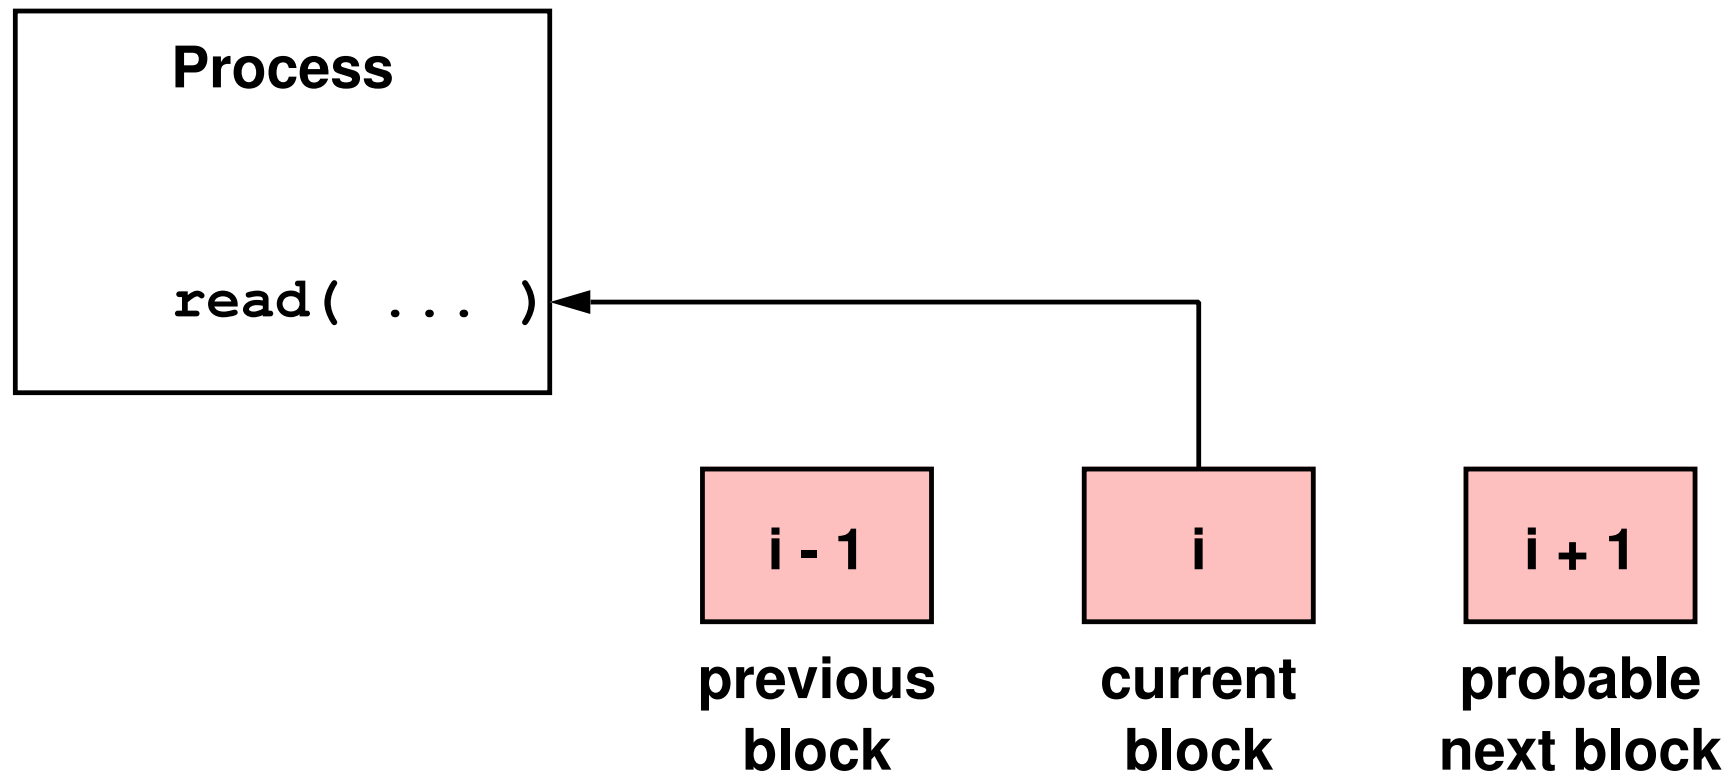
MFT
MFT Mirror
Log
Volume Info
Attribute Definitions
Root Directory
Free-Space Bitmap
Boot File
Bad-Cluster File
Quota Info
Expansion entries
User File 0
User File 1



The Buffer Cache



Multi-Buffered I/O



Maintaining the Cache

