


6.4 Multiple Disks



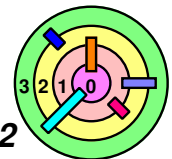
RAID

— what can be done if you lose an entire disk?





Benefits of Multiple Disks

- ➡ They hold more data than one disk does
- ➡ Data can be stored *redundantly* so that if one disk fails, they can be found on another
 - increase reliability
 - increase availability
- ➡ Data can be spread across multiple drives, allowing *parallel* access
 - increase effective access time




Logical Volume Manager


➡ Try not to modify existing file systems much



Logical Volume Manager

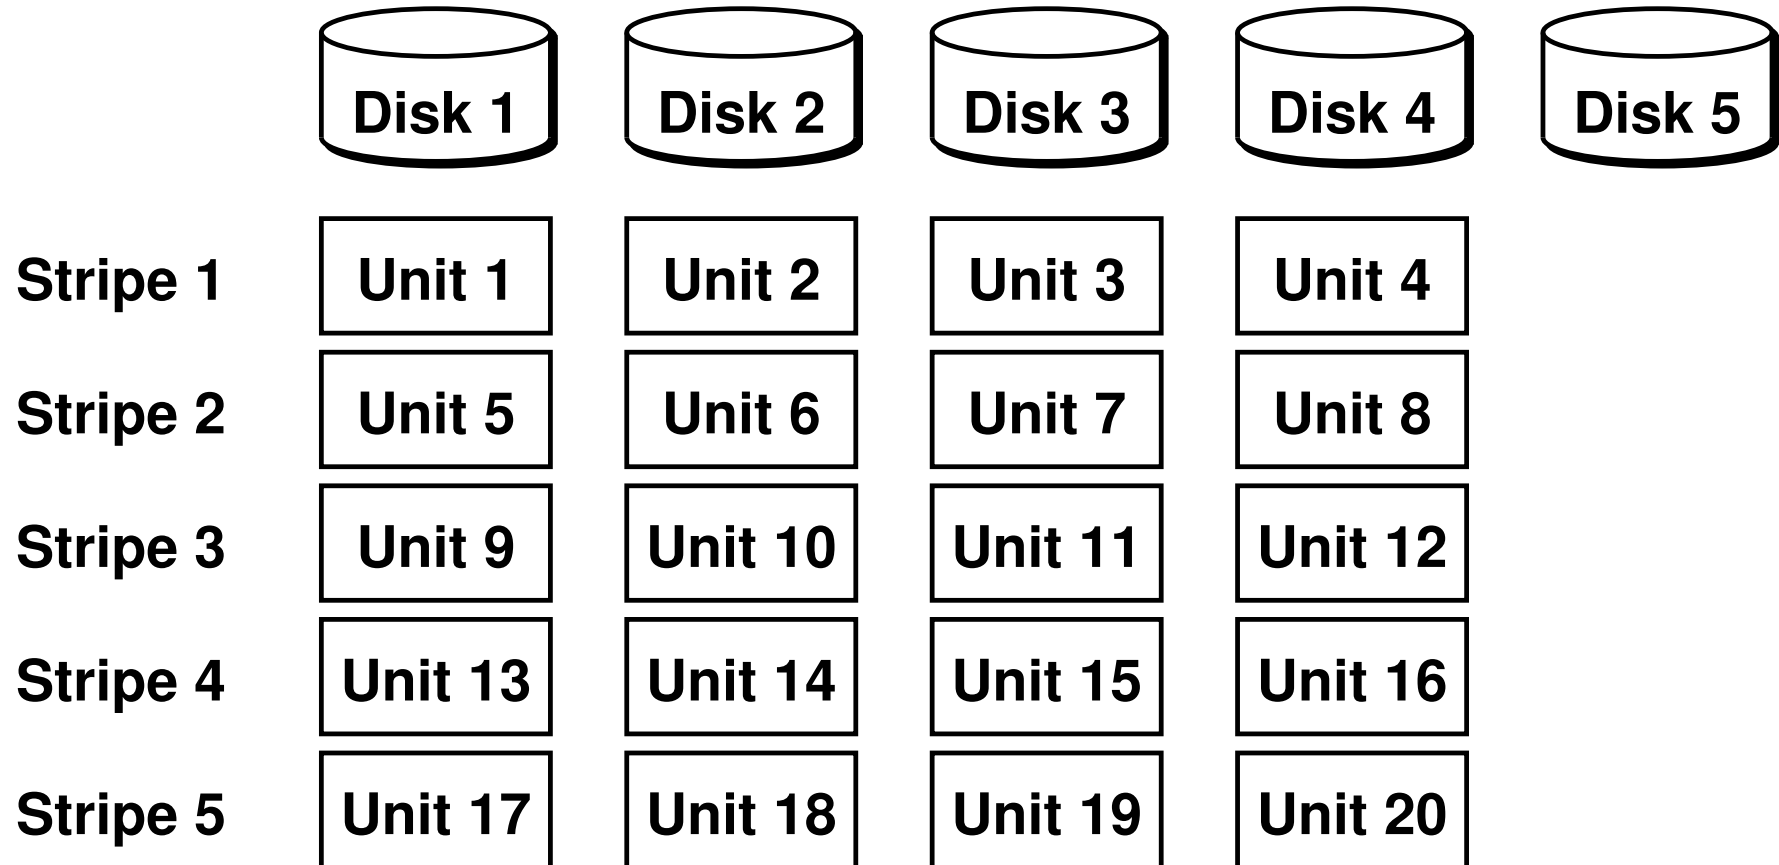


- ➡ **Spanning**
 - two real disks appear to file system as one large disk
- ➡ **Mirroring**
 - file system writes redundantly to both disks
 - reads from one




Striping

➡ Ex: *stripe width* = 4




- ➡ theoretically, a *striping unit* can be a bit (i.e., *bit-interleaving*)
 - pack these bits into disk blocks and store on disk



Parallel Disks

➡ Advantages


- increase parallelism
 - can retrieve blocks belonging to multiple files simultaneously if they are on different disks



- reduced access time if a block is spread over multiple disks
 - ◆ seek in parallel, same rotational latency on all disks

➡ Disadvantages

- higher variance
 - average is just part of the story
- worse reliability
- heterogenous disks



How To Stripe?




How to stripe?

- what's the best *stripe width* (i.e., across how many disks)?
- how large should be the *striping unit* (i.e., how much to put on one disk before moving to the next disk)?





Concurrency Factor: how many requests are available to be executed at once?


- one request in queue at a time
 - concurrency factor = 1
 - e.g., one single-threaded application placing one request at a time
- many requests in queue
 - concurrency factor > 1
 - e.g., multiple threads placing file-system requests
- the larger the concurrency factor, the less important striping is
 - in general, performance is better with *larger striping unit*





Striping Unit Size




- ➡ Bottom solution seems better
- data transfer time is 1/4 of the solution on the top



Striping Unit Size




➡ The above two cases have *same transfer time*
➡ can *reduce seek time* by using a *larger striping unit*




Striping: The Effective Disk

- ➡ Improved effective transfer speed
 - parallelism
- ➡ No improvement in seek and rotational delays
 - sometimes worse
- ➡ A system depending on N disks is much more likely to fail than one depending on one disk
 - if probability of one disk failing is f
 - probability of N -disk system failing is $(1 - (1 - f)^N)$
 - assumes failures are i.i.d., which is probably wrong ...
 - ◆ *i.i.d.:* independent and identically distributed





RAID to the Rescue


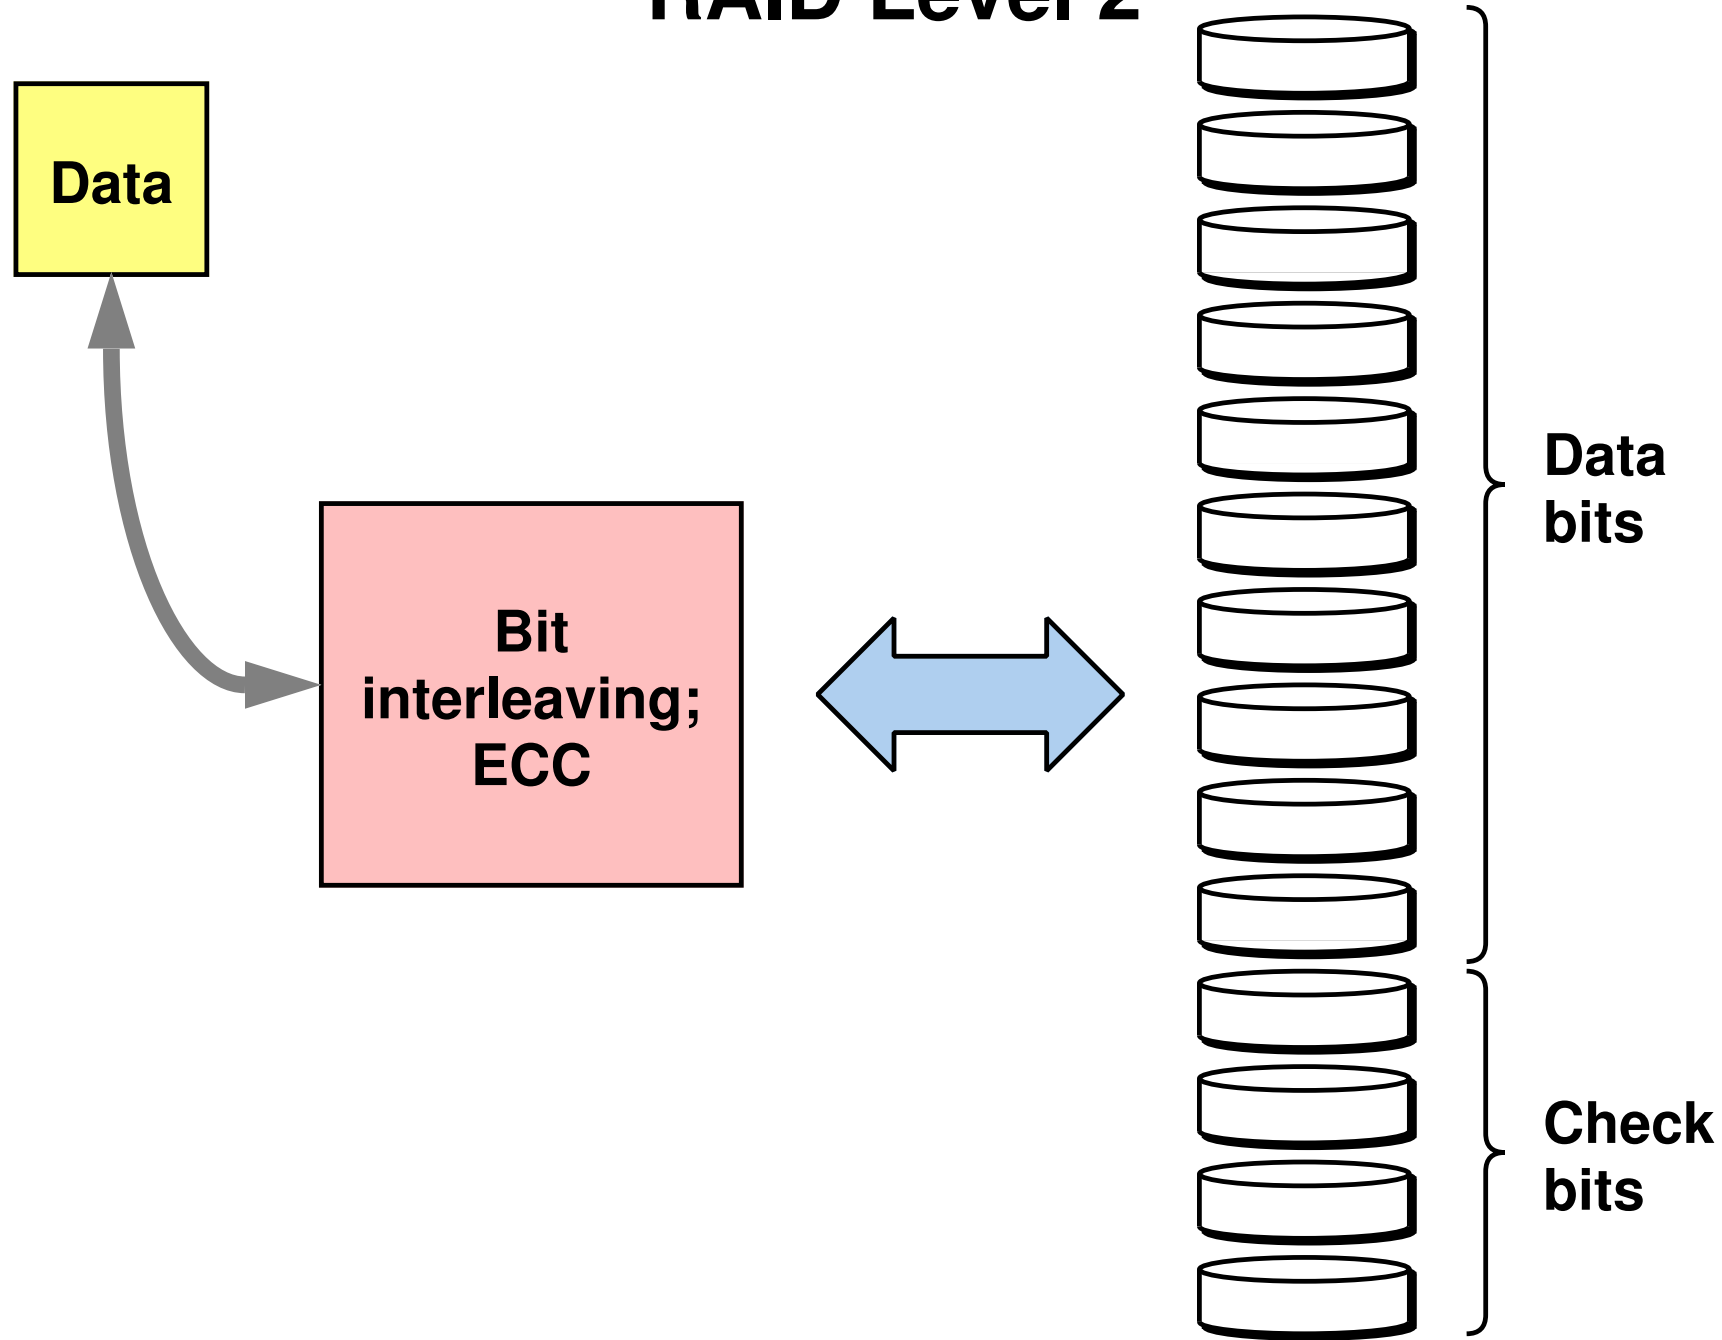
- ➡ **RAID:** Redundant Array of Inexpensive Disks
- (as opposed to Single Large Expensive Disk: SLED)
 - combine *striping* with *mirroring*
 - 5 different variations originally defined
 - RAID level 1 through RAID level 4 developed by IBM
 - RAID level 5 developed by UC Berkeley
 - ◆ RAID level 0: pure striping (numbering extended later)
 - RAID level 1: pure mirroring





RAID Level 1



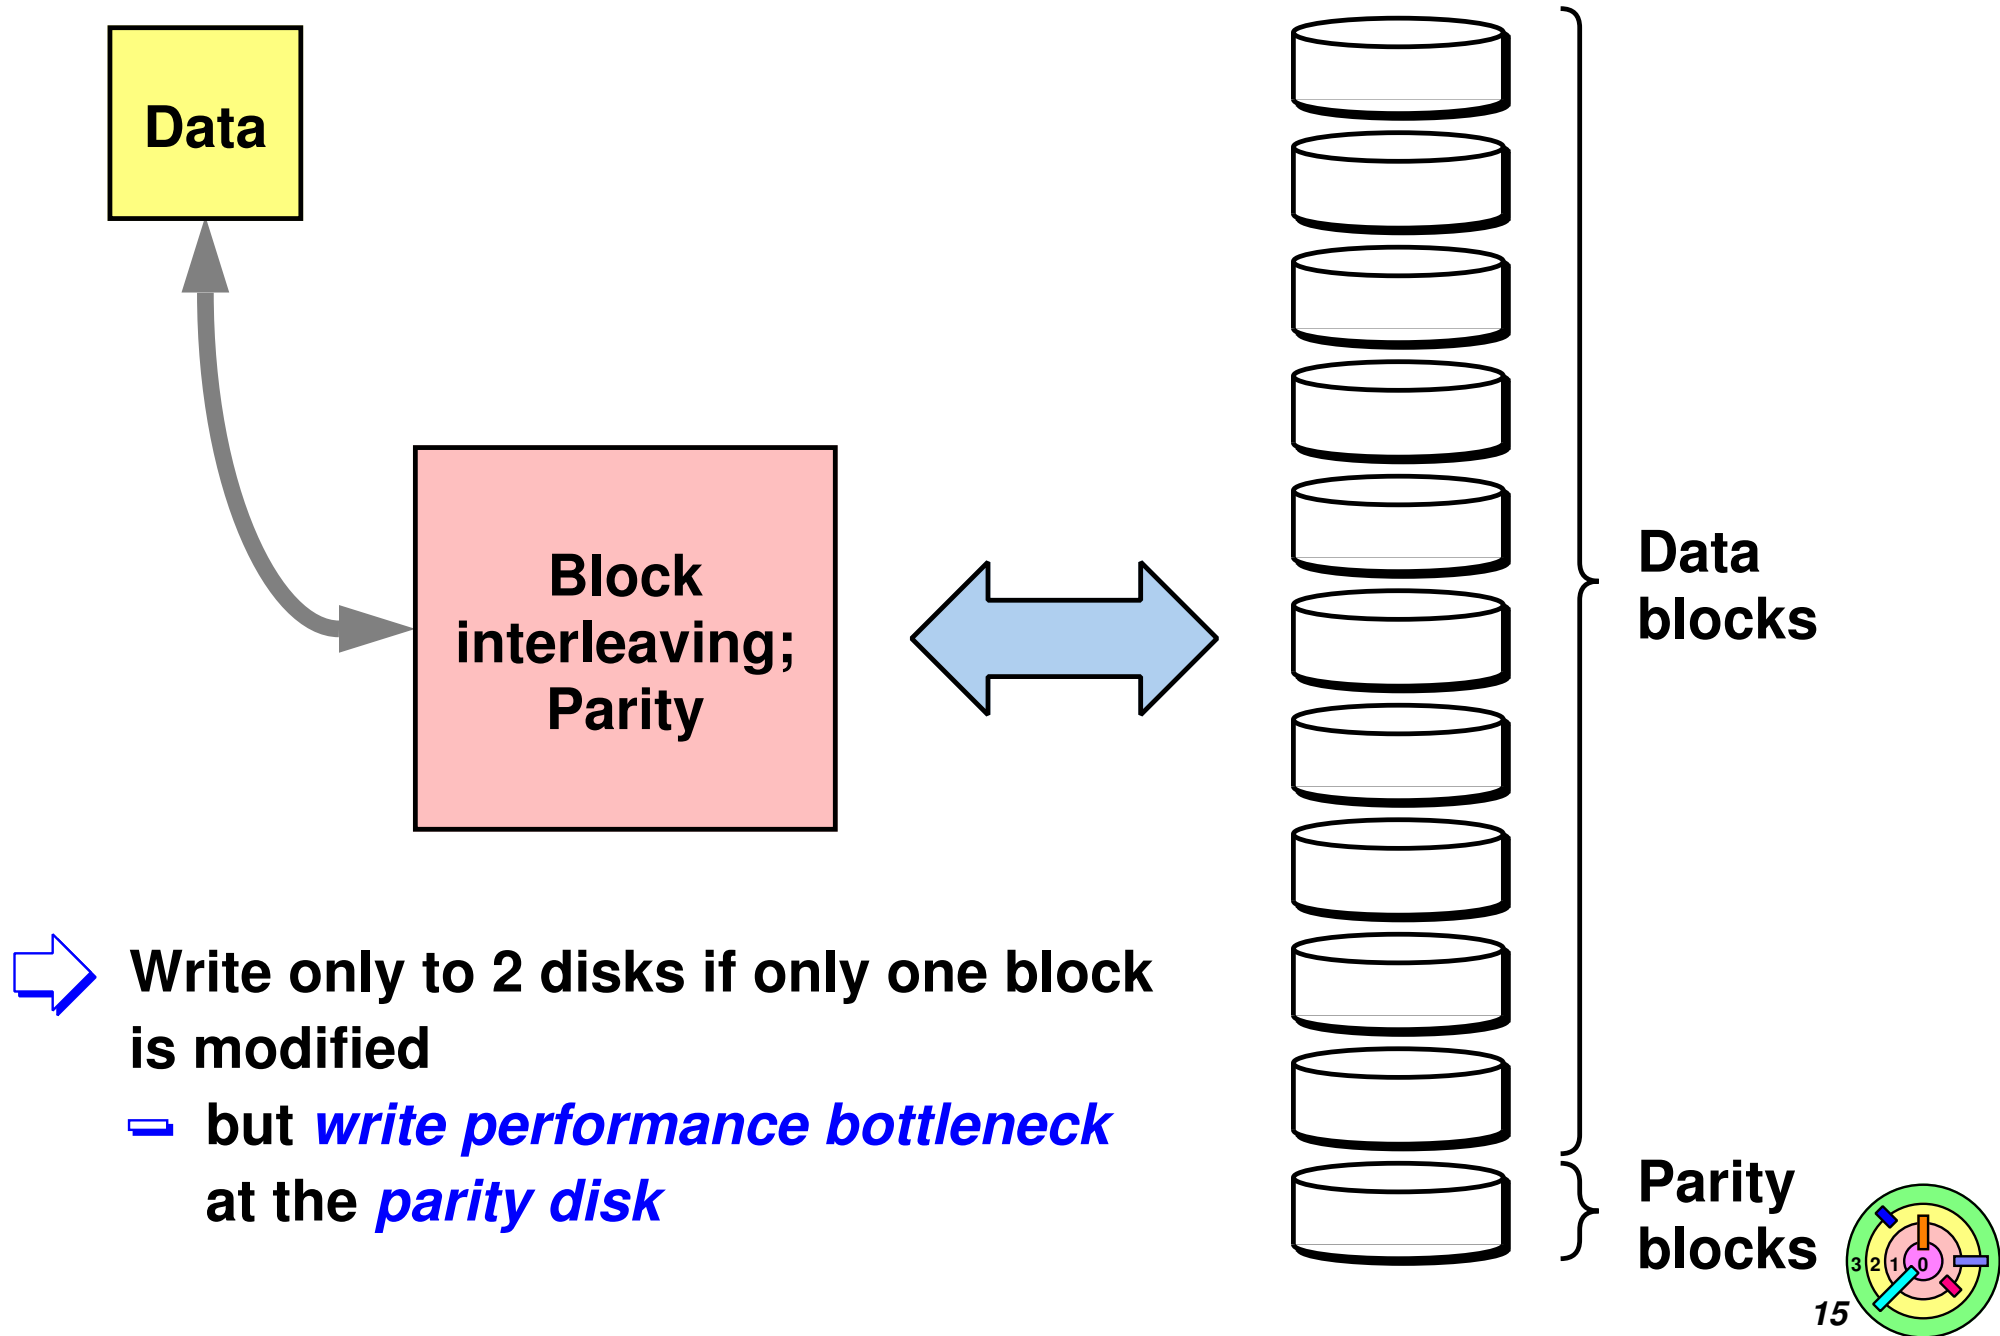
RAID Level 2



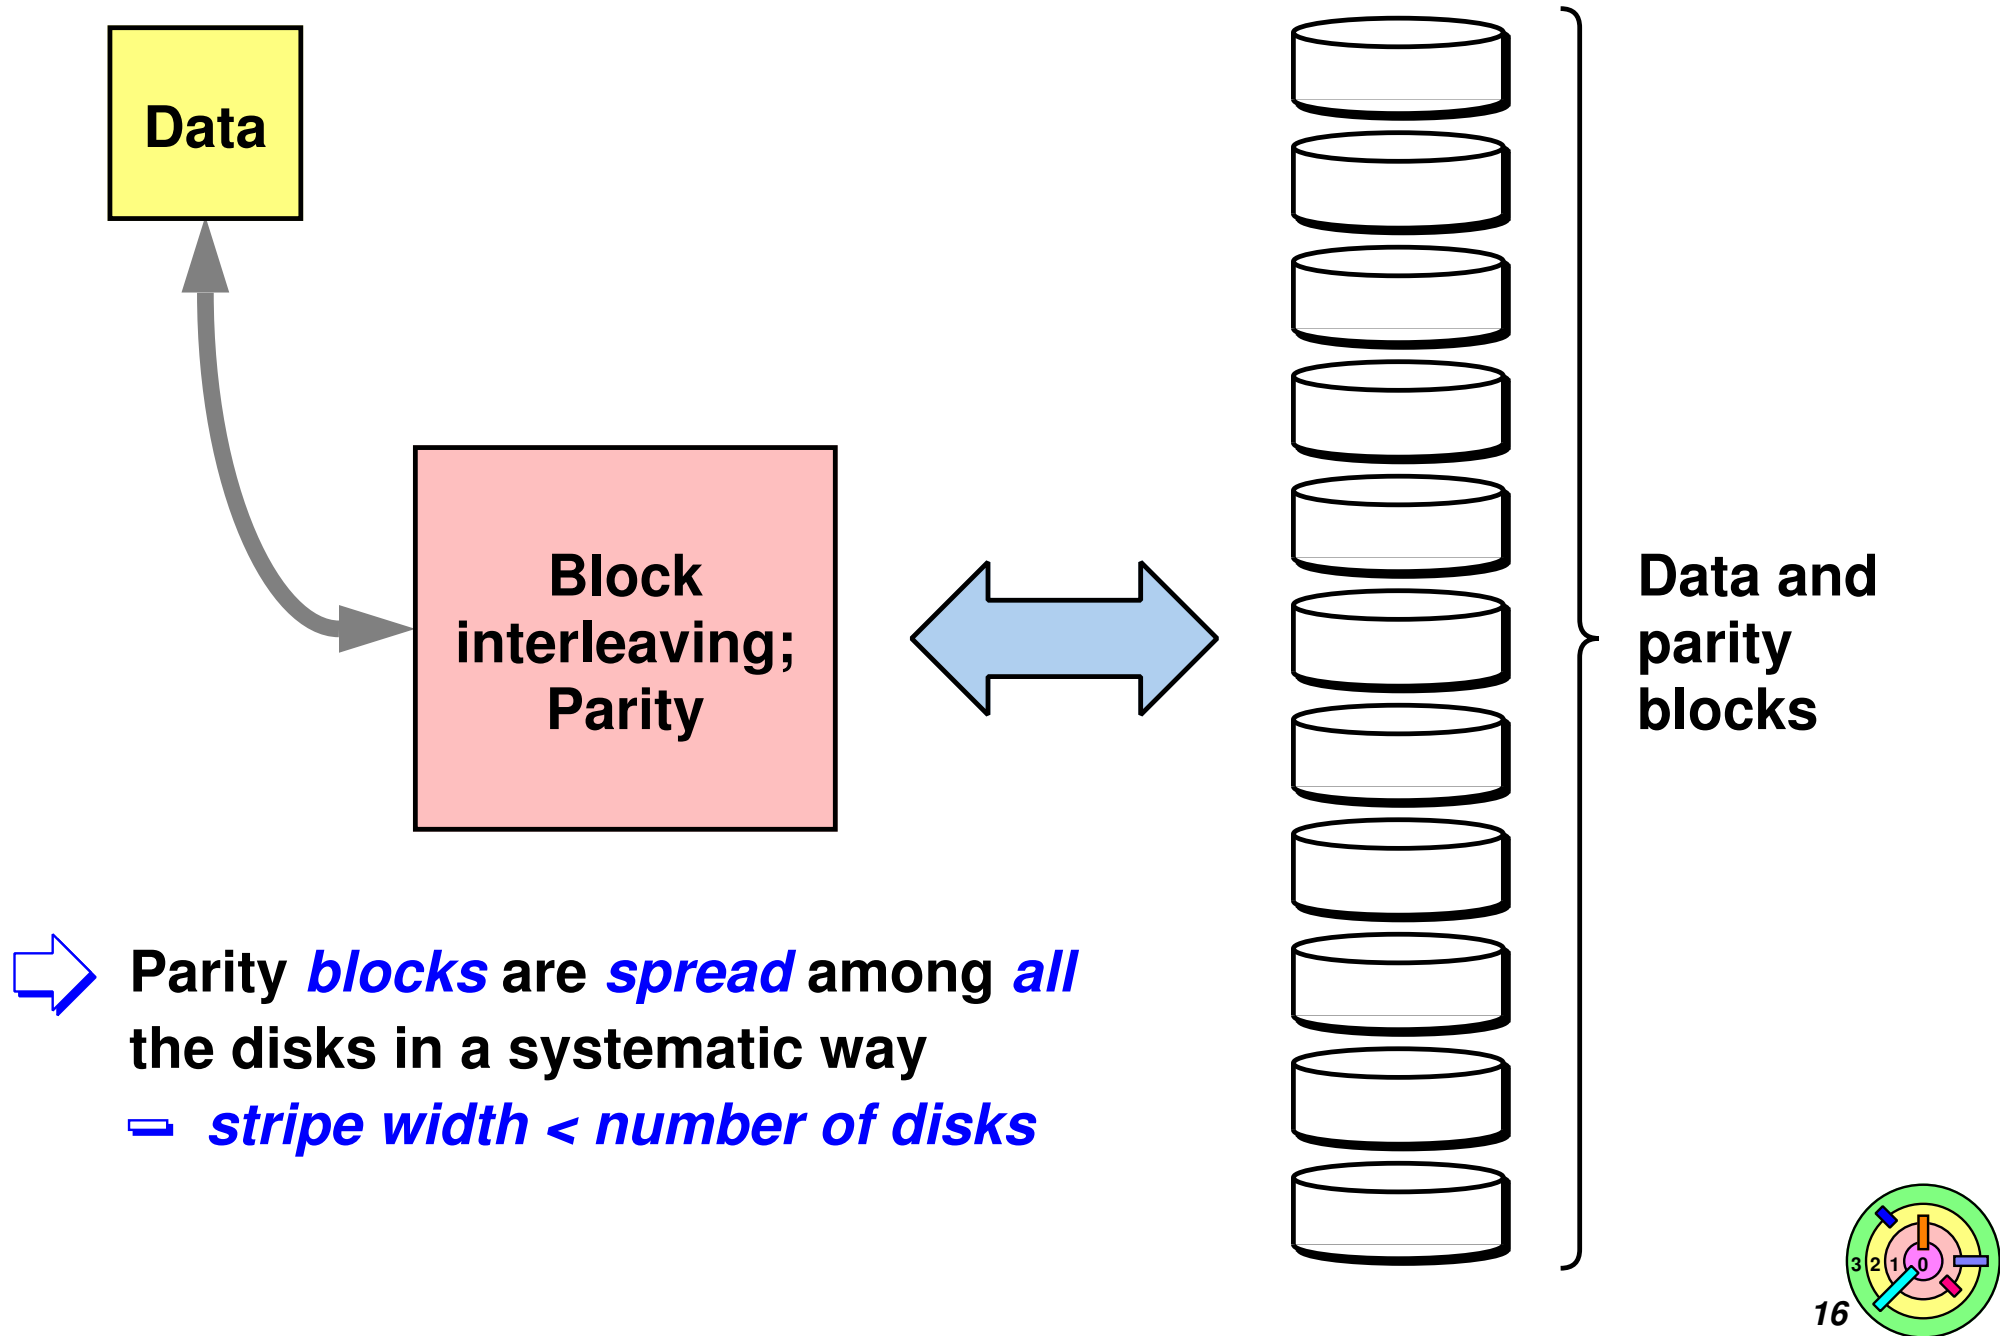
RAID Level 3



RAID Level 4

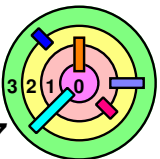


RAID Level 5



RAID 4 vs. RAID 5

- ➡ **Lots of small writes**
 - RAID 5 is best
- ➡ **Mostly large writes**
 - multiples of stripes
 - either is fine
- ➡ **Expansion**
 - add an additional disk or two
 - RAID 4: add them and recompute parity
 - RAID 5: add them, recompute parity, shuffle data blocks among all disks to reestablish check-block pattern
- ➡ **Write performance**
 - RAID 4: parity disk have workload multiple of other disks
 - RAID 5: same workload on all disks on the average
- ➡ **One disk failure**
 - RAID 4: parity disk have workload multiple of other disks
 - RAID 5: work load spread out more evenly



Beyond RAID 5




RAID 6

- like RAID 5, but additional parity
- handles two failures




Cascaded RAID

- RAID 1+0 (RAID 10)
 - striping across mirrored drives
- RAID 0+1
 - two striped sets, mirroring each other



6.5 Flash Memory

- ➡ **Flash Technology**
- ➡ **Flash-Aware File Systems**
- ➡ **Augmenting Disk Storage**



Beyond Disks: Flash




Pro

- Flash block \approx file-system block
- Random access
 - no seek, no rotational latency
- Low power
- Vibration-resistant



Con

- Limited lifetime
- Write is expensive
- Cost more than disks
 - 128GB SSD: ~\$300
 - 1TB disk: ~\$60



Flash Memory



Two technologies

— NOR

- byte addressable

— NAND

- page addressable (about 1-4KB per page and 512KB per block)
- cheaper
 - ◆ suitable for file systems use
- limit on P/E (program/erase) cycle, about 10,000




Writing

— newly "erased" block is all ones

— "programming" changes some ones to zeroes

- per byte in NOR; per page in NAND (multiple pages/block)
- to change zeroes to ones, must erase entire block
- can erase no more than ~100k times/block





Coping

➡ *Wear leveling*

- spread writes (erasures) across entire drive
- approach:
 - *flash translation layer (FTL)*
 - *log-structured file system*
 - ◆ blocks on the flash drive are used sequentially

➡ FTL: Flash translation layer (often on a separate device)

- specification from 1994
- provides disk-like block interface (firmware on device controller)
- maps disk blocks to flash blocks
 - mapping changed dynamically to effect wear-leveling





Coping

➡ *Wear leveling*

- spread writes (erasures) across entire drive
- approaches:
 - *flash translation layer (FTL)*
 - *log-structured file system*
 - ◆ blocks on the flash drive are used sequentially

➡ FTL: Flash translation layer (often on a separate device)

- specification from 1994
- provides disk-like block interface (firmware on device controller)
- maps disk blocks to flash blocks
 - mapping changed dynamically to effect wear-leveling





Coping

➡ *Wear leveling*

- spread writes (erasures) across entire drive
- approaches:
 - *flash translation layer (FTL)*
 - *log-structured file system*
 - ◆ blocks on the flash drive are used sequentially

➡ FTL: Flash translation layer (often on a separate device)

- specification from 1994
- provides disk-like block interface (firmware on device controller)
- maps disk blocks to flash blocks
 - mapping changed dynamically to effect wear-leveling





Coping

➡ *Wear leveling*

- spread writes (erasures) across entire drive
- approaches:
 - *flash translation layer (FTL)*
 - *log-structured file system*
 - ◆ blocks on the flash drive are used sequentially

➡ FTL: Flash translation layer (often on a separate device)

- specification from 1994
- provides disk-like block interface (firmware on device controller)
- maps disk blocks to flash blocks
 - mapping changed dynamically to effect wear-leveling





Coping

➡ *Wear leveling*

- spread writes (erasures) across entire drive
- approaches:
 - *flash translation layer (FTL)*
 - *log-structured file system*
 - ◆ blocks on the flash drive are used sequentially

➡ FTL: Flash translation layer (often on a separate device)

- specification from 1994
- provides disk-like block interface (firmware on device controller)
- maps disk blocks to flash blocks
 - mapping changed dynamically to effect wear-leveling





Coping

➡ *Wear leveling*

- spread writes (erasures) across entire drive
- approaches:
 - *flash translation layer (FTL)*
 - *log-structured file system*
 - ◆ blocks on the flash drive are used sequentially

➡ FTL: Flash translation layer (often on a separate device)

- specification from 1994
- provides disk-like block interface (firmware on device controller)
- maps disk blocks to flash blocks
 - mapping changed dynamically to effect wear-leveling





Coping

➡ *Wear leveling*

- spread writes (erasures) across entire drive
- approaches:
 - *flash translation layer (FTL)*
 - *log-structured file system*
 - ◆ blocks on the flash drive are used sequentially


➡ FTL: Flash translation layer (often on a separate device)

- specification from 1994
- provides disk-like block interface (firmware on device controller)
- maps disk blocks to flash blocks
 - mapping changed dynamically to effect wear-leveling




Flash with FTL

- ➡ Which file system?
- FAT32 (sort of like S5FS, but from Microsoft)
 - NTFS
 - FFS
 - Ext3
- ➡ All were designed to exploit disks
- much of what they do are irrelevant for flash




Flash without FTL

- ➡ Known as memory technology device (MTD)
 - software wear-leveling
 - perhaps other tricks



JFFS and JFFS2


- ➡ **Journaling flash file system**
 - ▬ **log-based: no journal!**
 - each log entry contains inode info and some data
 - garbage collection copies info out of partially obsoleted blocks, allowing block to be erased
 - complete index of inodes (i.e., meta-data) kept in RAM
 - ◆ entire file system must be read when mounted



UBI / UBIFS

- ➡ **UBI (unsorted block images)**
 - ▬ supports multiple logical volumes on one flash device
 - ▬ performs wear-leveling across entire device
 - ▬ handles bad blocks

- ➡ **UBIFS**
 - ▬ file system layered on UBI
 - ▬ it really has a journal (originally called JFFS3)
 - ▬ file map kept in flash as B+ tree
 - ▬ no need to scan entire file system when mounted




Flash as Part of the Hierarchy

- ➡ **Flash as log device**
 - ▬ aggregate write throughput sufficient, but latency is bad
 - ▬ augment with DRAM and a "super-capacitor"

- ➡ **Flash as cache**
 - ▬ large level-2 cache
 - integrated into ZFS
 - can use cheaper (slower) disks with no loss of performance
 - ◆ reduced power consumption

6.6 Case Studies

- ➡ FFS
- ➡ Ext3
- ➡ Reiser FS
- ➡ NTFS
- ➡ WAFL
- ➡ ZFS




Linux

➡ Linux had 57 file systems built for it to date!

➡ FFS

- ext2
- ext3 (journaling - crash resiliency)
 - ReiserFS (B-tree everywhere)
- ext4
 - extents (optimize read/write)
 - LVM
 - hash trees for directories
- BtrFS (Oracle)



Windows NT


➡ NTFS

- extents (optimize read/write)
- B-trees (optimize directory lookup)
- journaling (crash resiliency)

Mac OS X


➡ Mac OS X

- HFS+ (planned to use ZFS but dropped the idea)
 - extents (optimize read/write)
 - B*-trees (optimize directory lookup)
 - journaling (crash resiliency)

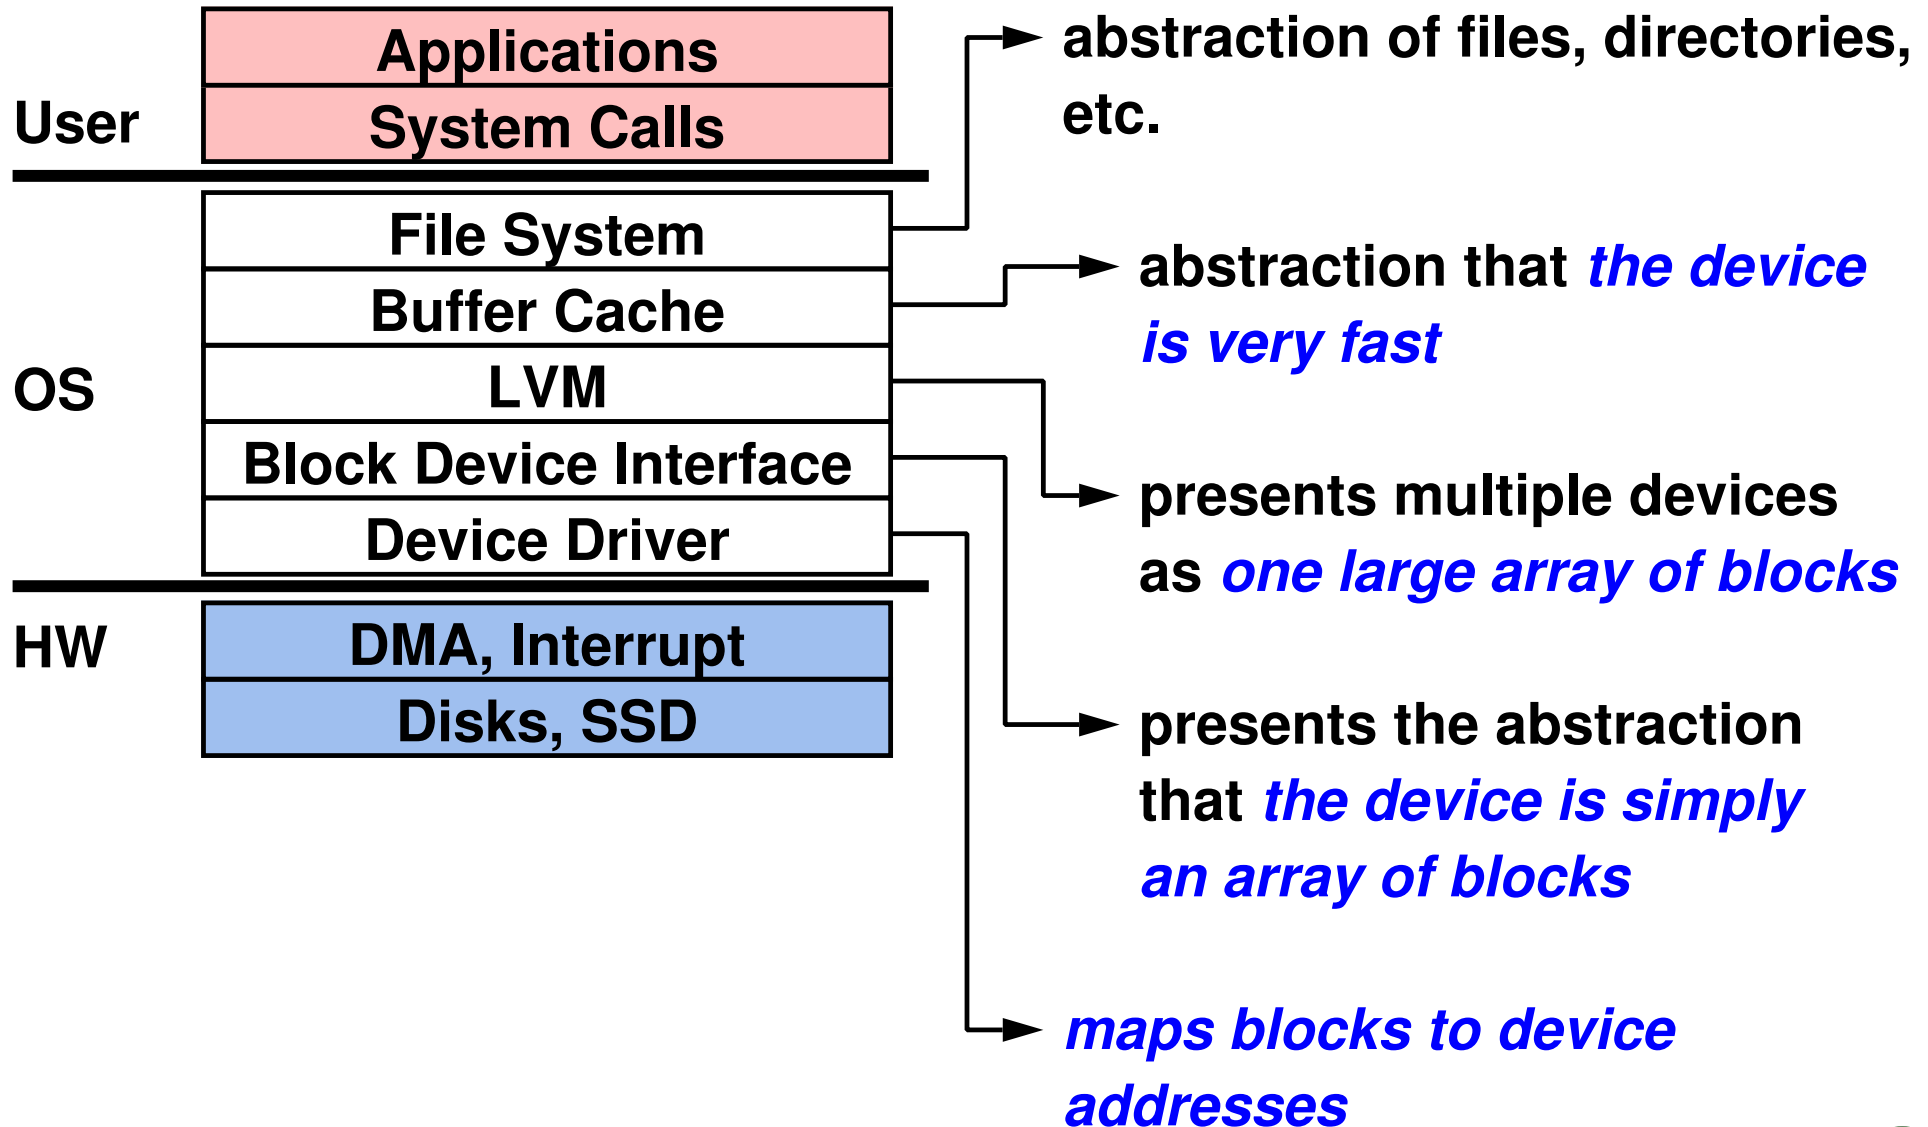


Journaling


- ➡ Why did everyone choose journaling and not shadow pages?
- journaling can be added to any existing file system



File System Summary




Extra Slides



NTFS

"Volume aggregation" options


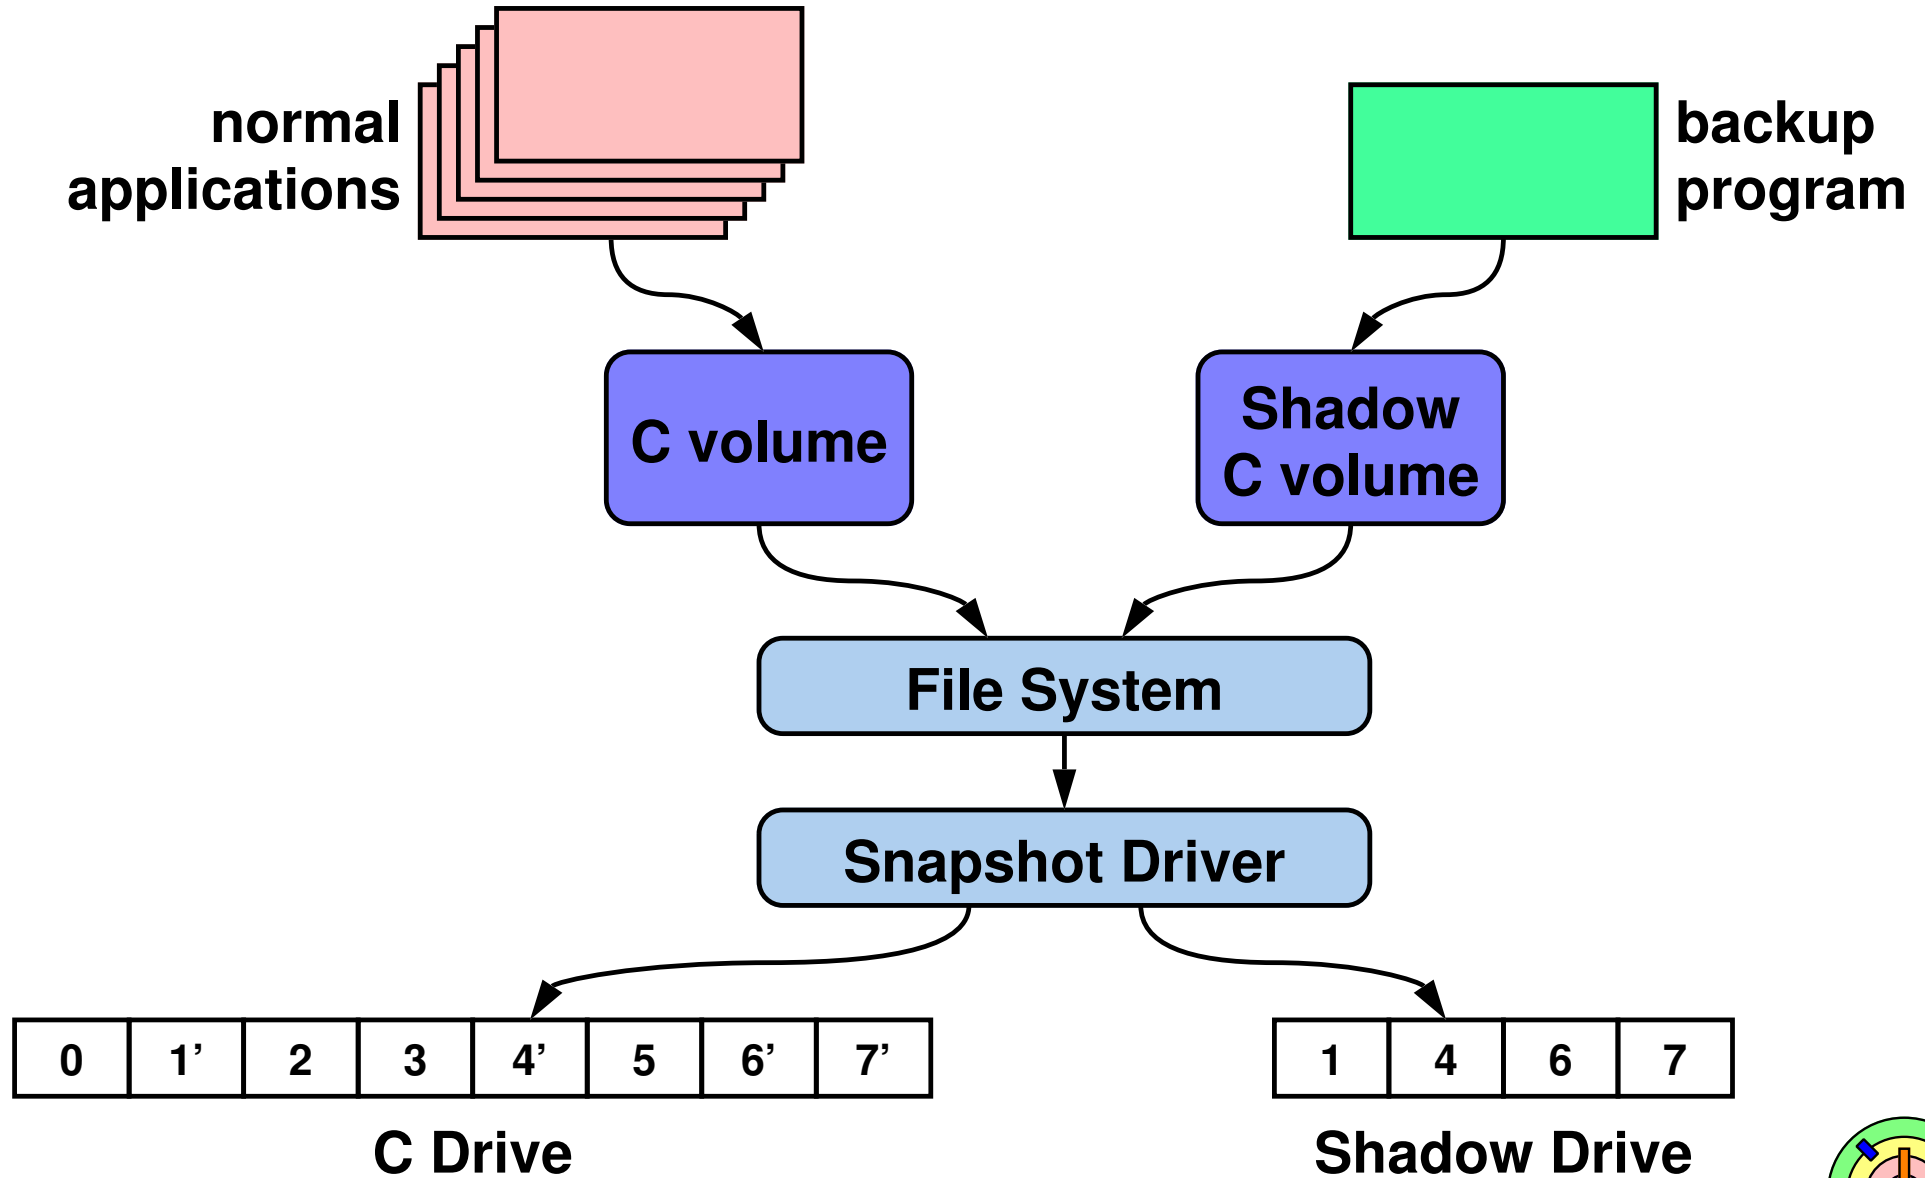
- spanned volumes
- RAID 0 (striping)
- RAID 1 (mirroring)
- RAID 5
- snapshots



Backups

- ➡ **Want to back up a file system**
 - while still using it
 - files are being modified while the backup takes place
 - applications may be in progress - files in inconsistent states
- ➡ **Solution**
 - have critical applications quickly reach a safe point and pause
 - snapshot the file system
 - resume applications
 - back up the snapshot


Windows Snapshots



NTFS File Records


Name	Standard attributes	Object ID	Data stream
------	---------------------	-----------	-------------

Name	Standard attributes	Object ID	Properties stream	Data stream
------	---------------------	-----------	-------------------	-------------




Additional NTFS Features

- ➡ **Data compression**
 - ▬ run-length encoding of zeroes
 - ▬ compressed blocks
- ➡ **Encrypted files**

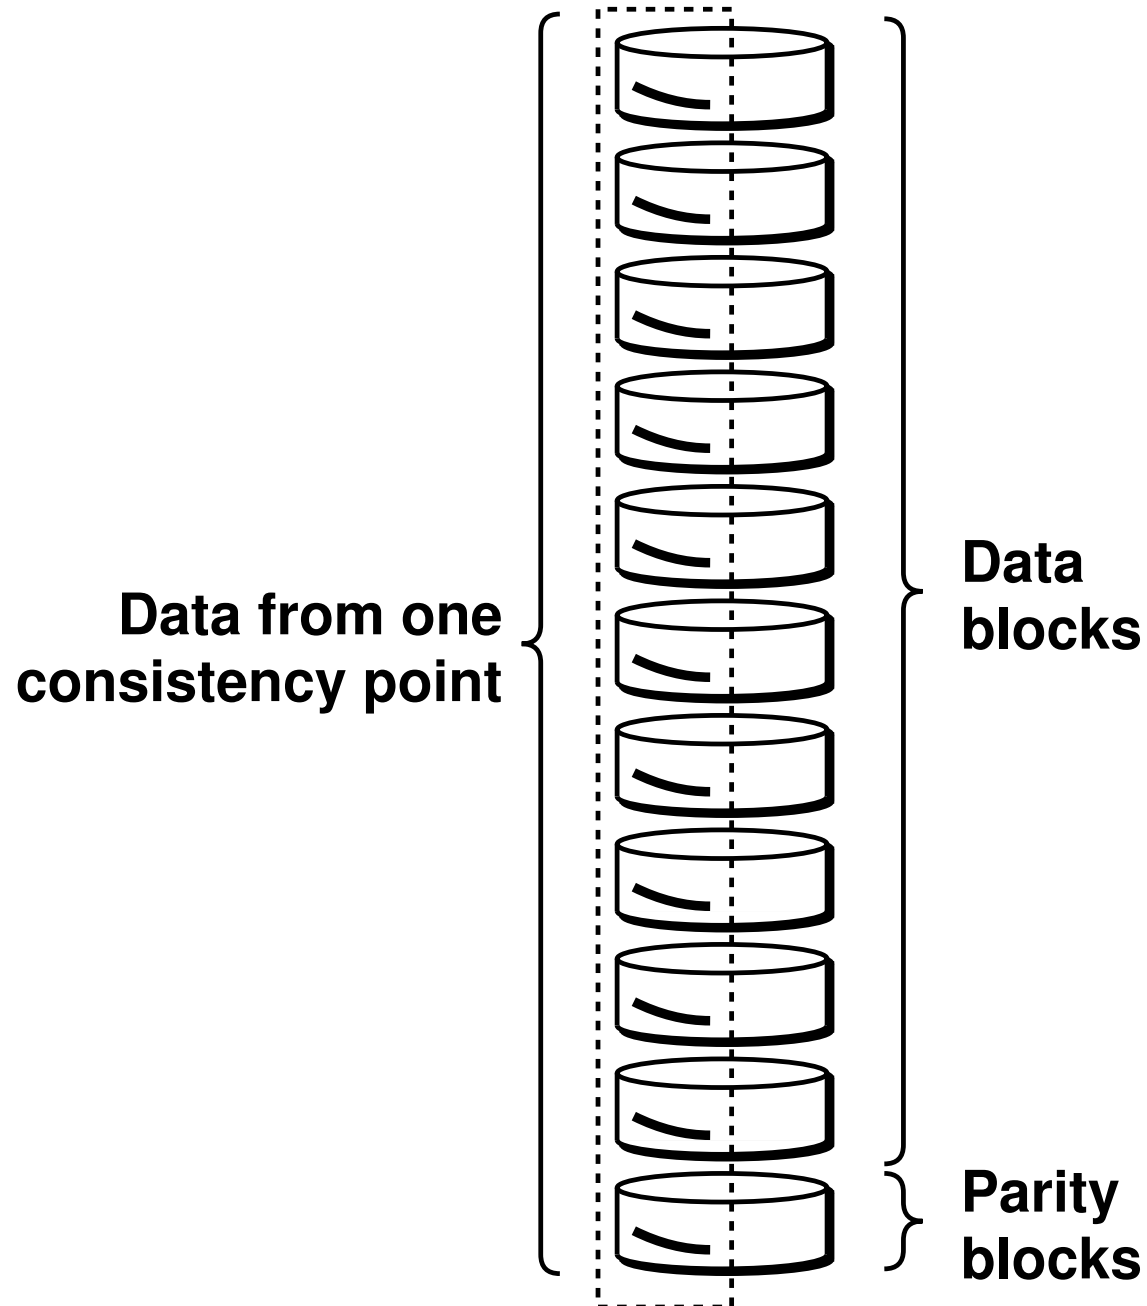


WAFL

- ➡ Runs on special-purpose OS
 - machine is dedicated to being a *filer*
 - handles both NFS and CIFS requests
- ➡ Utilizes shadow paging and log-structured writes
- ➡ Provides snapshots




WAFL and RAID



Consistency Points ... and Beyond

- ➡ **Consistency points taken every ~10 seconds**
 - ▬ too relaxed for many applications
 - ▬ NFS
 - ▬ databases
- ➡ **Solution ...**
 - ▬ battery-backed-up RAM
 - a.k.a. non-volatile RAM (NVRAM)



Snapshots


- ➡ Periodic snapshots kept of file system
 - made easy with shadow paging




Paranoia

➡ You think your files are safe simply because they're on a RAID-4 or RAID-5 system ...


- power failure at inopportune moment
- parity is irreparably wrong
- obscure bug in controller firmware or OS
- data is garbage (but with correct parity!)
- sysadmin accidentally scribbled on one drive
- (profuse apologies ...)
- out of disk space
- must restructure 4TB file system
- out of address space
- 264 isn't as big as it used to be




Partial Writes



Partial Writes





Partial Writes



Small Writes

Writing these:




Data
blocks

Parity
blocks



Small Writes

Writing these:



Data
blocks


Requires reading
then writing this:





Parity
blocks





Hardware RAID



Adding a Disk (1)




Adding a Disk (2)



ZFS

The Last (?!) Word in File Systems




ZFS Layers


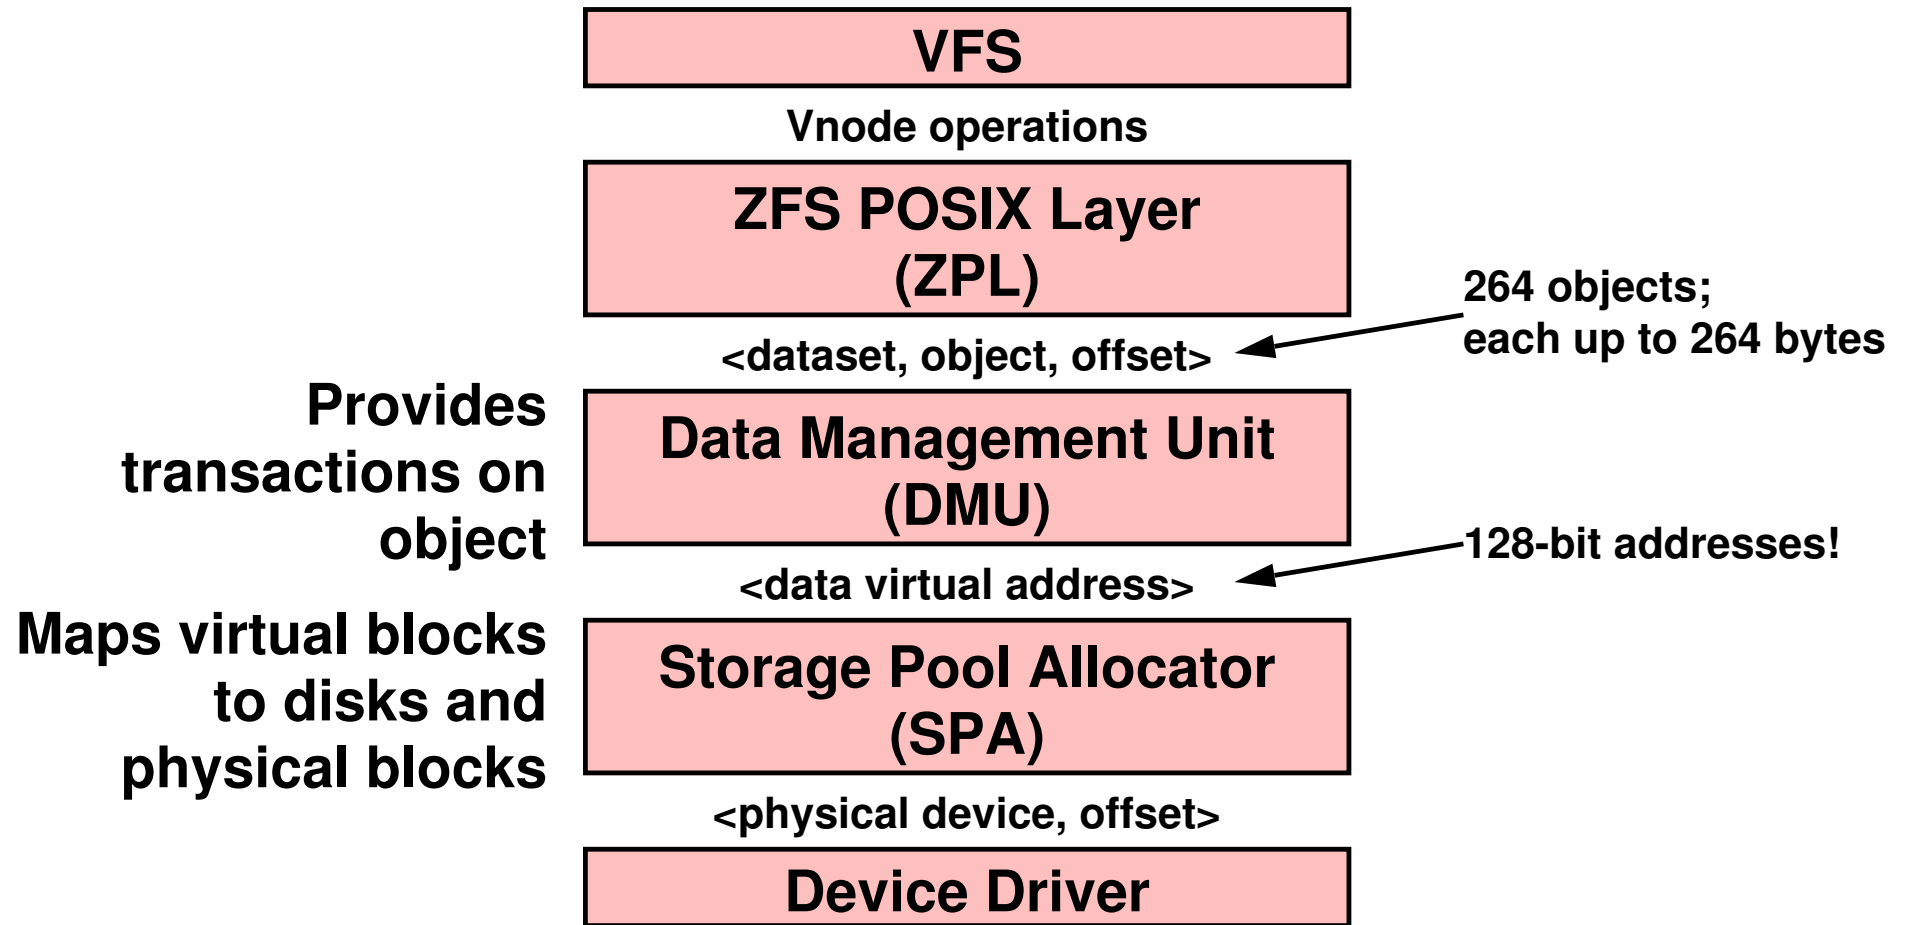


Data Management

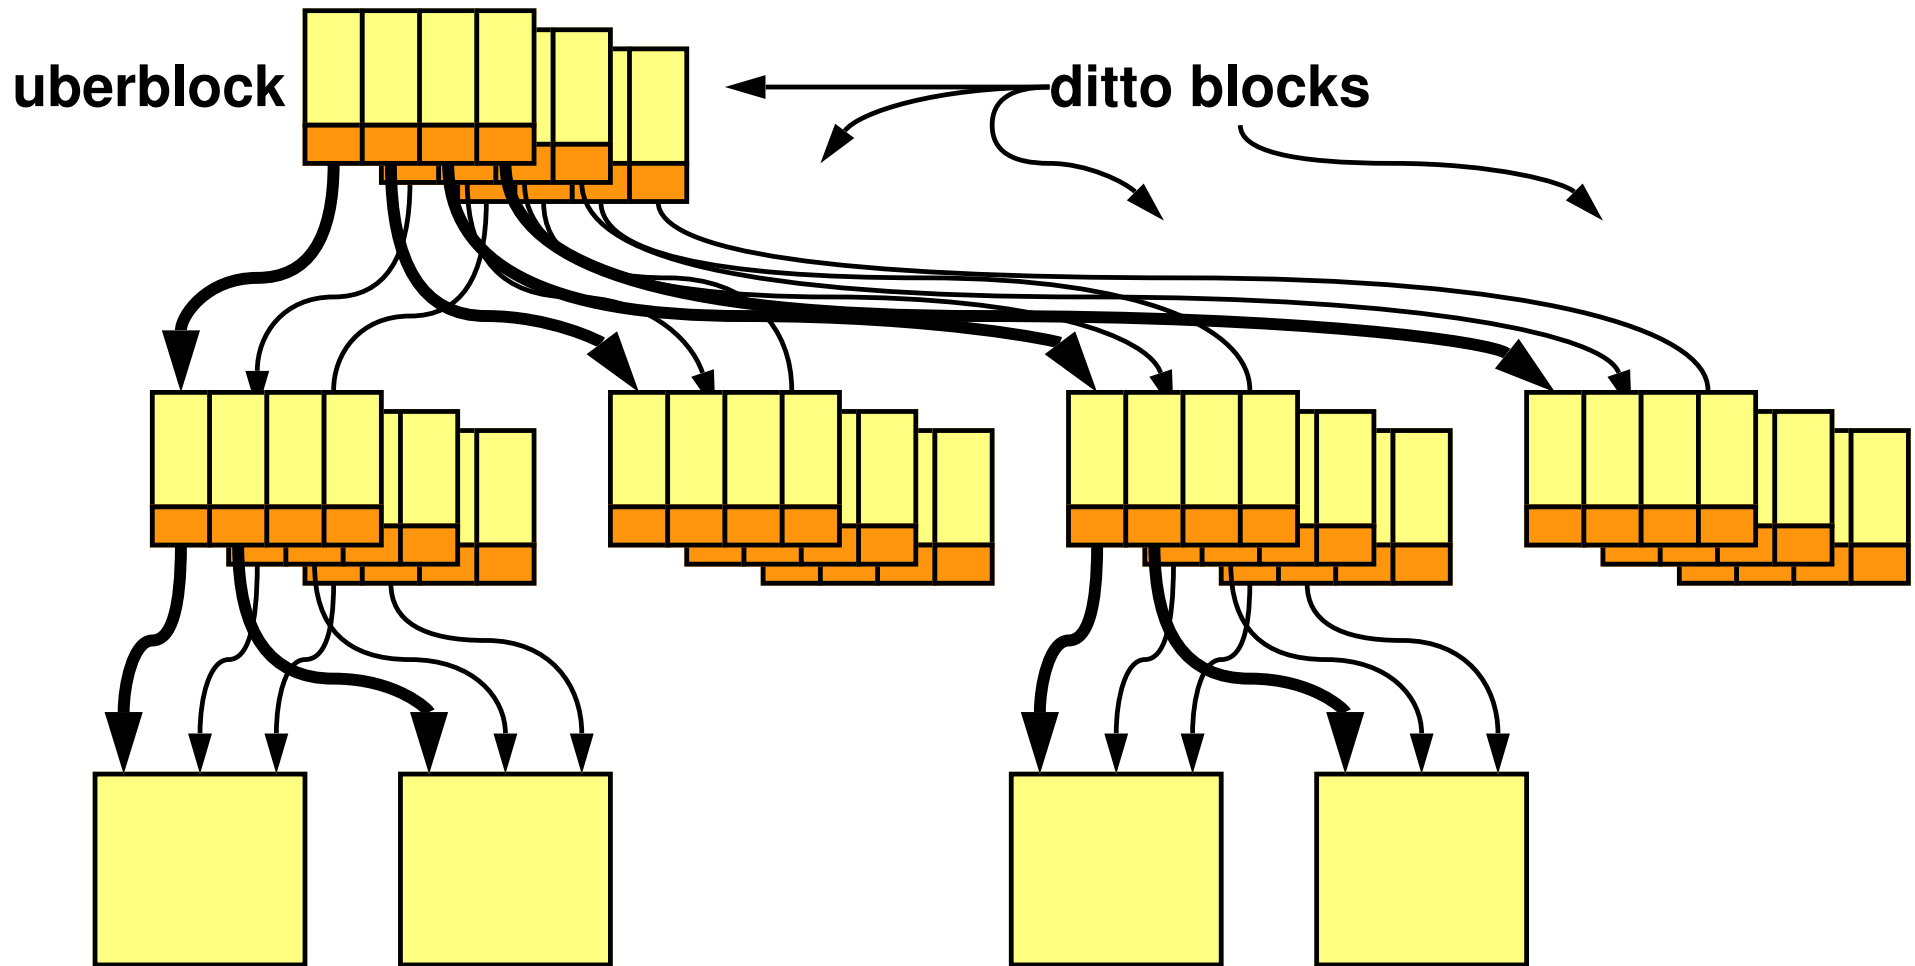
Storage Pool



Enter ZFS



Shadow-Page Tree (with a twist ...)



Storage Pool Allocator

**Data Management Unit
(DMU)**

Storage Pool Allocator (SPA)

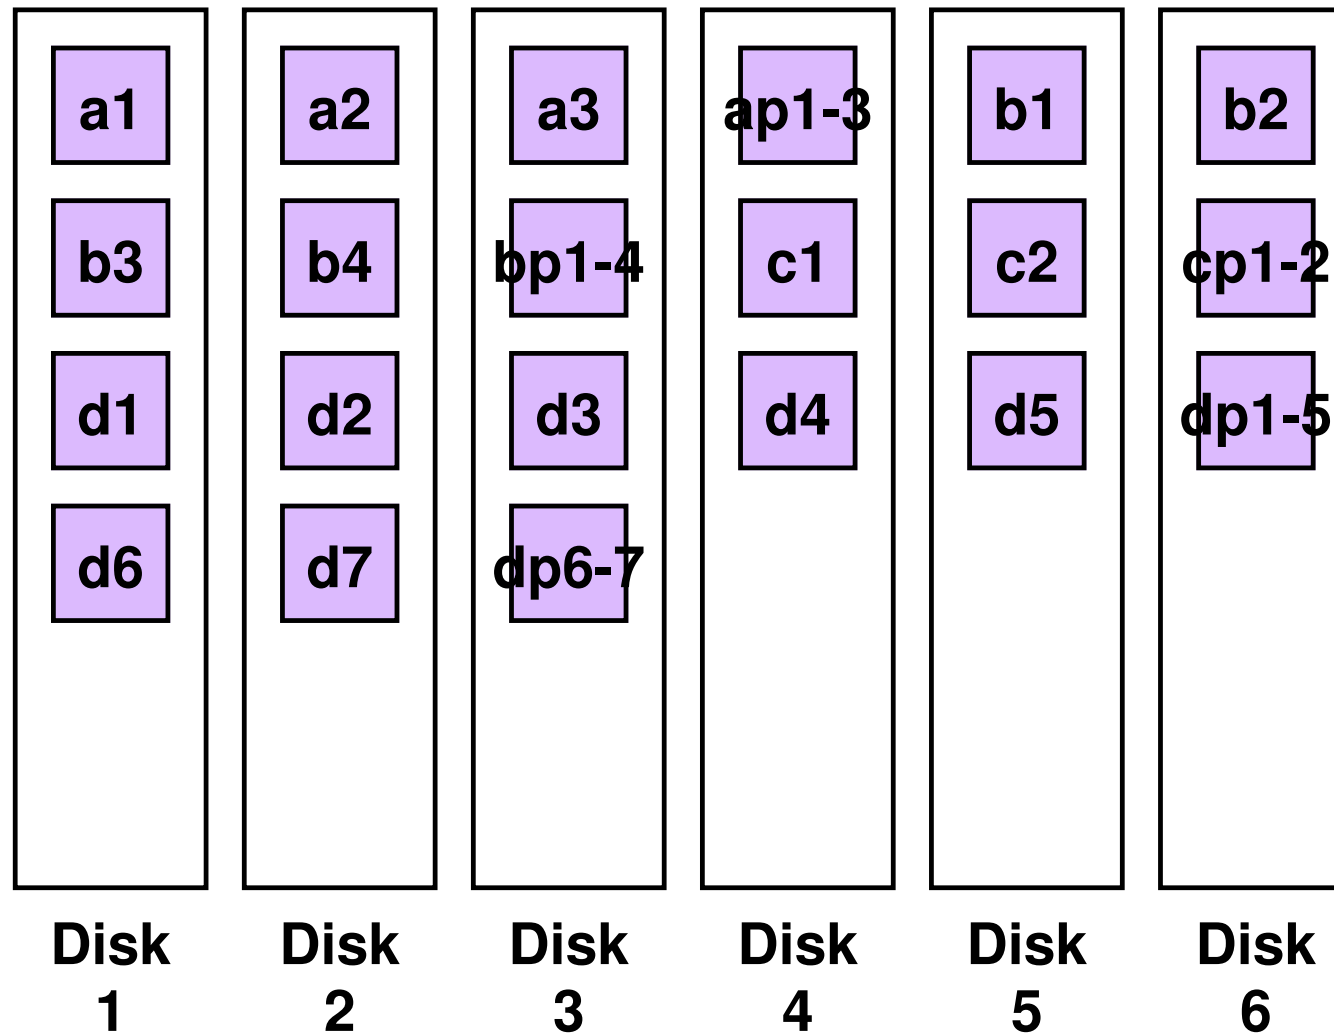


**Mirroring,
spanning, or
RAID**




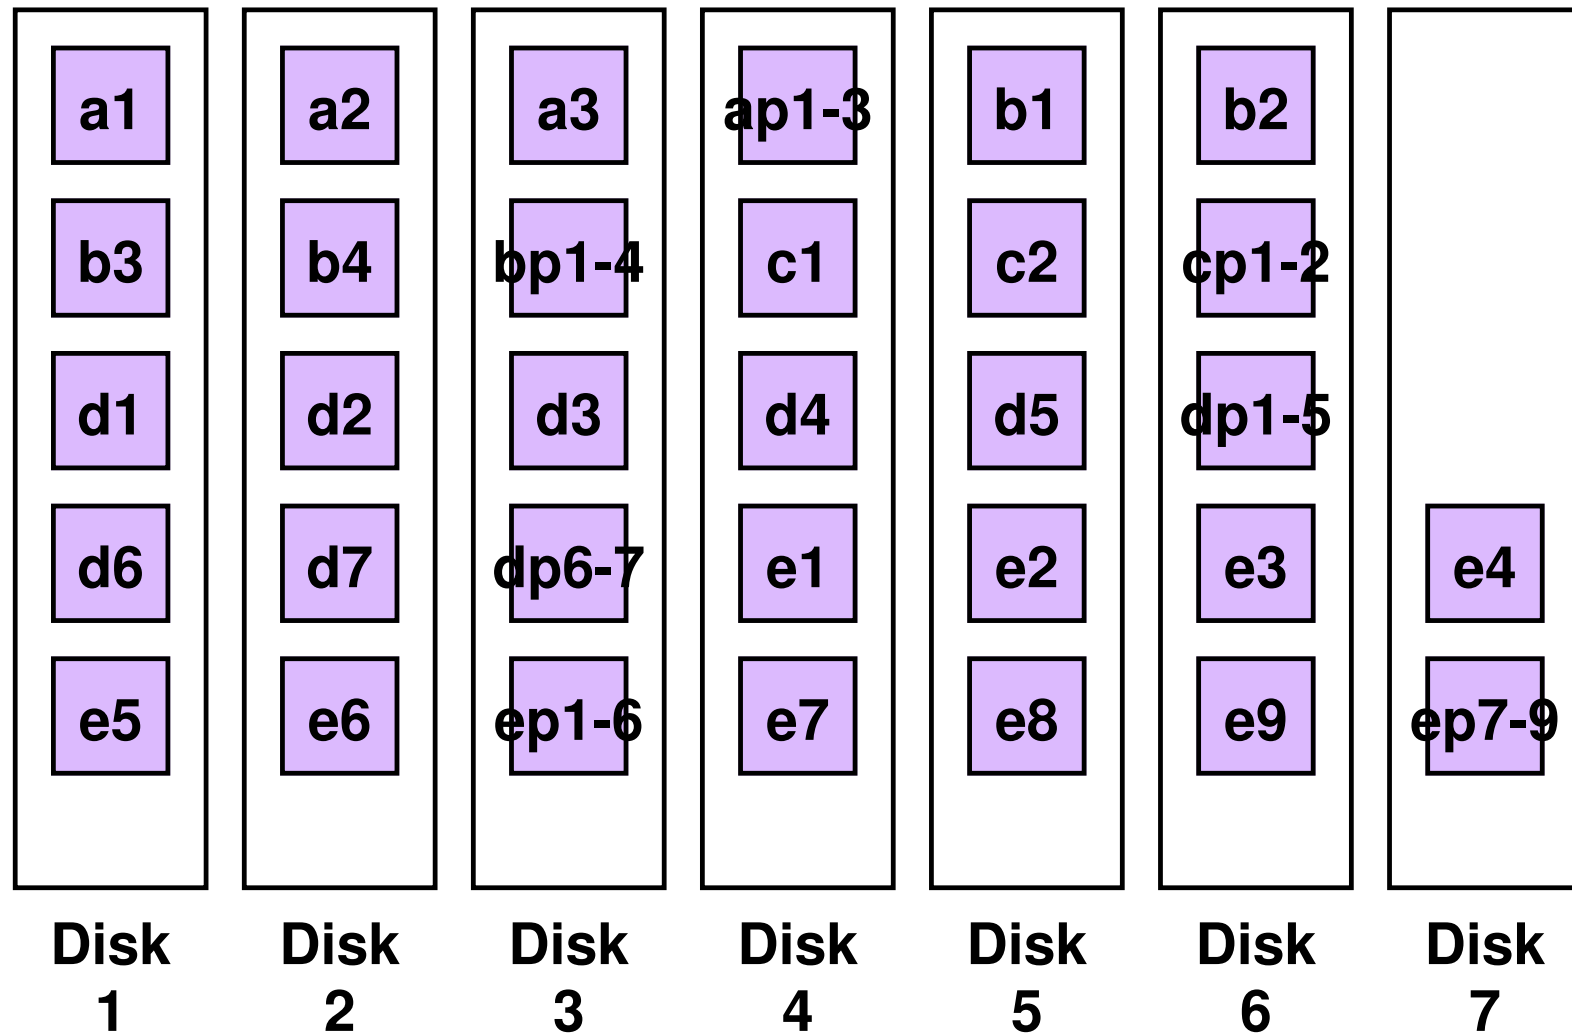
RAID-Z

➡ Software Dynamic Striping




RAID-Z

➡ Adding a Disk




Scenarios

- ➡ **Power failure at inopportune moment**
 - "live data" is not modified
 - single lost write can be recovered
- ➡ **Obscure bug in controller firmware or OS**
 - detected by checksum in pointer
- ➡ **Sysadmin accidentally scribbled on one drive**
 - detected and repaired
- ➡ **Out of disk space**
 - add to the pool; SPA will cope
- ➡ **Out of address space**
 - 2¹²⁸ is big
 - 1 address per cubic yard of a sphere bounded by the orbit of Neptune




And There's More ...

- ➡ Adaptive replacement cache
- ➡ Advanced prefetching




LRU Caching

- ➡ **LRU cache holds n least-recently-used disk blocks**
 - ▬ working sets of current processes
- ➡ **New process reads n -block file sequentially**
 - ▬ cache fills with this file's blocks
 - ▬ old contents flushed
 - ▬ new cache contents never accessed again




(Non-Adaptive) Solution

- ➡ Split cache in two
 - ▬ half of it is for blocks that have been referenced exactly once
 - ▬ half of it is for blocks that have been referenced more than once
- ➡ Is 50/50 split the right thing to do?



Adaptive Replacement Cache



t1 ; b1:

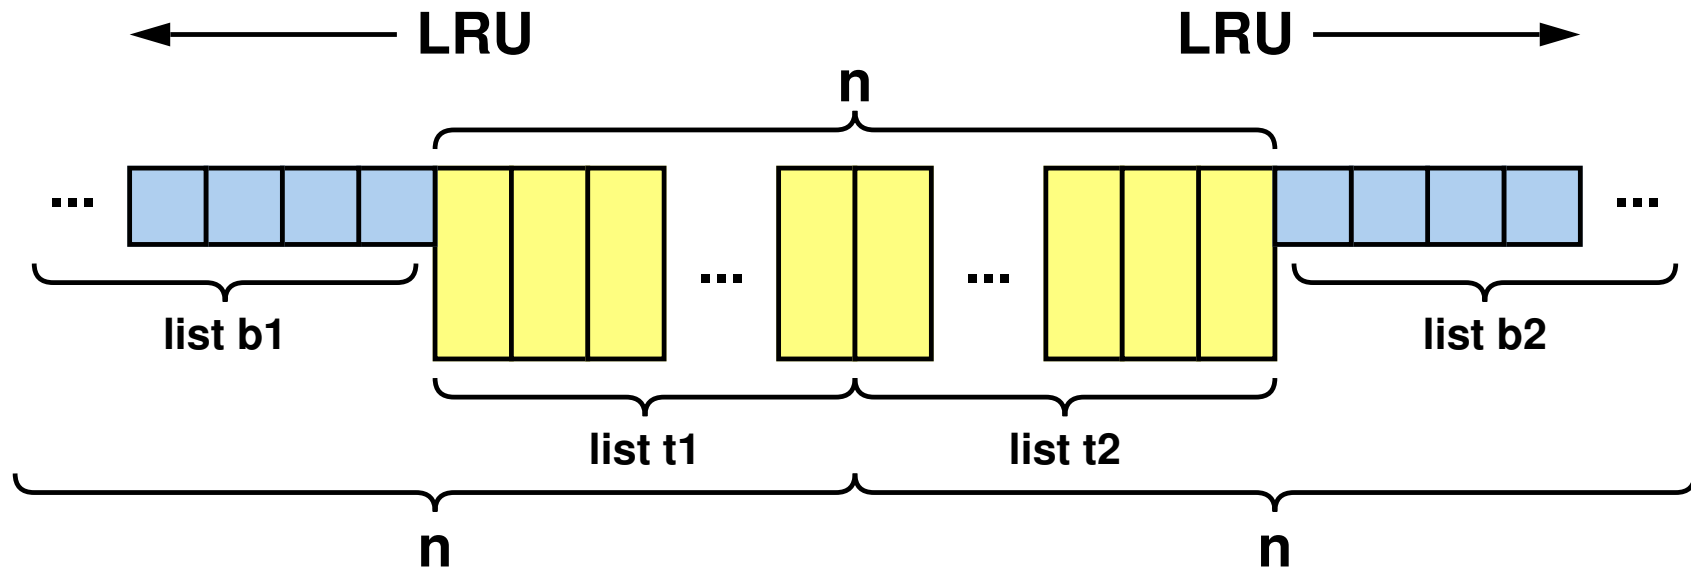
- ▢ LRU list of blocks referenced once
- ▢ t1 list (most recently used) contain contents
- ▢ b1 list (least recently used) contain just references

t2 ; b2:

- ▢ LRU list of blocks referenced more than once
- ▢ t2 list (most recently used) contain contents
- ▢ b2 list (least recently used) contain just references



Adaptive Replacement Cache

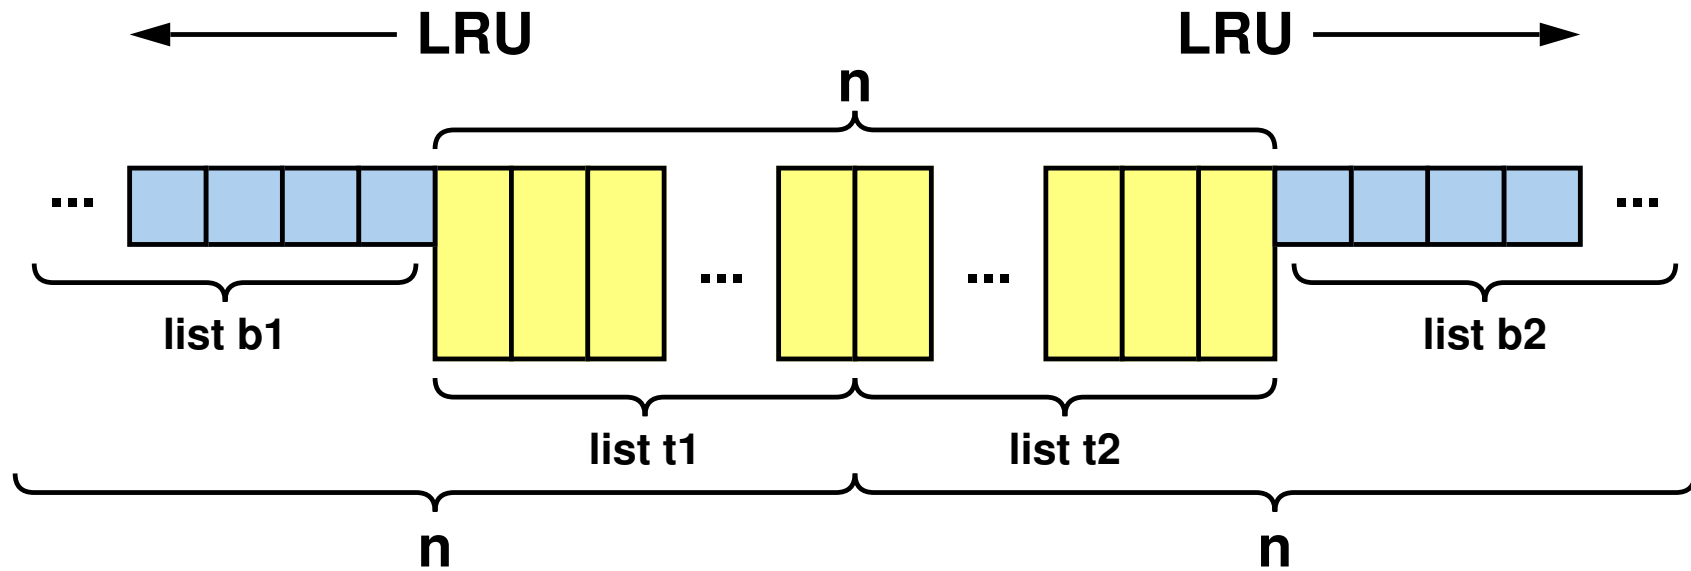


cache miss:

- ▢ if t1 is full
 - ▢ evict LRU(t1) and make it MRU(b1)
- ▢ referenced block becomes MRU(t1)



Adaptive Replacement Cache




cache hit:

- ▢ if in t1 or t2, block becomes MRU(t2)
- ▢ otherwise
 - ▢ if block is referred to by b1, increase t1 space at expense of t2
 - ▢ otherwise
 - increase t2 space at expense of t1
 - ▢ if t1 is full, evict LRU(t1) and make it MRU(b1)
 - ▢ if t2 is full, evict LRU(t2) and make it MRU(b2)
 - ▢ insert block as MRU(t2)



Prefetch

- ➡ FFS prefetch
- keeps track of last block read by each process
 - fetches block $i+1$ if current block is i and previous was $i-1$
 - chokes on
 - diff file1 file2



zfetch

- ➡ Tracks multiple prefetch streams
- ➡ Handles four patterns
 - ▢ forward sequential access
 - ▢ backward sequential access
 - ▢ forward strided access
 - iterating across columns of matrix stored by columns
 - ▢ backward strided access

