

Ch 7: Memory Management

Bill Cheng

<http://merlot.usc.edu/cs402-s16>





Copyright © William C. Cheng



Copyright © William C. Cheng

Memory Management



- Challenges
 - what to do when you run out of space?
 - protection

The Address-Space Concept

- Protect processes from one another
- Protect the OS from user processes
- Provide efficient management of available storage
- Illusion of large memory
- Sharing (code, data, communication)
- new abstraction (such as pipes, memory-mapped files)




Copyright © William C. Cheng




Copyright © William C. Cheng


Virtual Address



- Who uses *virtual address*?
 - user processes
 - kernel processes
 - pretty much every piece of software
- You would use a virtual address to address any memory location in the 32-bit address space
- Anything uses *physical address*?
 - nothing in OS
 - well, the hardware uses
 - physical address (and the processor is hardware)
 - the OS *manages* the



Virtual Address

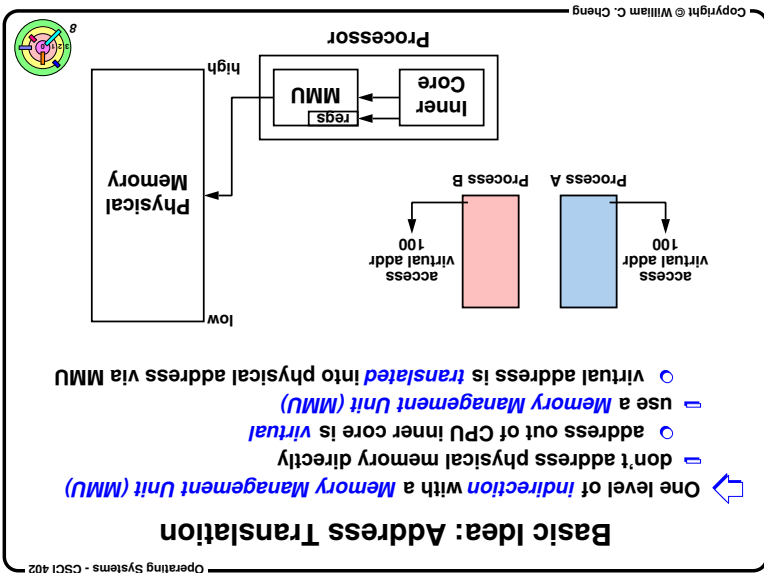
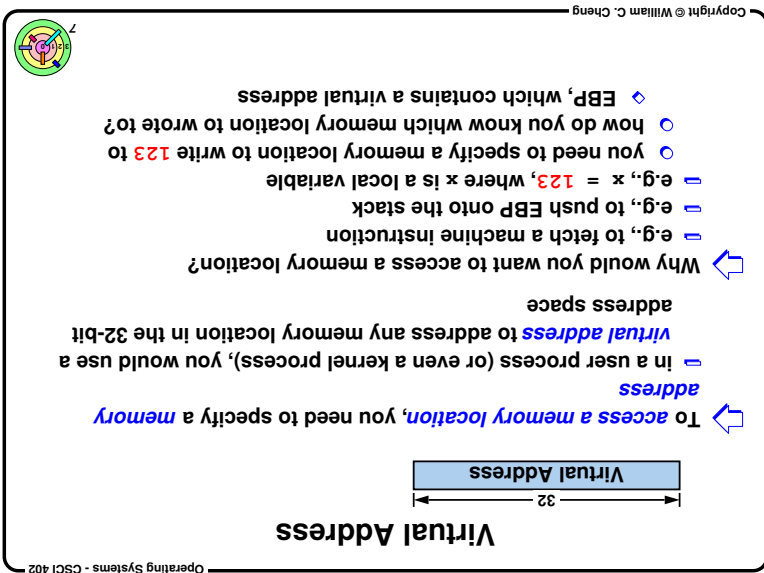
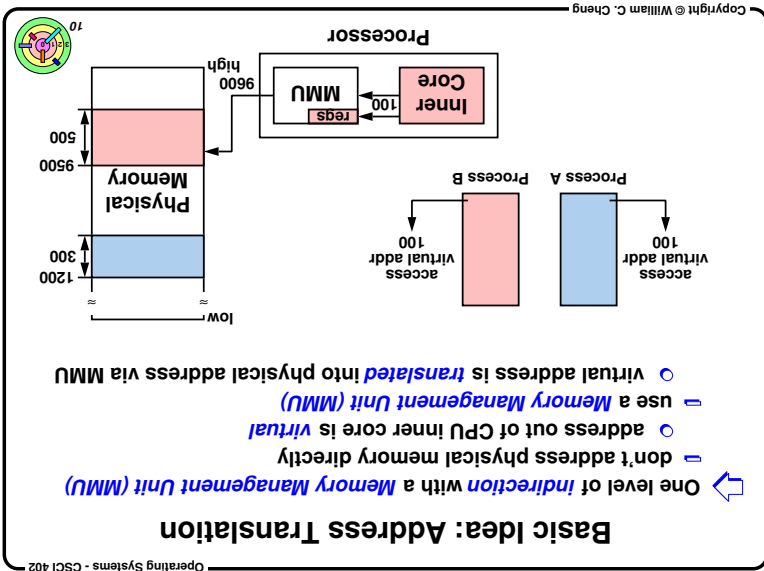
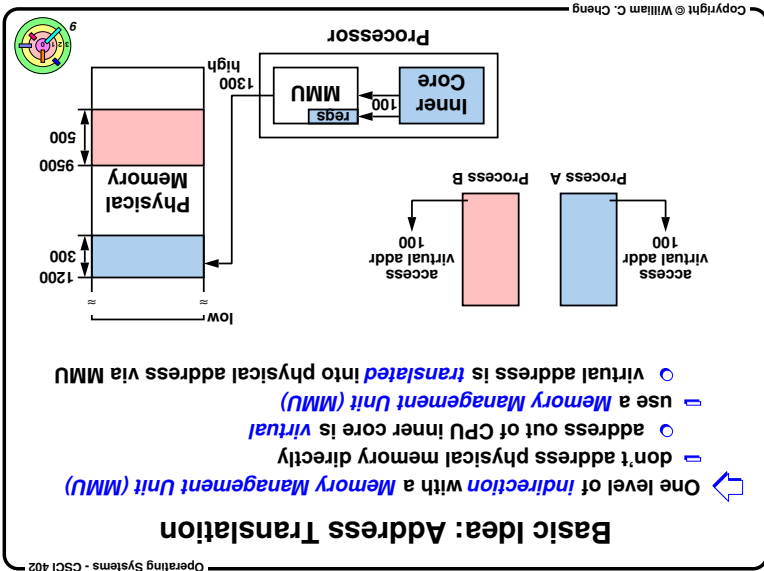
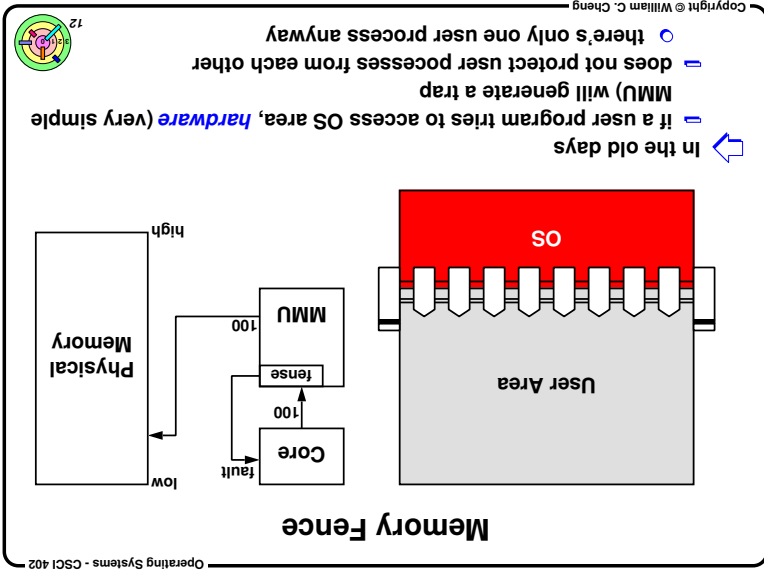
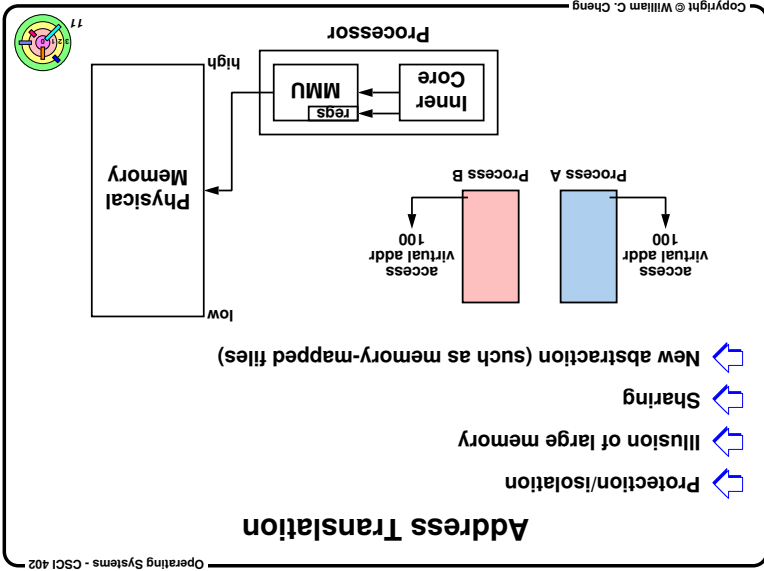


- To *access a memory location*, you need to specify a *memory address*
 - in a user process (or even a kernel process), you would use a *virtual address* to address any memory location in the 32-bit address space

- Why would you want to access a memory location?
 - e.g., to fetch a machine instruction
 - e.g., to push EBP onto the stack
 - you need to specify a memory location to store the content of EBP
 - how do you know which memory location to write to?
 - ESP, which contains a virtual address



Copyright © William C. Cheng



Copyright © William C. Cheng

17

Access Control With Segmentation

Access control / protection
→ read-only, read/write

MMU

Physical Memory

Core

Process A

Process B

Process C

stack

heap

data

code

shared

low

high

Can simply setup base and bounds

registers to share segments

Operating Systems - CSCI 402

Copyright © William C. Cheng

15

Base and Bounds Registers

Multiple user processes
→ OS maintains a pair of registers for each user process
→ **bounds register**: address space size of the user process
→ **base register**: start of physical memory for the user process
→ addresses **relative** to the **base register**
→ memory reference ≥ 0 and $<$ bounds, **independent of base** (this is known as "position independence")

MMU

Physical Memory

Core

Process A

Process B

Process C

stack

heap

data

code

shared

low

high

One pair of **base** and **bounds** registers **per segment**
→ code, data, heap, stack, and may be more
→ compiler compiles programs into segments

Operating Systems - CSCI 402

Copyright © William C. Cheng

13

Memory Fence and Overlays

What if the user program won't fit in memory?

→ use **overlays**
→ programmers (not the OS) have to keep track of which overlays is in physical memory and deal with the complexities of managing **overlays**

MMU

Physical Memory

Core

Process A

Process B

Process C

stack

heap

data

code

shared

low

high

Base and Bounds Registers

Operating Systems - CSCI 402

Copyright © William C. Cheng

16

Sharing Segments

Can simply setup base and bounds
→ registers to share segments

MMU

Physical Memory

Core

Process A

Process B

Process C

stack

heap

data

code

shared

low

high

Operating Systems - CSCI 402

Copyright © William C. Cheng

16

Generalization of Base and Bounds: Segmentation

One pair of **base** and **bounds** registers **per segment**
→ code, data, heap, stack, and may be more
→ compiler compiles programs into segments

MMU

Physical Memory

Core

Process A

Process B

Process C

stack

heap

data

code

shared

low

high

Operating Systems - CSCI 402

Copyright © William C. Cheng

14

Base and Bounds Registers

Multiple user processes
→ OS maintains a pair of registers for each user process
→ **bounds register**: address space size of the user process
→ **base register**: start of physical memory for the user process
→ addresses **relative** to the **base register**
→ memory reference ≥ 0 and $<$ bounds, **independent of base** (this is known as "position independence")

MMU

Physical Memory

Core

Process A

Process B

Process C

stack

heap

data

code

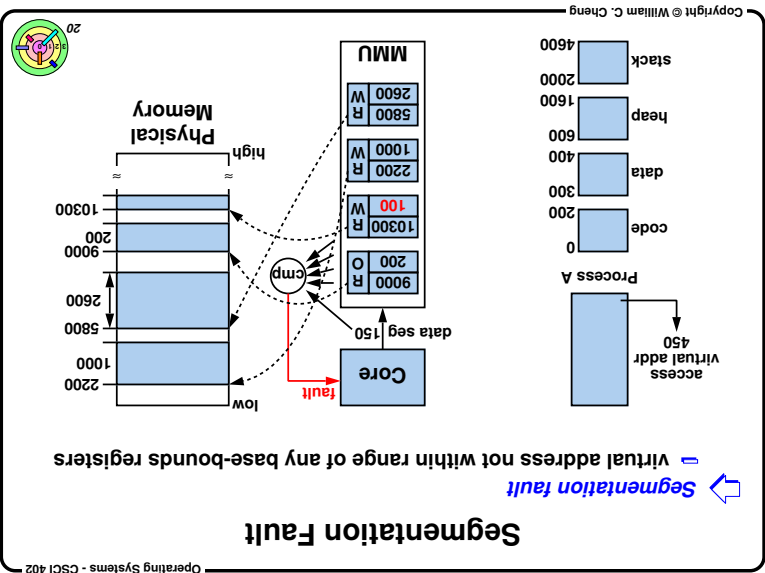
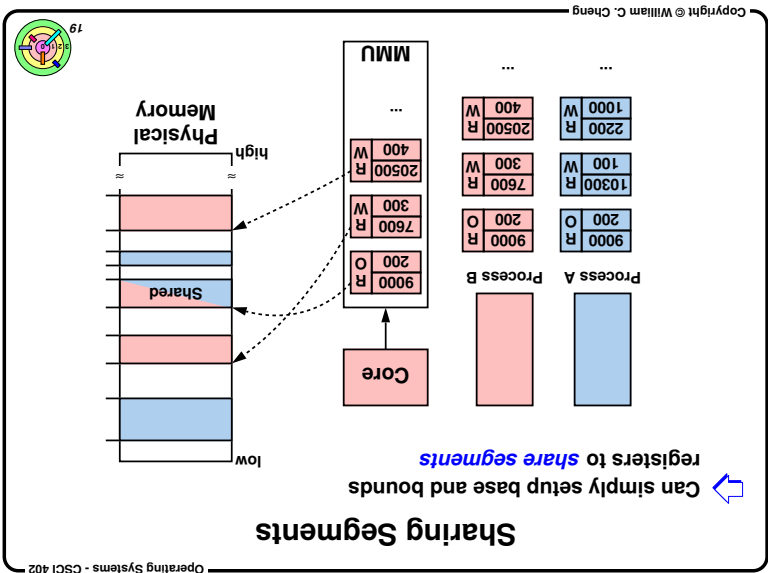
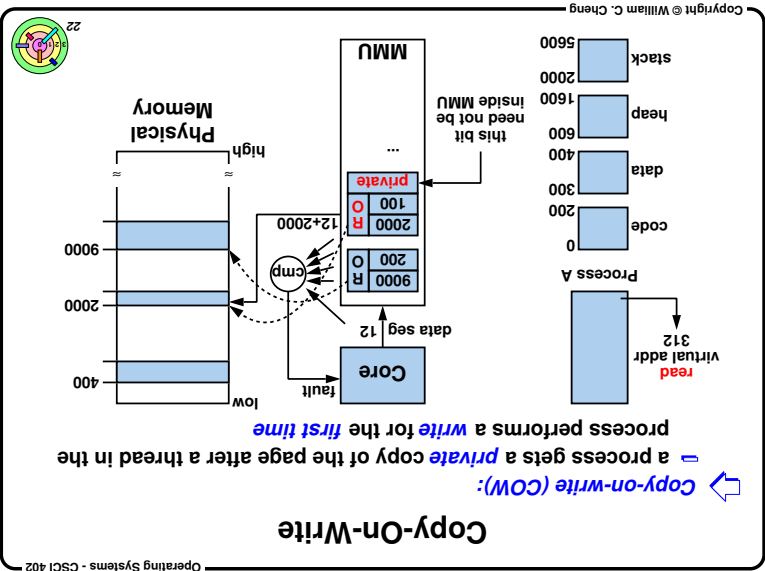
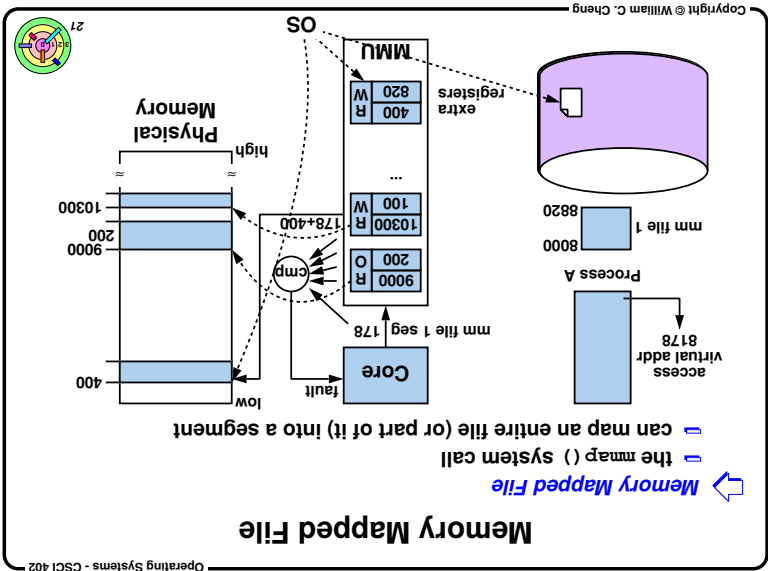
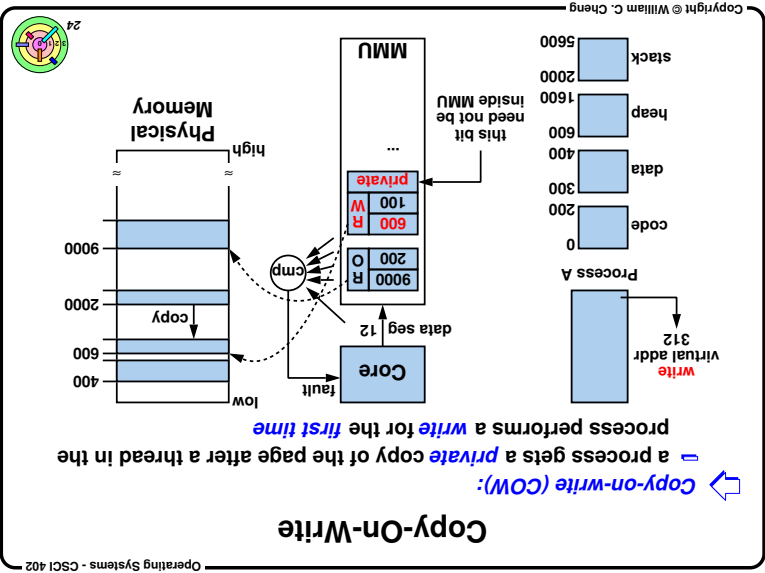
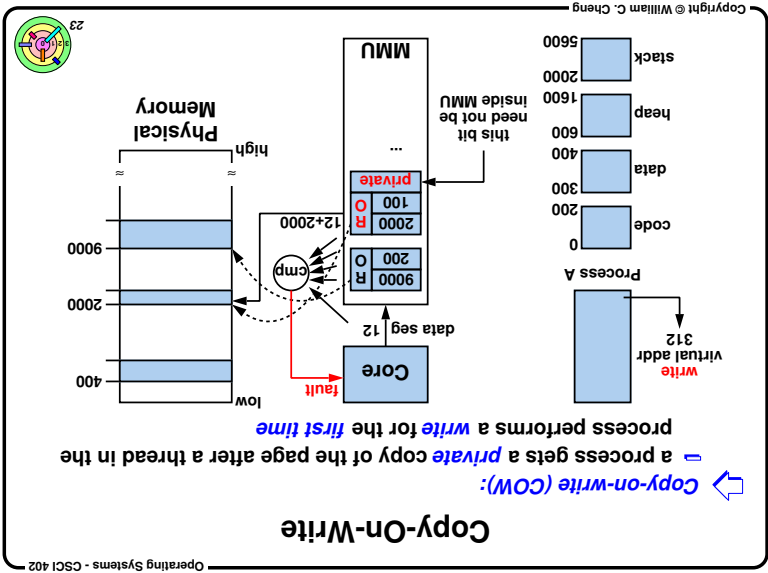
shared

low

high

Base and Bounds Registers

Operating Systems - CSCI 402



Copyright © William C. Cheng

29

stack 5600
2000
1600
heap 600
400
data 300
200
code 100
0
Process A
virtual addr 678
access
Core
heap seg 78
fault
cmp
MMU
V=1 5800
V=1 2600
V=0 1000
V=1 100
V=1 10300
V=1 200
V=1 9000
Physical Memory
high
low
10300
200
9000
2600
5800
1000

control bits

➡ No space for new segment, make room by swapping out a segment
➡ use a *validity* bit for each segment (in addition to access control bits)

Swapping / Backing Store

Operating Systems - CSCI 402

Copyright © William C. Cheng

27

stack 5600
2000
1600
heap 600
400
data 300
200
code 100
0
Process A
virtual addr 678
access
Core
heap seg 78
fault
cmp
MMU
V=1 5800
V=1 2600
V=1 1000
V=1 100
V=1 10300
V=1 200
V=1 9000
Physical Memory
high
low
10300
200
9000
2600
5800
1000

control bits

➡ No space for new segment, make room by swapping out a segment
➡ use a *validity* bit for each segment (in addition to access control bits)

Swapping / Backing Store

Operating Systems - CSCI 402

Copyright © William C. Cheng

25

stack 5600
2000
1600
heap 600
400
data 300
200
code 100
0
Process A
virtual addr 312
write
Core
data seg 12
fault
cmp
MMU
private 600
R 100
W 12+600
R 200
R 9000
Physical Memory
high
low
9000
2000
600
400

control bits

➡ a process gets a *private* copy of the page after a thread in the process performs a *write* for the *first time*
➡ use a *validity* bit for each segment (in addition to access control bits)

Copy-On-Write

Operating Systems - CSCI 402

Copyright © William C. Cheng

30

stack 5600
2000
1600
heap 600
400
data 300
200
code 100
0
Process A
virtual addr 678
access
Core
heap seg 78
fault
cmp
MMU
V=0 2600
V=0 1000
V=0 100
V=0 200
V=0 9000
Physical Memory
high
low
10300
200
9000
2600
5800
1000

control bits

➡ No space for new segment, make room by swapping out a segment
➡ use a *validity* bit for each segment (in addition to access control bits)

Swapping / Backing Store

Operating Systems - CSCI 402

Copyright © William C. Cheng

28

stack 5600
2000
1600
heap 600
400
data 300
200
code 100
0
Process A
virtual addr 678
access
Core
heap seg 78
fault
cmp
MMU
V=1 5800
V=1 2600
V=0 1000
V=1 100
V=1 10300
V=1 200
V=1 9000
Physical Memory
high
low
10300
200
9000
2600
5800
1000

control bits

➡ No space for new segment, make room by swapping out a segment
➡ use a *validity* bit for each segment (in addition to access control bits)

Swapping / Backing Store

Operating Systems - CSCI 402

Copyright © William C. Cheng

26

stack 5600
2000
1600
heap 600
400
data 300
200
code 100
0
Process A
virtual addr 678
access
Core
heap seg 78
fault
cmp
MMU
V=1 5800
V=1 2600
V=1 1000
V=1 100
V=1 10300
V=1 200
V=1 9000
Physical Memory
high
low
10300
200
9000
2600
5800
1000

control bits

➡ No space for new segment, make room by swapping out a segment
➡ use a *validity* bit for each segment (in addition to access control bits)

Swapping / Backing Store

Operating Systems - CSCI 402

7.2 Hardware Support for Virtual Memory

- Forward-Mapped Page Tables
- Linear Page Tables
- Hashes Page Tables
- Translation Lookaside Buffers
- 64-Bit Issues
- Virtualization



Copyright © William C. Cheng

Operating Systems - CSCI 402

Structuring Virtual Memory

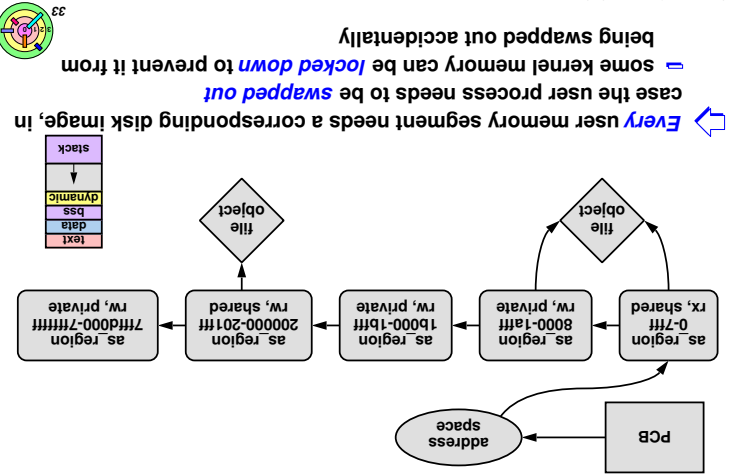
- Segmentation (just discussed)
 - divide the address space into variable-size segments (typically each corresponding to some logical unit of the program, such as a module or subroutine)
 - external fragmentation possible
 - "first-fit" is slow
 - not very common these days
- Paging
 - divide the address space into fixed-size pages
 - internal fragmentation possible



Copyright © William C. Cheng

Operating Systems - CSCI 402

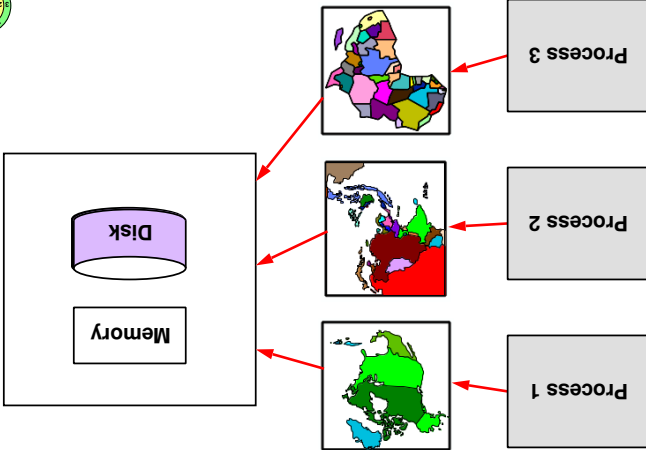
Swapping / Backing Store



Copyright © William C. Cheng

Operating Systems - CSCI 402

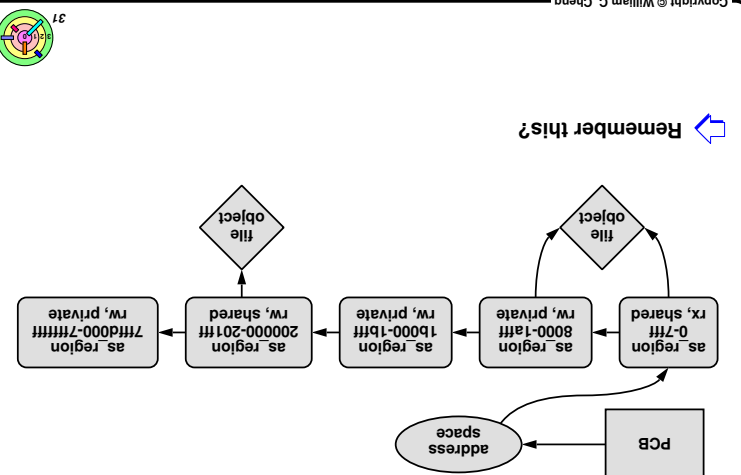
Virtual Memory



Copyright © William C. Cheng

Operating Systems - CSCI 402

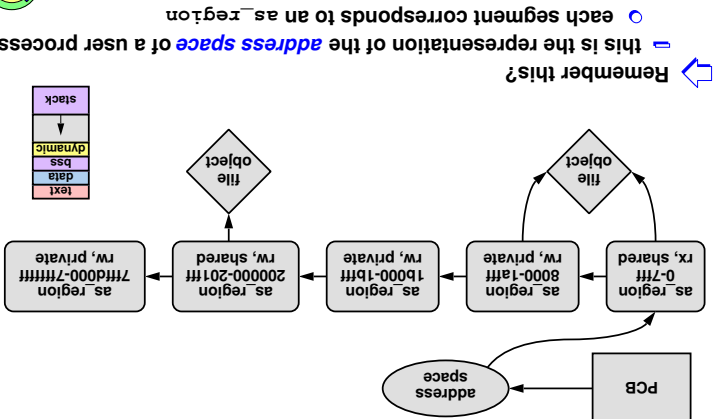
Swapping / Backing Store



Copyright © William C. Cheng

Operating Systems - CSCI 402

Swapping / Backing Store



Copyright © William C. Cheng

Operating Systems - CSCI 402

Copyright © William C. Cheng

41

Basic (Two-level) Page Tables

VA: 20 12 4KB Physical Page

Page # Offset

Page Table

vp=0 vp=1 vp=2

Page Table Entry

Physical Page #

Extension (PAE)

- in x86: **Physical Address** more than 20 bits long
- "physical page number" would be more than 4GB physical memory
- A **page table entry** is analogous to the **base and bounds registers** approach

Copyright © William C. Cheng

39

Page Frames

Physical Memory

4KB Pages

0 1 2 3 4

Objects

Page Frame

It is important to be able to perform both **forward lookup** and **reverse lookup**

- given a virtual address of a process, find page frame
- given a page frame, find processes and virtual addresses that uses this page frame
- we need x page frame
- data structure is a bit involved
- see kernel 3 FAQ
- the kernel must use a **virtual address** to write into a physical page or read the content of a physical page

Copyright © William C. Cheng

37

Paging

Map **fixed-size pages** into physical memory (**into physical pages**)

- address space is divided into pages
- indexed by virtual page number
- physical memory is divided into pages (of the same size)
- indexed by physical page number
- need a lookup table to map **virtual page numbers** to **physical page numbers**

Ex: 1GB of physical memory with 4KB pages

- 2^{18} physical pages
- an address (either physical or virtual) is **page-aligned** if its least significant 12 bits are all zero

Many **hardware** mapping techniques

1GB Physical Memory

page 0	4KB Pages
page 1	4KB Pages
page 2	4KB Pages
...	...
page 262141	4KB Pages
page 262142	4KB Pages
page 262143	4KB Pages

Many **hardware** mapping techniques

- MMU** and **page table** (mostly in software)
- translation lookaside buffers (TLB)**

Copyright © William C. Cheng

42

Basic (Two-level) Page Tables

VA: 20 12 4KB Physical Page

Page # Offset

Page Table

vp=0 vp=1 vp=2

Page Table Entry

Physical Page #

Location of page table is part of the **context** for a **process**

- if virtual address is used, need to translate it to a physical address before giving it to the CPU
- for x86, the **physical address** of the page table is stored in the CR3 register

Copyright © William C. Cheng

40

Basic (Two-level) Page Tables

VA: 20 12 4KB Physical Page

Page # Offset

Page Table

vp=0 vp=1 vp=2

Page Table Entry

Physical Page #

- V**: validity (or "present")
- M**: modified
- R**: reference (set when page is referenced)
- Prot**: what type of access is allowed in general, can have more bits and include **"no access"** in x86, only one bit
- no shared/private bit here

Copyright © William C. Cheng

38

Page Frames

Physical Memory

4KB Pages

0 1 2 3 4

Page Frame

Page

A **page frame** data structure / object is used to maintain information about physical pages and their association with important kernel data structures

- contains a **physical page number**
- there is a **one-to-one mapping** between page frames and physical pages
- we use "page frame" and "physical page" interchangeably

Page fault

- page table does not have the requested address
- OS finds a free page frame
- OS loads the requested page from disk
- OS adjusts MMU and *restarts* user memory reference

Page Table

Copyright © William C. Cheng 47

Page-Table Size

- Consider a full 2^{32} -byte address space
- assume 4096-byte (2^{12} -byte) pages
- 4 bytes per page table entry
- the page table would consist of $2^{32}/2^{12} (= 2^{20})$ entries
- its size would be 2^{22} bytes (or 4 megabytes)

This is a general *scaling* problem

solutions:

- hash page tables* or hierarchy (*forward-mapped page tables*)
- virtual linear page tables* (i.e., page tables in virtual memory)

Page-Table Size

Copyright © William C. Cheng 48

Page fault

- page table does not have the requested address
- OS finds a free page frame
- OS loads the requested page from disk

Page Table

Copyright © William C. Cheng 45

Page fault

- page table does not have the requested address
- OS finds a free page frame
- OS loads the requested page from disk

Page Table

Copyright © William C. Cheng 46

Page Table

- A page table (*usually sits in physical memory*) is associated with each *process*
- OS has its page table as well
- Memory Management Unit (MMU) maps virtual address to physical address
- MMU got turned on some time during boot

Page Table

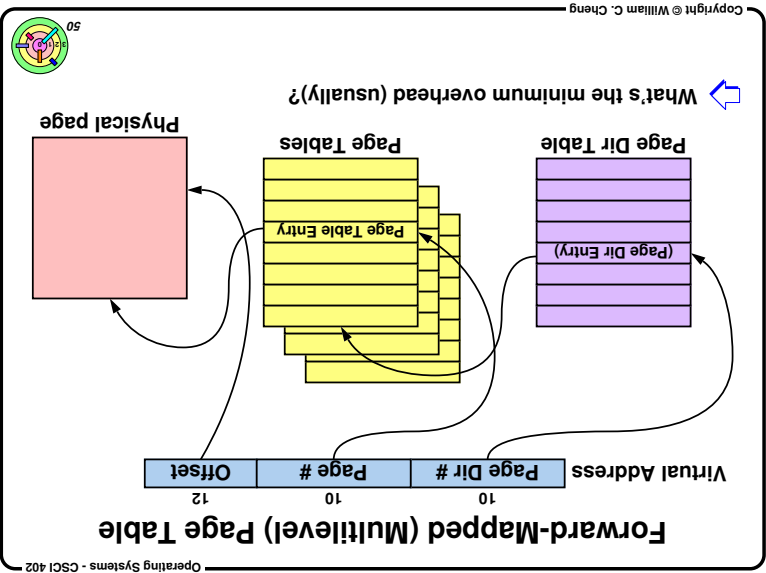
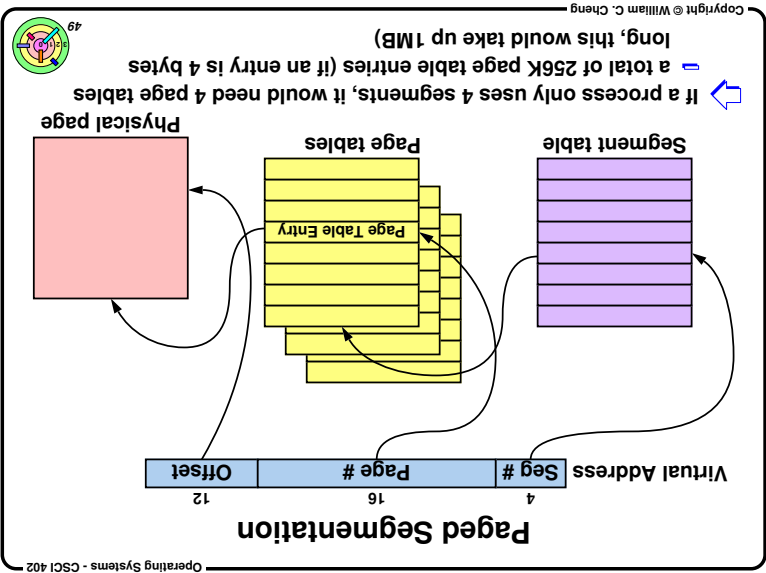
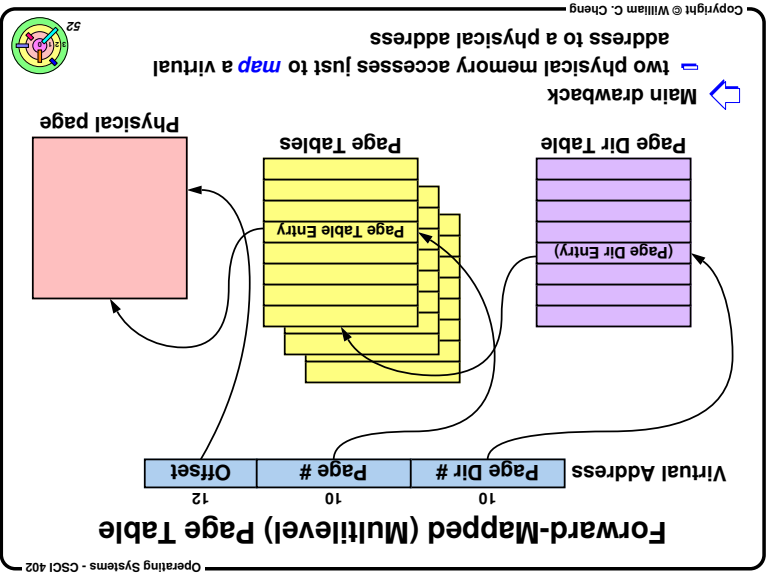
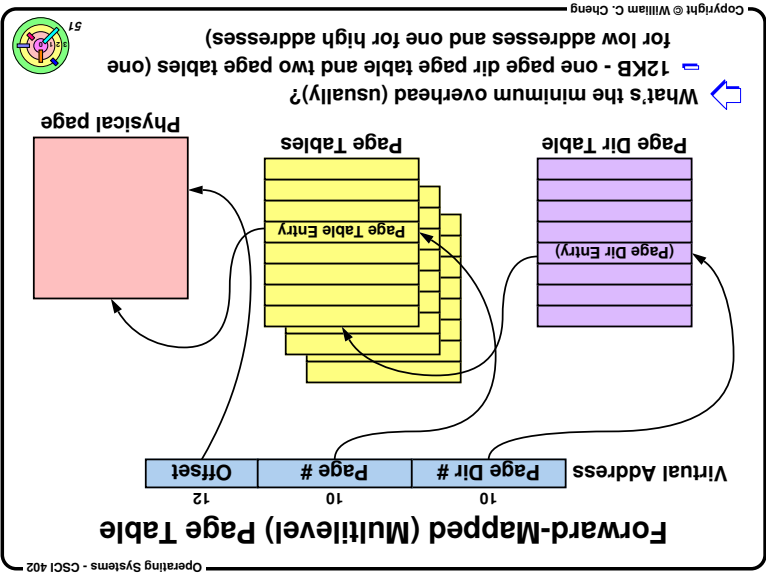
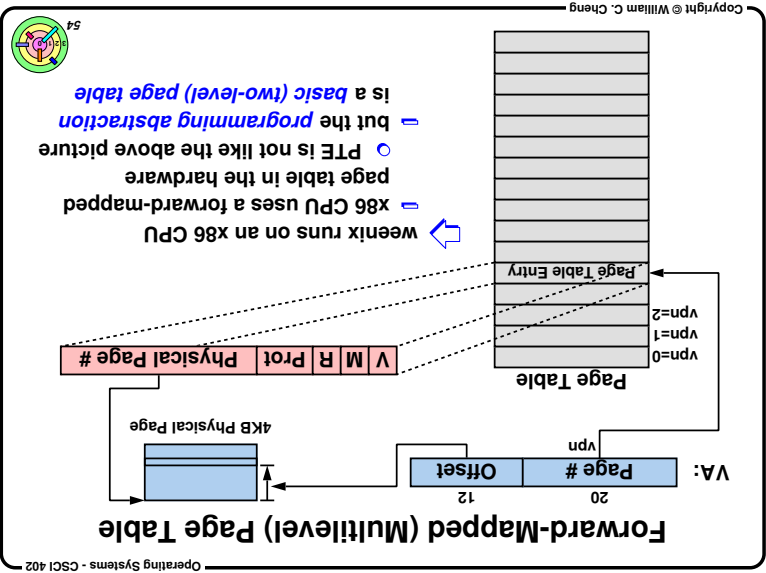
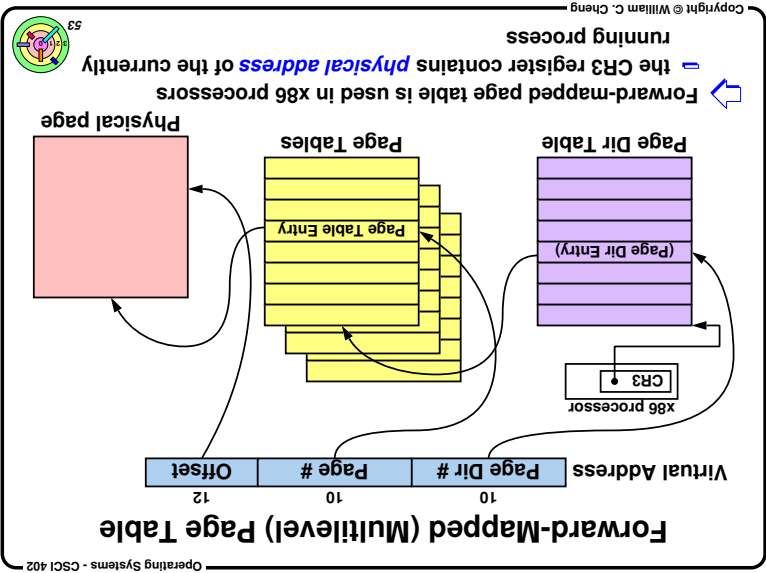
Copyright © William C. Cheng 43

Page Table

- page table does not have the requested address

Page Table

Copyright © William C. Cheng 44



7.2 Hardware Support for Virtual Memory


- Forward-Mapped Page Tables
- Linear Page Tables
- Hashes Page Tables
- Translation Lookaside Buffers
- 64-Bit Issues
- Virtualization



Copyright © William C. Cheng



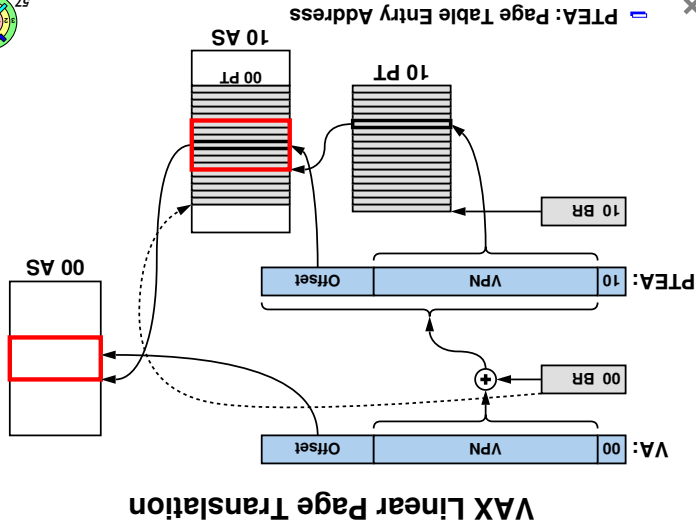
Copyright © William C. Cheng



Copyright © William C. Cheng



Copyright © William C. Cheng



Copyright © William C. Cheng



Copyright © William C. Cheng

Linear Page Table Management

- 00 and 01 page tables each require contiguous locations in 10 space
- with 512-byte pages, 8MB ($= 4\text{bytes} \times 2^{21}$) each:
- memory cost was \$40,000 per MB for the VAX
- maximum of 64 such page tables (to fill up 500MB of physical memory)
- (need room for other things, e.g. OS)
- Reduce size requirements with partial page tables
- length register* constrains size of each space




Copyright © William C. Cheng



Copyright © William C. Cheng

Traditional Unix with Linear PTs

- With traditional Unix, using a length register worked pretty well



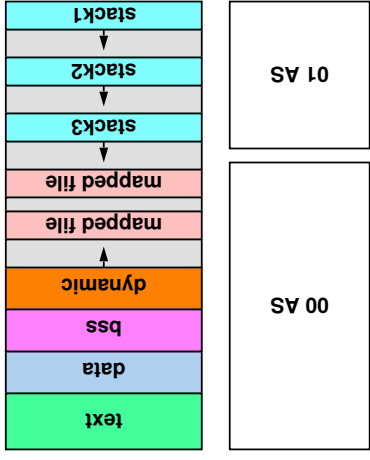
Copyright © William C. Cheng



Copyright © William C. Cheng

Modern Unix

- Not so well with modern Unix



Copyright © William C. Cheng



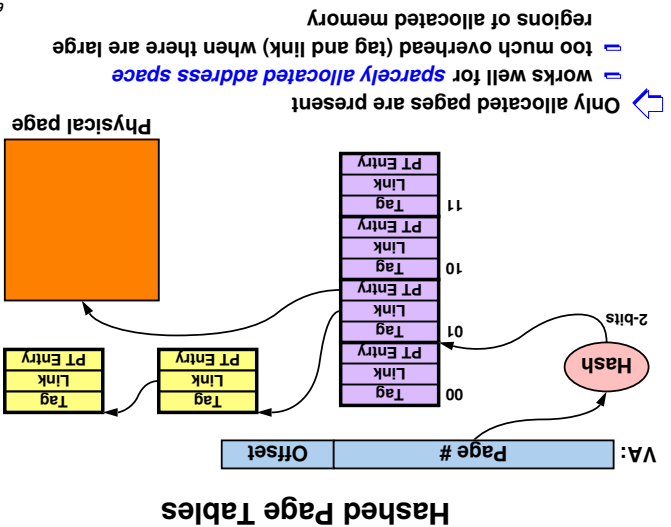
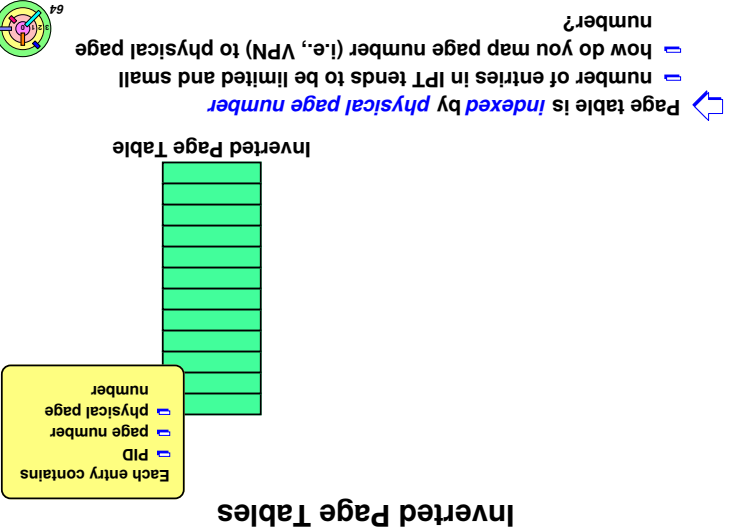

Copyright © William C. Cheng

7.2 Hardware Support for Virtual Memory

- Forward-Mapped Page Tables
- Linear Page Tables
- Hashes Page Tables*
- Translation Lookaside Buffers
- 64-Bit Issues
- Virtualization

7.2 Hardware Support for Virtual Memory

- Forward-Mapped Page Tables
- Linear Page Tables
- Hashes Page Tables
- Translation Lookaside Buffers*
- 64-Bit Issues
- Virtualization



Copyright © William C. Cheng

71

Translation Lookaside Buffers (TLB)

VA: Tag 14 Key 6 Offset 12

Tag Page Table Entry 0
Tag Page Table Entry 1
Tag Page Table Entry 2
⋮
Tag Page Table Entry 63

Ex: **two-way set-associative** cache (hardware cache) with 64 ($= 2^6$) lines

- number of bits in **key** tells you how many **lines** the TLB has
- the key is used as an index to look at it tells you which line to access the TLB
- it tells you which line to look at

amount of **set-associativity** is the "array size" in each line

the "**tag**" in the virtual address is compared against **all** tags in a line simultaneously (i.e., under the same key)

Copyright © William C. Cheng

72

Translation Lookaside Buffers (TLB)

VA: Tag 14 Key 6 Offset 12

Tag Page Table Entry 0
Tag Page Table Entry 1
Tag Page Table Entry 2
⋮
Tag Page Table Entry 63

Other TLB flavors

- direct mapping cache
- same as one-way set associative cache

Copyright © William C. Cheng

69

Translation Lookaside Buffers (TLB)

When a page table entry is modified, the OS must **flush** (invalidate) the corresponding TLB entry

Copyright © William C. Cheng

70

Translation Lookaside Buffers (TLB)

When a page table entry is modified, the OS must **flush** (invalidate) the corresponding TLB entry

When switching to a different **process**, must **flush** the **entire TLB** in x86, this can be achieved by setting the CR3 register

Copyright © William C. Cheng

67

Translation Lookaside Buffers (TLB)

Table lookup requires one memory access

- to access memory starting with a virtual address will take at least **two** memory accesses
- that's one access too many

If the processor has additional memory

- it can be used to **cache page table entries**
- Translation Lookaside Buffer** (or **TLB**)
- the TLB caches the **mapping** from virtual page number to physical page number (along with other information in the page table entry)
- hopefully, resolving a virtual address into a physical address will take one TLB lookup and one addition, inside the processor

TLB miss vs. page fault

- penalty for a TLB miss is $O(1)$ memory accesses
- penalty for a page fault is trap into the kernel

Copyright © William C. Cheng

68

Translation Lookaside Buffers (TLB)

Conceptually:

7.2 Hardware Support for Virtual Memory

- Forward-Mapped Page Tables
- Linear Page Tables
- Hashes Page Tables
- Translation Lookaside Buffers
- 64-Bit Issues
- Virtualization

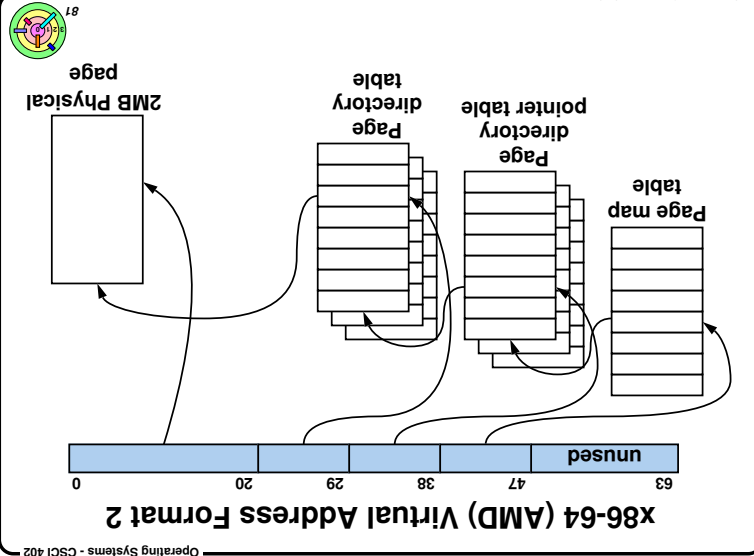


79



Copyright © William C. Cheng

Operating Systems - CSCI 402

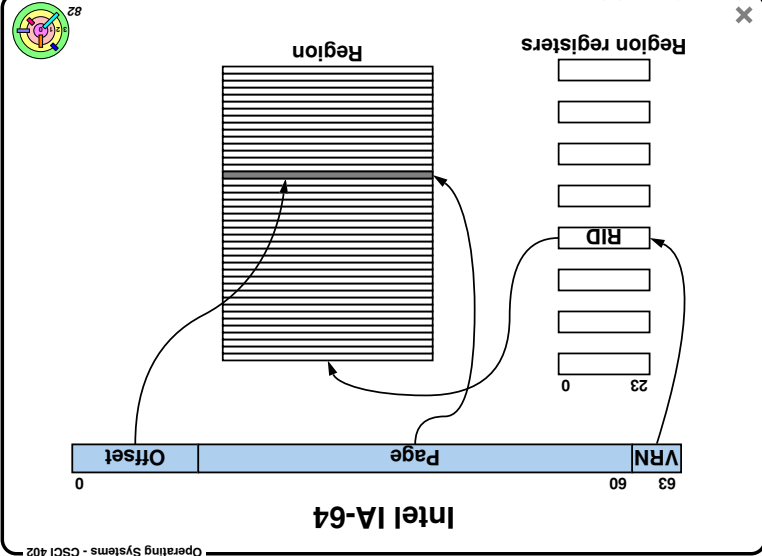


81



Copyright © William C. Cheng

Operating Systems - CSCI 402




82



Copyright © William C. Cheng

Operating Systems - CSCI 402

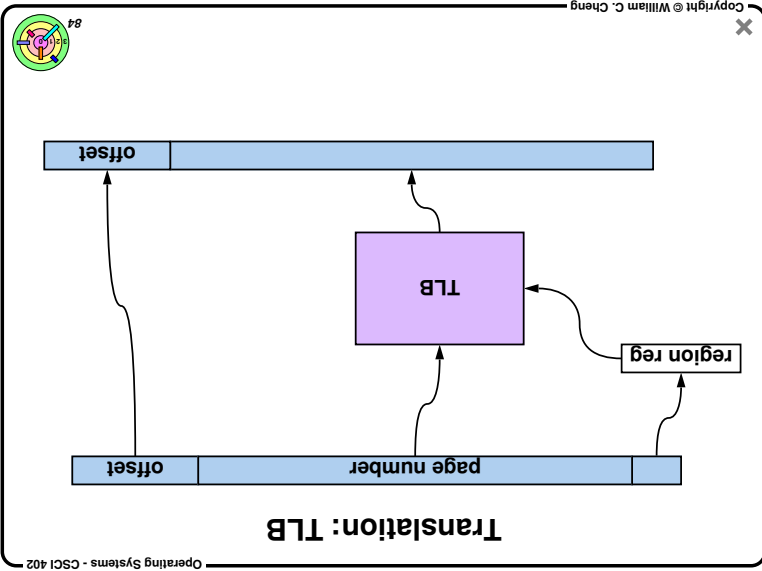


80



Copyright © William C. Cheng

Operating Systems - CSCI 402




84



Copyright © William C. Cheng

Operating Systems - CSCI 402



83



Copyright © William C. Cheng

Operating Systems - CSCI 402

