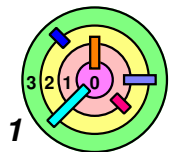



Housekeeping (Lecture 18 - 3/9,10/2016)

- ➡ Kernel 1 due at 11:45pm this Friday, 3/11/2016
 - if you have code from a previous semester, be very careful and **not copy any code from it**
 - it's best if you just get rid of it
- ➡ **Grading guidelines** is the only way we will grade
 - when running `faber_thread_test()`, you need to make sure that all the **exit codes** are correct
 - read the code to figure out what values to expect
 - you should be able to run commands after commands, etc.
 - if you are confused about "SELF-checks", please send me e-mail
- ➡ After submission, make sure you **Verify Your Kernel Submission**
 - tests in sections (C), (D), and (E) of the grading guidelines must run in the "foreground"
- ➡ This Friday, the TAs will give an introduction to Kernel 2
- ➡ By the way, **midterm** exam does cover **kernel 1**



Virtual Address





➡ Who uses *virtual address*?

- user processes
- kernel processes
- pretty much every piece of software


➡ You would use a virtual address to address any memory location in the 32-bit address space

➡ Anything uses *physical address*?


- nothing in OS
- well, the hardware uses physical address (and the processor is hardware)
- the OS *manages* the *physical address space*




Virtual Address



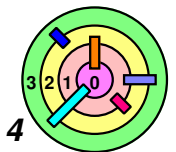
- ➡ To *access a memory location*, you need to specify a *memory address*
 - ▬ in a user process (or even a kernel process), you would use a *virtual address* to address any memory location in the 32-bit address space
- ➡ Why would you want to access a memory location?
 - ▬ e.g., to fetch a machine instruction
 - you need to specify a memory location to fetch from
 - how do you know which memory location to fetch from?
 - ◆ EIP (on an x86 machine), which contains a virtual address




Virtual Address




- ➡ To *access a memory location*, you need to specify a *memory address*
 - in a user process (or even a kernel process), you would use a *virtual address* to address any memory location in the 32-bit address space
- ➡ Why would you want to access a memory location?
 - e.g., to fetch a machine instruction
 - e.g., to push EBP onto the stack
 - you need to specify a memory location to store the content of EBP
 - how do you know which memory location to write to?
 - ◆ ESP, which contains a virtual address



Virtual Address

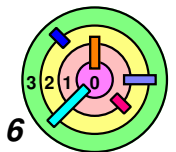
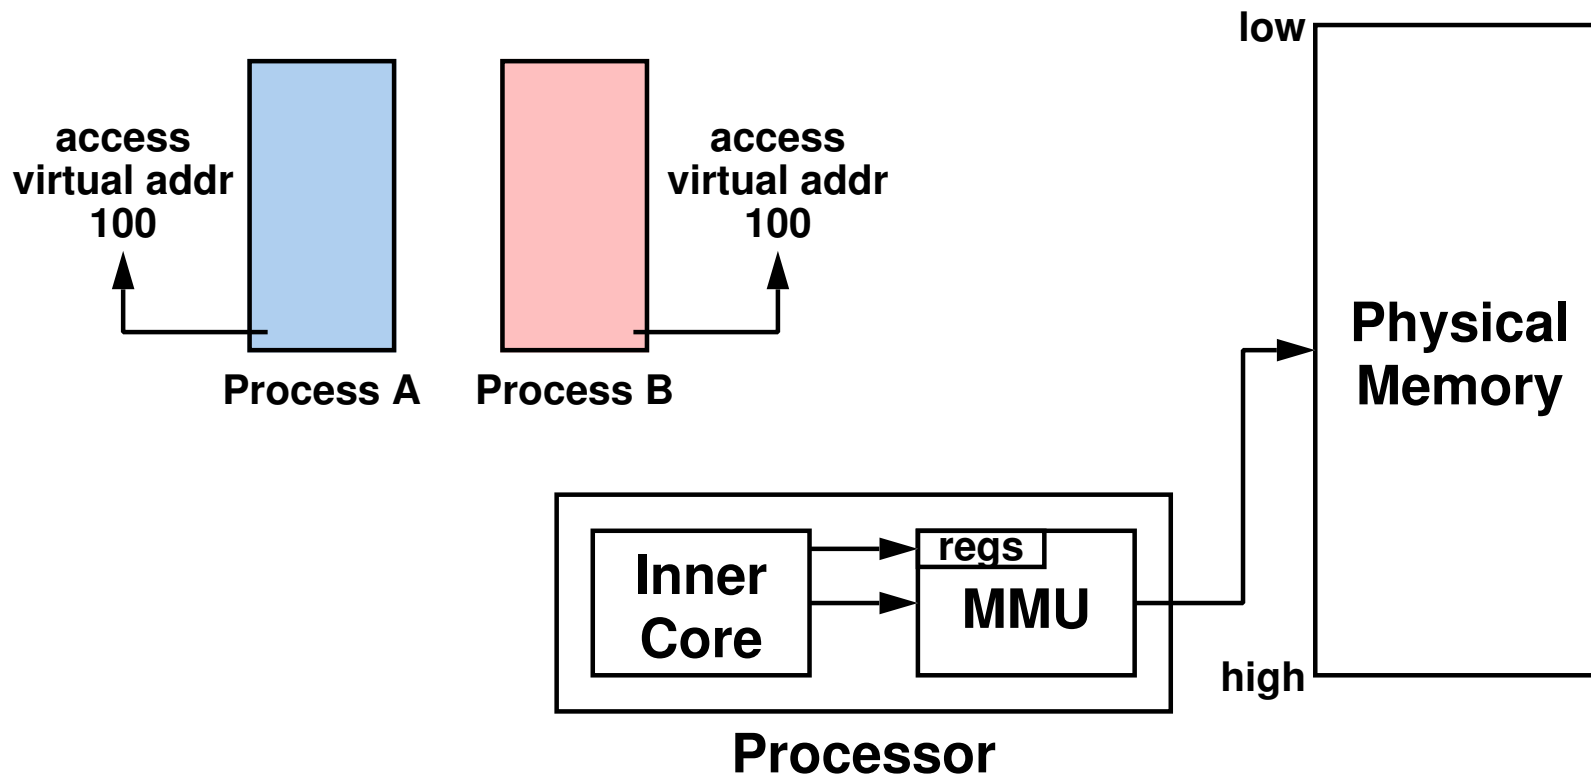


- ➡ To *access a memory location*, you need to specify a *memory address*
 - in a user process (or even a kernel process), you would use a *virtual address* to address any memory location in the 32-bit address space
- ➡ Why would you want to access a memory location?
 - e.g., to fetch a machine instruction
 - e.g., to push EBP onto the stack
 - e.g., $x = 123$, where x is a local variable
 - you need to specify a memory location to write **123** to
 - how do you know which memory location to write to?
 - ◆ EBP, which contains a virtual address



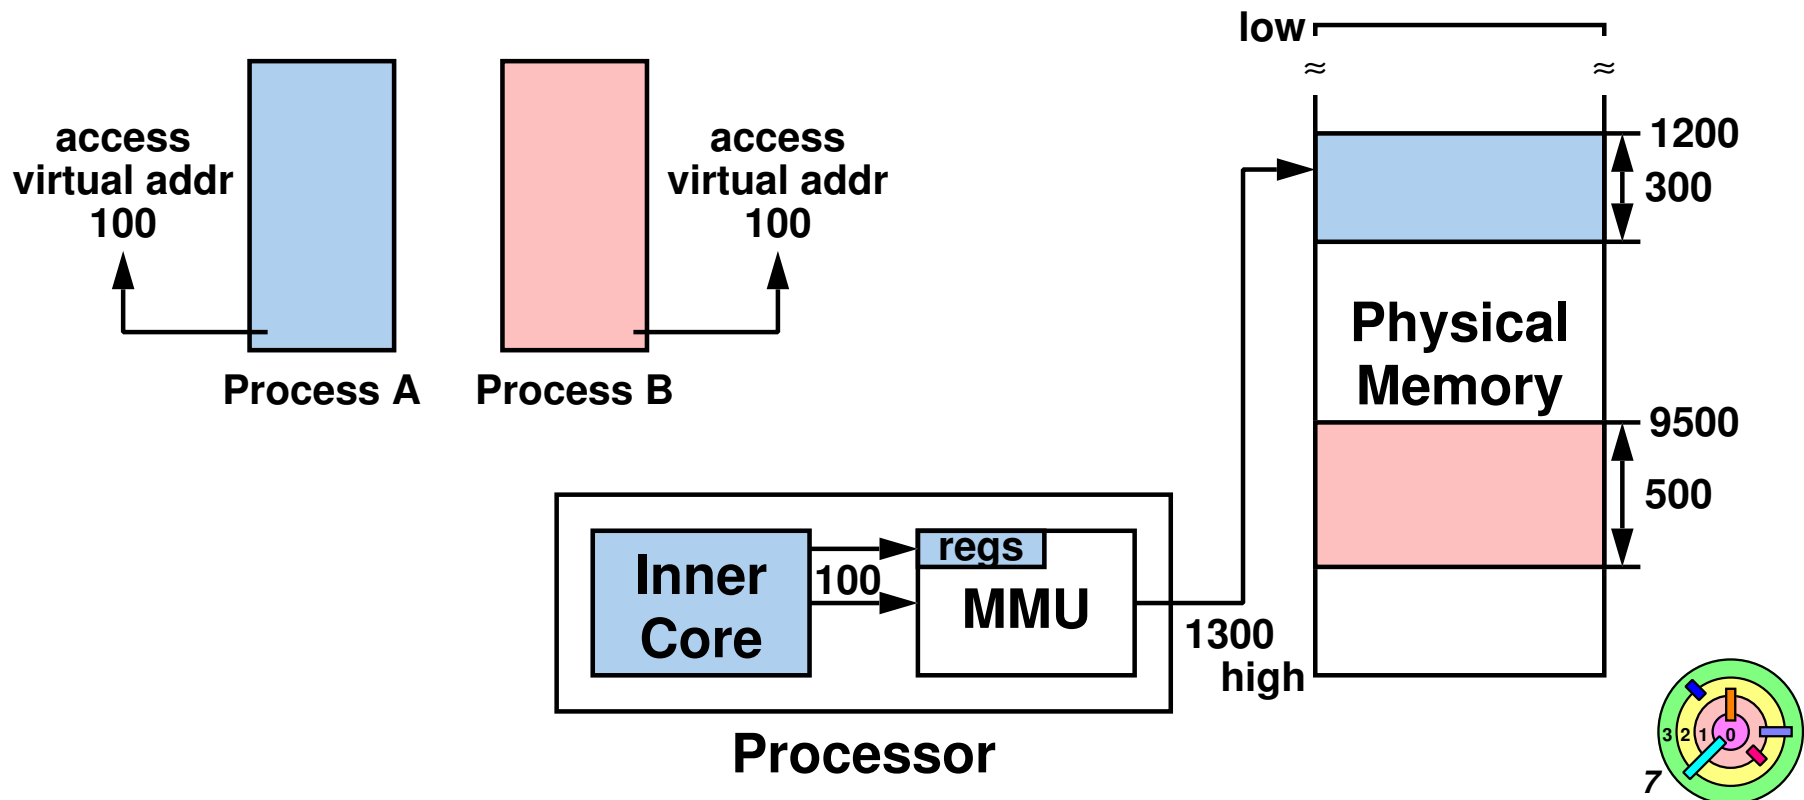
Basic Idea: Address Translation

- ➡ One level of *indirection* with a *Memory Management Unit (MMU)*
- don't address physical memory directly
 - address out of CPU inner core is *virtual*
 - use a *Memory Management Unit (MMU)*
 - virtual address is *translated* into physical address via MMU



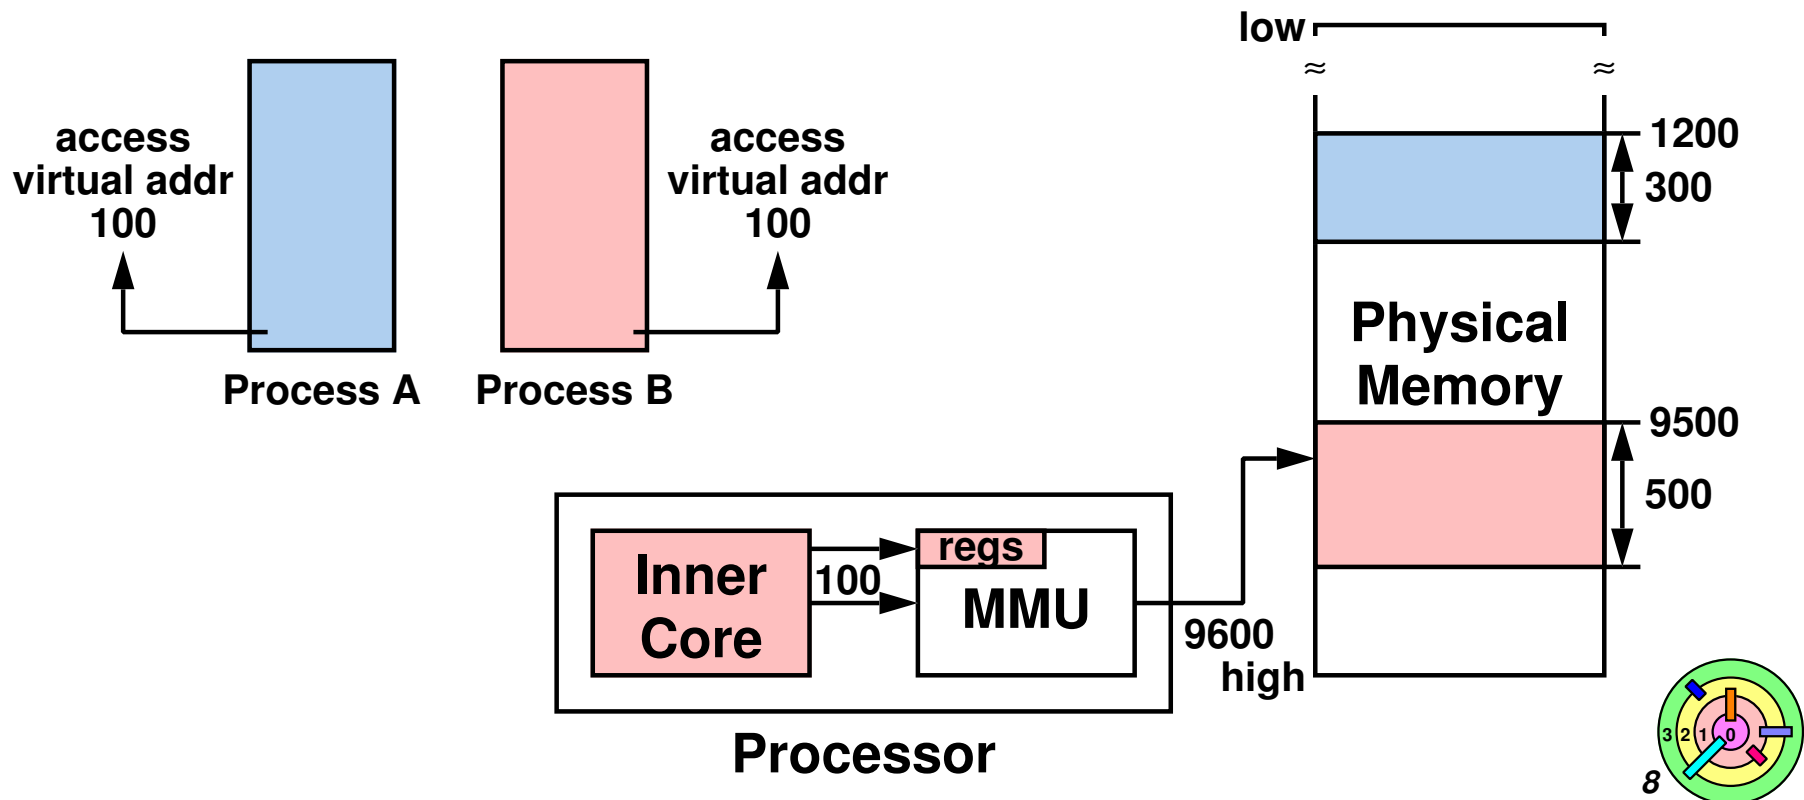
Basic Idea: Address Translation

- ➡ One level of *indirection* with a *Memory Management Unit (MMU)*
- don't address physical memory directly
 - address out of CPU inner core is *virtual*
 - use a *Memory Management Unit (MMU)*
 - virtual address is *translated* into physical address via MMU




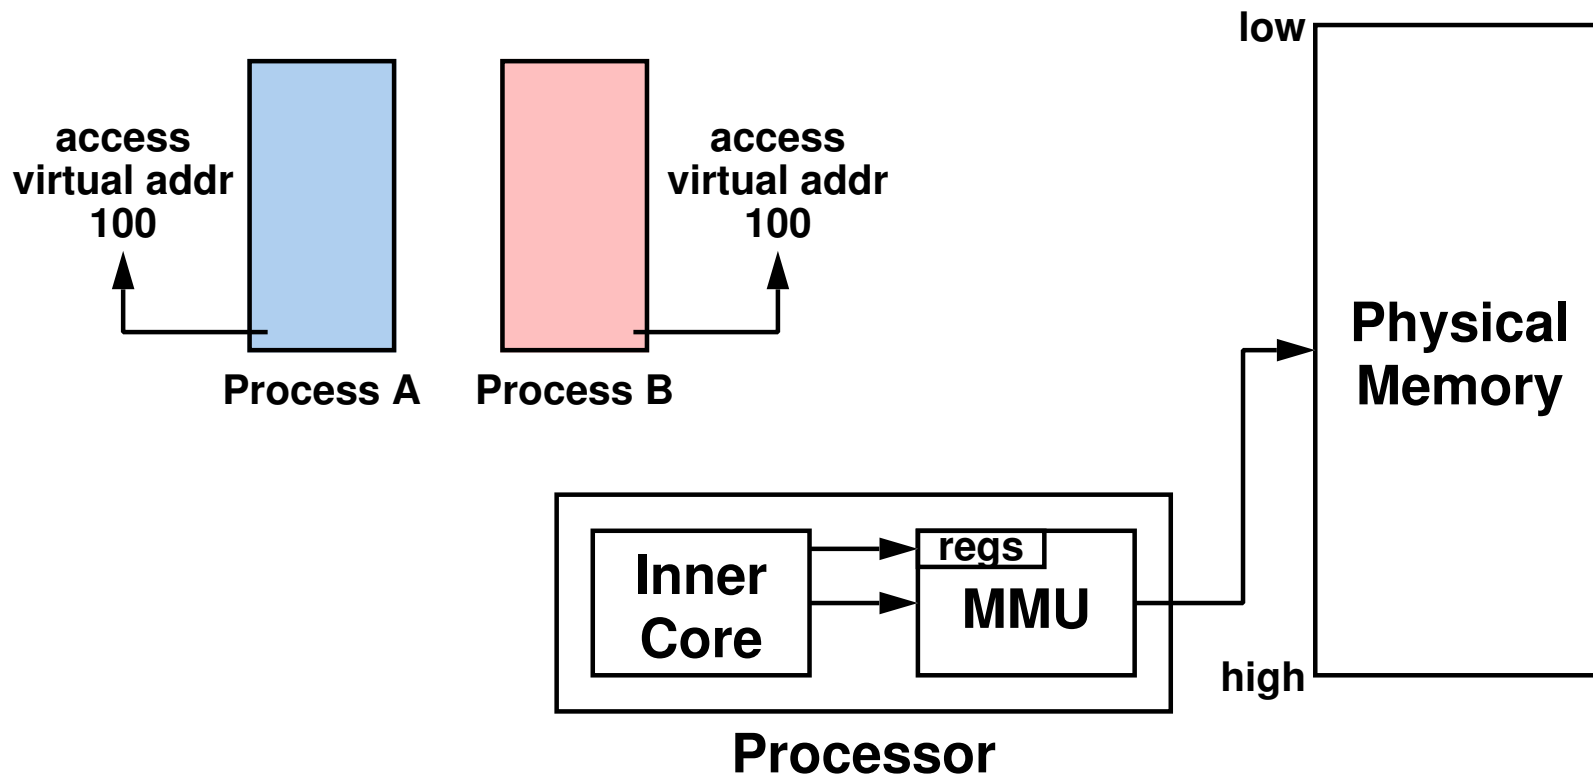
Basic Idea: Address Translation

- ➡ One level of *indirection* with a *Memory Management Unit (MMU)*
- don't address physical memory directly
 - address out of CPU inner core is *virtual*
 - use a *Memory Management Unit (MMU)*
 - virtual address is *translated* into physical address via MMU

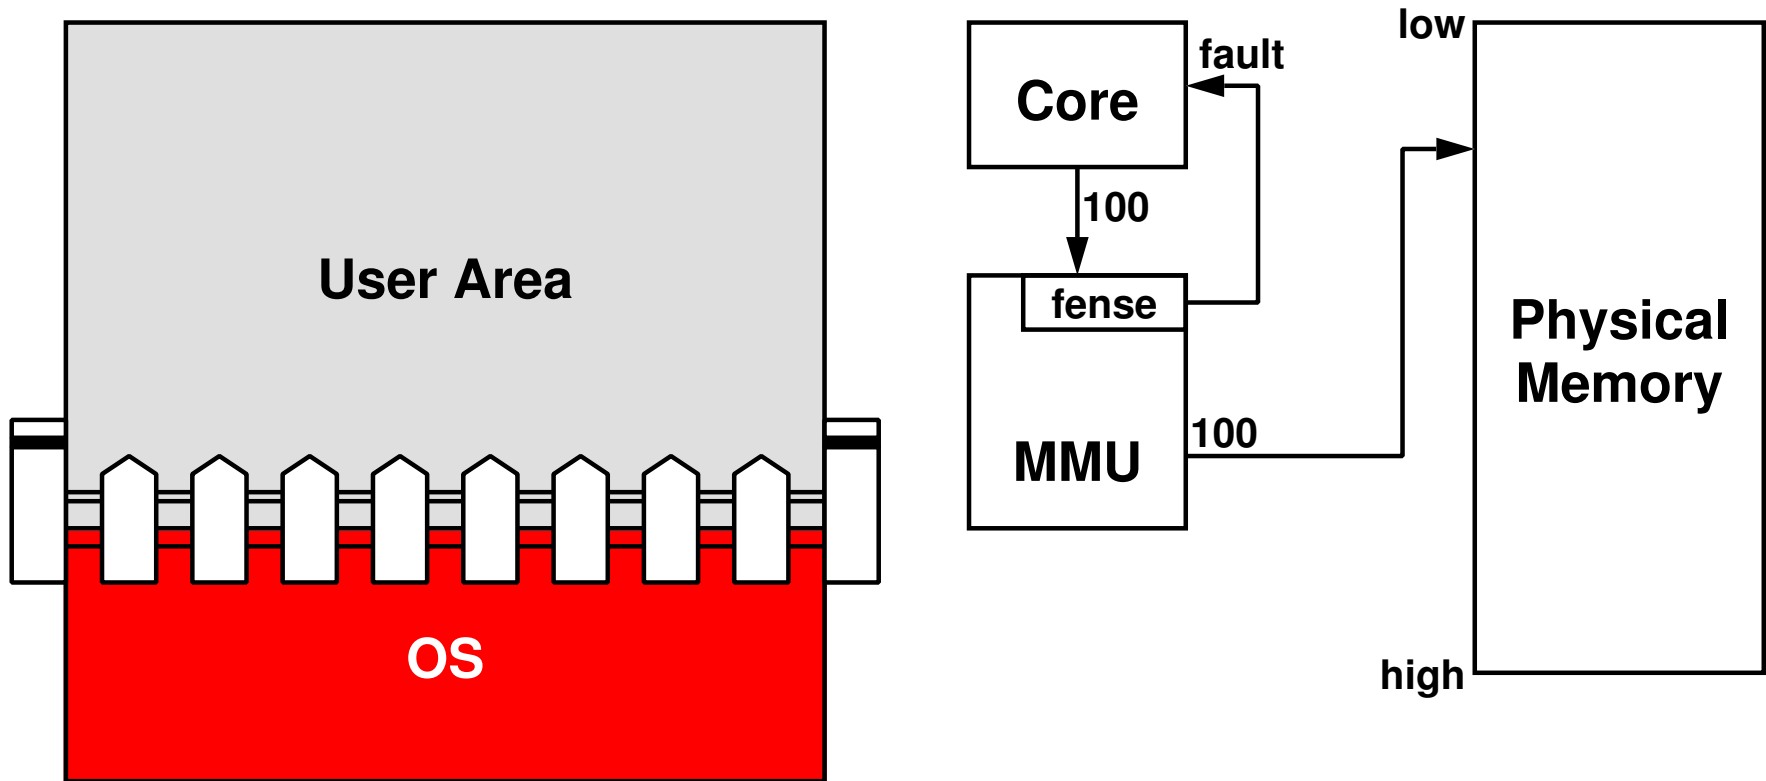


Address Translation

- ➡ Protection/isolation
- ➡ Illusion of large memory
- ➡ Sharing
- ➡ New abstraction (such as memory-mapped files)




Memory Fence




In the old days


- if a user program tries to access OS area, *hardware* (very simple MMU) will generate a trap
- does not protect user pocesses from each other
 - there's only one user process anyway



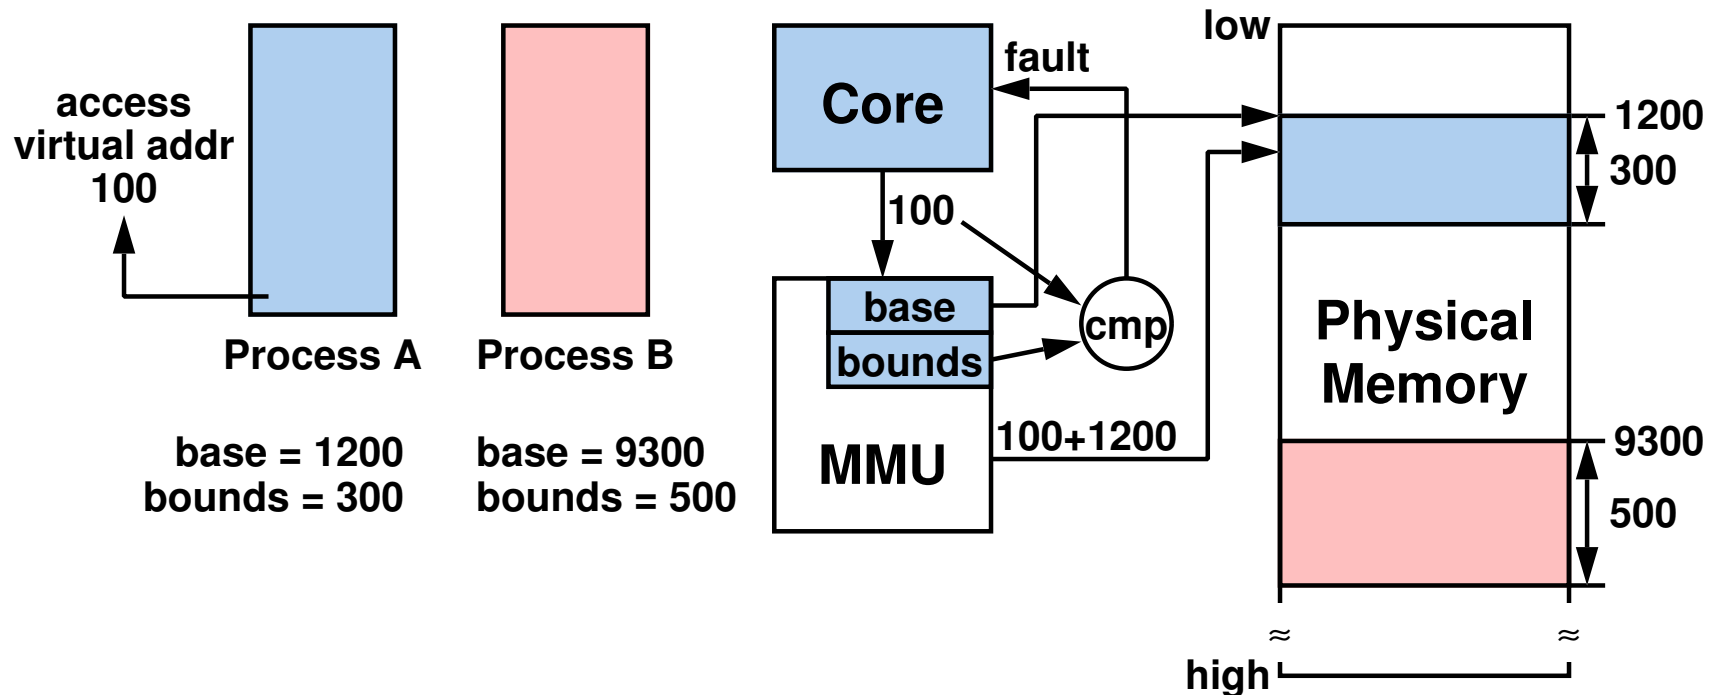
Memory Fence and Overlays



- ➡ What if the user program won't fit in memory?
- use *overlays*
 - programmers (not the OS) have to keep track of which overlay is in physical memory and deal with the complexities of managing *overlays*




Base and Bounds Registers




Multiple user processes

- OS maintains a pair of registers for each user process
 - **bounds register:** address space size of the user process
 - **base register:** start of physical memory for the user process
- addresses **relative** to the **base register**
- memory reference ≥ 0 and $< \text{bounds}$, **independent** of **base** (this is known as "position independence")




Base and Bounds Registers




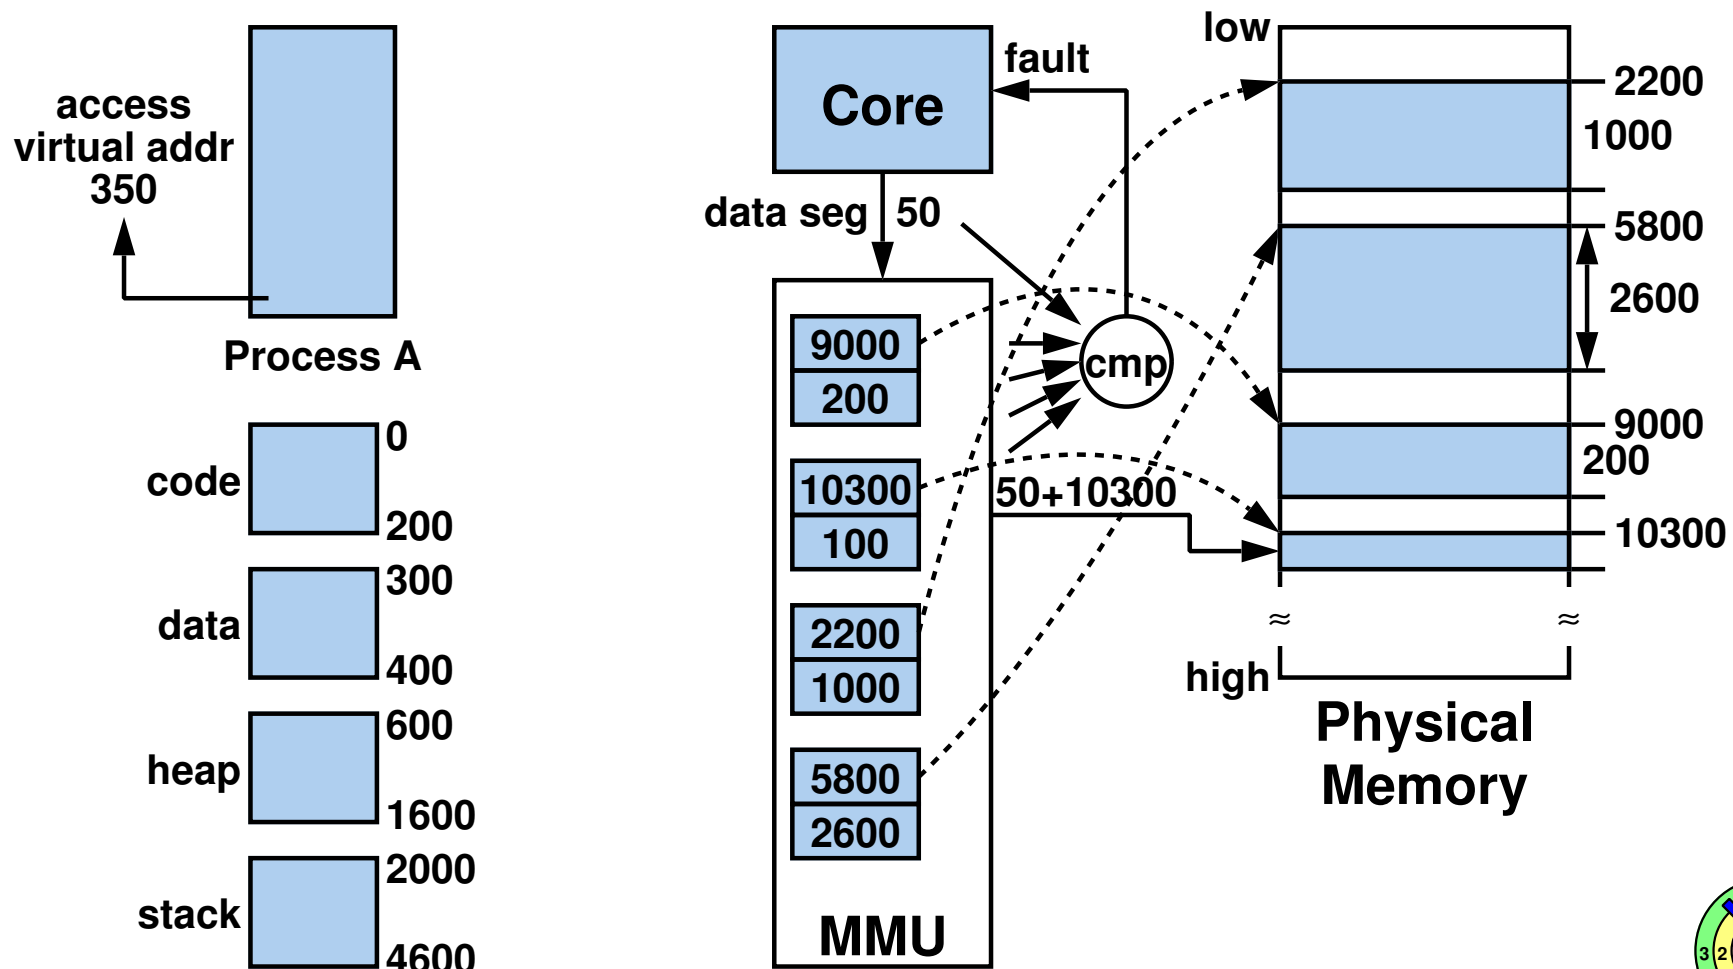
Multiple user processes

- OS maintains a pair of registers for each user process
 - **bounds register:** address space size of the user process
 - **base register:** start of physical memory for the user process
- addresses **relative** to the **base register**
- memory reference ≥ 0 and $<$ bounds, **independent** of **base** (this is known as "position independence")




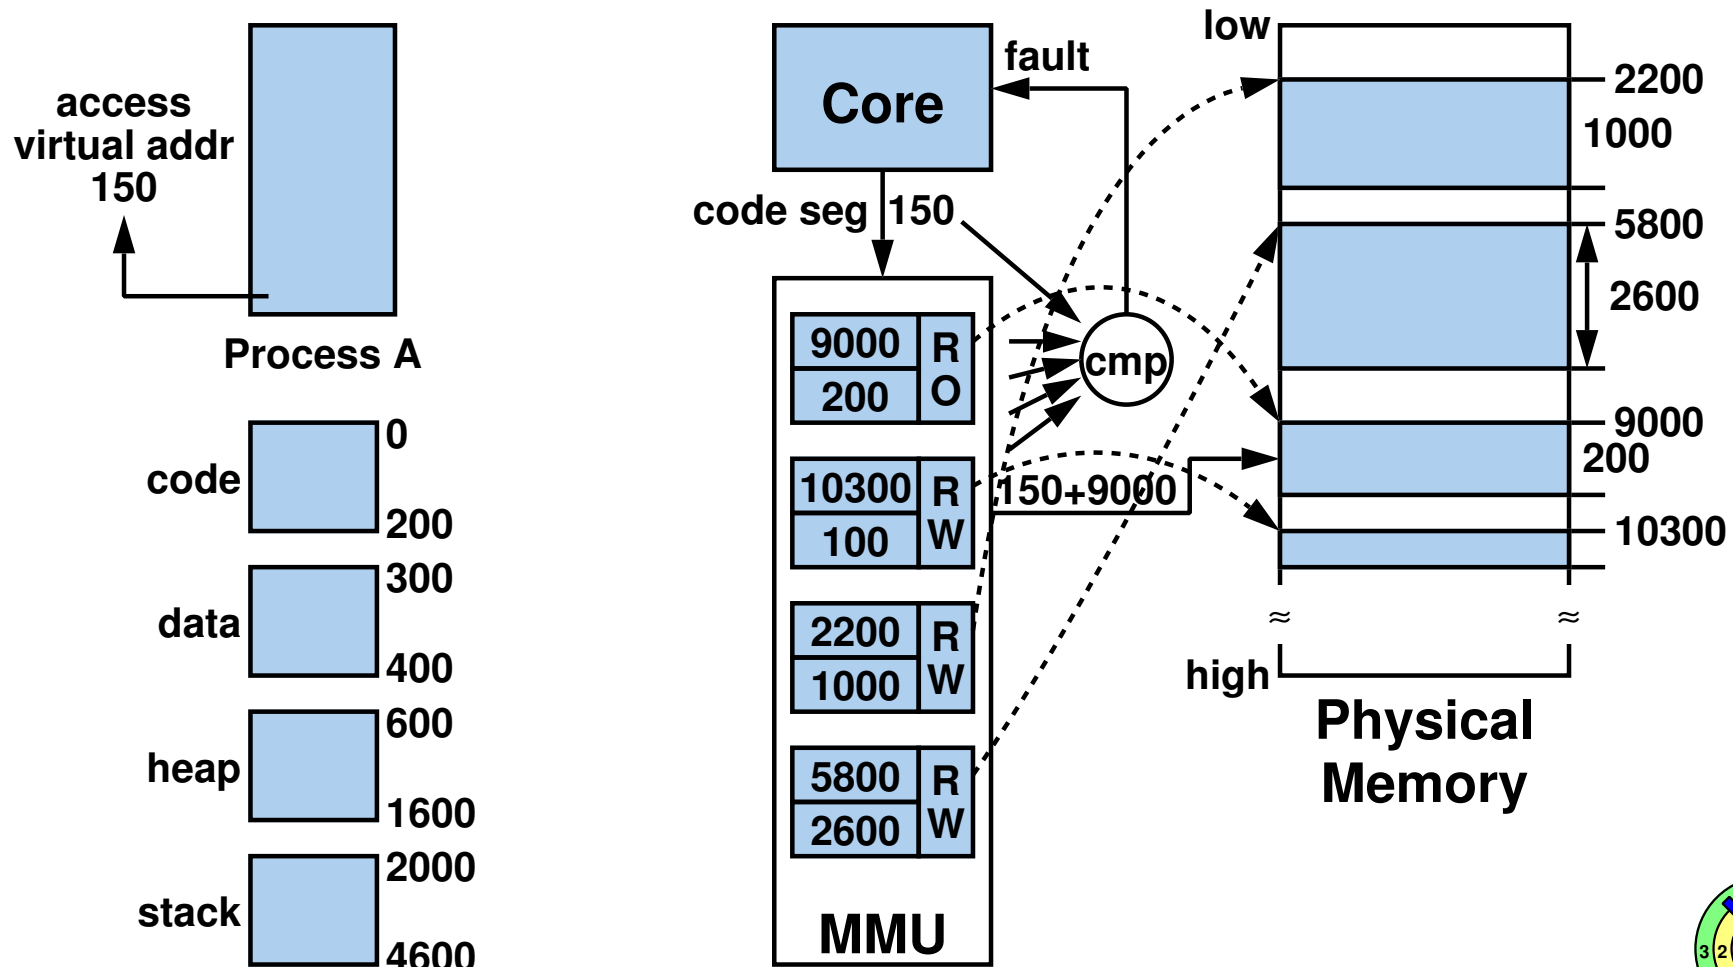
Generalization of Base and Bounds: Segmentation

- ➡ One pair of *base* and *bounds* registers *per segment*
- code, data, heap, stack, and may be more
 - compiler compiles programs into segments



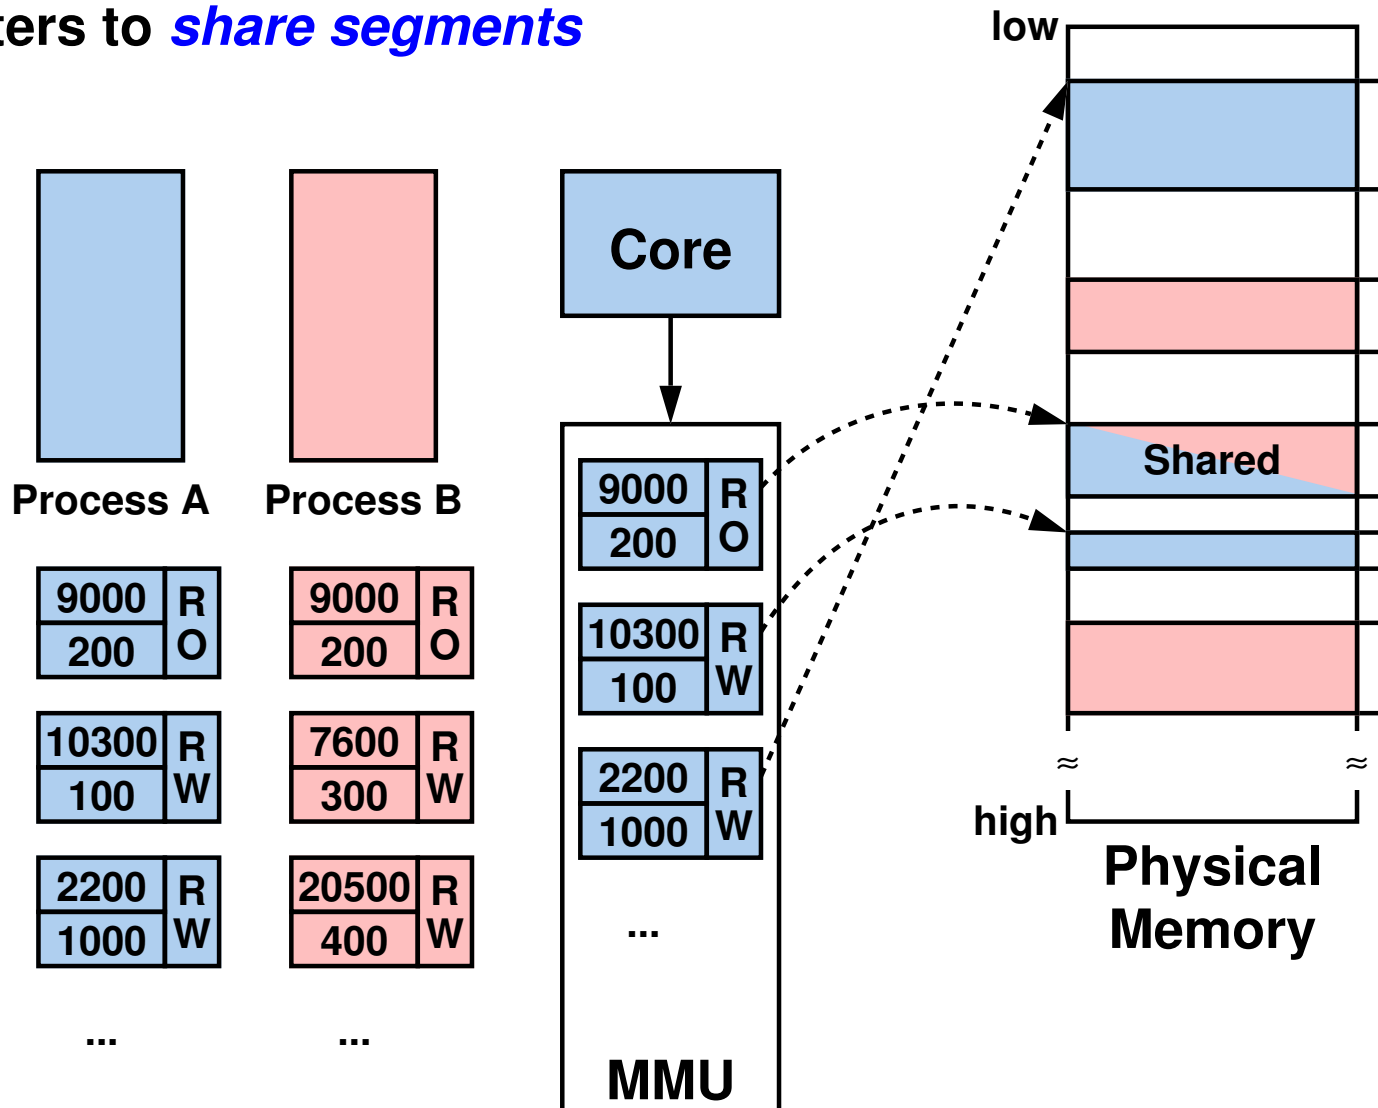
Access Control With Segmentation

➡ Access control / protection
 = *read-only*, *read/write*




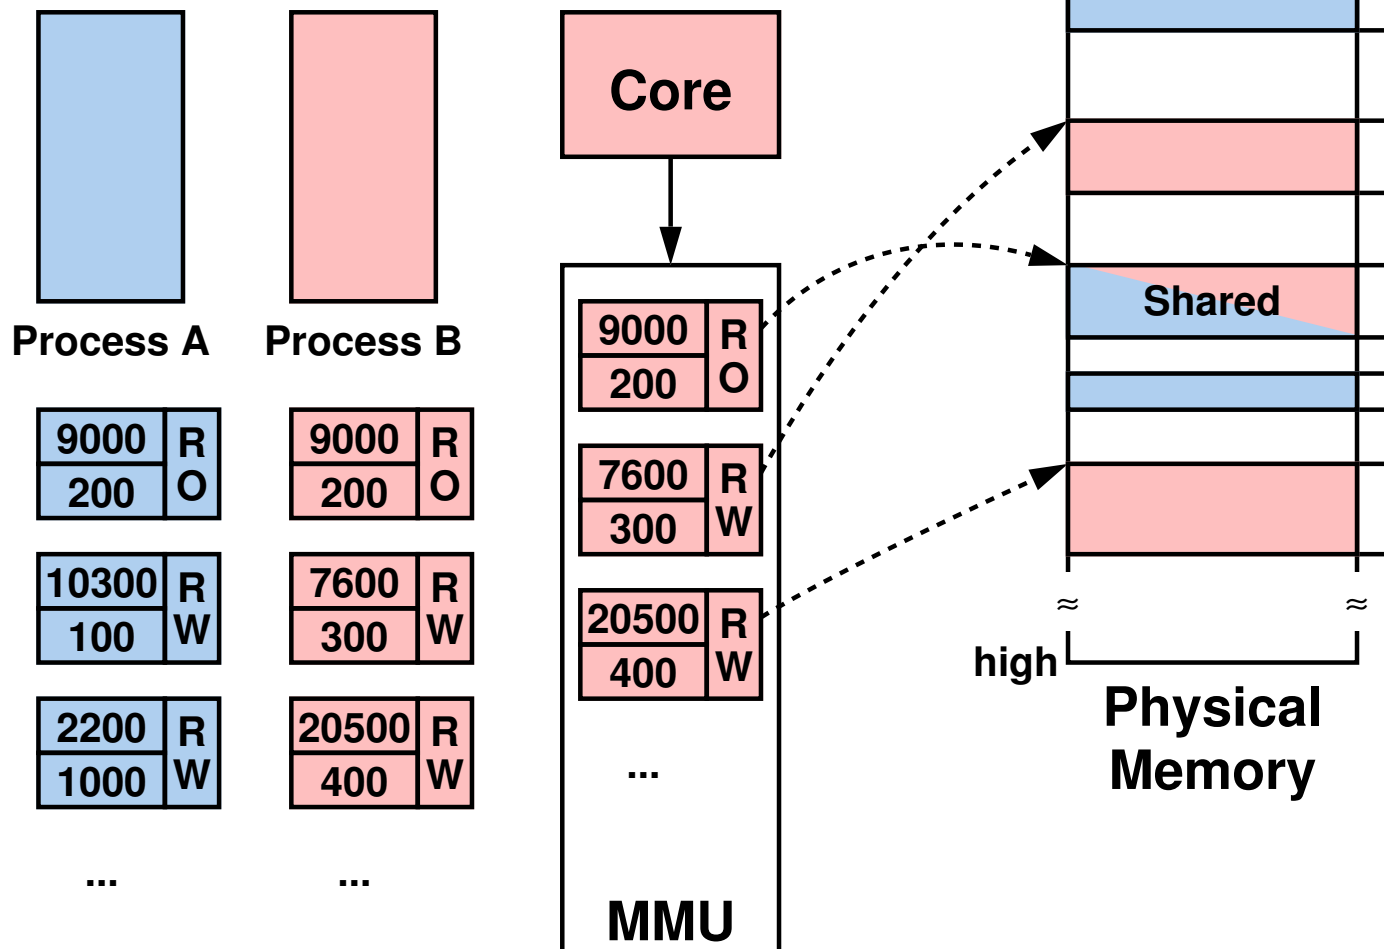
Sharing Segments

➡ Can simply setup base and bounds registers to *share segments*



Sharing Segments

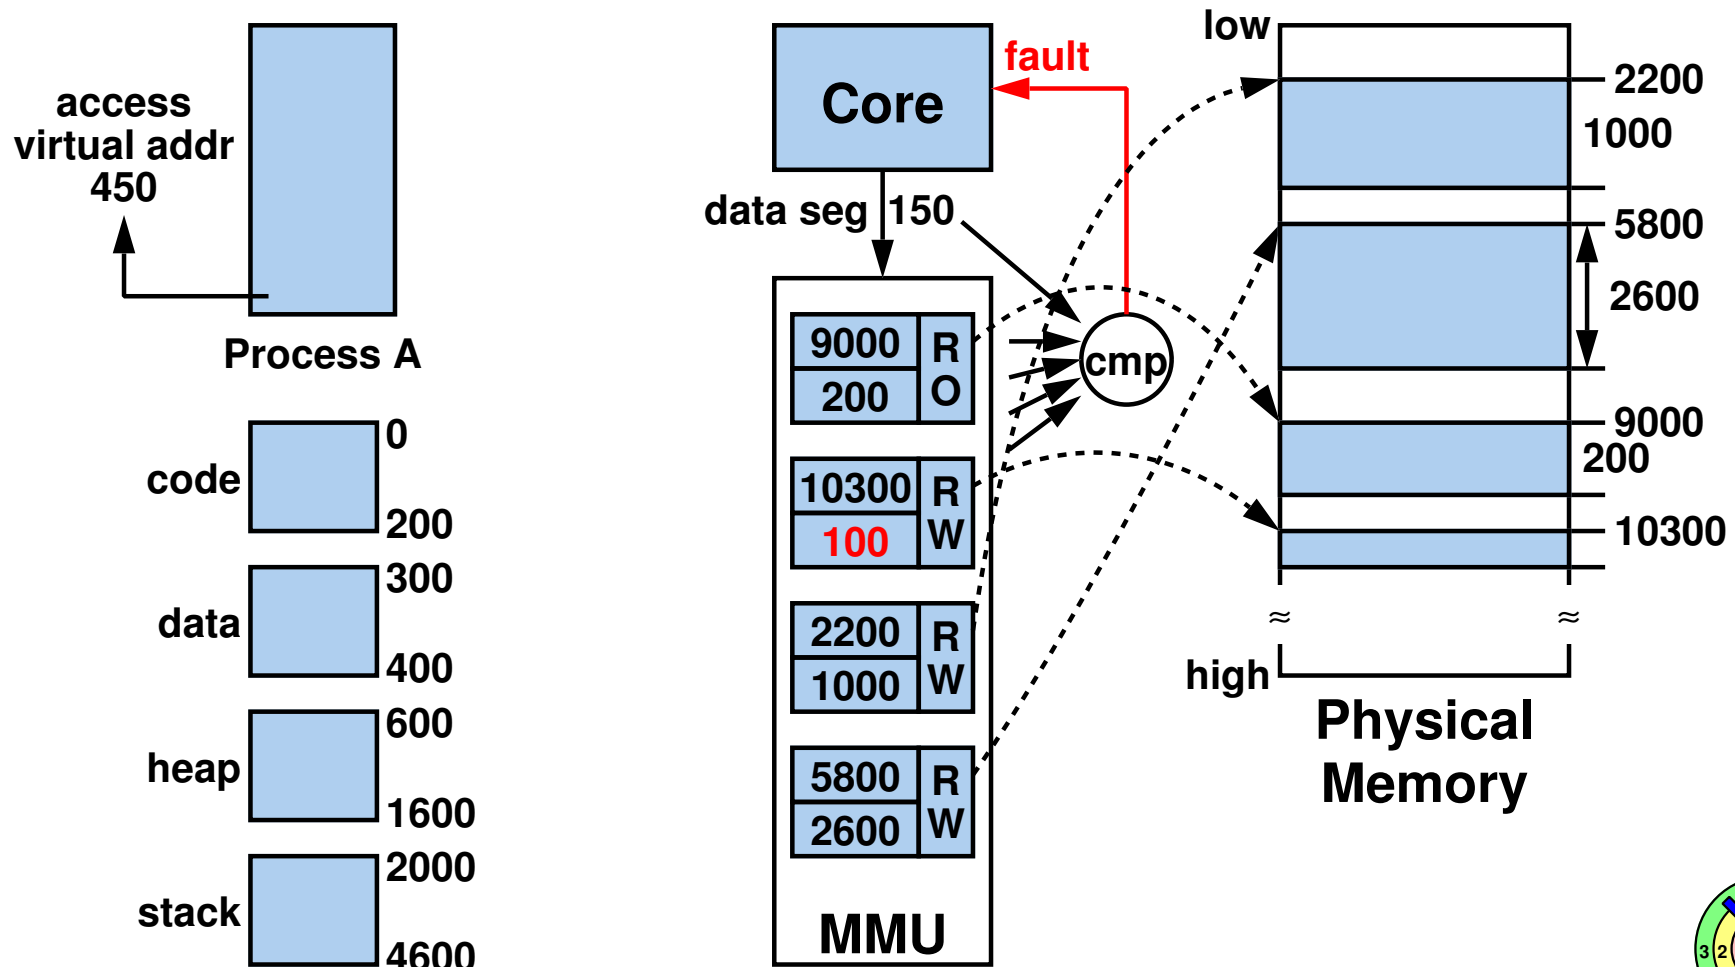
➡ Can simply setup base and bounds registers to *share segments*



Segmentation Fault

➔ *Segmentation fault*

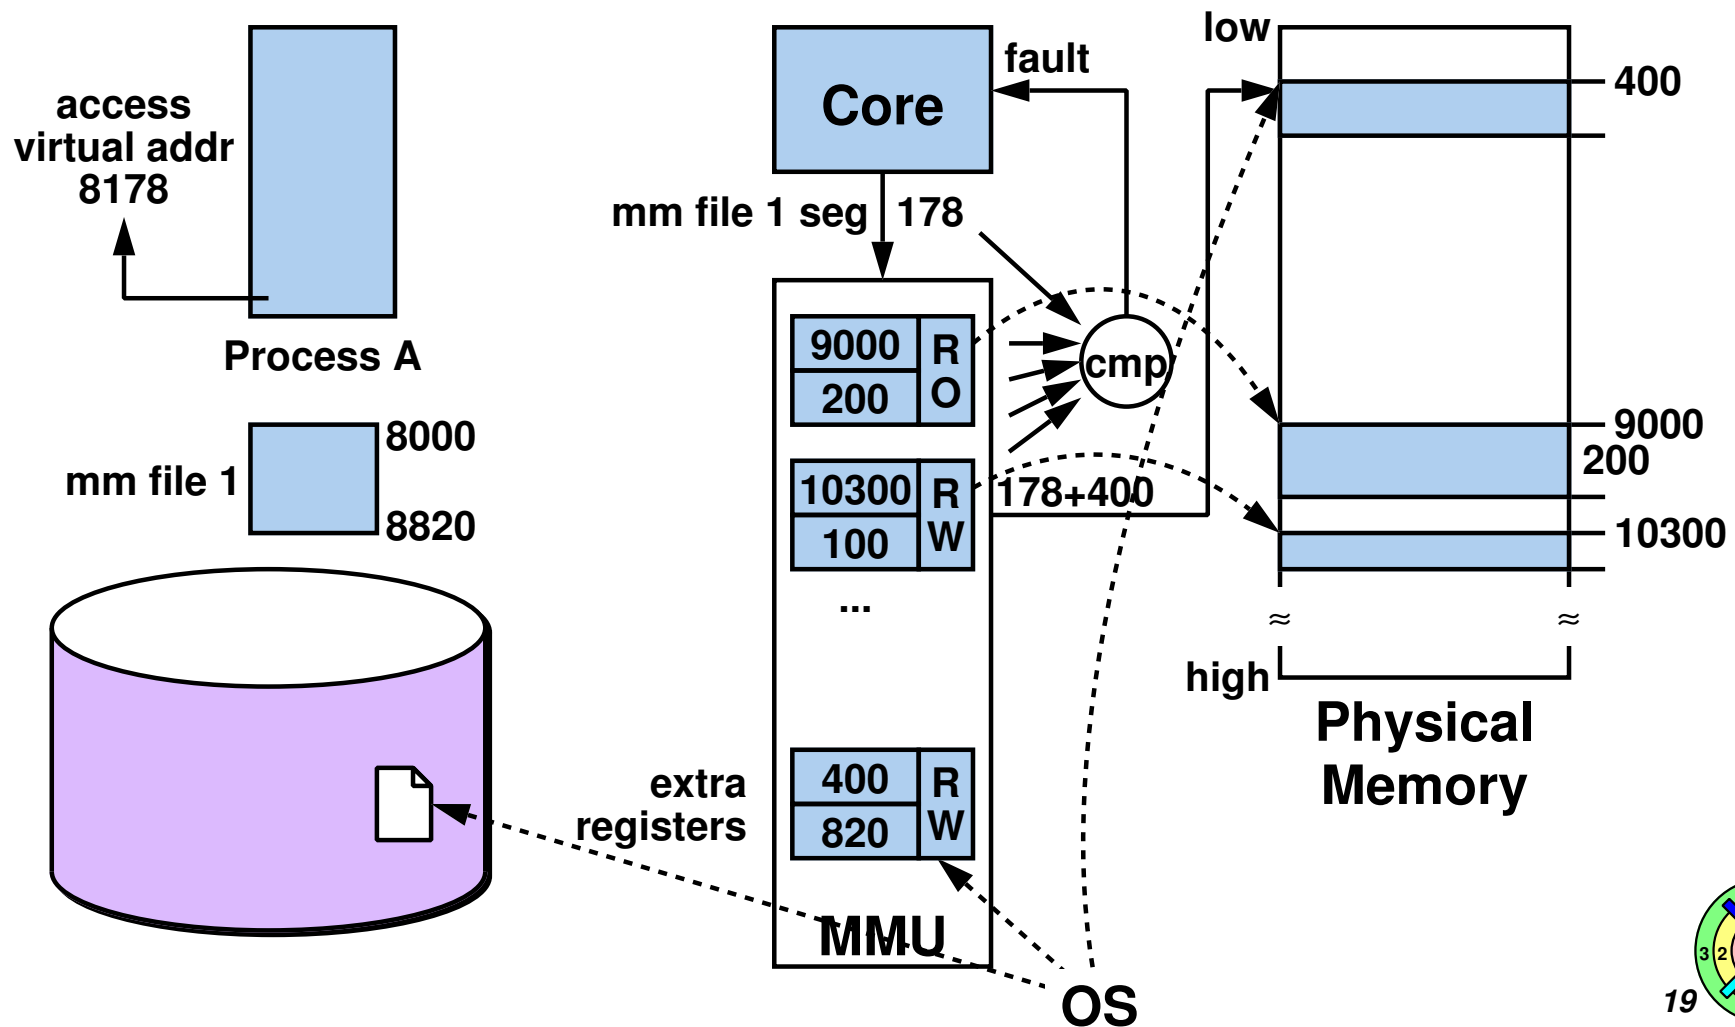
— virtual address not within range of any base-bounds registers



Memory Mapped File

➡ *Memory Mapped File*


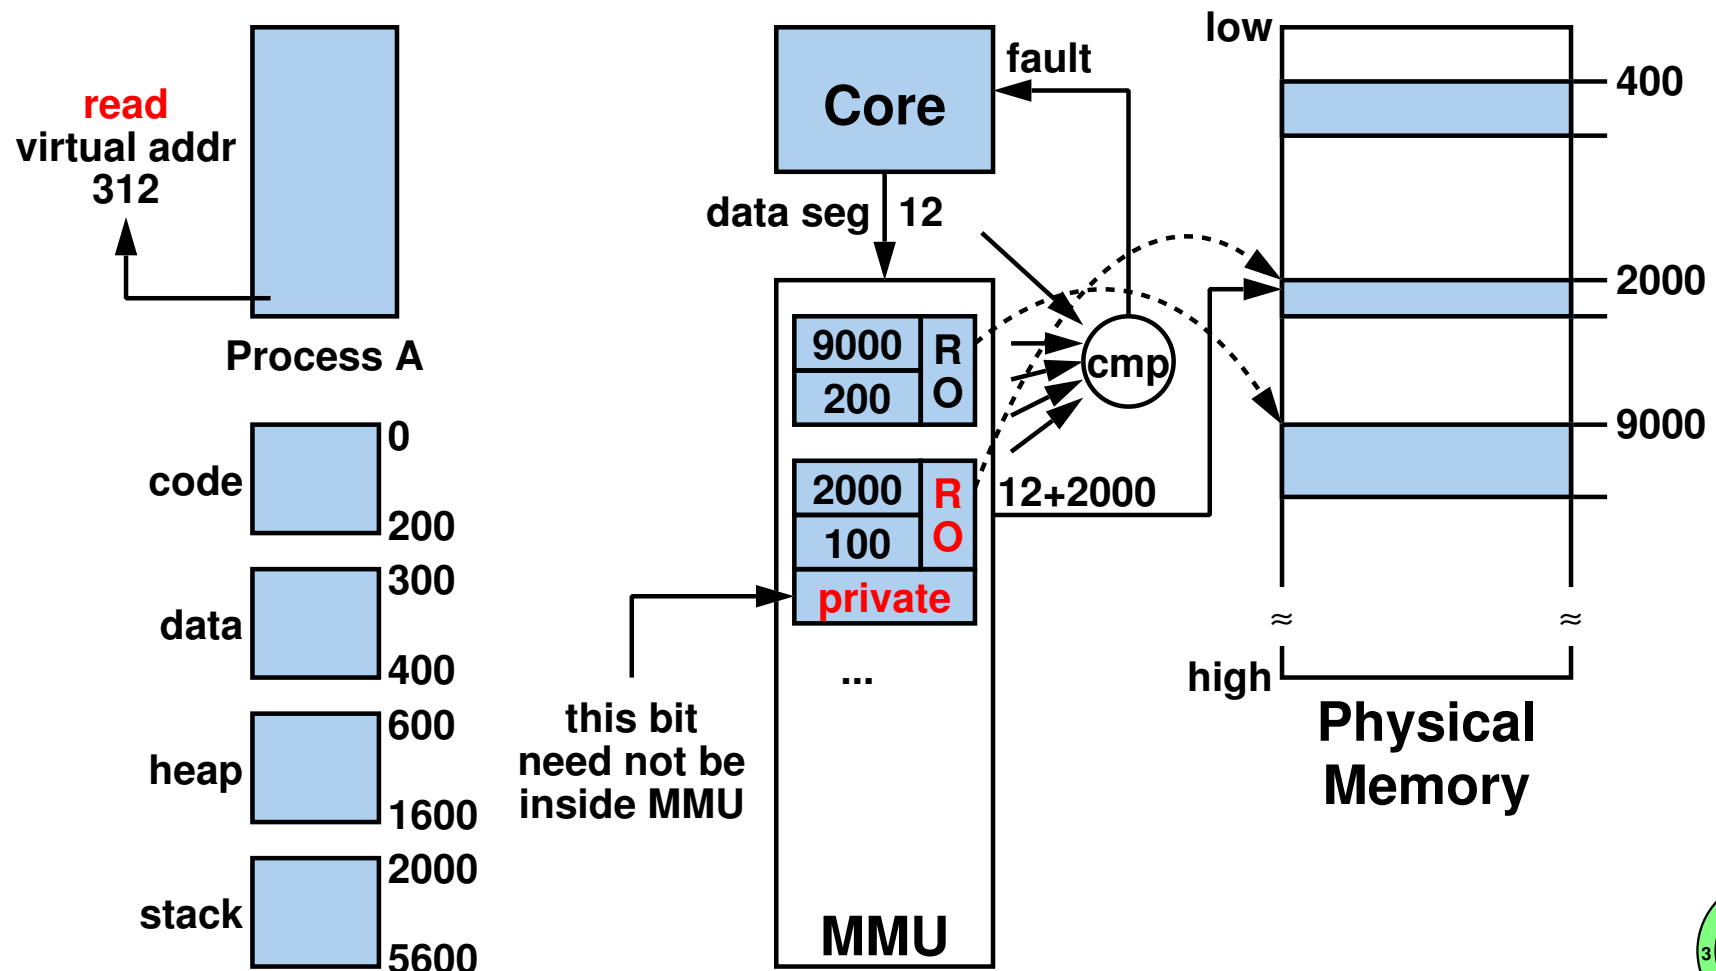
- the `mmap()` system call
- can map an entire file (or part of it) into a segment



Copy-On-Write

➡ **Copy-on-write (COW):**


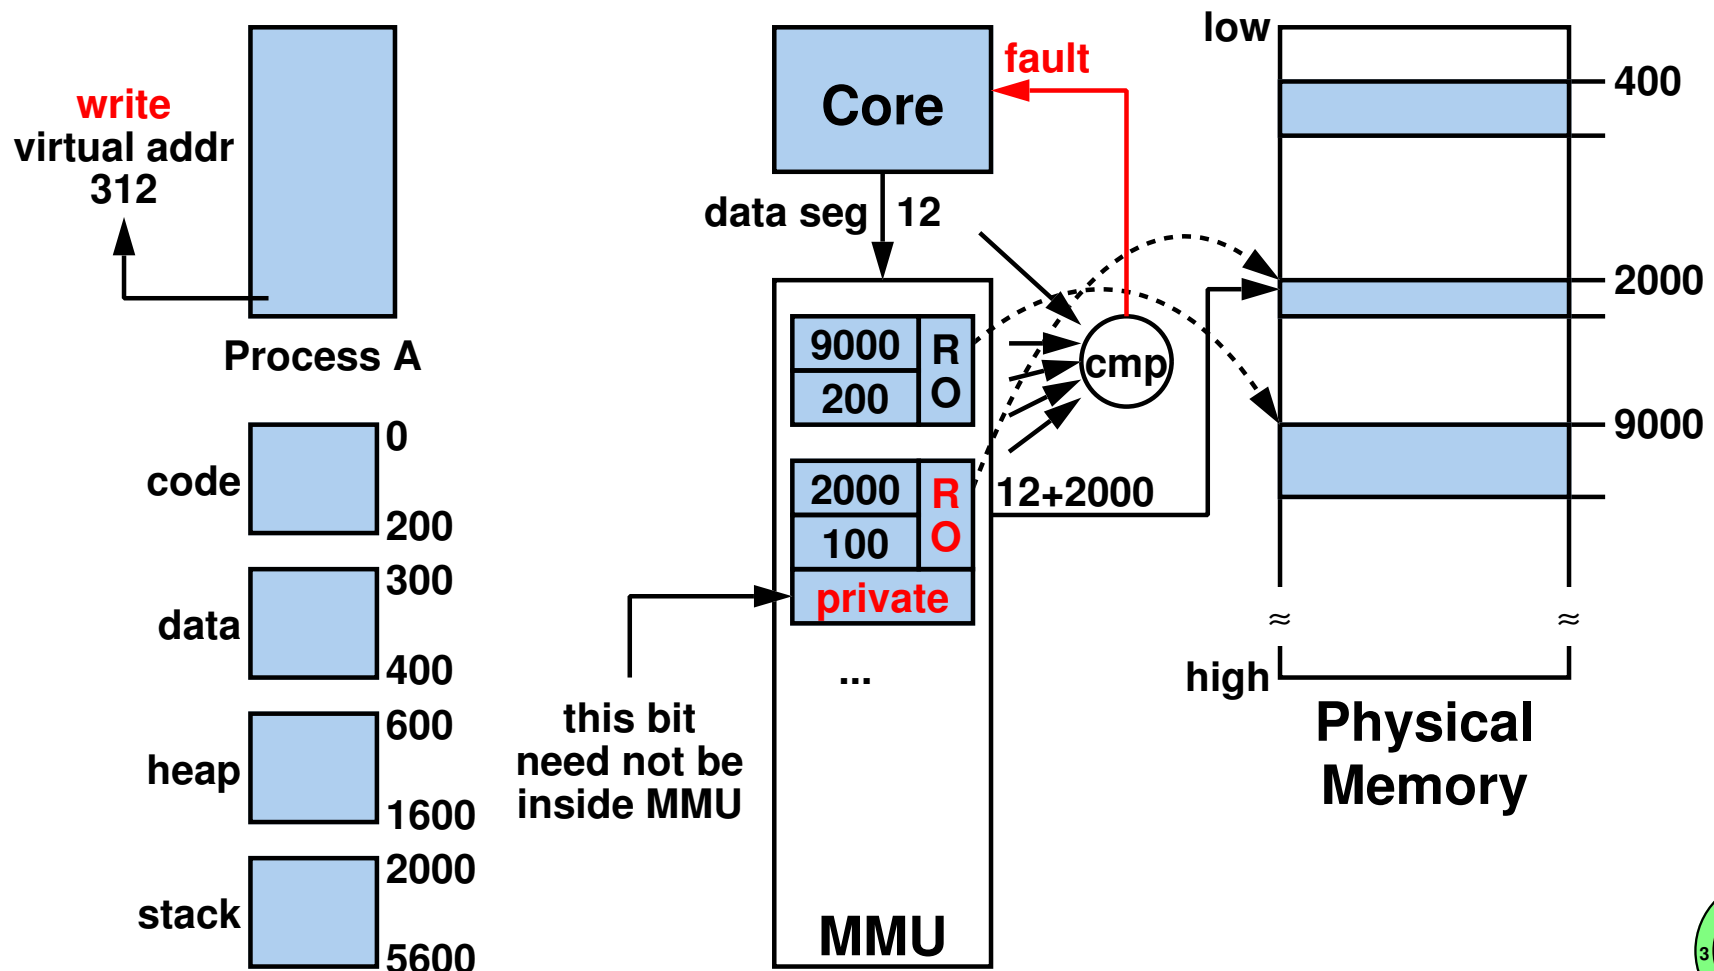
- a process gets a **private** copy of the page after a thread in the process performs a **write** for the **first time**



Copy-On-Write

➡ **Copy-on-write (COW):**


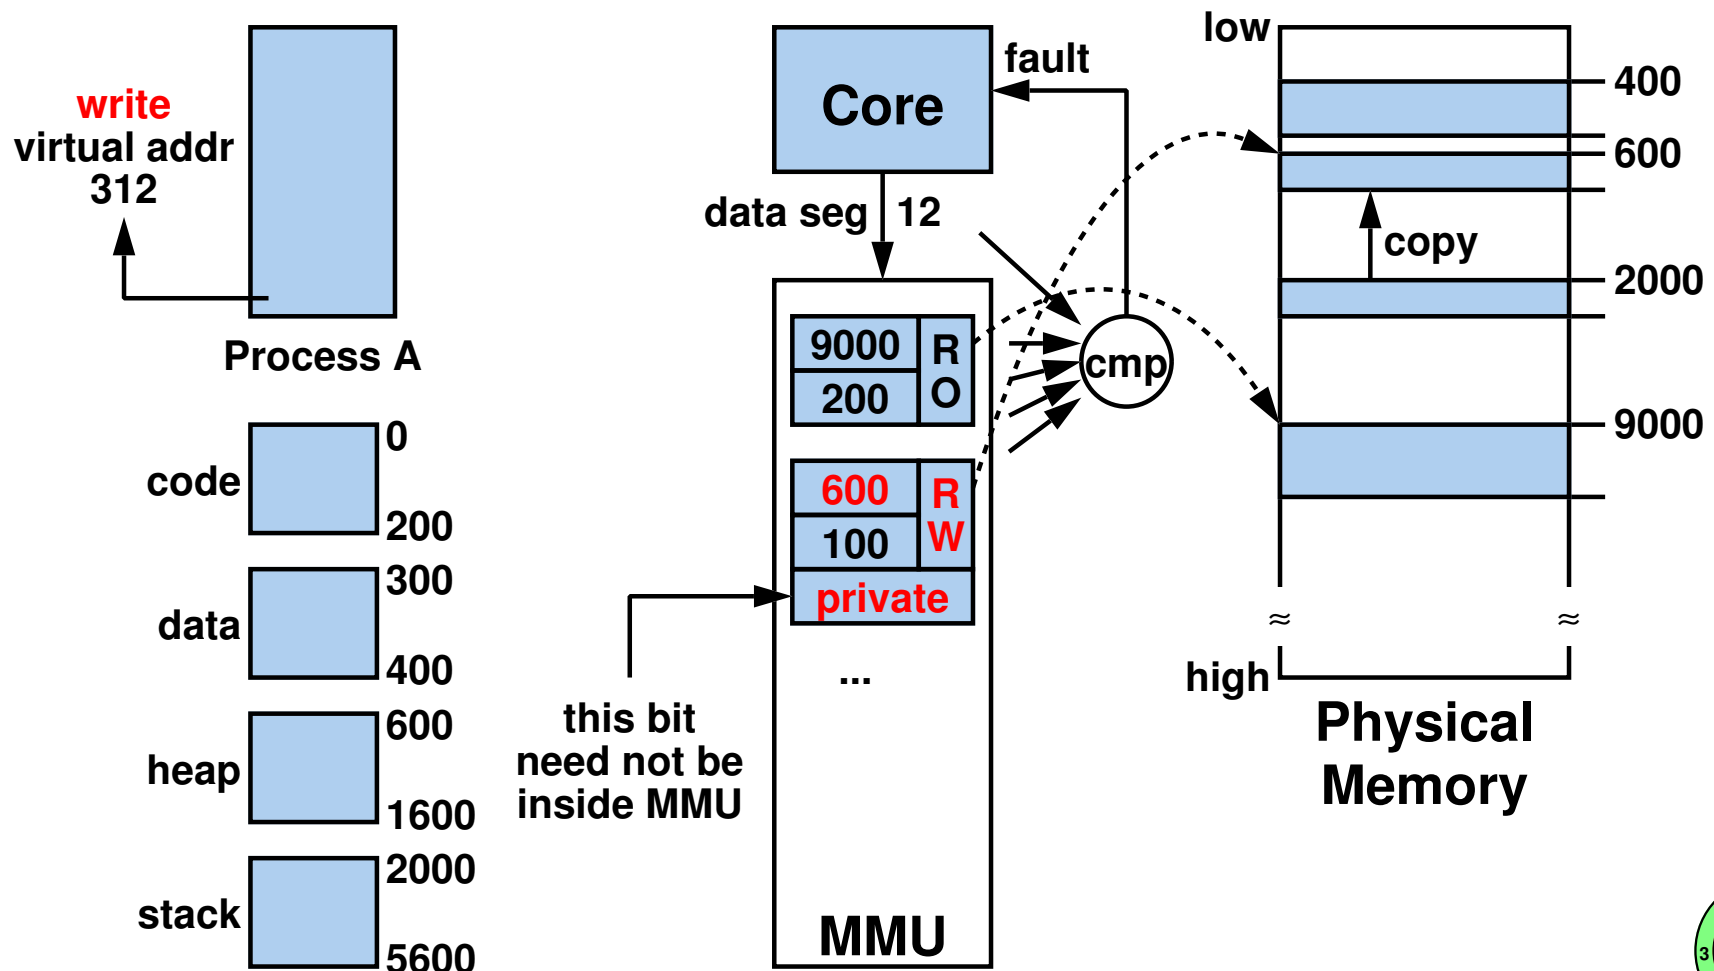
- a process gets a **private** copy of the page after a thread in the process performs a **write** for the **first time**



Copy-On-Write

➡ **Copy-on-write (COW):**

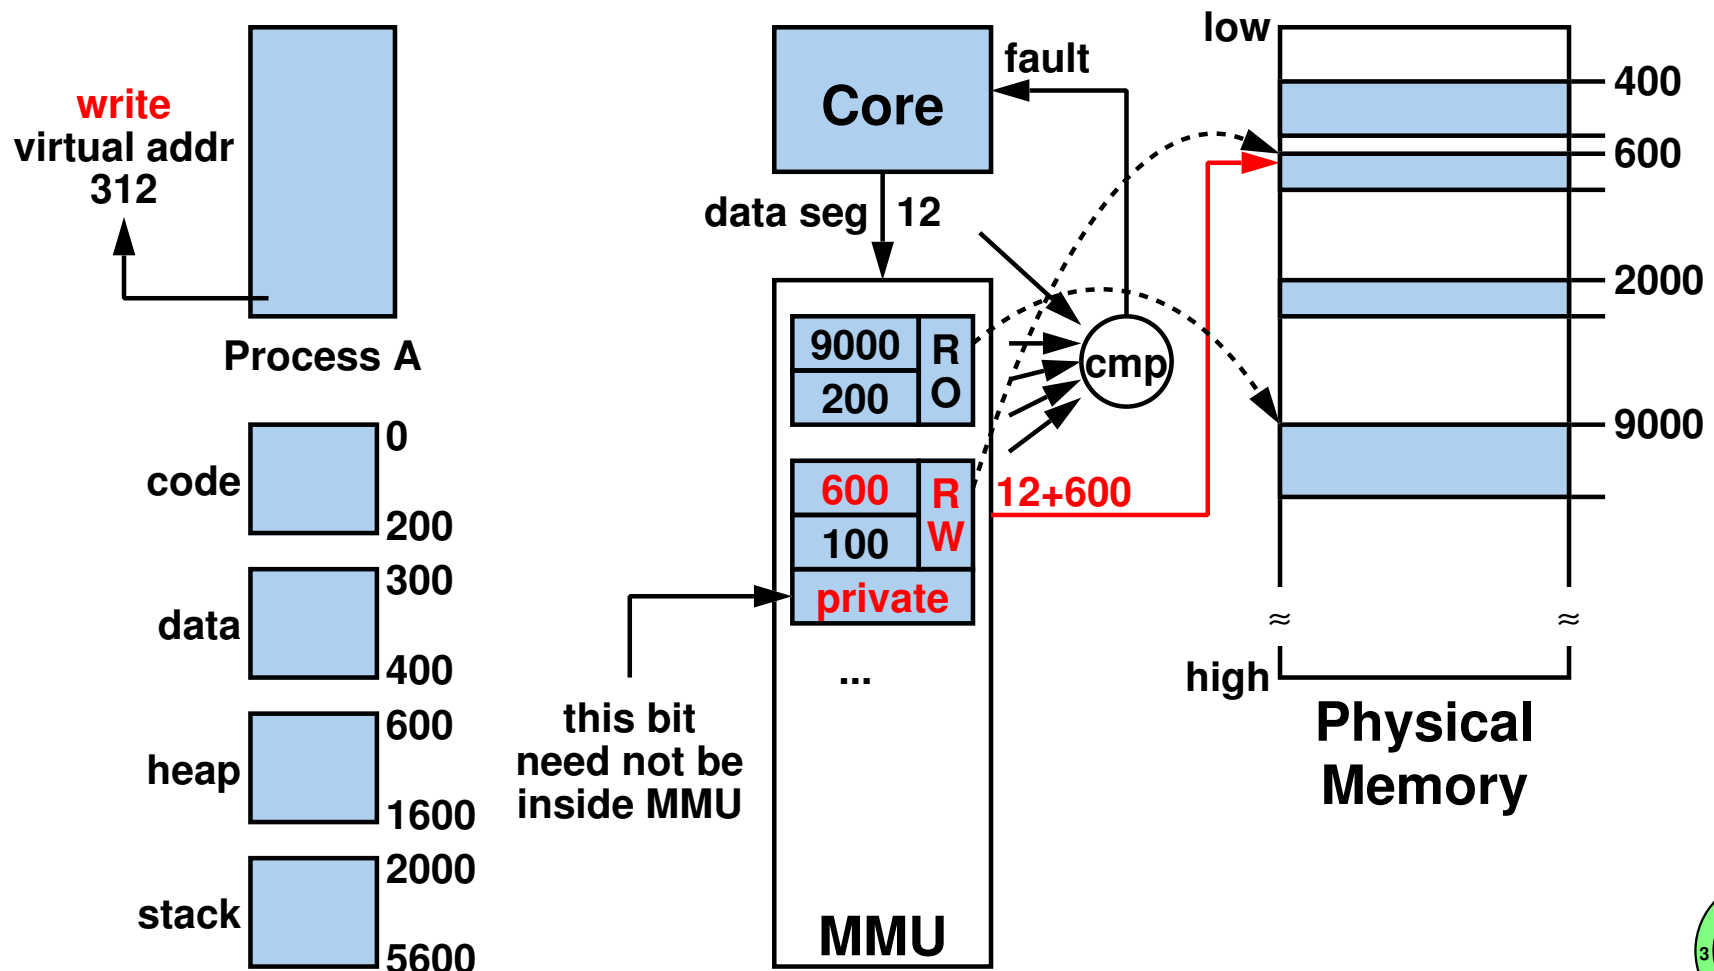
- a process gets a **private** copy of the page after a thread in the process performs a **write** for the **first time**



Copy-On-Write


➡ **Copy-on-write (COW):**

- a process gets a **private** copy of the page after a thread in the process performs a **write** for the **first time**



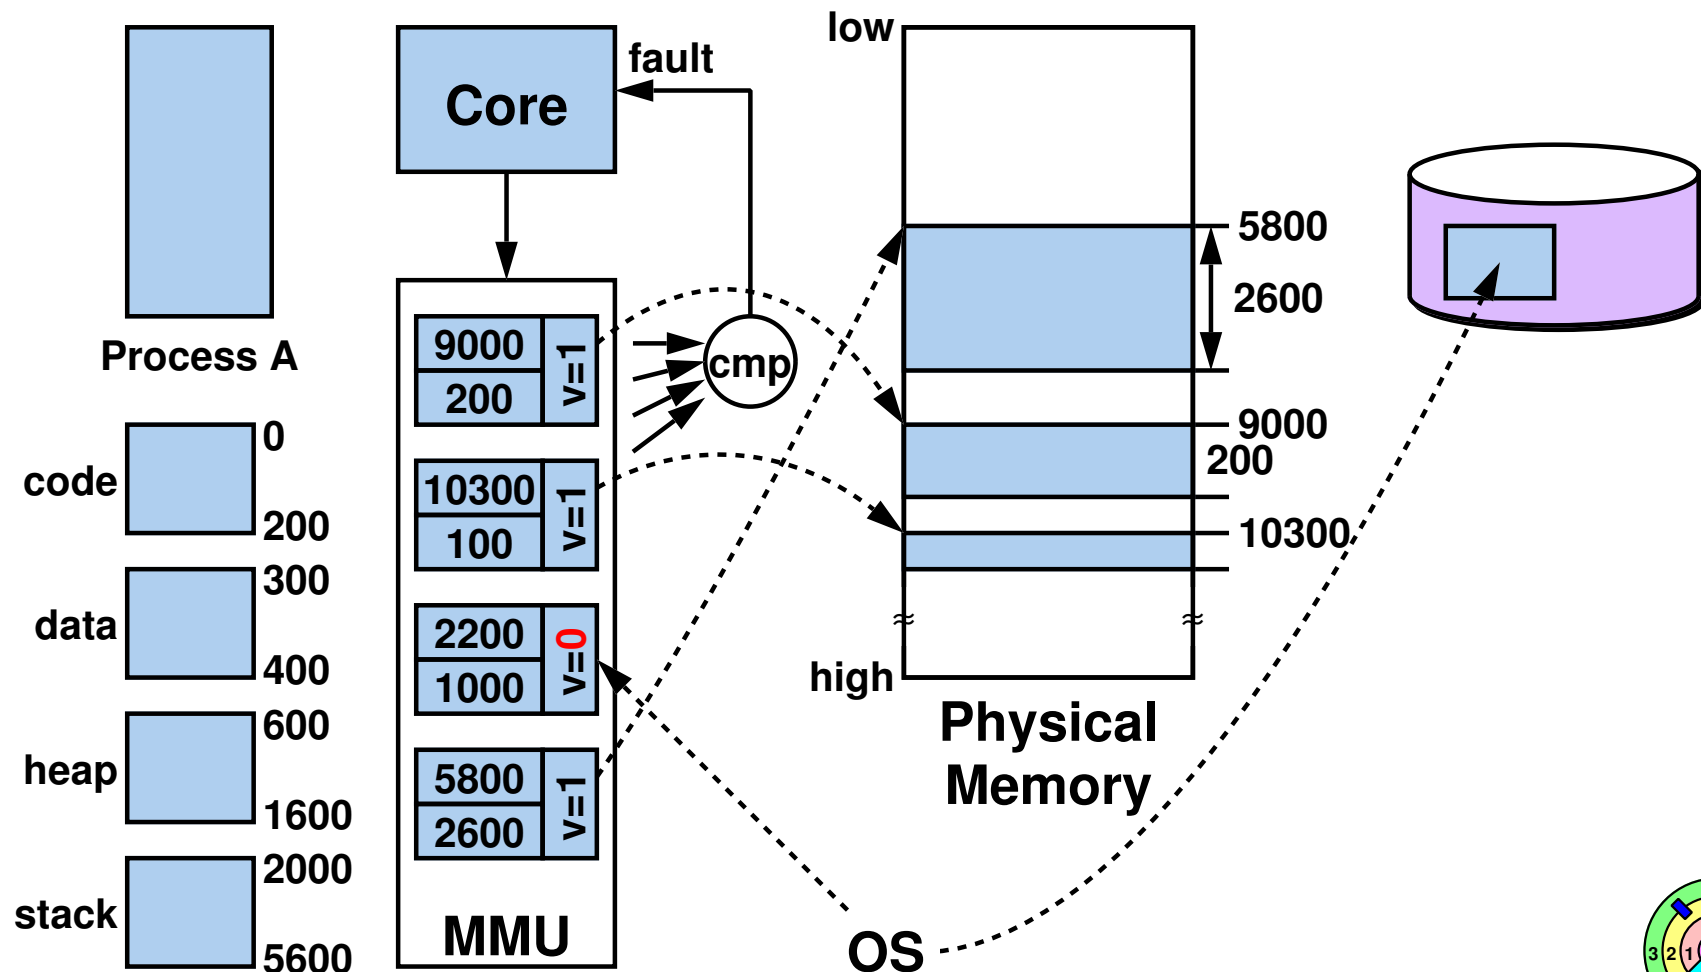
Swapping / Backing Store

- ➡ No space for new segment, make room by swapping out a segment
 — use a **validity** bit for each segment (in addition to access control bits)



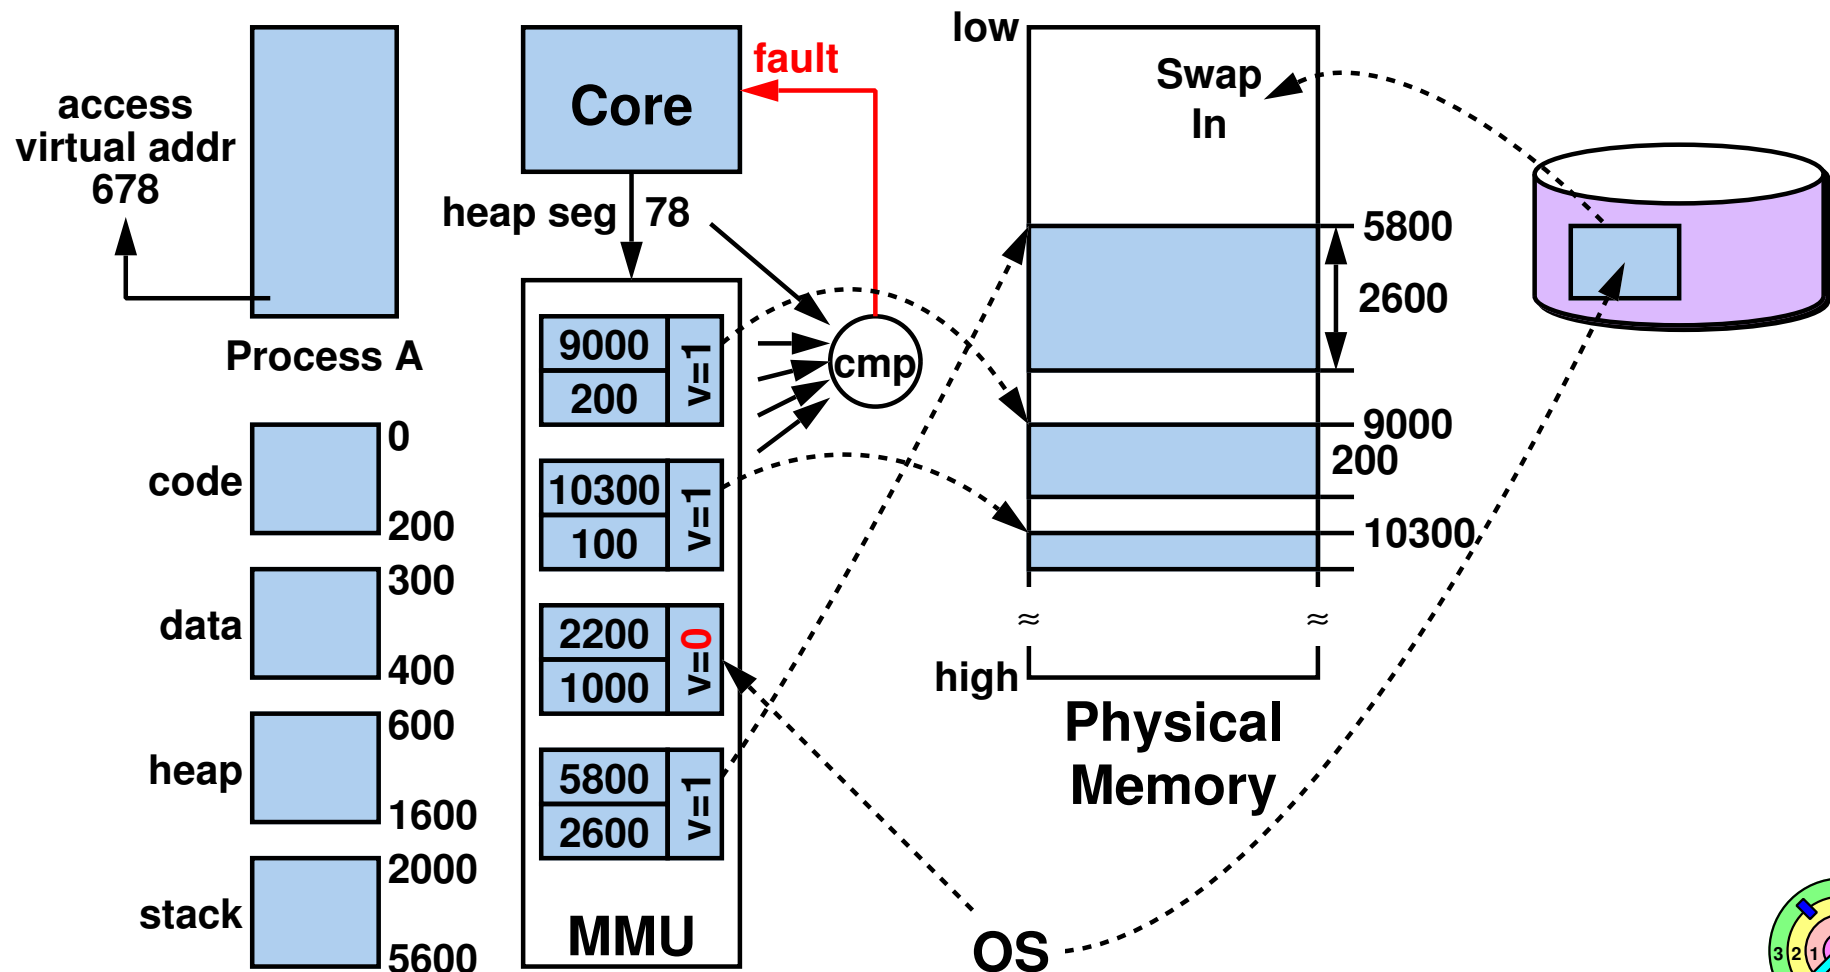
Swapping / Backing Store

- ➡ No space for new segment, make room by swapping out a segment
- use a *validity* bit for each segment (in addition to access control bits)



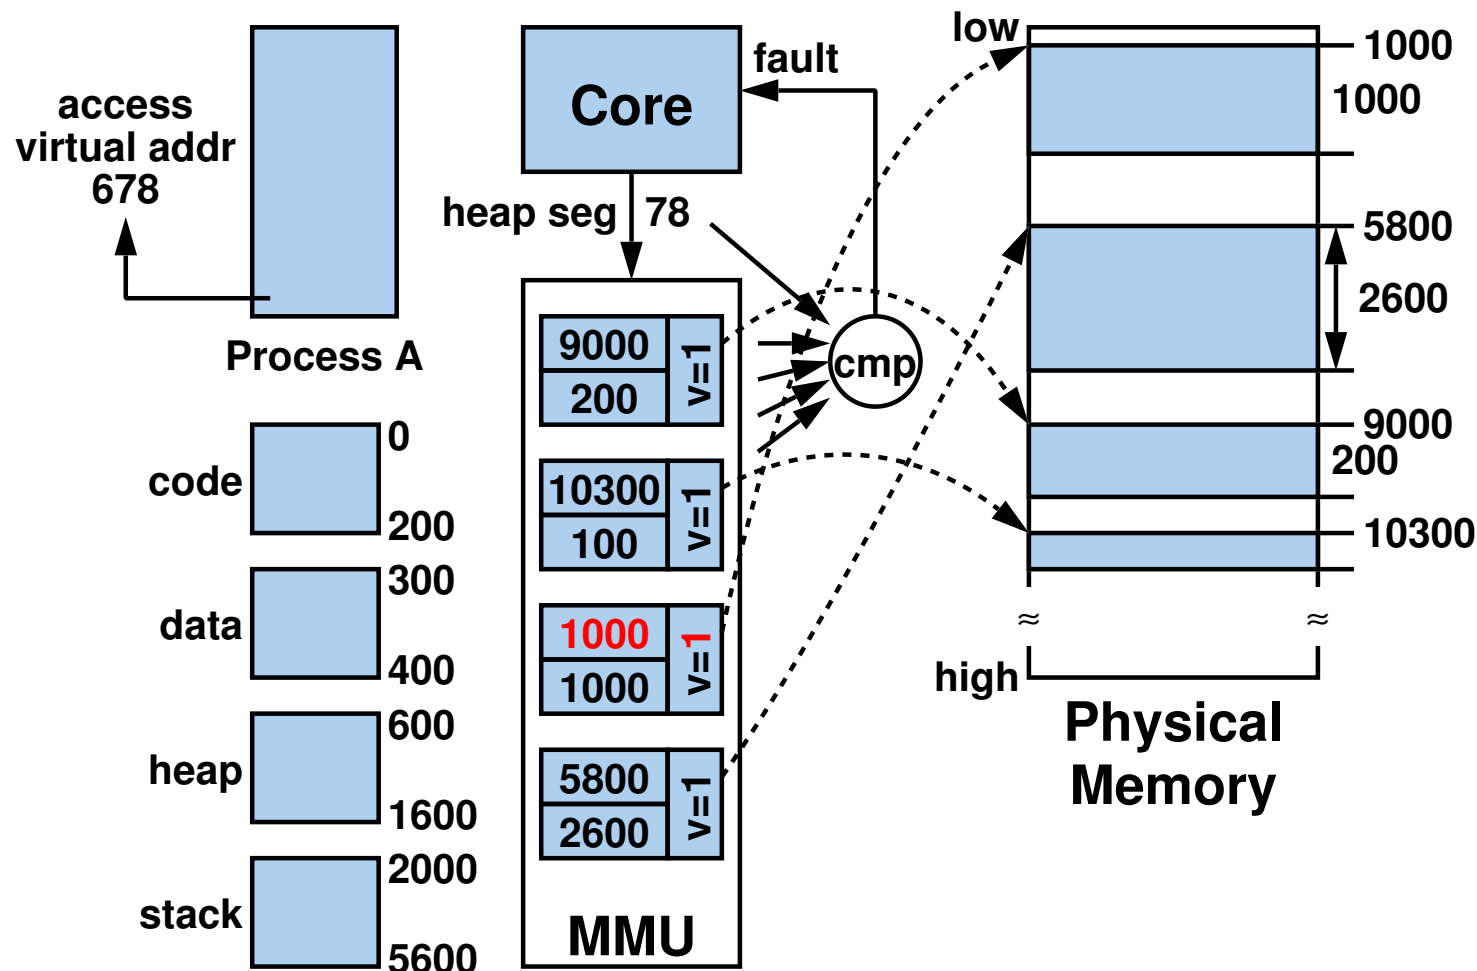
Swapping / Backing Store

- ➡ No space for new segment, make room by swapping out a segment
 — use a *validity* bit for each segment (in addition to access control bits)



Swapping / Backing Store

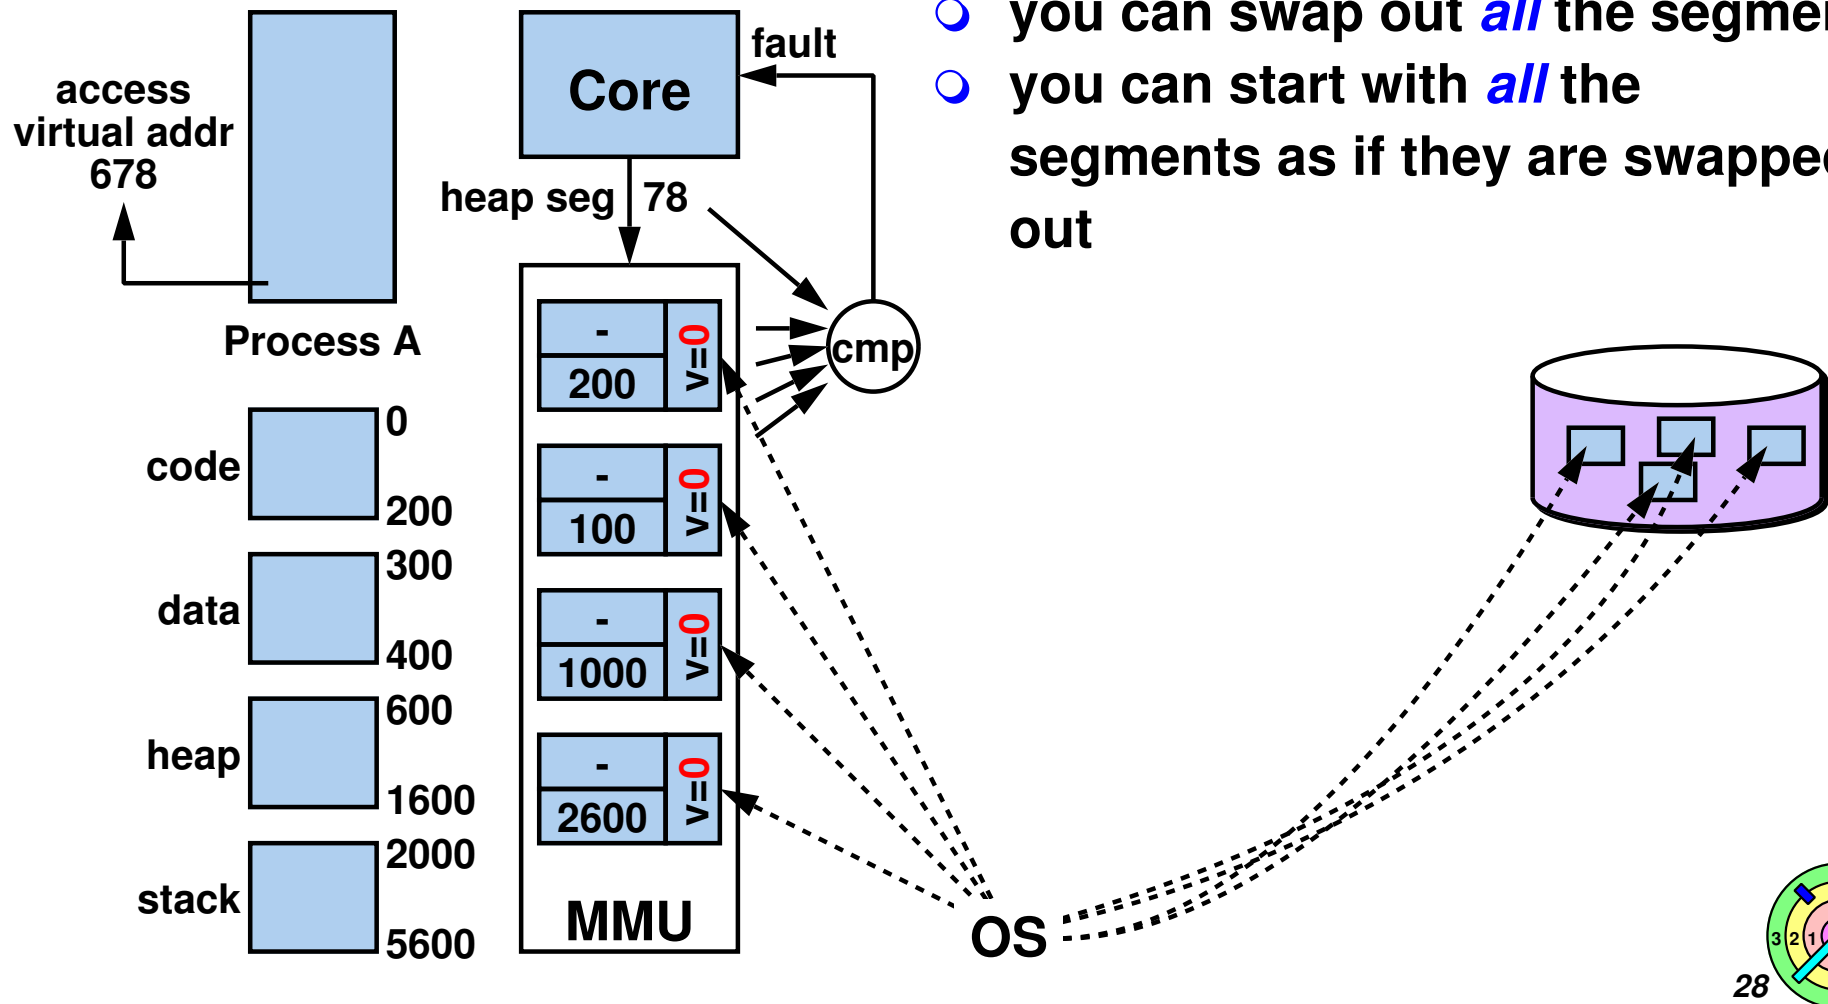
- ➡ No space for new segment, make room by swapping out a segment
- use a *validity* bit for each segment (in addition to access control bits)



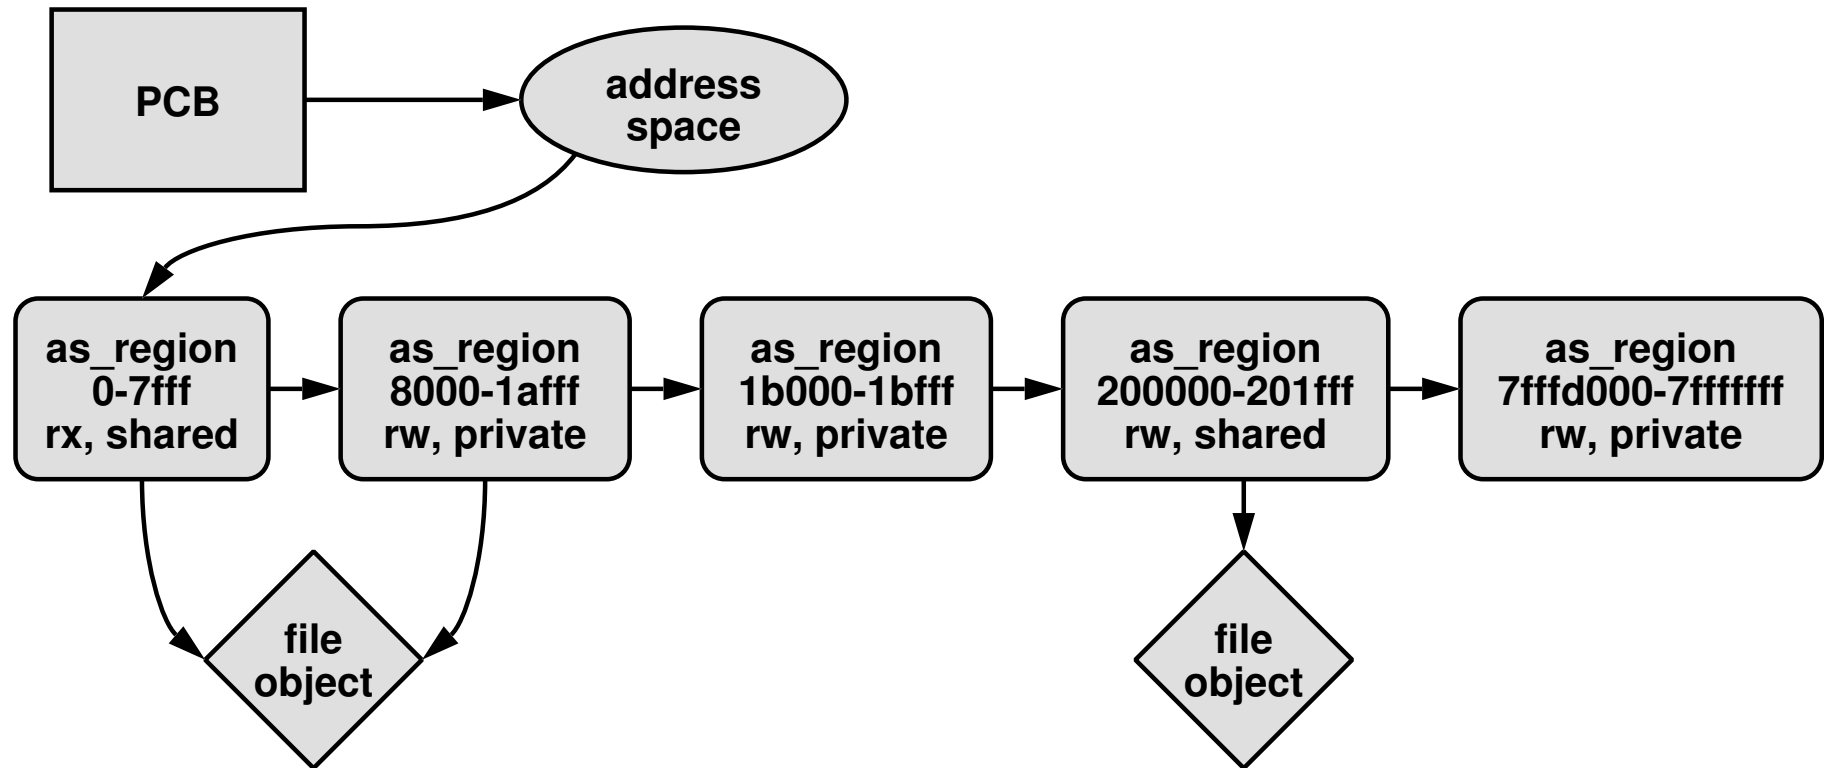
Swapping / Backing Store

- ➡ No space for new segment, make room by swapping out a segment
 — use a **validity** bit for each segment (in addition to access control bits)


- you can swap out **all** the segments
- you can start with **all** the segments as if they are swapped out



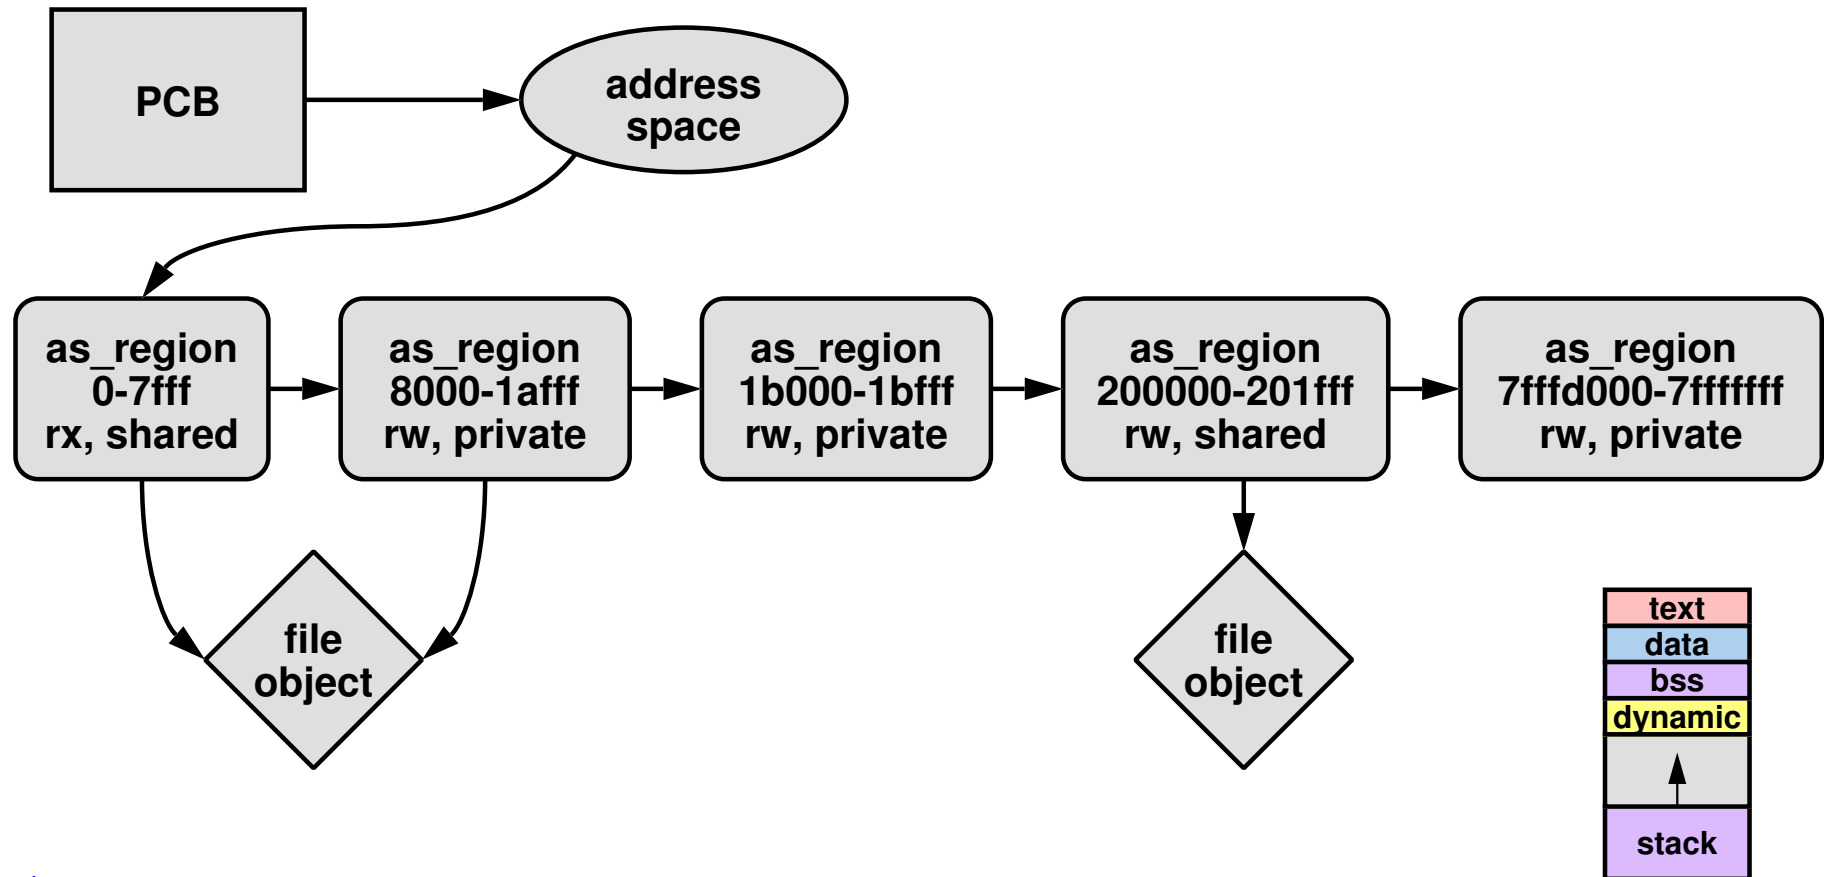
Swapping / Backing Store



➡ Remember this?




Swapping / Backing Store

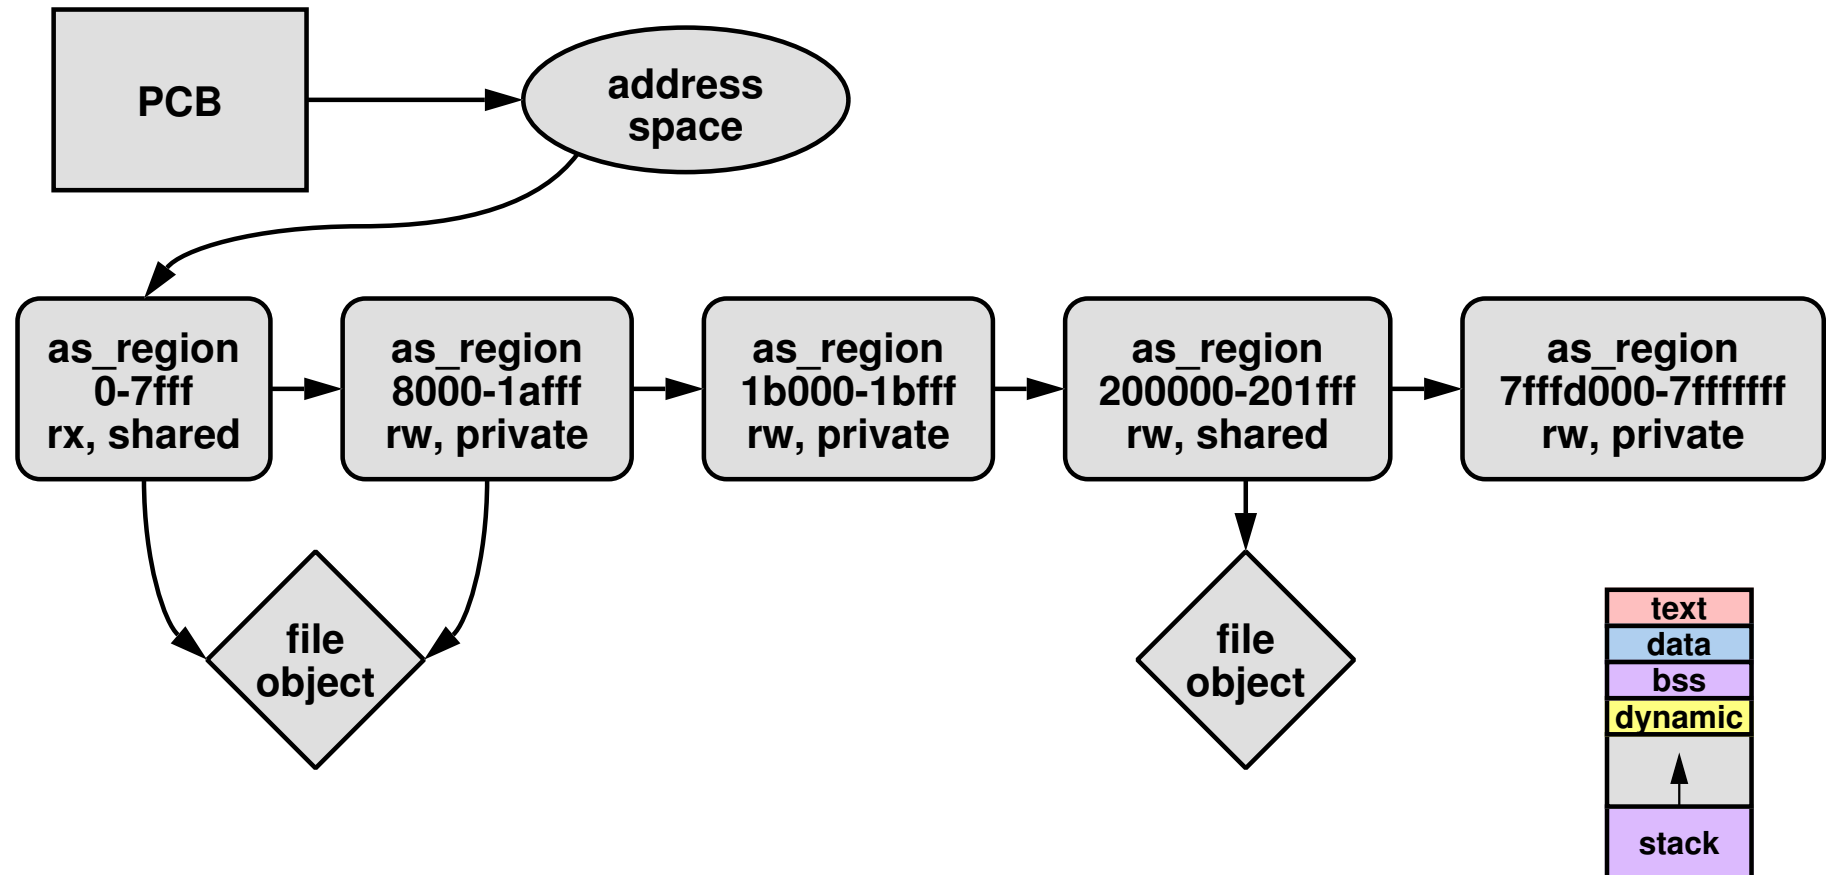


Remember this?


- this is the representation of the *address space* of a user process
 - each segment corresponds to an `as_region`



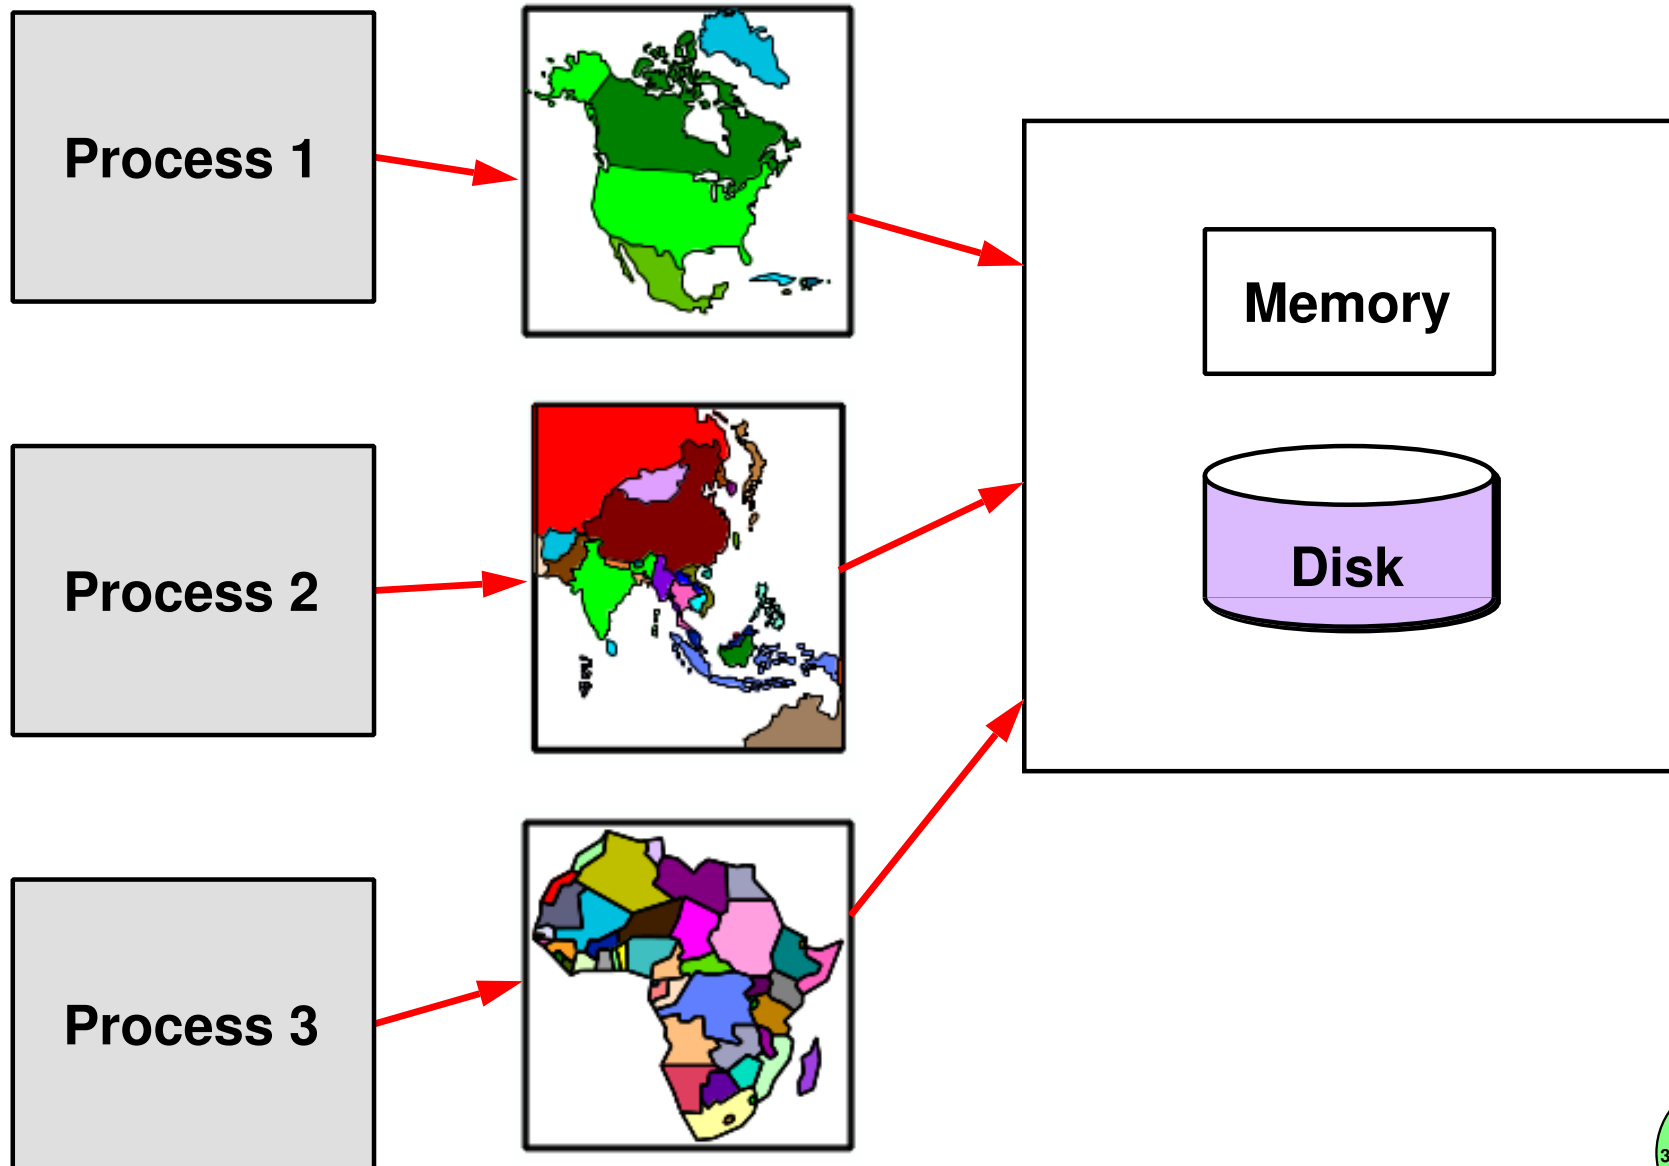
Swapping / Backing Store



- ➡ **Every** user memory segment needs a corresponding disk image, in case the user process needs to be **swapped out**
- ➡ some kernel memory can be **locked down** to prevent it from being swapped out accidentally




Virtual Memory



7.2 Hardware Support for Virtual Memory

- ➡ *Forward-Mapped Page Tables*
- ➡ Linear Page Tables
- ➡ Hashes Page Tables
- ➡ Translation Lookaside Buffers
- ➡ 64-Bit Issues
- ➡ Virtualization




Structuring Virtual Memory

➡ *Segmentation* (just discussed)

- divide the address space into variable-size segments (typically each corresponding to some logical unit of the program, such as a module or subroutine)
- external fragmentation possible
- "first-fit" is slow
- not very common these days

➡ *Paging*

- divide the address space into fixed-size pages
- internal fragmentation possible





Paging

- ➡ Map *fixed-size pages* into physical memory (*into physical pages*)
 - ▬ address space is divided into pages
 - indexed by virtual page number
 - ▬ physical memory is divided into pages (of the same size)
 - indexed by physical page number
 - ▬ need a lookup table to map *virtual page numbers* to *physical page numbers*

- ➡ Ex: 1GB of physical memory with 4KB pages
 - ▬ 2^{18} physical pages
 - ▬ an address (either physical or virtual) is *page-aligned* if its least significant 12 bits are all zero



- ➡ Many *hardware* mapping techniques
 - ▬ *MMU* and *page table* (mostly in software)
 - ▬ *translation lookaside buffers (TLB)*



Page Frames



➡ A **page frame** data structure / object is used to maintain information about physical pages and their association with important kernel data structures

- ➡ contains a **physical page number**
- ➡ there is a **one-to-one mapping** between page frames and physical pages
 - we use "page frame" and "physical page" interchangeably


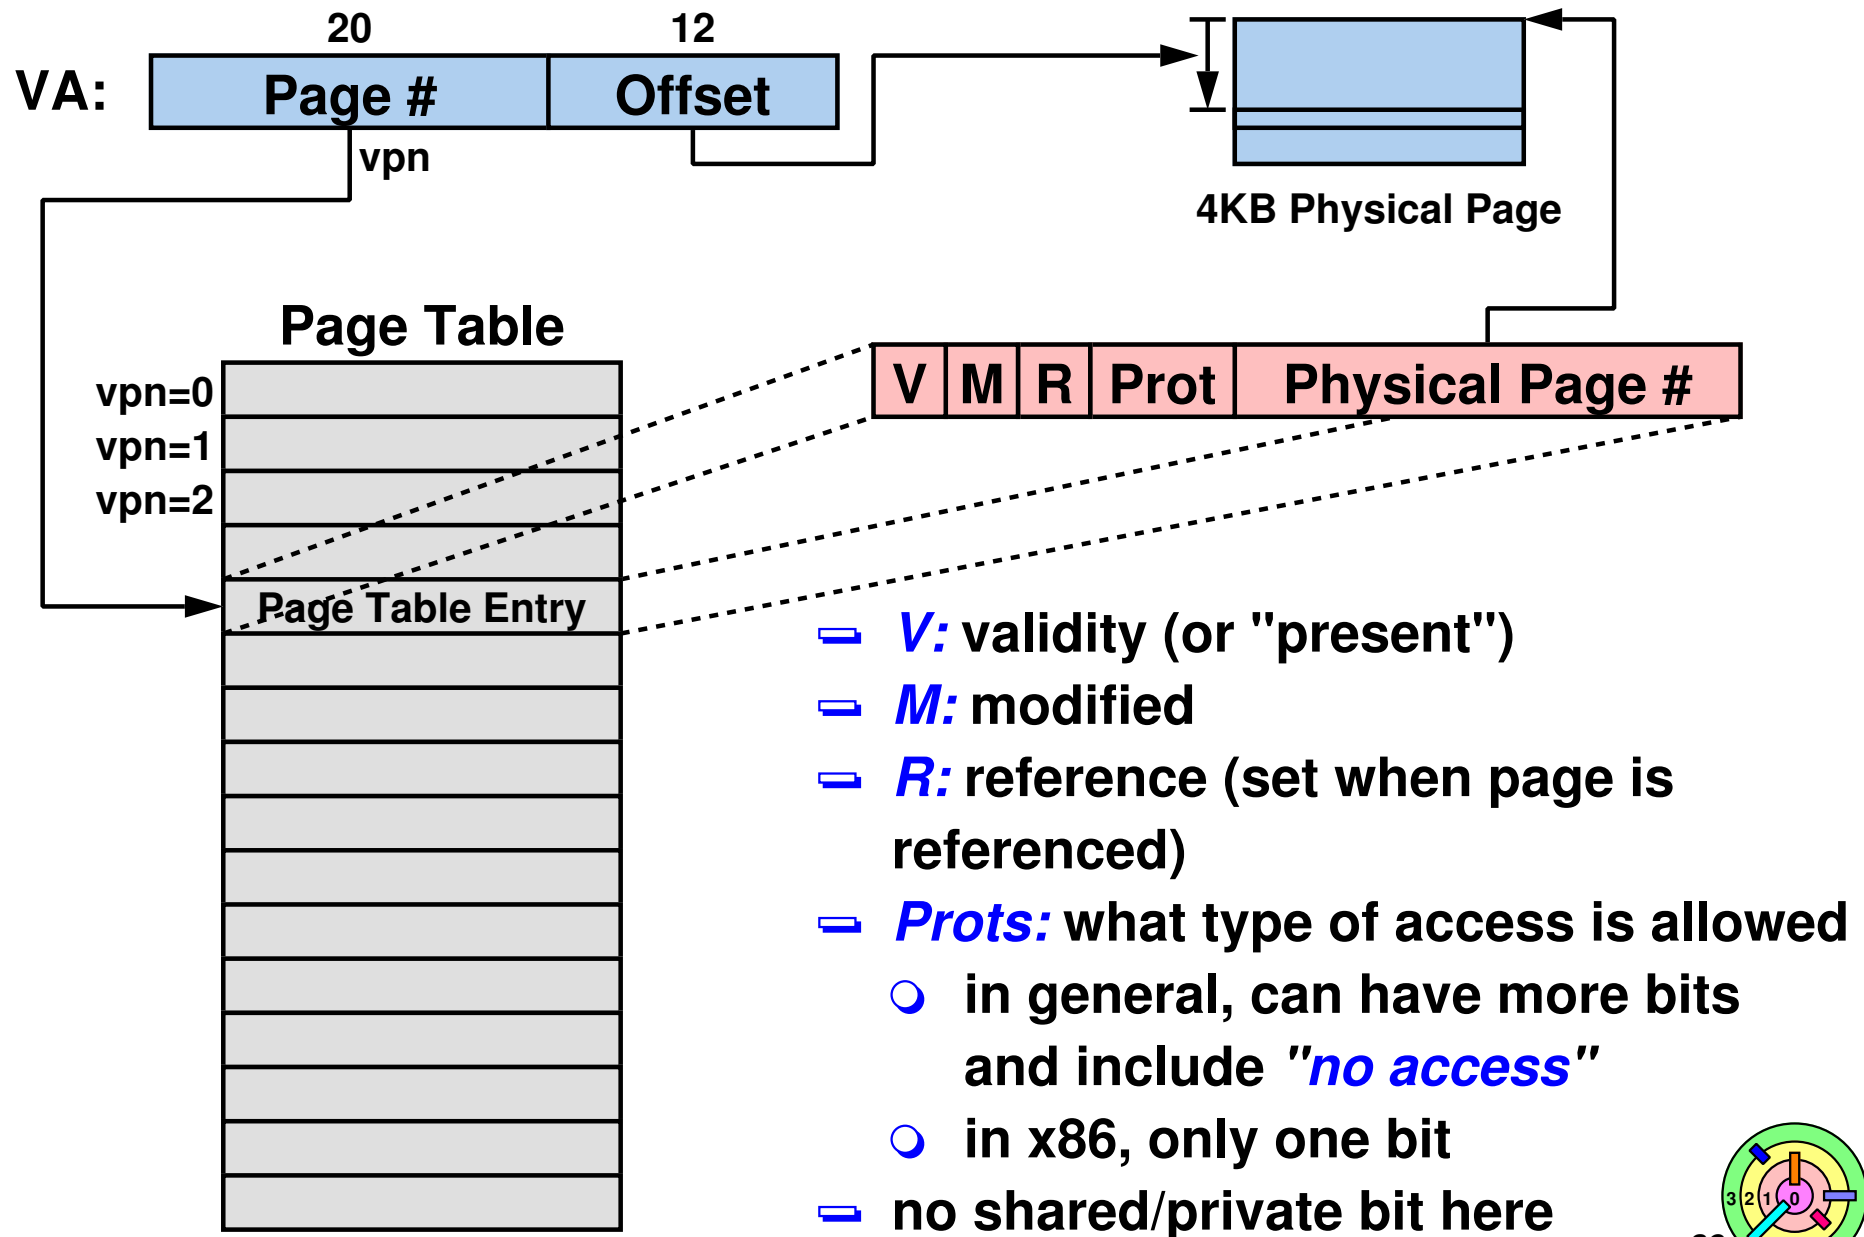


Page Frames

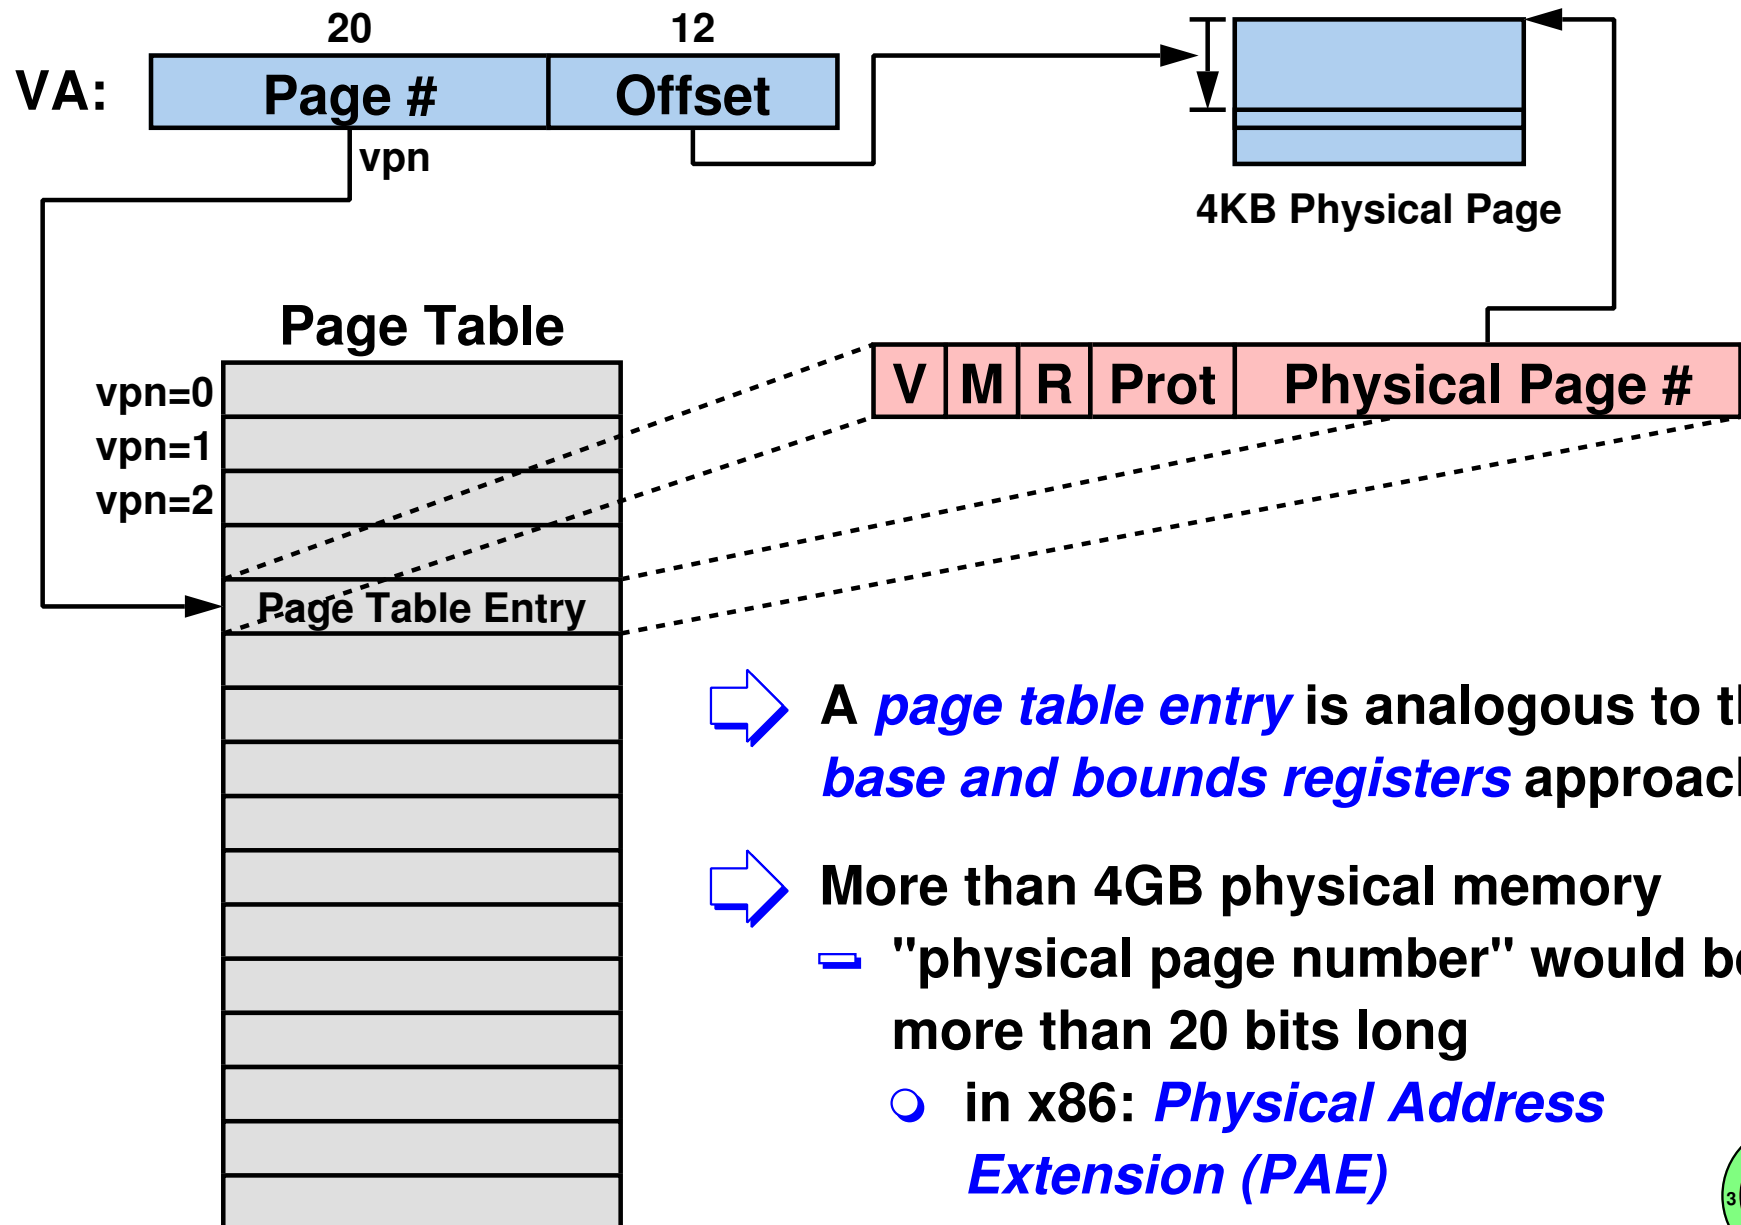
- ➡ It is important to be able to perform both *forward lookup* and *reverse lookup*
- given a virtual address of a process, find page frame
 - given a page frame, find processes and virtual addresses that uses this page frame
 - weenix page frame data structure is a bit involved
 - see kernel 3 FAQ
 - the kernel must use a *virtual address* to write into a physical page or read the content of a physical page



Basic (Two-level) Page Tables




Basic (Two-level) Page Tables


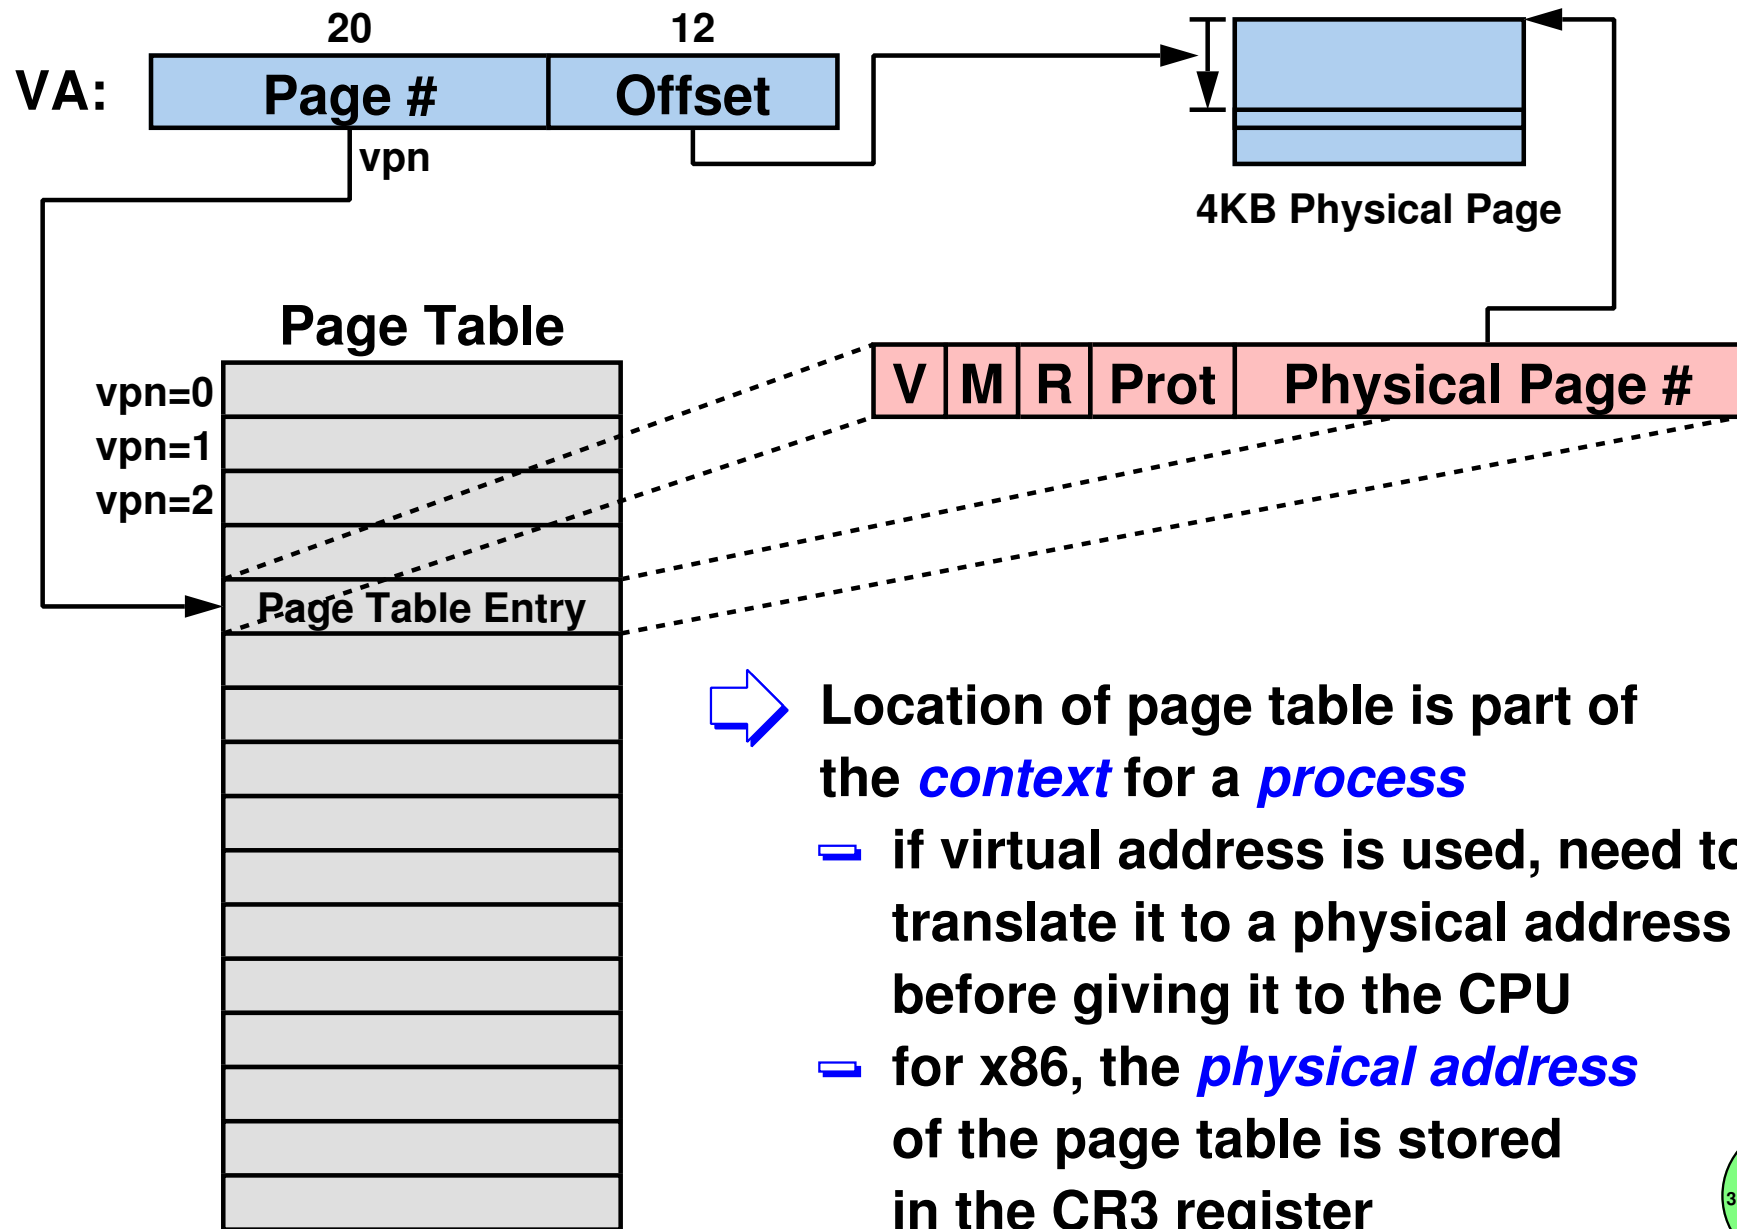


➡ A *page table entry* is analogous to the *base and bounds registers* approach

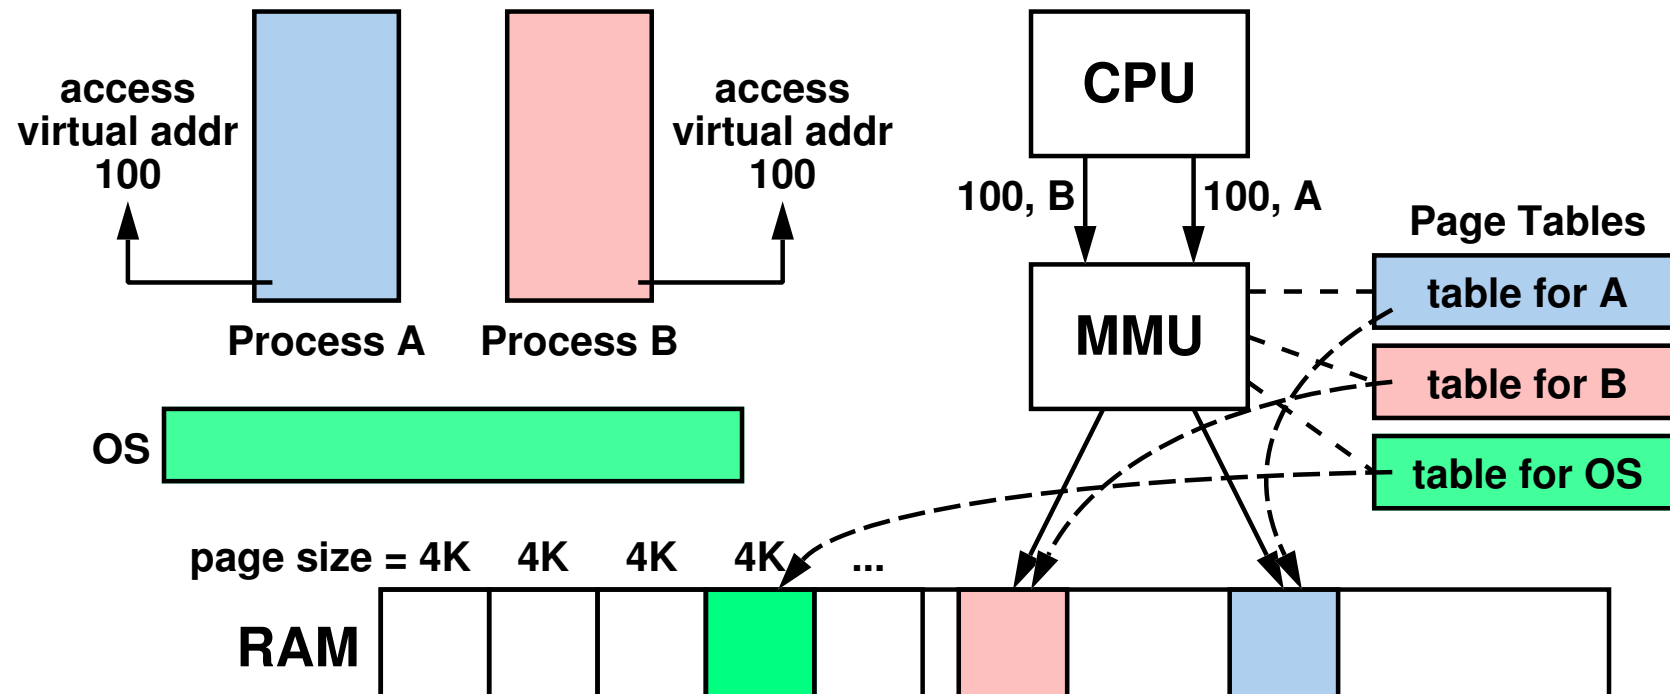
➡ More than 4GB physical memory
 = "physical page number" would be more than 20 bits long
 ○ in x86: *Physical Address Extension (PAE)*




Basic (Two-level) Page Tables



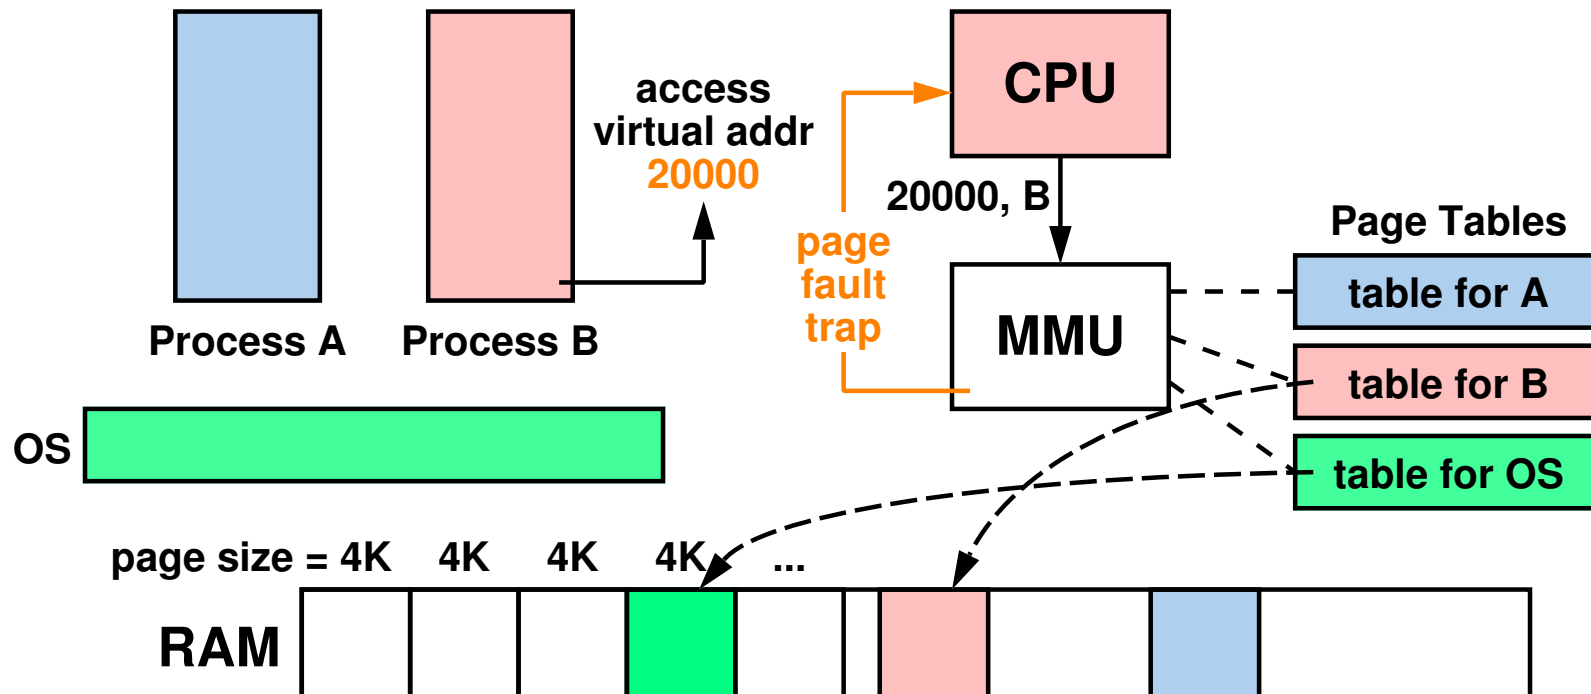
Page Table



- ➡ A page table (*usually sits in **physical memory***) is associated with each **process**
 - OS has its page table as well
- ➡ Memory Management Unit (MMU) maps virtual address to physical address
 - MMU got turned on some time during boot




Page Table

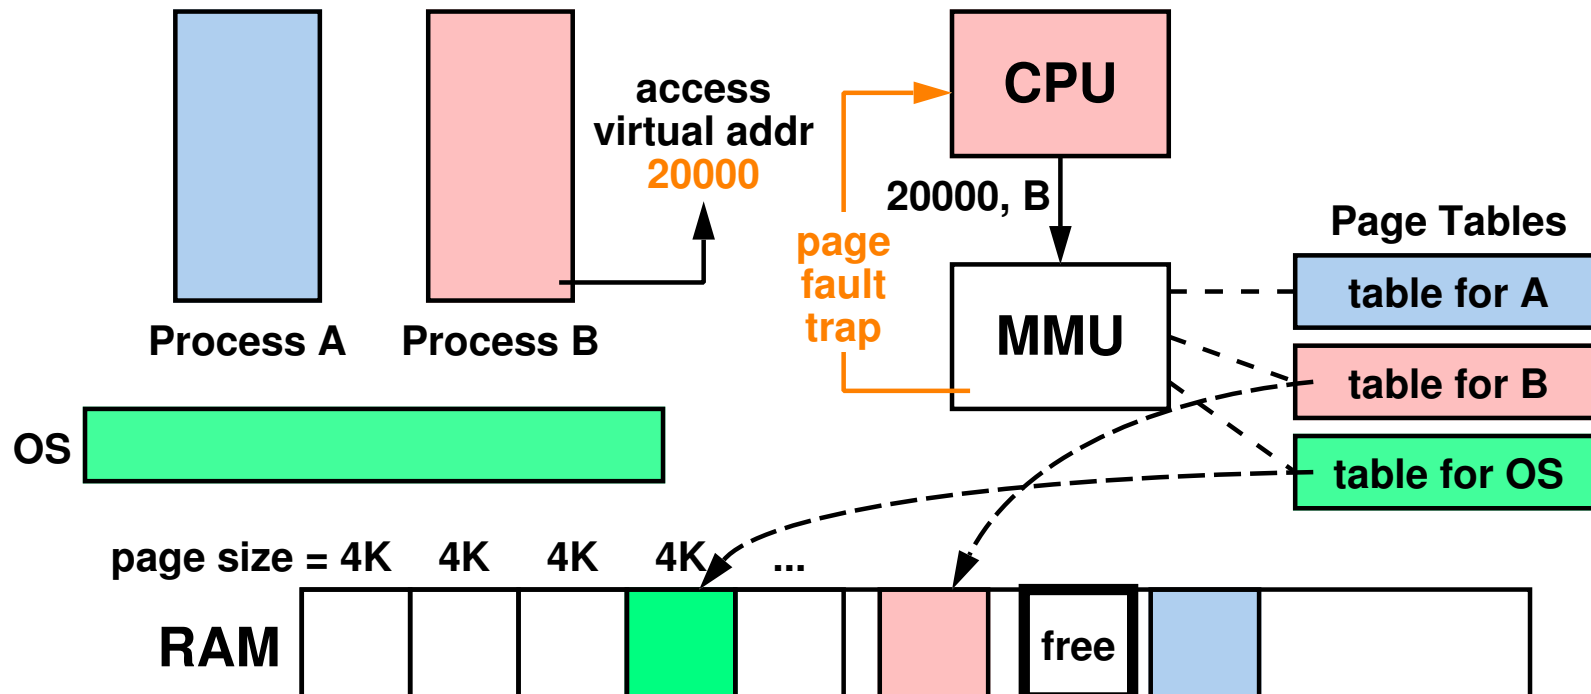


Page fault

= page table does not have the requested address




Page Table

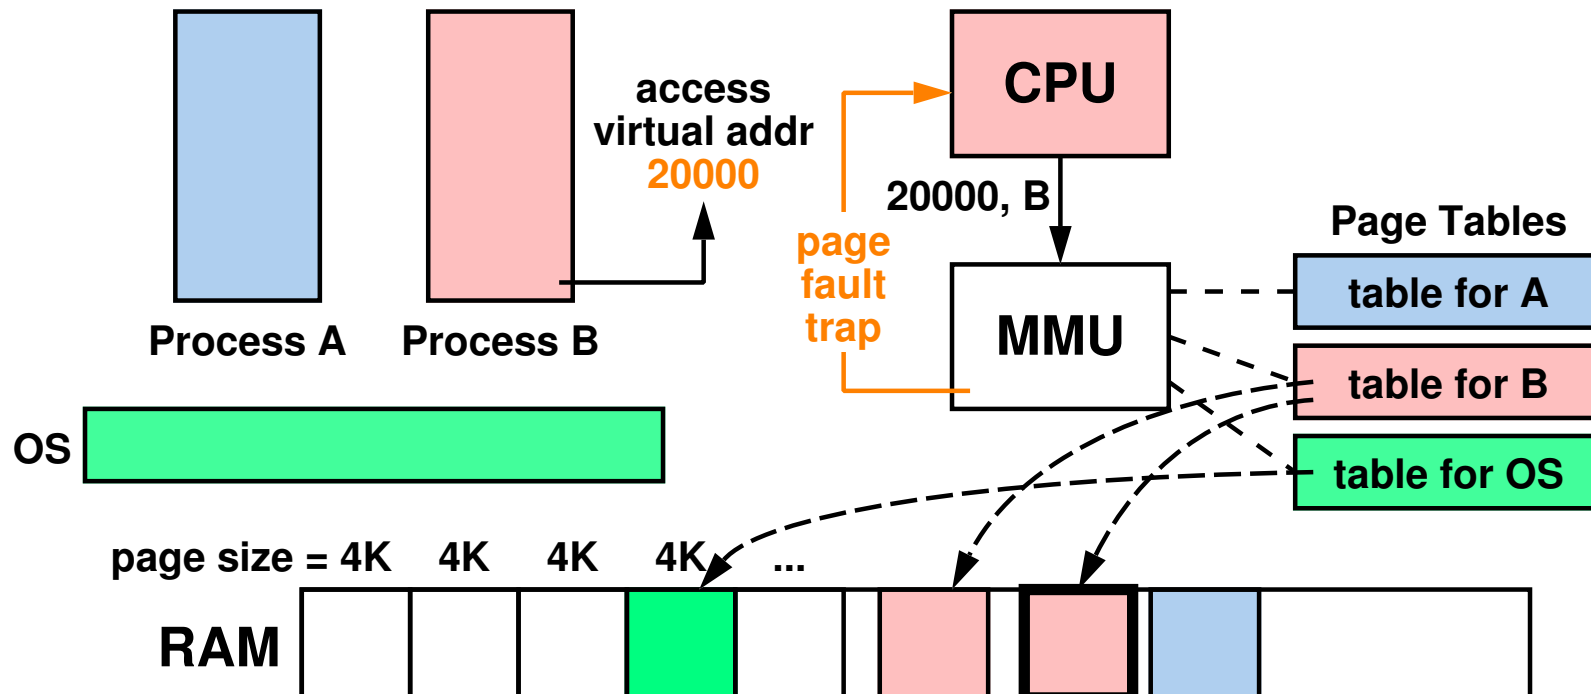


➡ Page fault

- page table does not have the requested address
- OS finds a free page frame
 - what if no free page frame is available?



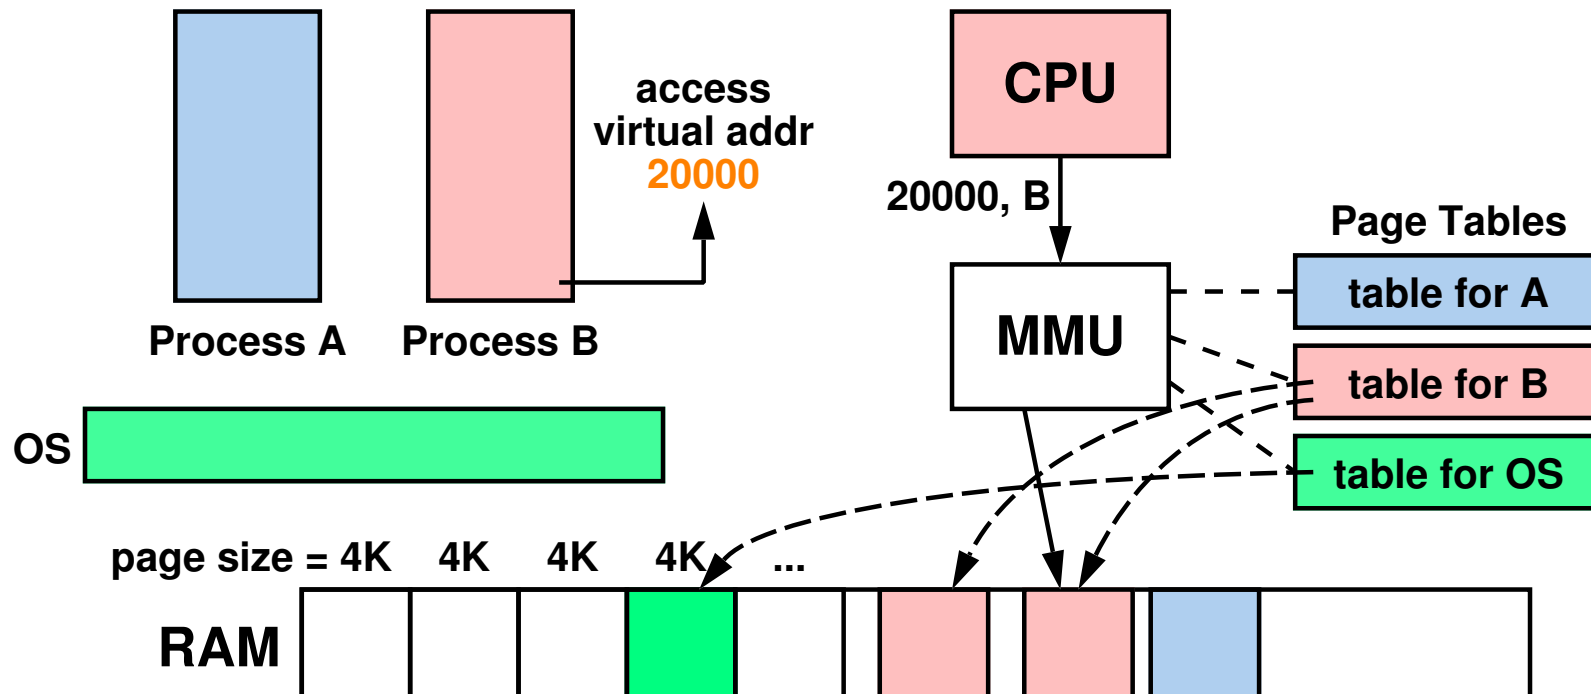
Page Table



➡ Page fault

- page table does not have the requested address
- OS finds a free page frame
 - what if no free page frame is available?
- OS loads the requested page from disk

Page Table



➡ Page fault

- page table does not have the requested address
- OS finds a free page frame
 - what if no free page frame is available?
- OS loads the requested page from disk
- OS adjusts MMU and *restarts* user memory reference

